



MS AZURE ADMINISTRATION

9963799240 / 7730997544
Ameerpet / Kondapur
Hyderabad

Contents

Introduction to Cloud Computing	8
Cloud computing	8
Types of cloud computing service models	9
Types of cloud deployment models	11
What is Microsoft Azure?	14
Overview of Microsoft Azure	15
Azure basics	15
Azure accounts versus Azure subscriptions	15
Azure Resource Manager (ARM) versus Azure Service Manager (ASM)	16
Azure global infrastructure	17
Availability Sets versus Availability Zones.....	19
Azure tools	19
Azure Portal.....	19
Azure command-line interface (Azure CLI).....	20
Azure Cloud Shell	21
Azure PowerShell.....	22
Azure SDK	23
Azure RESTful API	23
ARM templates.....	24
Azure developer tools	24
Overview of Microsoft Azure core services	25
Azure Compute services – IaaS versus PaaS	25
Azure Networking	26
Azure Storage	26
Data and analytics services	27
Backup services and disaster recovery	28
Administrative roles and role-based access control	28
Further reading.....	29
Implementing and Managing Azure Virtual Machines	30
The principles of Azure VMs	30
Planning and deploying Azure VMs	31

Identifying the workloads	31
Choosing the appropriate Azure VM sizing	31
Azure VM storage options	33
Managed disks versus unmanaged disks.....	33
Azure reserved VM instances (RIs) versus pay-as-you-go instances.....	34
Deploying an Azure VM	34
Creating Azure VMs.....	34
Connecting to Azure VMs.....	51
Connecting to a Windows Azure VM via Remote Desktop Protocol (RDP)	52
Connecting to a Linux Azure VM via Secure Shell (SSH)	52
Configuring Azure VMs in security.....	53
Restricting access to Azure VMs from the internet using NSG	53
Managing Azure VMs with VM Agent and VM extensions.....	55
Configuring the availability and scalability of Azure VMs	57
Scaling Azure VMs.....	57
Managing Azure VM's availability	67
Implementing and Managing Containers in Azure.....	73
The principle of containers and microservices	73
Containers versus container clusters.....	75
Docker basics	76
Container registry	77
Dockerizing your web application in Azure	78
Preparation work	78
Implementing Azure Container Registry	79
Pushing your Docker image in ACR.....	81
Deploying your Dockerized application with CI/CD capabilities	84
Clustering solutions with Azure ACS in Azure.....	84
An overview of container cluster solutions working with ACSs	85
Implementing three types of orchestrators of Azure ACS in Azure	89
Docker Swarm	89
Kubernetes	93
Implementing and managing a Kubernetes cluster with AKS	99

Implementing and Managing Azure Virtual Networks	105
Planning and designing Azure virtual networks	105
Analyze network requirements	105
Determine the type of connectivity	107
Determine the address space	107
Assigning static, public, and private IP addresses	108
Implementing Azure virtual networks	110
Creating an Azure virtual network	110
Updating the Azure virtual network	115
Delete Azure virtual network	117
Managing Azure virtual networking	118
Routing network traffic	118
Filtering the network traffic	120
Distributing network traffic	121
Design and implement cross-premise and multisite connectivity	129
Point-to-site virtual private network (VPN) over IKEv2 or SSTP	130
Site-to-site and multisite virtual private network (IPsec/IKE VPN tunnel)	132
VNet-to-VNet virtual private network (IPsec/IKE VPN tunnel)	133
Virtual network peering (VNet Peering)	134
ExpressRoute	140
Configuring Hybrid Connections for App Service	143
Configuring multi-region applications with Azure Traffic Manager.....	145
Creating a Traffic Manager profile	146
Adding endpoints to the Traffic Manager profile	148
Managing Traffic Manager profiles	150
Integrating Azure services with an Azure virtual network.....	151
Restrict network access to PaaS resources using a service endpoint	153
Integrating a web app in App Service with an Azure virtual network	153
Configuring accelerated networking to improve your networking performance	155
Managing Azure Identities	156
Implementing and managing Azure Active Directory (Azure AD)	157
Managing identities via Azure Active Directory admin center	157

Creating an Azure Active Directory via the Azure Portal	159
Creating and managing Azure AD users	162
Creating Azure AD groups and managing user groups	164
Enabling Multi-Factor Authentication for users	165
Using bulk update for custom user profile properties.....	167
Managing devices	168
Add a custom domain	169
Conditional access	171
Configuring self-service password reset.....	172
Configuring privileged identity management.....	174
Configuring Azure AD identity management.....	174
Leveraging Microsoft Graph other than Azure AD Graph API.....	175
Integrating applications with Azure AD	175
Creating an Azure AD B2C directory	176
Managing Azure AD B2C directory	178
Implementing Business to Business (B2B) collaboration	180
Integrating applications with Azure AD	181
Implementing federation and social identity provider authentication.....	183
Configuring SAML-based SSO for an application with Azure AD	186
Managing hybrid identities	186
Configuring Azure AD Connect and synchronization services	186
Managing domains with Azure AD domain services	188
Implementing SSO in hybrid scenarios	190
Monitoring on-premises identity infrastructure and synchronization services.....	190
Planning and Implementing Azure Storage, Backup, and Recovery Services	191
Implementing and managing Azure Storage.....	192
An overview of Azure Storage services	192
Implementing Azure Storage services	192
Managing Azure Storage services	214
Implementing hybrid storage solutions	222
Moving data to and from Azure Storage.....	222
Implementing data storage services	225

SQL Database	226
Azure Database for MySQL	226
Azure Database for PostgreSQL	227
Database-managed instances	227
Azure SQL Data Warehouse	228
Cosmos DB	228
Configuring Content Delivery Network	229
Implementing a business continuity and disaster recovery (BCDR) strategy in Azure	231
Planning a BCDR strategy	232
BCDR in Azure	232
Implementing Azure Backup	233
Planning and implementing Azure Site Recovery	243
Managing Azure Operations and Automation	247
Implementing Azure Automation.....	247
An overview of runbooks	248
Creating an Automation account	248
Creating or importing PowerShell runbooks	251
Managing PowerShell runbooks	257
An overview of Desired State Configuration	261
Implementing PowerShell Desired State Configurations	262
Managing PowerShell Desired State Configurations	263
Other excellent configuration management tools	264
Implementing Azure Automation-based cloud management	264
Integrating Azure Automation with Web Apps	265
Comparing Azure Automation with Azure Functions	265
Integrating with Event Grid	266
Integrating with Logic Apps	268
Runbook gallery	269
Implementing monitoring solutions in Azure	270
Core monitoring	271
Deep application monitoring	275
Deep infrastructure monitoring	276

Shared capabilities	277
Implementing Azure VMs monitoring solutions	278
Configuring ARM VM monitoring	278
Configuring alerts	279
Configuring diagnostic and monitoring storage location.....	280
Implementing Log Analytics (OMS) solutions	282
Creating an OMS workspace.....	282
Collecting and searching across data sources from multiple systems	284
Transforming Azure activity data and managed resource data	285
Building custom visualizations with view designer	287
Sending data to Log Analytics with the HTTP Data Collector API	288
IT Service Management Connector.....	288



Introduction to Cloud Computing

Information technologies have evolved significantly over the last decade. For those of us in the professions of creating, building, and developing within these changes, cloud computing is one of the fastest-growing areas in our world. **Cloud** is the keyword of our new age, and it will be a fundamental part of everything in our future.

As an IT professional, I believe that being successful in IT is really about being a lifelong student, which means that we must constantly learn new skills and different platforms. Thus, this exam guide is not only a tool to help IT professionals get certified, but it will also help them to get hands-on experience through practice labs. I would love to share my experience on how to build our knowledge of Microsoft Azure from the conceptual level to the operational level.

Cloud computing



Cloud computing has been a star since it was born; it appears with big data, **Internet of Things (IoT)**, and **Artificial Intelligence (AI)** in our conversations.

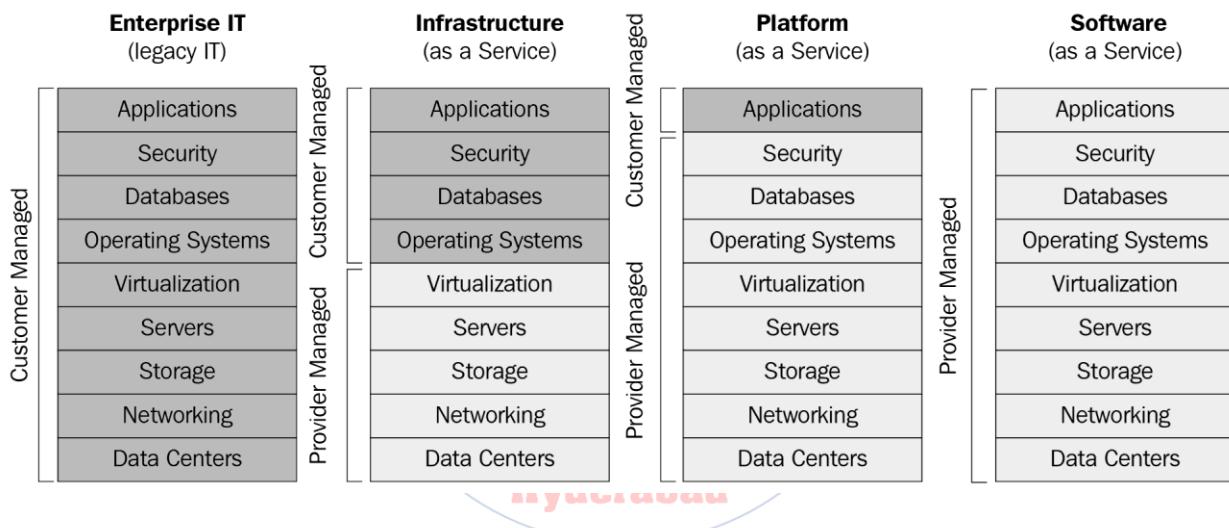
Ameerpet / Kondapur

Modern cloud computing services providers such as AWS, Microsoft Azure, and Google Cloud Platform, generally consist of four basic components: **compute, network, database, and storage**. These providers provide computer services such as virtual machine, storage services to store objects/files on the cloud, provide the different cloud-based databases to store data, and network service to deploy virtual networks. Nowadays, the popular cloud providers seem more ambitious than we thought they were. They don't limit themselves to act as infrastructure that supports application deployments but they also provide the management services to support DevOps, monitoring, logging, and alarming, backup and restore in time on the cloud and integration tools to build CI/CD pipelines. More and more, it provides advanced services support such as the **ETL (extract, transform and load)** processing and data analytics, **Machine Learning (ML)**, AI, and also IoT services to communicate

with IoT devices. Theoretically, cloud computing technology can help us do everything we want in the cloud.

Types of cloud computing service models

Predominantly, cloud computing is built into three types of cloud computing service models: **Infrastructure as a Service (IaaS)**, **Platform as a Service (PaaS)**, and **Software as a Service (SaaS)**. Each service offering provides a different level of virtualization and management responsibilities. The following is a screenshot showing which responsibilities each service model provides for:



Different cloud computing service models

IaaS – Infrastructure as a Service

As shown in the preceding screenshot, the IaaS service model provided is generally the capability of infrastructure level to the users; the cloud providers managed the hardware and infrastructure such as virtual servers, storage, networks, connectivity, operating systems and other fundamental computing resources. Based on the IaaS offering, users should manage them with administration works such as installing the patches, updates, and configurations. One of the best examples of this model is the virtual machine in the cloud.

PaaS – Platform as a Service

The PaaS services model provides the capability that comprises deployed and configured IT resources in a ready state based on the IaaS model. Users don't need to care about the infrastructure level and even the administration work they face when they're in the IaaS model. It directly provides the environment with the specified runtime. What users need to do is focus their work on the application level.

SaaS – Software as a Service

Compared to PaaS, SaaS has a more advanced virtualization level, which is widely accessed over the internet and directly used by users. The most common example is Google's **G Suite** and Microsoft's Visual Studio Team Services (which is also known as **Visual Studio Online** or **VSTS**).

X as a Service in the cloud

Based on these models, there are also some extension concepts which are known as **X as a service** such as:

- **Database as a Service (DBaaS):** A managed database service in the cloud that aims to offer the database layer to the applications; the cloud provider manages the complex database environments.
- **Container as a Service (CaaS):** A managed service model that provides the container-based virtualization technology to let users manage and deploy containers; and applications as well as container clusters in the cloud.
- **Messaging as a Service (MaaS):** A messaging service in the cloud that allows sending and receiving messages through a queue. Originally implemented for the purpose of resolving queue-based load-leveling problems for a service whose peaks in demand make services or applications in the cloud overload and therefore unable to respond to requests in a timely manner. The queue acts as a buffer, storing the message until it's retrieved by the service. Applications or services in the cloud retrieve the messages from the queue and process them.
- **Logic as a Service (LaaS):** Also known as serverless. It gives little control over the infrastructure, the related infrastructure is managed by the cloud providers, and users can focus themselves on coding and configuring settings. Great examples include Azure Functions and Azure Logic Apps.
- **Identity as a service (IDaaS):** Supplies cloud-based authentication or identity management to enterprises or organizations. The goal is to ensure if a user has access to

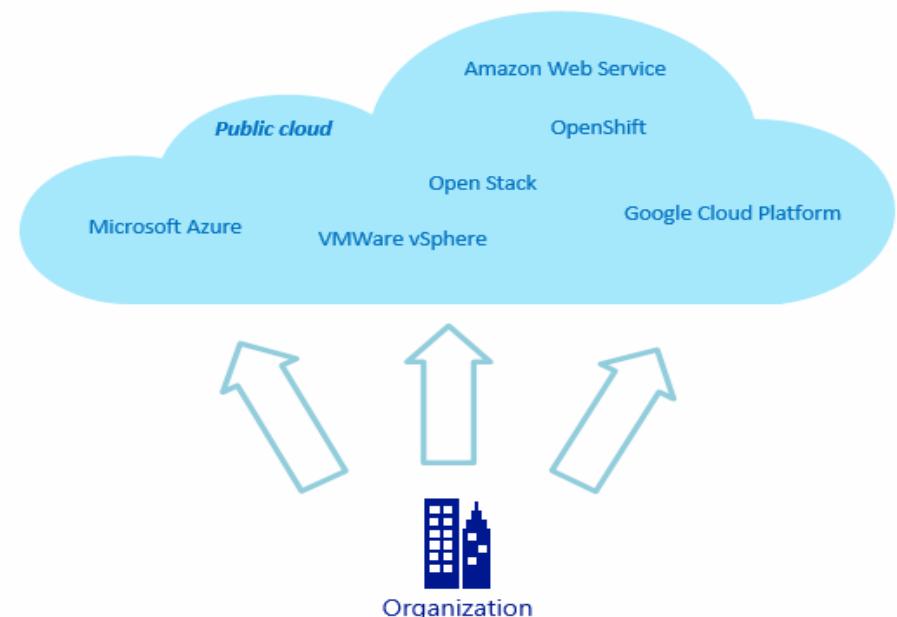
cloud applications or services and which type of access they could have to cloud applications or services.

There are some other service models such as **Disaster Recovery as a Service (DRaaS)**, **Data as a Service (DaaS)**, **Big Data as a Service (BDaaS)**, **Log as a Service (LaaS)**, and more than the mentioned models here. We believe, as cloud computing is one of the fastest-growing technologies, more and more services will appear and may serve together for future cloud computing platforms.

Types of cloud deployment models

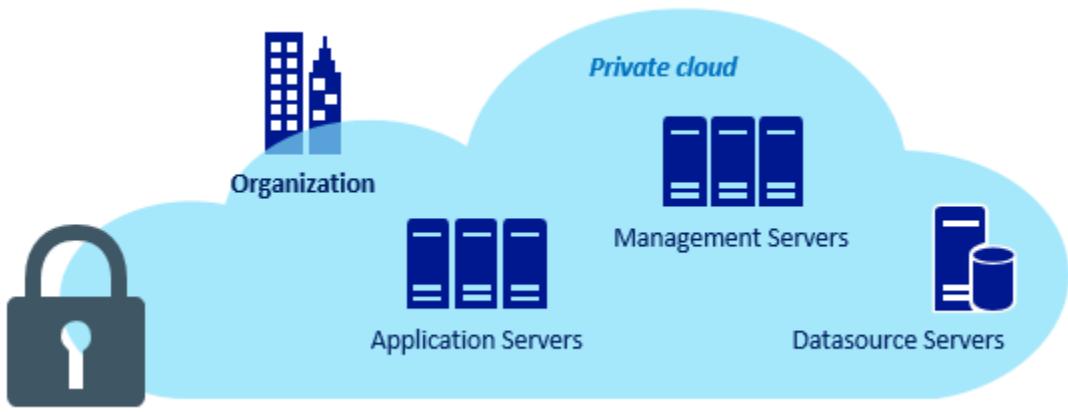
Typically, within cloud computing, it is possible to build our model as one of the four types of cloud deployment model: public cloud, private cloud, hybrid cloud, and community cloud. Each of them is defined for different levels of management, such as where the IT resource is located and security reasons.

The **public cloud** is a publicly accessible cloud environment provided by a cloud provider such as Microsoft Azure, **Amazon Web Service (AWS)**, **OpenShift**, **Google Cloud Platform**, **Microsoft Azure**, **Open Stack**, **VMWare vSphere**, **Cloud**, **Platform (GCP)**. These platforms manage the IT resources in their data center and are responsible for the security of these IT resources. With the help of the Internet, users can access cloud services on the **public cloud** from anywhere around the world. The following is an image showing how the **public cloud** works:



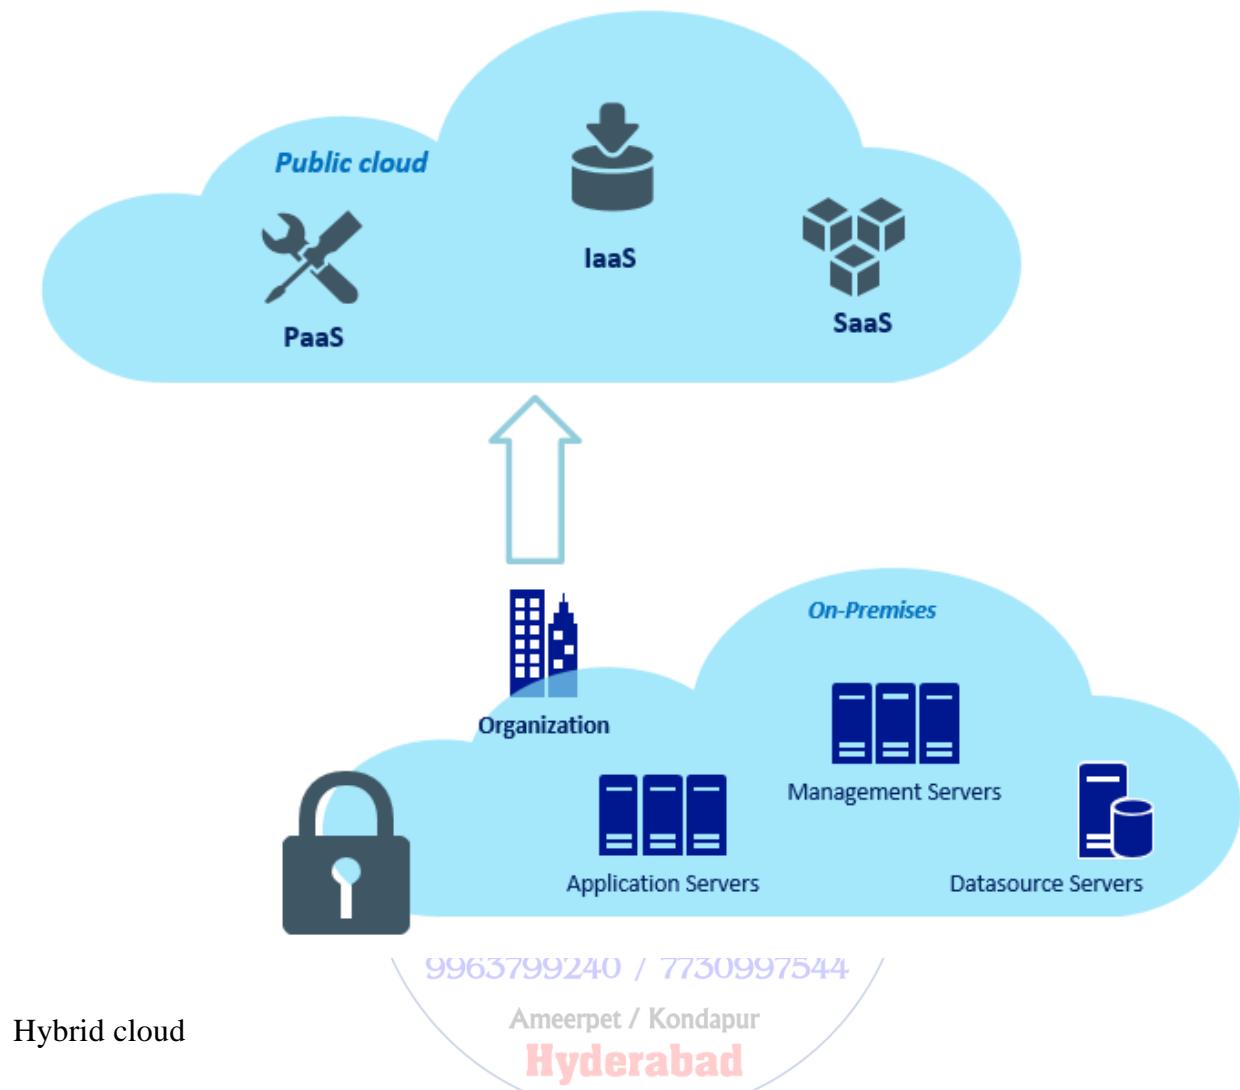
Public cloud

As compared to the **public cloud**, a **private cloud** is more critical in security and can only be accessed from the internal network where the infrastructure is hosted. IT resources in private clouds are managed by companies' or **organizations'** own data center. Users can access it only when they're on the internal network. The following is an image showing how a **private cloud** works:

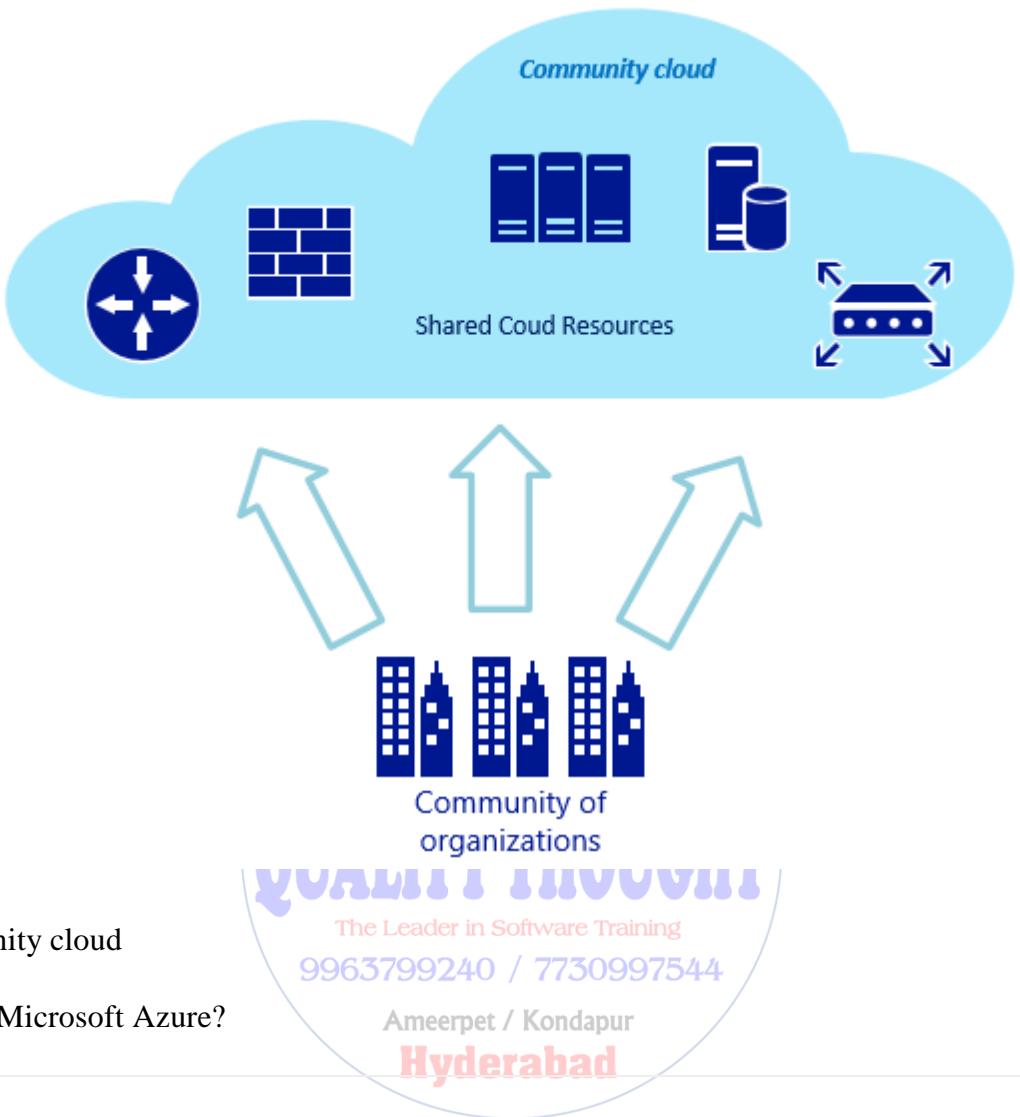


Private cloud

Based on the **public cloud** and **private cloud**, a **hybrid cloud** is intended to combine both of them in the same scenario. Generally, a network connection (dedicated or private) should be established between the **private cloud** and the **public cloud**. Hence, it is important to define which IT resources are on-premises and in the cloud and how do they work together. Be careful, as a hybrid cloud is intended to be for short-term configurations. If we are in a transition stage, a hybrid cloud is the most common cloud deployment model. The following is an image showing how a hybrid cloud works:



A **community cloud** is similar to a **public cloud**, but it is intended to be limited to a specific member of the community (or authorized access by the community); the IT resources can be located in the community's data center and shared by several members. The general scenario is several applications consume the common IT resources, but in this context, the **community cloud** usually ensures a dedicated connection between each consumer and IT resource, as illustrated in the following diagram:



Microsoft Azure was announced in October 2008 and then released on February 1, 2010 as Windows Azure. By March 25, 2014, it was renamed Microsoft Azure. The Microsoft Azure public cloud platform offers IaaS, PaaS, and SaaS services to enable businesses worldwide to create, deploy, and operate cloud-based applications and infrastructure services.

One of the reasons why Microsoft Azure is popular and fast-developing in the current market is because it is easy to work along with other Microsoft solutions such as Microsoft System Center, and can be leveraged together to extend an organization's current data center into a hybrid cloud that expands capacity and provides capabilities beyond what could be delivered solely from an on-premises standpoint.

Overview of Microsoft Azure

In this chapter, we'll introduce the core concepts of Microsoft Azure, including the different ways to access Microsoft Azure; introduce different Azure tools; indicate how to install and configure them; and provide an overview of Azure core services.

The following are the topics that we will cover in this chapter:

- Azure account and Azure subscription
- The classic model and Azure Resource Manager deployment model of Microsoft Azure
- Azure regions, global data centers, Availability Set, and Availability Zone
- Introduction to different tools to access Microsoft Azure
- An overview of Azure Compute, networking, storage, data and analytics, and backup and disaster recovery services
- Administrative roles and **role-based access control (RBAC)**

Azure basics



Microsoft Azure is a cloud platform launched by Microsoft that helps individuals and organizations provision, deploy, and operate cloud-based services and IT assessment

Azure accounts versus Azure subscriptions

When you're starting to use Azure, you'll create an Azure account. Microsoft lets users start Azure with a free Azure account. You can use the address given here to open your first Azure account: <https://azure.microsoft.com/en-us/free/>.

You can find details about all the service limits when you're using the Azure free account here at <https://azure.microsoft.com/en-us/free/free-account-faq/>.

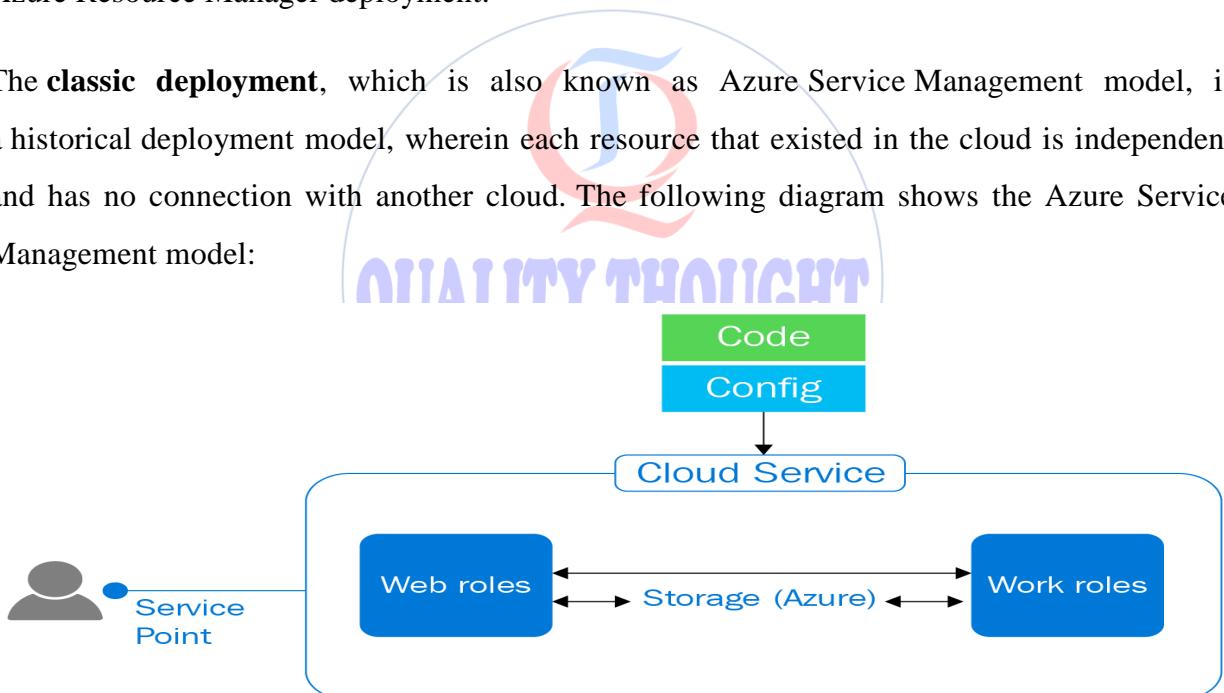
An Azure account contains one or more subscriptions. A subscription contains the details of services and billing use within an account. You can check the subscription in your account at <https://account.azure.com/subscriptions/>.

In Azure, the maximum number of services and resources is applied on per-subscription and per-region levels.

Azure Resource Manager (ARM) versus Azure Service Manager (ASM)

Microsoft Azure also has two different deployment models: Azure classic deployment and Azure Resource Manager deployment.

The **classic deployment**, which is also known as Azure Service Management model, is a historical deployment model, wherein each resource that existed in the cloud is independent and has no connection with another cloud. The following diagram shows the Azure Service Management model:



Azure classic deployment model

The IT resources for hosting virtual machines are provided by a **Cloud Service**, which is required in this model so that it acts as a container to host these VMs. There is also a **network interface card (NIC)** and an IP address, which are allocated by Azure linked with this VM.

Another deployment model is the **Azure Resource Manager** deployment model, ARM, which is different from classic deployment model. It lets you deploy, manage, and monitor all of the IT resources in the cloud such as virtual machines, storage accounts, virtual networks, or a database with a logical group, which is known as a **resource group**. The advantage of resource groups is organizing the resources in a logical way, and all the resources in the same resource group share the same life cycle, which means you can deploy, update, or delete them in a **one-click** way with only one single operation. Another great advantage is that resources deployed with ARM model can also be provisioned by a JSON-based template, which defines the dependencies between the deployed resources and the connection with the different resource groups. ARM deployment mode was inspired by **Infrastructure as Code (IaC)** which will be explained in the coming section. While interacting with the ARM, you can use command-line tools such as Azure PowerShell or Azure CLI, and you can also use ARM RESTful APIs.

Azure Resource Manager is not only a new deployment technique but also provides a consistent management layer for the different tasks you perform through different Azure tools, which we'll discuss in the upcoming section of this chapter.

When we start to deploy each a new resource in Azure, it will be necessary to specify the deployment model if this resource exists in two models. Take note of the fact that the Resource Manager deployment model and classic deployment model are not completely compatible with each other. ARM model is strongly recommended by Microsoft while deploying new IT solutions.

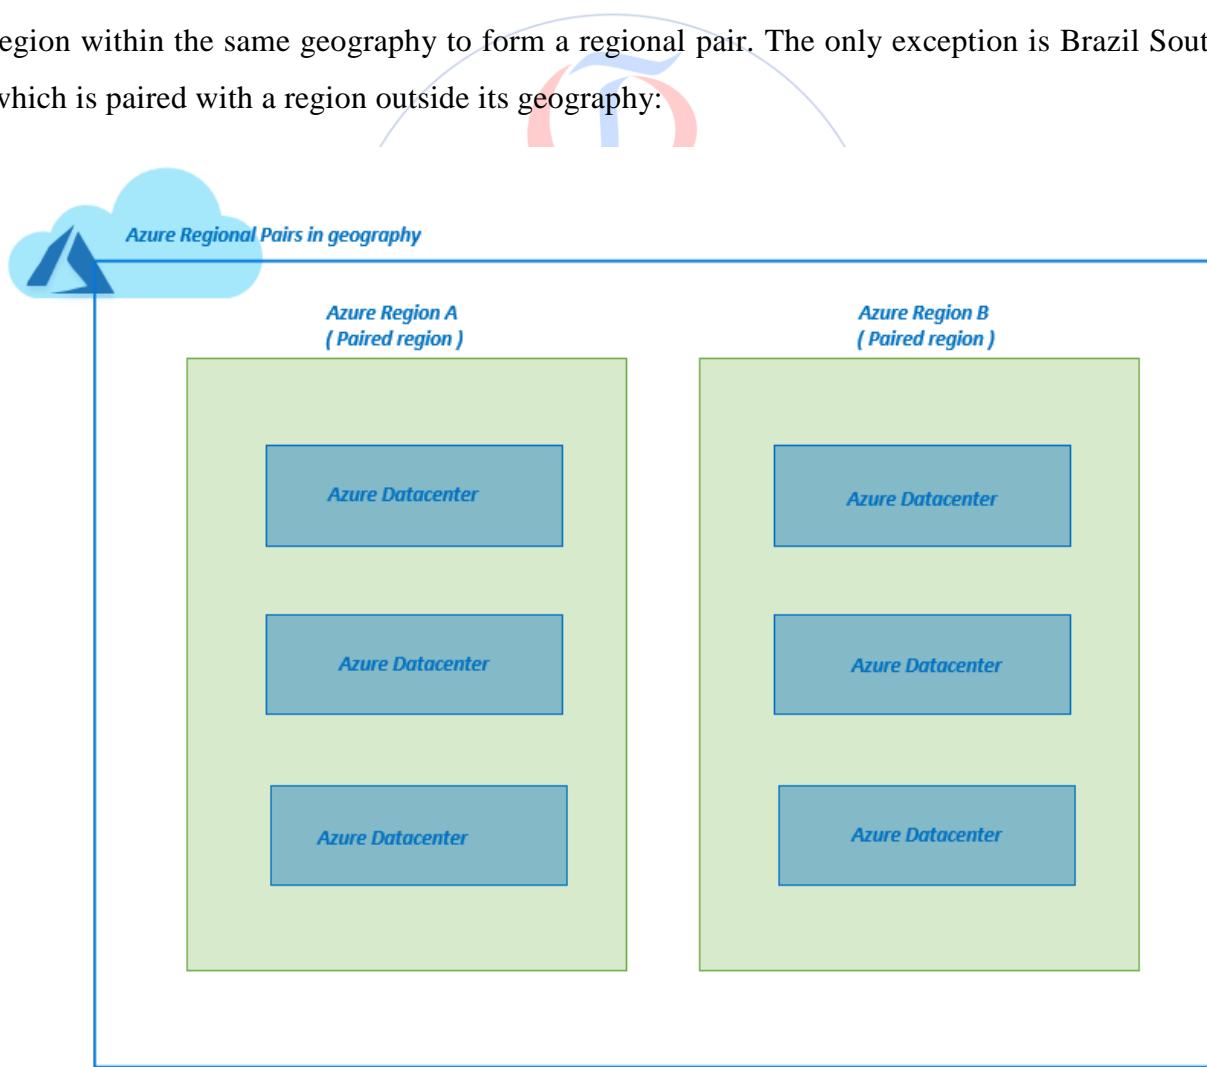
Azure global infrastructure

Azure services are hosted in physical Microsoft-managed data centers throughout the world. As of the time writing this , Azure was generally available in 54 regions and in over 140 countries around the world. To view the latest available regions of Azure global infrastructures, please check the following page at <https://azure.microsoft.com/en-us/global-infrastructure/regions/>.

Azure operates in multiple geographies around the world. An Azure geography is a defined area of the world that typically contains two or more regions and preserves data residency and compliance boundaries.

Whenever you create a new Azure resource, you must select an Azure region to determine the data center where the service will run. In Azure, a region is a set of data centers. Microsoft deployed their data centers within a latency-defined perimeter and then connected through a dedicated regional low-latency network.

One of the greatest difference between Azure regions and AWS regions is that the data centers are located in multiple geographic areas. Each Azure region is paired with another region within the same geography to form a regional pair. The only exception is Brazil South, which is paired with a region outside its geography:



Azure paired regions

Take note that you can specify the region where you want to host deployed resources in Azure, but not all Azure services are available from every region. Check <https://azure.microsoft.com/en-gb/global-infrastructure/services/> to know more about Azure products available by region.

Availability Sets versus Availability Zones

In Azure, there are two concepts when we talk about availability, namely, which is the type of availability that is set and the zone which it is available.

An **Availability Set** is used to make sure the deployed VMs in Azure are distributed across multiple isolated hardware nodes in a cluster so that only a subset of your VMs is impacted in the case of failure, but your overall solution remains available and operational, it can provide SLA of 99.95%.

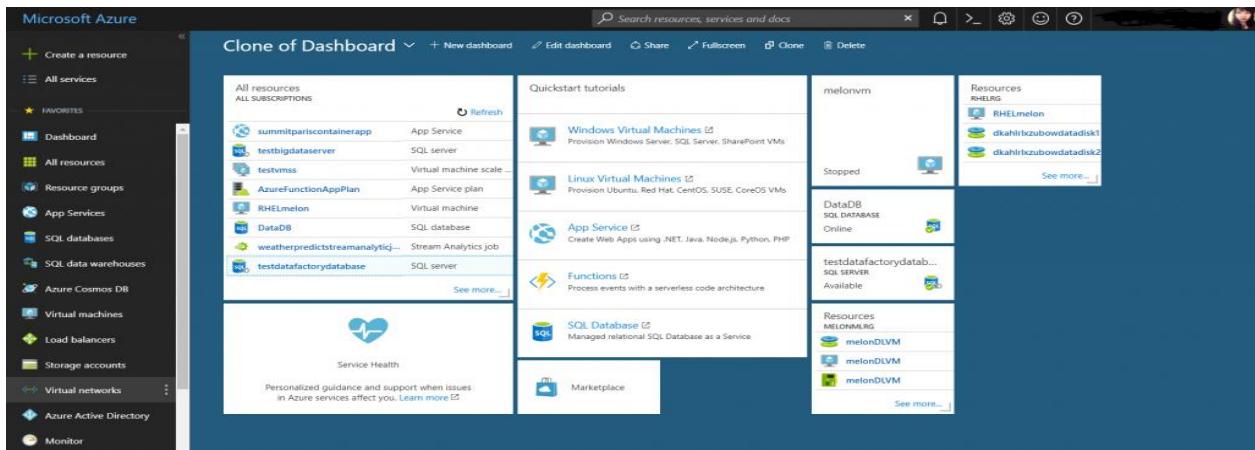
An **Availability Zone** consists of fault-isolated locations within an Azure region. One Availability Zone can consist of one or more data centers, it provides redundant power, cooling, and networking. Availability Zones are a great fit for mission-critical applications with requirements for intra-region resilience and fault tolerance in the case of failures of the data center., it can provide SLA of 99.99%.

Azure tools

Microsoft Azure provides different command-line tools and development tools to facilitate building, debugging, deploying, diagnosing, and especially managing scalable and elastic apps in Azure. Let's take a look at each of them.

Azure Portal

We can use <https://portal.azure.com> to navigate to Azure's new portal as shown as next:



Azure Portal

Azure command-line interface (Azure CLI)

Azure CLI, which was previously named Azure xPlat CLI, is an open-source, cross-platform, shell-based command-line interface designed for scripting and automating the creation and management of resources in Azure. Azure CLI works on Windows, Linux, macOS, and Docker containers. To install and configure Azure CLI under different OS, check the following page at <https://docs.microsoft.com/en-us/cli/azure/install-azure-cli?view=azure-cli-latest>.

You can test whether the installation was successful using the following command from your command line:

```
az --help
```

The following screenshot shows sample output:

```
Windows PowerShell
Copyright (C) Microsoft Corporation. Tous droits réservés.

PS C:\Users\mqin> az

Welcome to Azure CLI!
-----
Use `az -h` to see available commands or go to https://aka.ms/cli.

Telemetry
-----
The Azure CLI collects usage data in order to improve your experience.
The data is anonymous and does not include commandline argument values.
The data is collected by Microsoft.

You can change your telemetry settings with `az configure`.

AZURE

Welcome to the cool new Azure CLI!

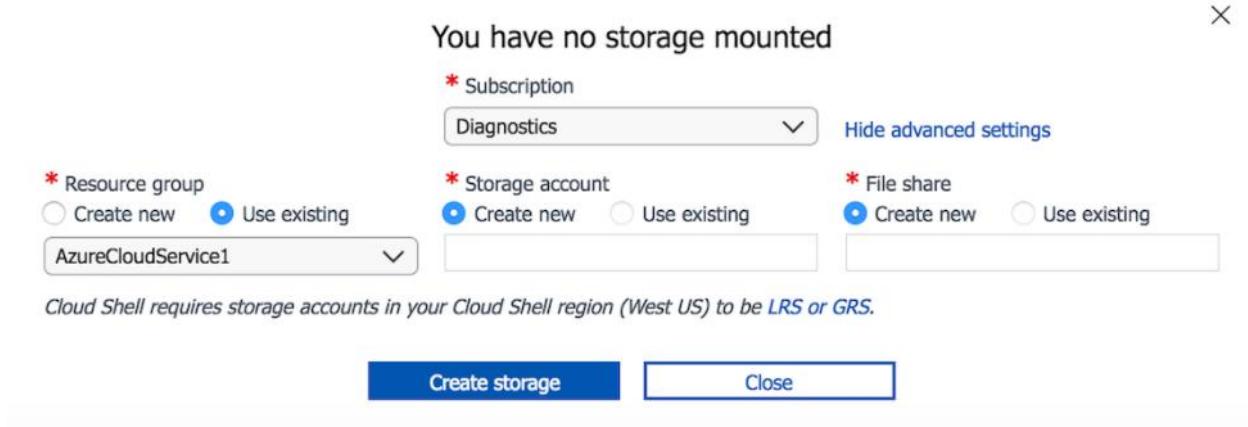
Here are the base commands:

account      : Manage Azure subscription information.
acr          : Manage Azure Container Registries.
acs          : Manage Azure Container Services.
ad           : Manage Azure Active Directory Graph entities needed for Role Based Access Control.
advisor      : Manage Azure Advisor.
aks          : Manage Azure Kubernetes Services.
appservice   : Manage App Service plans.
backup       : Commands to manage Azure Backups.
batch        : Manage Azure Batch.
batchai      : Batch AI.
billing      : Manage Azure Billing.
cdn          : Manage Azure Content Delivery Networks (CDNs).
cloud         : Manage registered Azure clouds.
cognitiveservices: Manage Azure Cognitive Services accounts.
configure    : Display and manage the Azure CLI 2.0 configuration. This command is interactive.
consumption  : Manage consumption of Azure resources.
container    : (PREVIEW) Manage Azure Container Instances.
cosmosdb     : Manage Azure Cosmos DB database accounts.
disk         : Manage Azure Managed Disks.
```

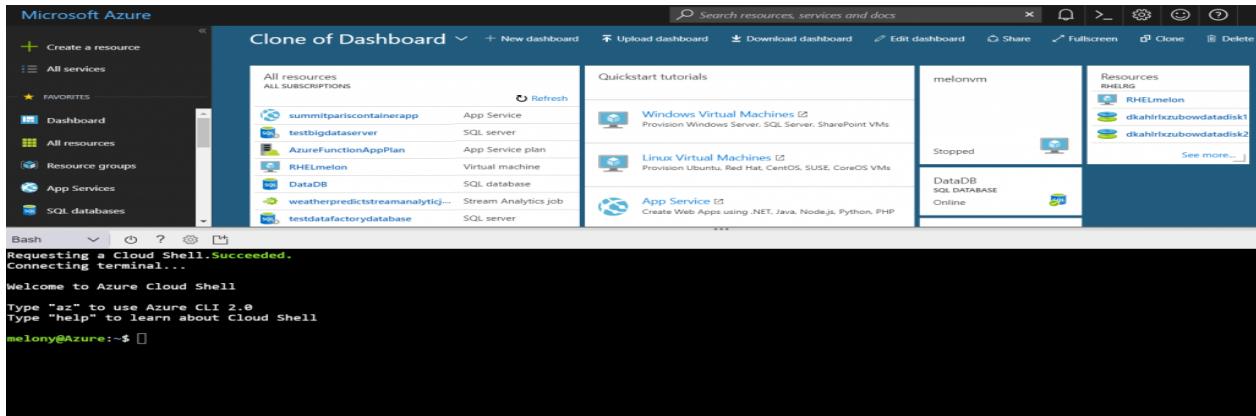
Azure Cloud Shell

Azure Cloud Shell is an interactive browser-based shell command line to manage cloud resources in Azure. Cloud Shell comes preinstalled with some popular command-line tools and language support such as Azure CLI, Bash, npm, mvn, git, and Docker.

The first time you launch Cloud Shell from the Azure Portal, it will create a resource group, storage account, and file share on your behalf; this is a one-time step and will be automatically attached for all sessions:



You can access Cloud Shell from shell.azure.com or via the Azure Portal (the following screenshot shows accessing Cloud Shell via Azure Portal):



Azure PowerShell

Azure PowerShell is one of the most powerful tools developed by Microsoft, and contains a set of modules within PowerShell that provide cmdlets to manage cloud resources in Azure. Azure PowerShell contains two different modes, which are defined using the Azure Service Manager or Azure Resources Manager modules, and which we explained previously. Azure PowerShell works on Windows, macOS, and Linux. To install and configure Azure PowerShell check the address: <https://docs.microsoft.com/en-us/powershell/azure/install-azurerm-ps?view=azurermps-5.5.0>.

Azure SDK

Azure SDK helps developers to deploy infinitely scalable applications and APIs, configure diagnostics, create and manage app service resources, and integrate data from Visual Studio. Currently, Azure SDK is available in many popular development languages such as .NET, Java, Node.js, Python, Ruby, PHP, JavaScript, and Swift.

To download Azure SDK in different programming languages, check <https://azure.microsoft.com/en-au/downloads/>.

Microsoft Azure provides also some Azure code samples at <https://azure.microsoft.com/en-us/resources/samples>.

Azure RESTful API

Most Azure service REST APIs have client libraries that provide a native interface for using Azure services. **Representational State Transfer (REST)** APIs are service endpoints that support sets of HTTP operations (methods), which provide create, retrieve, update, or delete access to the service's resources. The Azure RESTful API is available in .NET, Java, Node.js, Python, and Azure CLI 2.0 SDK. Azure also provides a great way to secure your REST requests by registering your client application with **Azure Active Directory (Azure AD)**. You can refer to the page of REST API browser which is currently in the preview state from <https://docs.microsoft.com/en-us/rest/api/>.

Similar to Azure PowerShell, Access APIs endpoint of Azure Resource Manager uses <https://management.azure.com/> and the Azure classic deployment model uses <https://management.core.windows.net/>. Both of them should add an api-version query-string parameter at the end, which is mandatory. Below is an example of the Get VM scale set in current subscription:

GET

<https://management.azure.com/subscriptions/{subscriptionId}/resourceGroups/{resourceGroupName}/providers/Microsoft.Compute/virtualMachineScaleSets/{vmScaleSetName}/virtualmachines/{instanceId}?api-version=2017-12-01>

ARM templates

An ARM template is a JSON file that defines the required resources in a declarative format to deploy IT solutions in a quick way in Azure, which is an excellent implementation of Infrastructure as a Code. It uses the definition file to define the infrastructure for your IT resource in the cloud. Similar to CloudFormation for AWS, ARM templates for Microsoft Azure make the infrastructure version possible. It is a simple and effective way to manage infrastructure resources in the cloud. The best practice recommended by Microsoft is to implement different ARM templates for different environments such as test, staging, and production.

You can find many useful sample ARM templates at: <https://azure.microsoft.com/en-us/resources/templates/>.

Or go to GitHub at: <https://github.com/Azure/azure-quickstart-templates>.

Azure developer tools

There are some other related Azure developer tools to support development in Azure with Microsoft solution or other open source solutions such as Visual Studio Tools for Azure, PowerShell Tools for Visual Studio, Storage Explorer, Docker Tools, and Azure Service Fabric Tools. Check the following address to know more about installation and configuration of these tools: <https://azure.microsoft.com/en-in/tools/>.

Overview of Microsoft Azure core services

Microsoft Azure is broken down into several high-level groupings of services. So far, there are more than 100 services in Microsoft Azure. They're generally grouped as Azure Compute, Azure Networking, Azure Storage, Azure Data and Analytics services, Azure Backup, and Azure Disaster Recovery.

You can take a quick look at the following link to take a global view of the latest Azure services: <https://azure.microsoft.com/en-us/resources/infographics/azure/>.

Microsoft also provides a search page so that users can use it to browse the latest products in different Azure categories as follows: <https://azure.microsoft.com/en-us/services/>.

Azure Compute services – IaaS versus PaaS

Azure provides different hosting models such as running applications on virtual servers or containers, or in a serverless computing environment, and each provides a different set of services. The following are some of the hosting models:

- **Virtual Machines** are the most important IaaS offering provided by Microsoft Azure. Different from physical machines, a virtual machine is a machine based on virtualization technology; they can be Windows-based or Linux-based VMs.
- **Virtual Machine Scale Sets** is a managed VM pool that contains a set of identical VMs. All VMs in VM scale sets with the same configuration is designed to improve scalability and availability.
- **App Service** contains PaaS offerings such as web apps, mobile apps, API apps, and logic apps in the same app service plan to provide a managed hosted environment.
- **Cloud Services** is a deployment solution with more control of the OS than App Service, there are two versions, IaaS cloud services; and PaaS cloud services.
- **Service Fabric** is a PaaS service which is designed for building packaging, deploying, and managing scalable and reliable microservices.

- **Azure Functions** is a **Function as a Service (FaaS)** service which provides serverless functions in the cloud. They are designed for deploying small pieces of code in different languages and in the cloud without managing the infrastructure.
- **Azure Batch** allows running large-scale parallel and high-performance computing applications efficiently in the Azure cloud.
- **Azure Container Service (ACS)** is designed for deploying and managing fully portable application containers and container orchestrators solutions such as Docker Swarm, DC-OS, as well as Kubernetes in the Cloud. Microsoft Azure also has a managed Kubernetes managed cluster which is known as **Azure Kubernetes Service (AKS)**.

Azure Networking

Azure Networking provides the following connections to connect your virtual machines, PaaS cloud services, and on-premise infrastructures:

- **Azure Virtual Networks**: This enables you to deploy isolated networks in the cloud to securely connect Azure resources to each other
- **Azure ExpressRoute**: This helps you to create a dedicated high-speed connection from your on-premise data center to Azure
- **VPN Gateway**: The virtual private network gateway is used to send network traffic between Azure virtual networks and on-premise locations, and also between virtual networks within Azure
- **Traffic Manager**: This provides DNS-level load balancing for applications that need to be high-availability
- **Load Balancer**: This provides level 4 load balancing features distributing network traffic to your applications
- **Azure DNS**: This is a domain name resolution service to manage DNS records for both Azure services and external resources

Azure Storage

Azure Storage is a cloud storage service that is durable, available, and scalable. Azure storage provides the following four types of storage:

- **Blob storage:** This is an object-based storage used to store documents, media files, or even application installers
- **Table storage:** This is a NoSQL key-attribute data store that is designed for semi-structured data, which allows for rapid development and fast access to large quantities of data
- **Queue storage:** This provides reliable messaging for workflow processing and for communication between components of Azure services
- **File storage:** This offers shared storage for legacy applications using the standard SMB protocol

Data and analytics services

Azure provides different databases and analytics services to help migrate your data to the cloud. Azure takes care of scalability, backup, and high availability through the following services:

- **SQL Database:** This is a managed database service, which is different from AWS RDS, which is a container-based service with a different database engine.
- **Azure SQL Data Warehouse:** This is a cloud-based, scaled-out data warehouse in the cloud, which can process massive volumes of data, both relational and nonrelational.
- **CosmosDB:** This is designed as a globally distributed database, which allows you to use key-value, graph, column-family, and document data in one service. Multi-model and globally distributed is the most important aspect of CosmosDB. It independently and elastically scales storage and throughput at any time, anywhere across the globe, making it perfect for your serverless applications.
- **Azure Redis Cache:** This is a distributed, managed cache designed for building highly scalable and responsive applications by providing super-fast access to your data.
- **Azure Machine Learning:** This enables you to apply statistical models to data and perform predictive analytics in the cloud.
- **Azure Search:** This provides a fully managed search service in the cloud.
- **Azure Data Factory:** This is a cloud-based data integration service that orchestrates and automates the movement and transformation of data within the data pipeline.
- **Azure Data Lake Store:** This is an enterprise-wide hyper-scale repository for big data analytics workloads, which is naturally integrated with HDFS to support Hadoop-based analytics.

Backup services and disaster recovery

Azure provides a business continuity plan that includes disaster recovery for all your major IT systems without the expense of secondary infrastructure. The following are the core services:

- **Azure StorSimple:** This is an integrated storage solution that manages storage tasks between on-premise devices and Azure cloud storage in the case of failure
- **Azure Backup:** This provides a cloud-based back up and works with ASR to restore your data in the Azure cloud
- **Azure Site Recovery (ASR):** This orchestrates replication of on-premise virtual machines and physical servers in the Azure cloud and restores the backup

Administrative roles and role-based access control

Microsoft uses RBAC to let users manage permission levels on resources in Azure. It is strongly recommended to apply the least privilege principle for each role, which defines a set of permitted actions. For more information refer to: <https://docs.microsoft.com/en-us/azure/billing/billing-add-change-azure-subscription-administrator>.

Azure provides three subscription-level administrative roles which have basic access management permissions:

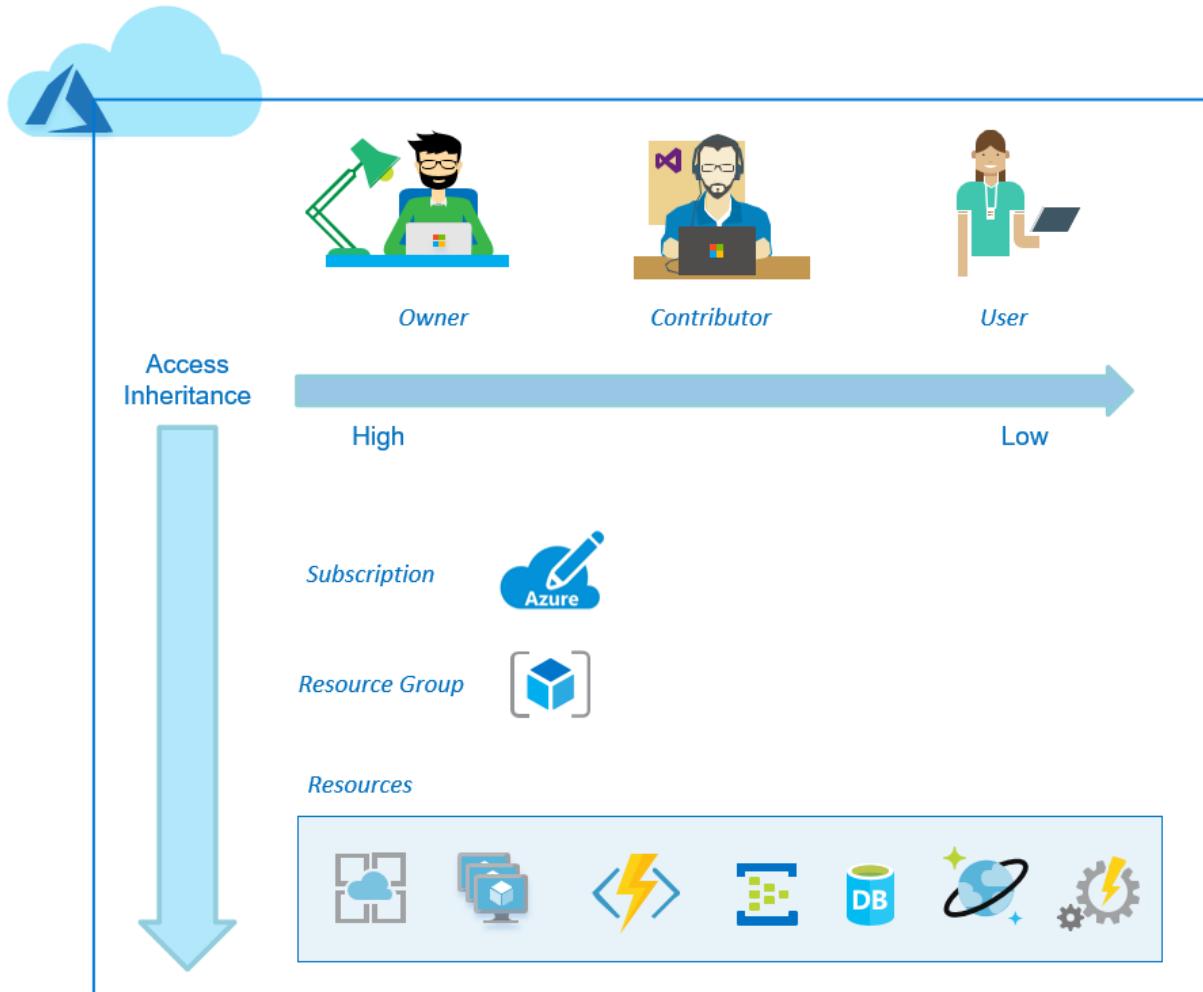
- Account administrator
- Service administrator
- Co-administrator

Azure RBAC is the latest access control system, and is recommended by Microsoft. It offers fine-grained access management, which has three basic roles that apply to all resource types in Azure:

- The **Owner** has full access to the defined resources, including the right to delegate access to others for these resources
- The **Contributor** can create and manage defined Azure resources, but can't grant access to others for these resources

- The **User** (reader) can view existing defined Azure resources

Each of these three roles has respected the scope, as shown in the following image:



The level and scope of management access with RBAC

Further reading

For more on the topics we have covered in this chapter, read the following links:

Adopting Azure is the first stage in organizational maturity for an enterprise. By the end of this stage, people in your organization can deploy simple workloads to

Azure: <https://docs.microsoft.com/en-us/azure/architecture/cloud-adoption-guide/adoption-intro/overview>

When deploying enterprise workloads to the cloud, Azure Virtual Datacenter is a great approach that helps IT organizations and business units balance governance but with developer agility. I recommend you read: <https://azure.microsoft.com/en-us/resources/azure-virtual-datacenter/en-us/>

The following is a great guide on how to identify and plan the migration of applications and servers to Azure using the lift and shift method, minimizing any additional development costs while optimizing cloud hosting options in Microsoft Azure: <https://azure.microsoft.com/en-us/resources/azure-virtual-datacenter-lift-and-shift-guide/en-us/>

The naming rules and restrictions for Azure resources and a baseline set of recommendations for naming conventions: <https://docs.microsoft.com/en-us/azure/architecture/best-practices/naming-conventions#naming-rules-and-restrictions>

Implementing and Managing Azure Virtual Machines

Azure VMs, which are **virtual machines (VMs)** provided by Microsoft Azure, have many advantages over traditional on-premise physical computers. In Azure, we can choose the deployment model, which can be classic or ARM VMs, before deploying them. Since Microsoft retired all the resources in the ASM model in the latest updates of the exam and has become more and more focused on ARM, in this book, all the resources that we deploy in Azure are with the **Azure Resource Manager (ARM)** model.

In the following chapter, we'll learn the basics of Azure VMs—that is, how to plan and deploy them with the Azure Portal—and the ARM template via Azure Portal, Visual Studio, Azure PowerShell, and Azure CLI.

The principles of Azure VMs

Microsoft Azure provides cloud-based VMs to help users deploy their workloads with more control of the system, such as to host custom services and applications, which is more flexible

than other Azure services. This module introduces how to plan, deploy, and monitor Azure VMs in different ways.

Planning and deploying Azure VMs

Compared to a traditional on-premises virtual machine, a cloud-based VM allows you to quickly scale up and down as requested; you don't need to buy and maintain the physical hardware that runs the virtual machine. Another great advantage is that you only pay for what you use in the cloud. In practice you, still need to maintain the virtual machine by configuring, patching, and maintaining the software that runs on the VM. However, its flexible pricing makes it more efficient in terms of time and costs.

Identifying the workloads

Before deploying a virtual machine in Azure, we will always start by identifying the workloads, whether the best deployment solution for the target workloads is on Azure VM or maybe on other Azure offerings.

In real life, organizations thinking about migrating their existing application to Azure quickly should not only take into account the technical concerns but also the financial aspects. Besides, certain types of workloads are a great fit for hosting in an Azure IaaS environment, for example, when you need a high flexibility to control your OS and don't mind higher administration efforts than other PaaS offerings in Azure.

Hyderabad

However, not every application is always a suitable fit for the cloud, as the following case :

- Certain low volumes or limited growth workloads where it might be cheaper to run the service or applications on commodity hardware on-premises
- Certain regulated environment workloads where the type of data is more sensitive or credentials requested by an organization needs to be kept on-premises or using other private cloud platforms such as Open Stack or the extension of the public cloud such as Azure Stack or VMware cloud on AWS.

Choosing the appropriate Azure VM sizing

If you decide to go further with Azure VM, you should choose the SKU of the virtual machine or the size of virtual machine of the Azure virtual machine provided by Microsoft Azure. In

Azure, the SKU or sizing is based on a variety of options for the number and speed of its processors (VCPU), amount of memory (RAM), the number of data disks you can attach to it, the maximum size of a temporary disk, IOPS, and the type of disks for the operating system. Generally, when the VM sizes support premium storage, which uses **solid-state drive (SSD)**, the maximum aggregate disk I/O performance would be better than standard storage with a **hard disk drive (HDD)**.

Virtual machines are available in several different sizes. When your requirements change, it is easy to resize the VM, which means you can use more advanced VM configurations, such as a more powerful CPU or larger RAM.

You can choose the appropriate size depending on your technical requirements. Try to balance the appropriate size of VMs and the number of VMs in your project. In real life, very often, the final decision on the size of VMs and number of instances for Dev/Test or the production environment would be made after a period of workload testing.

The following are the available categories of Azure virtual machines that are available so far:

Series	VM Size Family	User case
A Series	Entry-level economical VMs for dev/test	Development and test servers, low traffic web servers, small to medium databases, servers for proof-of-concepts, and code repositories.
B Series	Base level performance for workload, ability to burst CPU performance up to 100% of the CPU	Ideal for workloads that do not need the full performance of the CPU continuously, like web servers, small databases and development and test environments
D Series	General purpose compute	most applications, relational databases, in-memory caching, and analytics
Dv2 Series	Next generation general purpose compute	Ideal for applications that demand faster CPUs, better local disk performance or higher memories, and offer a powerful combination for many enterprise-grade applications
E Series	Optimized for in-memory hyper-threaded applications	SAP HANA, SAP S/4 HANA, SQL Hekaton and other large in-memory business critical workloads
F Series	Compute optimized virtual machines	Batch processing, web servers, analytics, and gaming
G Series	Memory and storage optimized virtual machines	Large SQL and NoSQL databases, ERP, SAP, and data warehousing solutions
H Series	High performance virtual machines	High performance computing, batch processing, analytics, molecular modeling, and fluid dynamics
L Series	Storage optimized virtual machines	NoSQL databases such as Cassandra, MongoDB, Cloudera, and Redis. Data warehousing applications and large transactional databases are great use cases as well
M Series	Largest memory optimized virtual machines	SAP HANA, SAP S/4 HANA, SQL Hekaton and other large in-memory business critical workloads requiring massive parallel compute power
N Series	GPU enabled virtual machines	Graphics rendering, video editing, remote visualization, high performance computing, and analytics

Azure VM storage options

Every Azure VM generally has two disks:

- An **operating system disk**, which is registered as a SATA drive and labeled as the C: drive for Windows machines and the repository /dev/sda for Linux machines, by default.
- A **temporary disk**, which is labeled as the D: drive for Windows that stores the pagefile.sys file. For a Linux machine, the repository /dev/sdb is the temporary disk.

Besides these two disks, Azure VMs can be attached to a number of **data disks**. The operating system disk is created from a VM image; the operating system disk and data disks are **virtual hard disks (VHDs)** stored in a page blob in an Azure storage account.

Managed disks versus unmanaged disks

Azure **unmanaged disks** are the disks created and managed by service administrators.

Azure **Managed Disks** are the disks that allow Microsoft Azure to handle the disk management of the IaaS VMs.

Compared to unmanaged disks, Azure-managed disks provide better scalability while scaling the VMs with VMSS (scale sets) and breaks the limit of IOPS per storage account, that is, Azure has a limit of 20,000 IOPS per storage account which will impact the number of VMs that can be created per storage account. Azure Managed Disks are recommended by Microsoft for storing persistent storage of data while creating Azure VMs.

Both of them have standard and premium pricing tiers. A standard tier is based on HDD. A premium tier is based on high-performance SSD to support I/O intensive workloads.

Unmanaged disks are available for **locally redundant storage (LRS)**, **zone-redundant storage (ZRS)**, **geo-redundant storage (GRS)**, and **read-access geo-redundant storage (RA-GRS)**. At the time of writing this book managed disks are only available for LRS.

Azure reserved VM instances (RIs) versus pay-as-you-go instances

Compared to pay-as-you-go prices, Azure also provides a way to cut down VM costs in a significant way by purchasing the **reserved VM instances (RIs)**, which link to a 1-year or 3-year engagement on both Windows and Linux VMs.

Purchasing Azure-reserved VM instances can help users save up to 72% VMs costs. Additionally, it accumulates the **Azure Hybrid Benefit**, which can help users save up to 82% VM costs.

Deploying an Azure VM

To facilitate the deployments of different workloads using an Azure virtual machine, Microsoft Azure offers different ways to release the deployment. Users can deploy their virtual machines via Azure Portal, Azure PowerShell, Azure CLI, Azure Cloud Shell, or ARM Template.

The deployment usually starts from choosing an OS image from the Azure Marketplace. The Azure Marketplace provides images of various Microsoft and Linux operating systems, such as CentOS, Debian, Ubuntu, and so on, and also provides preconfigured products with a **ready-to-use** image. Microsoft and Microsoft's third-party partner provides various popular image solutions, such as Windows Server 2016 Data Center, Red Hat Enterprise, SUSE Linux Enterprise, the Data Science Virtual Machine, and so on.

Creating Azure VMs

There are many ways to create an Azure VM, for example, via Azure Portal, Azure PowerShell, or Azure CLI.

Creating Azure VMs via the Azure Portal

Let's take a look at how to create an Azure VM via the Azure Portal with a Windows image and a Linux image in the Azure Marketplace.

Creating an Azure VM with a Windows image

Via the Azure Portal, click on Create a resource and then on the Compute option, then choose the appropriate Windows image (**Windows Server 2016 Datacenter**, which is the latest one). The following screenshot shows the first page that will appear when you start to deploy an Azure Windows VM; be careful to choose the ARM as a deployment model, which will register your resource with Azure Resource Manager:



Windows Server 2016 is a comprehensive server operating system designed to run the applications and infrastructure that power your business. It includes built-in layers of security and innovation to help you run traditional and cloud-native applications with confidence. This Server with Desktop Experience image includes all roles including the graphical user interface (GUI).

This image can be used with [Azure Hybrid Benefit for Windows Server](#).

Legal Terms

By clicking the Create button, I acknowledge that I am getting this software from Microsoft and that the [legal terms](#) of Microsoft apply to it. Microsoft does not provide rights for third-party software. Also see the [privacy statement](#) from Microsoft.



PUBLISHER

Microsoft

[Documentation](#)

Select a deployment model ⓘ

Resource Manager

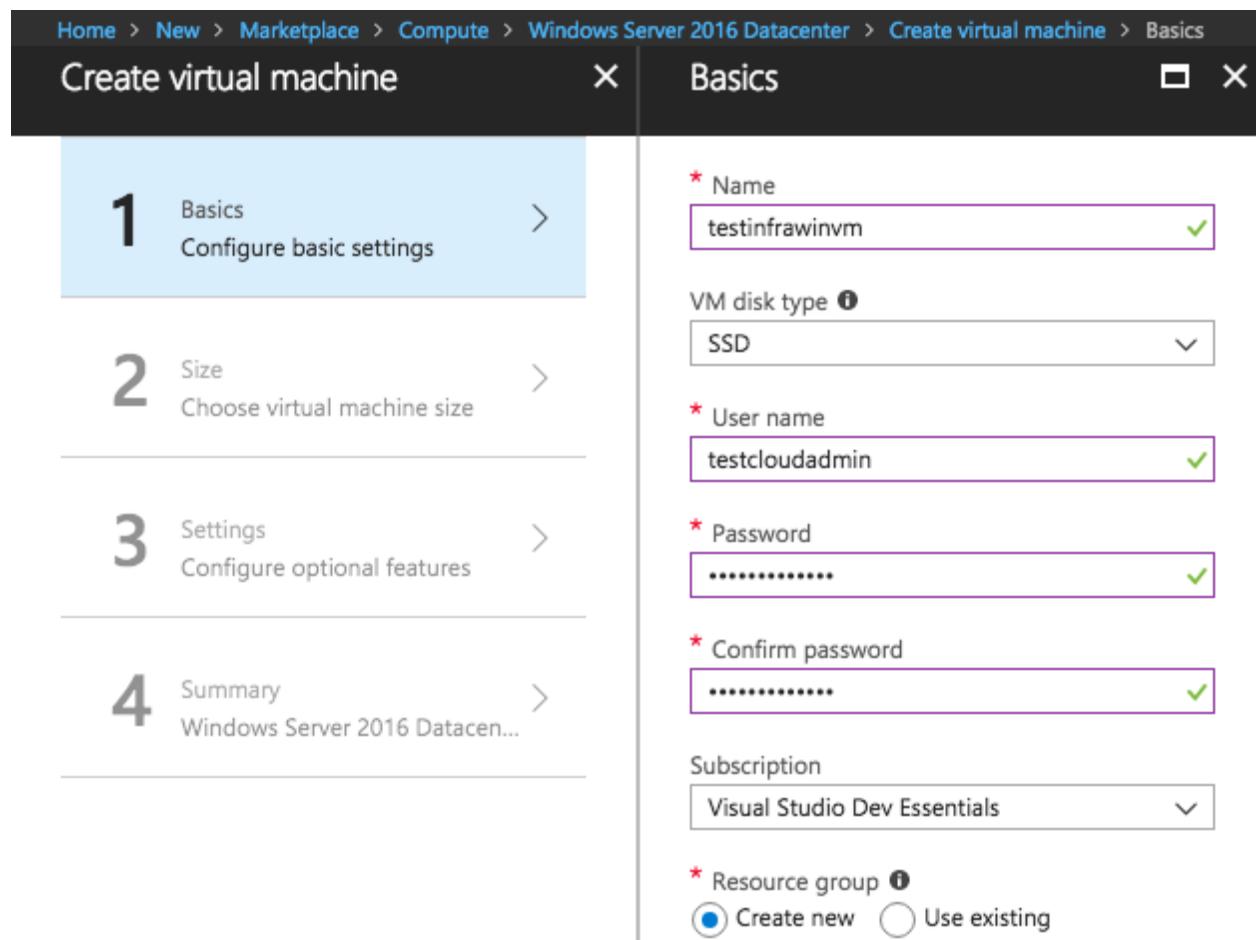
Create

Choosing the appropriate deployment model to start deploying a windows VM

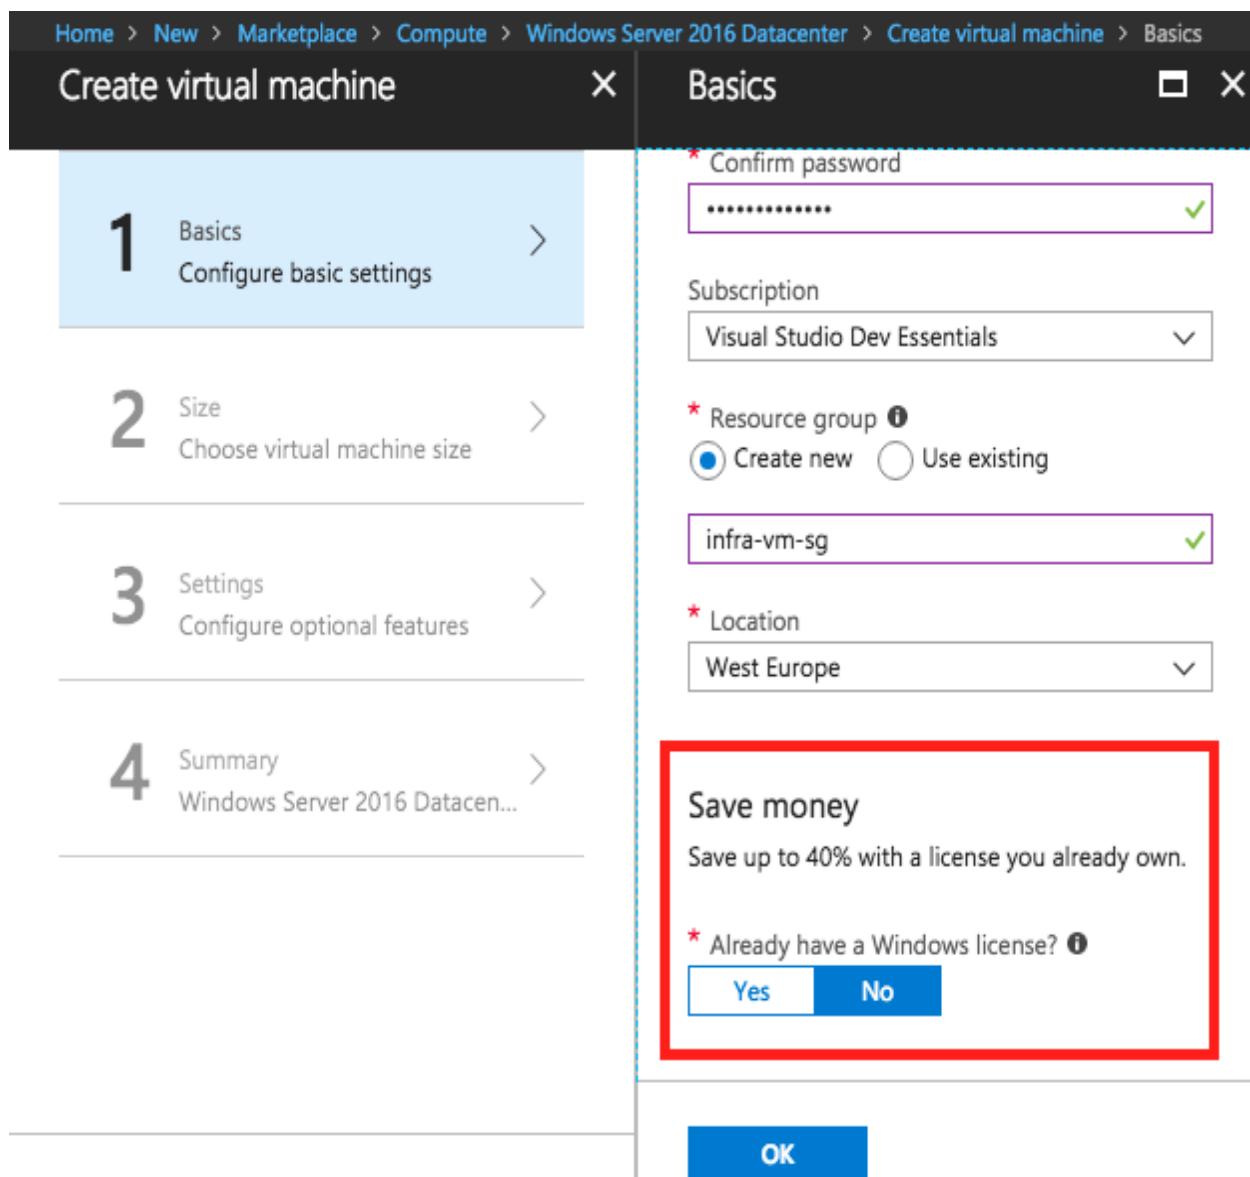
Then, you should fill in the necessary information in the **Basics** blade. The VM Disk type will affect the proposed pricing plan in the next step, which lets users choose the appropriate size of the VM. Azure provides the following types of disks:

- The **Premium disks (SSD)** are backed by SSDs, provide consistent, low-latency performance, and are ideal for I/O-intensive applications and production workloads
- The **Standard disks (HDD)** are backed by magnetic drives and are designed for applications with infrequently accessed data

The **User name** and **Password** will be used to connect to the virtual machine. Usernames can be a maximum of 20-characters long, and the password must be at least 12-characters long and should have lower characters, upper characters, at least a digit, and a special character. While creating Azure VMs, users should pay attention to how to choose the most appropriate resource group, the subscription of the organization, and the location closest to your geography to reduce latency. The following screenshot is an example of the information in the **Basics** blade:



All Microsoft software installed in the Azure virtual machine environment must be licensed correctly. Microsoft Azure provides Azure Hybrid Benefit for the Windows server, which allows users to use on-premises Windows Server licenses and run Windows virtual machines on Azure at a reduced cost—this offer allows users to save up to 40% in costs. To obtain the benefits of Azure Hybrid, just confirm that you already have an on-premise license, as indicated in the following screenshot:



The screenshot shows the 'Create virtual machine' wizard in the Azure portal. The current step is 'Basics'. The left sidebar lists four steps: 1. Basics (Configure basic settings), 2. Size (Choose virtual machine size), 3. Settings (Configure optional features), and 4. Summary (Windows Server 2016 Datacenter...).

In the 'Basics' step, the following fields are filled:

- Confirm password:** A masked password is entered.
- Subscription:** Visual Studio Dev Essentials.
- Resource group:** infra-vm-sg.
- Location:** West Europe.

A red box highlights a section titled 'Save money' with the subtext 'Save up to 40% with a license you already own.' It contains a question 'Already have a Windows license?' with 'Yes' and 'No' buttons. A blue 'OK' button is at the bottom right of the main form.

Microsoft Azure provides different pricing solutions with a range of predefined configuration options that correspond to different VM sizes. The different VM sizes indicate the different numbers and speed of its processors, different amounts of memory, a maximum number of network adapters or data disks that users can attach to it, and the maximum size of a temporary disk. As shown in the following screenshot, users should choose an initial VM size while deploying a new VM in Azure:

Home > New > Marketplace > Compute > Windows Server 2016 Datacenter > Create virtual machine > Choose a size

Create virtual machine

Choose a size

Browse the available sizes and their features

Prices presented are estimates in your local currency that include only Azure infrastructure costs and any discounts for the subscription and location. The prices don't include any applicable software costs. Recommended sizes are determined by the publisher of the selected image based on hardware and software requirements.

Supported disk type: SSD Minimum vCPUs: 1 Minimum memory (GiB): 0

★ Recommended View all		
DS1_V2 Standard	DS1 Standard	B1S Standard
1 vCPU	1 vCPU	1 vCPU
3.5 GB	3.5 GB	1 GB
4 Data disks	4 Data disks	2 Data disks
3200 Max IOPS	3200 Max IOPS	800 Max IOPS
7 GB Local SSD	7 GB Local SSD	4 GB Local SSD
Premium disk support	Premium disk support	Premium disk support

Select

Choosing an Azure VM initial size

Here is a list view of the same:

Choose a size

Browse the available sizes and their features

Search: Compute type: Show all compute types Disk type: SSD only vCPUs: 1

SKU	Type	Compute ...	vCPUs	GB RAM	Data Disks	Max IOPS	Local SSD	Premium ...	Addition...	EUR/Mo
B1s	Standard	General purpose	1	1	2	800	4 GB	SSD		€7.91
B1ms	Standard	General purpose	1	2	2	1600	4 GB	SSD		€15.25
B2s	Standard	General purpose	2	4	4	3200	8 GB	SSD		€30.49
B2ms	Standard	General purpose	2	8	4	4800	16 GB	SSD		€60.98
B4ms	Standard	General purpose	4	16	8	7200	32 GB	SSD		€121.97
B8ms	Standard	General purpose	8	32	16	10800	64 GB	SSD		€243.94
D2s_v3	Standard	General purpose	2	8	4	4000	16 GB	SSD		€75.29

Prices presented are estimates in your local currency that include only Azure infrastructure costs and any discounts for the subscription and location. The prices don't include any applicable software costs. Recommended sizes are determined by the publisher of the selected image based on hardware and software requirements.

Select

Choosing an Azure VM initial size (list view)

In the **settings** step, if you're not going to use the managed disk, you should specify a storage account to store disks, as shown in the following screenshot:

The screenshot shows the 'Storage' configuration step in the Azure portal. A red box highlights the 'Use managed disks' section, which contains a 'No' button (selected) and a 'Yes' button. Below this, another red box highlights the 'Storage account' section, which lists '(new) testmanageddisks'. The 'Network' section below it includes configurations for a virtual network, subnet, public IP address, and network security group.

Storage

Use managed disks ⓘ

No Yes

* Storage account ⓘ
(new) testmanageddisks >

Network

* Virtual network ⓘ >
(new) testloTRG-vnet

* Subnet ⓘ >
default (10.1.0.0/24)

* Public IP address ⓘ >
(new) test-ip

* Network security group (firewall) ⓘ >
(new) test-nsg

Choosing storage account while using non-managed disks

Select your storage option:

Home > New > Marketplace > Compute > Windows Server 2016 Datacenter > Create virtual machine > Settings

Create virtual machine X

Settings □ X

1 Basics Done ✓

2 Size Done ✓

3 Settings Configure optional features >

4 Summary Windows Server 2016 Datacen... >

High availability

Availability zone i
None

Availability zones are not available for the chosen location and size.

* Availability set i
None >

Storage

Use managed disks i
Yes No

Network

* Virtual network i
(new) infra-vm-sg-vnet >

* Subnet i
default (10.1.0.0/24) >

OK

Choosing whether to use managed disks in Settings step

There are many options for creating a new Azure VM, such as **Virtual Network (VNet)**, **Subnet**, and **Network Security Group (NSG)**. Microsoft Azure manages a default configuration for a predefined VM template that is ready to use. You can change these options while creating a new VM, depending on your intentions and requirements.

After filling in all the necessary information, Microsoft Azure will summarize and validate these details, as follows:

Home > New > Marketplace > Compute > Windows Server 2016 Datacenter > Create virtual machine > Create

Create virtual machine Create

1 Basics Done Validation passed

2 Size Done

3 Settings Done

4 Summary Windows Server 2016 Datacenter...

Offer details
 Standard DS1 v2 by Microsoft 0.1138 EUR/hr Pricing for other VM sizes
[Terms of use](#) | [privacy policy](#)

Azure resource
 You may use your Azure monetary commitment funds or subscription credits for these purchases. Prices presented are retail prices and may not reflect discounts associated with your subscription.

Terms of use
 By clicking "Create", I (a) agree to the legal terms and privacy statement(s) associated with each Marketplace offering above, (b) authorize Microsoft to charge or bill my current payment method for the fees associated with my use of the offering(s), including applicable taxes, with the same billing frequency as my Azure subscription, until I discontinue use of the offering(s), and (c) agree

I give Microsoft permission to use and share my contact information so that Microsoft or the Provider can contact me regarding this product and related products.

Create Download template and parameters

Azure VM deployment validation

After deploying a Windows-based VM in Azure, the related resources such as vnet, nsg, and NIC are shown:

infra-vm-sg Resource group

Search (Ctrl+/
)

Overview Activity log Access control (IAM) Tags

Subscription (change) Visual Studio Dev Essentials Subscription ID: 5efcd8d-8da8-4993-9c36-4cc553... Deployments: 1 Succeeded

Tags (change)
Click here to add tags

Filter by name... All types All locations No grouping

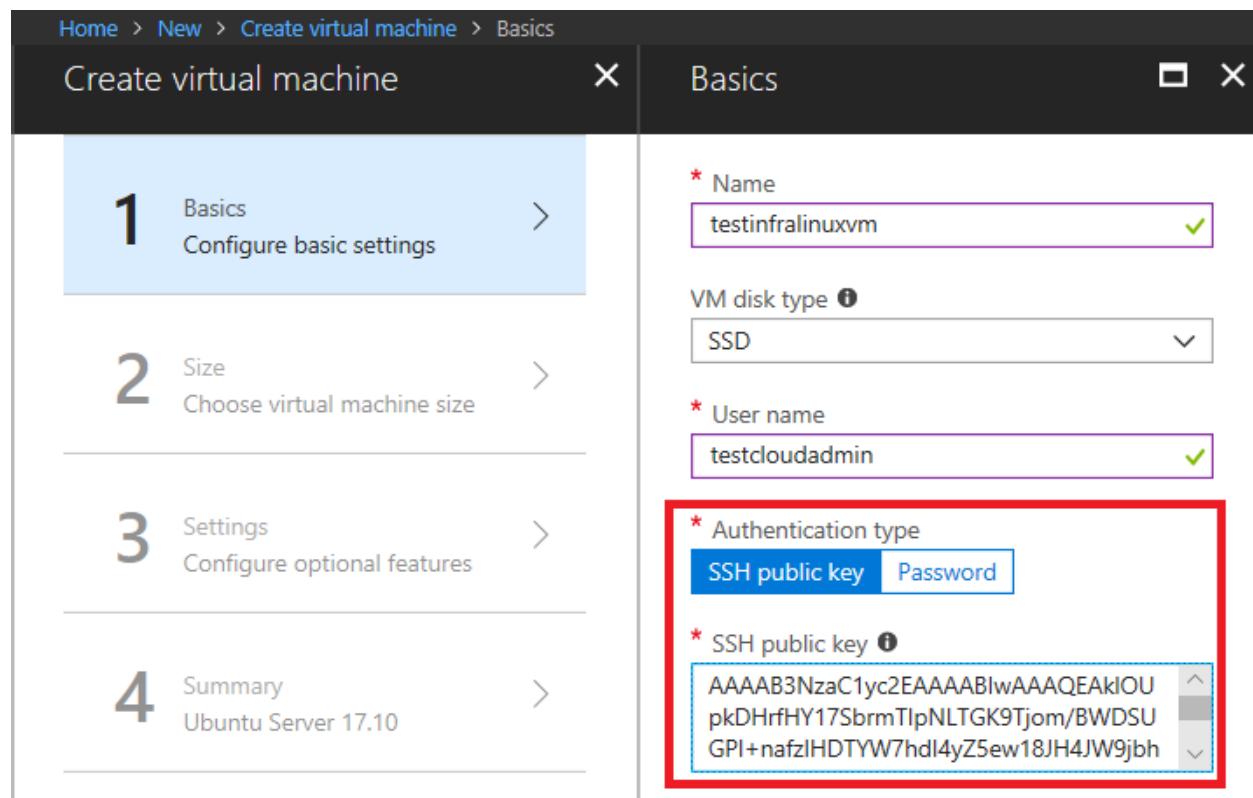
7 items Show hidden types

NAME	TYPE	LOCATION
infravmsgdiag469	Storage account	West Europe
infra-vm-sg-vnet	Virtual network	West Europe
testinfrawinvm	Virtual machine	West Europe
testinfrawinvn_OsDisk_1_a1ff150110e85436d9b621dbae14feed2	Disk	West Europe
testinfrawinvn459	Network interface	West Europe
testinfrawinvn-ip	Public IP address	West Europe
testinfrawinvn-nsg	Network security group	West Europe

Related resources after a successful Azure VM deployment

Creating an Azure VM with a Linux distribution image

While creating a Linux-based VM, Microsoft provides almost the same options. A little different from the approach to create a Windows-based VM, a Linux-based VM is the authentication type, which means that Azure allows users to choose between the password-based and SSH public key-based authentication types while creating Linux-based Azure VMs, as follows:



Choose between the password-based and SSH public key-based authentication

Users should provide an RSA public key in the single line format (starting with `ssh-rsa`) or multiline PEM format (the multiline SSH key must begin with `---- BEGIN SSH2 PUBLIC KEY ----` and end with `---- END SSH2 PUBLIC KEY ----`). You can generate SSH keys using `ssh-keygen` on Linux and macOS X, or PuTTYGen on Windows.

If you're using macOS or Linux, go to the following link to create your RSA key:

<https://docs.microsoft.com/en-us/azure/virtual-machines/linux/mac-create-ssh-keys>

If you're a Windows users, don't worry, go to the following link to get your RSA public key:

<https://docs.microsoft.com/en-us/azure/virtual-machines/linux/ssh-from-windows>

We will use the Azure Service Manager to deploy resources in Azure with the classic deployment model and the Azure Resource Manager to deploy resources in Azure with the ARM deployment model. That's why all the PowerShell commands in the ARM model are prefixed with RM.

Let's begin now:

1. To sign in to Azure via Azure PowerShell or Windows PowerShell ISE, use the following command:

Login-AzureRmAccount

Running the preceding command will display a popup to let you put in your Azure account name and password to pursue the authentication in Azure.

2. After logging into Azure with Azure PowerShell successfully, you'll get the following output:

```
PS C:\WINDOWS\system32> Login-AzureRmAccount
```

```
Account          : [REDACTED]
SubscriptionName : [REDACTED] - MPN
SubscriptionId   : f38e1d90-3a11-460c-a4d2-186e1660d993
TenantId         : 6494460e-8600-4edc-850f-528e8faad290
Environment      : AzureCloud
```

Output after login in Azure with Azure PowerShell successfully

3. Now, to get the list of Azure subscriptions associated with your account, use the following command:

Get-AzureRmSubscription

You will get the following output after executing the preceding command:

```
PS C:\WINDOWS\system32> Get-AzureRmSubscription

Name      : Visual Studio Dev Essentials
Id        : 5efcdf8d-8da8-4993-9c36-4cc5534b9563
TenantId  : 6494460e-8600-4edc-850f-528e8faad290
State     : Disabled

Name      : [REDACTED]
Id        : f38e1d90-3a11-460c-a4d2-186e1660d993
TenantId  : 6494460e-8600-4edc-850f-528e8faad290
State     : Enabled

Name      : Microsoft Azure [REDACTED]
Id        : 659714ac-e6c2-4fc2-84b8-e2cf5d1a6eba
TenantId  : 8265bf70-d0b6-438b-b01d-fa623087ca84
State     : Enabled

Name      : [REDACTED]
Id        : d9c16e03-bbea-45ed-b211-c99049536d51
TenantId  : 933c9cbe-35d3-4416-abbd-ddd1bca5879c
State     : Enabled

Name      : [REDACTED]
Id        : 0e95efb5-2252-441b-953e-9754cc7d5433
TenantId  : 933c9cbe-35d3-4416-abbd-ddd1bca5879c
State     : Enabled
```

List of all subscriptions in the current tenant/ [7730997544](tel:7730997544)

Ameerpet / Kondapur

Hyderabad

4. If it is not the target subscription, use the following command to choose the target subscription:

Select-AzureRmSubscription -SubscriptionName "<subscription name>"

5. To set the subscription context, using the following command:

Set-AzureRmContext -SubscriptionId "<subscription id>"

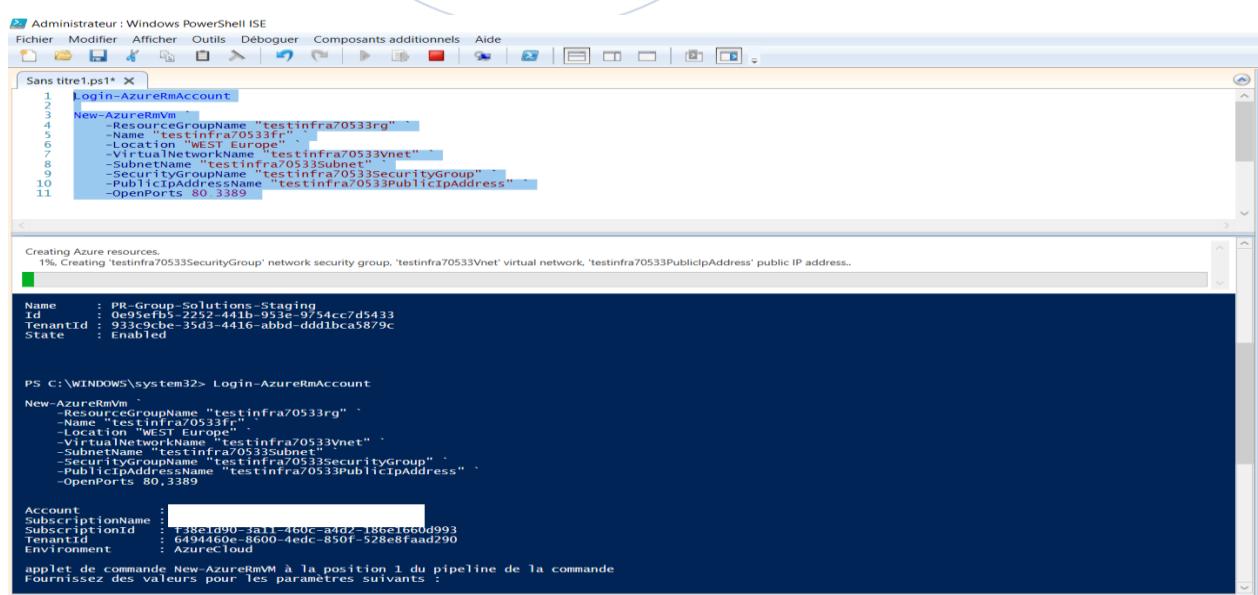
6. To start the deployment of the Windows image, we should collect the related information regarding the deployment:
 1. The virtual network and its subnet
 2. Public IP address (optional)
 3. Network interface card (NIC)
 4. NSG with a rule allowing inbound RDP traffic (open inbound traffic for the port 3389 of your Azure VM)
 5. OS admin credentials (it is recommended to store them in a variable)

7. Then, create the Azure VM using the following command:

New-AzureRmVm

```
-ResourceGroupName "testinfra70533rg"
-Name "testinfra"
-Location "WEST Europe"
-VirtualNetworkName "testinfra70533Vnet"
-SubnetName "testinfra70533Subnet"
-SecurityGroupName "testinfra70533SecurityGroup"
-PublicIpAddressName "testinfra70533PublicIpAddress"
-OpenPorts 80,3389
```

After authentication, it will start to create resources in Azure, as shown in the following screenshot:



```
PS C:\WINDOWS\system32> Login-AzureRmAccount
New-AzureRmVm
-ResourceGroupName "testinfra70533rg" -
-Name "testinfra" -
-Location "WEST Europe" -
-VirtualNetworkName "testinfra70533Vnet" -
-SubnetName "testinfra70533Subnet" -
-SecurityGroupName "testinfra70533SecurityGroup" -
-PublicIpAddressName "testinfra70533PublicIpAddress" -
-OpenPorts 80,3389

Creating Azure resources.
1% Creating 'testinfra70533SecurityGroup' network security group, 'testinfra70533Vnet' virtual network, 'testinfra70533PublicIpAddress' public IP address.

Name          : PR-Group-Solutions-Staging
Id           : 0f939cb5-2252-441b-953e-9754cc7d5433
TenantId     : 0319c9c2-35d3-4416-abbd-dd1bcfa5879c
State        : Enabled

PS C:\WINDOWS\system32> Login-AzureRmAccount
New-AzureRmVm
-ResourceGroupName "testinfra70533rg" -
-Name "testinfra70533fr" -
-Location "WEST Europe" -
-VirtualNetworkName "testinfra70533Vnet" -
-SubnetName "testinfra70533subnet" -
-SecurityGroupName "testinfra70533SecurityGroup" -
-PublicIpAddressName "testinfra70533PublicIpAddress" -
-OpenPorts 80,3389

Account       : 
SubscriptionName : 
SubscriptionId  : f38cfd90-3a11-460c-a4d2-18be1660d993
TenantId      : e49448de-8600-4edc-850f-528e8faad290
Environment   : AzureCloud

applet de commande New-AzureRmVm à la position 1 du pipeline de la commande
Fournissez des valeurs pour les paramètres suivants :
```

Creating Azure VMs via an Azure CLI

To create an Azure VM using a Windows image, for example, a Windows Server 2016 image via the Azure CLI, similar to Azure PowerShell, we should configure the login and subscription that we will use in Azure. You can use the following command:

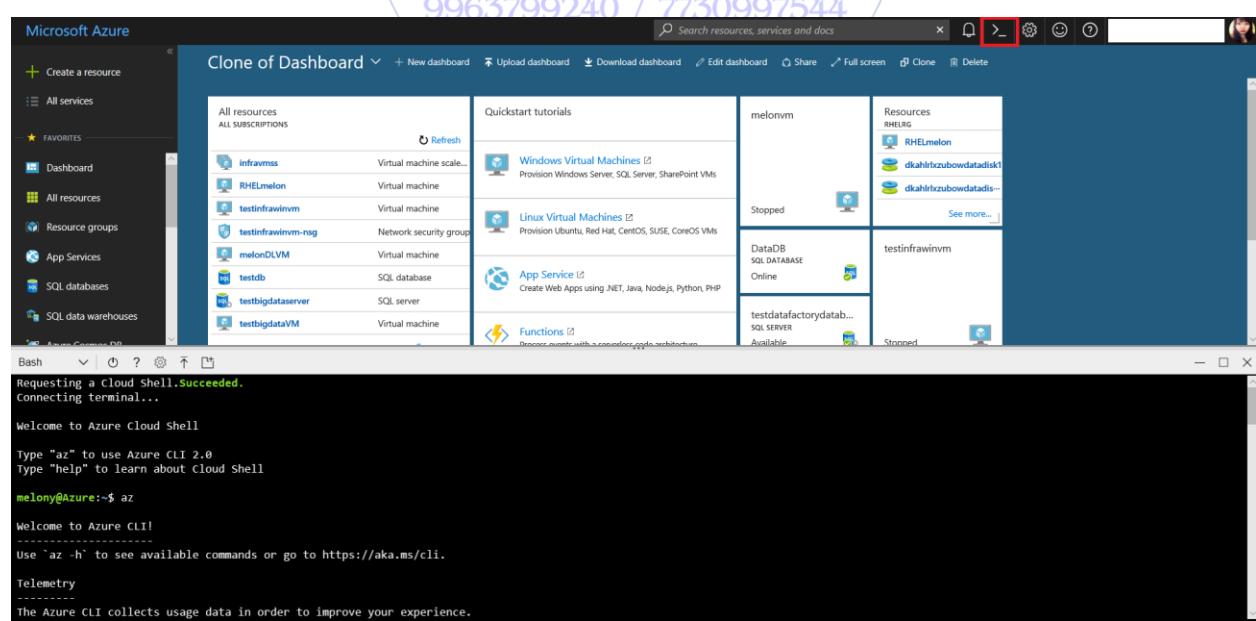
az login

Use the following command to set your subscription, and where you want, to deploy your Azure VM:

az account set –subscription <subscription name>

A great way to access Azure CLI is to use Azure Cloud Shell, which will let you always work with the latest Azure CLI commands without worrying about installing updates.

Launching the Cloud Shell via Azure Portal, as shown in the following screenshot:



Launching a Cloud Shell via Azure Portal

After launching Cloud Shell successfully, create a resource group using the following command:

```
az group create --name <resource groupname> --location <Azure region>
```

Specify the resource group that will host the Azure VM, its location following your intention. The example output will look like what's shown here in the following screenshot:

```
melony@Azure:~$ az group create --name testinfravmazurecli --location 'West Europe'
{
  "id": "/subscriptions/f38e1d90-3a11-460c-a4d2-186e1660d993/resourceGroups/testinfravmazurecli",
  "location": "westeurope",
  "managedBy": null,
  "name": "testinfravmazurecli",
  "properties": {
    "provisioningState": "succeeded"
  },
  "tags": null
}
```

Creating resource group via Azure CLI

To create the Azure VM, use the following command:

```
az vm create --resource-group <resource groupname> --name <VM name> --image
<Azure MarketplaceImage> --generate-ssh-keys
```

Currently, the valid images contain CentOS, CoreOS, Debian, openSUSE-Leap, RHEL, SLES, UbuntuLTS, Win2016Datacenter, Win2012R2Datacenter, Win2012Datacenter, and Win2008R2SP1.

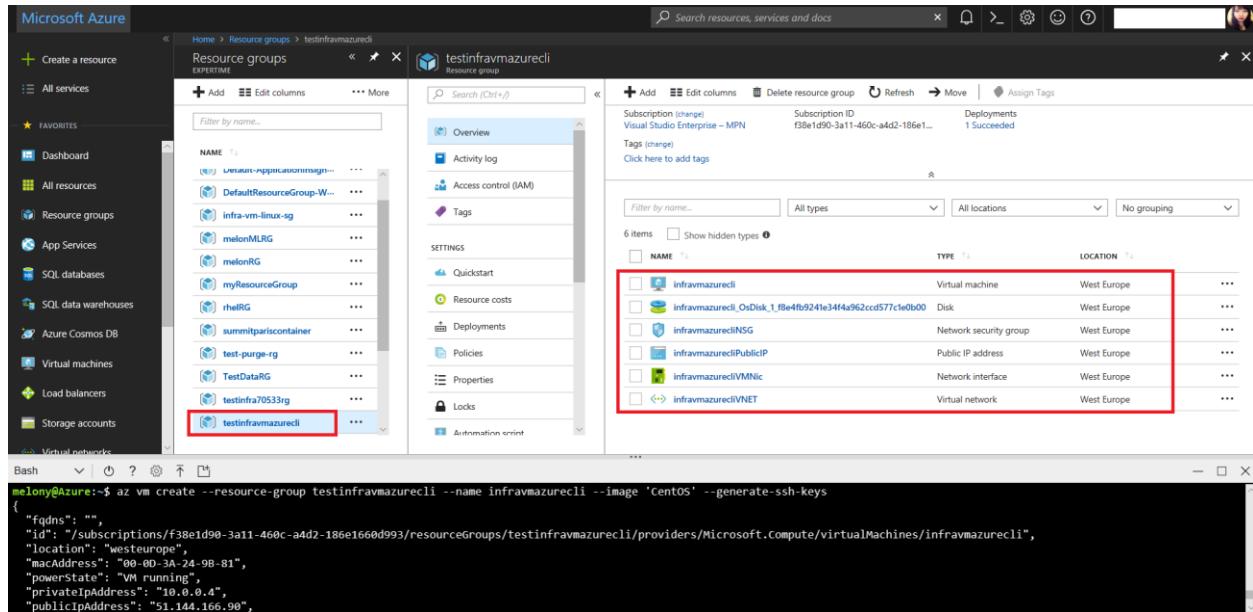
The following is an example command I launched to create a Azure VM by using the Azure CLI:

```
az vm create --resource-group testinfravmazurecli --name infravmazurecli --image
'CentOS' --generate-ssh-keys
```

The output of the preceding commands will be as follows, which means we have created a Linux VM via Azure CLI successfully:

```
melony@Azure:~$ az vm create --resource-group testinfravmazurecli --name infravmazurecli --image 'CentOS' --generate-ssh-keys
{
  "fqdns": "",
  "id": "/subscriptions/f38e1d90-3a11-460c-a4d2-186e1660d993/resourceGroups/testinfravmazurecli/providers/Microsoft.Compute/virtualMachines/infravmazurecli",
  "location": "westeurope",
  "macAddress": "00-0D-3A-24-9B-81",
  "powerState": "VM running",
  "privateIpAddress": "10.0.0.4",
  "publicIpAddress": "51.144.166.90",
  "resourceGroup": "testinfravmazurecli",
```

When returning to Azure Portal, we can find the resources that have been deployed in Azure successfully, as shown in the following screenshot:



```
melony@Azure:~$ az vm create --resource-group testinfravmazurecli --name infravmazurecli --image 'CentOS' --generate-ssh-keys
{
  "fqdns": "",
  "id": "/subscriptions/f38e1d90-3a11-460c-a4d2-186e1660d993/resourceGroups/testinfravmazurecli/providers/Microsoft.Compute/virtualMachines/infravmazurecli",
  "location": "westeurope",
  "macAddress": "00-0D-3A-24-9B-81",
  "powerState": "VM running",
  "privateIpAddress": "10.0.0.4",
  "publicIpAddress": "51.144.166.90",
```

Creating Azure VMs via the ARM template

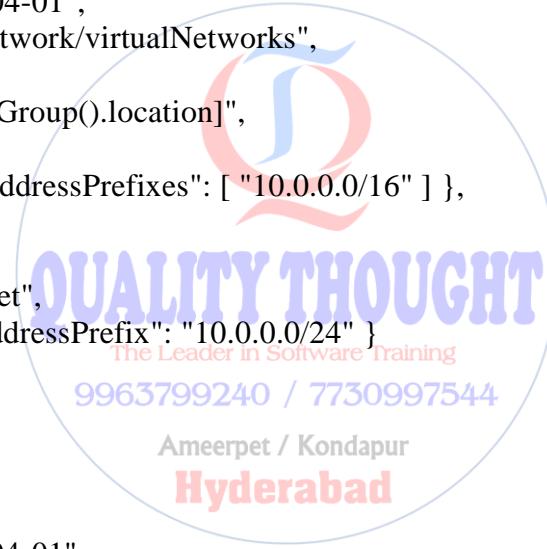
You can also create VMs using Azure Resource Manager templates, which facilitate the deployment process. Microsoft Azure provides different ways to deploy ARM templates, such as Azure Portal, Visual Studio, and Visual Studio Code. This capability is provided by Azure Resource Manager, which makes it possible to use a formatted **JSON file** and include definitions of all the Azure Resource Manager resources that are part of the deployment.

An ARM Template contains \$schema, contentVersion, parameters, variables, resources, and outputs. \$schema is defined by Microsoft Azure which is mandatory; other elements, except contentVersion, and resources, are optional. The following is a sample template to create a Linux-based Azure VM using managed disks:

```

"resources": [
    {
        "apiVersion": "2018-06-01",
        "type": "Microsoft.Network/publicIPAddresses",
        "name": "myPublicIPAddress",
        "location": "[resourceGroup().location]",
        "properties": {
            "publicIPAllocationMethod": "Dynamic",
            "dnsSettings": {
                "domainNameLabel": "testinfradns"
            }
        }
    },
    {
        "apiVersion": "2018-04-01",
        "type": "Microsoft.Network/virtualNetworks",
        "name": "myVNet",
        "location": "[resourceGroup().location]",
        "properties": {
            "addressSpace": { "addressPrefixes": [ "10.0.0.0/16" ] },
            "subnets": [
                {
                    "name": "mySubnet",
                    "properties": { "addressPrefix": "10.0.0.0/24" }
                }
            ]
        }
    },
    {
        "apiVersion": "2018-04-01",
        "type": "Microsoft.Network/networkInterfaces",
        "name": "myNic",
        "location": "[resourceGroup().location]",
        "dependsOn": [
            "[resourceId('Microsoft.Network/publicIPAddresses/', 'myPublicIPAddress')]",
            "[resourceId('Microsoft.Network/virtualNetworks/', 'myVNet')]"
        ],
        "properties": {
            "ipConfigurations": [
                {
                    "name": "ipconfig1",
                    "properties": {
                        "privateIPAllocationMethod": "Dynamic",
                        "publicIPAddress": {
                            "id": "[resourceId('Microsoft.Network/publicIPAddresses','myPublicIPAddress')]"
                        }
                    }
                }
            ]
        }
    }
]

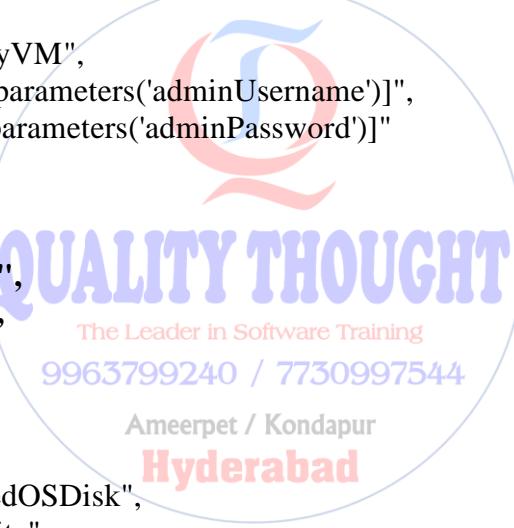
```



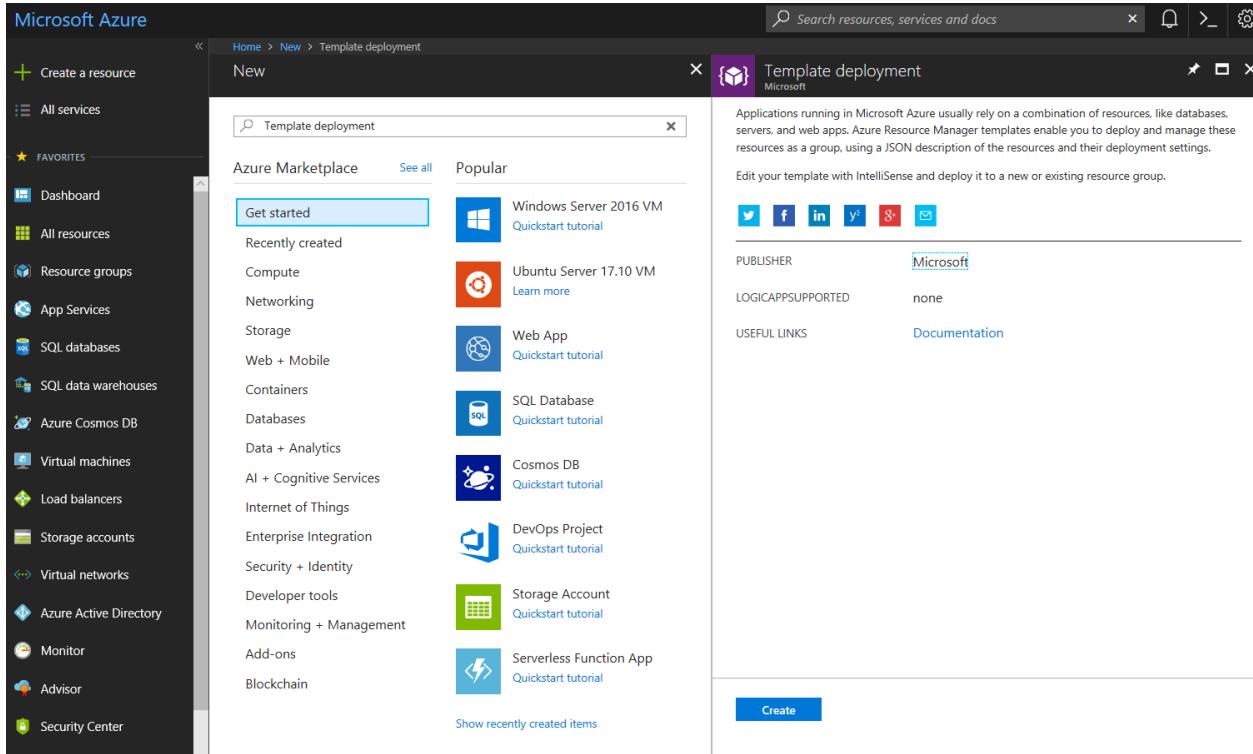
```

        "subnet": { "id": "[variables('subnetRef')]" }
    }
}
],
},
{
"apiVersion": "2018-04-01",
"type": "Microsoft.Compute/virtualMachines",
"name": "myVM",
"location": "[resourceGroup().location]",
"dependsOn": [
    "[resourceId('Microsoft.Network/networkInterfaces/', 'myNic')]"
],
"properties": {
    "hardwareProfile": { "vmSize": "Standard_DS1" },
    "osProfile": {
        "computerName": "myVM",
        "adminUsername": "[parameters('adminUsername')]",
        "adminPassword": "[parameters('adminPassword')]"
    },
    "storageProfile": {
        "imageReference": {
            "publisher": "OpenLogic",
            "offer": "CentOS",
            "sku": "7.4",
            "version": "latest"
        },
        "osDisk": {
            "name": "myManagedOSDisk",
            "caching": "ReadWrite",
            "createOption": "FromImage"
        }
    },
    "networkProfile": {
        "networkInterfaces": [
            {
                "id": "[resourceId('Microsoft.Network/networkInterfaces','myNic')]"
            }
        ]
    }
}
]

```



With the ARM Template, you can deploy it directly via Azure Portal. Start from **Create a resource** and then search **Template deployment** terms, and start to deploy ARM Template, as follows:



Ameerpet / Kondapur
Hyderabad

To deploy an ARM Template via Azure CLI, use the following commands:

```
az login

az group create --name TestInfra70533RG --location "West Europe"
az group deployment create \
--name TestInfra70533Deployment \
--resource-group TestInfra70533RG \
--template-file storage.json \
--parameters storageAccountType=Standard_GRS
```

Connecting to Azure VMs

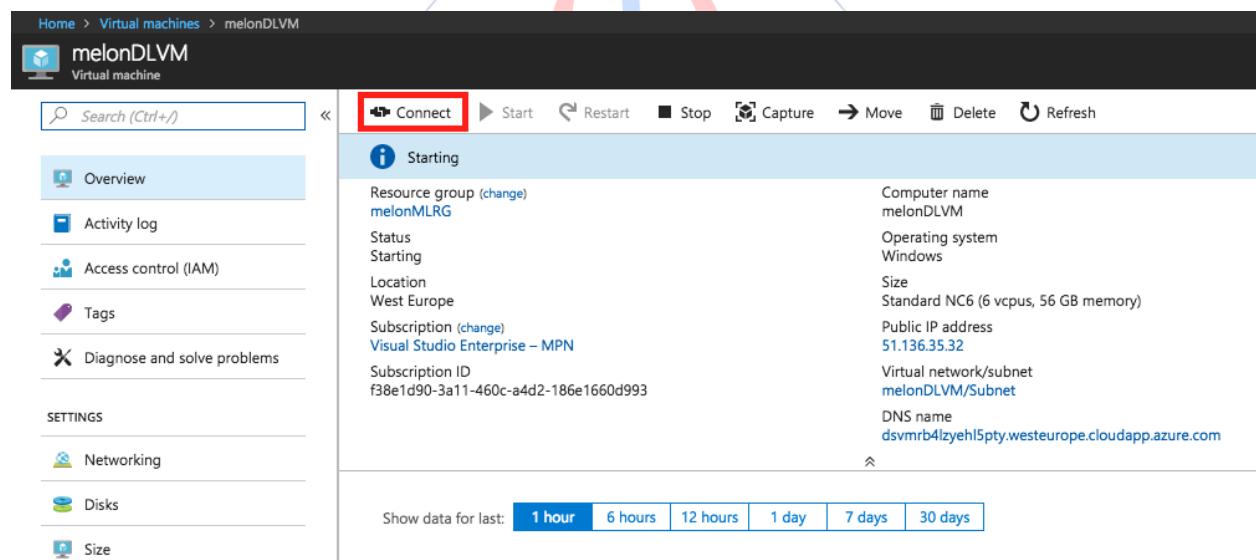
The default network security group that is attached to the virtual network of the deployed VM has the following NAT rules:

- Regarding Windows-based VMs: A rule which allows connectivity from the internet to port 3389
- Regarding Linux-based VMs: A rule which allows connectivity from the internet to port 22

Microsoft Azure provides three ways to connect to an Azure VM. The possible approaches are mentioned in the upcoming sections.

Connecting to a Windows Azure VM via Remote Desktop Protocol (RDP)

In the **Overview** blade, there is a **Connect** menu, as shown in the following screenshot, which allows you to download an RDP file after clicking on it; you can use it to connect to Windows Azure:



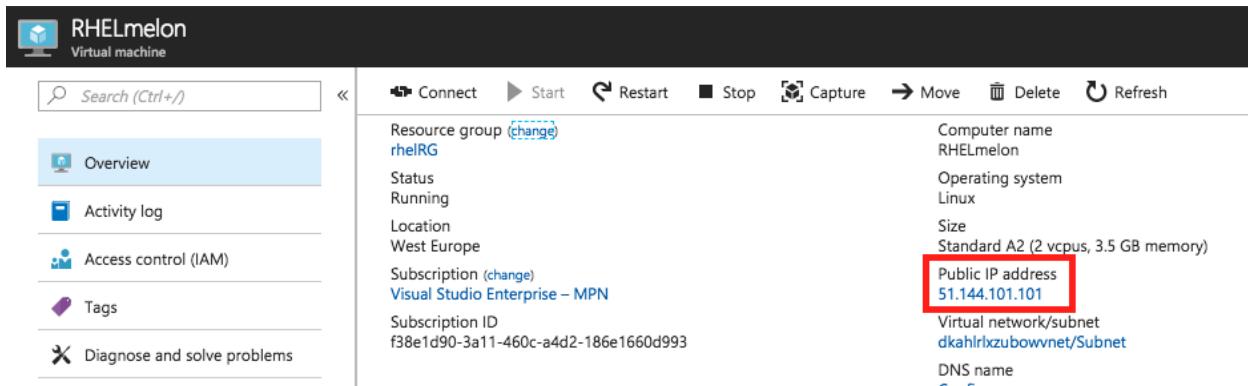
The screenshot shows the Azure portal's Overview blade for a virtual machine named "melonDLVM". The "Connect" button, which is highlighted with a red box, is located at the top right of the main content area. The main content area displays various details about the VM, including its resource group, status, location, subscription, and public IP address. The "Starting" status is indicated by a blue progress bar icon. The "Public IP address" listed is 51.136.35.32. The "DNS name" listed is dsvmrb4lzyehl5pty.westeurope.cloudapp.azure.com. At the bottom of the blade, there is a "Show data for last:" dropdown menu with options for 1 hour, 6 hours, 12 hours, 1 day, 7 days, and 30 days.

Connecting to a Linux Azure VM via Secure Shell (SSH)

Using the command line, you can connect to an Azure Linux VM:

```
ssh <yourAdminUsername>@<PublicIPOfYourVM>
```

You can find the **Public IP address** of your VM in the **Overview** blade:



RHELmelon
Virtual machine

Resource group: rheLRG
Status: Running
Location: West Europe
Subscription: Visual Studio Enterprise – MPN
Subscription ID: f38e1d90-3a11-460c-a4d2-186e1660d993

Computer name: RHELmelon
Operating system: Linux
Size: Standard A2 (2 vcpus, 3.5 GB memory)
Public IP address: **51.144.101.101**
Virtual network/subnet: dkahrlxzbubowvnet/Subnet
DNS name: [Configure](#)

The following is an example regarding how to connect to an Azure Linux VM using the command line:

```
melony — testcloudadmin@RHELmelon:~ — ssh testcloudadmin@51.144.101.101 — 88x...
[MelonyQins-MacBook-Pro:~ melony$ ssh testcloudadmin@51.144.101.101
The authenticity of host '51.144.101.101 (51.144.101.101)' can't be established.
ECDSA key fingerprint is SHA256:t+L03cEKtu5mJ0sq6SVrVTUEo+PjXBRWTzsBdoAUfhI.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '51.144.101.101' (ECDSA) to the list of known hosts.
>Password:
Last login: Thu Mar  8 20:36:23 2018
[testcloudadmin@RHELmelon ~]$ ]
```

Configuring Azure VMs in security

Microsoft Azure is engaged to guarantee data privacy and sovereignty, and enables their customers to control Azure-hosted data securely. The solutions provided by Microsoft take into account the potential business needs of their customers and give their customers the flexibility to choose the solution that fits best. To make sure that the Azure VM is secure, there are two aspects we always need to take care of. We will explain these in the following two subsections.

Restricting access to Azure VMs from the internet using NSG

An NSG, which is also known as a **network security group**, based on the Azure Virtual network, contains a list of security rules that allow or deny network traffic to resources connected to the same VNet. Besides, NSGs can be associated with subnets, individual VMs in the classic deployment model, or even individual network interfaces attached to VMs, such as

the **Resource Manager**. An example is that when an NSG is associated to a subnet, the rules apply to all resources connected to the subnet.

As we explained in the **Connecting to Azure VM** section of this chapter, to connect to an Azure VM successfully, there is a rule to allow inbound traffic in port 3389 for Windows-based VMs as follows:

Network Interface: **testinfrawinvnvm459**
Effective security rules
Topology ⓘ

Virtual network/subnet: infra-vm-sg-vnet/default
Public IP: **testinfrawinvnvm-ip**
Private IP: **10.1.0.4**

INBOUND PORT RULES ⓘ

PRIORITY	NAME	PORT	PROTOCOL	SOURCE	DESTINATION	ACTION	...
1000	⚠ default-allow-rdp	3389	TCP	Any	Any	Allow	...
65000	AllowVnetInBound	Any	Any	VirtualNetwork	VirtualNetwork	Allow	...
65001	AllowAzureLoadBalancerInBound	Any	Any	AzureLoadBalancer	Any	Allow	...
65500	DenyAllInBound	Any	Any	Any	Any	Deny	...

OUTBOUND PORT RULES ⓘ

PRIORITY	NAME	PORT	PROTOCOL	SOURCE	DESTINATION	ACTION	...
65000	AllowVnetOutBound	Any	Any	VirtualNetwork	VirtualNetwork	Allow	...
65001	AllowInternetOutBound	Any	Any	Any	Internet	Allow	...
65500	DenyAllOutBound	Any	Any	Any	Any	Deny	...

Inbound port rule in NSG of Windows-based Azure VM

For port 22 Linux-based VMs, there is also an inbound port rule for port 22, as follows:

Network Interface: [infravmazurecliVMNic](#) Effective security rules Topology 

 Virtual network/subnet: [infravmazurecliVNET/infravmazurecliSubnet](#) Public IP: **51.144.166.90** Private IP: **10.0.0.4**

INBOUND PORT RULES

 Network security group [infravmazurecliNSG](#) (attached to network interface: [infravmazurecliVMNic](#))
 Impacts 0 subnets, 1 network interfaces

[Add inbound port rule](#)

PRIORITY	NAME	PORT	PROTOCOL	SOURCE	DESTINATION	ACTION	...
1000	 default-allow-ssh	22	TCP	Any	Any	 Allow	
65000	AllowVnetInBound	Any	Any	VirtualNetwork	VirtualNetwork	 Allow	
65001	AllowAzureLoadBalancerInBound	Any	Any	AzureLoadBalancer	Any	 Allow	
65500	DenyAllInBound	Any	Any	Any	Any	 Deny	

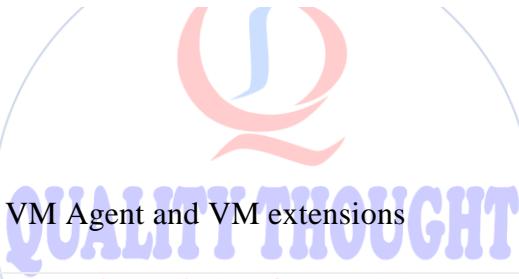
OUTBOUND PORT RULES

 Network security group [infravmazurecliNSG](#) (attached to network interface: [infravmazurecliVMNic](#))
 Impacts 0 subnets, 1 network interfaces

[Add outbound port rule](#)

PRIORITY	NAME	PORT	PROTOCOL	SOURCE	DESTINATION	ACTION	...
65000	AllowVnetOutBound	Any	Any	VirtualNetwork	VirtualNetwork	 Allow	
65001	AllowInternetOutBound	Any	Any	Any	Internet	 Allow	
65500	DenyAllOutBound	Any	Any	Any	Any	 Deny	

Managing Azure VMs with VM Agent and VM extensions



The Leader in Software Training

9963799240 / 7730997544

Ameerpet / Kondapur

Hyderabad

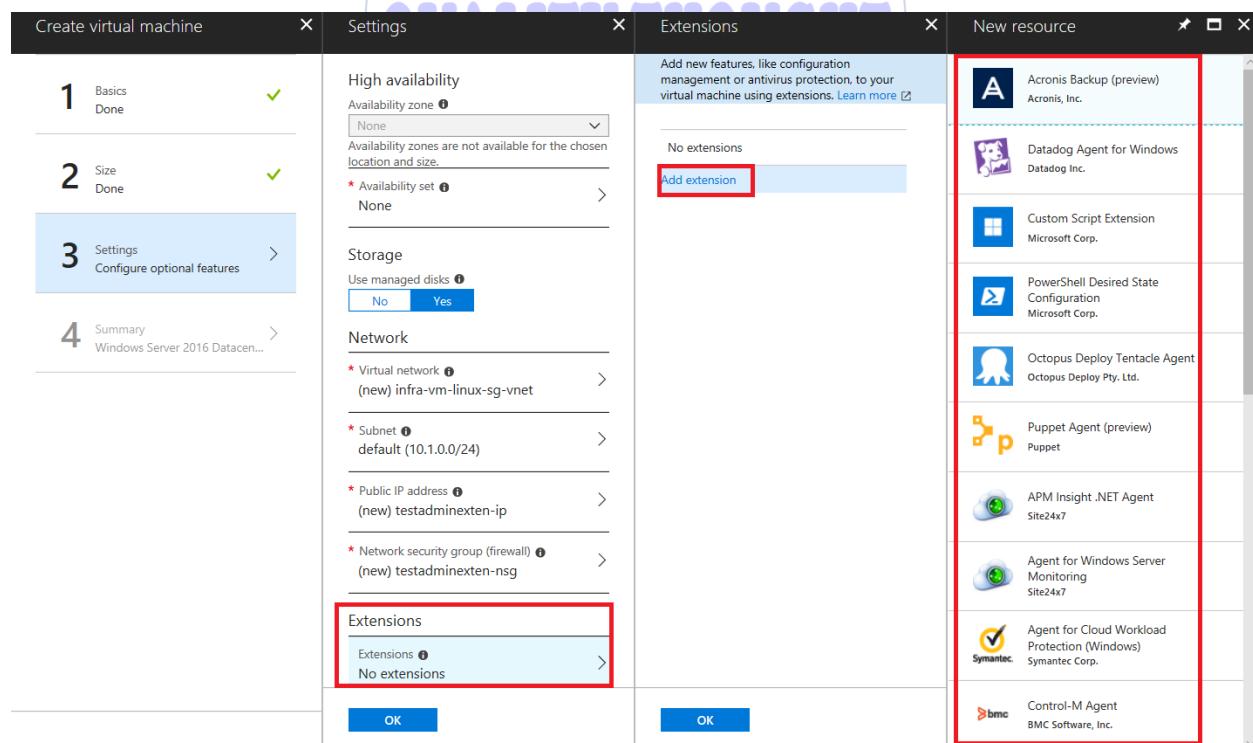
Azure offers a couple of methods that simplify and enhance the management of both Windows and Linux Azure VMs. Users can manage an Azure VM via Azure Portal, RESTful API, Azure PowerShell, Azure CLI, and so on. It is also possible to connect to Azure VMs when it is necessary to interact with an OS running within the VM.

When deploying Azure VMs, sometimes, we need to implement platform-specific configurations in VM. Azure provides a way to configure the OS and workloads running in the VM using a software component, which is known as the **Azure VM Agent**. The VM Agent supports **VM extensions**, which implement an additional functionality, especially in the areas of management, monitoring, and security terms.

Based on the VM Agent, users can add VM extensions. The following are some commonly used VM extensions:

- **Azure VM Access extension** enables you to reset local administrative credentials, fix misconfigured RDP settings on Windows VM, reset the admin password or SSH key, fix misconfigured SSH settings, create a new sudo user account, and check disk consistency on Linux VM.
- **Chef Client and Puppet Enterprise Agent** integrate Windows and Linux VMs into cross-platform Chef and Puppet enterprise management solutions.
- **Custom Script extension for Windows and Linux** makes it possible to run custom scripts within Azure VMs to apply custom configuration settings during VM provisioning. The extension supports any scripting language that the OS supports, such as Python or Bash.
- **DSC extension for Windows and Linux** implements a script-based or template-based configuration of OS components and applications.
- **Docker extension** facilitates automatic installation of Docker components, including the Docker Daemon, Docker Client, and Docker Compose on Linux VMs, and simplifies the process of implementing and managing containerized workloads in a significant way.

You can find out how to install the extension for your Azure VM by choosing **Extensions** while deploying a new VM; alternatively, after deploying, you can add the extension, as described in the following screenshot:



If your Azure VM has been already created, you can also go to the **Extension** blade and try to deploy a new extension by clicking **Add extension**.

Configuring the availability and scalability of Azure VMs

Let's take a look at managing Azure VMs regarding two aspects: availability and scalability. While implementing the Azure VM, it is important to make sure that workloads based on Azure are resilient and deal with all possible hardware failures.

Scaling Azure VMs

Generally, there are two kinds of scaling in the cloud:

- **Vertical scaling** is also called as **scale up**. Vertical scaling increases the capacity of existing hardware or software by adding compute resources, such as CPU memory-based processing power to a server, to make it faster. In this context, it means that users can scale by changing the VM's size.
- **Horizontal scaling** is also called **scale out**. Horizontal scaling is used to increase the number of multiple entities so that they can handle more incoming requests while scaling in the case of peak time, such as Black Friday. In this context, it means that users can scale by increasing or decreasing the number of VMs that reside in the same Availability Set and share their load through internal or external load balancing. To implement horizontal scaling of Azure VMs, Azure virtual machine scale sets (scale sets) would be a great choice.

In terms of improving the availability of Azure VMs, horizontal scaling is more desirable than vertical scaling because changing the VMs' size will cause the shutdown of Azure VMs.

A possible approach to improve Azure VM's scalability is presented in subsequent sections.

Configuring scale up by resizing Azure VMs

After deploying an Azure VM, it is possible to resize it when needed via Azure Portal, Azure CLI, or Azure PowerShell. At the Azure Portal, go to the Azure VM that you've deployed. There is a **Size** option in the blade and click on it. You'll note the potential Azure VMs size for your consideration.

After choosing the size you want, you can click on **Select** to start a resize deployment. You can achieve the same results via Azure CLI and Azure PowerShell.

You can check the link <https://docs.microsoft.com/en-us/azure/virtual-machines/windows/resize-vm> regarding the Azure documentation to learn how to resize a Windows Azure VM by using Azure PowerShell. In addition, you can learn to resize a Linux Azure VM by using Azure CLI. This is provided at the following link: <https://docs.microsoft.com/en-us/azure/virtual-machines/linux/change-vm-size>.

Configuring scale out by deploying ARM VM scale sets (VMSS) and configuring ARM VMSS auto-scale

Microsoft Azure provides facilities with the workloads, such as deploying a set of identical VMs while they have identical configurations and deliver the same functionality to support a service or application using **virtual machine scale sets (VMSS)**, or VM scale sets, for short. With VM scale sets, users can manage the scalability of VMSS by increasing or decreasing the number of VMs or resizing the VMSS which will resize the instances in VMSS, in an easy way. Another consideration is to facilitate the management of the availability of VMs in the VM pool which can deal with incoming requests in a flexible manner.

There are two basic ways to configure VMs deployed in a scale set:

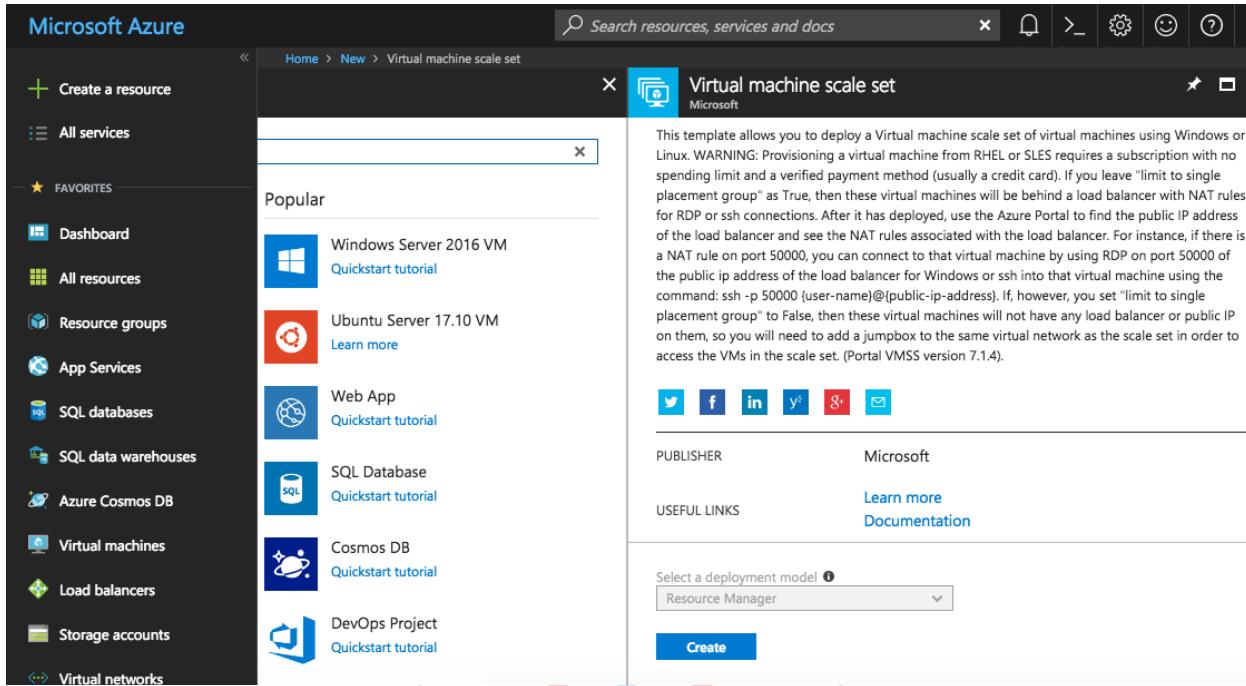
- Building your custom image before provisioning VMs with VMSS
- Configuring the VM using VM extensions after provisioning them

VMSS integrates the Azure load balancer or application gateway to handle dynamic distribution of network traffic across multiple VMs. It supports NAT rules so that users can connect to individual instances in VMSS.

Deploying VMSS via Azure Portal

To create a VMSS via Azure Portal perform the following steps:

1. You can search Virtual machine scale set, then click on **Create**. As shown in the following screenshot, VMSS is only available in the ARM model, which means you can only select **Resource Manager** as the deployment model:



This template allows you to deploy a Virtual machine scale set of virtual machines using Windows or Linux. **WARNING:** Provisioning a virtual machine from RHEL or SLES requires a subscription with no spending limit and a verified payment method (usually a credit card). If you leave "limit to single placement group" as True, then these virtual machines will be behind a load balancer with NAT rules for RDP or ssh connections. After it has deployed, use the Azure Portal to find the public IP address of the load balancer and see the NAT rules associated with the load balancer. For instance, if there is a NAT rule on port 50000, you can connect to that virtual machine by using RDP on port 50000 of the public ip address of the load balancer for Windows or ssh into that virtual machine using the command: ssh -p 50000 (user-name)@public-ip-address. If, however, you set "limit to single placement group" to False, then these virtual machines will not have any load balancer or public IP on them, so you will need to add a jumpbox to the same virtual network as the scale set in order to access the VMs in the scale set. (Portal VMSS version 7.1.4).

2. While filling in the information in the **Basic** blade, you can choose the operating system disk image you want to deploy in your dedicated virtual machine.
3. Microsoft Azure provides thousands of OS images in the Azure Marketplace. Similar to the user name and password, it will be applied to every deployed instance in VMSS:

Create virtual machine scale set

BASICS

* Virtual machine scale set name	infravmss
* Operating system disk image	Windows Server 2016 Datacenter
Browse all images	
* Subscription	Visual Studio Enterprise – MPN
* Resource group	<input checked="" type="radio"/> Create new <input type="radio"/> Use existing testinfra70533
* Location	West Europe
Availability zone	None
* User name	testcloudadmin
* Password	*****

4. Click on **Browse all images**—you'll note a list of OS images from the Azure Marketplace.
5. In the **INSTANCES** section, you can deploy your scale set as low priority (as shown in the following screenshot), which can help you save up to 80% over usual on-demand costs, especially while deploying stateless workloads, batch processing jobs, or large compute workloads, because the VMs in the scale set may be evicted when they're not in use:

INSTANCES

* Instance count

* Instance size ([View full pricing details](#)) ▼

Deploy as low priority

Use managed disks

[- Hide advanced settings](#)

Enable scaling beyond 100 instances

Note the instance count, which is the number of virtual machines in the scale set. It ranges from 1 to 100, which is much less than the maximum size of VMSS with the capacity of 1,000 VMs because, by default, the placement Group (not to scale more than 100 instances) was set to No. This means that the scale set will be limited to one placement group with a maximum capacity of 100.

6. Choosing Yes allows the scale set to span Placement Groups. This will enable the scaling beyond 100 and changes the availability guarantees of the scale set at the same time:

INSTANCES

* Instance count

* Instance size ([View full pricing details](#)) ▼

Deploy as low priority

Use managed disks

[- Hide advanced settings](#)

Enable scaling beyond 100 instances

7. In the networking section, there are two options to manage web traffic while deploying VMSS: **Application Gateway** and **Load balancer**.
8. Choose **Azure Load balancer** for VMSS as a load balancer, which allows you to scale web applications and improve high availability. In this case, you should fill in the public IP address name and FQDN for the load balancer in front of the scale set.

FQDN, **fully qualified domain name**, is a domain name that specifies its exact location in the tree hierarchy of the **Domain Name System (DNS)**, which must be unique across all of Azure. You should enter an appropriate name that adapts to the individual situation. The following is an example of this:

Choose Load balancing options	<input type="radio"/> Application Gateway <input checked="" type="radio"/> Load balancer
* Public IP address name <small>(1)</small>	<input type="text" value="testinfra70533"/> ✓
* Domain name label <small>(1)</small>	<input type="text" value="testinfra70533"/> ✓
.westeurope.cloudapp.azure.com	

Azure Application Gateway also acts as a web traffic load balancer. It is a dedicated virtual appliance providing the **application delivery controller (ADC)** as a service. It is an OSI layer 7 (application layer) load balancer, which can perform more specific functions than the traditional OS Layer 4 load balancer.

Once you've confirmed that all the information you have entered on the **Summary** blade is correct, you can click on **OK** to start the scale set deployment. The deployment will last about a couple of minutes. After the deployment, we can see that there are some related resources in the same resource group of the VMSS. We have created a public IP for the load balancer and a virtual network for our VMSS, as shown in the following screenshot:

Home > Resource groups > infra-vmss-sg

infra-vmss-sg

Resource group

Search (Ctrl+ /)

- Overview
- Activity log
- Access control (IAM)
- Tags

SETTINGS

- Quickstart
- Resource costs
- Deployments
- Policies

Add Edit columns Delete resource group Refresh Move Assign Tags

Subscription (change) Visual Studio Dev Essentials Subscription ID 5efcd8d-8da8-4993-9c36-4cc553... Deployments 1 Succeeded

Filter by name... All types All locations No grouping

5 items Show hidden types

NAME	TYPE	LOCATION
infravmss	Virtual machine scale set	West Europe
infravmss	Public IP address	West Europe
infravmsslb	Load balancer	West Europe
infravmssnsg	Network security group	West Europe
infravmssVnet	Virtual network	West Europe

Deploying VMSS using the ARM Template

You can also deploy the VMSS using the ARM Template. You can define your VMSS as follows:

```
{
  "type": "Microsoft.Compute/virtualMachineScaleSets",
  "name": "[variables('namingInfix')]",
  "location": "[resourceGroup().location]",
  "apiVersion": "2018-04-01",
  "dependsOn": [
    "[concat('Microsoft.Network/loadBalancers/', variables('loadBalancerName'))]",
    "[concat('Microsoft.Network/virtualNetworks/', variables('virtualNetworkName'))]"
  ],
  "sku": {
    "name": "[parameters('vmSku')]",
    "tier": "Standard",
    "capacity": "[parameters('instanceCount')]"
  },
  "properties": {
    "overprovision": "true",
    "upgradePolicy": {
      "mode": "Manual"
    },
    "virtualMachineProfile": {
      "storageProfile": {
        "osDisk": {
          "createOption": "FromImage",
          "caching": "ReadWrite"
        }
      }
    }
  }
}
```


{}

Deploying VMSS using Azure CLI

You can also deploy VMSS using the following Azure CLI command (replace the words between # with your own):

```
az vmss create -n #vmssname# -g #resourcegroupname# --instance-count  
#instancenumber# --image #vmimage# --data-disk-sizes-gb #gbsize#
```

To learn more about how to manage VMSS using Azure CLI, check out the following link: <https://docs.microsoft.com/en-us/azure/virtual-machine-scale-sets/virtual-machine-scale-sets-manage-cli>.

Deploying VMSS using Azure PowerShell

You can also deploy VMSS using the following Azure PowerShell command:

```
AzureRmVmss `  
-ResourceGroupName #resourcegroupname# ` 7730997544  
-Location #location# ` Ameerpet / Kondapur  
-VMScaleSetName #vmssname# ` Hyderabad  
-VirtualNetworkName #VnetName`  
-SubnetName #subnetname# `  
-PublicIpAddressName #publicIpAddressName# `  
-LoadBalancerName #lbname# `  
-UpgradePolicyMode #upgrademode#
```

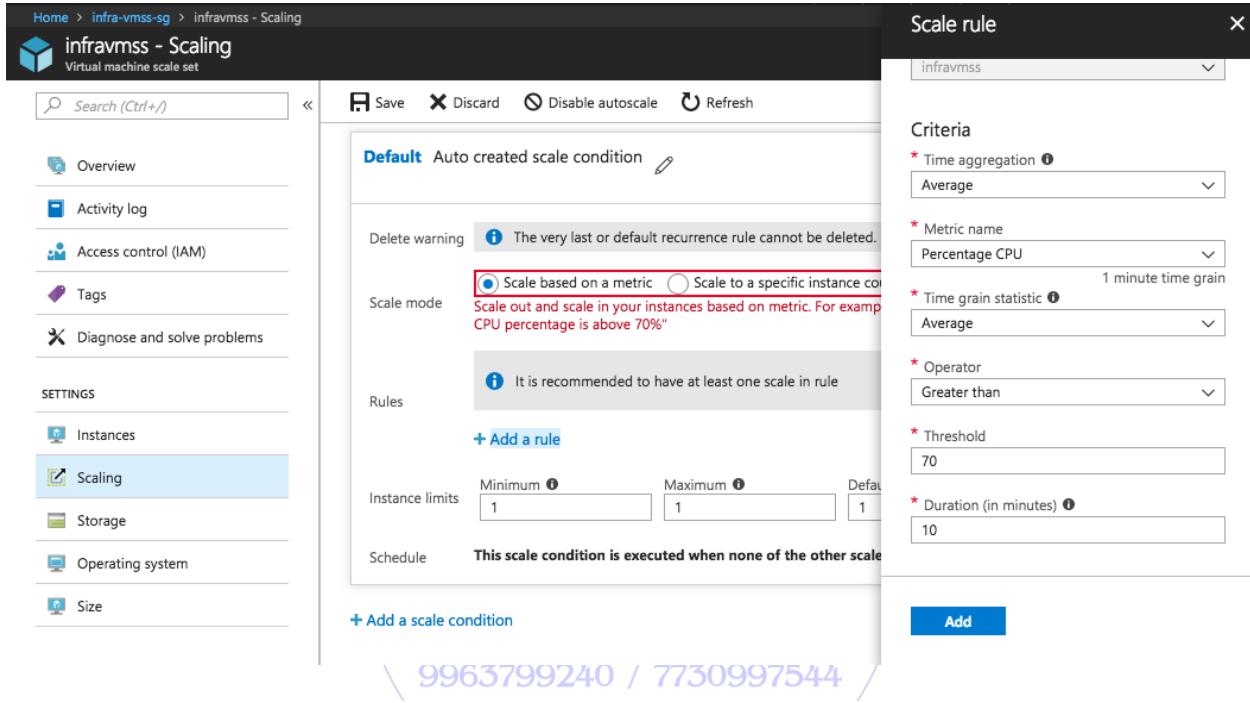
To learn more about how to manage VMSS using PowerShell, check out the following link: <https://docs.microsoft.com/en-us/azure/virtual-machine-scale-sets/virtual-machine-scale-sets-manage-powershell>

Configuring ARM VMSS autoscale

Even after the creation of VMSS, it is possible to adjust the autoscale condition for the VM in the deployed VMSS. You can also enable autoscale. There are two scale modes, as follows:

- Scale based on a metric
- Scale to a specific instance count

Configure the scaling rules, such as the **Minimum** or **Maximum** number of VMs and the CPU percentage threshold, which will take effect while specifying condition matching so that VMSS achieves scaling out or down. Make sure that every condition meets your intentions and finally click on **OK**:



It is also possible to use Azure Resource Explorer to preview your autoscaling condition ARM Template. You should add an autoscaling setting in the template, as follows:

```
{
  "type": "Microsoft.Insights/autoscaleSettings",
  "apiVersion": "2015-04-01",
  "name": "autoscalewad",
  "location": "[resourceGroup().location]",
  "dependsOn": [
    "[concat('Microsoft.Compute/virtualMachineScaleSets/', variables('namingInfix'))]"
  ],
  "properties": {
    "name": "autoscalewad",
    "targetResourceUri": "[concat('/subscriptions/', subscription().subscriptionId, '/resourceGroups/', resourceGroup().name, '/providers/Microsoft.Compute/virtualMachineScaleSets/', variables('namingInfix'))]",
    "enabled": true,
    "profiles": [
      {
        "scaleRules": [
          {
            "scaleCondition": {
              "timeAggregation": "Average",
              "metricName": "Percentage CPU",
              "timeGrain": "PT1M",
              "operator": "Greater Than",
              "threshold": 70,
              "duration": 10
            }
          }
        ]
      }
    ]
  }
}
```

```

"name":"Profile1",
"capacity":{
"minimum":"1",
"maximum":"10",
"default":"1"
},
"rules":[
{
"metricTrigger":{
"metricName":"Percentage CPU",
"metricNamespace":"",
"metricResourceUri":"[concat('/subscriptions/',subscription().subscriptionId,
'/resourceGroups/',
/resourceGroup().name,
'/providers/Microsoft.Compute/virtualMachineScaleSets/', variables('namingInfix'))]",
"timeGrain":"PT1M",
"statistic":"Average",
"timeWindow":"PT5M",
"timeAggregation":"Average",
"operator":"GreaterThan",
"threshold":60
},
"scaleAction":{
"direction":"Increase",
"type":"ChangeCount",
"value":1,
"cooldown":"PT1M"
}
},
{
"metricTrigger":{
"metricName":"Percentage CPU",
"metricNamespace":"",
"metricResourceUri":"[concat('/subscriptions/',subscription().subscriptionId,
'/resourceGroups/',
/resourceGroup().name,
'/providers/Microsoft.Compute/virtualMachineScaleSets/', variables('namingInfix'))]",
"timeGrain":"PT1M",
"statistic":"Average",
"timeWindow":"PT5M",
"timeAggregation":"Average",
"operator":"LessThan",
"threshold":30
},
"scaleAction":{
"direction":"Decrease",
"type":"ChangeCount",
"value":1,
}
}
]
}

```

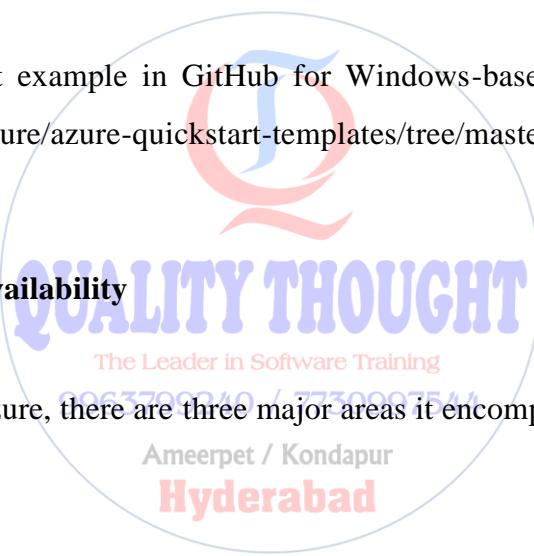


```
"cooldown":"PT5M"
}
}
]
}
]
}
}
```

You can find a great example in GitHub for Linux-based VMSS in the following link: <https://github.com/Azure/azure-quickstart-templates/tree/master/201-vmss-ubuntu-autoscale>

You can also find a great example in GitHub for Windows-based VMSS in the following link: <https://github.com/Azure/azure-quickstart-templates/tree/master/201-vmss-windows-autoscale>

Managing Azure VM's availability



In terms of resilience in Azure, there are three major areas it encompasses:

- High availability
- Disaster recovery
- Backup

Some readers may be wondering, we usually talk about availability as being 99.95%, 99.9% — what does it mean? Actually, these percentages of availability are the major cause of failure that Microsoft Azure can guarantee. After calculation, they may mean that the downtime of a Azure VM may be affected in order to meet its **Service Level Agreements (SLA)** as shown in the following table:

Availability	Downtime per year	Downtime per month	Downtime per week	Downtime per day
90%	36.52 days	3.04 days	16.80 hours	2.40 hours
95%	18.26 days	1.52 days	8.40 hours	1.20 hours
99%	3.65 days	7.30 hours	1.68 hours	14.40 minutes
99.5%	1.83 days	3.65 hours	50.40 minutes	7.20 minutes
99.9%	8.77 hours	43.83 minutes	10.08 minutes	1.44 minutes
99.95%	4.38 hours	21.91 minutes	5.04 minutes	43.20 seconds
99.99%	52.59 minutes	4.38 minutes	1.01 minutes	8.64 seconds
100.000%	5.26 minutes	26.30 seconds	6.05 seconds	0.86 seconds

You can see that when availability is at 99.95%, we'll have over 4 hours of downtime per year; when availability is 99.99%, the downtime reduces within 1 hour per year, which may be acceptable by most applications. When availability is 99%, which seems like a good number but after calculation, it means over 3 days per year when the application won't run. This situation cannot be accepted by mission-critical applications. When there is a problem, there is a solution, and Microsoft is striking to improve the resilience of applications by using major levels to manage the SLA (service-level agreement) of VMs:

- Single VM (based on premium storage): SLA 99.9%
- Availability Set: SLA 99.95%
- Availability Zone: SLA 99.99%
meerpet / Kondapur
Hyderabad
- Region Pairs

Here is Microsoft's latest infographic regarding different levels of availability provided by Azure services: https://azurecomcdn.azureedge.net/mediahandler/files/resourcefiles/azure-resiliency-infographic/Azure_resiliency_infographic.pdf

There are several general approaches to achieving high availability across regions pairs:

- **Active/passive with hot standby** means that when the traffic goes to the primary region, the VMs in the secondary region are allocated and running at all times.
- **Active/passive with cold standby** means that when the traffic goes to the primary region, but the VMs in the secondary region is not allocated until needed for failover and it will take time to be allocated in case of fail-over.

- **Active/active** means the primary and secondary region are both active, and requests can be distributed by load balancing between them. The healthy status will be determined by healthy

There are two types of events, planned maintenance and unplanned maintenance, in Azure that will affect the availability of Azure virtual machines.

Planned maintenance is when VMs are restarted due to Microsoft updates on the underlying platform. **Unplanned maintenance** is when there is a hardware failure.

To learn more about the SLA of VMs, check out the following link: https://azure.microsoft.com/en-us/support/legal/sla/virtual-machines/v1_8/

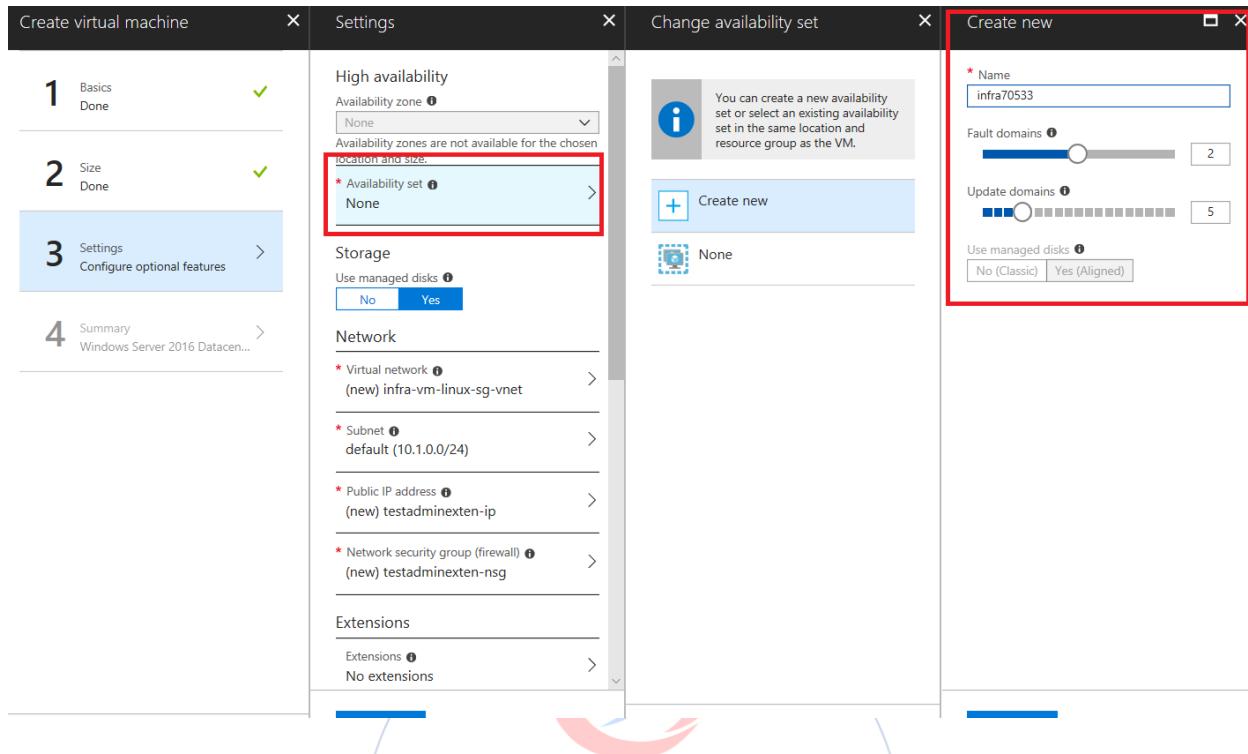
Now, in the subsequent sections, we'll be putting theory into practice.

Improving Azure VM's availability using Availability Sets or Availability Zones

To provide redundancy to your application, Microsoft recommends that you group two or more virtual machines in an Availability Set, which is a logical grouping of two or more virtual machines. This configuration ensures that, during a planned or unplanned maintenance event, at least 2 virtual machines in the Availability Set will meet the 99.95% Azure SLA.

When creating Availability Sets, Microsoft recommends the following best practices:

- For redundancy, configure multiple virtual machines in an Availability Set
- Configure each application tier into separate Availability Sets
- Combine a load balancer with Availability Sets



Another fact is that if you have two or more instances deployed across two or more Availability Zones (AZ is only available in some regions and for some Azure services mentioned in the following link, <https://docs.microsoft.com/en-us/azure/availability-zones/az-overview> at the moment) in the same Azure region, the SLA will be at least 99.99 %. You can choose the Availability Zone while creating a new VM from the drop-down menu.

Looking into the VM's availability by converting a Windows virtual machine from unmanaged disks to managed disks

As we explained previously, since managed disks compared to unmanaged disks break the limits of IOPS per storage account, Microsoft recommends that you convert the VMs to use managed disks through the Azure Managed Disks service. This means that the best practice is to convert both the OS disk and any attached data disks of an Azure VM. This approach provides better availability than unmanaged disks. In Azure, for any single instance of Azure VM using premium storage (SSD) for all Operating System Disks and Data Disks, the will meet SLA at least 99.9%.

Improving Azure VM's availability by combining a load balancer with Availability Sets

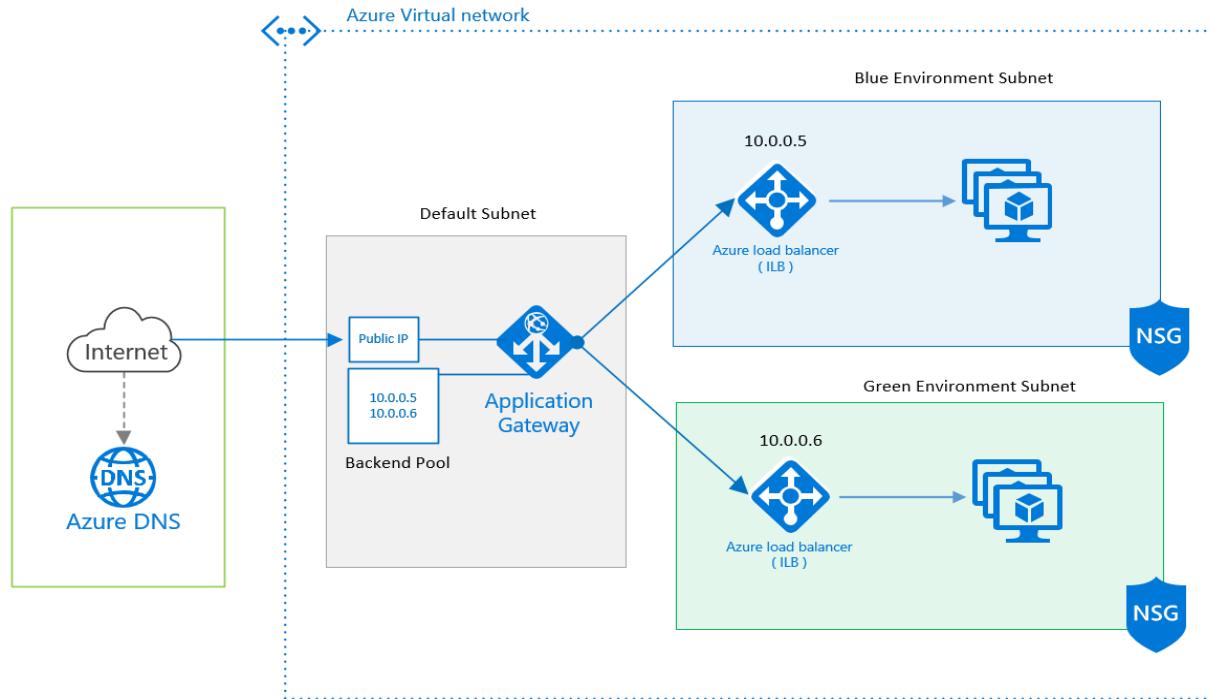
While deploying multiple Azure VMs to improve availability purposes, we usually combine a load balancer that allows distributing traffic between multiple virtual machines. The load

balancer is usually integrated with a health check process, which means it attempts connections or sends requests to test the VMs periodically. It routes requests to the VMs that are available. Microsoft recommends combining the Azure load balancer with an Availability Set to get the most application resiliency.

Improving Azure VM's availability by implementing blue-green deployment in Azure

The main principle of blue/green deployments, which are also known as A/B deployments, is to deploy two identical environments, which are configured in the same way. Generally, while one environment is live and in use by users, the other environment stays idle. When downtime occurs, this architecture allows you to redirect the incoming traffics to the idle configuration, which runs the original version with the help of a load balancer. The aim is to reduce downtime during production deployments.

The general B/G architecture in Azure contains one resource group for the green environment, which usually contains an application in the old version deployed in a VMSS, and one resource group for the blue environment, with an application in the newer version deployed in a VMSS. All the resources are in the same virtual network (VNet); the green and blue environment is on a different subnet. The Application Gateway receives all incoming traffic and distributes it to the backend load balancers. Application Gateway contains two addresses in its backend pool, which are the frontend of two load balancers. Each VMSS has an internal load balancer with a private frontend IP address so that it can distribute the incoming traffic across the backend VMs. The following is an example of designing a B/G deployment in Azure with VMSS; an Azure Load Balancer, which is a Level 4 load balancer; and an Application Gateway, which is a Level 7 load balancer.



Improving Azure VM's availability by implementing multi-region deployments with ARM Templates

Microsoft also recommends managing availability by deploying your infrastructure to a multi-region. In the case of regional outage that affects the primary region, a DNS-level load balancer, which is a Traffic Manager in Azure, provides capabilities to fail over to the secondary region. The multi-region architecture helps if an individual subsystem of the application fails. To learn more regarding the multi-region architecture, refer to the examples in the Azure reference architecture center.

Run Linux VMs in multiple regions for high availability using the following link:

<https://docs.microsoft.com/en-us/azure/architecture/reference-architectures/virtual-machines-linux/multi-region-application>

Run Windows VMs in multiple regions for high availability using the following link:

<https://docs.microsoft.com/en-us/azure/architecture/reference-architectures/virtual-machines-windows/multi-region-application>

Implementing and Managing Containers in Azure

As cloud computing technology has become the new norm, it has helped in the quick growth of virtualization technology, which has changed the IT industry in a significant way. Containers and container clusters in the cloud are becoming more widely spoken of today.

In this chapter, we'll cover the following topics:

- Implementing the Azure Container Registry in Azure
- Creating and managing container images with Docker
- Deploying and managing clusters of containers using **Azure Container Service (ACS)**, along with open source and non-Microsoft container orchestration solutions, such as Docker Swarm, Kubernetes, and DC/OS
- Deploying a Kubernetes cluster with **Azure Kubernetes Service (AKS)** in Azure
- Migrating container workloads to, and from, Azure
- Monitoring Kubernetes using Microsoft **Operations Management Suite (OMS)**



The principle of containers and microservices Software Training

9963799240 / 7730997544

Ameerpet / Kondapur

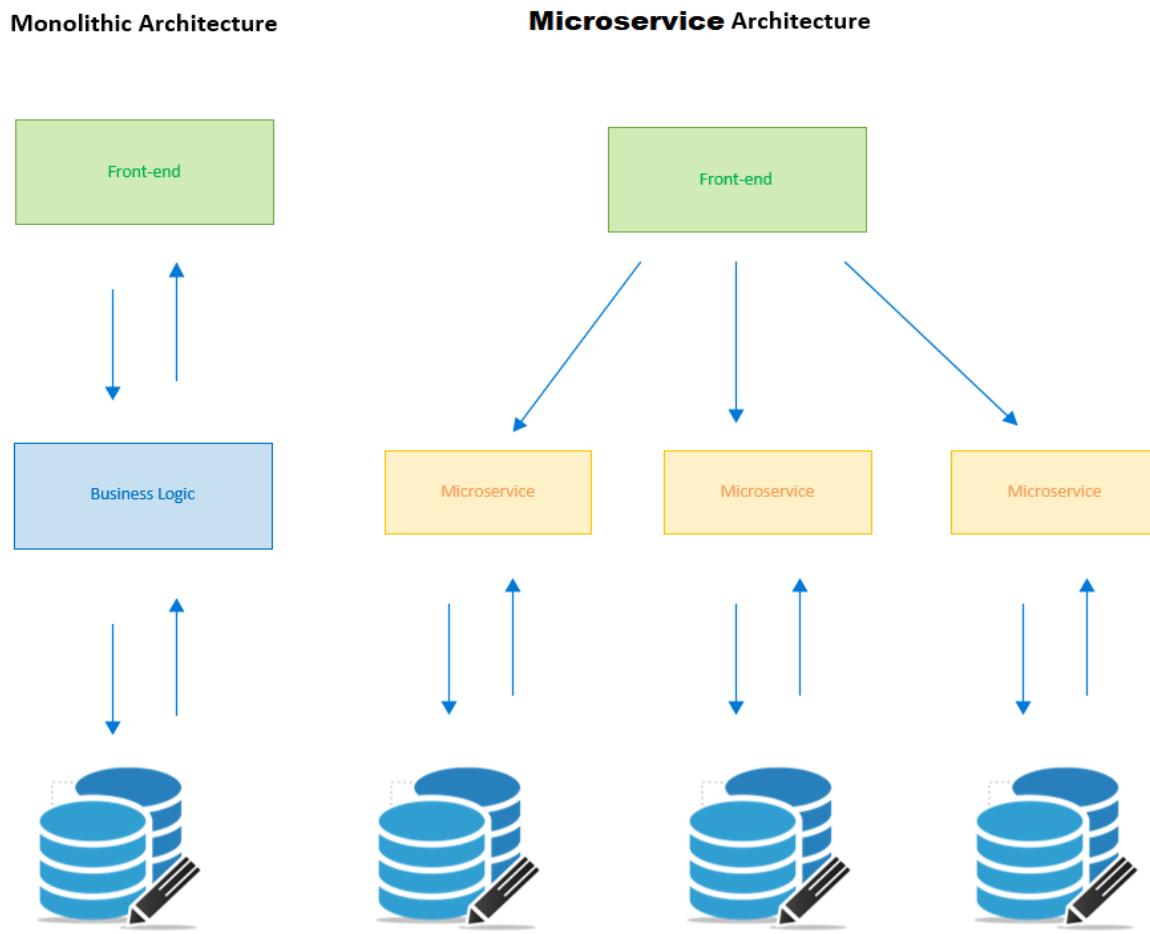
Hyderabad

Before discussing containers, we cannot ignore the term **microservice**. Microsoft defined microservices architecture as a whole system that contains a collection of small, autonomous services.

The following are the three main reasons for building the modern applications with the microservice architecture:

- With independent, autonomous modules, it is possible make each module scale at their own space
- Using different technologies in the same application has become possible, for example, a RESTful API service can be developed both in Node.js and in the .NET web API in the same application
- High backwards compatibility makes client-side applications evolve at their own pace

The difference between a traditional monolithic architecture, or N-tier applications (such as the front web tier, middleware business logic, and backend data tier), and a microservice architecture can be explained with the following schema:



As we can see in the preceding schema, the microservice has classified an application into individual modules, and each of them has its own data tier that makes sure that a microservice can work as an independent module. Each module can be scaled at their own pace. In case of failure of one microservice, other microservices won't be affected—the system will be still working, all we need to do is to restore the failed modules.

With the help of containers that allow us to run the applications in an isolated virtualized environment, there will be no more challenges such as administering the patch of the operating system, and managing the dependencies of applications. Containers aim to enforce the

portability and agility of applications; that is why it is a powerful technology to build applications in a microservice architecture.

Containers versus container clusters

Containers is definitely a popular keyword in the IT industry, and is a higher level of virtualization technology compared to virtual machines. Containers act as a **recipe** in the kitchen, and they manage to virtualize the operating system and the related infrastructure. We can add applications with anything related to it, such as all dependencies, libraries, and other binaries and configuration files needed to run the application inside, bundled into one package. Definitely, containers enforce the portability of an application, it is a great to resolve the problem of how to run the application while moving from one computing environment to another.

Containers act as a deployment unit when we need to deploy multiple container clusters. It was initially designed for the development and staging environment and will be ready for production environments soon. One of the greatest examples of containerization is Docker.

We'll explain the basics of Docker and provide a demonstrative example while working with Azure in the following section.

Containers provide an effective way to orchestrate your software, operating system, and hardware configurations to provide a typical running environment to make your application run the way you want it to. However, while running an application with numerous instances, such as hundreds, thousands, or even more container clusters, managing all these clusters becomes a really challenging question. Notable examples of container-clustering solutions are Kubernetes, Docker Swarm, and DO/OS.

We'll explain the basics of container-clustering solutions and practice demos while working with Azure in the following section.

Docker basics

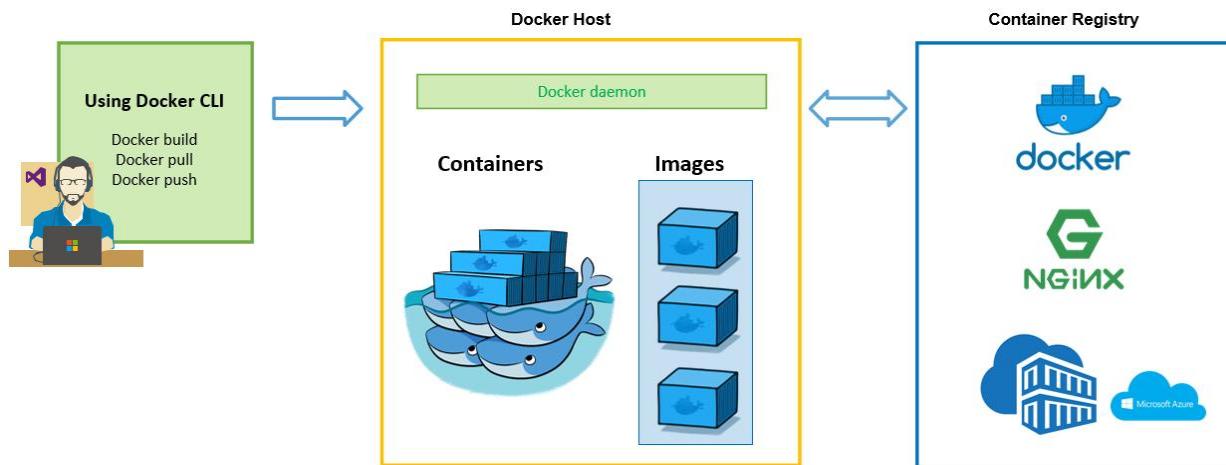
Docker is the world's leading software container platform available for developers, DevOps, and businesses to help them build, deliver, and run any application on independent infrastructures. Docker uses a client-server architecture. Docker can be built into three important parts: **Docker Client**, **Docker Host** (with Docker Daemon), and **Docker Registry**. Each part has its own responsibilities:

- **Docker Client** is where a Docker environment should be installed to build Docker images with the target application
- **Docker Host** is a managed host with Docker Daemon (also known as dockerd, which is the persistent process that manages containers)
- **Docker Registry** provides or stores different Docker images

The Docker Client can communicate with Docker Daemon using the RESTful API over UNIX sockets or a network interface in the following ways:

- Docker Client can run on the same system with Docker Daemon
- Docker Client connects to a remote Docker Daemon

This can be done as shown in the following picture (from **Docker Documentation**):



Container registry

The container registry provides different Docker images in the marketplace. We will provide examples of some famous open communities for Docker images in this section.

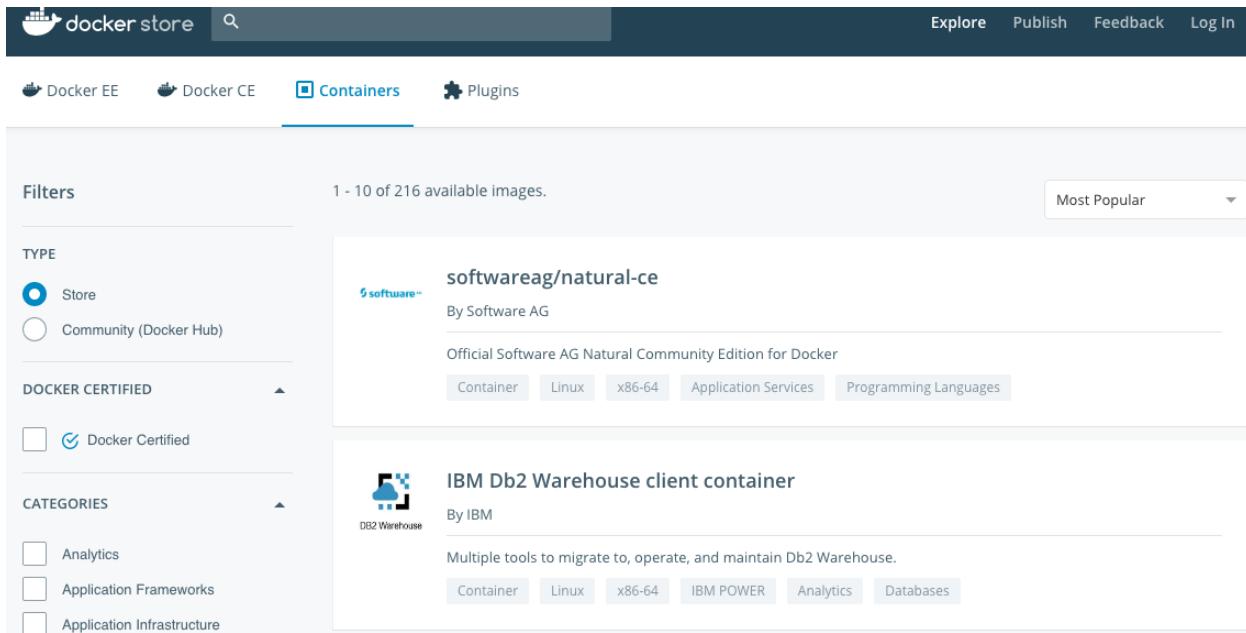
Docker Hub

Docker Hub provides public and private registries for Docker images that are built by other communities. Using Docker Hub, users can also upload their own Docker-built images to Docker Hub. It provides a webhook to support dev-test pipeline automation. Navigate to the official site of Docker Hub at <https://hub.docker.com/> to find what you need:



Docker Store

The Docker Store is a common Docker registry. It provides especially trusted and enterprise-ready containers, plugins, and Docker editions. The official site of Docker Store is <https://store.docker.com/>, as shown in the following screenshot:



The screenshot shows the Docker store interface. On the left, there are filters for Type (Store selected), Docker Certified (unchecked), and Categories (Analytics, Application Frameworks, Application Infrastructure). The main area displays 1 - 10 of 216 available images. Two images are listed:

- softwareag/natural-ce** by Software AG: Official Software AG Natural Community Edition for Docker. Tags: Container, Linux, x86-64, Application Services, Programming Languages.
- IBM Db2 Warehouse client container** by IBM: Multiple tools to migrate to, operate, and maintain Db2 Warehouse. Tags: Container, Linux, x86-64, IBM POWER, Analytics, Databases.

Nginx

Another repository that contains official Docker images for Nginx is on GitHub. You can find some official Nginx Dockerfiles at the following repository: <https://github.com/nginxinc/docker-nginx>.

Dockerizing your web application in Azure

9963799240 / 7730997544

Ameerpet / Kondapur

Hyderabad

Microsoft Azure provides a couple of capabilities to help developers and organizations to Dockerize their application.

Preparation work

Before working with Azure, users should make sure that they have installed the Docker environment and Docker CLI tools correctly:

1. Users can download Docker and install it using the following useful links that have guidance for installation:
 - Install Docker for Windows: <https://docs.docker.com/docker-for-windows/install/>
 - Install Docker for Mac: <https://docs.docker.com/docker-for-mac/install/>

If you have an older version of Windows, you may need to install Docker Toolbox before installing Docker from <https://docs.docker.com/toolbox/overview/>.

After installing Docker, a whale will be displayed in the notification area, which shows that Docker has started, and users can access Docker from a terminal or console from your OS. The following is a simple command to verify that Docker is ready in your host.

Run the following command to list all the available Docker commands:

```
docker --version
```

Here is a screenshot of the output (in macOS):

```
MelonyQins-MacBook-Pro:~ melony$ docker -v
Docker version 17.12.0-ce, build c97c6d6
MelonyQins-MacBook-Pro:~ melony$ docker --version
Docker version 17.12.0-ce, build c97c6d6
MelonyQins-MacBook-Pro:~ melony$ docker

Usage: docker COMMAND

A self-sufficient runtime for containers

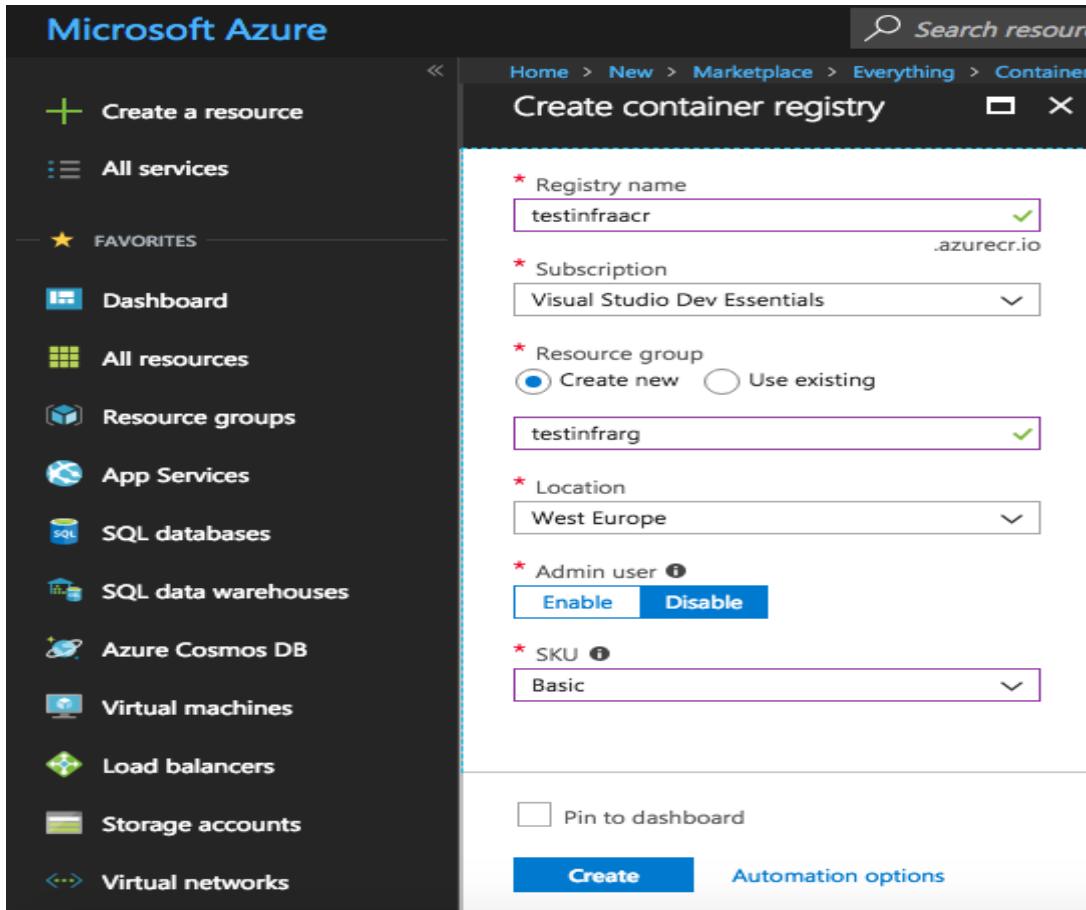
Options:
  --config string      Location of client config files (default "/Users/melony/.docker")
  -D, --debug          Enable debug mode
  -H, --host list      Daemon socket(s) to connect to
  -l, --log-level string
                       Set the logging level ("debug"|"info"|"warn"|"error"|"fatal") (default "info")
  --tls               Use TLS; implied by --tlsverify
  --tlscacert string  Trust certs signed only by this CA (default "/Users/melony/.docker/ca.pem")
```

Implementing Azure Container Registry

Azure Container Registry is a private registry performed by Microsoft Azure to help developers host Docker-formatted images.

Azure Container Registry integrates well with orchestrators hosted in Azure Container Service, such as Docker Swarm, DC/OS, and Kubernetes. Users can benefit from using open source CLI tools, and this made it possible to maintain Windows and Linux container images in a single Docker Registry.

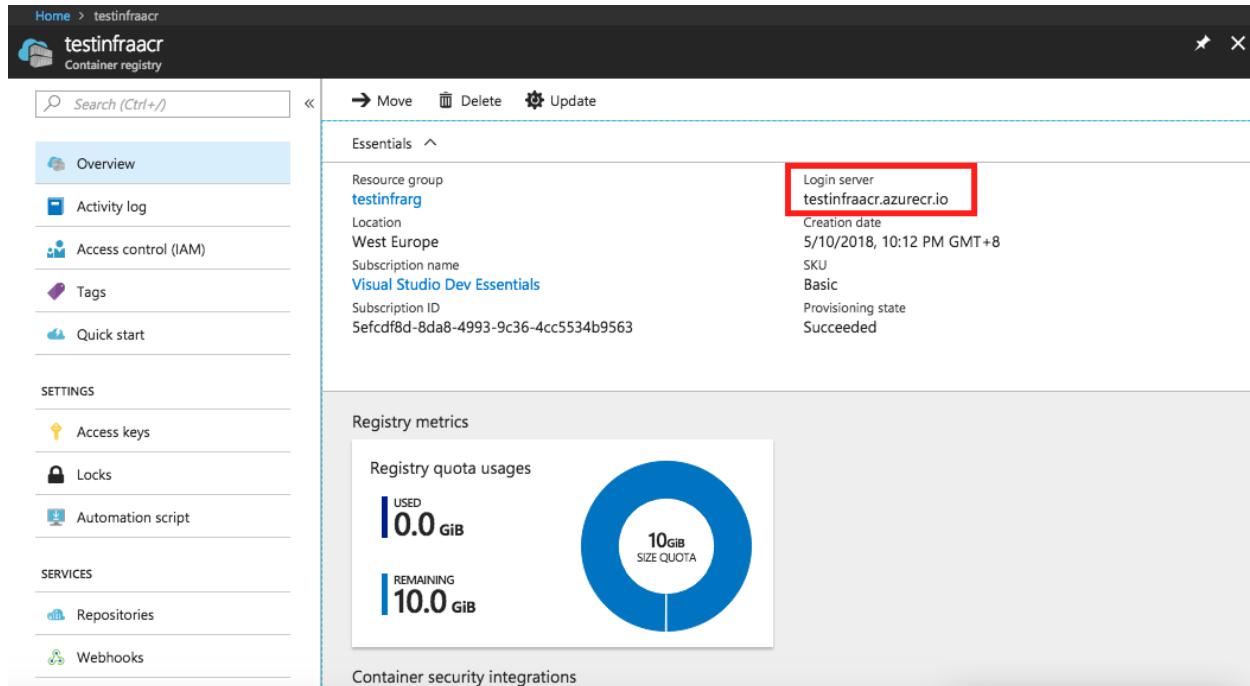
To create an Azure Container Registry, go to Azure portal, click on **Create a resource**, then search for Container Registry and choose it. Then, fill in the necessary information in the blade:



Creating the Container Registry

There are currently three types of capabilities that is, Basic, Standard, and Premium. All SKUs provide the same programmatic capabilities, but a higher SKU will provide more available storage, total webhooks, and a geo-replication feature (which is only available for Premium SKU), for example.

After creating ACR successfully via Azure Portal, you can return to the resource group that you've created and search for the created container registry, as follows:



Creating the Azure Container Registry

The following is the URL of the login server that you'll use to log in:

#nameofregistry#.azurecr.io

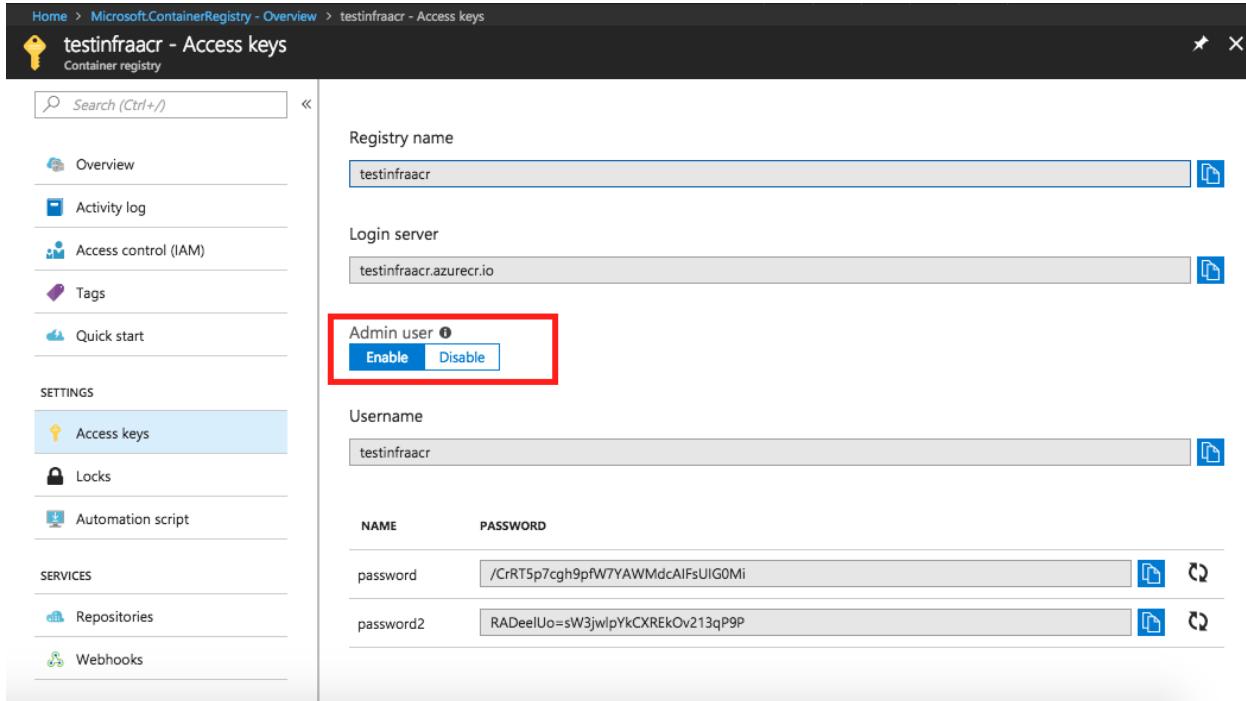
9963799240 / 7730997544

Ameerpet / Kondapur

Hyderabad

Pushing your Docker image in ACR

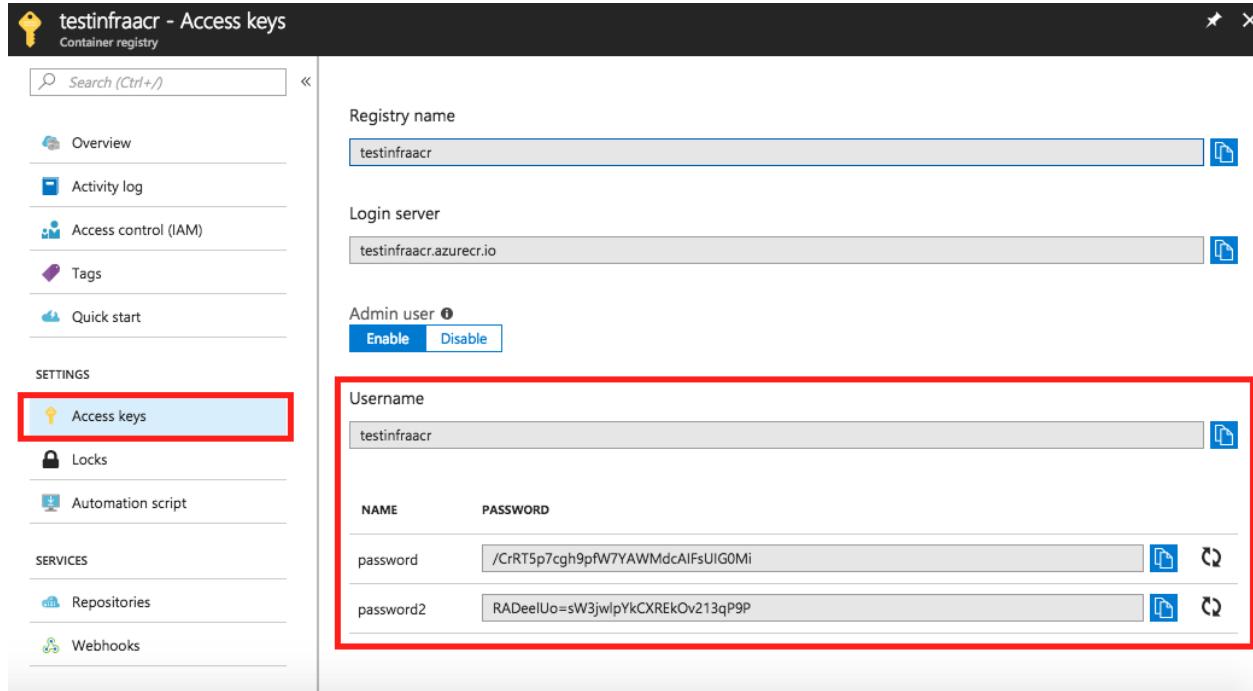
If you haven't enabled the **Admin user** while creating an ACR, you can navigate to the **Access keys** blade and click on **Enable** (as shown in the following screenshot):



The screenshot shows the Microsoft Container Registry - Overview page for a registry named "testinfraacr". The left sidebar has links for Overview, Activity log, Access control (IAM), Tags, Quick start, Settings (with Access keys selected), Locks, Automation script, Services (Repositories, Webhooks), and Support. The main area shows the Registry name as "testinfraacr" and the Login server as "testinfraacr.azurecr.io". Under the "Admin user" section, there is a "Enable" button which is highlighted with a red box. Below it are fields for Username ("testinfraacr") and two password fields ("password" and "password2").

Enabling the admin user in the Access Key blade

Then, you can use your username and password on the **Admin user** section to log in to your ACR registry and manage your credentials here (as shown in the following screenshot). If you lose your password, you can regenerate it by clicking on the refresh icon, and your current password will immediately become invalid and not recoverable:



Before pushing your local Docker image to ACR, use the following command to log in to Docker (replace the words between # with your own):

```
docker login myContainerRegistry.azurecr.io -u #username# -p #password#
```

The Leader in Software Training

9963486280 / 9963799240

Ameerpet / Kondapur

If you log in to Azure container registry, you'll get the following output:

```
melonyqins-macbook-pro:Github melony$ docker login testinfraacr.azurecr.io -u testinfraacr -p /CrRT5p7cgh9pfW7YAWdcaIFsUIG0Mi
WARNING! Using --password via the CLI is insecure. Use --password-stdin.
Warning: failed to get default registry endpoint from daemon (Cannot connect to the Docker daemon at unix:///var/un/docker.sock. Is the docker daemon running?). Using system default: https://index.docker.io/v1/
Cannot connect to the Docker daemon at unix:///var/run/docker.sock. Is the docker daemon running?
melonyqins-macbook-pro:Github melony$ docker login testinfraacr.azurecr.io -u testinfraacr -p /CrRT5p7cgh9pfW7YAWdcaIFsUIG0Mi
WARNING! Using --password via the CLI is insecure. Use --password-stdin.
Login Succeeded
```

Tag the locally built image to the ACR repository:

```
docker tag #yourImageId
yourContainerRegistry.azurecr.io/starterapp:latest
```

Then, push it to ACR using the following commands:

```
docker push yourContainerRegistry.azurecr.io/starterapp:latest
```

If you pushed your Docker images to ACR successfully, you'll get a message stating that it is Pushed in the output, as shown in the following example:

```
[melonyqins-macbook-pro:Github melony$ docker push testinfraacr.azurecr.io/starterapp:latest
The push refers to repository [testinfraacr.azurecr.io/starterapp]
f999ae22f308: Pushed
latest: digest: sha256:8072a54ebb3bc136150e2f2860f00a7bf45f13eeb917cca2430fc0054c8e51b size: 524
```

If you still want to verify that your images are in the registry, you can also use the following commands to verify it in the cloud shell or your local machine if you've already installed Azure CLI:

```
az acr repository list -n yourContainerRegistry
```

Deploying your Dockerized application with CI/CD capabilities

Currently, Microsoft Azure provides two ways to deploy the Dockerized application:

- Web App for Containers
- Azure Containers instances

Clustering solutions with Azure ACS in Azure

Nowadays, there are a couple of container orchestrators that help us to simplify the management of container clusters, in order to improve an application's scalability and resilience. The common objective of these tools is to let users handle the entire cluster as a

single deployment, which also extends the life cycle management capabilities to complex workloads with multiple containers deployed on a cluster of machines.

The open source container orchestrators are popular in the market, such as Docker Swarm, Kubernetes, and Mesosphere's DC/OS.

As a fast-growing cloud provider, Microsoft Azure implemented a container service known as ACS for all of these popular container orchestrators in Azure. This makes the containerization application fully portable by leveraging Docker images. These applications can be scaled to thousands or even more containers.

To implement any container cluster, Microsoft Azure provides a cloud-based service, which is known as Azure Container Service.

An overview of container cluster solutions working with ACSs

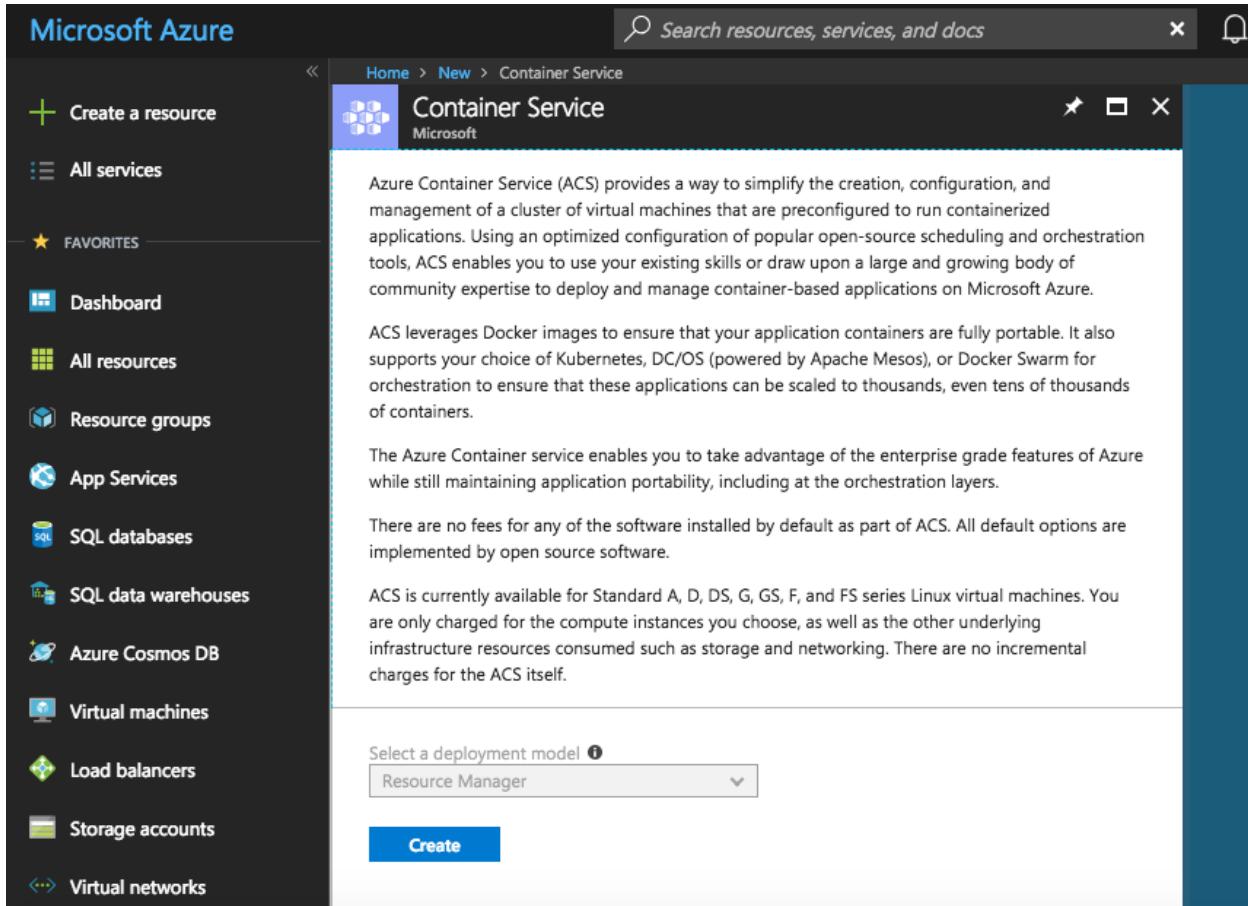
The goal of ACS is to simplify the creation, configuration, and management of a container cluster in the Azure cloud using an optimized configuration of popular open source scheduling and orchestration tools. ACS implements three kinds of popular open source orchestrators, such as Kubernetes, DC/OS (datacenter and operating system, which is powered by Apache Mesos), and Docker Swarm. You can use them for orchestration in Azure.

When you're using ACS, Microsoft Azure only charges for the compute instances and the underlying infrastructure resources consumed, such as storage or networking. There are no fees for any of the software installed by default as part of ACS.

We can directly deploy an ACS cluster via the portal, use the Azure CLI, or deploy an ARM (Azure Resource Manager) template.

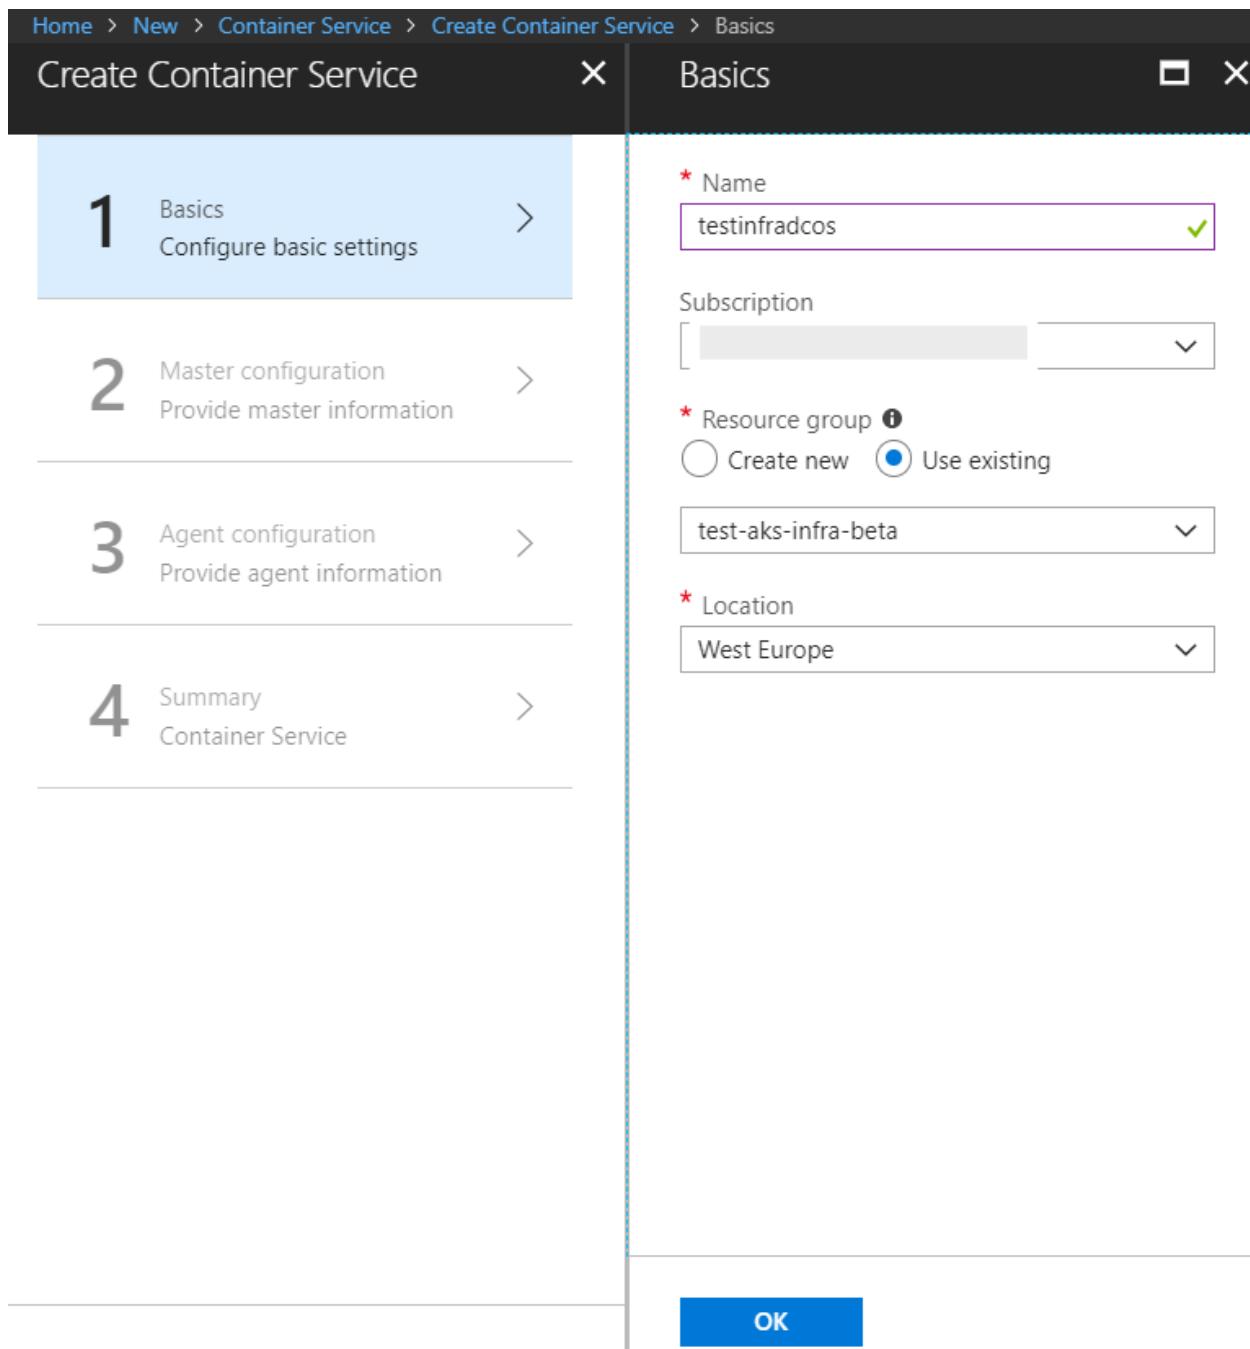
Creating ACS via the Azure Portal

In the Azure Portal, click on **Create a resource**. You can search for container services and then choose it to start to create a container service in Azure, as shown in the following screenshot:



The screenshot shows the Microsoft Azure portal interface. On the left, there is a sidebar with various service icons and links: Create a resource, All services, Favorites, Dashboard, All resources, Resource groups, App Services, SQL databases, SQL data warehouses, Azure Cosmos DB, Virtual machines, Load balancers, Storage accounts, and Virtual networks. The main content area has a title bar "Container Service" with the Microsoft logo. Below the title bar, there is a descriptive text about Azure Container Service (ACS) and its features. At the bottom of the main content area, there is a "Select a deployment model" dropdown set to "Resource Manager" and a "Create" button.

In the **Basic** blade, you can specify the name of the cluster and subscription that you want to use for this resource and define the right resource group and resource location. Then, go to **Master configuration**, which is important to identify the type of your orchestrator:



The screenshot shows the 'Create Container Service' wizard with four steps:

- Step 1: Basics** (selected): Configure basic settings.
- Step 2: Master configuration**: Provide master information.
- Step 3: Agent configuration**: Provide agent information.
- Step 4: Summary**: Container Service.

The 'Basics' step is currently active, showing the configuration options:

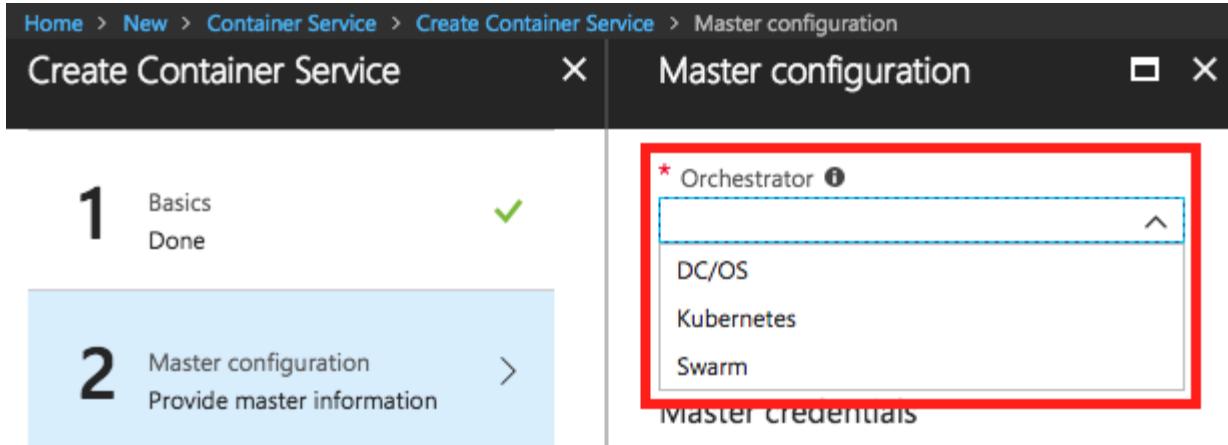
- Name**: testinfradcos (highlighted with a green checkmark)
- Subscription**: A dropdown menu.
- Resource group**:
 - Create new
 - Use existing
- Resource group selection**: test-aks-infra-beta
- Location**: West Europe

A blue 'OK' button is at the bottom right of the configuration pane.

Choosing your target orchestrator

To choose **Master configuration**, you can select the type of container cluster that you want to deploy in ACS. You can choose from among the following three types of orchestrator supported by Microsoft Azure: Kubernetes, Docker Swarm, and DC/OS (datacenter operating system). As they have different architectures, your choice will change the type of credentials

you need in the **Master configuration**. All the orchestrators will need a RSA key, and you may also need a service principal if you are going to deploy a Kubernetes cluster. We'll explain this in the next section:



The screenshot shows the 'Create Container Service' wizard with two steps visible:

- Step 1: Basics** (Done) - A green checkmark is present.
- Step 2: Master configuration** (Provide master information) - This step is currently selected.

In the 'Master configuration' panel, under the heading 'Master credentials', there is a dropdown menu labeled 'Orchestrator' with three options: DC/OS, Kubernetes, and Swarm.

Generating your keygen

In **Master configuration**, under the section **Master credentials**, you have several ways to generate your keygen. Take a look at a way to generate keygen using the following command via cloud shell:

```
ssh-keygen -t rsa -b 2048
```

9963799240 / 7730997544

Ameerpet / Kondapur

The preceding command will generate a public key and a private key with the name that you specified while executing the commands. The following is a sample output to inform the system that you have created your key pairs successfully:

```
melony@Azure:~$ ssh-keygen -t rsa -b 2048
Generating public/private rsa key pair.
Enter file in which to save the key (/home/melony/.ssh/id_rsa): testkeygenk8s
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in testkeygenk8s.
Your public key has been saved in testkeygenk8s.pub.
The key fingerprint is:
SHA256:B6Iu3YjKpLwTOstbBqrGnjeKSy4RQxC0N2Cow0GKZi8 melony@cc-cdaace85-2228111606-9zfhv
The key's randomart image is:
+--[RSA 2048]----+
|B*|
|=oo|
|*+.o . .
|B.o . . .
| E .. S .
| o o+ o .
| o++++ .
| %-*=
| %&=..
+---[SHA256]----+
melony@Azure:~$
```

The cloud shell uses an Azure file storage to persist files across sessions, which was specified when you were starting it the first time. You can use Bash commands, such as ls, to display your files and folders in the current repository as follows:

```
melony@Azure:~$ ls
clouddrive  testkeygenk8s  testkeygenk8s.pub
```

Use the cat command if you want to show the content of your private key:

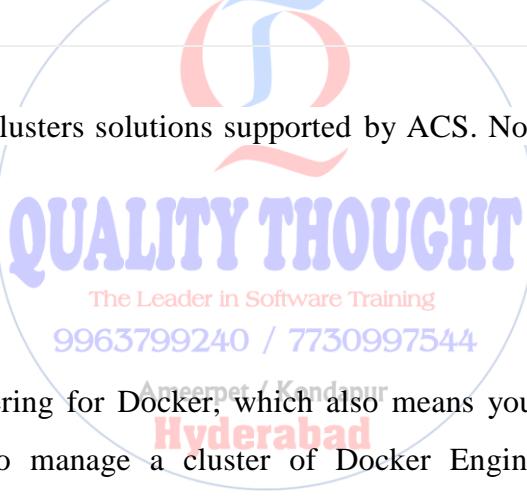
Check public key

Use the cat command if you want to show the content of your public key

Implementing three types of orchestrators of Azure ACS in Azure

There are three container clusters solutions supported by ACS. Now, let's get an overview on each of them.

Docker Swarm



Swarm is the native clustering for Docker, which also means you can use Docker in swarm mode, which is a way to manage a cluster of Docker Engines. Based on the **docker-native** principle, any tools or containers that work with Docker run equally well in Docker Swarm. Originally, Docker Swarm performed a resilient zero single-point-of-failure architecture, secured by default with automatically generated certificates and backward compatibility with existing components.

As an excellent orchestrator, Docker Swarm can be installed and configured in an easy way. As a **docker-native** orchestrator, Docker Swarm can deploy container clusters faster than Kubernetes or other orchestrators, especially in very large clusters or contexts, which requires fast reaction times to scaling on demand. However, for every plus, there is a minus. As it is naturally designed to extend Docker support, the functionalities are limited by the Docker API

that works with the core Docker Engine. That is why it can't support specific complex operations that aren't supported by Docker.

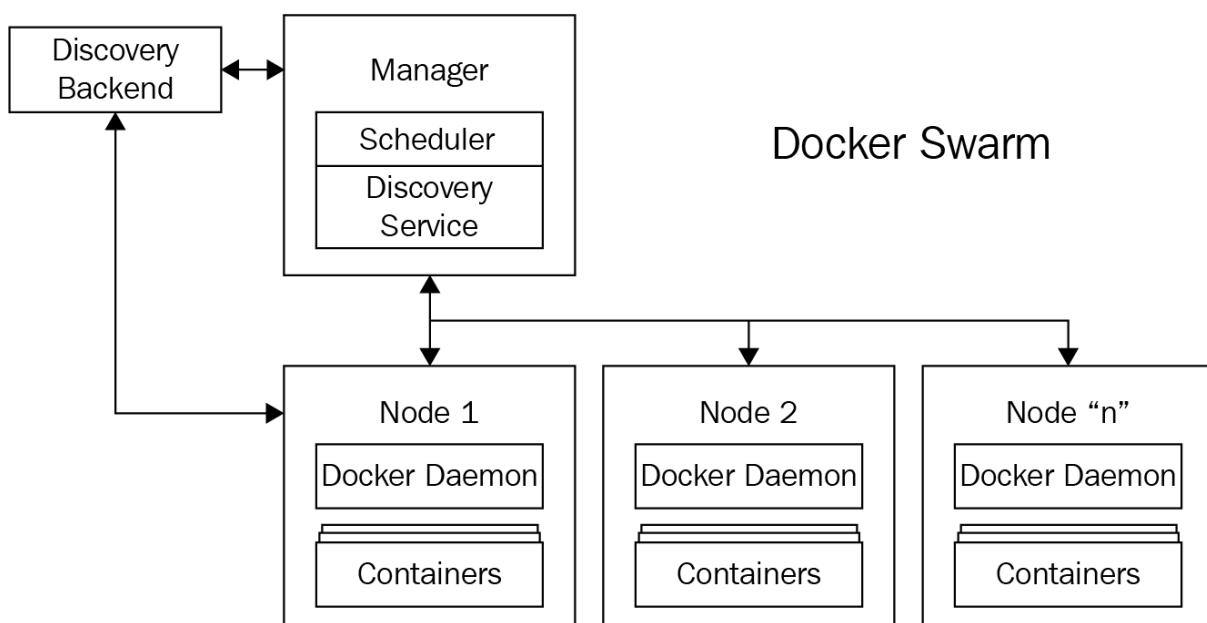
Docker Swarm architecture

There are two types of nodes in Docker Swarm: the manager node and worker node. The concept of a node is an instance of the Docker Engine in the swarm. It is possible to run one or more nodes on a single physical computer or virtual machine in the cloud. It is also possible to run distributed nodes across multiple physical machines and VMs in the cloud.

The two nodes have the following different roles in swarm mode:

- **The manager** node is in charge of dispatching tasks to the worker node. A task is a unit of work that carries a Docker container and the commands to run inside the container.
- **The worker** node runs swarm services and receives and execute tasks dispatched from manager nodes.

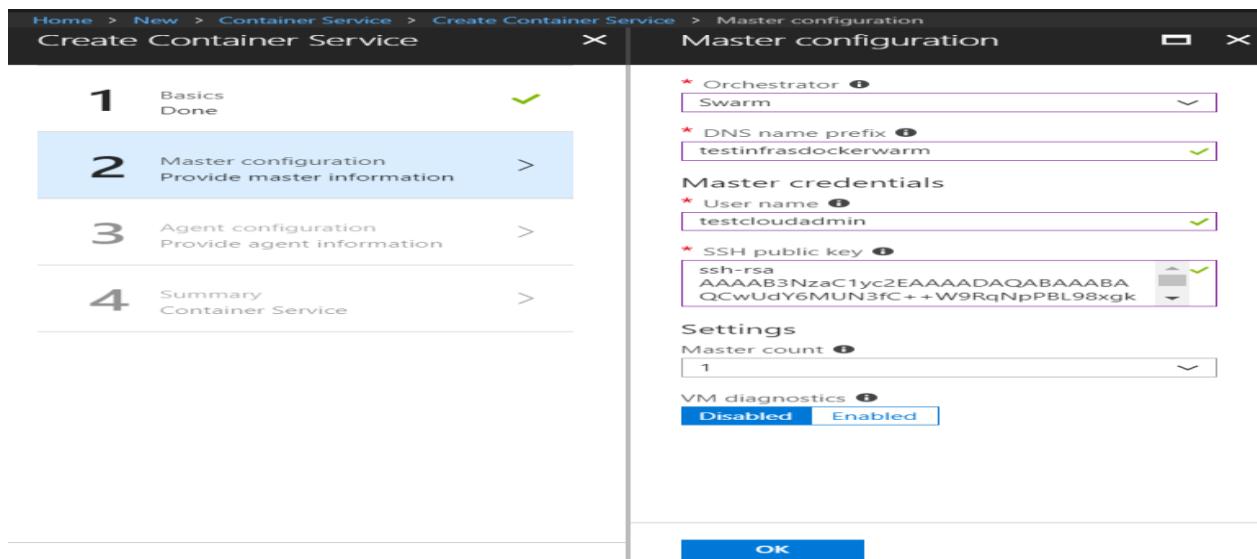
A swarm consists of multiple Docker Hosts, which run in swarm mode and act as managers and workers. A Docker Host can perform as a manager, worker, or both. If a worker node becomes unavailable, Docker is in charge of scheduling that node's tasks to other nodes. This architecture also ensures availability while deploying an application with Docker Swarm. The following diagram shows the **Docker Swarm** architecture:



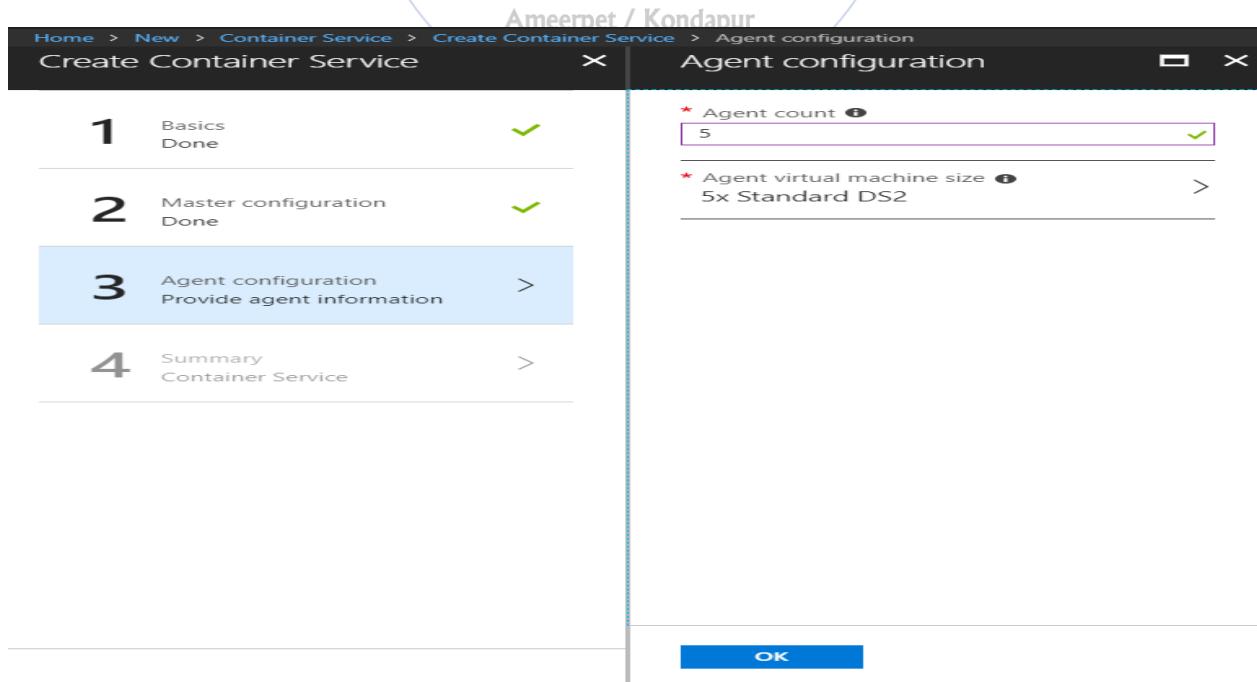
In action – implementing and managing an ACS Docker Swarm cluster

If you're creating an ACS Docker Swarm cluster via the Azure Portal, follow these steps:

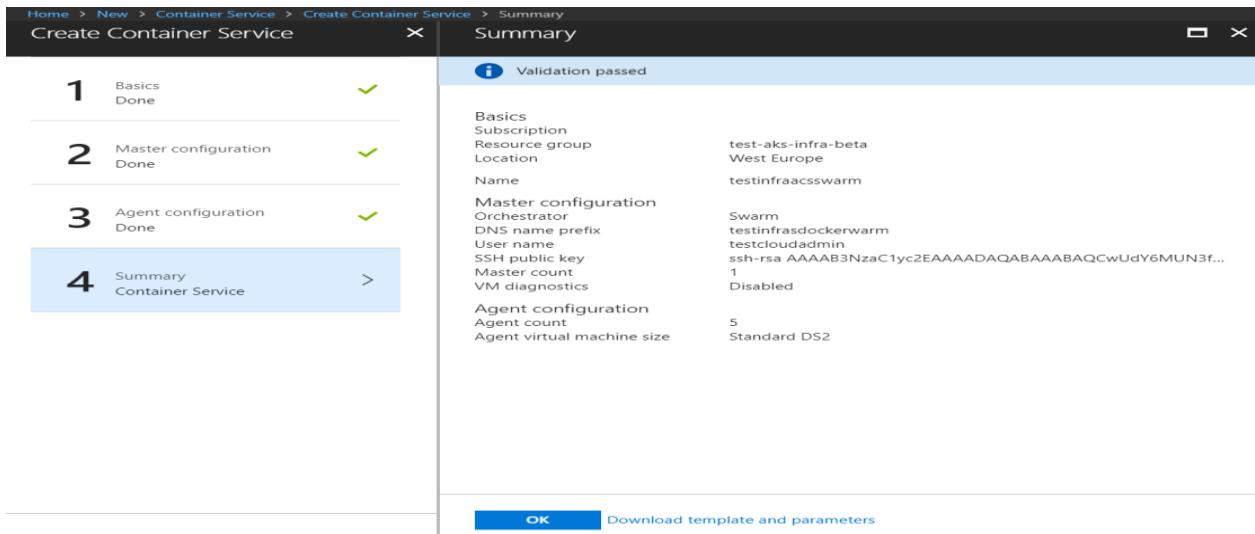
1. Start by creating a **Container Service**, then choose your **Orchestrator** as Swarm in the **Master configuration** step.
2. In the **Master configuration** page, you can specify the **VM size** and the number of nodes you need as the master node:



3. Similarly, you can also specify the **Agent count** and the VM size in the **Agent configuration** page, as shown in the following screenshot:



4. The following is a **Summary** page that is shown after all the information is filled in:



Step	Description	Status
1	Basics Done	✓
2	Master configuration Done	✓
3	Agent configuration Done	✓
4	Summary Container Service	>

Validation passed

Basics

- Subscription: test-aks-infra-beta
- Resource group: West Europe
- Name: testinfraacsswarm

Master configuration

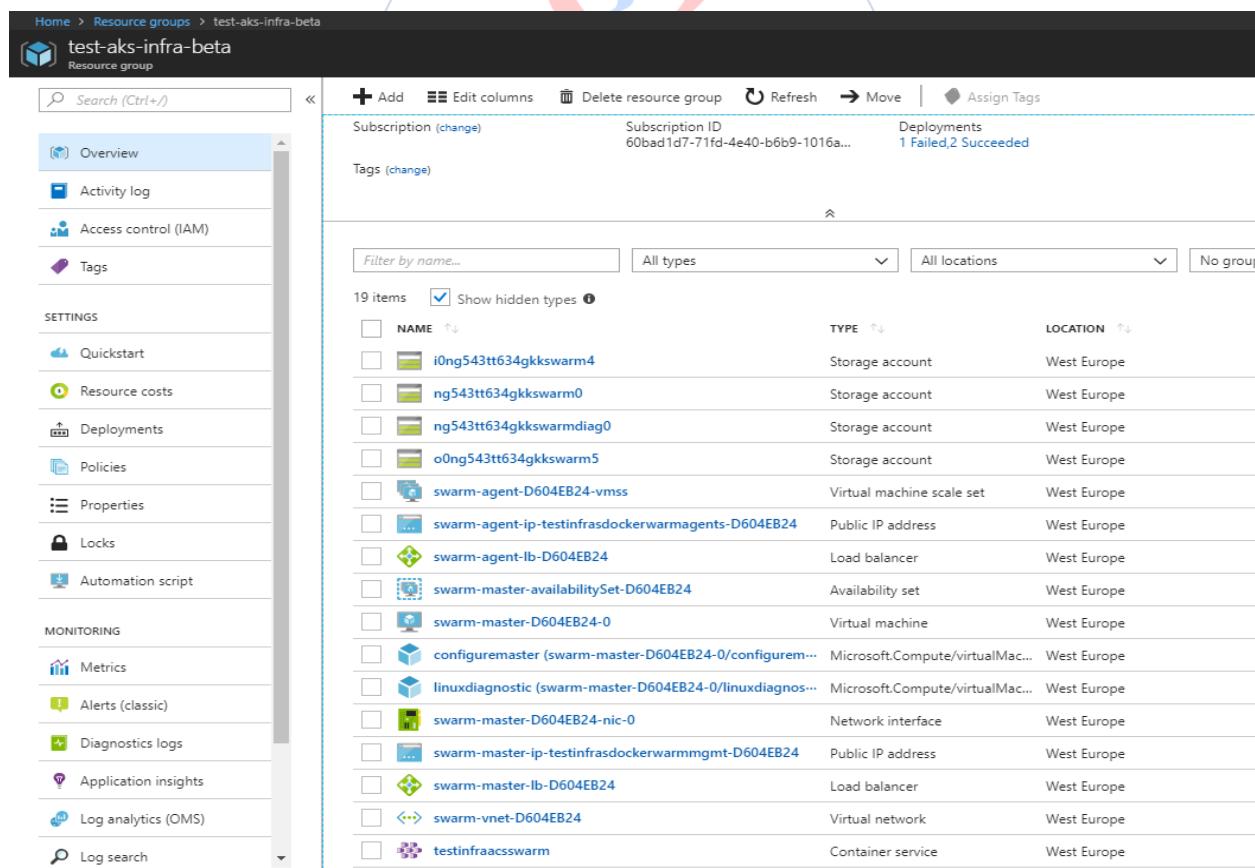
- Orchestrator: Swarm
- DNS name prefix: testinfrasdockerm...
- User name: testcloudadmin
- SSH public key: ssh-rsa AAAAB3NzaC1yc2EAAAQABAAQ...
1
- Master count: 1
- VM diagnostics: Disabled

Agent configuration

- Agent count: 5
- Agent virtual machine size: Standard DS2

OK Download template and parameters

5. The deployment usually takes several minutes. After creating an ACS Docker Swarm cluster successfully, go to the resource group; you can note the deployed resource, as follows:



test-aks-infra-beta Resource group

Overview

Activity log

Access control (IAM)

Tags

SETTINGS

Quickstart

Resource costs

Deployments

Policies

Properties

Locks

Automation script

MONITORING

Metrics

Alerts (classic)

Diagnostics logs

Application insights

Log analytics (OMS)

Log search

Subscription (change): 60bad1d7-71fd-4e40-b6b9-1016a...

Subscription ID: 60bad1d7-71fd-4e40-b6b9-1016a...

Deployments: 1 Failed, 2 Succeeded

NAME	TYPE	LOCATION
i0ng543tt634gkkswarm4	Storage account	West Europe
ng543tt634gkkswarm0	Storage account	West Europe
ng543tt634gkkswarmdiag0	Storage account	West Europe
o0ng543tt634gkkswarm5	Storage account	West Europe
swarm-agent-D604EB24-vmss	Virtual machine scale set	West Europe
swarm-agent-ip-testinfrasdockermagents-D604EB24	Public IP address	West Europe
swarm-agent-lb-D604EB24	Load balancer	West Europe
swarm-master-availabilitySet-D604EB24	Availability set	West Europe
swarm-master-D604EB24-0	Virtual machine	West Europe
configuremaster (swarm-master-D604EB24-0/configurem...)	Microsoft.Compute/virtualMac...	West Europe
linuxdiagnostic (swarm-master-D604EB24-0/linuxdiagnos...)	Microsoft.Compute/virtualMac...	West Europe
swarm-master-D604EB24-nic-0	Network interface	West Europe
swarm-master-ip-testinfrasdockermgmt-D604EB24	Public IP address	West Europe
swarm-master-lb-D604EB24	Load balancer	West Europe
swarm-vnet-D604EB24	Virtual network	West Europe
testinfraacsswarm	Container service	West Europe

6. To create an ACS Docker Swarm, you can use the following command:

```
az acs create --name #yourSwarmCluster# --orchestrator-type dockerce --resource-group  
#yourResourceGroup# --generate-ssh-keys
```

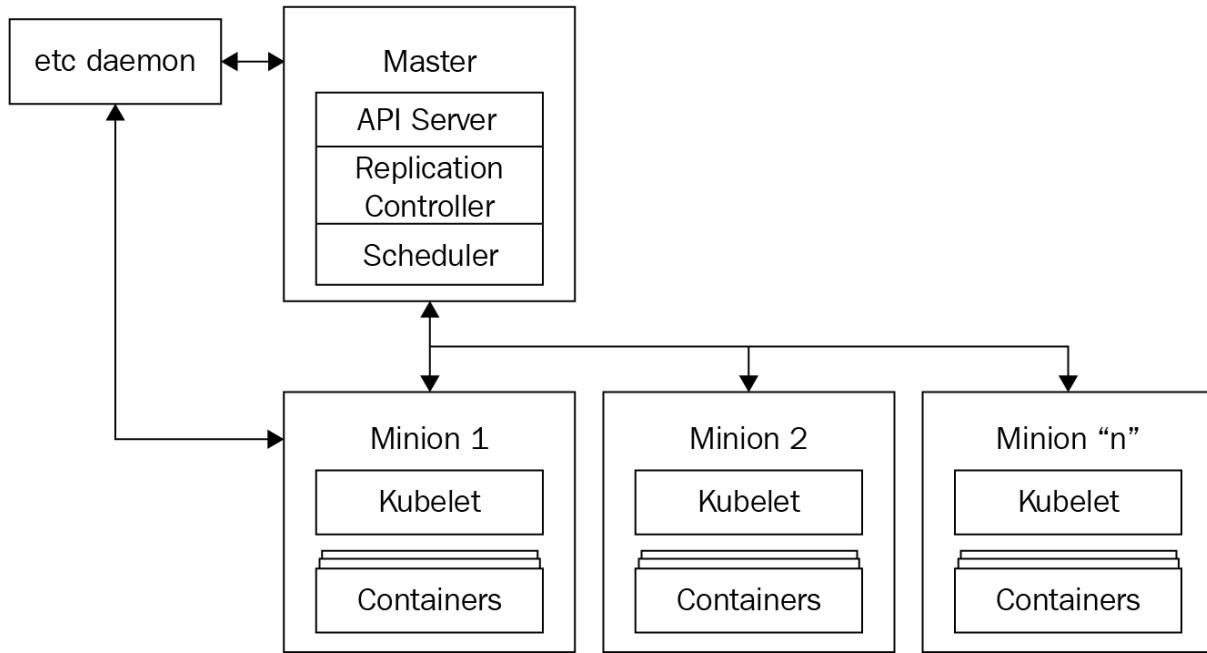
Kubernetes

Kubernetes is an open source platform for container deployment automation, scaling, and operations across clusters of hosts. It aims to provide the components and tools to relieve the burden of running applications in public and private clouds by grouping containers into logical units. The advantage of Kubernetes is flexibility, environment agnostic portability, and easy scaling.

Kubernetes requires a configuration file to configure such as etcd, flannel, the **Docker Engine**, the **cluster configuration** such as the IP addresses of the nodes, which determines which role each node is going to take, and how many nodes there are in total before starting the deployment.

The Kubernetes architecture

Traditionally, the Kubernetes architecture is as follows:



The orchestrators are designed to track and monitor the health of the containers and hosts. In the event of a node failure, orchestrators launch a replacement. We call the mechanism to detect whether the application is operating correctly a **health check**.

Kubernetes supports a user-implemented application health check, which is performed by using kubelet.

The following two ways are used to implement Kubernetes in Azure:

- Deploying the Kubernetes cluster using ACS.
- Deploying the Kubernetes cluster using Azure Kubernetes Service. This service was in preview at the time of writing this book.

In action – implementing an ACS Kubernetes cluster

To create an ACS Kubernetes, start by creating a resource group using the following command:

```
az group create --name test-infra70533 --location westeurope
```

Creating a new AKS Kubernetes cluster

To create an AKS Kubernetes cluster, use the following commands:

```
az acs create --orchestrator-type kubernetes --resource-group test-infra70533 --name infrak8scluster --generate-ssh-keys
```

Alternatively, you can use the following command:

```
az aks create -g test-infra70533 -n infrak8scluster --generate-ssh-keys
```

The `az acs create` command is used to create a Kubernetes cluster in ACS in the resource group `test-infra70533`. The `--generate-ssh-keys` parameter is used to generate new SSH keys. If you want to use your own SSH keys, you can replace the preceding command with the following command:

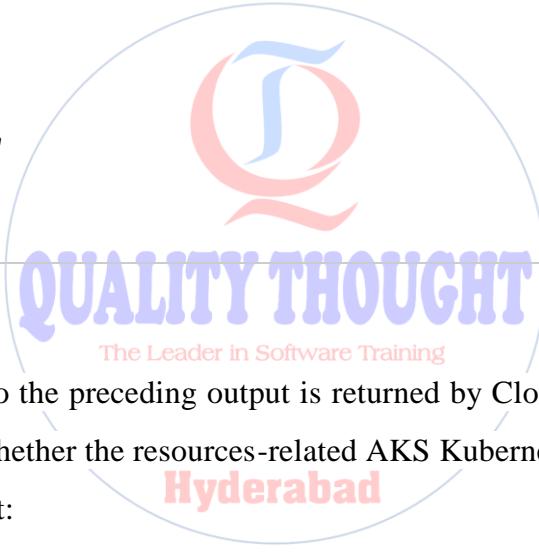
```
az aks create -g test-infra70533 -n infrak8scluster --ssh-key-value /path/to/publickey
```

After several minutes, the command completes and returns **JSON formatted** information about the cluster, which means our Kubernetes cluster has been created successfully. The following is an output information example:

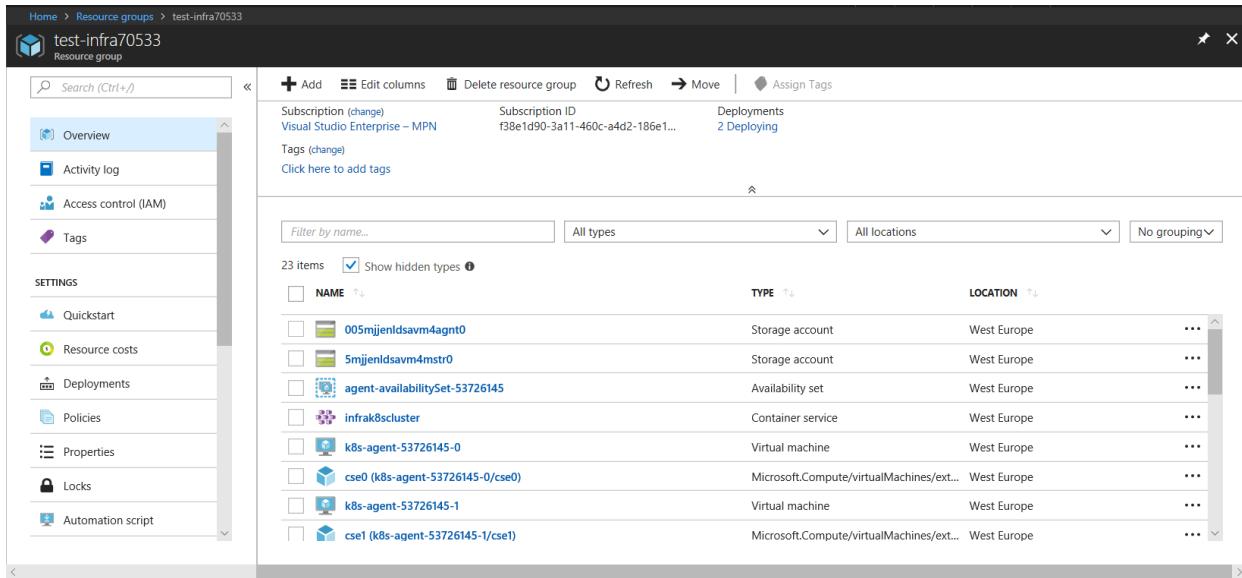
```
{
  "id": "/subscriptions/f38e1d90-3a11-460c-a4d2-186e1660d993/resourceGroups/test-infra70533/providers/Microsoft.Resources/deployments/azurecli1526992751.50577269674",
  "name": "azurecli1526992751.50577269674",
  "properties": {
    "additionalProperties": {
      "duration": "PT13M12.0422159S",
      "outputResources": [
        {
          "id": "/subscriptions/f38e1d90-3a11-460c-a4d2-186e1660d993/resourceGroups/test-infra70533/providers/Microsoft.ContainerService/containerServices/infrak8scluster",
          "resourceGroup": "test-infra70533"
        }
      ],
    }
  }
}
```

```

    "templateHash": "15580770358025216932"
},
"correlationId": "96ba0389-9d85-4685-b95f-c6b7bb216af8",
"debugSetting": null,
"dependencies": [],
"mode": "Incremental",
"outputs": {
    "masterFQDN": {
        "type": "String",
        "value": "infrak8scl-test-infra70533-f38e1dmgmt.westeurope.cloudapp.azure.com"
    },
    "sshMaster0": {
        "type": "String",
        "value": "ssh
azurereuser@infrak8scl-test-infra70533-
f38e1dmgmt.westeurope.cloudapp.azure.com -A -p 22"
    }
},
"parameters": {
    "clientSecret": {
        "type": "SecureString"
    }
}
,
```



When something similar to the preceding output is returned by Cloud Shell, you can go to the resource group to check whether the resources-related AKS Kubernetes are available, as shown in the following screenshot:



NAME	TYPE	LOCATION
005mijenldsavm4agn0	Storage account	West Europe
5mijenldsavm4mstr0	Storage account	West Europe
agent-availabilitySet-53726145	Availability set	West Europe
infrak8scluster	Container service	West Europe
k8s-agent-53726145-0	Virtual machine	West Europe
cse0 (k8s-agent-53726145-0/cse0)	Microsoft.Compute/virtualMachines/ext...	West Europe
k8s-agent-53726145-1	Virtual machine	West Europe
cse1 (k8s-agent-53726145-1/cse1)	Microsoft.Compute/virtualMachines/ext...	West Europe

In action – managing an ACS Kubernetes cluster

To connect to the Kubernetes cluster, use the `kubectl` command, which is Kubernetes, command-line client.

Connecting to a Kubernetes cluster

To configure `kubectl` to connect to your Kubernetes cluster, you should run the `az aks kubernetes get-credentials` command:

```
az aks kubernetes get-credentials --resource-group=test-infra70533 --  
name=infrak8scluster
```

The preceding step allows you to download credentials and configures the Kubernetes CLI to use them. Then, you can use the `kubectl get` command (as shown in the following screenshot) to return a list of the cluster nodes so that you can verify the connection to your cluster:

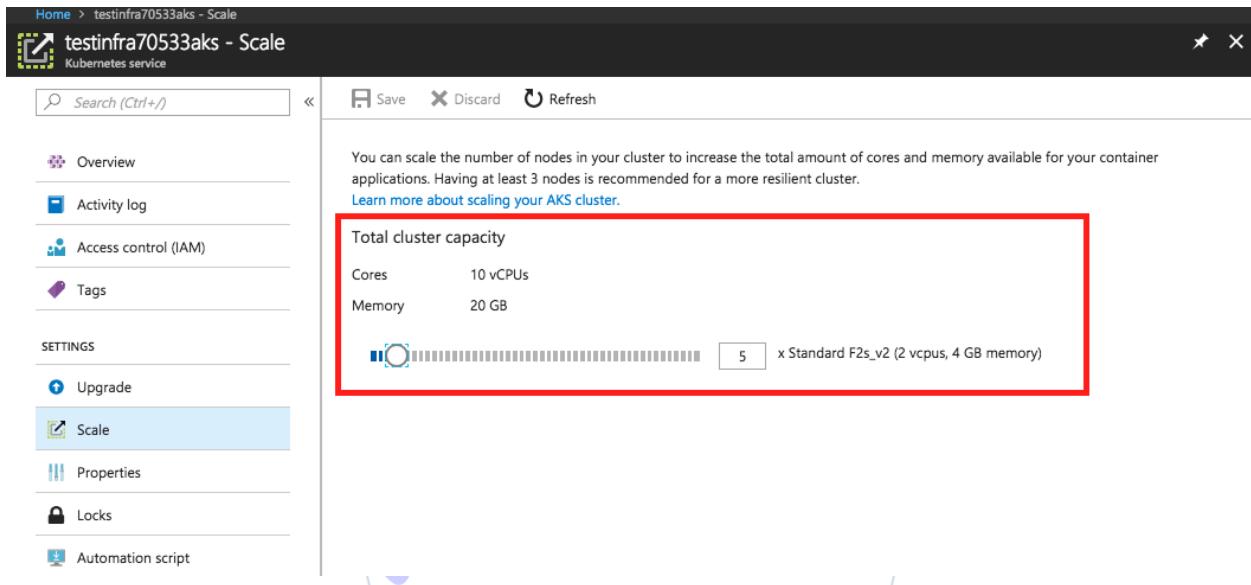
```
kubectl get nodes
```

Scaling the cluster nodes

The `az aks scale` command is used to scale the cluster nodes, as follows:

```
az aks scale --name#yourAKSCluster# --resource-group #yourResourceGroup# --node-count 1
```

Note that the number of the agent node of Kubernetes in the Azure Portal increased from three to five nodes, as shown in the following screenshot:



Home > testinfra70533aks - Scale
testinfra70533aks - Scale Kubernetes service

Save Discard Refresh

You can scale the number of nodes in your cluster to increase the total amount of cores and memory available for your container applications. Having at least 3 nodes is recommended for a more resilient cluster.
[Learn more about scaling your AKS cluster.](#)

Total cluster capacity

Cores	10 vCPUs
Memory	20 GB

5 x Standard F2s_v2 (2 vcpus, 4 GB memory)

Overview Activity log Access control (IAM) Tags

SETTINGS

- Upgrade
- Scale**
- Properties
- Locks
- Automation script

Upgrading an AKS cluster

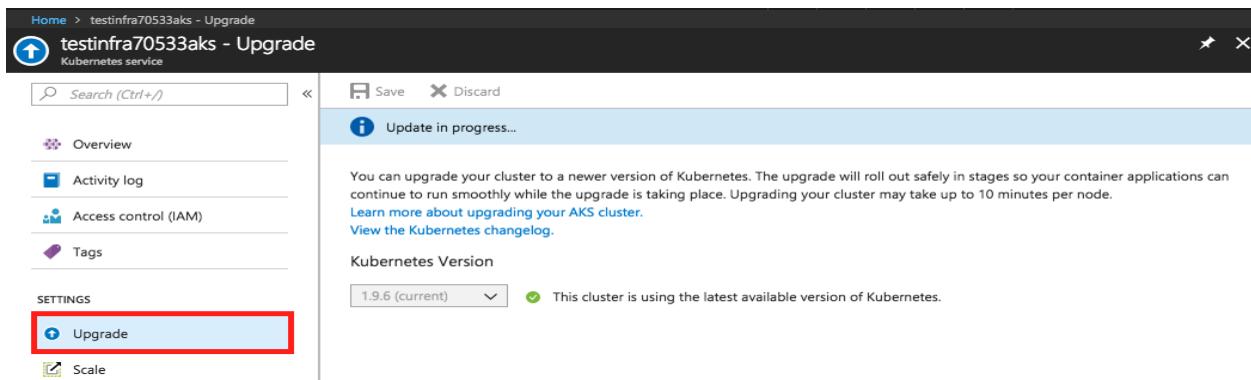
The Leader in Software Training
9963799240 / 7730997544

Ameerpet / Kondapur
Hyderabad

The az aks get-upgrades command is used to scale the cluster nodes, as follows:

```
az aks get-upgrades --name #yourAKSCluster# --resource-group #yourResourceGroup# -o table
```

Another way to upgrade your Kubernetes cluster is via Azure Portal:



Home > testinfra70533aks - Upgrade
testinfra70533aks - Upgrade Kubernetes service

Save Discard

Update in progress...

You can upgrade your cluster to a newer version of Kubernetes. The upgrade will roll out safely in stages so your container applications can continue to run smoothly while the upgrade is taking place. Upgrading your cluster may take up to 10 minutes per node.
[Learn more about upgrading your AKS cluster.](#)
[View the Kubernetes changelog.](#)

Kubernetes Version

1.9.6 (current) This cluster is using the latest available version of Kubernetes.

Overview Activity log Access control (IAM) Tags

SETTINGS

- Upgrade**
- Scale

Deploying applications to a Kubernetes cluster

You can run your application using the kubectl create command, as follows:

```
kubectl create -f #yourapplicationmanifestfile#.yml
```

Deleting Kubernetes clusters

When the cluster is no longer needed, you can use the az group delete command to remove the resource group, container service, and all related resources:

```
az group delete --name test-infra70533 --yes --no-wait
```

Implementing and managing a Kubernetes cluster with AKS

Another way to deploy a Kubernetes cluster is to create a Kubernetes cluster using AKS. ACS implemented the popular container orchestrator-managed Kubernetes in Azure. It simplifies how to create, configure, and manage a container cluster and many useful features, such as the following:

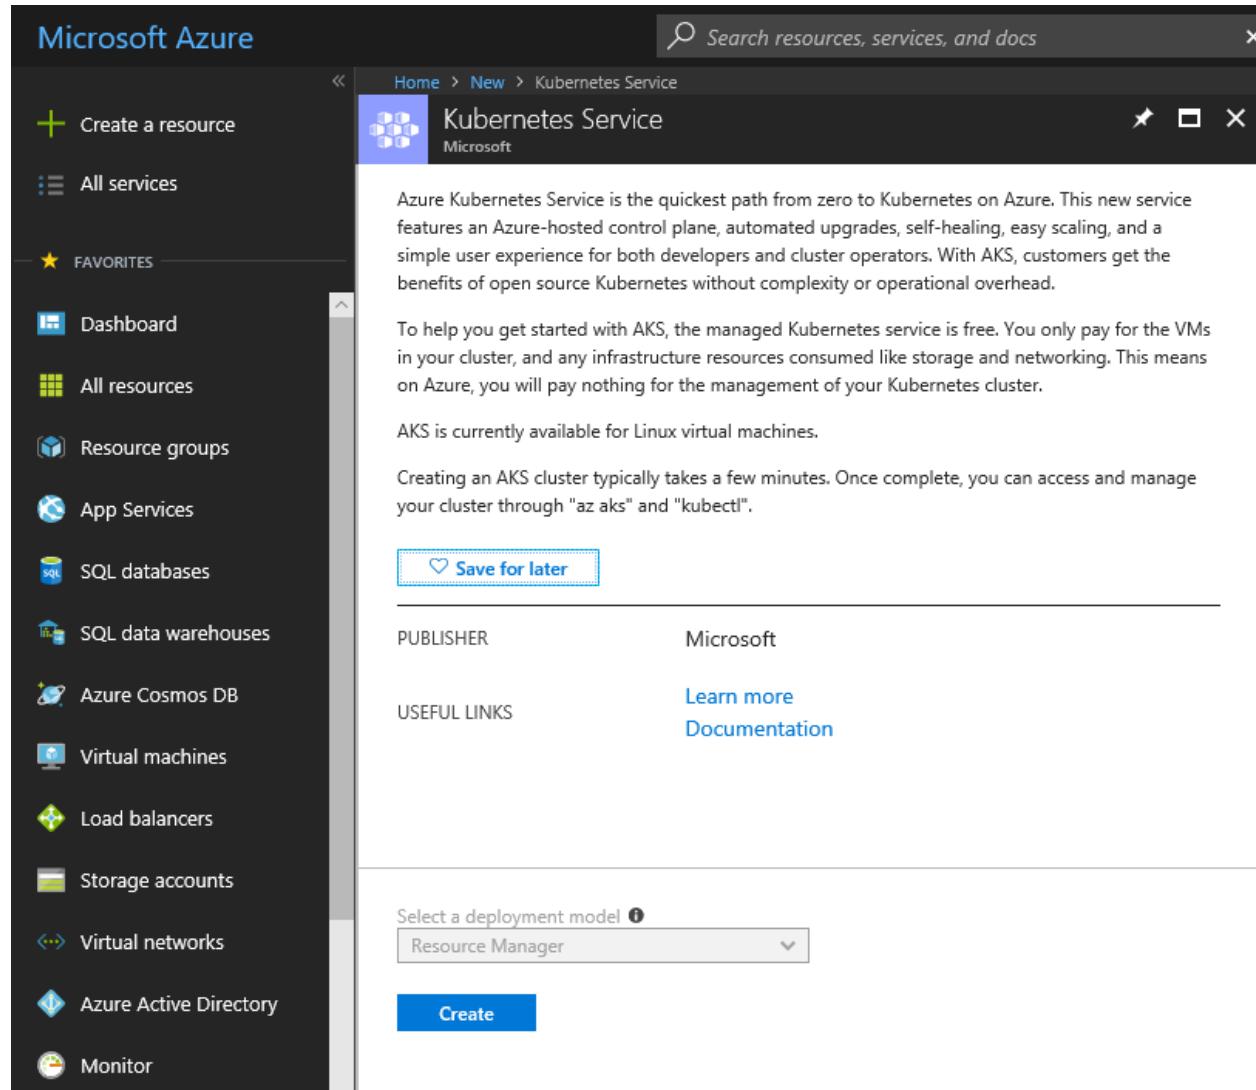
- Easy management of containers even when there are more than 100 instances
- Easy scaling
- Supports popular operating systems, such as Linux and Windows
- Easy rollout and rollback
- Can be combined with batch processes or cron jobs
- Automatic bin packing (depends on GPU / CPU usage, for example)

Using AKS, you can maintain application portability through Kubernetes and the Docker image format, and focus on building a containerized application. Azure will handle the rest of your work, such as container deployment, cluster configuration, and health monitoring.

Regarding pricing in AKS, users pay only for the agent nodes within the clusters, but not for masters, which is in charge of controlling tasks.

Creating AKS via the Azure Portal

From Azure Portal, click on **Create a resource** and search for kubernetes services. Choose it to start creating an AKS service in Azure, as follows:



Microsoft Azure

Search resources, services, and docs

Home > New > Kubernetes Service

Kubernetes Service
Microsoft

Azure Kubernetes Service is the quickest path from zero to Kubernetes on Azure. This new service features an Azure-hosted control plane, automated upgrades, self-healing, easy scaling, and a simple user experience for both developers and cluster operators. With AKS, customers get the benefits of open source Kubernetes without complexity or operational overhead.

To help you get started with AKS, the managed Kubernetes service is free. You only pay for the VMs in your cluster, and any infrastructure resources consumed like storage and networking. This means on Azure, you will pay nothing for the management of your Kubernetes cluster.

AKS is currently available for Linux virtual machines.

Creating an AKS cluster typically takes a few minutes. Once complete, you can access and manage your cluster through "az aks" and "kubectl".

[Save for later](#)

PUBLISHER	Microsoft
USEFUL LINKS	Learn more Documentation

Select a deployment model ?

Resource Manager

Create

From the **Basics** blade, you can choose the version of Kubernetes, the size of machine you'll deploy with Kubernetes, and the number of nodes. Microsoft recommends that you deploy at least three nodes to improve the resilience of your application in production. You can deploy

only one node for test purposes or in the development environments as described in the following screenshot:

Home > New > Kubernetes Service > Create Kubernetes cluster

Create Kubernetes cluster

PROJECT DETAILS

Select a subscription to manage deployed resources and costs. Use resource groups to organize and manage all your resources.

* Subscription [?](#) Visual Studio Enterprise – MPN

* Resource group [?](#) Create new Use existing
testinfraKubernetesRG ✓

CLUSTER DETAILS

* Kubernetes cluster name [?](#) testinfra70533aks ✓

* Region [?](#) West Europe

* Kubernetes version [?](#) 1.9.6

* DNS name prefix [?](#) testinfra70533aksdns ✓

AUTHENTICATION

* Service principal [?](#) (new) default service principal

[Config my service principal](#)

SCALE

The number and size of nodes in your cluster. For production workloads, at least 3 nodes are recommended for resiliency. For development or test workloads, only one node is required. You will not be able to change the node size after cluster creation, but you will be able to change the number of nodes in your cluster after creation. [Learn more about scaling in Azure Kubernetes Service](#)

* Node size [?](#) Standard F2s_v2 (2 vcpus, 4 GB memory) ✓

[Change size](#)

* Node count [?](#) 3

[Review + create](#) [Next: Networking >](#) [Download a template for automation](#)

AKS supports two kinds of network configuration: **basics** and **advanced**. The difference between them is the advanced networking places your pods in an Azure Virtual Network (VNet) that you can configure, providing them with an automatic connectivity to the VNet resources:

[Home](#) > [New](#) > [Kubernetes Service](#) > Create Kubernetes cluster

Create Kubernetes cluster

[Basics](#) [Networking](#) **Monitoring** [Tags](#) [Review + create](#)

You can enable HTTP application routing and choose between two networking options: "Basic" or "Advanced".

- "Basic" networking creates a new VNet for your cluster using default values.
- "Advanced" networking allows clusters to use a new or existing VNet with customizable addresses. Application pods are connected directly to the VNet, which allows for native integration with VNet features.

[Learn more about networking in Azure Kubernetes Service](#)

HTTP application routing 

No Yes

Network configuration 

Basic Advanced

Kubernetes in Azure supports working with OMS (**Log Analytics**) to perform the infrastructure-level monitoring strategy. You can get some basic metrics to monitor the nodes, such as CPU and memory usage, and the health of each node:

[Home](#) > [New](#) > [Kubernetes Service](#) > Create Kubernetes cluster

Create Kubernetes cluster

[Basics](#) **Networking** [Monitoring](#) [Tags](#) [Review + create](#)

With Azure Kubernetes Service, you will get CPU and memory usage metrics for each node. In addition, you can enable container monitoring capabilities and get insights into the performance and health of your entire Kubernetes cluster. You will be billed based on the amount of data ingested and your data retention settings.

[Learn more about container performance and health monitoring](#)

[Learn more about pricing](#)

AZURE MONITOR

Enable container monitoring

Log Analytics workspace 

(new) testinfra70533OMSworkspace 

[Create new](#)

Finally, there is a summary of the information that you've filled in before deploying. All the information will be validated by Microsoft Azure:

[Home](#) > [New](#) > [Kubernetes Service](#) > Create Kubernetes cluster

Create Kubernetes cluster



Validation passed

[Basics](#) [Networking](#) [Monitoring](#) [Tags](#) [Review + create](#)**BASICS**

Subscription	Visual Studio Enterprise – MPN
Resource group	testinfraKubernetesRG
Kubernetes cluster name	testinfra70533aks
Region	West Europe
Kubernetes version	1.9.6
DNS name prefix	testinfra70533aksdns
Node count	3
Node size	Standard_F2s_v2

NETWORKING

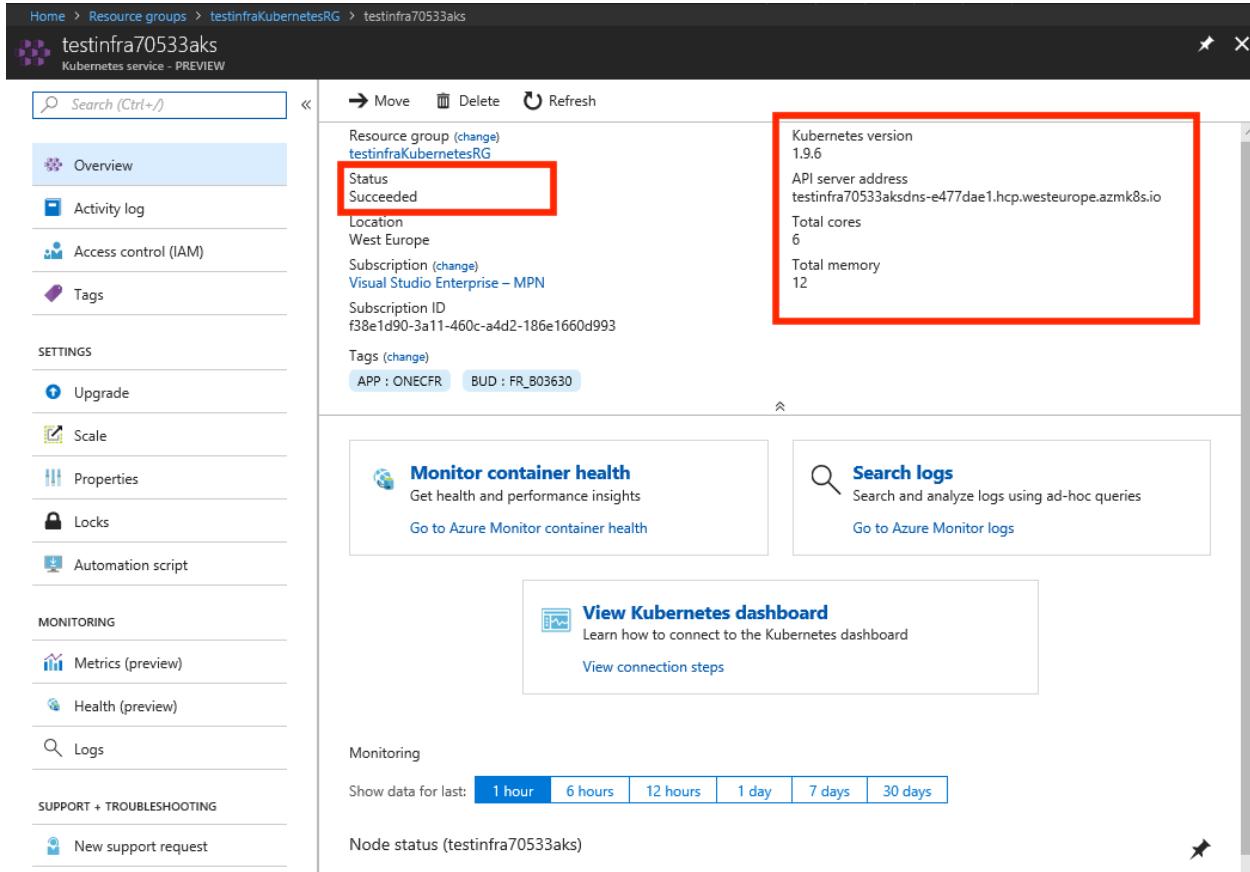
HTTP application routing	No
Network configuration	Basic

MONITORING

Enable container monitoring	Yes
Log Analytics workspace	(new) testinfra70533OMSworkspace

[Create](#)[« Previous: Tags](#)[Download a template for automation](#)

Usually, the deployment will take 3 - 5 minutes to complete. After a successful deployment, you can go to the **Overview** blade of Kubernetes Service and note that the status is **succeeded**. You can also check some other information of your deployed Kubernetes Cluster, such as Kubernetes version and API server address:



Home > Resource groups > testinfraKubernetesRG > testinfra70533aks

testinfra70533aks
Kubernetes service - PREVIEW

Search (Ctrl+ /)

Move Delete Refresh

Resource group (change)
testinfraKubernetesRG

Status
Succeeded

Location
West Europe

Subscription (change)
Visual Studio Enterprise – MPN

Subscription ID
f38e1d90-3a11-460c-a4d2-186e1660d993

Tags (change)
APP : ONECFR BUD : FR_B03630

Kubernetes version
1.9.6

API server address
testinfra70533aksdns-e477dae1.hcp.westeurope.azurek8s.io

Total cores
6

Total memory
12

Monitor container health
Get health and performance insights
Go to Azure Monitor container health

Search logs
Search and analyze logs using ad-hoc queries
Go to Azure Monitor logs

View Kubernetes dashboard
Learn how to connect to the Kubernetes dashboard
View connection steps

Metrics (preview)

Health (preview)

Logs

Monitoring
Show data for last: 1 hour 6 hours 12 hours 1 day 7 days 30 days

Node status (testinfra70533aks)

At the moment, ACS and AKS coexist in Azure. However, Kubernetes is winning over all of these competitors. For many reasons, AKS is becoming more and more important in Microsoft's roadmap.

Implementing and Managing Azure Virtual Networks

Azure networking provides different components in the cloud to help customers create and manage virtual private networks in Azure; it also enables connecting to other virtual networks or their own existing on-premise networks.

In this chapter, we'll cover the following topics:

- The basic concepts and components of Azure networking
- How to implement Azure virtual networks
- How to manage Azure virtual networks
- Traffic filtering and routing
- How to configure a multi-region strategy with Azure Traffic Manager
- How to configure different types of Azure virtual network hybrid and multisite connectivity
- How to monitor Azure networking services
- How to secure Azure virtual networks

Planning and designing Azure virtual networks

9963799240 / 7730997544

Ameerpet / Kondapur

Networking management is a critical and important topic within an organization. Networking connects people, applications, and services. These resources keep businesses running.

Analyze network requirements

Designing a network can be a challenging task. The first step is to understand networking requirements. The main considerations in planning and designing networking solutions are as follows:

- **Connectivity:** This determines the types of connection, public cloud, private cloud, or cross-premise connectivity. The common elements of Microsoft Cloud connectivity will help you to do a check before starting. Here is the link: <https://docs.microsoft.com/en-us/office365/enterprise/common-elements-of-microsoft-cloud-connectivity>.

- **Scalability:** This determines whether the designed network can grow to involve new users, new services, and new applications without affecting the existing services.
- **Availability:** This determines whether the designed network is consistent with reliable performance and offers reasonable response times from and to any host within the network.
- **Security:** This determines the location of security devices, filters, and firewall as well as compatibility with security requirements within the organization.
- **Manageability:** This determines whether the network can be managed effectively and efficiently.

These considerations are becoming less challenging when working with Azure virtual networking in Microsoft Azure. The greatest advantage when working with Azure is that we can make our concepts and design a reality in a "one-click" way. **Go cloud** made a significant difference for organizations which are on the road to digital transformation. When they're at the transition stage, while moving to the cloud customers usually need networking functionality similar to when they were using an on-premise deployment. Microsoft Azure networking components offer a range of functionalities and services that can help organizations design and manage their cloud networking resource. Among all the Azure networking services, Azure virtual network plays a key role.

An **Azure virtual network**, also called a VNet, defines an organization's network in the cloud. It is a logical isolation of the Azure cloud. Within each VNet, there are one or more subnets. The subnets facilitate segmentation of networks, providing a way of controlling communication between network resources.

Different from the traditional understanding of networks, a subnet acts as an address range within a VNet. They can be secured by Network Security Groups (NSGs), which we'll cover in this chapter. So, it is very important to define the address space of a VNET. Each VNet that you connect to another VNet must have a unique address space. Each VNet can have one or more public or private address ranges assigned to its address space.

You can get more inspiration from the Azure documentation about how to plan and design virtual networks in Azure as follows:

<https://docs.microsoft.com/en-us/azure/virtual-network/virtual-network-vnet-plan-design-arm>

Determine the type of connectivity

Before we begin designing a virtual network within the organization, it is important to know whether the transformation will happen in-cloud only, interconnected cloud-only, or cross-premise and interconnected cross-premise as well as which services are in the cloud or on-premise. A cross-premise network should have hybrid network connections between Azure and the on-premise network. These connections can be realized through an Azure gateway, ExpressRoute, and more. We'll cover it later in this chapter.

Determine the address space

In Azure, the address space is defined by the network prefix notation or **classless Inter-Domain Routing (CIDR)**, which is compact for allocating IP addresses and IP routing prefix. The notation is constructed with an IP address, a slash ('/'), and a decimal number: Here is an example of CIDR: 10.0.0.0/16.

In Azure, you can use any range in RFC 1918. In Azure, the smallest supported CIDR is /29, and the largest is /8.

In a single Subnet, Azure reserves the first and last IP addresses of each subnet for protocol conformance, as well as the x.x.x.1-x.x.x.3 addresses of each subnet.

You can consider the following recommendations to define the size of your subnet:

Number of virtual machines	Subnet Size
Up to 3	/29
Up to 11	/28
Up to 27	/27
Up to 59	/26
Up to 123	/25

For a number of technical reasons, Azure doesn't recommend users add the following address ranges:

- Multicast: 224.0.0.0/4
- Broadcast: 255.255.255.255/32
- Loopback: 127.0.0.0/8
- Link-local: 169.254.0.0/16
- Internal DNS: 168.63.129.16/32

Assigning static, public, and private IP addresses

The Leader in Software Training

9963799240 / 7730997544

IP addresses are like our name at home or at work, but are an identity on the internet.

Assigning these IP addresses can be done by using Azure Portal, Azure CLI, or AzurePowerShell.

Public IP versus private IP

We can take an example as public IP is your official name and private IP is our nickname. The translate job between your official name and your nickname is done by a Network Address Translation (NAT), which is a process in which the router translates the private IP Address into a public IP so that you'll be eligible to enter the internet through your **Internet Service Provider (ISP)**.

The Internet Assigned Numbers Authority (IANA) reserves some IP address blocks for usage as private IP addresses, as follows:

- 192.168.0.0 - 192.168.255.255 (about 65,536 IP addresses)
- 172.16.0.0 - 172.31.255.255 (1,048,576 IP addresses)
- 10.0.0.0 - 10.255.255.255 (16,777,216 IP addresses)

To know more about how to create, change, or delete a public IP address, check here: <https://docs.microsoft.com/en-us/azure/virtual-network/virtual-network-public-ip-address>.

Dynamic IP versus static IP

A dynamic IP is an IP address that is constantly changing. A static IP, on the other hand, is one that remains the same. In most cases, dynamic IP is used thanks to the Dynamic Host Control Protocol (DHCP), which is a protocol used to provide automatic, rapid, and central management for the distribution of IP addresses within a network.

DNS for resources in Azure VNets

DNS, which is the abbreviation of **Domain Name System**, is responsible for translating a public hostname, such as a website, public portal, or internal service name such as intranet portal, to its IP address. A simple example may help you understand better; www.pack.com can be translated by DNS to a public address such as 191.239.213.197. Within the Azure virtual network, it is possible to use custom DNS as well as Azure DNS.

Azure DNS provides reliable and secure name resolution using the Microsoft Azure infrastructure. Azure DNS supports all the common record types such as A, AAAA, CNAME, MX, NS, PTR, SOA, SRV, and TXT records.

Azure has public Azure-provided hostnames as well as DNS Private Zones, which provides name resolution both within a virtual network and between virtual networks. To know more about Azure DNS Private Zones scenarios, check the following link: <https://docs.microsoft.com/en-us/azure/dns/private-dns-scenarios>.

You can manage your DNS records using the same credentials, APIs, tools, and billing as your other Azure services via the Azure Portal, Azure PowerShell cmdlets, and Azure CLI.

Applications requiring automatic DNS management can integrate with the service via the REST API and SDKs. You can know more about it by going to the following link: <https://docs.microsoft.com/EN-US/azure/virtual-network/virtual-networks-name-resolution-for-vms-and-role-instances>.

Implementing Azure virtual networks

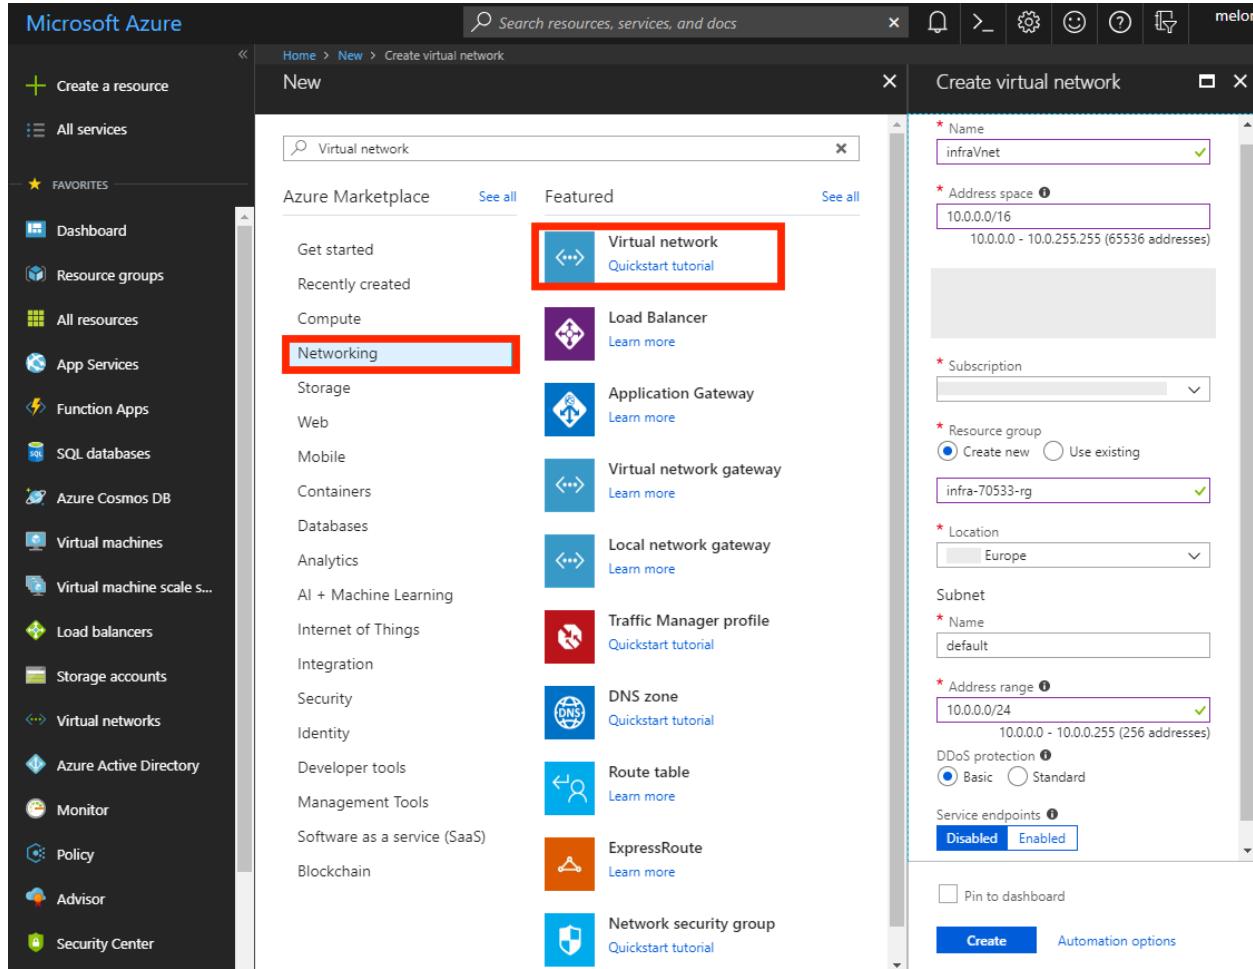
To create an Azure virtual network, you can use the different ways covered in upcoming subsections.

Creating an Azure virtual network

You can create a virtual network via Azure Portal, Azure CLI, PowerShell as well as ARM template.

Via Azure Portal

To create a virtual network, go to Azure Portal and click on **Create a resource**. Then, in the **Networking** category, click on **Virtual network**. You can start filling in the basic information, as shown in the following screenshot:

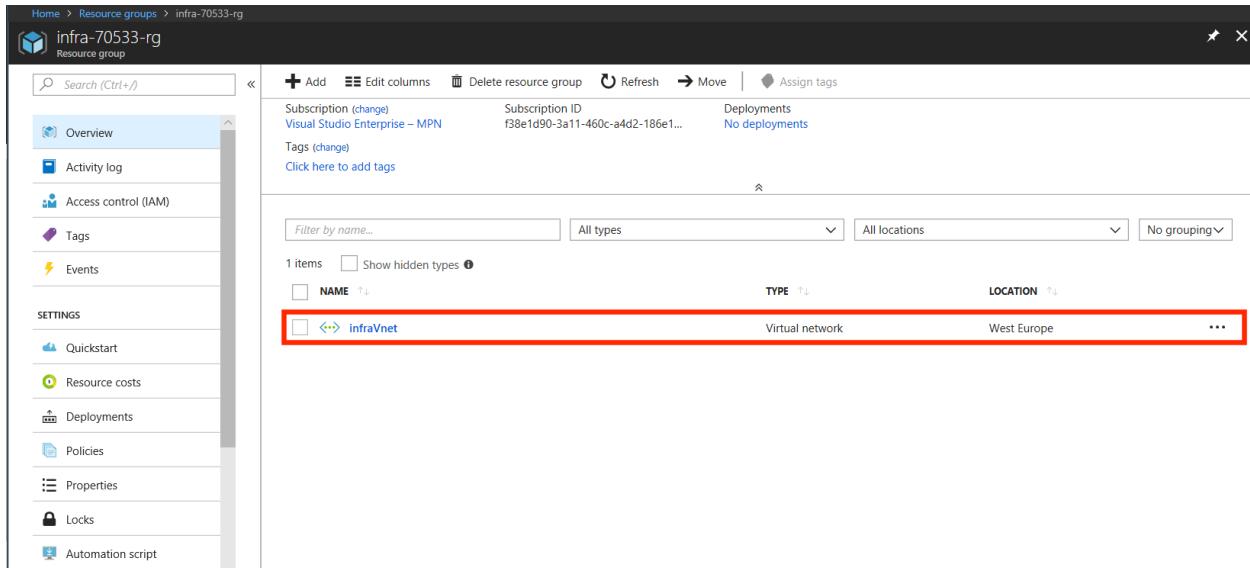


The screenshot shows the Microsoft Azure portal interface. On the left, there's a navigation bar with various service icons. In the center, a search bar at the top says "Search resources, services, and docs". Below it, the breadcrumb navigation shows "Home > New > Create virtual network". A red box highlights the "Virtual network" option in the main content area, which is part of the "Networking" category. To the right, a detailed configuration pane for creating a virtual network is open. It includes fields for "Name" (infraVnet), "Address space" (10.0.0.0/16), "Subscription" (infra-70533-rg), "Resource group" (Create new), "Location" (Europe), "Subnet" (default), "Address range" (10.0.0.0/24), "DDoS protection" (Basic), and "Service endpoints" (Disabled). At the bottom of this pane are "Create" and "Automation options" buttons.

9963799240 / 7730997544

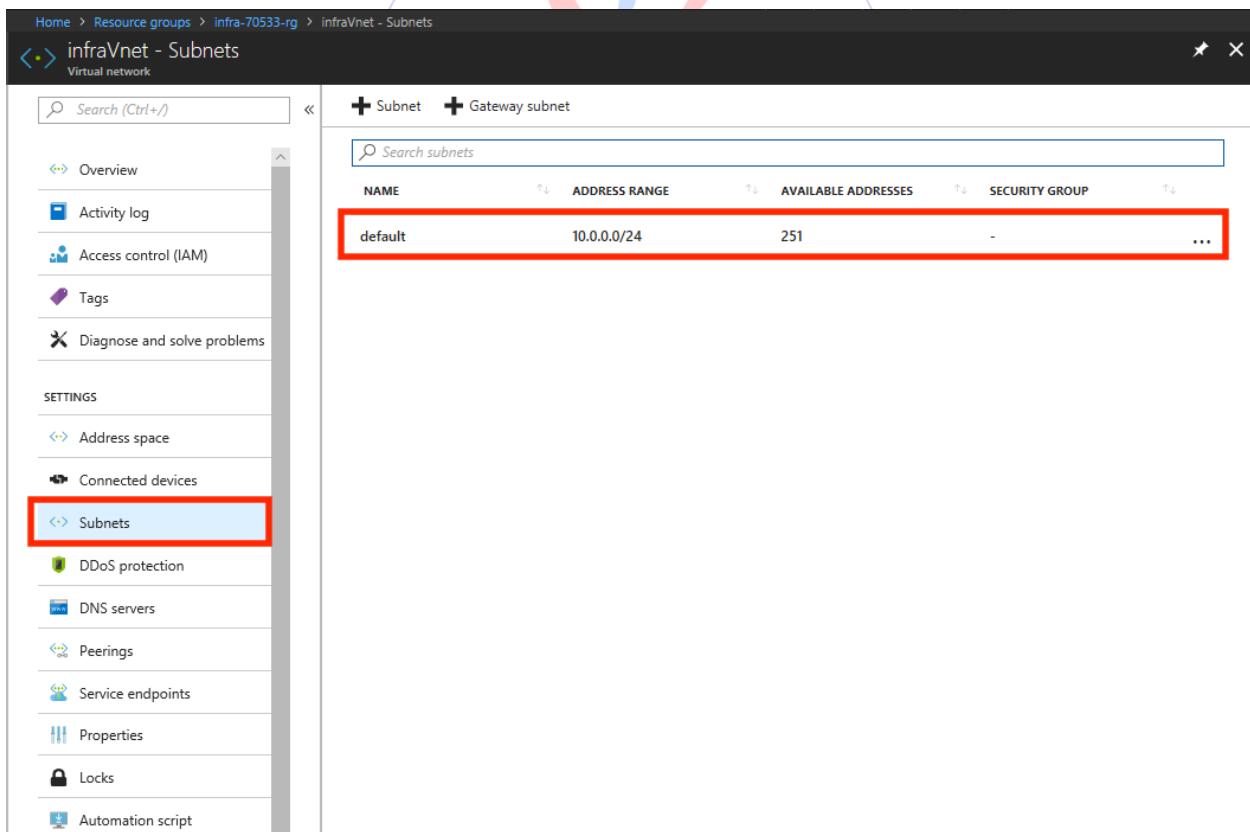
It is very important to specify the address range of your VNet and Subnets. The address range for the subnets or the whole virtual network must be specified with the CIDR notation, defined in RFC1918; within the same Azure virtual network, the address range of the subnets cannot overlap with each other. Every time you create a new Azure virtual network, it will create a default subnet. Finally, you can click on **Create**. The deployment will take just a few minutes.

After creating a VNet successfully, you can go to the **Resource group**. You can see that your VNet is deployed as shown in the following screenshot. You can find information regarding the virtual network such as address space and DNS servers in the **Overview** blade:



A screenshot of the Azure Resource Groups blade for the resource group 'infra-70533-rg'. The left sidebar shows navigation options like Overview, Activity log, Access control (IAM), Tags, Events, and various settings. The main area displays a single item: 'infraVnet' (Virtual network) located in West Europe.

If you go to the VNet, then click on the **Subnets** blade. You will see that the default subnet is in the VNet, as shown in the following screenshot:



A screenshot of the Azure Subnets blade for the 'infraVnet' virtual network. The left sidebar highlights the 'Subnets' option. The main area shows a table with one row for the 'default' subnet, which has an address range of 10.0.0.0/24 and 251 available addresses.

NAME	ADDRESS RANGE	AVAILABLE ADDRESSES	SECURITY GROUP
default	10.0.0.0/24	251	-

Via Azure PowerShell

You can also create a virtual network using PowerShell. The commands are as follows:

```
# Create a resource group.
$resgroupname = "infra-70533-rg"
$location = "West Europe"
```

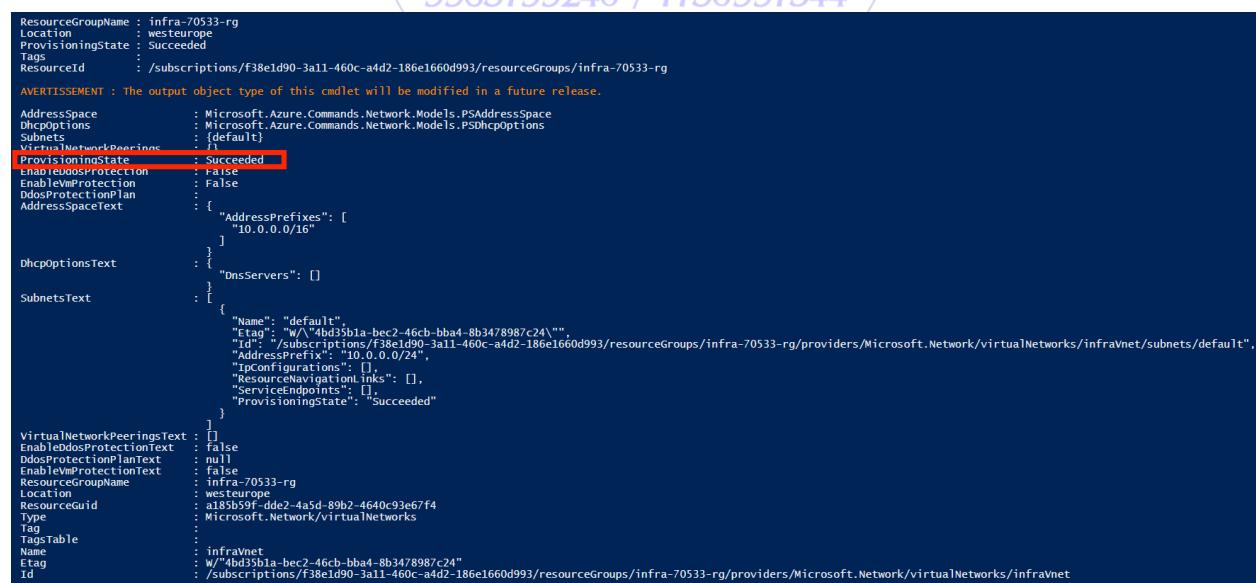
```
New-AzureRmResourceGroup -Name $resgroupname -Location $location
```

```
# Create a new virtual network
$virtualNetwork = New-AzureRmVirtualNetwork -ResourceGroupName $resgroupname -Location $location -Name infraVnet -AddressPrefix 10.0.0.0/16
```

```
# Define subnet configuration
$subnetConfig = Add-AzureRmVirtualNetworkSubnetConfig -Name default -AddressPrefix 10.0.0.0/24 -VirtualNetwork $virtualNetwork
```

```
#Apply the designed configuration to new virtual network
$virtualNetwork | Set-AzureRmVirtualNetwork
```

If everything is okay, you will see output similar to the following screenshot:



```
ResourceGroupName : infra-70533-rg
Location         : westeurope
ProvisioningState : Succeeded
Tags             :
ResourceId       : /subscriptions/f38e1d90-3a11-460c-a4d2-186e1660d993/resourceGroups/infra-70533-rg
AVERTISSEMENT : The output object type of this cmdlet will be modified in a future release.
AddressSpace     : Microsoft.Azure.Commands.Network.Models.PSAddressSpace
DhcpOptions      : Microsoft.Azure.Commands.Network.Models.PSDhcpOptions
Subnets          : [{}]
VirtualNetworkPeeringsText:
ProvisioningState : Succeeded
EnableDdosProtection : false
EnableVmProtection : false
DdosProtectionPlanText:
AddressSpaceText : {
    "AddressPrefixes": [
        "10.0.0.0/16"
    ]
}
DhcpOptionsText : {
    "DnsServers": []
}
SubnetsText : [
    {
        "Name": "default",
        "Id": "/subscriptions/f38e1d90-3a11-460c-a4d2-186e1660d993/resourceGroups/infra-70533-rg/providers/Microsoft.Network/virtualNetworks/infraVnet/subnets/default",
        "Etag": "W/\"4bd35b1a-bea2-46cb-bba4-8b3478987c24\"",
        "Type": "Microsoft.Network/virtualNetworks",
        "Location": "westeurope",
        "ResourceGroup": "infra-70533-rg",
        "AddressPrefix": "10.0.0.0/24",
        "IpConfigurations": [],
        "ResourceNavigationLinks": [],
        "ServiceEndpoints": [],
        "ProvisioningState": "Succeeded"
    }
]
VirtualNetworkPeeringsText:
EnableDdosProtectionText : false
DdosProtectionPlanText :
EnableVmProtectionText :
ResourceGroupText:
Location           : westeurope
ResourceGuid       : a185b59f-dde2-4a5d-89b2-4640c93e67f4
Type              : Microsoft.Network/virtualNetworks
Tags              :
Name              : infraVnet
Etag              : W/\"4bd35b1a-bea2-46cb-bba4-8b3478987c24"
Id                : /subscriptions/f38e1d90-3a11-460c-a4d2-186e1660d993/resourceGroups/infra-70533-rg/providers/Microsoft.Network/virtualNetworks/infraVnet
```

Via Azure CLI

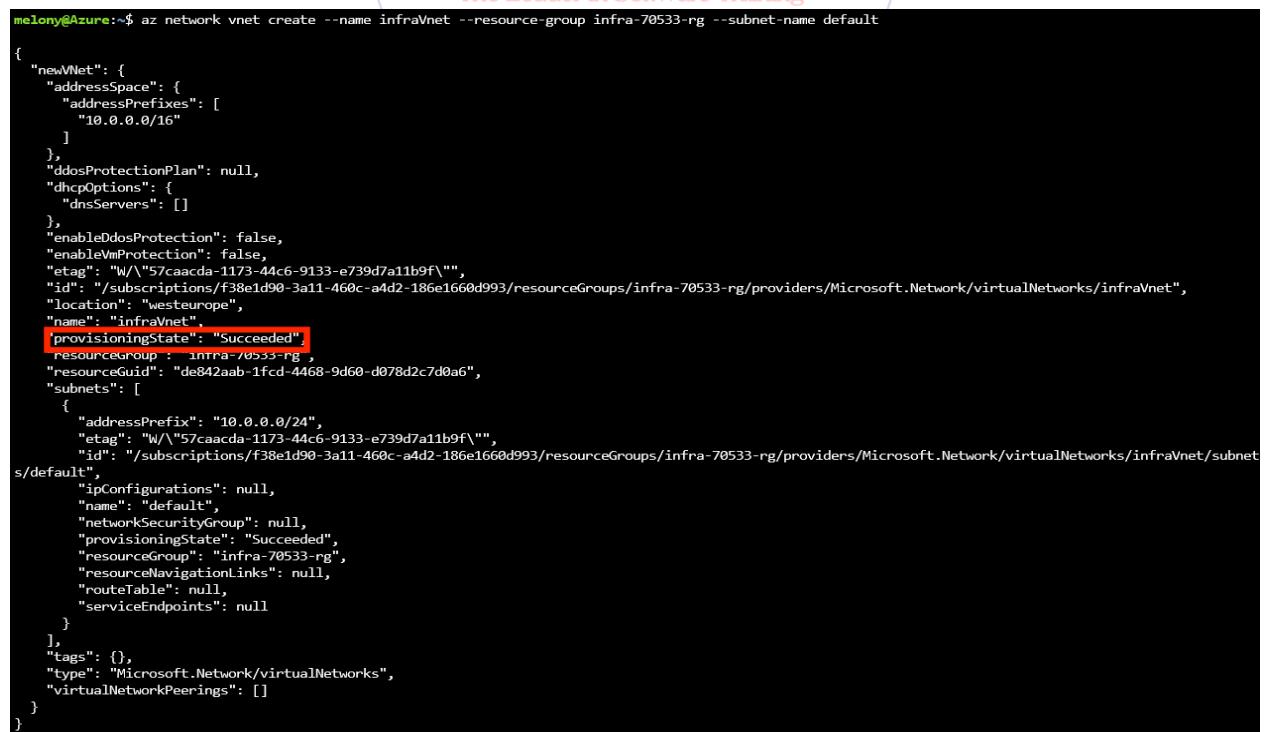
To achieve the same result, you can also use Azure CLI as shown in the following code. You can start by creating a new resource group ():

```
az group create --name #yourresourcegroup# --location #yourlocation#
```

You can create a new Azure VNet by the following command:

```
az network vnet create \
--name #youvnetname# \
--resource-group #yourresourcegroup# \
--subnet-name #yoursubnetname,here is default#
```

If everything is okay, you will see an output similar to the following screenshot. If ProvisioningState is marked Succeeded as shown below, we have created a Azure VNet by using Azure CLI:



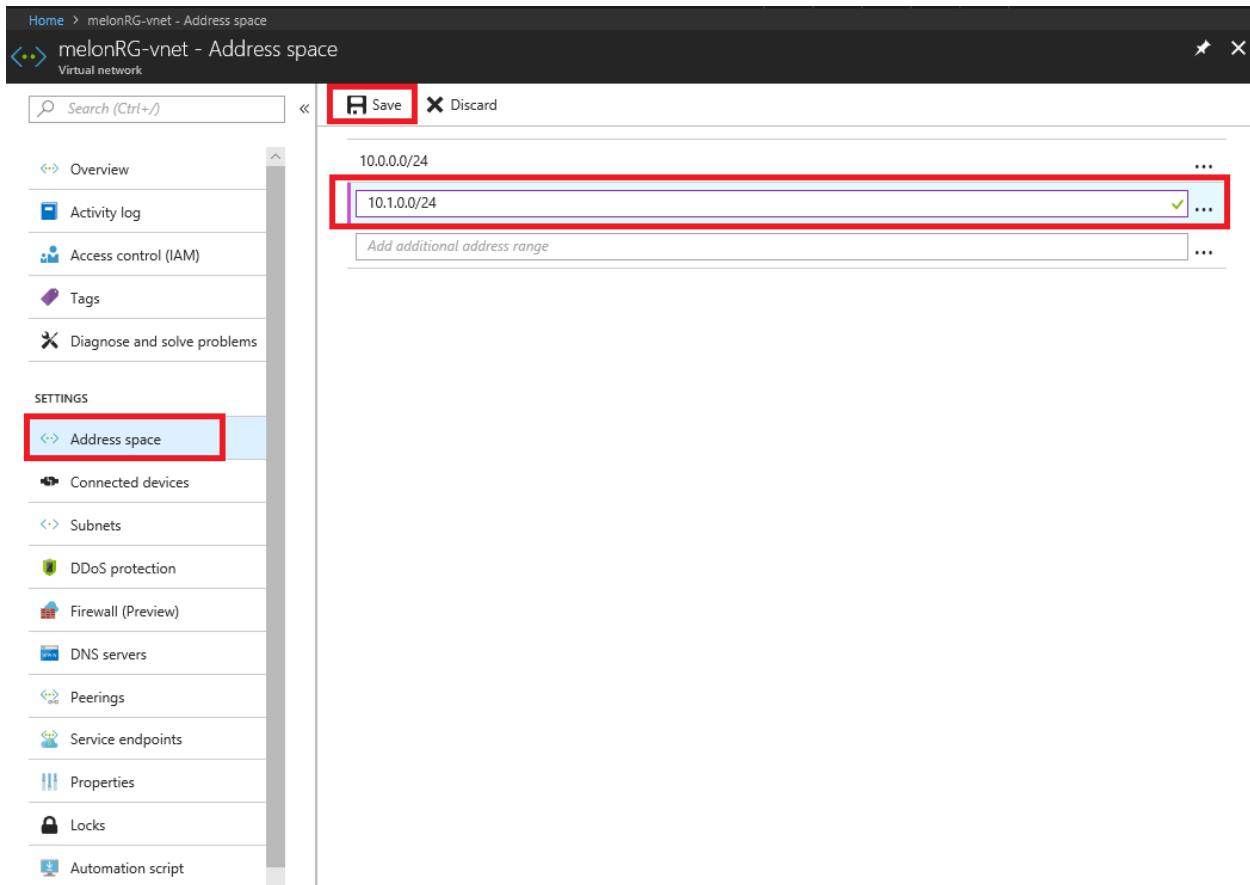
```
melony@Azure:~$ az network vnet create --name infraVNet --resource-group infra-70533-rg --subnet-name default
{
  "newVNet": {
    "addressSpace": {
      "addressPrefixes": [
        "10.0.0.0/16"
      ]
    },
    "ddosProtectionPlan": null,
    "dhcpOptions": {
      "dnsServers": []
    },
    "enableDdosProtection": false,
    "enableVmProtection": false,
    "etag": "W/"57caacda-1173-44c6-9133-e739d7a11b9f"",
    "id": "/subscriptions/f38e1d90-3a11-460c-a4d2-186e1660d993/resourceGroups/infra-70533-rg/providers/Microsoft.Network/virtualNetworks/infraVNet",
    "location": "westeurope",
    "name": "infraVNet",
    "provisioningState": "Succeeded",
    "resourceGroup": "infra-70533-rg",
    "resourceGuid": "de842aab-1fcf-4468-9d60-d078d2c7d0a6",
    "subnets": [
      {
        "addressPrefix": "10.0.0.0/24",
        "etag": "W/"57caacda-1173-44c6-9133-e739d7a11b9f"",
        "id": "/subscriptions/f38e1d90-3a11-460c-a4d2-186e1660d993/resourceGroups/infra-70533-rg/providers/Microsoft.Network/virtualNetworks/infraVNet/subnet
s/default",
        "ipConfigurations": null,
        "name": "default",
        "networkSecurityGroup": null,
        "provisioningState": "Succeeded",
        "resourceGroup": "infra-70533-rg",
        "resourceNavigationLinks": null,
        "routeTable": null,
        "serviceEndpoints": null
      }
    ],
    "tags": {},
    "type": "Microsoft.Network/virtualNetworks",
    "virtualNetworkPeerings": []
  }
}
```

Updating the Azure virtual network

It is possible to add or remove address ranges and change DNS servers after creating the Azure virtual network. You can do that using Azure Portal, Azure CLI, or PowerShell.

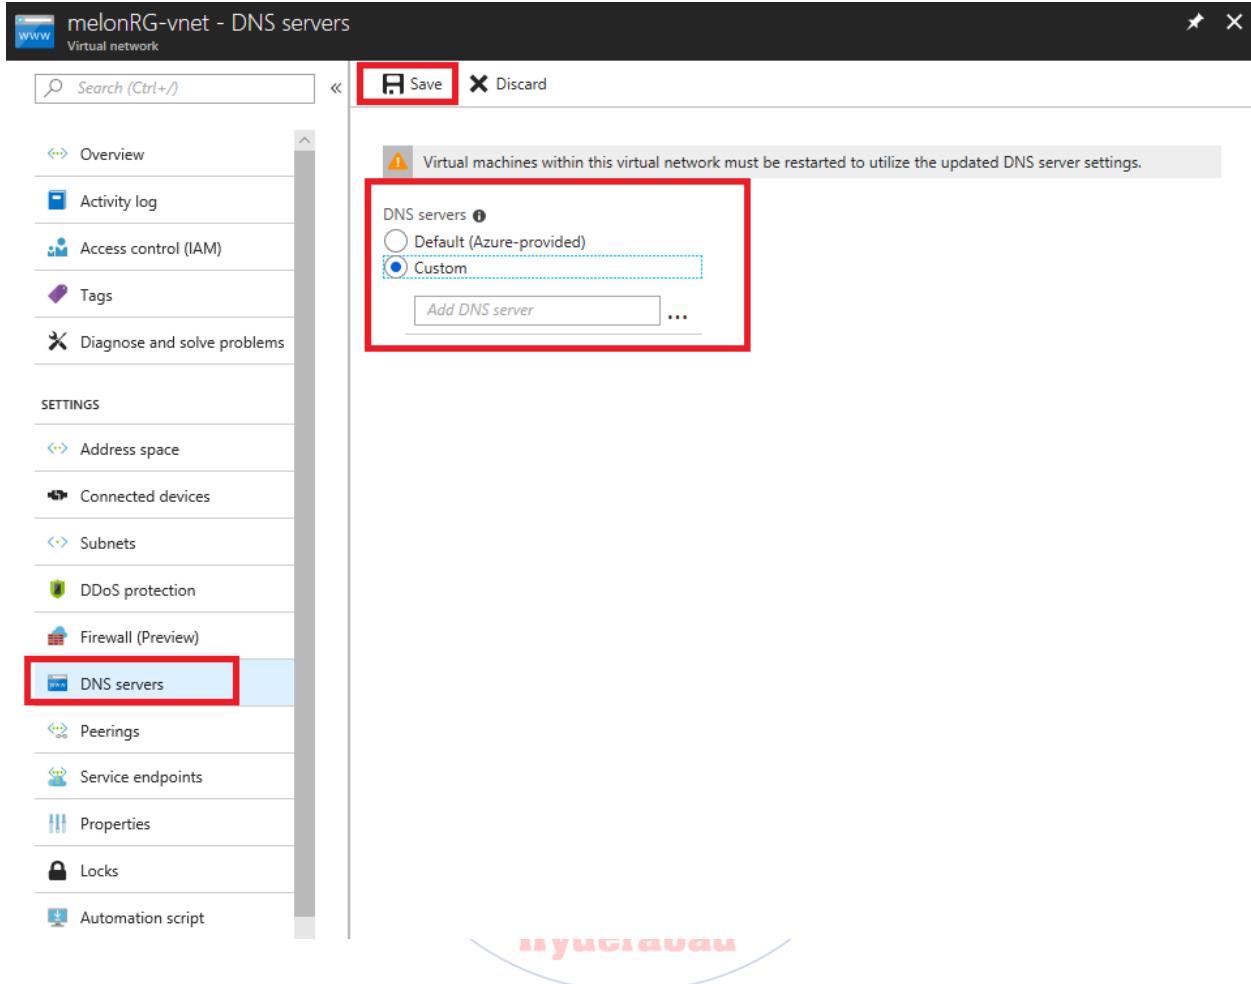
Via Azure Portal

To add or remove the address range, you can go to the **Address space** blade of the Azure virtual network and add or remove a space range, as shown in the following screenshot:



Name resolution for the devices connected to the current virtual network is managed by Azure DNS by default. Users can decide whether they want to choose the Azure internal DNS server or a custom DNS server.

To set a DNS server, you can go to the **DNS server** blade to use Azure default DNS or a custom server, as shown in the following screenshot:



Via Azure CLI

You can use the following Azure CLI cmdlet to update the virtual network:

```
az network vnet update -g #yourresourcegroupname# -n #VNetname@ --dns-servers 10.1.0.5
```

There are some other settings as well that can be updated using the same command. You can check the following link for more information: <https://docs.microsoft.com/en-us/cli/azure/network/vnet?view=azure-cli-latest#az-network-vnet-update>.

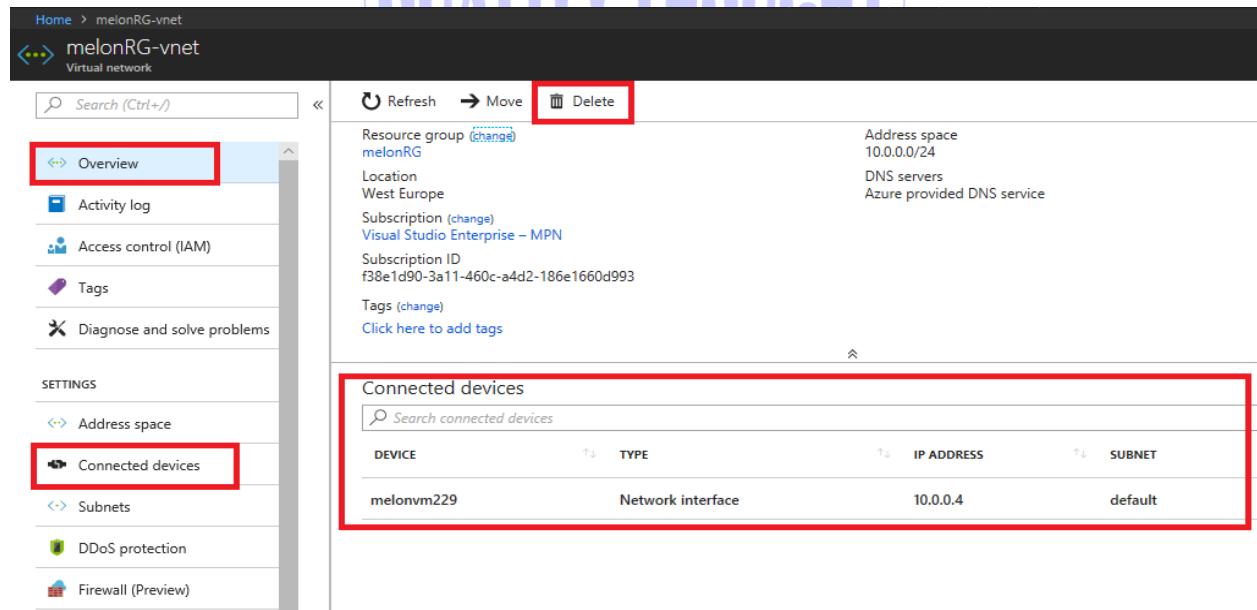
Via Azure PowerShell

The same result can be achieved using the Set-AzureRmVirtualNetwork cmdlet, set the expected state of a virtual network:

```
$virtualNetwork = New-AzureRmVirtualNetwork -Name MyVirtualNetwork -  
ResourceGroupName  
TestResourceGroup -Location westeurope -AddressPrefix "10.0.0.0/16"  
$virtualNetwork | Set-AzureRmVirtualNetwork
```

Delete Azure virtual network

When a virtual network is no longer needed, you can delete it via Azure Portal by clicking on **Delete** in the **Overview** blade. Just to remind you, when we said a VNet is no longer in use we meant there are no devices connected with this VNet. You can check the connected devices in the **Overview** or **Connected devices** blades in the virtual network (as shown in the following screenshot):



DEVICE	TYPE	IP ADDRESS	SUBNET
melonvm229	Network interface	10.0.0.4	default

You can also use the following Azure CLI cmdlet to delete a virtual network:

```
az network vnet delete -g MyResourceGroup -n myVNet
```

Alternatively, you can delete the current virtual network using Azure PowerShell:

```
Remove-AzureRmVirtualNetwork -Name #virtualnetworkname# -ResourceGroupName  
#resourcegroupname#
```

Managing Azure virtual networking

Network traffic from the internet, on-premise, or any other cloud providers to Azure can be routed, filtered, and distributed thanks to the different PaaS services of Azure networking. In this section, we'll discuss different Azure networking components to route, filter, and distribute networking traffic.

Routing network traffic

In Azure, there are a couple of options to route network traffic between subnets or connected VNets in Azure, on-premise, and the internet. Typically, you can optionally use **User-defined routes (UDR)** to override Azure's default routing or using **Border gateway protocol (BGP)** routes through a network gateway.

User-defined routes

Azure routes traffic between Azure, on-premise, and the internet. Azure automatically creates a route table containing a set of routes, which specifies how to route traffic within a virtual network for each subnet within an Azure virtual network and the system default routes will be added to the table.

User-defined routes will be very useful when you want to override some of Azure's system routes with custom routes and add additional custom routes to route tables. Azure routes outbound traffic from a subnet based on the routes in the route table of the subnet. The relationship between the route table and subnet is **one to many**, which means each route table

can be associated to multiple subnets, but a subnet can only be associated to a single route table.

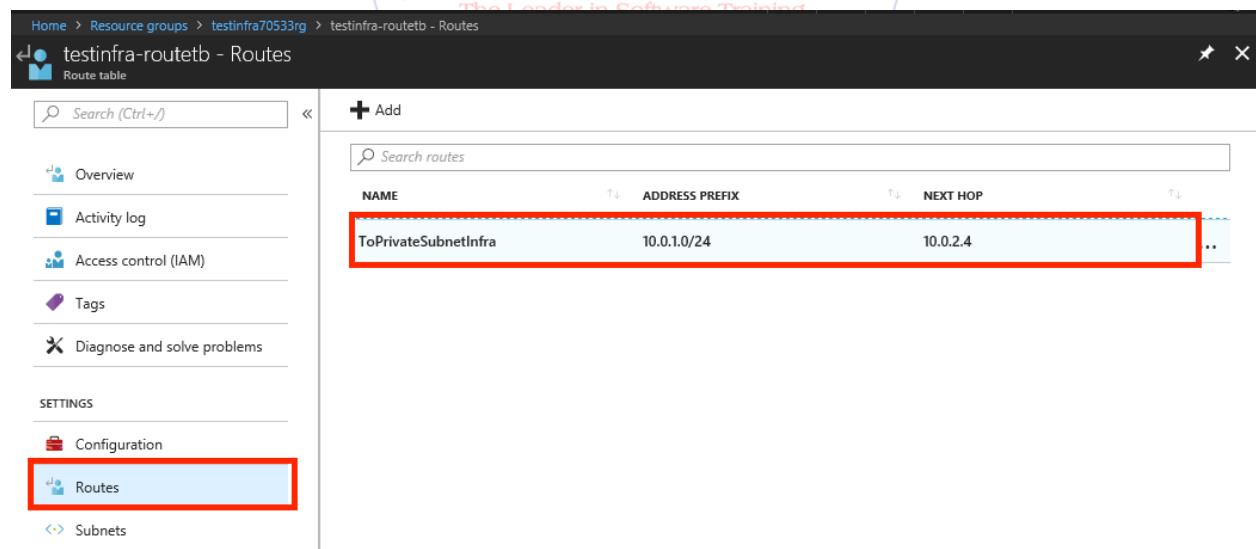
Border gateway protocol

Border gateway protocol (BGP) is a standard routing protocol commonly used on the internet to exchange routing and reachability information between two or more VNets. Users can connect the virtual network to an on-premise network using an Azure VPN Gateway or ExpressRoute connection. You can know more about BGP by checking: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-bgp-overview>.

Managing routes in a route table

To know more about how to create a route table, delete a route table, and manage routes in a route table, you can go to the following link: <https://docs.microsoft.com/en-us/azure/virtual-network/manage-route-table>.

You can use the Azure Portal to check an existing user-defined route by clicking on the **Routes** blade of the created route table, as shown in the following screenshot:



NAME	ADDRESS PREFIX	NEXT HOP
ToPrivateSubnetInfra	10.0.1.0/24	10.0.2.4

You can also use the following Azure CLI `az network route-table route show` cmdlet to consult the routes in the route table:

```
az network route-table route list -g testinfra70533rg --route-table-name testinfra-routetb
```

In Azure PowerShell, the `Get-AzureRmRouteConfig` cmdlet is also able to check the route config in the route table. The following is a sample command:

```
Get-AzureRmRouteTable -ResourceGroupName "testinfra70533rg" -Name "testinfra-routetb" | Get-AzureRmRouteConfig -Name "ToPrivateSubnetInfra"
```

Filtering the network traffic

Inbound and outbound network traffic between subnets in Azure can be filtered by source IP address and port, destination IP address and port, and protocol using **Network security groups (NSG)** or **Virtual Network Appliances(VNA)**.

Network security groups

A network security group (NSG) contains a list of security rules that allow or deny network traffic access to resources connected to Azure Virtual Networks (VNet). NSGs can be associated to subnets, individual VMs (classic), or individual network interfaces (NIC) attached to VMs (Resource Manager). When an NSG is associated to a subnet, the rules apply to all resources connected to the subnet. Traffic can further be restricted by also associating an NSG to a VM or NIC.

Virtual Network Appliances

A **Virtual Network Appliance (VNA)** is a VM running software that performs a network function, such as a firewall. VNAs are also available that provide WAN optimization and other network traffic functions. VNAs are typically used with user-defined or BGP routes. You can also use a NVA to filter traffic between VNets.

Distributing network traffic

As you probably know, there are different options such as Azure Load Balancer, Traffic Manager, and Application Gateway to distribute network traffic using Microsoft Azure.

These three options work on different layers of the OSI model. They have different feature sets. Users can use these services individually or combine their methods depending on their needs to build the optimal solution. The three options are explained as follows:

- **Azure Load Balancer:** This works at the transport layer, which is level 4 of the OSI model. It provides network-level distribution of traffic across instances generally in the same Azure region. Users can configure public and internal load-balanced endpoints and define rules to map inbound connections to backend pool destinations using TCP and HTTP health-probing options to manage service availability.
- **Traffic Manager:** This is another load-balancing solution that is included within Azure. It works at the DNS level. You can use Traffic Manager to load-balance between endpoints that are located in different Azure regions. Users can configure this load-balancing service to use different traffic distribution methods such as priority, weighted performance, or geographic routing methods.
- **Application Gateway:** This is a load balancer that works at level 7 of the OSI model, which is also known as the application layer. It provides load-balanced solutions for network traffic that is based on the HTTP protocol. It uses routing rules as application-level policies that can offload Secure Sockets Layer (SSL) processing from load-balanced VMs. Similar to ALB, the Application Gateway can be configured as an internet-facing gateway, an internal-only gateway, or a combination of both.

These three options work differently from each other, have different feature sets, and support different scenarios. You can use these services individually or combine their methods, depending on your needs, to build the optimal solution.

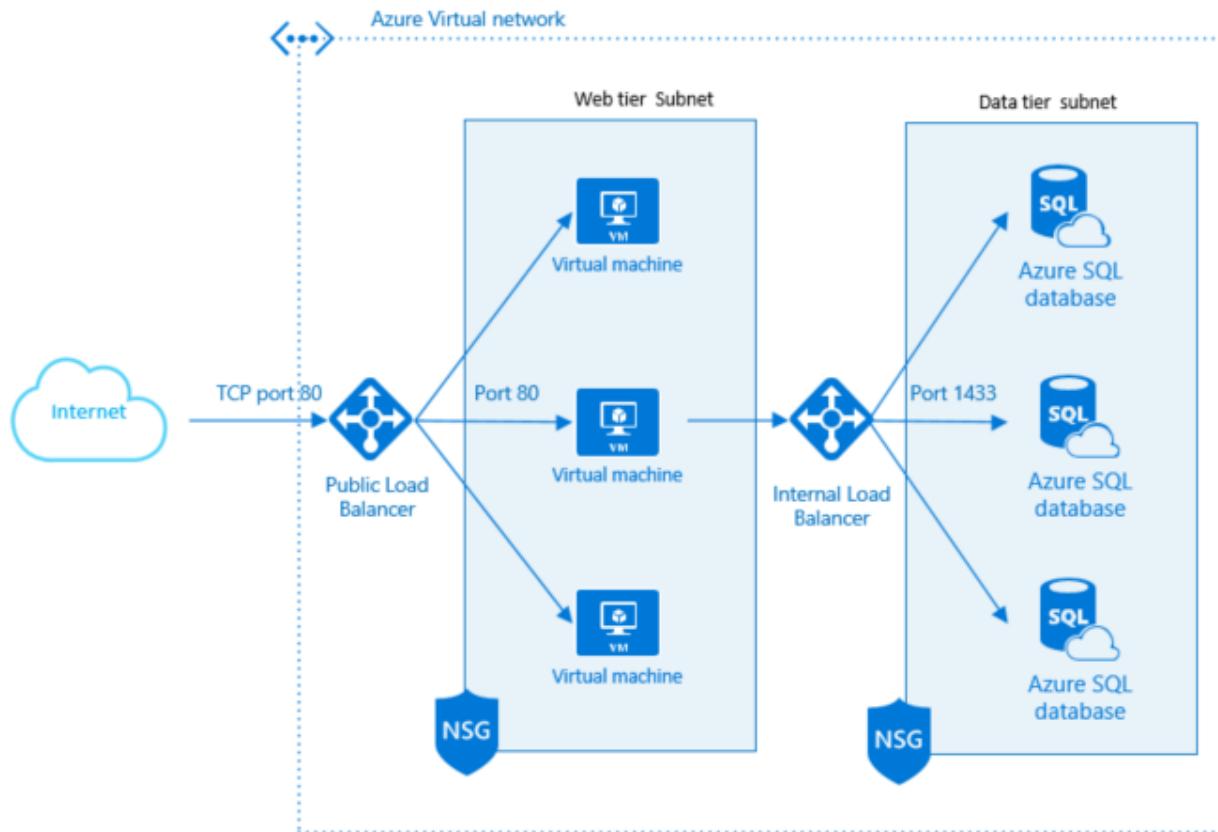
Configure external and internal load balancing

Load Balancer distributes new inbound traffic to arrive at the load balancer's frontend to the backend pool instances based on rules and health probes.

There are two kinds of load balancer:

- **Internet-facing load balancer or public load balancer:** This helps to distribute the incoming internet traffic to web front VMs
- **Internal load balancer:** This helps to distribute traffic across VMs inside a virtual network such as data tier VMs

You can combine both types of load balancer using the following schema:



Combine Public Load Balancer and Internal Load Balancer in the same scenario

Azure Load Balancer is available in two SKUs, Basic and Standard, based on scalability, availability, pricing, and other features.

To know more about how to create a load balancer in the basic tier, you can check the following links:

To use the Azure Portal: <https://docs.microsoft.com/en-us/azure/load-balancer/quickstart-create-basic-load-balancer-portal>.

To use Azure CLI: <https://docs.microsoft.com/en-us/azure/load-balancer/quickstart-create-basic-load-balancer-cli>.

To use Azure PowerShell: <https://docs.microsoft.com/en-us/azure/load-balancer/quickstart-create-basic-load-balancer-powershell>.

Load balancing in the standard tier provides a higher level of availability and scalability. It can distribute incoming requests across multiple Azure VMs. There is some more configuration work to do while creating the Standard Load balancer. You can check the following links for more information.

Creating a standard load balancer via the Azure Portal: <https://docs.microsoft.com/en-us/azure/load-balancer/quickstart-load-balancer-standard-public-portal>.

Creating a standard load balancer via Azure CLI: <https://docs.microsoft.com/en-us/azure/load-balancer/quickstart-load-balancer-standard-public-cli>

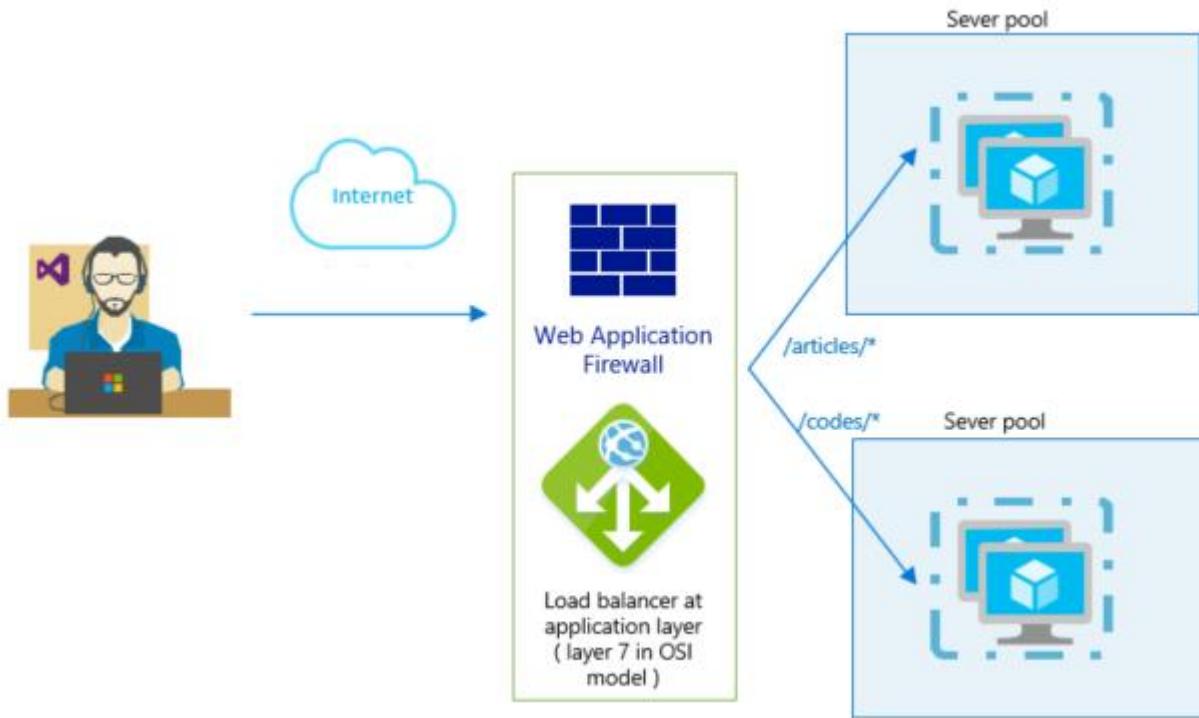
9963799240 / 7730997544

Creating a standard load balancer via Azure PowerShell: <https://docs.microsoft.com/en-us/azure/load-balancer/quickstart-create-standard-load-balancer-powershell>.

Implementing an Application Gateway

We can create an Application Gateway with a **web application firewall (WAF)** feature using Azure Portal. The WAF uses rules from the **Open Web Application Security Project (OWASP)** core rule sets to protect your application. These rules include protection against attacks such as SQL injection, cross-site scripting attacks, and session hijacks.

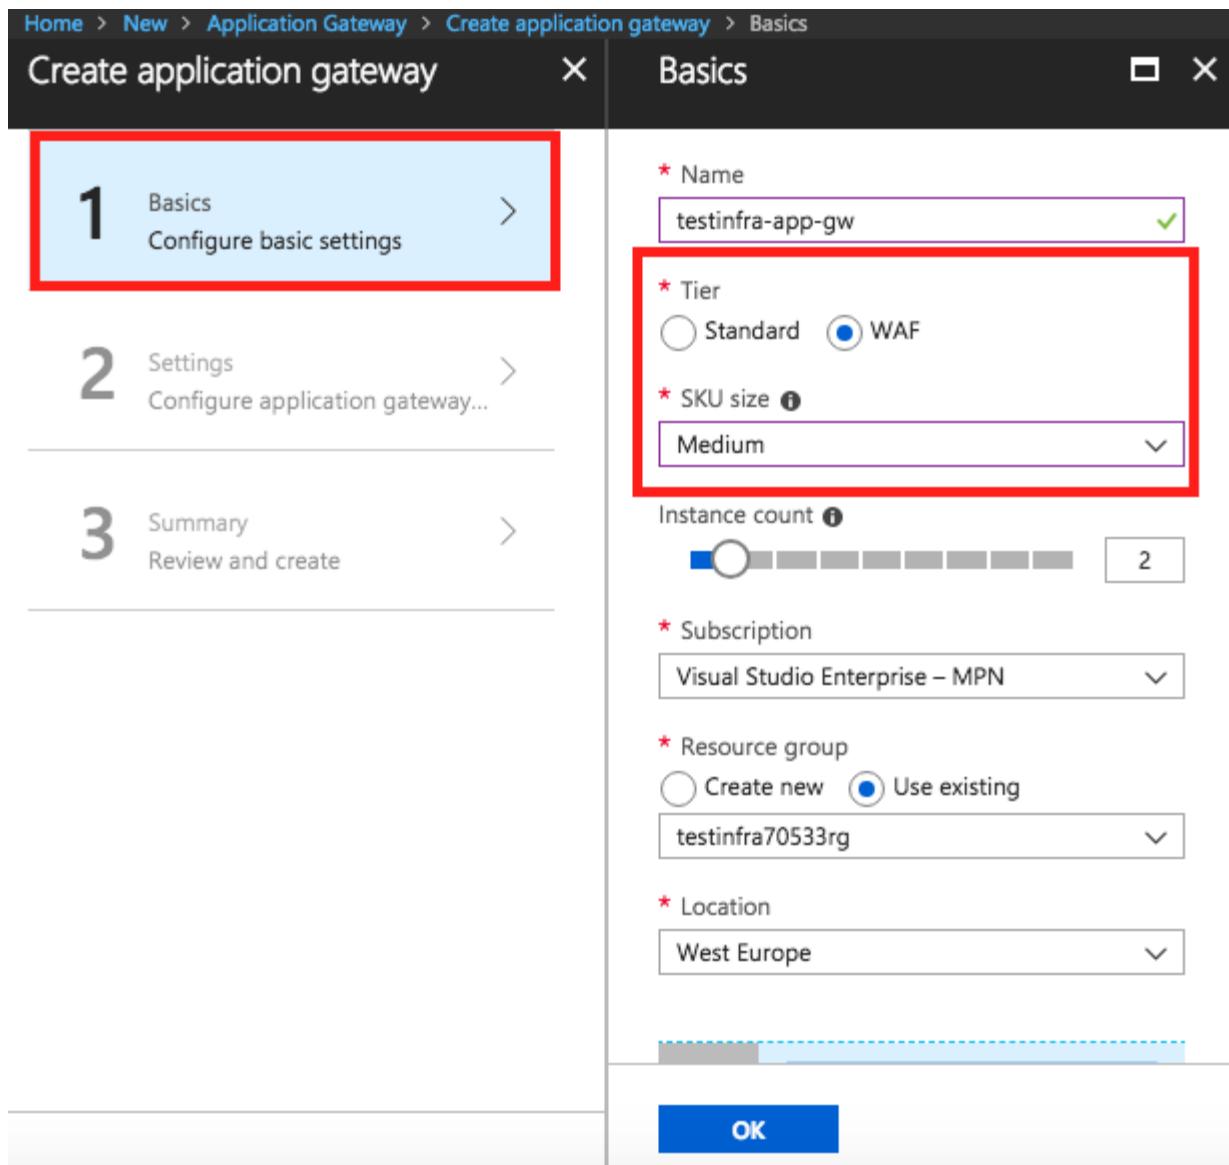
You can see how the Application Gateway works in the following schema. It provides **URL Path-Based Routing**, which allows us to route traffic to the backend server pools based on the URL paths of the request.



Application Gateway with URL Path-Based Routing and WAF feature

To create an **Application Gateway** via Azure Portal, you can click on **Create a resource** and find **Application Gateway** in the Networking category. After clicking on **Create**, you'll see a form, as shown in the following screenshot:

QUALITY THOUGHT
 The Leader in Software Training
 9963799240 / 7730997544
 Ameerpet / Kondapur
Hyderabad



Home > New > Application Gateway > Create application gateway > Basics

Create application gateway X

Basics X

1 Basics >
Configure basic settings

2 Settings >
Configure application gateway...

3 Summary >
Review and create

* Name
testinfra-app-gw ✓

* Tier
 Standard WAF

* SKU size ⓘ
Medium

Instance count ⓘ
2

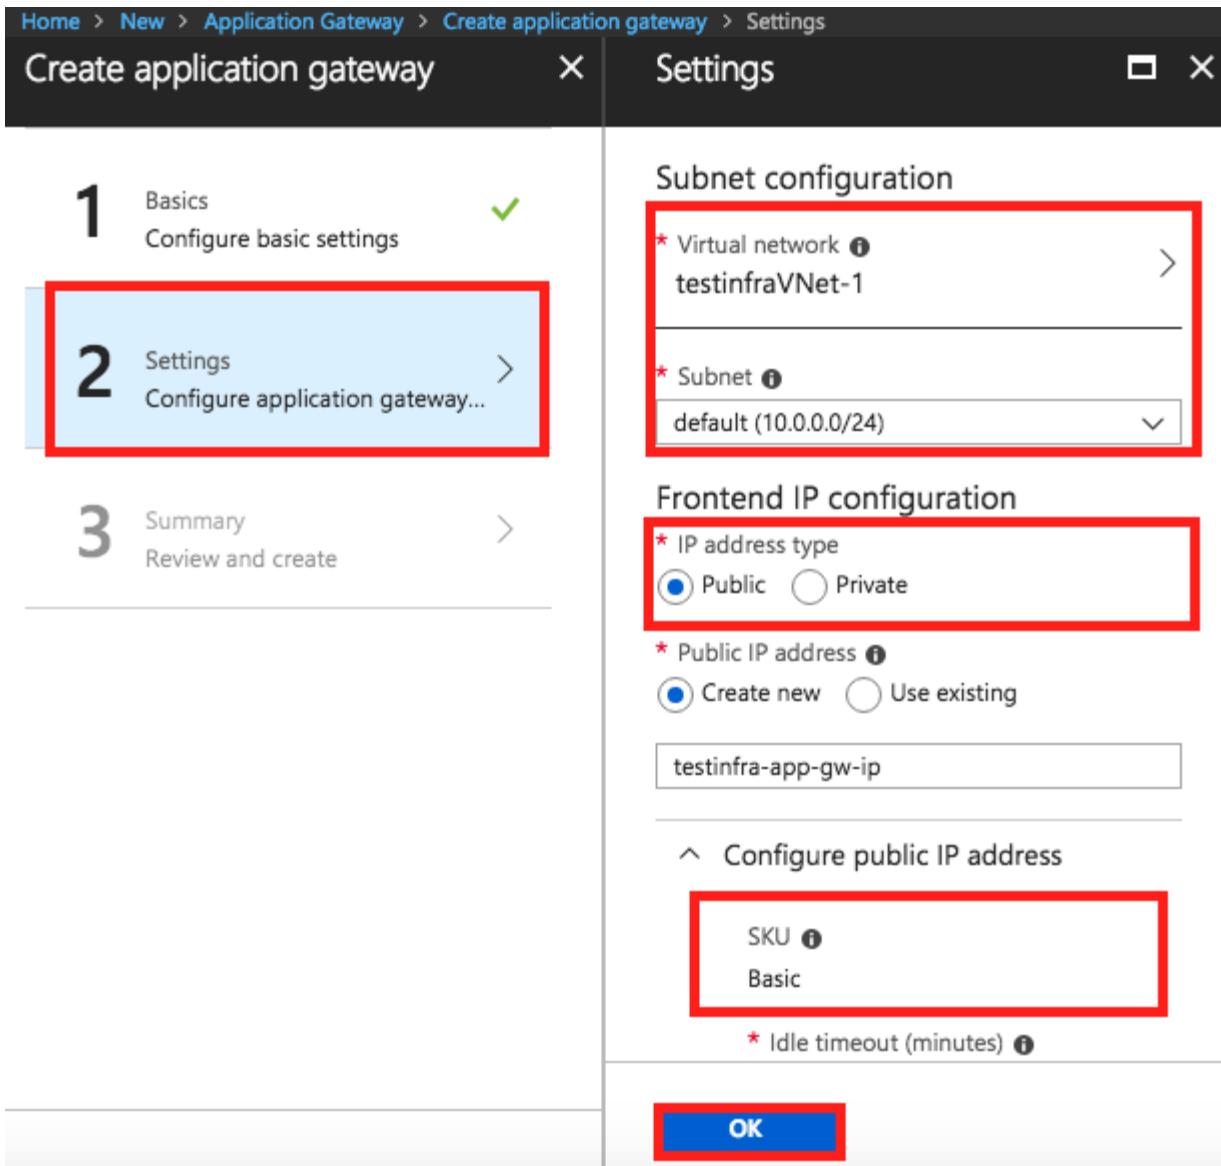
* Subscription
Visual Studio Enterprise – MPN

* Resource group
 Create new Use existing
testinfra70533rg

* Location
West Europe

OK

Choose the tier with **WAF** and **Medium** size. Note that the existing virtual network you'll choose in the next step and the public IP address must be in the same location as your Application Gateway. You can click on **OK** if everything looks good and go to the 2nd step to configure the Application Gateway. You'll then see the following screenshot:



The screenshot shows the 'Create application gateway' wizard in the Azure portal. The current step is 'Settings'. The process consists of three steps:

- 1 Basics: Configure basic settings (Completed)
- 2 Settings: Configure application gateway... (Currently selected)
- 3 Summary: Review and create

Subnet configuration:

- * Virtual network: testinfraVNet-1
- * Subnet: default (10.0.0.0/24)

Frontend IP configuration:

- * IP address type: Public (selected)
- * Public IP address: Create new (selected)
- SKU: Basic
- * Idle timeout (minutes): (not specified)

OK button is visible at the bottom.

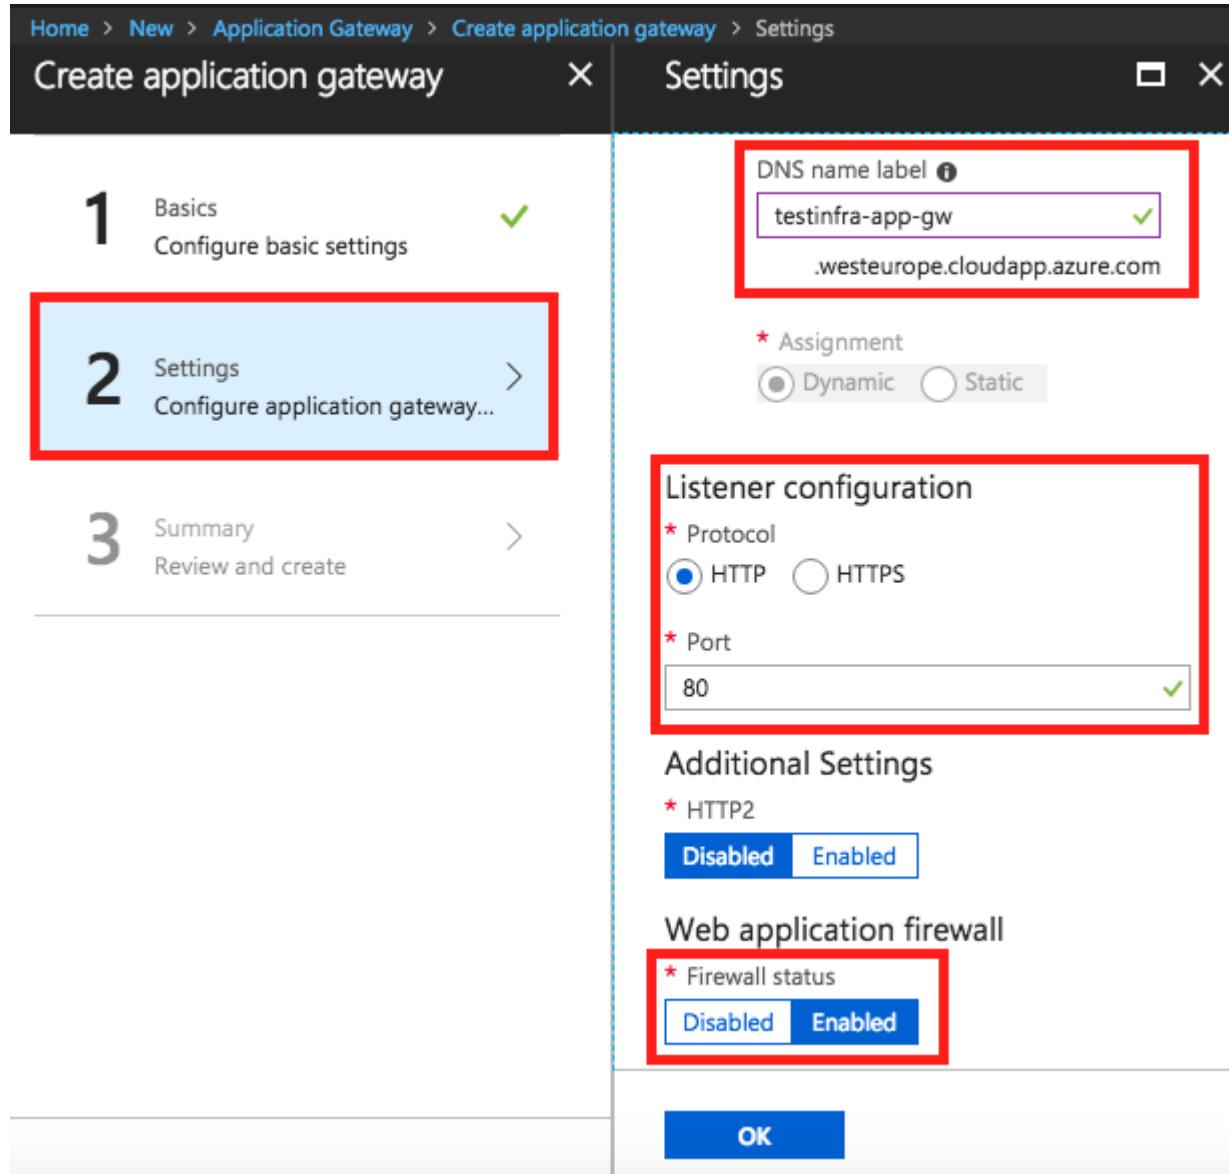
In this step, choose an existing virtual network or create a new virtual network which is in the same location as the Application Gateway. Then choose your frontend IP type. We'll choose Public IP in our case, since the traffic is coming from the internet. As indicated in the preceding screenshot, the SKU for public IP addresses will be defined as BASIC since only the basic tier can be used with an Application Gateway.

There are some other interesting ones in the same step. You should enter a DNS name label for your Application Gateway, which is actually an A record for your public IP address that will

be registered with Azure-provided DNS servers. In our case, the FQDN of our Application Gateway will look like this:

testinfra-app-gw.westeurope.cloudapp.azure.com.

In this step, you should also choose the protocol of your Application Gateway listener, **HTTP** or **HTTPS**. Make sure that you have **Enabled** the WAF, as shown in the following screenshot:



The screenshot shows the 'Create application gateway' settings page. On the left, a navigation bar lists three steps: 1. Basics (Completed), 2. Settings (selected and highlighted with a red box), and 3. Summary. The 'Settings' tab is open, displaying configuration options for a listener.

Listener configuration:

- Protocol:** HTTP (radio button selected)
- Port:** 80

Additional Settings:

- HTTP2:** Enabled (button highlighted with a blue border)

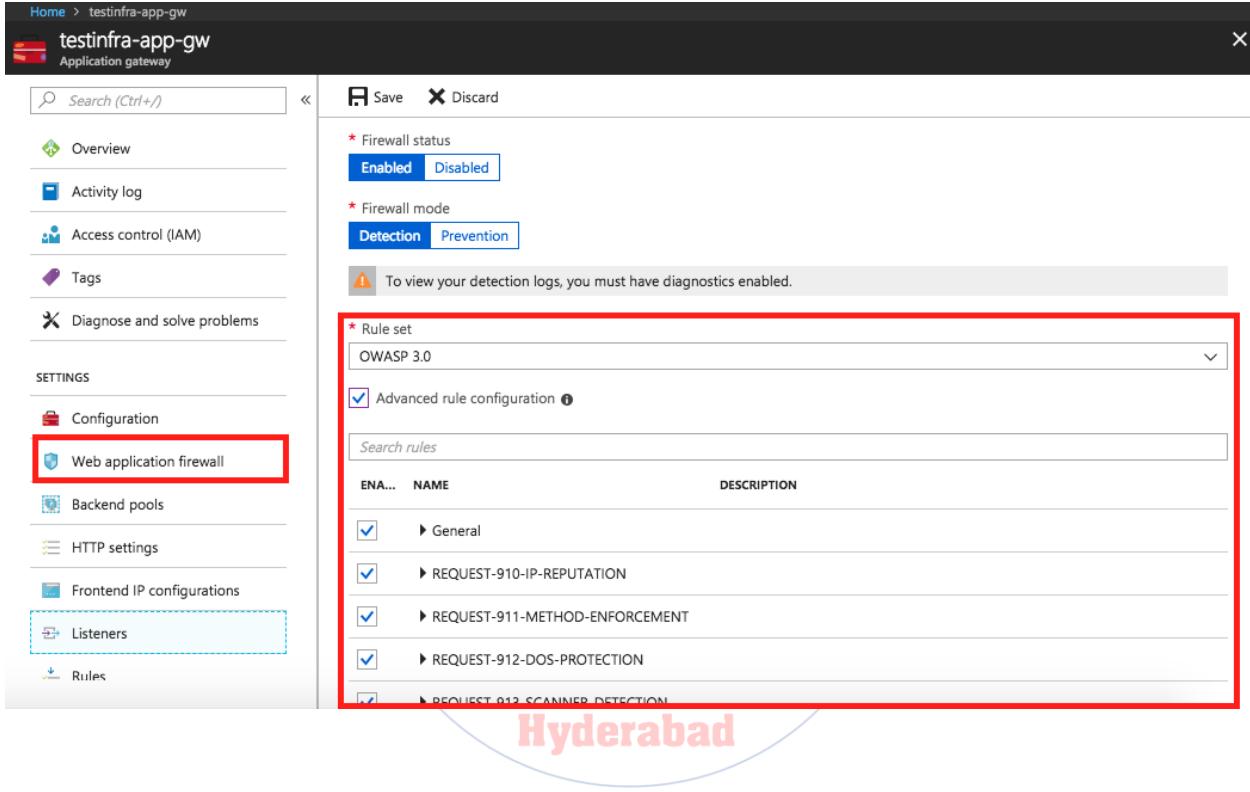
Web application firewall:

- Firewall status:** Enabled (button highlighted with a blue border)

A large red box highlights the 'DNS name label' input field and its value, 'testinfra-app-gw.westeurope.cloudapp.azure.com'. Another red box highlights the 'Assignment' section, specifically the 'Dynamic' radio button.

Finally, you'll have a summary page with everything you have entered into the form. Then you can click on OK. The deployment of the Application Gateway will take a couple of minutes.

After creating the Application Gateway successfully, you can go to the **Web application firewall** blade to choose the rule set that you want to use or do some other advanced rule configuration:



Home > testinfra-app-gw

testinfra-app-gw

Application gateway

Search (Ctrl+/)

Save Discard

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

SETTINGS

Configuration

Web application firewall

Backend pools

HTTP settings

Frontend IP configurations

Listeners

Rules

* Firewall status
Enabled Disabled

* Firewall mode
Detection Prevention

To view your detection logs, you must have diagnostics enabled.

* Rule set
OWASP 3.0

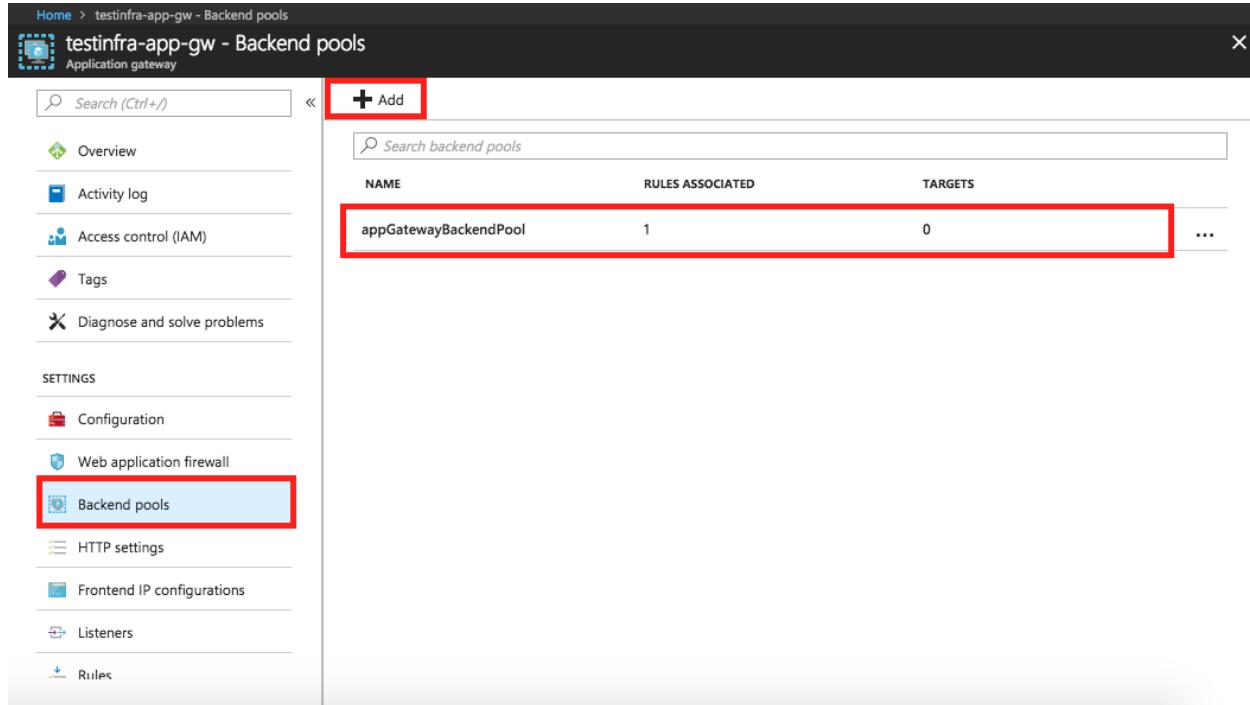
Advanced rule configuration

Search rules

ENA...	NAME	DESCRIPTION
<input checked="" type="checkbox"/>	▶ General	
<input checked="" type="checkbox"/>	▶ REQUEST-910-IP-REPUTATION	
<input checked="" type="checkbox"/>	▶ REQUEST-911-METHOD-ENFORCEMENT	
<input checked="" type="checkbox"/>	▶ REQUEST-912-DOS-PROTECTION	
<input checked="" type="checkbox"/>	▶ REQUEST-913-SCANNER-DETECTION	

Hyderabad

You can also go to **Backend pools** blade to add, modify, or delete the Azure VM, VM Scale Sets, IP address, or FQDN in the backend configuration, as shown in the following screenshot:



NAME	RULES ASSOCIATED	TARGETS
appGatewayBackendPool	1	0

To create an Application Gateway via Azure PowerShell, you can refer to the following link: <https://docs.microsoft.com/en-us/azure/application-gateway/tutorial-restrict-web-traffic-powershell>.

9963799240 / 7730997544

Combining Azure load balancing services

Microsoft Azure provides three types of load balancing service to manage network traffic that is distributed. They can be used individually or the methods can be combined depending on your needs. To know more about how to combine Azure Load balancing solution, you can check the following link: <https://docs.microsoft.com/en-us/azure/traffic-manager/traffic-manager-load-balancing-azure>.

Design and implement cross-premise and multisite connectivity

Microsoft Azure provides capabilities to work as an extension in the cloud of on-premise datacenter for organizations by configuring the network connectivity between the on-premise existing environment and Azure virtual networks.

Microsoft Azure provides the following connectivity options to implement cross-premise and multisite connectivity with Azure virtual networks:

- Point-to-site VPN
- Site-to-site VPNs
- VNet-to-VNet
- VNet peering
- ExpressRoute

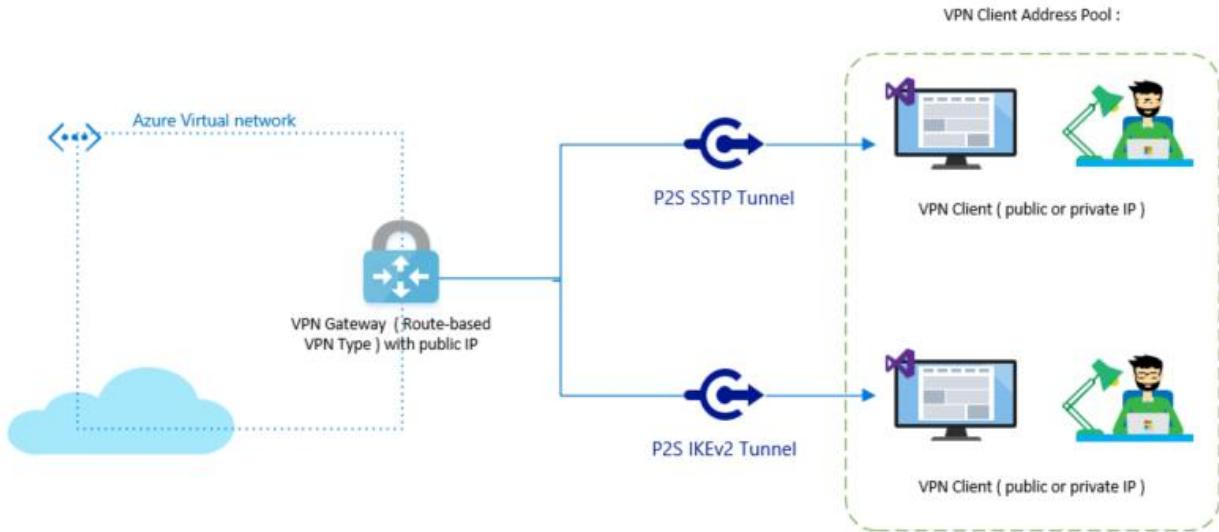
Point-to-site virtual private network (VPN) over IKEv2 or SSTP

Point-to-site virtual private network (VPN) is a solution used to establish a secure connection between an Azure virtual network and a single client computer. The general use case for P2S is to establish connections between individual computers to connect to Azure Virtual Network from a remote location, such as some offsite employer who wants to work from home.

There are two kinds of protocol that can be used to establish a P2S connection:

- **Secure Socket Tunneling Protocol (SSTP):** This is a proprietary SSL-based VPN protocol, which can be used to connect from Windows devices (Windows 7 and above)
- **IKEv2 VPN:** This is a standard-based IPsec VPN solution, which can be used to connect from Mac devices (MAC OS X10.11 and above)

The following schema shows what P2S looks like:



As we can see from the schema, P2S also needs a VPN Gateway, which is a virtual network gateway in Azure with route-based VPNs. Each client compute needs to use self-signed certificates (root and client certificate) before connecting to Azure.

Users can generate a certificate using Azure PowerShell or using make cert. You can see how to generate and export certificate for P2S at the following links.

Using Azure PowerShell: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-certificates-point-to-site>.

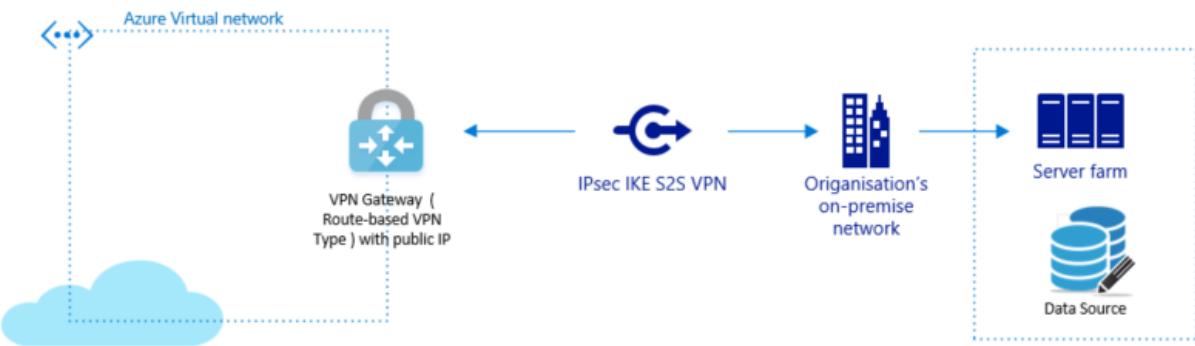
Using make cert: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-certificates-point-to-site-makecert>.

For more information on how to implement P2S via Azure Portal, click on the following link: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-howto-point-to-site-resource-manager-portal>.

You can also implement P2S with Azure PowerShell as indicated in the following link: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-howto-point-to-site-rm-ps>.

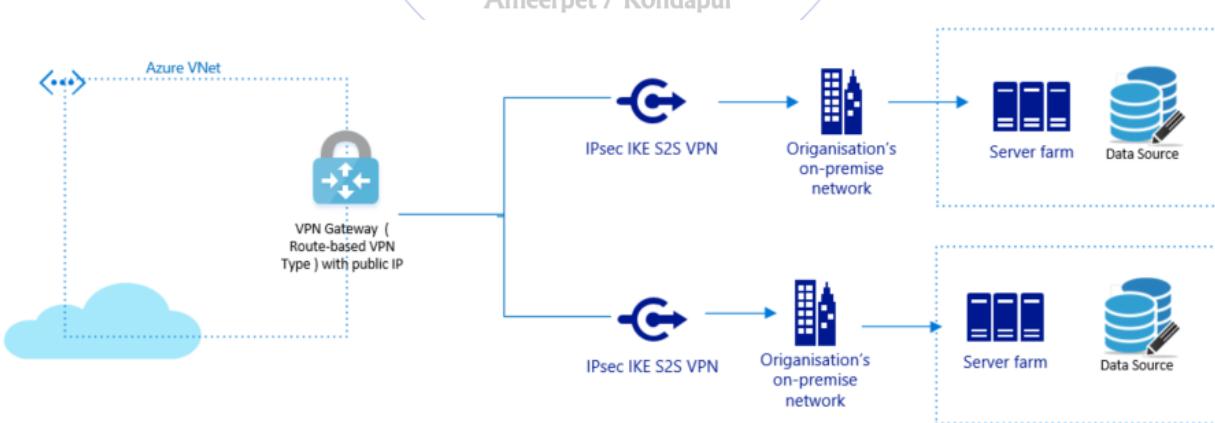
Site-to-site and multisite virtual private network (IPsec/IKE VPN tunnel)

A Site-to-Site VPN gateway connection is a solution used to connect on-premise to an Azure virtual network. This connection is over an IPsec/IKE (IKEv1 or IKEv2) VPN tunnel. To establish this connection, you should have a VPN gateway located on-premise with an externally facing public IP address. This type of connection is shown in the following schema:



Site-to-Site VPN in Azure

While creating more than one VPN connection in Route-Based VPN type so that users can connect to multiple on-premise sites, we called it multisite VPN this is a variation on the S2S connection. This type of connection is shown in the following schema:



Multisite VPN in Azure

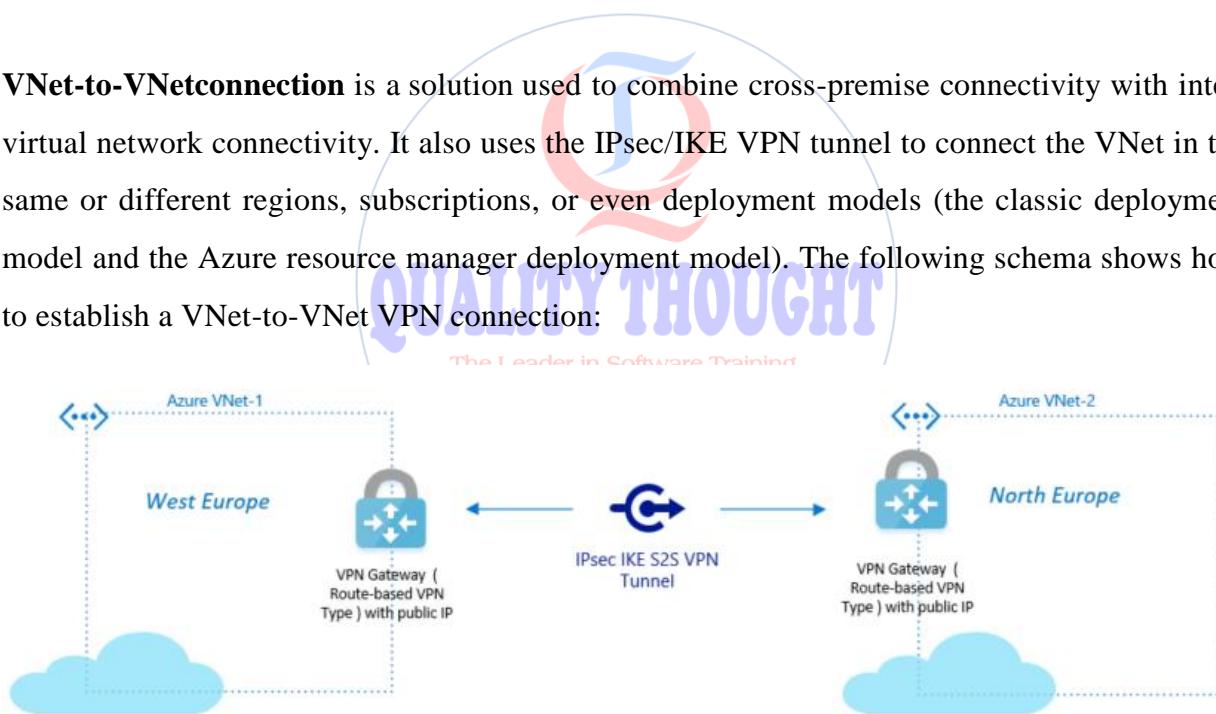
For more information on how to implement S2S via the Azure Portal, click on the following link: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-howto-site-to-site-resource-manager-portal>.

You can also do the same using Azure PowerShell: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-create-site-to-site-rm-powershell>.

Another great way is to use Azure CLI: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-howto-site-to-site-resource-manager-cli>.

VNet-to-VNet virtual private network (IPsec/IKE VPN tunnel)

VNet-to-VNetconnection is a solution used to combine cross-premise connectivity with inter-virtual network connectivity. It also uses the IPsec/IKE VPN tunnel to connect the VNet in the same or different regions, subscriptions, or even deployment models (the classic deployment model and the Azure resource manager deployment model). The following schema shows how to establish a VNet-to-VNet VPN connection:



VNet-to-VNet connection in Azure

For more information on how to implement VNet-to-VNet via the Azure Portal, click on the following link: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-howto-vnet-vnet-resource-manager-portal>.

You can also do this using Azure PowerShell: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-vnet-rm-ps>.

Another great way is to use Azure CLI: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-howto-vnet-vnet-cli>.

To know how to connect classic VNets to Resource Manager VNets to allow resources located in the separate deployment models via Azure Portal, you can check the following link: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-connect-different-deployment-models-portal>.

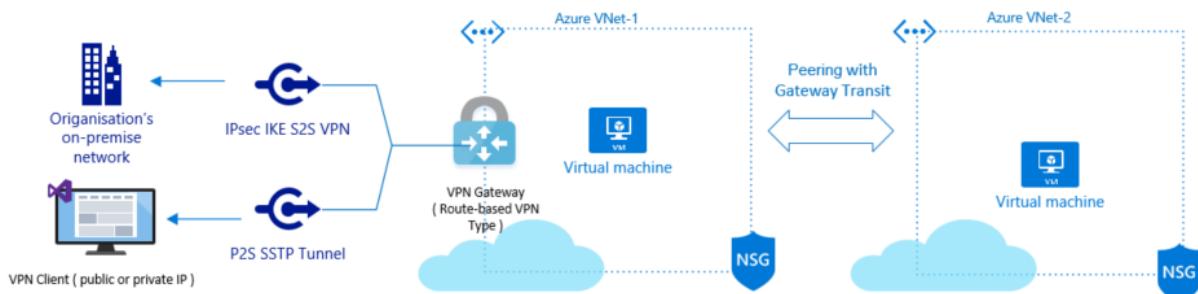
It is also possible to use Azure PowerShell to do that. You can go to the following link for more information: <https://docs.microsoft.com/en-us/azure/vpn-gateway/vpn-gateway-connect-different-deployment-models-powershell>.

Virtual network peering (VNet Peering)

Virtual network peering (VNet Peering) is an option provided by Microsoft Azure to connect two Azure virtual networks without using VPN gateways. There are two types of virtual network peering:

- **VNet peering:** This is a way to connect VNets within the same Azure region
- **Global VNet peering:** This is a way to connect VNets across different Azure regions

An example of VNet peering is described in the following schema:



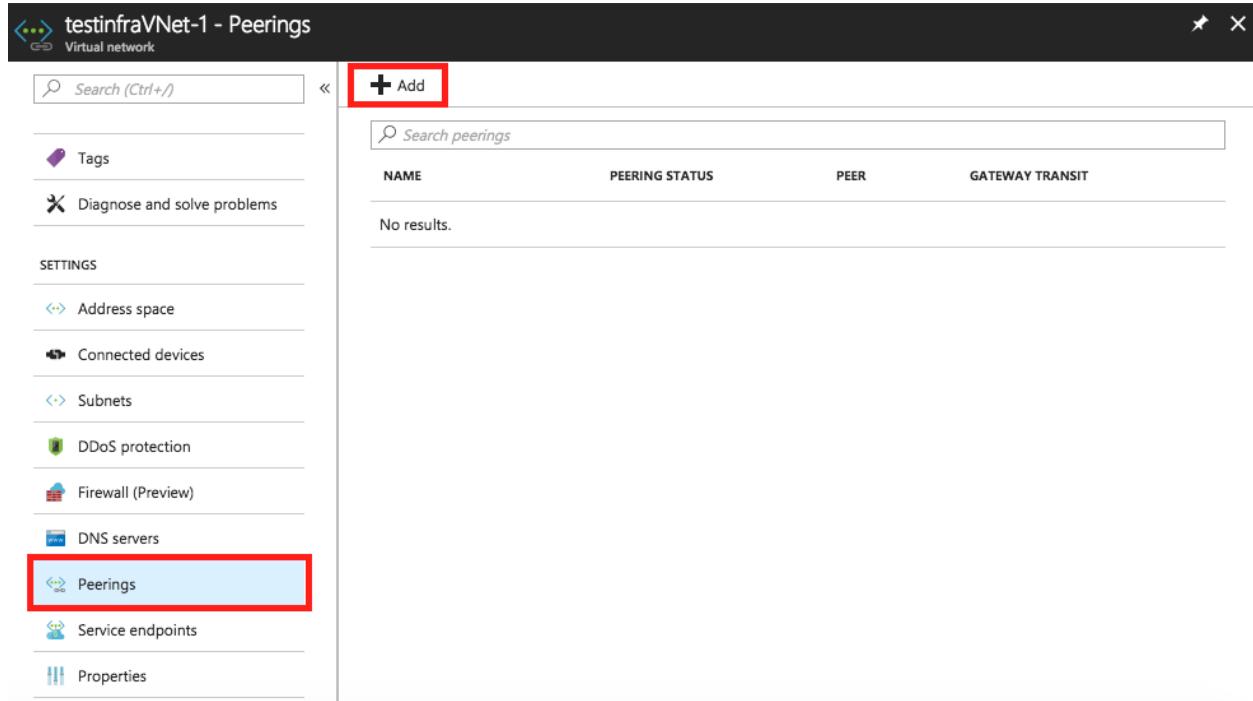
As we can see from this schema, **VNet-1** and **VNet-2** have been connected by a VNet peering. **VNet-1** has a **VPN Gateway** to connect with another organization using S2S VPN and another client compute using **P2SVPN**. The **VPN Gateway** has enabled the Gateway transit peering property, which enables **VNet-2** to use the **VPN Gateway** in the peered virtual network, which is **VNet-1** for cross-premise or VNet-to-VNet connectivity. Connectivity applies to both **VNet -1** and **VNet-2**.

The greatest advantage of VNet peering is that Network traffic between peered virtual networks is private and routed through the Microsoft backbone infrastructure. The bandwidth and latency across the VNets are the same in the same region as if the resources were connected to the same VNet. When creating the peering, there is no downtime to resources in either virtual network. The peering is also possible to happen to peer one virtual network created through Azure Resource Manager to a virtual network created through the classic deployment model that exists in the same or different subscriptions.

Creating VNet peering via Azure Portal

Now, let's implement a peering connection between VNet-1 and VNet-2. To add a VNet peering via Azure Portal, you can go to the **Peerings** blade of VNet-1, and then click on **Add** to add a new peering:

Ameerpet / Kondapur
Hyderabad



The screenshot shows the Azure portal interface for managing peerings in a virtual network. The main area displays a table with columns: NAME, PEERING STATUS, PEER, and GATEWAY TRANSIT. A search bar at the top says 'Search peerings'. The left sidebar has a 'Tags' section and several 'SETTINGS' options: Address space, Connected devices, Subnets, DDoS protection, Firewall (Preview), DNS servers, and Peerings. The 'Peerings' option is highlighted with a red box. At the top right, there are 'Add' and 'X' buttons.

After clicking on Add, the Add peering page will be displayed. In this creation form, we should create the peering between VNet-1 to VNet-2, and select Enable gateway transit in the configuration section, as shown in the following screenshot. Make sure that we already have a virtual network gateway was in the gateway Subnet of the current VNet so that we can choose to use the remote gateway in the VNet-2 when we create the peering between VNet-2 to VNet-1:

Home > Resource groups > testinfra70533rg > testinfraVNet-1 - Peerings > Add peering

Add peering

testinfraVNet-1

* Name
peering-vnet1tovnet2 ✓

Peer details

Virtual network deployment model i
 Resource manager Classic

I know my resource ID i

* Subscription i
Visual Studio Enterprise – MPN

* Virtual network
testinfravnet-2 (testinfra70533rg)

Configuration

Allow virtual network access i

Allow forwarded traffic i

Allow gateway transit i

Use remote gateways i

i Virtual network 'testinfraVNet-1' has a gateway; peerings created from this virtual network can't enable 'use remote gateways'.

As explained in the previous step, we can go to the Peering blade for VNet - 2 and add a peering in the same way, as shown in the following screenshot. Enabling **Allow forwarded traffic** allows the traffic go to other peered VNets in a transitive way:

Home > Resource groups > testinfra70533rg > testinfravnet-2 - Peerings > Add peering

Add peering

testinfravnet-2

*** Name**
peerings-vnet2tovnet1

Peer details

Virtual network deployment model [i](#)
 Resource manager Classic

I know my resource ID [i](#)

*** Subscription** [i](#)
Visual Studio Enterprise – MPN

*** Virtual network**
testinfraVNet-1 (testinfra70533rg)

Configuration

Allow virtual network access [i](#)

Allow forwarded traffic [i](#)

Allow gateway transit [i](#)

Use remote gateways [i](#)

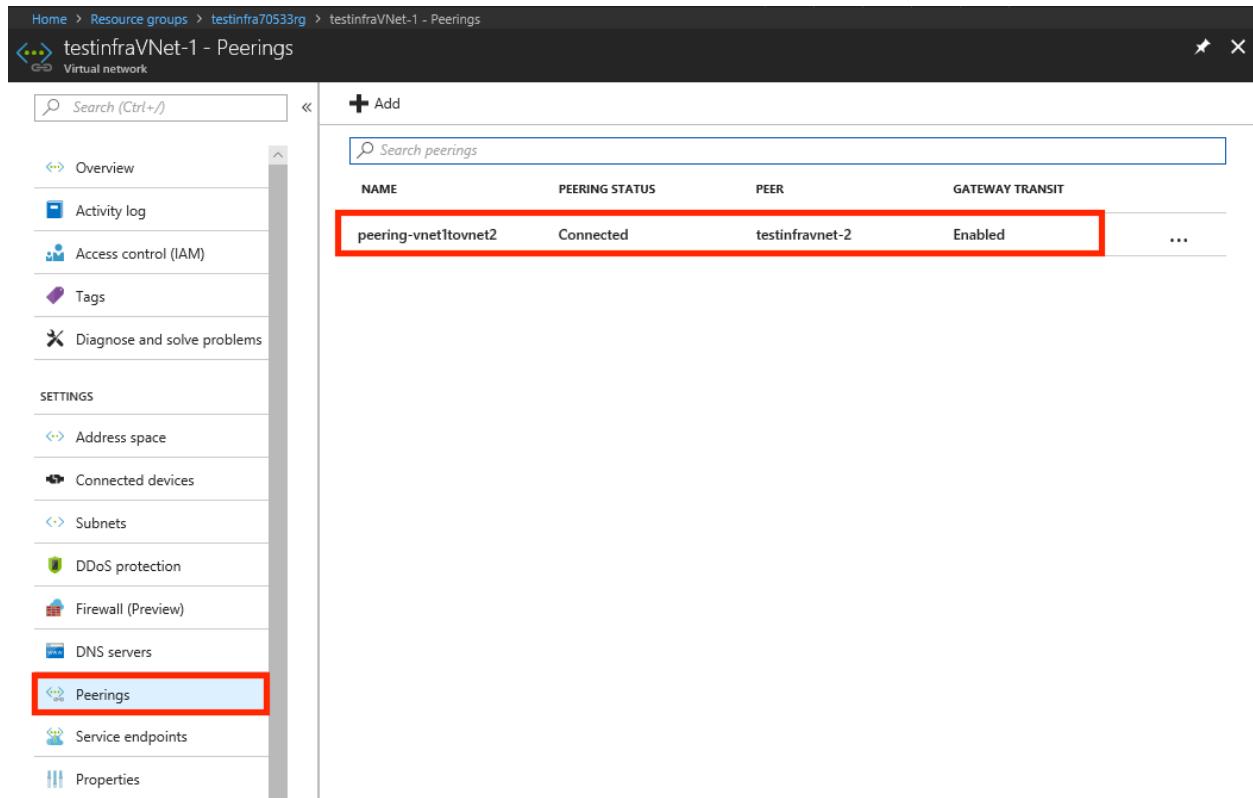
i To use a remote gateway, the peer virtual network must use the resource manager deployment model, have a gateway, and enable 'allow gateway transit'. Ensure this setting is enabled on your peer before selecting this setting.

OK

Verify the peering connection

Finally, after clicking on OK and waiting for two peering connections to be created, you can go to the Peering blade of each VNet to verify whether two VNets have been connected

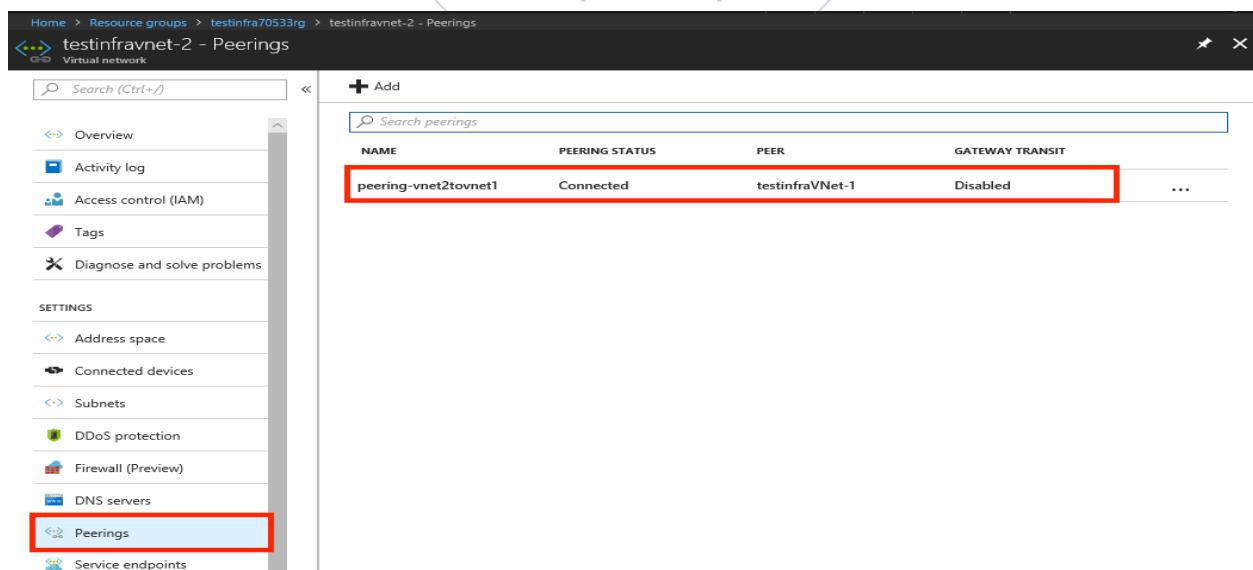
successfully. If everything is going well, you'll be able to see the peering status of VNet- 1, as shown in the following screenshot:



The screenshot shows the Azure portal interface for a virtual network named "testinfraVNet-1". The left sidebar menu is open, with "Peerings" highlighted by a red box. The main content area displays a table of peerings. One entry is highlighted with a red box: "peering-vnet1tovnet2" with "Connected" status, "testinfravnet-2" as the peer, and "Enabled" gateway transit.

NAME	PEERING STATUS	PEER	GATEWAY TRANSIT
peering-vnet1tovnet2	Connected	testinfravnet-2	Enabled

The peering status of VNet-2 is shown in the following screenshot:
The Leader in Software Training
9963799740 / 7730997541
Ameerpet / Kondapur



The screenshot shows the Azure portal interface for a virtual network named "testinfravnet-2". The left sidebar menu is open, with "Peerings" highlighted by a red box. The main content area displays a table of peerings. One entry is highlighted with a red box: "peering-vnet2tovnet1" with "Connected" status, "testinfraVNet-1" as the peer, and "Disabled" gateway transit.

NAME	PEERING STATUS	PEER	GATEWAY TRANSIT
peering-vnet2tovnet1	Connected	testinfraVNet-1	Disabled

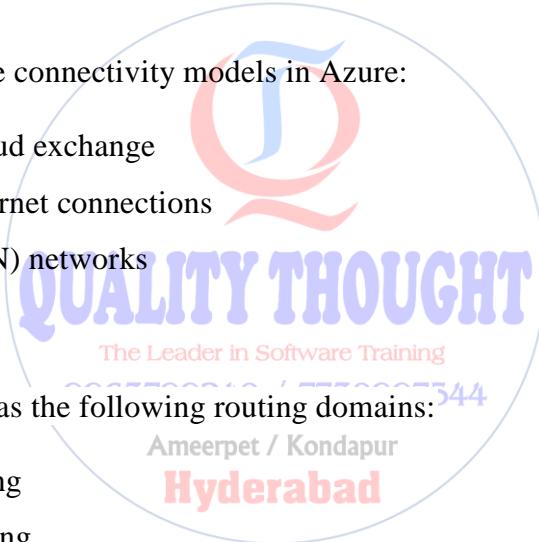
ExpressRoute

ExpressRoute is an Azure service that lets you create private dedicated connections that do not go over the public internet between Microsoft Clouds such as **Microsoft Azure**, **Microsoft Office 365**, **Microsoft Dynamics 365**, and the organization's IT environment.

To connect an on-premise infrastructure to Microsoft Cloud, users should order an ExpressRoute circuit through a connectivity provider. In Azure, a single ExpressRoute circuit represents a logical connection; it is also possible to order multiple ExpressRoute circuits and each connection can be in the same or different regions, and can be connected to your premises through different connectivity providers.

The following are the three connectivity models in Azure:

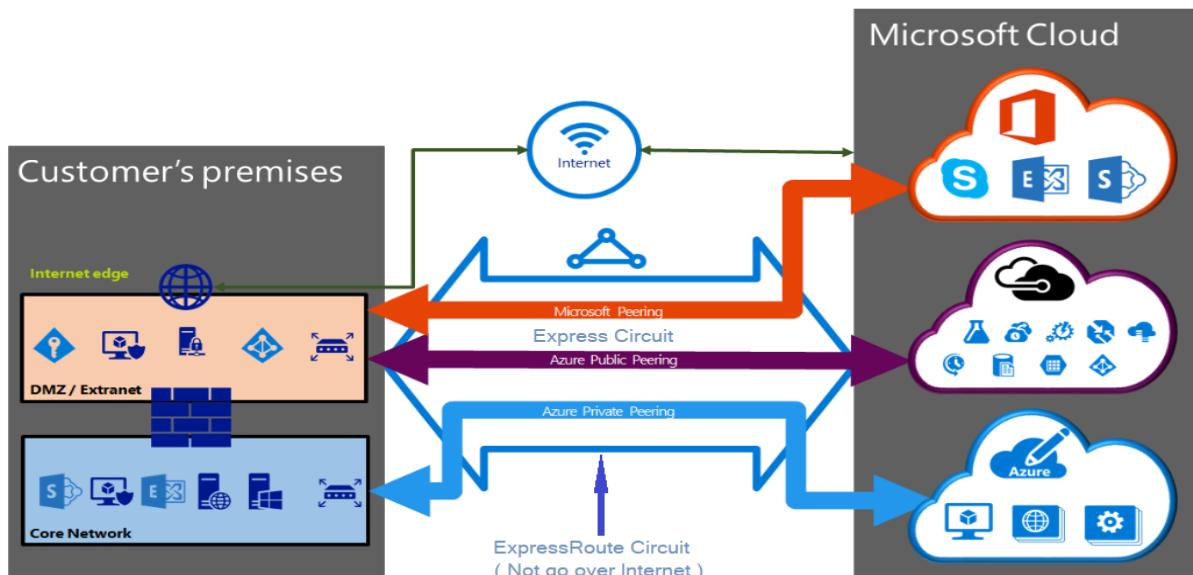
- Co-located at a cloud exchange
- Point-to-point Ethernet connections
- Any-to-any (IPVPN) networks



An ExpressRoute circuit has the following routing domains:

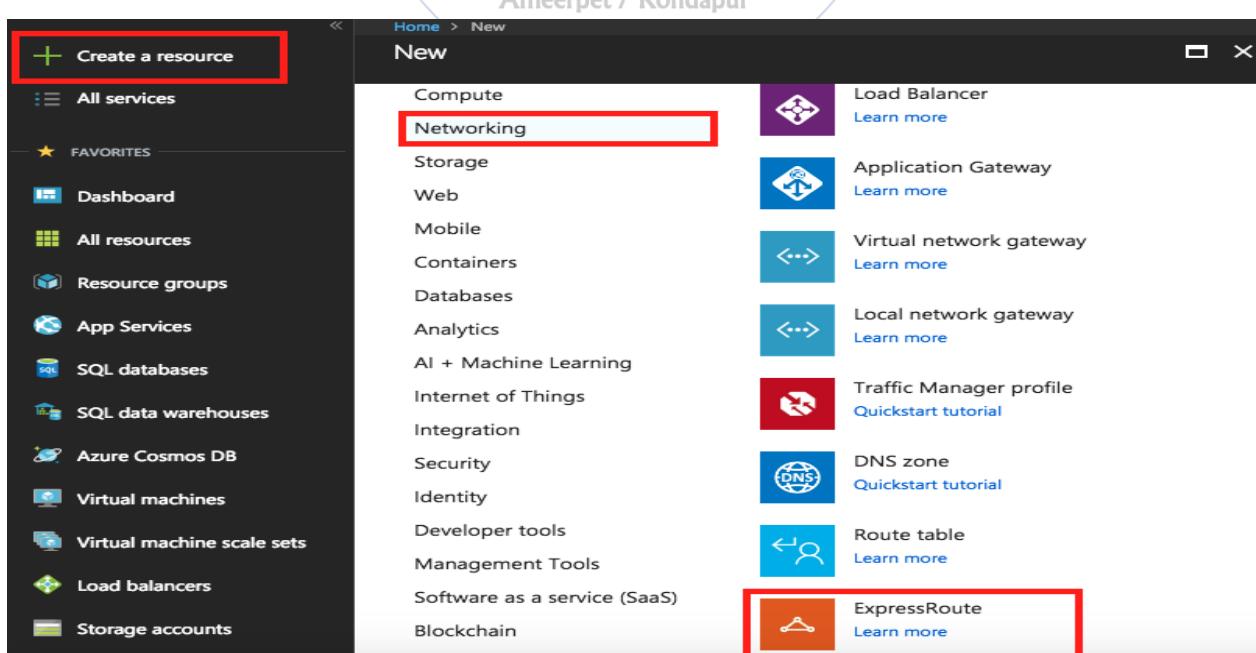
- Azure public peering
- Azure private peering
- Microsoft peering

The following schema shows the connectivity between on-premise and Microsoft Cloud through multiple routing domains:



ExpressRoute circuit connecting on-premise and Microsoft Cloud
[\(https://github.com/MicrosoftDocs/azure-docs/blob/master/articles/expressroute/expressroute-circuit-peerings.md\)](https://github.com/MicrosoftDocs/azure-docs/blob/master/articles/expressroute/expressroute-circuit-peerings.md)

To create an **ExpressRoute** circuit, you can go to the Azure Portal and click on **Create a resource**; you'll find **ExpressRoute** in the **Networking** category, as shown in the following screenshot:



After clicking on ExpressRoute, you'll see the **Create ExpressRoute circuit** page where there is a creation form. You should choose a connectivity provider and a peering location, which is the physical location from which you are peering with Microsoft. You can also choose the available bandwidth based on your previous selection. If everything is okay, you can click on **Create**; the deployment will last for a couple minutes:

Home > New > Create ExpressRoute circuit

Create ExpressRoute circuit

Create new or import from classic ⓘ

Create new **Import**

* Circuit name
testinfra-expressroute

* Provider ⓘ
AT&T

* Peering location ⓘ
Silicon Valley

* Bandwidth ⓘ
50Mbps

* SKU ⓘ
Standard **Premium**

* Billing model ⓘ
Unlimited **Metered**

Allow classic operations ⓘ

* Subscription
Visual Studio Enterprise – MPN

* Resource group
 Create new Use existing
testinfra70533rg

* Location
West Europe

Create **Automation options**

By clicking the create button, you understand that billing will start immediately upon creation of the ExpressRoute and you agree to accept the charges.

After creating an ExpressRoute circuit successfully, the next step is to create the routing configuration, and finally to link a VNet to an ExpressRoute circuit. For more information

on how to do that, you can check the following links to know different ways to achieve your objective:

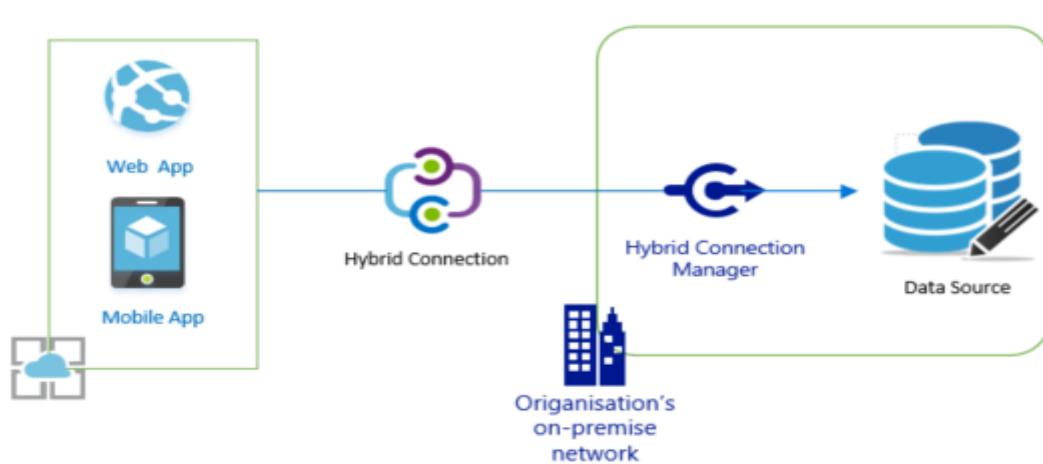
- Via Azure Portal: <https://docs.microsoft.com/en-us/azure/expressroute/expressroute-howto-routing-portal-resource-manager>
- Via Azure PowerShell: <https://docs.microsoft.com/en-us/azure/expressroute/expressroute-howto-routing-arm>
- Via Azure CLI: <https://docs.microsoft.com/en-us/azure/expressroute/howto-routing-cli>

For more information on how to link a VNet to a ExpressRoute circuit, you can check the different ways to do that at the following links:

- Via Azure Portal: <https://docs.microsoft.com/en-us/azure/expressroute/expressroute-howto-linkvnet-portal-resource-manager>
- Via Azure PowerShell: <https://docs.microsoft.com/en-us/azure/expressroute/expressroute-howto-linkvnet-arm>
- Via Azure CLI: <https://docs.microsoft.com/en-us/azure/expressroute/howto-linkvnet-cli>

Configuring Hybrid Connections for App Service

Hybrid Connections is a capability provided by Azure App Service to let web apps and mobile apps in App Service access on-premise systems and services securely. The following schema shows how Hybrid Connections works in Azure: [7730997544](#)



To create a hybrid connection, you can go to the Azure Portal and select your web app. Then, go to the **Networking** blade and click on **Configure your Hybrid Connection endpoints**; you will be provided with a screen as shown in the following screenshot:

Home > Resource groups > testinfra-webapp > testinfra-webapp - Networking > Hybrid connections

Hybrid connections
testinfra-webapp

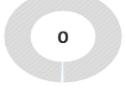
Refresh

 **Hybrid connections**

App Service integration with hybrid connections enables your app to access a single TCP endpoint per hybrid connection. Here you can manage the new and classic hybrid connections used by your app. [Learn more](#)

App service plan (pricing tier) : Connections used
testinfraAdvancedPlan (PremiumV2)

Location : West Europe



CONNECTIONS USED
0
CONNECTIONS QUOTA
220

 Download connection manager

 **Add hybrid connection**

NAME	STATUS	ENDPOINT	NAMESPACE
------	--------	----------	-----------

Click on **Add hybrid connection**; you should fill in the endpoint and endpoint port. Each hybrid connection is attached to a service bus namespace in one Azure region. Microsoft recommends you use an existing service bus namespace in the same region, or create a new service bus namespace as the target web app, to reduce network latency, as shown in the following screenshot:

Home > Resource groups > testinfra-webapp > testinfra-webapp - Networking > Hybrid connections > Add hybrid connection

Add hybrid connection
testinfra-webapp

 Create new hybrid connection  Add selected hybrid connection  Refresh

 **Add hybrid connection**

To use a hybrid connection with your app select it from the list and click on 'Add selected hybrid connection'. Click on 'Create new' to create a new one.

NAME	HOST	PORT	NAMESPACE
No results			

* Hybrid connection Name  **testinfrahybridhost**

* Endpoint Host  **testinfrahybridhost.com**

* Endpoint Port  **1433**

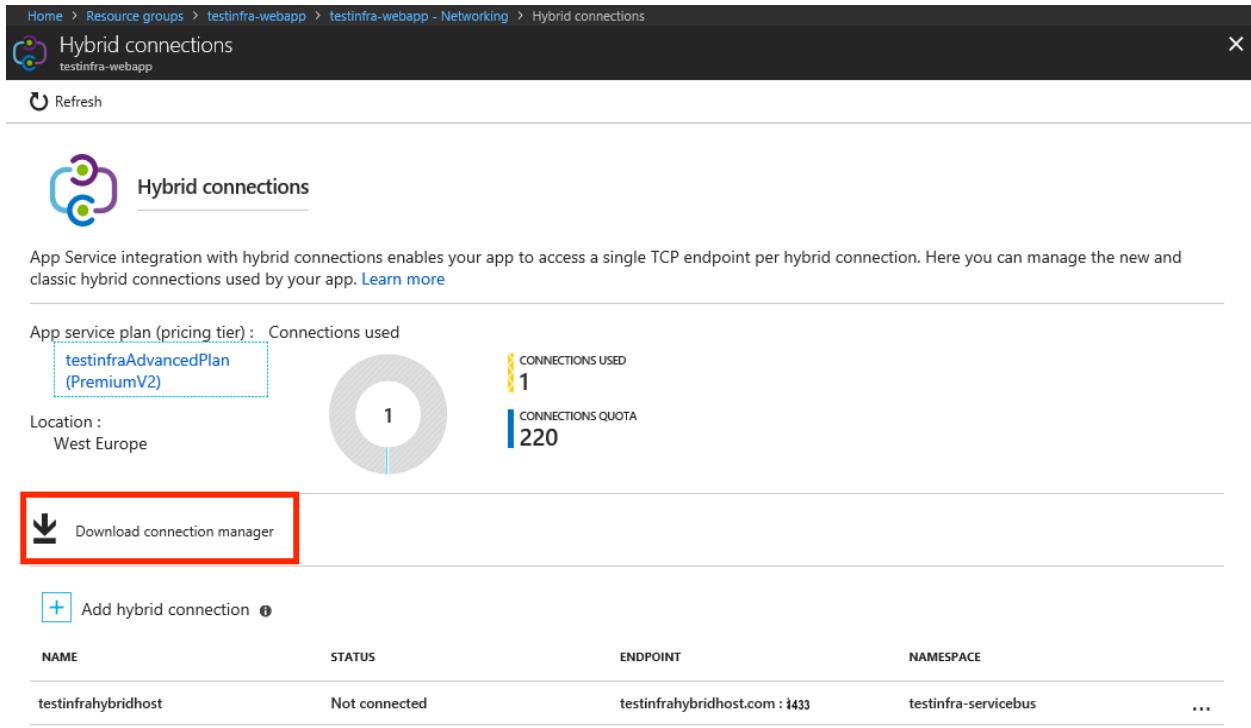
* Servicebus namespace 
 Create new Select existing

* Location
 West Europe

* Name
testinfra-servicebus

OK

After a hybrid connection has been created successfully, you can go to the **Networking** blade and then click on Hybrid Connections. You should download the Hybrid connection Manager and install it on the on-premise resource so that the Hybrid Connections are established:



Home > Resource groups > testinfra-webapp > testinfra-webapp - Networking > Hybrid connections

Hybrid connections
testinfra-webapp

Refresh

Hybrid connections

App Service integration with hybrid connections enables your app to access a single TCP endpoint per hybrid connection. Here you can manage the new and classic hybrid connections used by your app. [Learn more](#)

App service plan (pricing tier) : Connections used
testinfraAdvancedPlan (PremiumV2)

Location : West Europe

CONNECTIONS USED
1

CONNECTIONS QUOTA
220

Download connection manager

Add hybrid connection

NAME	STATUS	ENDPOINT	NAMESPACE	...
testinfrahybridhost	Not connected	testinfrahybridhost.com :443	testinfra-servicebus	...

Configuring multi-region applications with Azure Traffic Manager

9963799240 / 7730997544

Traffic Manager is a DNS-level load-balancing solution that is included within Azure. It uses the following four traffic-routing methods to direct client requests to the most suitable service endpoint:

- **Priority:** This is a method to distribute traffic to the primary location but it will direct the traffic to a secondary location in the case of failure of the primary region.
- **Weighted:** This is a method to distribute traffic across a set of endpoints according to weights. It can be evenly distributed as well depending on the user's configuration.
- **Performance:** This is a method to distribute traffic depending on the location with the lowest network latency.
- **Geographic:** This is a method to distribute traffic depending on which geographic location the DNS query originates from.

In Azure, all the Traffic Manager profiles perform the following two features:

- Monitoring the health of endpoints

- Handling endpoint failover automatically

In the multi-region scenario, the Traffic Manager can be configured with priority methods and it routes incoming requests to the primary region and fails over to the secondary region in the event a failure occurs in the primary region; for example, the application running in the primary region becomes unavailable.

Another setting is about Health probes. Traffic Manager uses an HTTP probe to monitor the availability of configured endpoint linked to the Traffic Manager profile. The main responsibility of a health probe is to check the availability of each region. It can send a request to a specified URL path via the defined protocol (HTTP or HTTPS) and port to check for uptime and determine whether the instances in this region (current endpoint) are healthy by getting a 200 response or unhealthy if they get a non-200 response within a determined period of time (failing these, it will throw a timeout). After several retries, if the requests still fail, the Traffic Manager will consider the current endpoint as a failure and will fail over to the other endpoint.

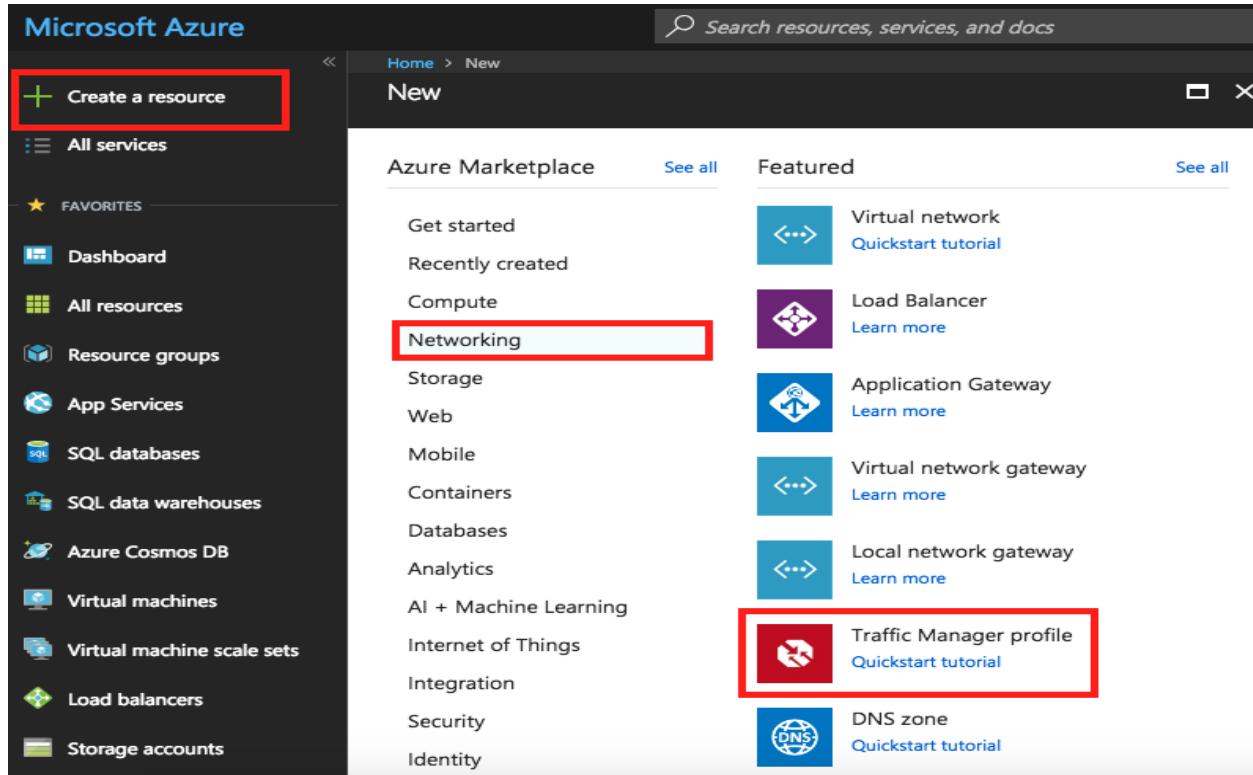
Creating a Traffic Manager profile

2016799240 / 7730997544

Ameerpet / Kondapur

Hyderabad

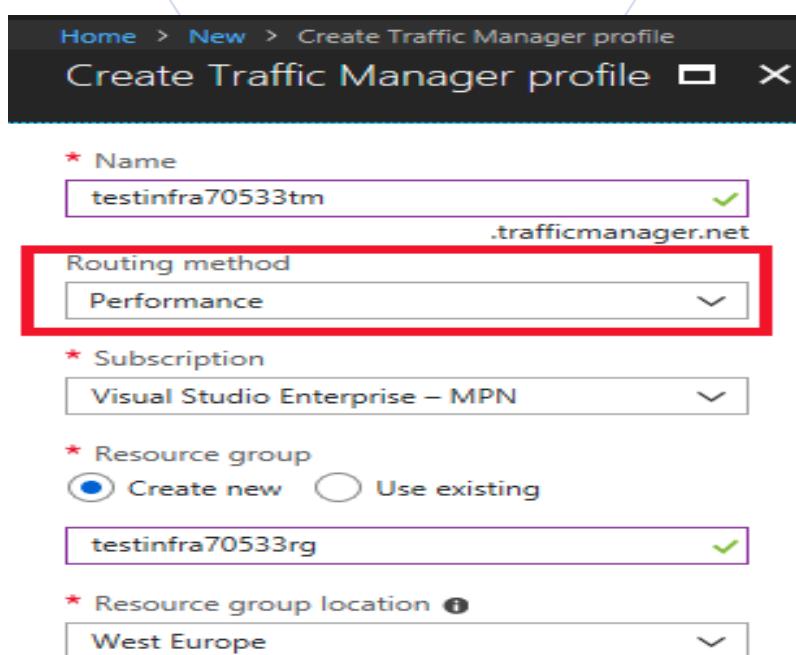
You can create a new Traffic Manager by clicking on **Create a resource** via the Azure Portal. Then, in the **Networking** category, you'll find **Traffic Manager profile**, as shown in the following screenshot:



The screenshot shows the Microsoft Azure portal interface. On the left sidebar, under the 'New' section, the 'Create a resource' button is highlighted with a red box. In the main content area, the 'Networking' category is selected from the 'Azure Marketplace' list, also highlighted with a red box. On the right, a 'Featured' section lists various networking services, with the 'Traffic Manager profile' item highlighted with a red box. The URL in the browser is [https://portal.azure.com/#create/NewResourceGroup/Networking/TrafficManagerProfile](#).

Creating Traffic Manager profile via Azure Portal

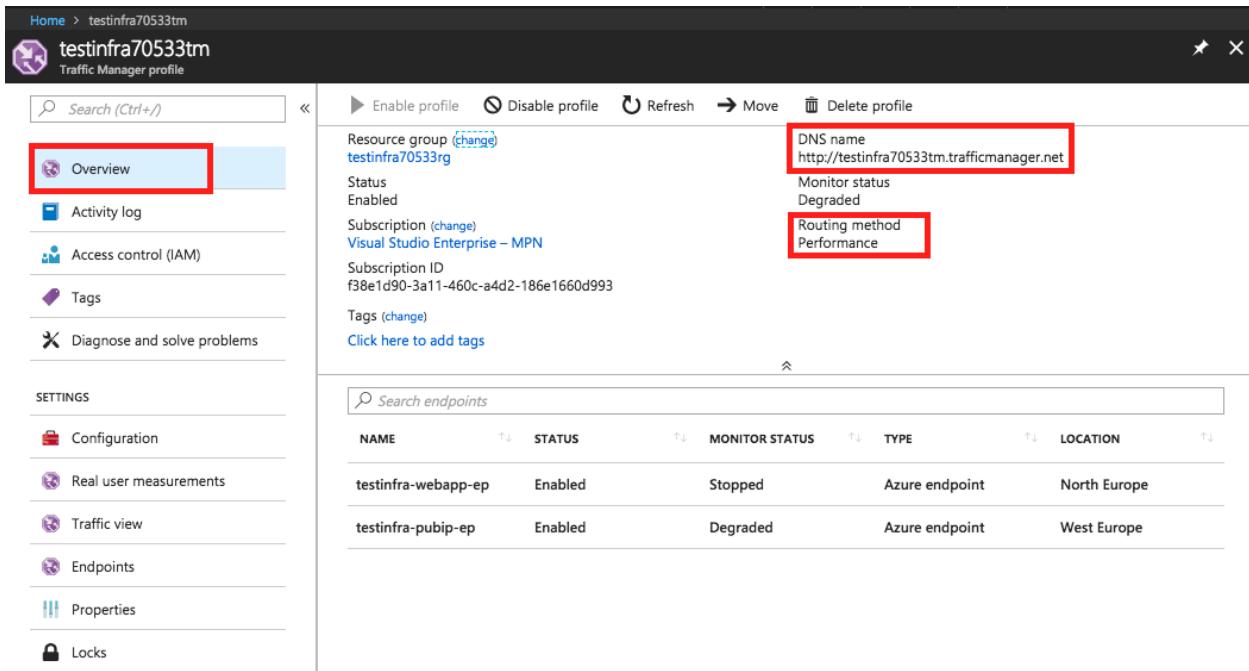
In the basic information form, you can enter the name of the Traffic Manager profile and choose a routing method for the Traffic Manager:



The screenshot shows the 'Create Traffic Manager profile' form. The 'Name' field is set to 'testinfra70533tm'. The 'Routing method' dropdown is set to 'Performance'. The 'Subscription' dropdown is set to 'Visual Studio Enterprise – MPN'. Under 'Resource group', the 'Create new' option is selected, and the new resource group is named 'testinfra70533rg'. The 'Resource group location' is set to 'West Europe'. The URL in the browser is [https://portal.azure.com/#create/NewResourceGroup/TrafficManagerProfile](#).

After clicking on **Create**, the **Traffic Manager profile** will be created in a few minutes. After creating the a Traffic Manager profile, you can go to the **Overview** blade and share information such as routing methods you have chosen previously. The URL of a Traffic Manager profile named testinfra70533 looks like this: <http://testinfra70533tm.trafficmanager.net>.

The following screenshot shows all the available information that you can see via the **Overview** blade:



NAME	STATUS	MONITOR STATUS	TYPE	LOCATION
testinfra-webapp-ep	Enabled	Stopped	Azure endpoint	North Europe
testinfra-pubip-ep	Enabled	Degraded	Azure endpoint	West Europe

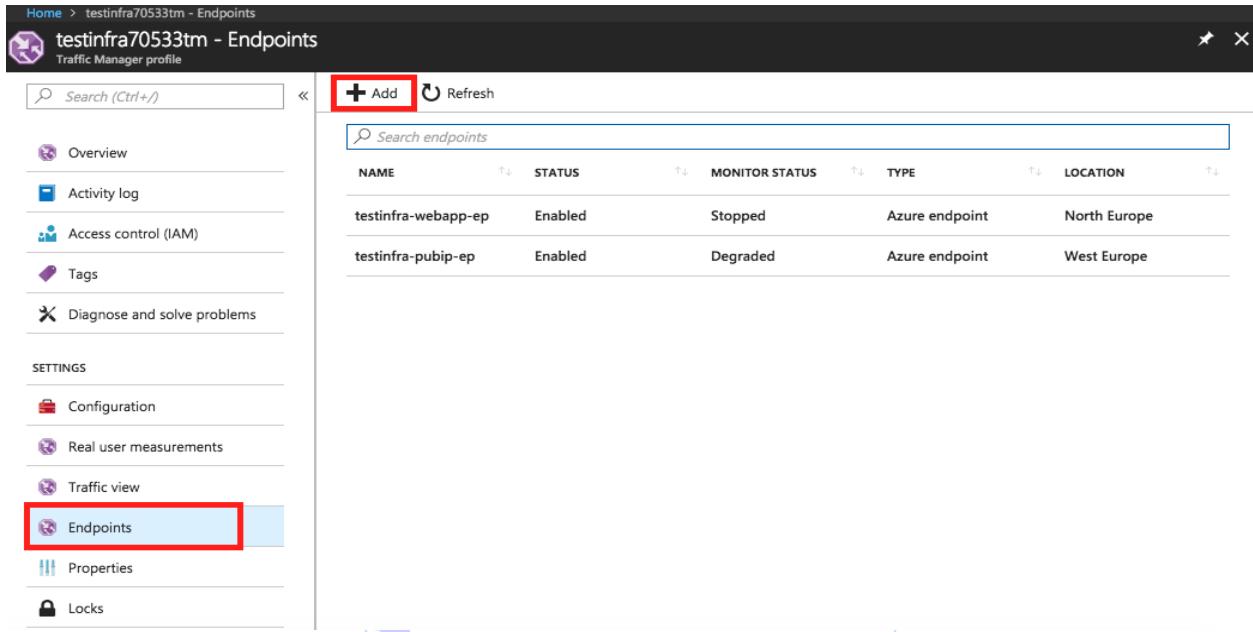
Adding endpoints to the Traffic Manager profile

After creating a Traffic Manager profile, you'll be able to add endpoints to the Traffic Manager. Currently, there are the following three types of endpoint supported by Traffic Manager:

- **Azure endpoints:** These are used for different IaaS services, PaaS services, or Public IP addresses within Azure.
- **External endpoints:** These are used for different services hosted on-premise or by other hosting providers to resolve any fully-qualified domain name (FQDN) outside Azure.

- Nested endpoints:** This is an endpoint type combining a parent endpoint and child endpoint in the same scenario, which is a more complex deployment scenarios.

While working with Azure Traffic Manager, it is possible to combine different types of endpoint in a single Traffic Manager profile as shown in the following screenshot. Additionally, each profile can contain multiple endpoint types. You can add an endpoint by clicking on **Add**:

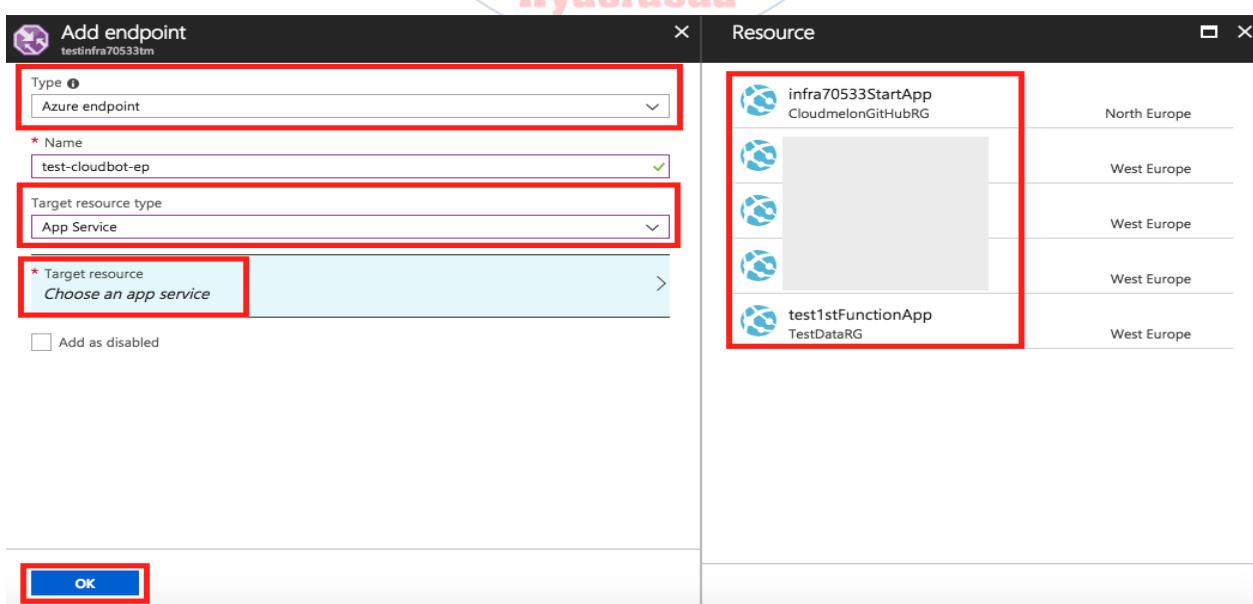


The screenshot shows the Azure Traffic Manager Endpoints page for the profile "testinfra70533tm". On the left, there's a sidebar with options like Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Settings (Configuration, Real user measurements, Traffic view, Endpoints), Properties, and Locks. The "Endpoints" option is highlighted with a red box. In the main area, there's a table of existing endpoints:

NAME	STATUS	MONITOR STATUS	TYPE	LOCATION
testinfra-webapp-ep	Enabled	Stopped	Azure endpoint	North Europe
testinfra-pubip-ep	Enabled	Degraded	Azure endpoint	West Europe

A red box also surrounds the "Add" button at the top of the table.

For example, while creating an Azure endpoint, you can choose a type such as cloud service, public IP address, or App Service, and then specify the target source, as shown in the following screenshot:



The screenshot shows the "Add endpoint" dialog for the profile "testinfra70533tm". The "Type" dropdown is set to "Azure endpoint" (highlighted with a red box). The "Name" field contains "test-cloudbot-ep" (with a green checkmark). The "Target resource type" dropdown is set to "App Service" (highlighted with a red box). The "Target resource" section shows a list of available app services:

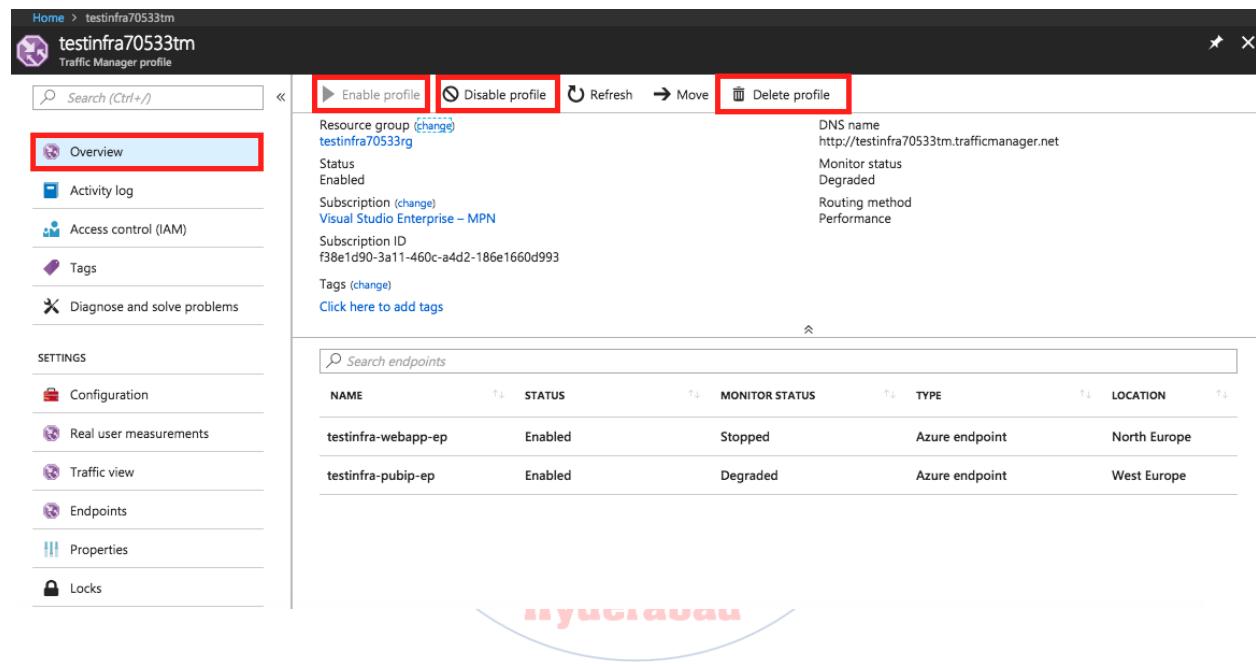
Resource	Location
infra70533StartApp CloudmelonGitHubRG	North Europe
(empty placeholder)	West Europe
(empty placeholder)	West Europe
(empty placeholder)	West Europe
test1stFunctionApp TestDataRG	West Europe

A red box highlights the entire "Resource" list. At the bottom of the dialog, there's an "OK" button (highlighted with a red box).

After you click on **OK**, within a few seconds an endpoint will be added in the Traffic Manager profile.

Managing Traffic Manager profiles

To manage created Traffic Manager profiles, let's go back to the **Overview** blade. Here you can select **Enable profile**, **Disable profile**, or **Delete profile** for current Traffic Manager profile:



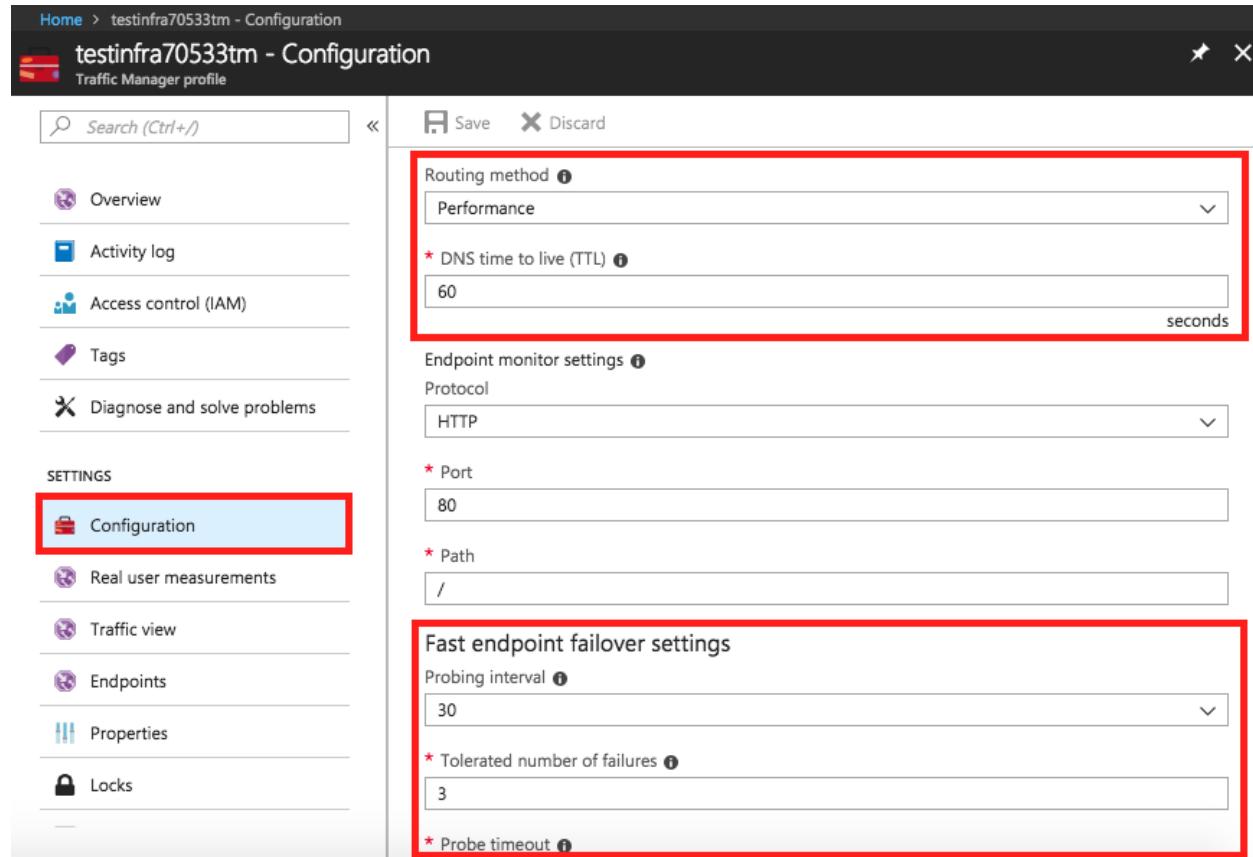
NAME	STATUS	MONITOR STATUS	TYPE	LOCATION
testinfra-webapp-ep	Enabled	Stopped	Azure endpoint	North Europe
testinfra-pubip-ep	Enabled	Degraded	Azure endpoint	West Europe

If you go to the **Configuration** blade in the Traffic Manager, you can also modify the configured routing method while creating the Traffic Manager. You can also define the value of the DNS **time to live (TTL)** as the live time of the client's local caching name server. After this period, the local caching name server will query the Traffic Manager system to update DNS entries.

Additionally, if you want to improve application resilience using a multi-region configuration with Traffic Manager, some interesting settings in the Fast endpoint failover settings section will definitely help you:

- **Probing interval:** This represents the time interval between endpoint health probes.
- **Tolerated number of failures:** This defines the number of health probe failures tolerated before an endpoint failure is triggered. It can be any number from 0 to 9.
- **Probe timeout:** This defines the time required before an endpoint health probe times out. This value must start from 5 and should be smaller than the probing interval value.

The following screenshot shows which settings are included in the **Configuration** blade for the Traffic Manager profile:

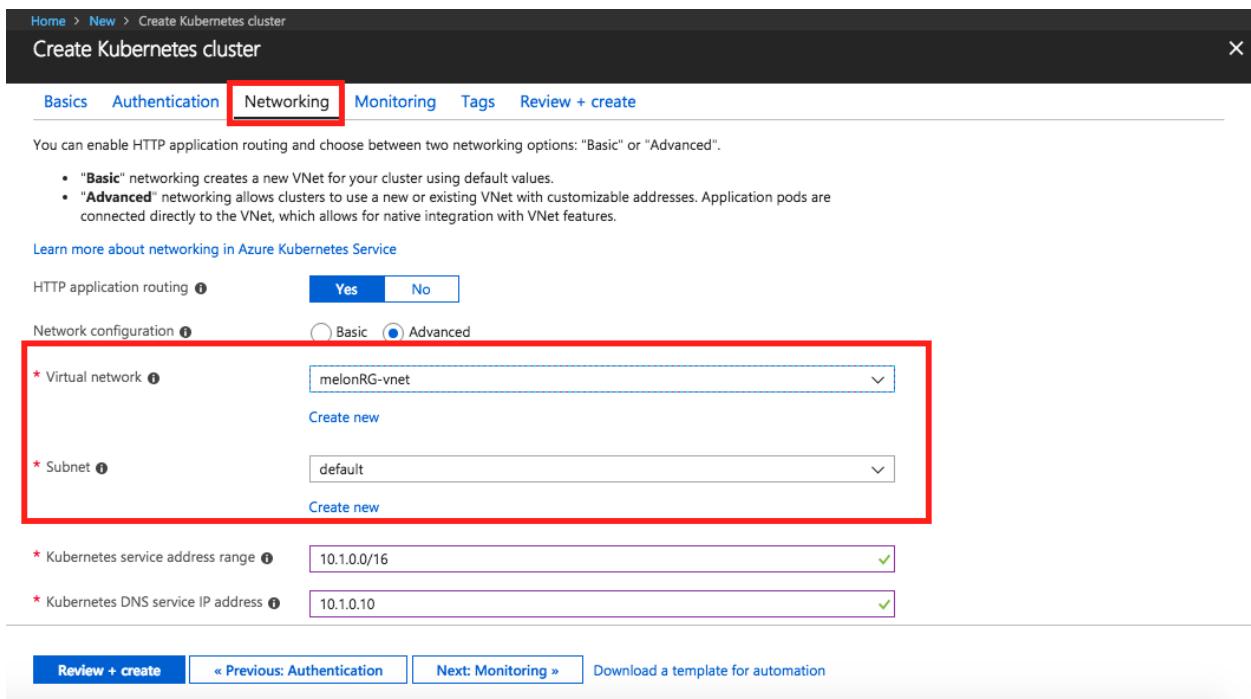


Integrating Azure services with an Azure virtual network

For most customers, a top concern when migrating their data and services to the cloud is how to restrict access control at the network level. Microsoft Azure provides a way to allow private access from instances of an Azure service deployed in the virtual network, by integrating Azure services with Azure Virtual Network. By definition, to integrate an Azure service means to guarantee communication between Azure services in the following two ways:

- Deploying dedicated instances of the Azure service into a **virtual network** while creation so that these services can be guaranteed private access within the virtual network
- Extending a virtual network to the Azure service through **Azure Network service endpoints** to allow access to individual service instances

There is a wide range of Azure services that can be deployed in Azure virtual networks, such as Azure Virtual machines, Virtual machine scale sets, Azure Kubernetes Services (AKS), and Azure Batch. A service such as AKS can be integrated into a virtual network. You can configure it to an existing VNet or create a new VNet for it during the first creation, as shown in the following screenshot:



Configuring Azure Virtual Network while creating a new AKS cluster

It is possible to deploy an Azure service into a subnet within an Azure virtual network that has also other Azure-integrated Azure Services and secure that service in the subnet with NSG.

Additionally, to reach out to these Azure services from on-premise, we can use different cross-premise connectivity options, and use a public IP address from the internet.

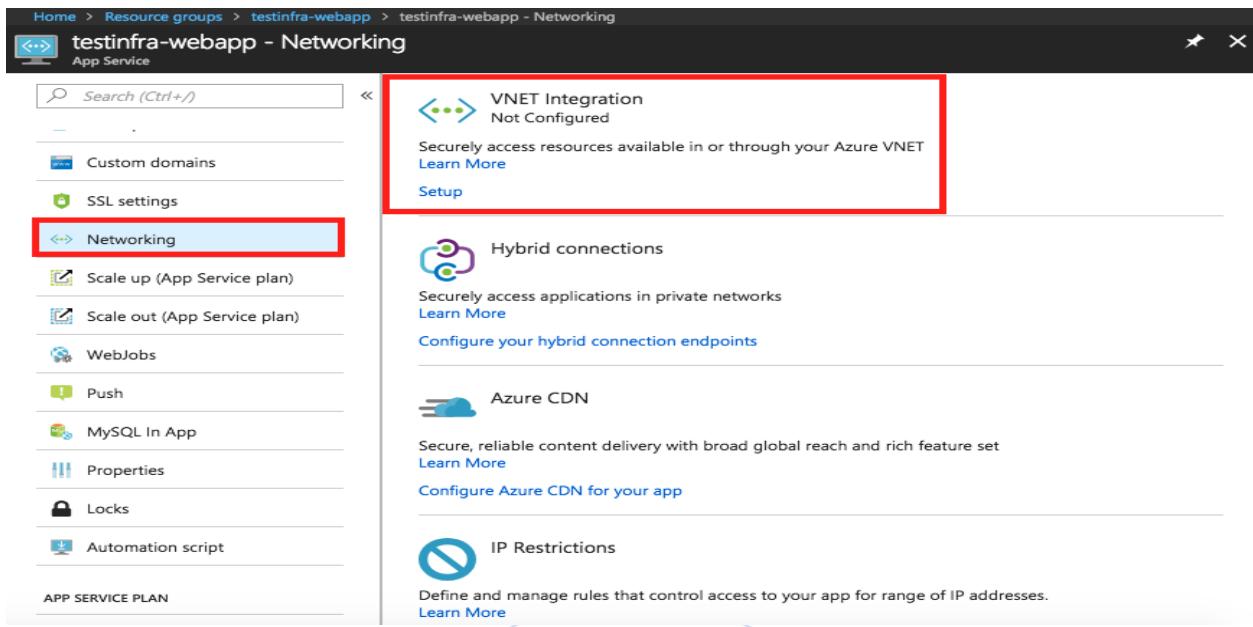
Restrict network access to PaaS resources using a service endpoint

Virtual network service endpoints is an amazing option in Azure to limit network access to some Azure service resources to a virtual network subnet in Azure. It allows you to use the private address space of the virtual network to access Azure services over a direct connection. Azure guarantees your traffic from the VNet to the Azure service always remains on the Microsoft Azure backbone network. This option is available to PaaS Azure services such as Azure Storage and Azure SQL Database for all Azure regions. For more information about this feature, you can go to the following link: <https://azure.microsoft.com/en-us/updates/?product=virtual-network>.

Integrating a web app in App Service with an Azure virtual network

VNet integration allows a web app hosted in the App Service plan to have access to resources in an Azure virtual network; users can host multiple Azure resources in an Azure virtual network with the control of access from the internet or on-premise using a variety of VPN connectivities.

For a web app that has been created, you can use the VNet Integration options to connect it to a new or existing Azure virtual network. You can go to the **Networking** blade for the web app and choose the VNet Integration option, as shown in the following screenshot:



Home > Resource groups > testinfra-webapp > testinfra-webapp - Networking

testinfra-webapp - Networking App Service

- Custom domains
- SSL settings
- Networking**
- Scale up (App Service plan)
- Scale out (App Service plan)
- WebJobs
- Push
- MySQL In App
- Properties
- Locks
- Automation script

APP SERVICE PLAN

VNET Integration
Not Configured
Securely access resources available in or through your Azure VNET [Learn More](#)

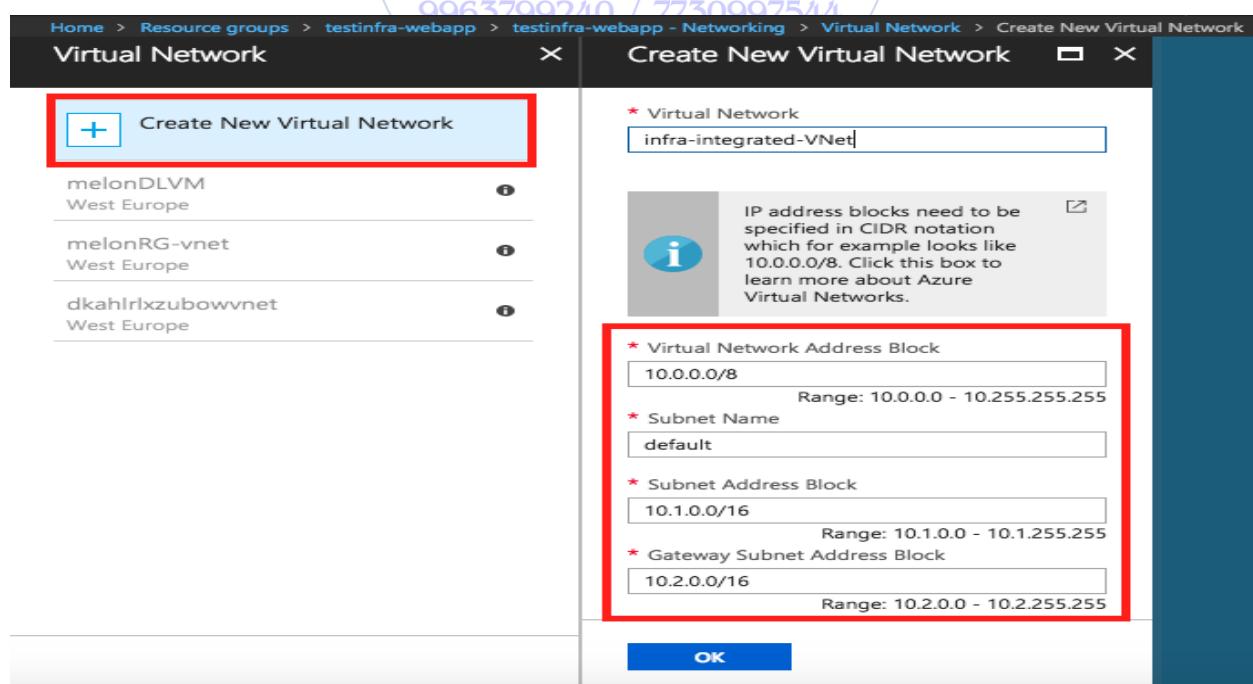
Setup

Hybrid connections
Securely access applications in private networks [Learn More](#)
[Configure your hybrid connection endpoints](#)

Azure CDN
Secure, reliable content delivery with broad global reach and rich feature set [Learn More](#)
[Configure Azure CDN for your app](#)

IP Restrictions
Define and manage rules that control access to your app for range of IP addresses. [Learn More](#)

Note that this feature is only available in the Standard and, Premium, and Isolated Pricing Plans. So if you're using the Standard or Premium plans, you can click on Setup to start to configure a VNet for your web app. If you want to integrate your web app with an existing VNet, the VNet would have a point-to-site VPN enabled with a Dynamic routing gateway so that it can be connected to an app; in the other case that P2S VPN with a static routing gateway or without any gateway, you cannot integrate it with your web app. In this case, you should create a new VNet. As shown in the following screenshot, you can create a new virtual network for your web app by specifying the address block and subnet for this VNet:



Home > Resource groups > testinfra-webapp > testinfra-webapp - Networking > Virtual Network > Create New Virtual Network

Virtual Network

Create New Virtual Network

Create New Virtual Network

melonDLVM
West Europe

melonRG-vnet
West Europe

dkahrlxzubowvnet
West Europe

* **Virtual Network**: infra-integrated-VNet

IP address blocks need to be specified in CIDR notation which for example looks like 10.0.0.0/8. Click this box to learn more about Azure Virtual Networks.

* **Virtual Network Address Block**: 10.0.0.0/8 (Range: 10.0.0.0 - 10.255.255.255)

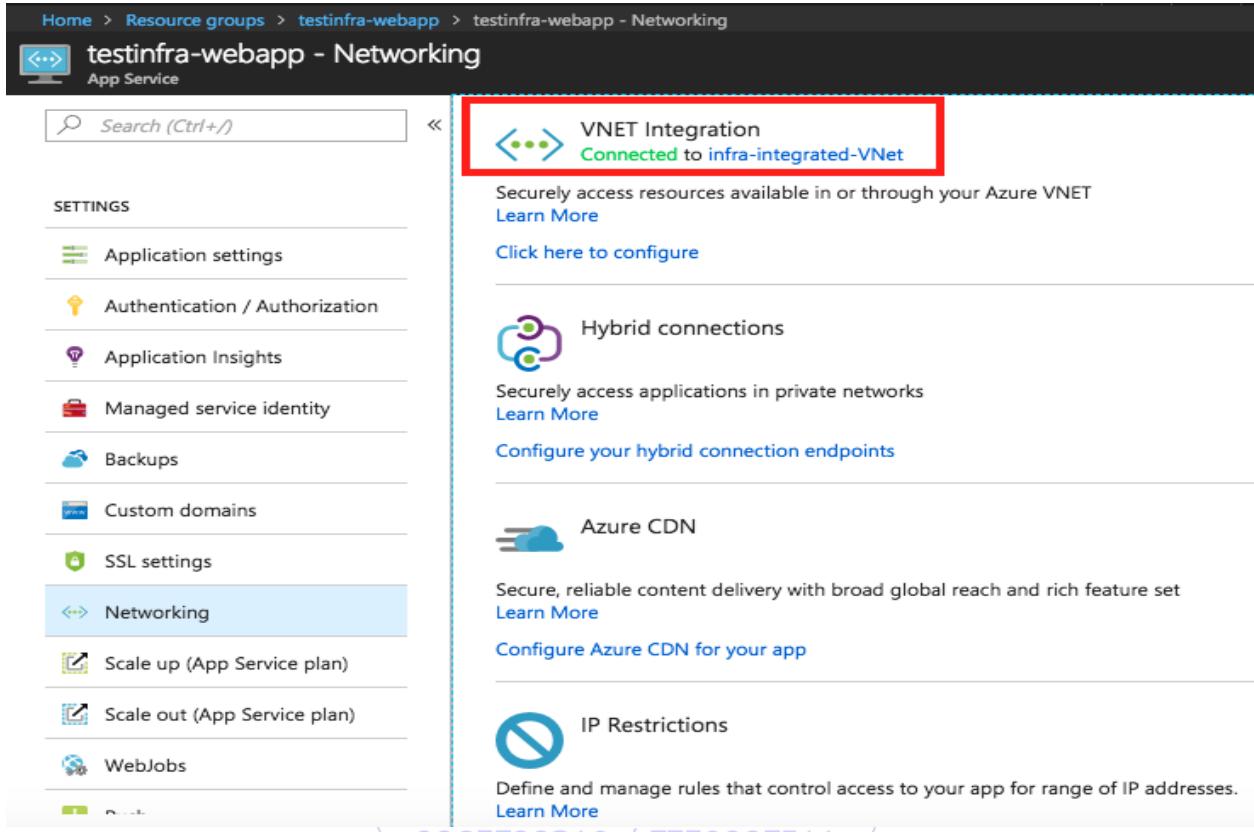
* **Subnet Name**: default

* **Subnet Address Block**: 10.1.0.0/16 (Range: 10.1.0.0 - 10.1.255.255)

* **Gateway Subnet Address Block**: 10.2.0.0/16 (Range: 10.2.0.0 - 10.2.255.255)

OK

The deployment usually takes a couple minutes. If your web app has been integrated with the newly created VNet, you'll see something similar to the following screenshot, while checking the **Networking** blade of your web app:



Home > Resource groups > testinfra-webapp > testinfra-webapp - Networking

testinfra-webapp - Networking

App Service

Search (Ctrl+ /)

SETTINGS

- Application settings
- Authentication / Authorization
- Application Insights
- Managed service identity
- Backups
- Custom domains
- SSL settings
- Networking**
- Scale up (App Service plan)
- Scale out (App Service plan)
- WebJobs

VNET Integration
Connected to **infra-integrated-VNet**

Securely access resources available in or through your Azure VNET
[Learn More](#)

[Click here to configure](#)

Hybrid connections

Securely access applications in private networks
[Learn More](#)

[Configure your hybrid connection endpoints](#)

Azure CDN

Secure, reliable content delivery with broad global reach and rich feature set
[Learn More](#)

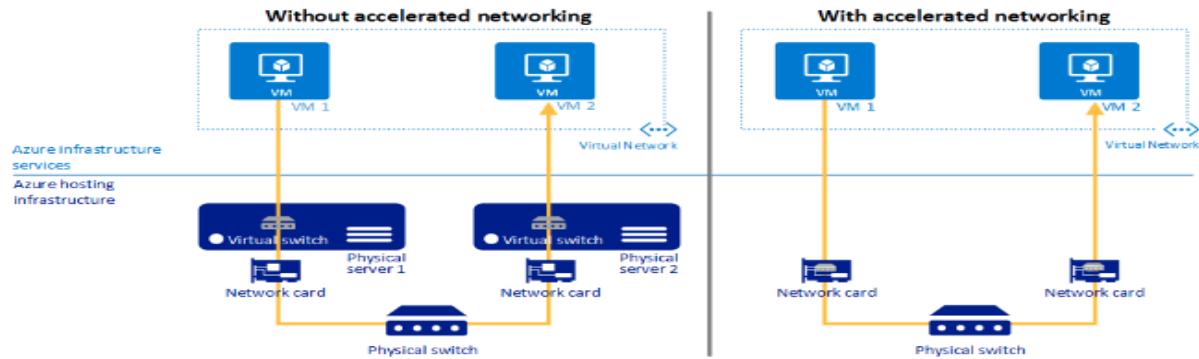
[Configure Azure CDN for your app](#)

IP Restrictions

Define and manage rules that control access to your app for range of IP addresses.
[Learn More](#)

Configuring accelerated networking to improve your networking performance

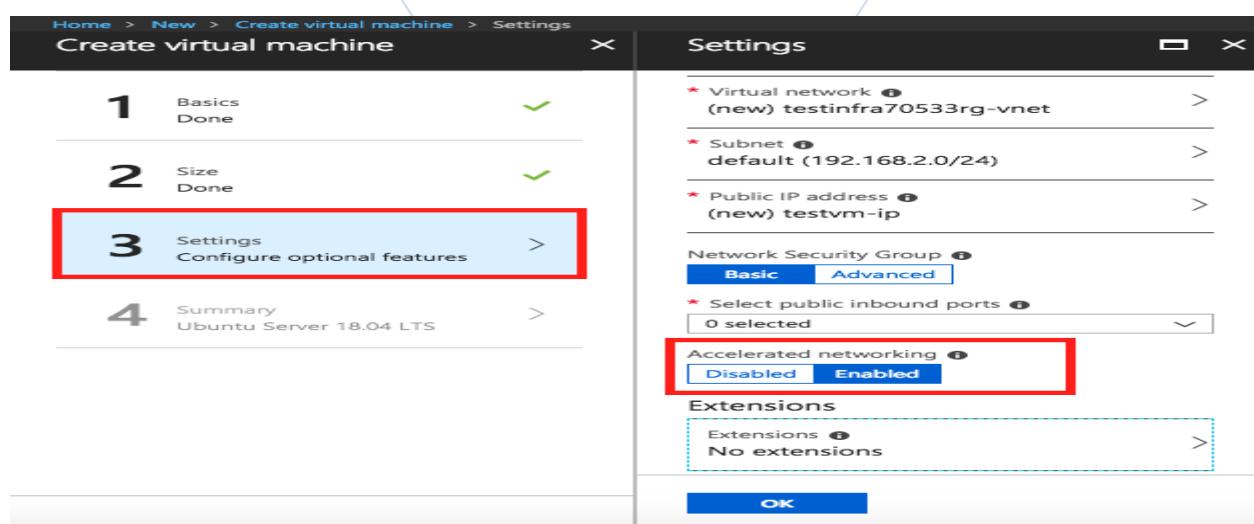
Accelerated networking is a pretty cool, free feature for Azure VMs and VM Scale Sets. It enables **single root I/O virtualization** (SR-IOV) for a VM or VMSS, which will greatly improve its networking performance. The reason why this feature is pretty cool is because it prevents all the inbound and outbound traffic from traversing the host and the virtual switch. With accelerated networking, network traffic arrives at the **Azure Network interface (NIC)** of the VMs and instances in VM Scale Sets then forwarded directly to the VM because all the network policies that the virtual switch applies are now offloaded and applied in hardware, as shown in the following screenshot:



The benefits from accelerated networking are reduced latency and CPU utilization, which significantly improves network performance.

Accelerated networking is supported on most general-purpose and compute-optimized such as **D/DSv2, F/Fs, D/DSv3, E/ESv3, Fsv2, and Ms/Mms** for Azure virtual machines and instances in VM scale sets.

To use this free feature, you can enable it while creating a new Azure VM only if you choose the VM size support this feature previously. If that is the case, you can enable this feature via the Azure Portal as shown in the following screenshot:



Managing Azure Identities

Azure AD is a cloud-based identity service provided by Microsoft Azure that allows you to secure access to cloud-based and on-premises applications and services.

In this chapter, we'll cover the following topics:

- Implementing and managing Azure AD
- Integrating applications with **Azure Active Directory (Azure AD)**
- Integrating **Active Directory Domain Services (ADDS)**
- Implementing **Active Directory Federation Services (ADFS)**

Implementing and managing Azure Active Directory (Azure AD)

Azure AD is a cloud-based multi-tenant identity solution that provides the following capabilities:

- Managing users, groups, and roles within organizations
- Providing multifactor authentication for both on-premises and cloud-resident resources
- Providing control access functionality control to enterprise-level applications
- Providing identity protection and managing identity devices
- Integrating **Single Sign-On (SSO)** functionality with a cloud-based **Software as a Service (SaaS)** applications and on-premises apps

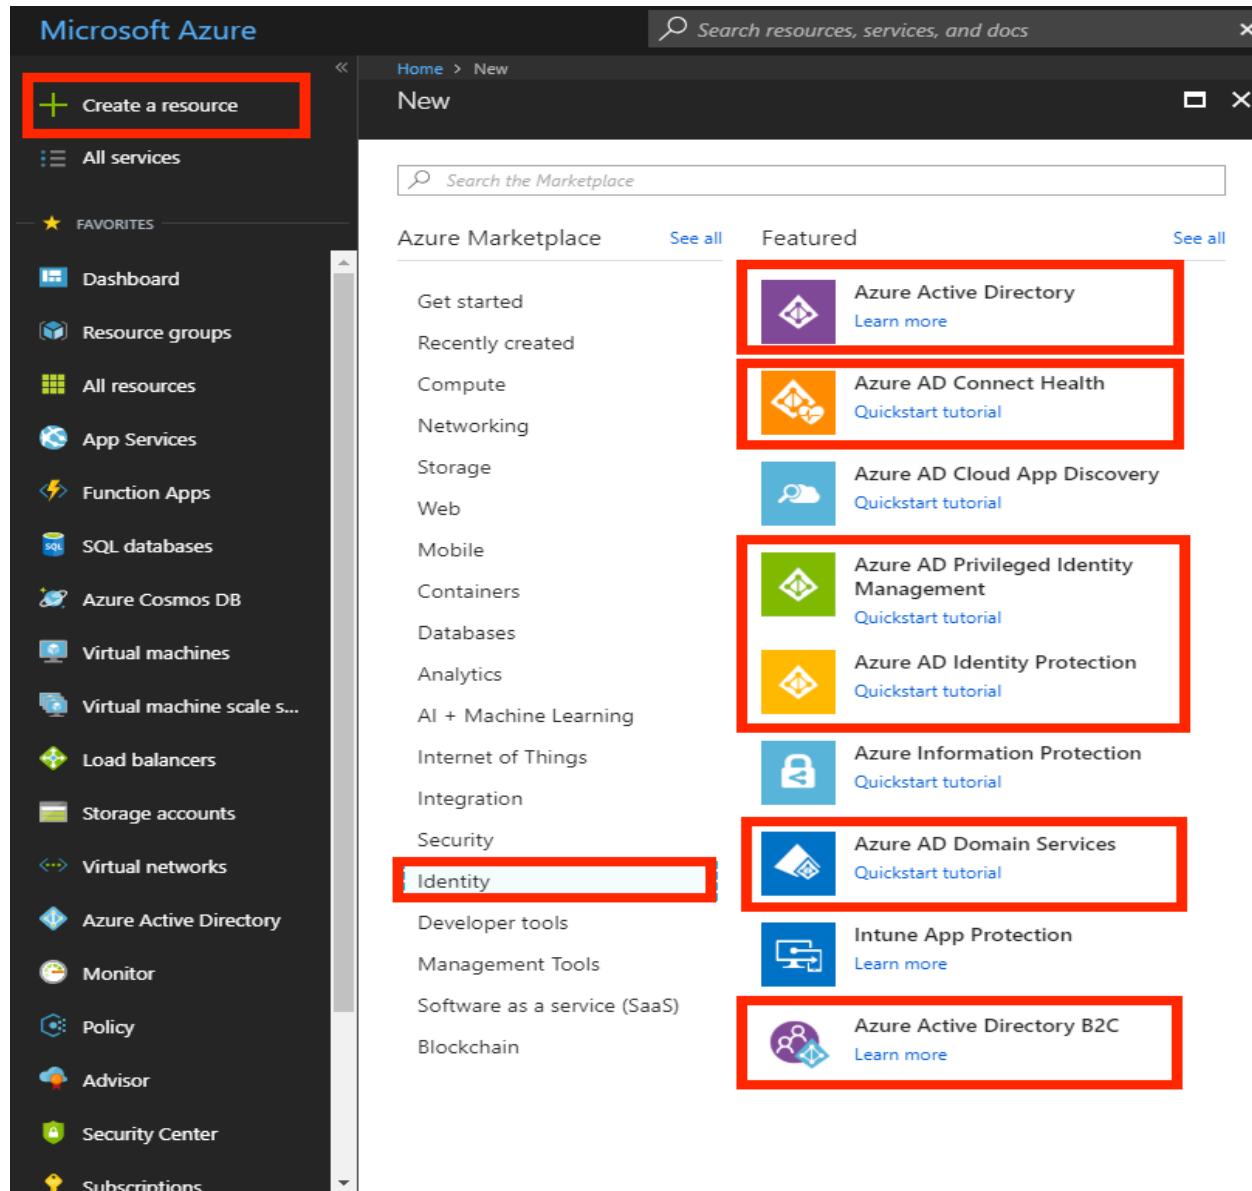
9963799240 / 7730997544

Ameerpet / Kondapur

Azure AD offers a rich, standards-based platform with many capabilities to help enterprises and organizations build cloud-based IDaaS solution in Azure.

Managing identities via Azure Active Directory admin center

If you go to the Azure Portal, click on **Create a resource**, and then on the **Identity** category, you'll find a couple of Azure AD-related services (as shown in the following screenshot):



The screenshot shows the Microsoft Azure Marketplace interface. On the left, there's a sidebar with a 'Create a resource' button highlighted by a red box. The main area has a search bar at the top. Below it, there are sections for 'Azure Marketplace' and 'Featured'. The 'Featured' section lists several services, each with a thumbnail icon and a 'Quickstart tutorial' link. Services shown include Azure Active Directory, Azure AD Connect Health, Azure AD Cloud App Discovery, Azure AD Privileged Identity Management, Azure AD Identity Protection, Azure Information Protection, Azure AD Domain Services, Intune App Protection, and Azure Active Directory B2C. The 'Identity' service is also highlighted with a red box.

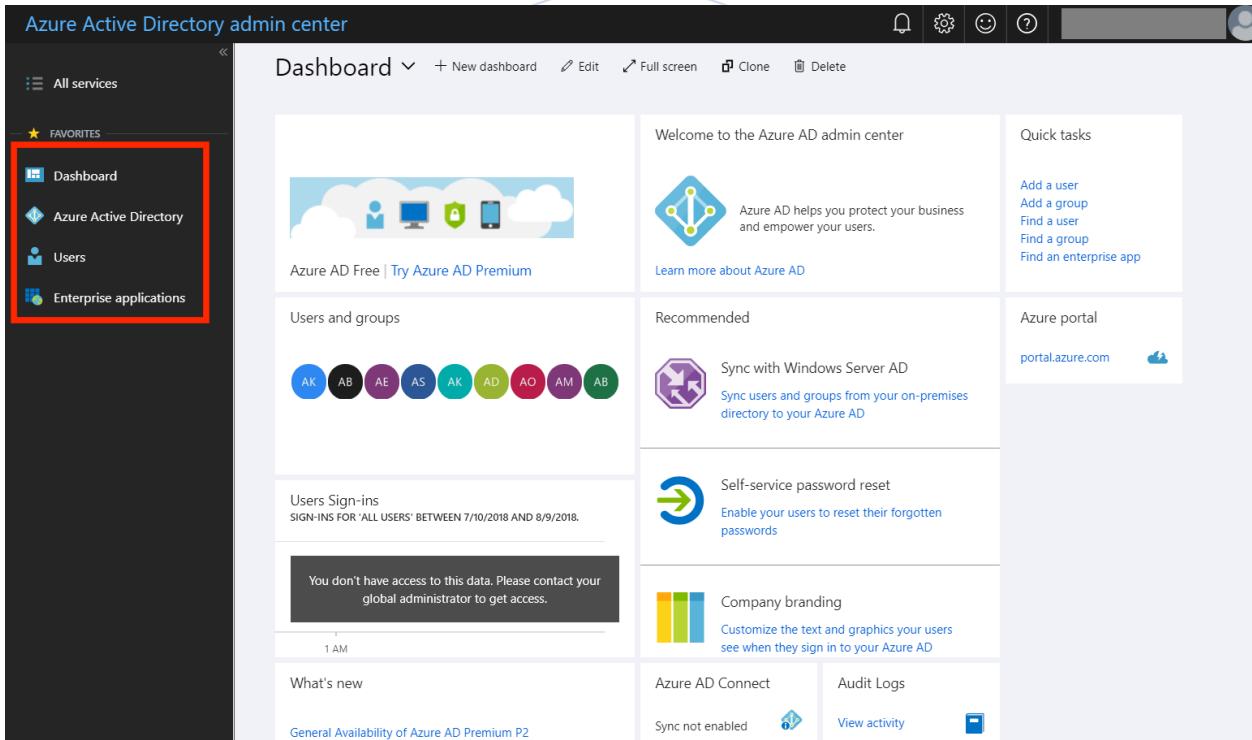
The following services are different capabilities provided by the Azure IDaaS solution:

- **Azure Active Directory:** It is a cloud-based directory, which is multi-tenant and supports identity management, application access management, and identity protection in a single solution.
- **Azure AD Identity Protection:** It is a feature provided by Azure that allows us to detect suspicious actions and configure responses automatically to protect the organization's identities.
- **Azure AD Privileged Identity:** It is an Azure Service to manage, control, and monitor access in Azure AD, Azure Resources, Office 365, Microsoft Intune, or other Microsoft online services within organization.

- **Azure AD Domain Services:** It provides managed domain features, such as domain join, LDAP, and Kerberos/NTLM authentication, which can extend the traditional Windows Server Active Directory.
- **Azure AD Connect Health:** It is a Azure service to monitor on-premises identity infrastructure and synchronization services.
- **Azure Information Protection (AIP):** It is a cloud-based solution that protects an organization's documents and emails using labels based on rules and conditions.

In Azure, there is a portal that was specially created for Azure AD, and you can input the following link into your browser to access it: <https://aad.portal.azure.com/>

Have a look at the following Azure AD portal:



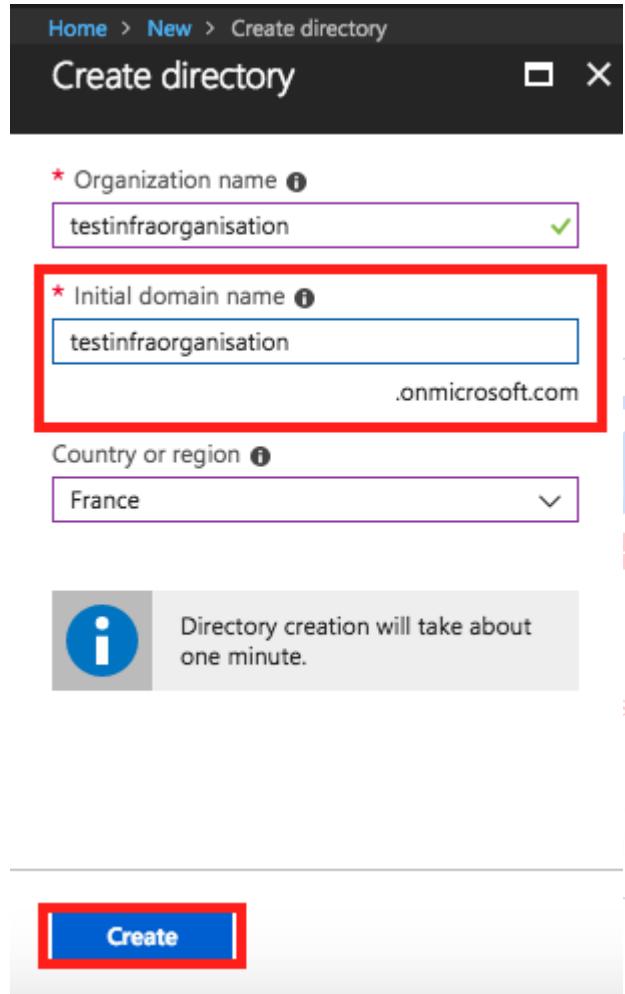
Creating an Azure Active Directory via the Azure Portal

Creating an Azure active directory is very simple. Go to the Azure Portal, search for **Azure active directory**, and then click on **Create**. You'll see the **Create directory** form, then fill in your organization's name and the initial domain name (as shown in the following screenshot).

Generally, your directory will look like the following

URL: #yourorganisationname#.onmicrosoft.com.

Then, click on **Create**, and it will take up to 1 minute to create a new directory successfully:



Home > New > Create directory

Create directory

* Organization name i

* Initial domain name i

.onmicrosoft.com

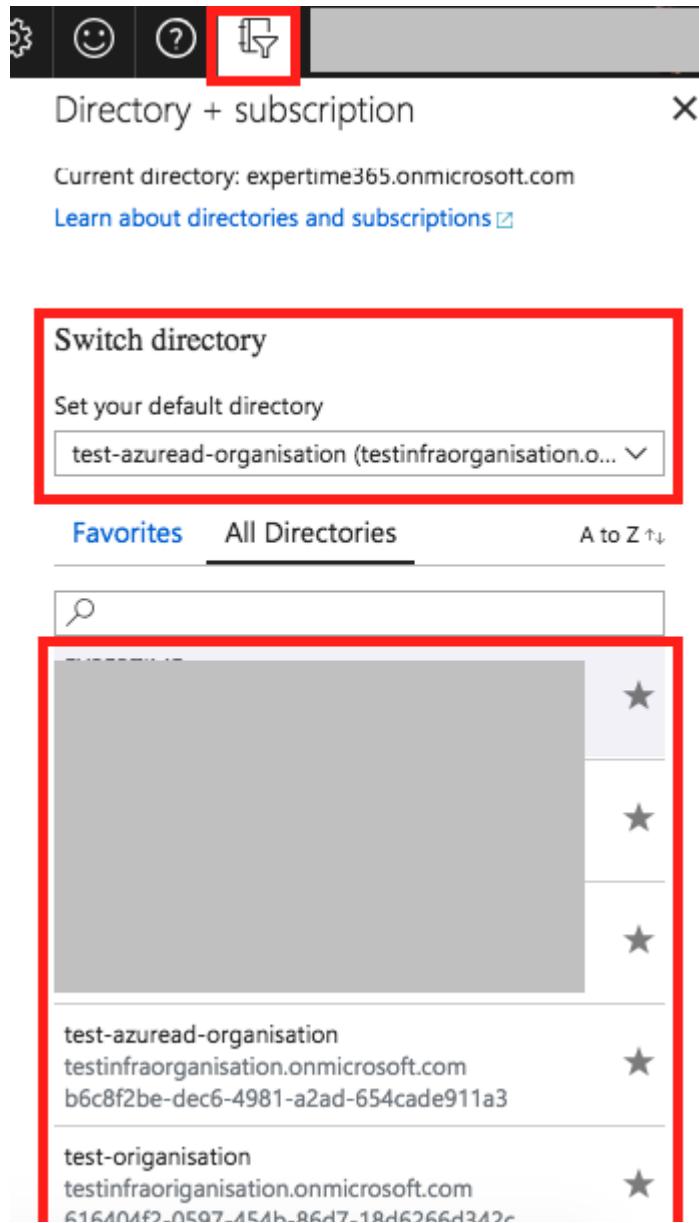
Country or region i

i Directory creation will take about one minute.

Create



After creating an Azure AD directory, go to the top of the Azure Portal and find the **Switch directory** button (as shown in the following screenshot) to switch to the target directory:



Directory + subscription X

Current directory: expertime365.onmicrosoft.com

[Learn about directories and subscriptions](#)

Switch directory

Set your default directory

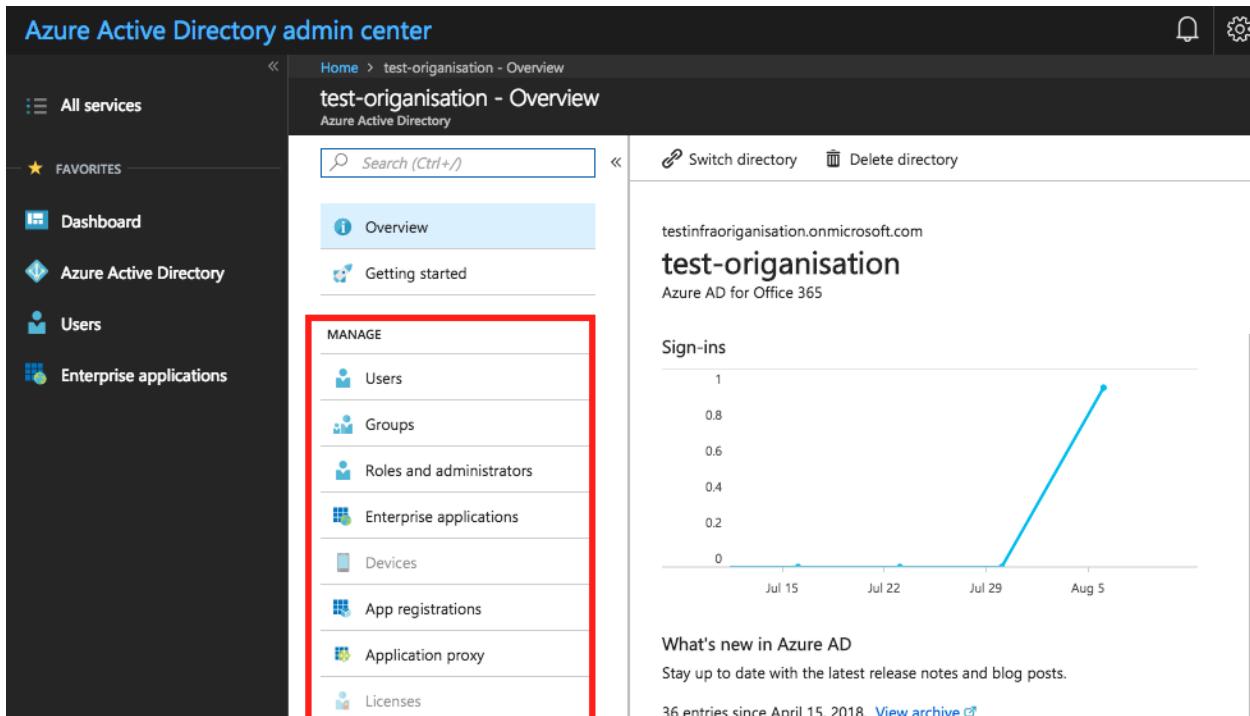
test-azuread-organisation (testinfraorganisation.o...)

Favorites [All Directories](#) [A to Z ↑↓](#)

test-azuread-organisation
testinfraorganisation.onmicrosoft.com
b6c8f2be-dec6-4981-a2ad-654cade911a3

test-organisation
testinfraorganisation.onmicrosoft.com
616A0AF9_0E07_AE1H_86A7_19A6C6642A0r

After switching a directory successfully, you can go to the Azure AD portal of this directory, shown in the following screenshot. You can manage the users, groups, and roles of current Azure AD and additionally integrate applications with the Azure AD directory:



Azure Active Directory admin center

Home > test-organisation - Overview

test-organisation - Overview
Azure Active Directory

Switch directory Delete directory

testinfraorganisation.onmicrosoft.com
test-organisation
Azure AD for Office 365

Sign-ins

Date	Sign-ins
Jul 15	0.0
Jul 22	0.0
Jul 29	0.0
Aug 5	1.0

What's new in Azure AD
Stay up to date with the latest release notes and blog posts.
36 entries since April 15, 2018. [View archive](#)

Creating and managing Azure AD users

In the Azure AD portal, we can take advantage of the following capabilities:

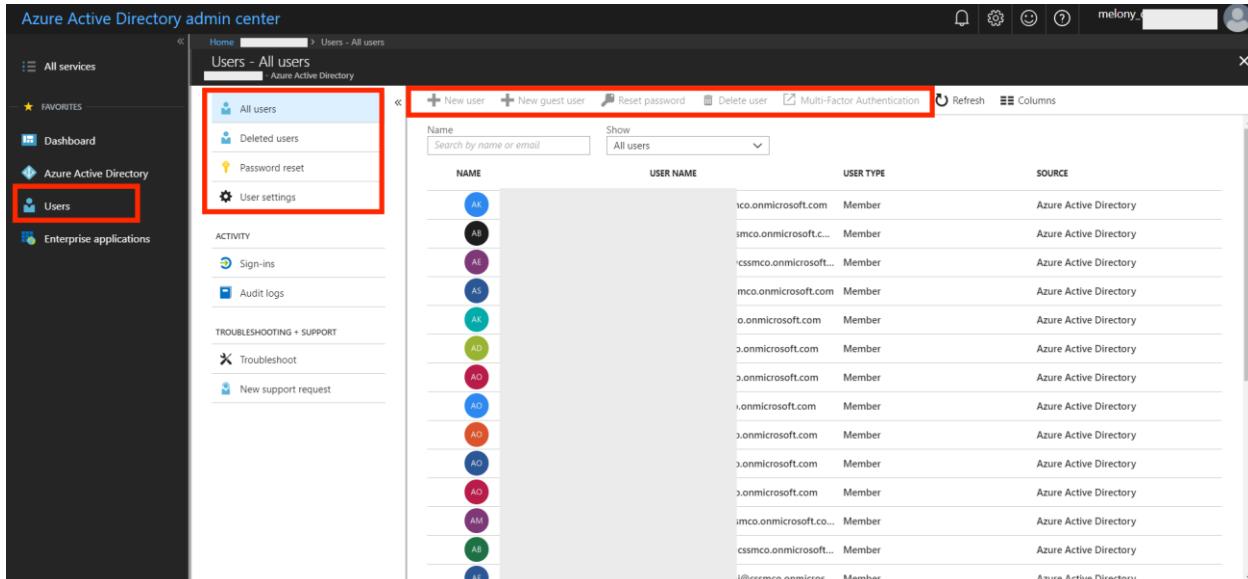
9963799240 / 7730997544

Ameerpet / Kondapur

Hyderabad

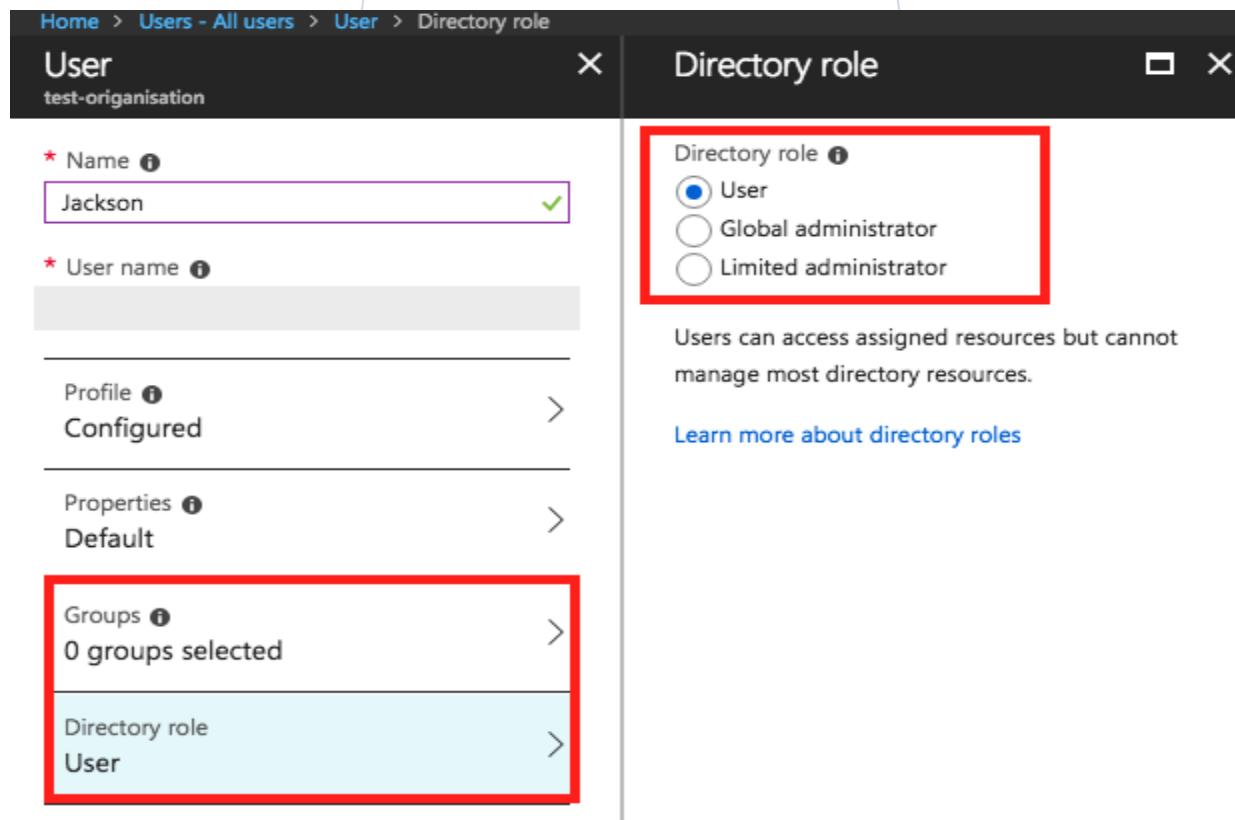
- Creating a new user and guest user
- Deleting a user
- Resetting the password
- Configuring MFA for an existing user

After clicking on **All users**, you can see as shown in the following screenshot with all the users in the current directory. At the top of the users list, you can find all these buttons to help you manage the users in your organization:



The screenshot shows the Azure Active Directory admin center interface. On the left, there's a navigation sidebar with 'All services', 'Dashboard', 'Azure Active Directory' (which has 'Users' selected), and 'Enterprise applications'. The main area is titled 'Users - All users' under 'Azure Active Directory'. It includes a toolbar with 'New user', 'New guest user', 'Reset password', 'Delete user', and 'Multi-Factor Authentication' buttons. A search bar at the top right allows searching by name or email. Below the toolbar is a table with columns for 'NAME', 'USER NAME', 'USER TYPE', and 'SOURCE'. The table lists numerous users from 'mco.onmicrosoft.com' to 'jilbresen.onmicrosoft.com', all marked as 'Member' in 'Azure Active Directory'.

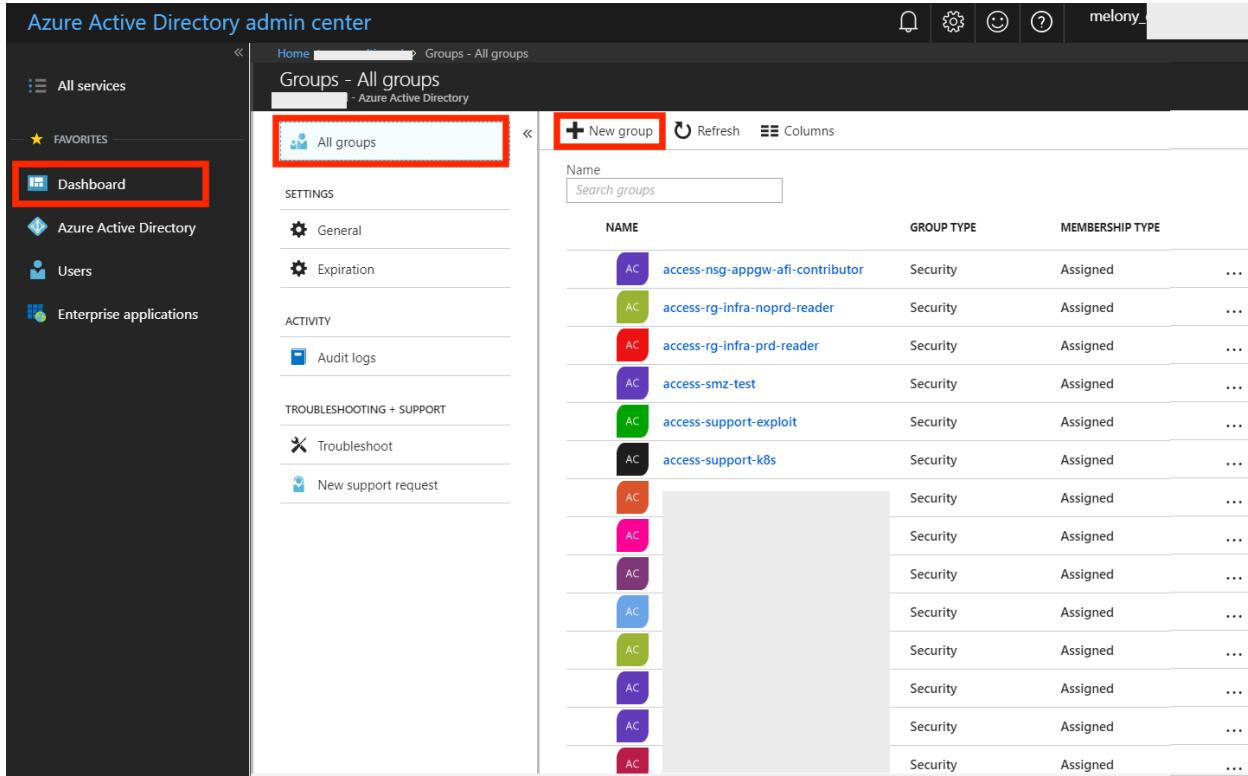
To add a new user, click on **+ New user** and then you'll see the page for user creation. On this page, you can configure the user group and its role. By default, the profile is only available as **Azure Active Directory**:



The screenshot shows the 'User' configuration page. The URL in the address bar is 'Home > Users - All users > User > Directory role'. The left pane shows a user named 'Jackson' with a 'test-organisation' profile. It includes sections for 'Profile' (Configured), 'Properties' (Default), 'Groups' (0 groups selected), and 'Directory role' (User). The 'Groups' and 'Directory role' sections are highlighted with red boxes. The right pane is titled 'Directory role' and contains a section for 'Directory role' with three radio button options: 'User' (selected), 'Global administrator', and 'Limited administrator'. A note below states: 'Users can access assigned resources but cannot manage most directory resources.' and a link to 'Learn more about directory roles'.

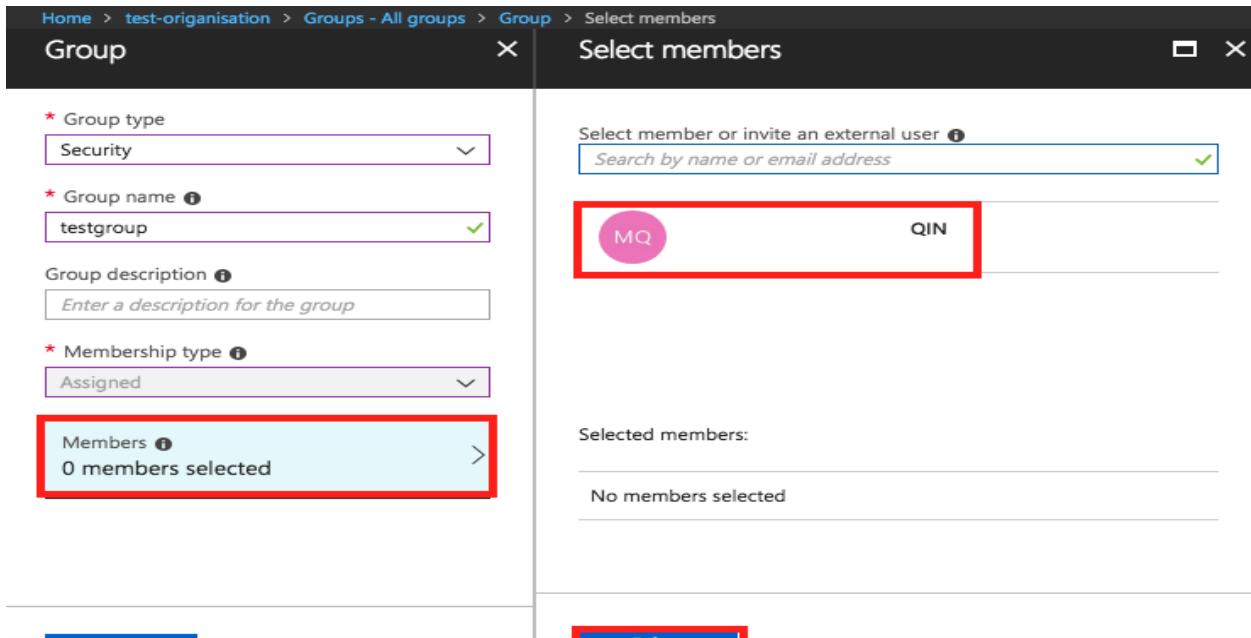
Creating Azure AD groups and managing user groups

You can go to the Azure AD portal and click on **Azure active directory**, then click on the **All groups** blade. Go to a new page, which has all the available groups within the organization (as shown in the following screenshot; you can manage your users in different group types, such as users in security part) and click on **+ New group** to create a new group:



Name	Group Type	Membership Type	Actions
access-nsg-appgw-afi-contributor	Security	Assigned	...
access-rg-infra-noprd-reader	Security	Assigned	...
access-rg-infra-prd-reader	Security	Assigned	...
access-smz-test	Security	Assigned	...
access-support-exploit	Security	Assigned	...
access-support-k8s	Security	Assigned	...
(empty)	Security	Assigned	...
(empty)	Security	Assigned	...
(empty)	Security	Assigned	...
(empty)	Security	Assigned	...
(empty)	Security	Assigned	...
(empty)	Security	Assigned	...
(empty)	Security	Assigned	...
(empty)	Security	Assigned	...
(empty)	Security	Assigned	...

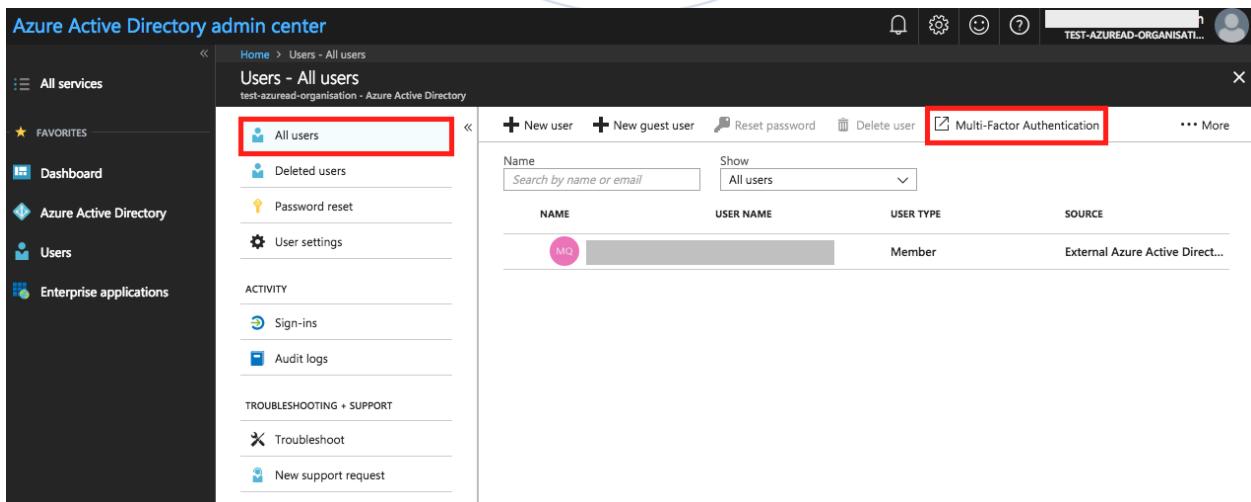
Then, fill in the group type that has office 365 or security type and group name. You can choose members in the member list or invite an external user to the group by clicking on **Members** and then **Select**, as shown in the following screenshot. Finally, click on **Create** to create a new Azure AD group:



Enabling Multi-Factor Authentication for users

To enable the MFA feature, you should have **Azure AD Premium** licenses so that you have access to full-featured use of MFA in the Azure cloud or the on-premises Azure MFA server.

Make sure you meet all the prerequisites, then you can click on **Multi-Factor Authentication** to configure it (as shown here):

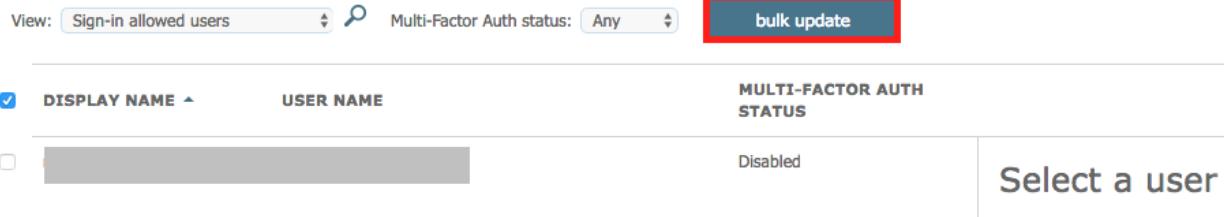


You'll see another page in another tab of your browser. Click on **service settings**:

multi-factor authentication

users **service settings**

Note: only users licensed to use Microsoft Online Services are eligible for Multi-Factor Authentication. Learn more about how to license other users. Before you begin, take a look at the [multi-factor auth deployment guide](#).



DISPLAY NAME	USER NAME	MULTI-FACTOR AUTH STATUS
<input type="checkbox"/> [REDACTED]	[REDACTED]	Disabled

Select a user

In the **service settings** tab, you can configure all the information related to MFA, choose the available ways to users such as text message, phone, notification through push-up or verification code (as follows):

multi-factor authentication

users **service settings**

app passwords

- Allow users to create app passwords to sign in to non-browser apps
- Do not allow users to create app passwords to sign in to non-browser apps

verification options

Methods available to users:

- Call to phone
- Text message to phone
- Notification through mobile app
- Verification code from mobile app or hardware token

remember multi-factor authentication

- Allow users to remember multi-factor authentication on devices they trust
Days before a device must re-authenticate (1-60):

save

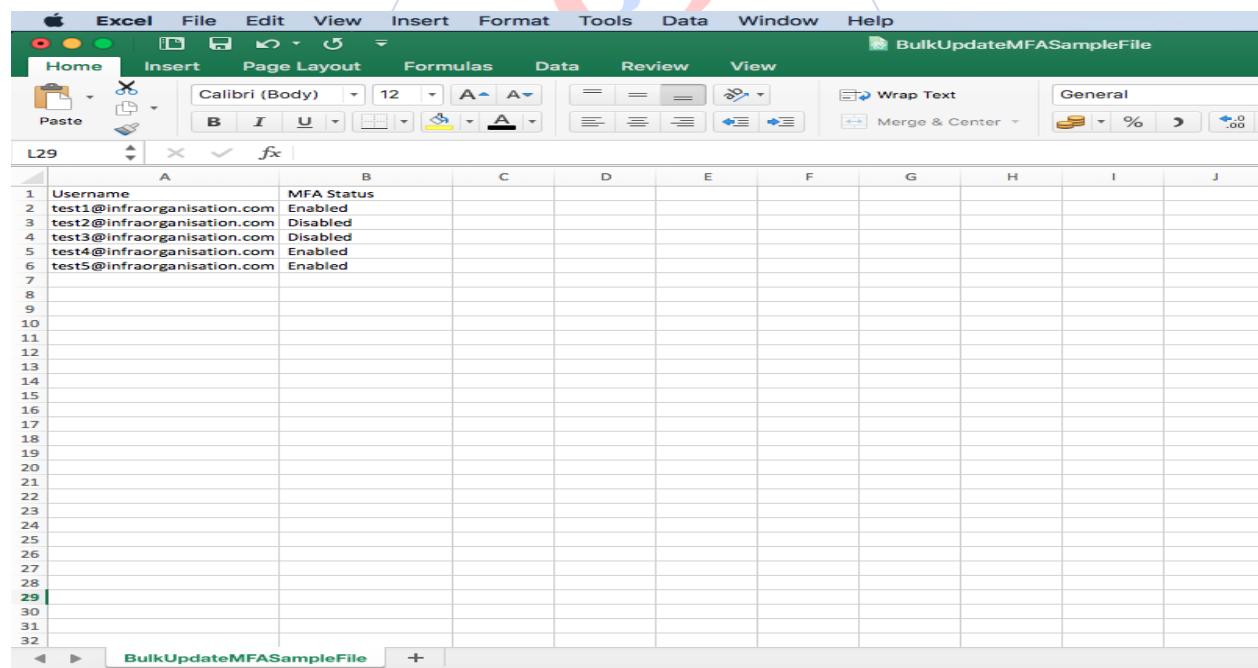
Using bulk update for custom user profile properties

When you want to configure the MFA feature for multiple users, it is possible to use the **bulk update** feature. To do this, you should click on the **bulk update** button in the users tab of the previous page. Then, you'll see a popup, shown here, where you can upload a .csv file:



a CSV file to enable MFA for multiple users

The example .csv file you'll upload is as follows:

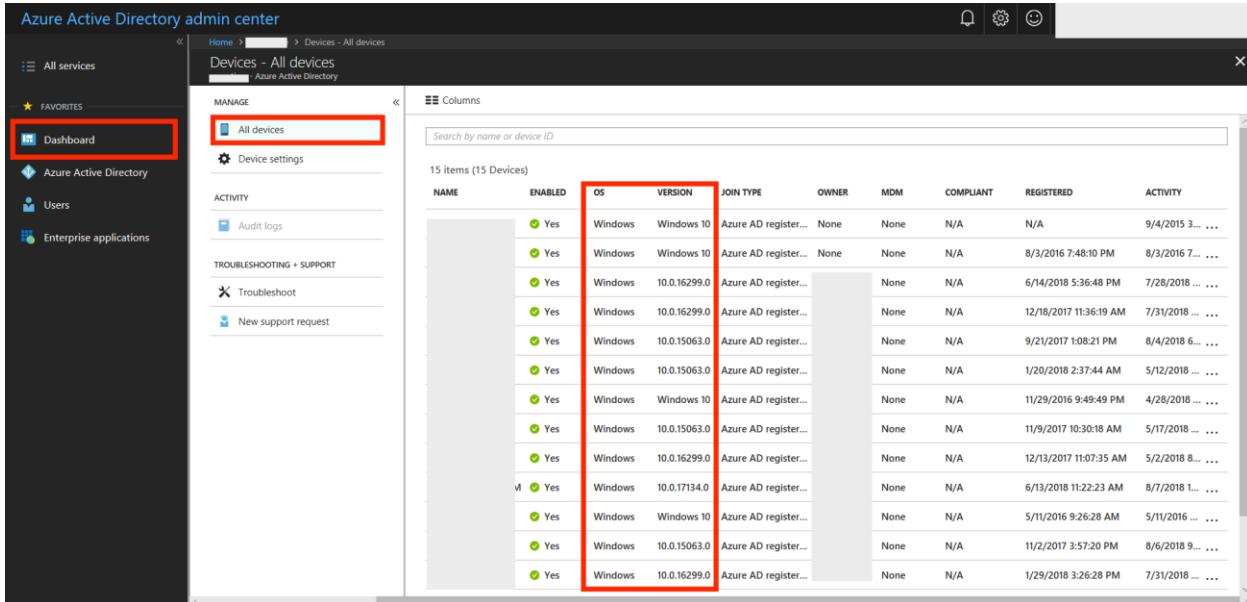


	A	B	C	D	E	F	G	H	I	J
1	Username	MFA Status								
2	test1@infraorganisation.com	Enabled								
3	test2@infraorganisation.com	Disabled								
4	test3@infraorganisation.com	Disabled								
5	test4@infraorganisation.com	Enabled								
6	test5@infraorganisation.com	Enabled								
7										
8										
9										
10										
11										
12										
13										
14										
15										
16										
17										
18										
19										
20										
21										
22										
23										
24										
25										
26										
27										
28										
29										
30										
31										
32										

How to use a CSV file to configure MFA for multiple users

Managing devices

Azure AD allows you to manage single sign-on for devices, apps, and services from anywhere. Go to the Azure AD portal and click on **Devices** to manage all the devices that have been registered in the repository (as shown here):



Name	Enabled	OS	Version	Join Type	Owner	MDM	Compliant	Registered	Activity
	Yes	Windows	Windows 10	Azure AD register...	None	None	N/A	N/A	9/4/2015 3...
	Yes	Windows	Windows 10	Azure AD register...	None	None	N/A	8/3/2016 7:48:10 PM	8/3/2016 7...
	Yes	Windows	10.0.16299.0	Azure AD register...	None	N/A	6/14/2018 5:36:48 PM	7/28/2018 ...	
	Yes	Windows	10.0.16299.0	Azure AD register...	None	N/A	12/18/2017 11:36:19 AM	7/31/2018 ...	
	Yes	Windows	10.0.15063.0	Azure AD register...	None	N/A	9/21/2017 1:08:21 PM	8/4/2018 6...	
	Yes	Windows	10.0.15063.0	Azure AD register...	None	N/A	1/20/2018 2:37:44 AM	5/12/2018 ...	
	Yes	Windows	Windows 10	Azure AD register...	None	N/A	11/29/2016 9:49:49 PM	4/28/2018 ...	
	Yes	Windows	10.0.15063.0	Azure AD register...	None	N/A	11/9/2017 10:30:18 AM	5/17/2018 ...	
	Yes	Windows	10.0.16299.0	Azure AD register...	None	N/A	12/13/2017 11:07:35 AM	5/2/2018 8...	
V	Yes	Windows	10.0.17134.0	Azure AD register...	None	N/A	6/13/2018 11:22:23 AM	8/7/2018 1...	
	Yes	Windows	Windows 10	Azure AD register...	None	N/A	5/11/2016 9:26:28 AM	5/11/2016 ...	
	Yes	Windows	10.0.15063.0	Azure AD register...	None	N/A	11/2/2017 3:57:20 PM	8/6/2018 9...	
	Yes	Windows	10.0.16299.0	Azure AD register...	None	N/A	1/29/2018 3:26:28 PM	7/31/2018 ...	

The Leader in Software Training
 9963799240 / 7730997544

Azure AD provides support for these devices for the **Bring Your Own Device (BYOD)** scenario so that the users in the directory that is with work or school account can work with different devices, such as on laptop, tablet, or mobile, and across different OSes, such as Windows 10, iOS, Android, and macOS.

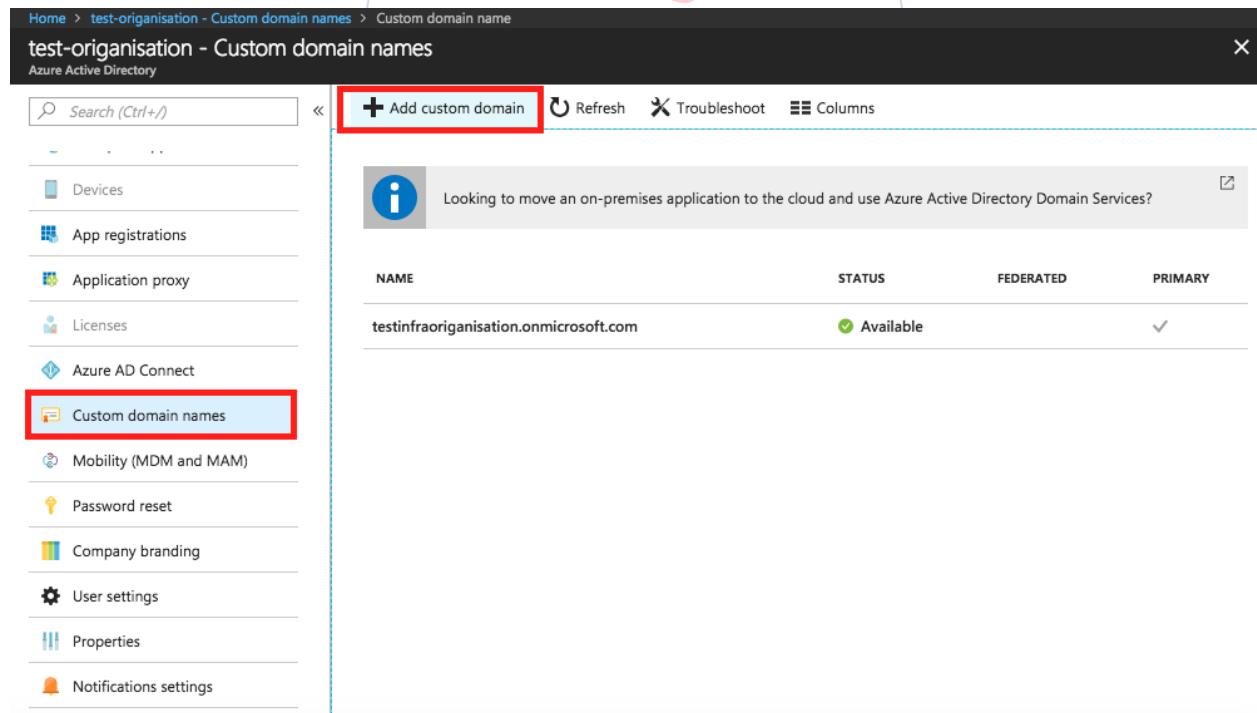
You may be wondering how did this device list come, if you want to know more about how to configure a new Windows 10 device joining Azure AD, you can refer to the following link: <https://docs.microsoft.com/en-us/azure/active-directory/devices/azuread-joined-devices-frx>.

Add a custom domain

When creating a new Azure AD directory, we configured an initial domain that is in the form of organisationname.onmicrosoft.com; the domain name cannot be changed or deleted, but the major user may not be familiar with this domain. Azure provides a friendly feature to let an administrator add a custom domain name for the Azure AD directory. After adding the custom domain name, the organisationname.onmicrosoft.com URL becomes the following:

youname.organisationname.com

An example is test@qualitythought.com. Here, test is your name or the name of anyone else, and qualitythought is the domain name; it can be another domain name that you expected. To add a custom domain name, you need to go to the Azure active directory page and click on the **Custom domain names** blade, then click on the **+ Add custom domain** button (shown in the following screenshot):



NAME	STATUS	FEDERATED	PRIMARY
testinfraorganisation.onmicrosoft.com	Available		✓

Then, you'll get a create **Custom domain name** page, and you can fill in the page **Custom domain name** and add the target domain in this page (as shown in the following screenshot):

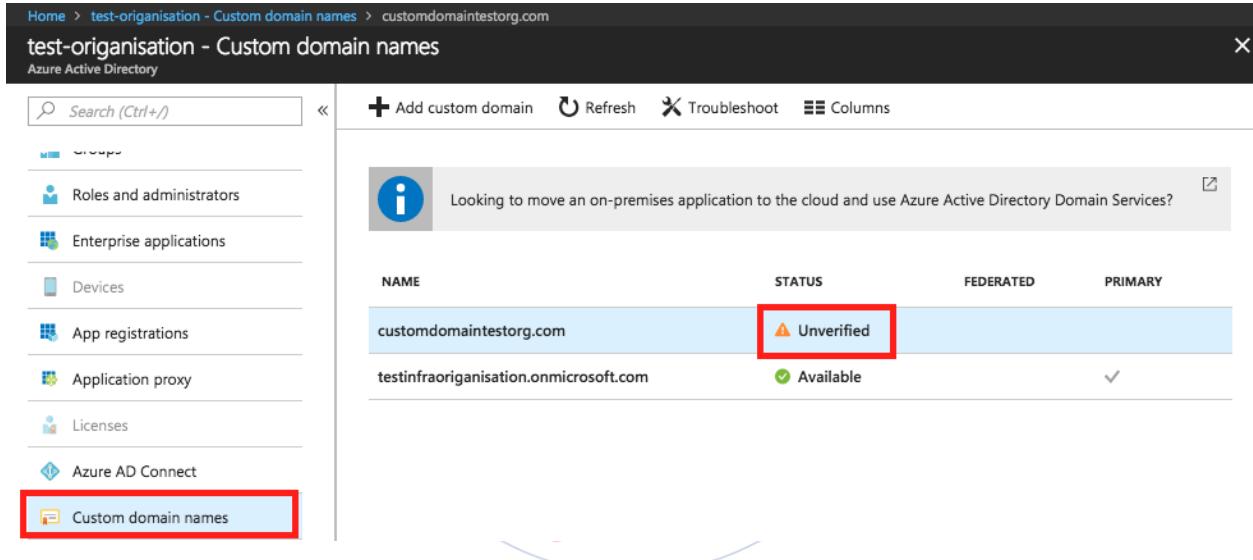
Custom domain name □ X

test-organisation

* Custom domain name i
 ✓

Add Domain

After creating this domain, it may not be operational yet. Go back to the **Custom domain names** page. As you can see as the following screenshot, the status of this domain is marked as **Unverified**. This means that this domain name should be verified by Azure:



NAME	STATUS	FEDERATED	PRIMARY
customdomaintestorg.com	⚠ Unverified		
testinfraorganisation.onmicrosoft.com	✓ Available		✓

Click on this domain label, then go to the verify page. Choose both the types of DNS records, TXT or MX, as record types here. The MX record specifies where the emails for your domain should be delivered, and the TXT record is used to store text-based information related to your domain. We choose the TXT record here and click on **Verify**, as shown in the following screenshot. If everything goes well, after a few minutes or so, the newly created custom domain will be operational:

[Home](#) > [test-organisation - Custom domain names](#) > [customdomaintestorg.com](#)

customdomaintestorg.com

Custom domain name



Delete



To use customdomaintestorg.com with your Azure AD, create a new TXT record with your domain name registrar using the info below.

RECORD TYPE
TXT
MX

ALIAS OR HOST NAME

@



DESTINATION OR POINTS TO ADDRESS

MS=ms37589618



TTL

3600


[Share these settings via email](#)
Verify domain

Verification will not succeed until you have configured your domain with your registrar as described above.

Verify

Conditional access

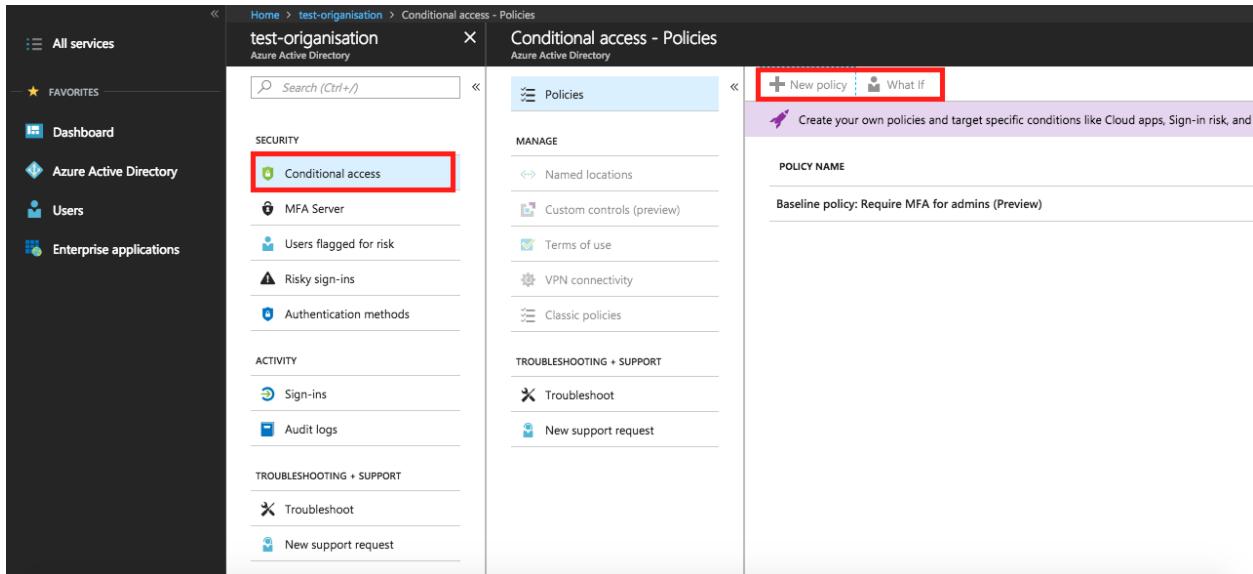
The Leader in Software Training
 9963799240 / 7730997544

Ameerpet / Kondapur

Hyderabad

Azure Active Directory also provides a very cool ability to control access to cloud native applications based on conditions. The conditions defined in the condition policy that defined **when does this happen and what to do when this happens**.

You can define a new condition policy via the Azure Portal; go to the Azure AD portal and click on the **Conditional access** blade (as shown in the following screenshot):



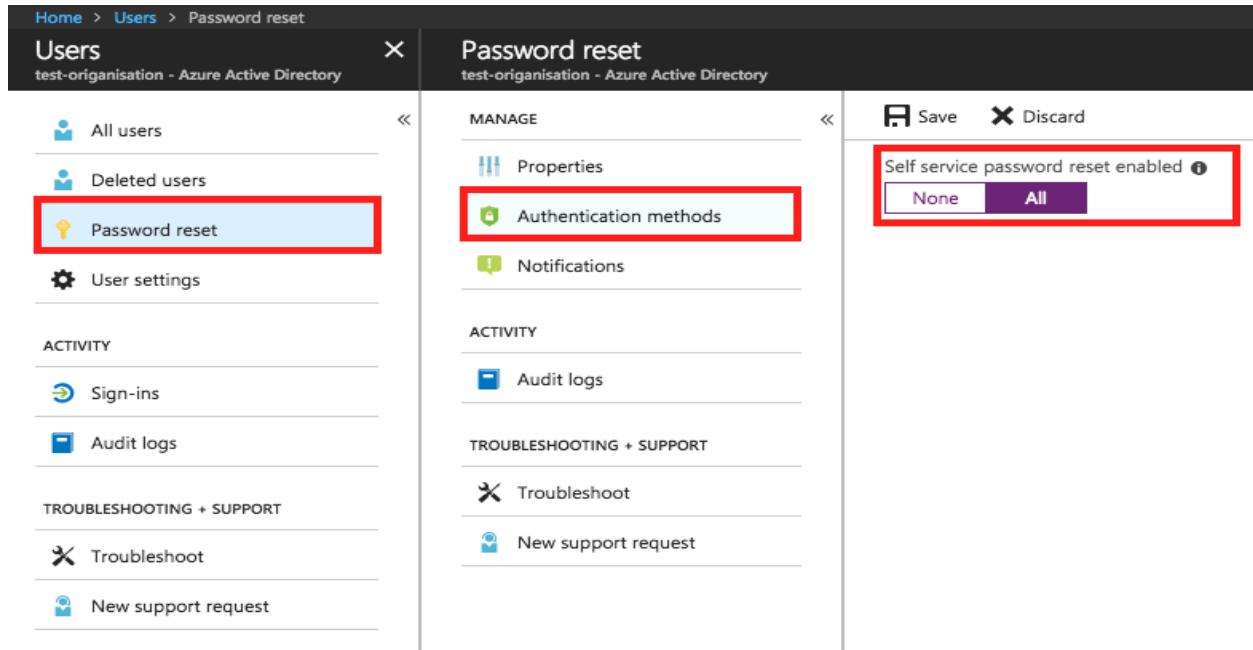
The screenshot shows the Azure Active Directory Conditional access - Policies blade. On the left, there's a sidebar with 'All services', 'Favorites' (Dashboard, Azure Active Directory, Users, Enterprise applications), and a search bar. Under 'SECURITY', 'Conditional access' is highlighted with a red box. Other options include MFA Server, Users flagged for risk, Risky sign-ins, Authentication methods, Sign-ins, Audit logs, Troubleshoot, and New support request. The main area has sections for 'MANAGE' (Named locations, Custom controls (preview), Terms of use, VPN connectivity, Classic policies) and 'TROUBLESHOOTING + SUPPORT' (Troubleshoot, New support request). At the top right, there's a 'New policy' button and a 'What If' link. A purple banner at the top says 'Create your own policies and target specific conditions like Cloud apps, Sign-in risk, and more'.

To find out more about how to define conditions, refer to the following link: <https://docs.microsoft.com/en-us/azure/active-directory/conditional-access/conditions>.

Configuring self-service password reset

Within an organization, IT administrators usually have a couple of regular and important tasks, such as managing users and their identities. One of the best practices for managing cloud security is to enable users to reset their passwords or unlock their accounts by configuring **self-service password reset (SSPR)**.

To enable this feature, go to the Azure AD portal and click on **Users**, where you can see the **Password reset** blade. Set **Self-service password rest enable** as all value in the authentication methods blade and then click on **Save**, as shown in the following screenshot:



Home > Users > Password reset

Users

test-organisation - Azure Active Directory

All users
Deleted users
Password reset
User settings

ACTIVITY

Sign-ins
Audit logs

TROUBLESHOOTING + SUPPORT

Troubleshoot
New support request

Manage

Properties
Authentication methods
Notifications

ACTIVITY

Audit logs

TROUBLESHOOTING + SUPPORT

Troubleshoot
New support request

Save Discard

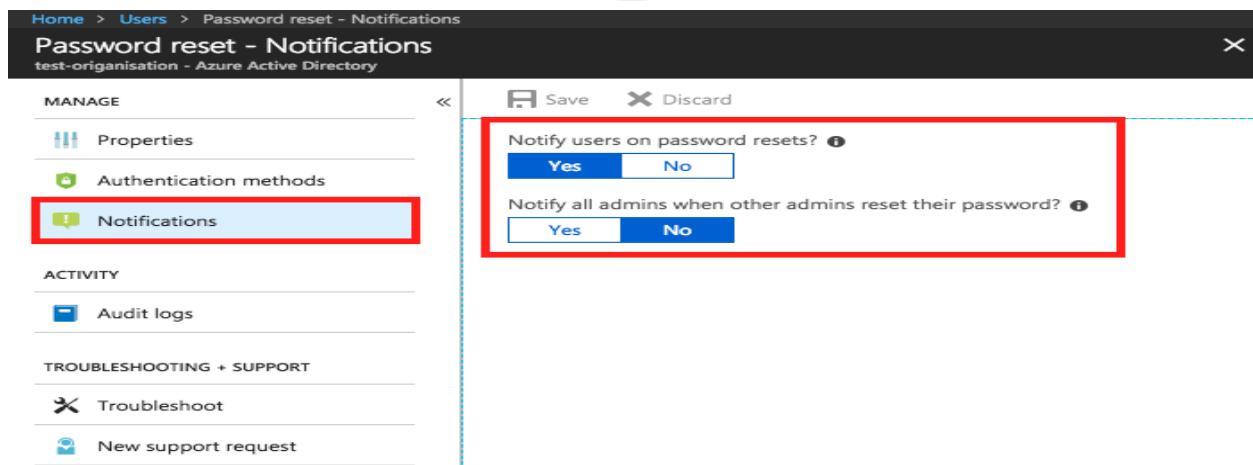
Self-service password reset enabled

None All

After enabling the self-service password reset feature, Go back to the authentication methods blade to configure self-service details such as the number of methods required to reset and the methods available to users, which include mobile phone, office phone, or email. Note that it would be great to also set the notifications to inform administrators and users so that they know the password has been reset successfully.

9963799240 / 7730997544

Go to the **Notifications** blade to configure the notification to users and administrators, as shown in the following screenshot:



Home > Users > Password reset - Notifications

Password reset - Notifications

test-organisation - Azure Active Directory

Manage

Properties
Authentication methods
Notifications

ACTIVITY

Audit logs

TROUBLESHOOTING + SUPPORT

Troubleshoot
New support request

Save Discard

Notify users on password resets?

Yes No

Notify all admins when other admins reset their password?

Yes No

Set notifications to notify users and administrators on password reset

You may be wondering, if we still have our existing on-premises directory, how do we get both Azure AD and our on-premises Active Directory Domain Services (AD DS) environment to be managed in a centralization way? Azure provides a very cool feature, **password writeback**, to help you synchronize password changes in this kind of hybrid scenario. After password writeback has been enabled, as it is a part of Azure AD Connect, it will send password changes back to an existing on-premises directory from Azure AD in a secure way.

To find out more about how to enable password writeback for your hybrid environment, check the following link: <https://docs.microsoft.com/en-us/azure/active-directory/authentication/tutorial-enable-writeback>

Configuring privileged identity management

The Azure AD also provides the Azure AD privileged identity management capability. With this feature, you can manage, control, and monitor on-demand and just-in-time administrative access in Azure AD, Azure Resources, Office 365, Microsoft Intune, or other Microsoft Online Services within an organization. To find out more about how to configure it, checkout the following link: <https://docs.microsoft.com/en-us/azure/active-directory/privileged-identity-management/pim-configure>.

Configuring Azure AD identity management

Azure AD identity protection is a great feature allows you to detect suspicious actions and configure responses automatically to protect the organization's identities. You can enable it by creating Azure AD Identity Protection via Azure Portal.

Refer to the following link to find out more about how to configure Azure AD Identity Management in Azure: <https://docs.microsoft.com/en-us/azure/active-directory/identity-protection/overview>.

Leveraging Microsoft Graph other than Azure AD Graph API

The **Azure Active Directory Graph API (Azure AD Graph API)** provides a programmatic way to access Azure AD using RESTful APIs. Developers can program call the the Azure AD Graph API to perform **create, read, update, and delete(CRUD)** operations on Azure AD data and objects, such as creating a new user in a directory and then getting the target user's detailed properties. Microsoft recommends us to use Microsoft Graph instead of the Azure AD Graph API to access Azure Active Directory objects. **Microsoft Graph** is the API in **Microsoft 365**, which is a more powerful identity API than the AAD Graph API,. It allows to connect to **Office 365, Windows 10, and Enterprise Mobility** in a secure way. To know more about Microsoft Graph, Checkout the following link: <https://developer.microsoft.com/en-us/graph/>.

Integrating applications with Azure AD

Azure AD **Business-to-Consumer (B2C)** is a dedicated Azure AD tenant provided by Microsoft Azure that enables individual customers working with organizations to get access to custom web applications, mobile apps, and desktop applications.

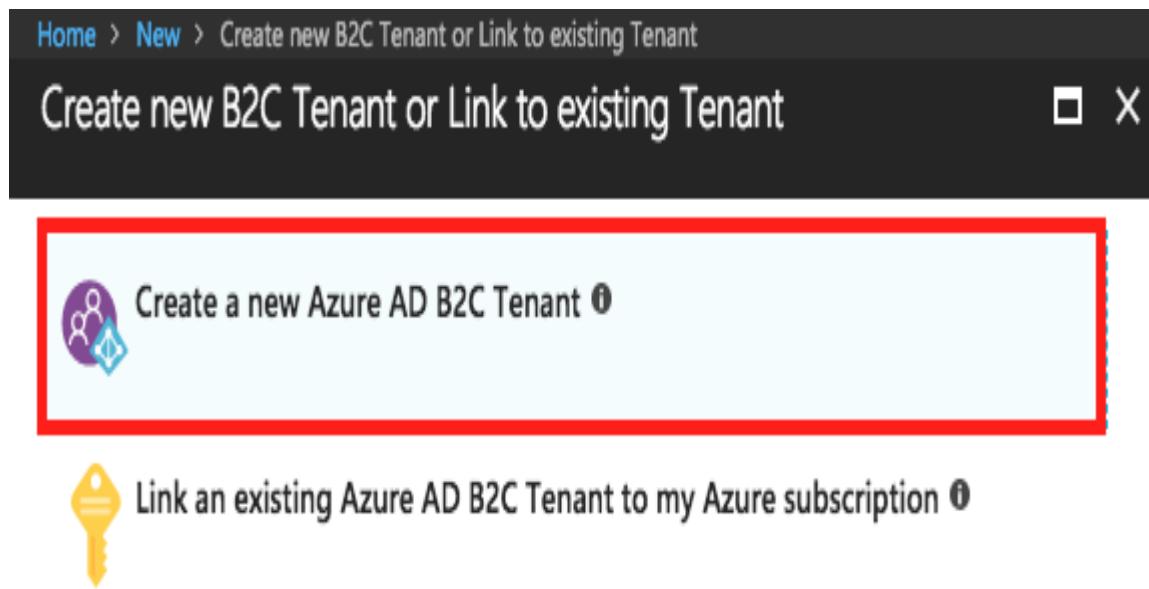
Azure Active Directory B2C helps organizations and enterprises worldwide to connect to their customers and serve their applications with a high level of cloud-based identity protection.

Azure AD B2C supports popular protocols such as **OpenID Connect, OAuth 2.0, and SAML**. The accounts used by Azure AD B2C can be created directly in the Azure B2C tenant or provided by

famous social media identities, such as Facebook, Google, Amazon, LinkedIn, and Twitter.

Creating an Azure AD B2C directory

To create a Azure AD B2C directory, go to the Azure Portal and click on **Create a new resource**, then choose **Identity repository** and **Azure Active Directory B2C** and click on **Create**; you can create a new Azure AD B2C Tenant or link an existing Azure AD B2C Tenant to a Azure Subscription, as shown here:



Here, we'll create a new Azure AD B2C tenant. To create a new Azure AD directory, fill in the organization name and initial domain name, then choose a location before clicking on **Create**, as shown in the following screenshot:

Azure AD B2C Create Tenant □ X

* Organization name i
 ✓

* Initial domain name i
 .onmicrosoft.com

Country or region i

 Directory creation will take about one minute.

Create

After creating an Azure AD B2C tenant, you can link it to an Azure subscription by choosing it and filling in other information such as resource group, subscription, and location, then click on **Create**. Azure will manage the remaining work:

Home > New > Create new B2C Tenant or Link to existing Tenant > Azure AD B2C Resource

Create new B2C Tenant or Link to existing Tenant X

 Create a new Azure AD B2C Tenant i

 Link an existing Azure AD B2C Tenant to my Azure subscription i

Azure AD B2C Resource X

* Azure AD B2C Tenant i

Azure AD B2C Resource name

* Subscription

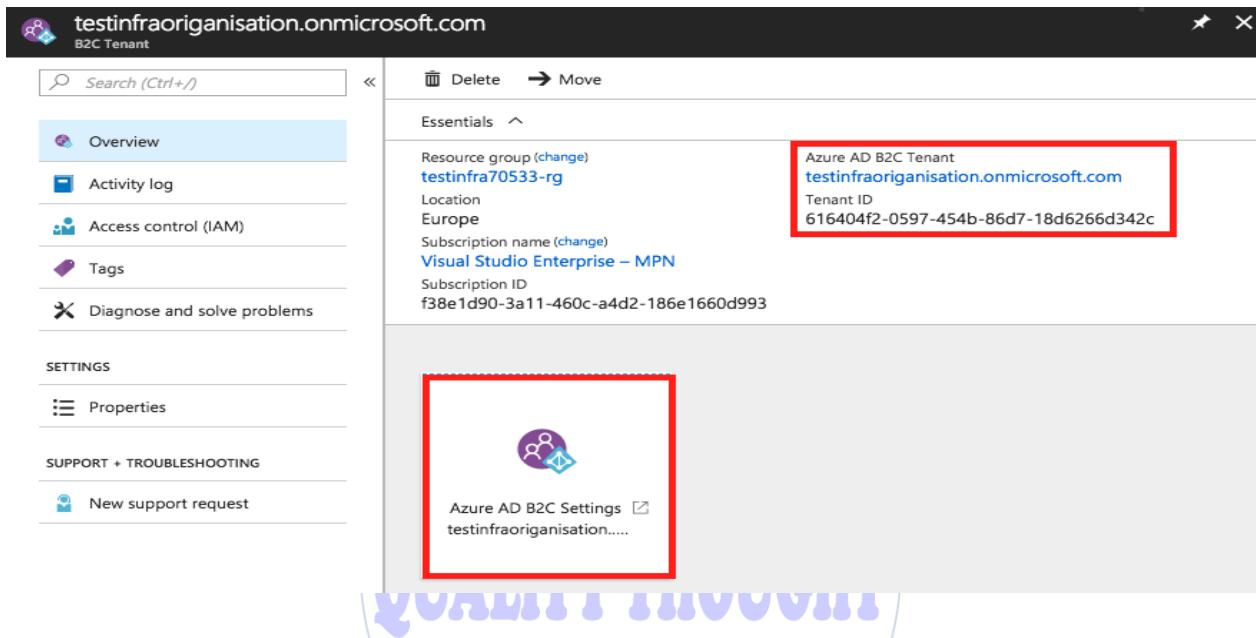
* Resource group
 Create new Use existing
 ✓

Resource group location

Create

Managing Azure AD B2C directory

After creating a B2C directory, go to the resource and check the **Overview** blade, where you can get overall information on the B2C tenant and click on **Azure AD B2C Settings** to manage this tenant, as shown in the following:



testinfraorganisation.onmicrosoft.com

B2C Tenant

Search (Ctrl+ /)

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

SETTINGS

Properties

SUPPORT + TROUBLESHOOTING

New support request

Delete Move

Essentials

Resource group (change)
testinfra70533-rg

Location
Europe

Subscription name (change)
Visual Studio Enterprise – MPN

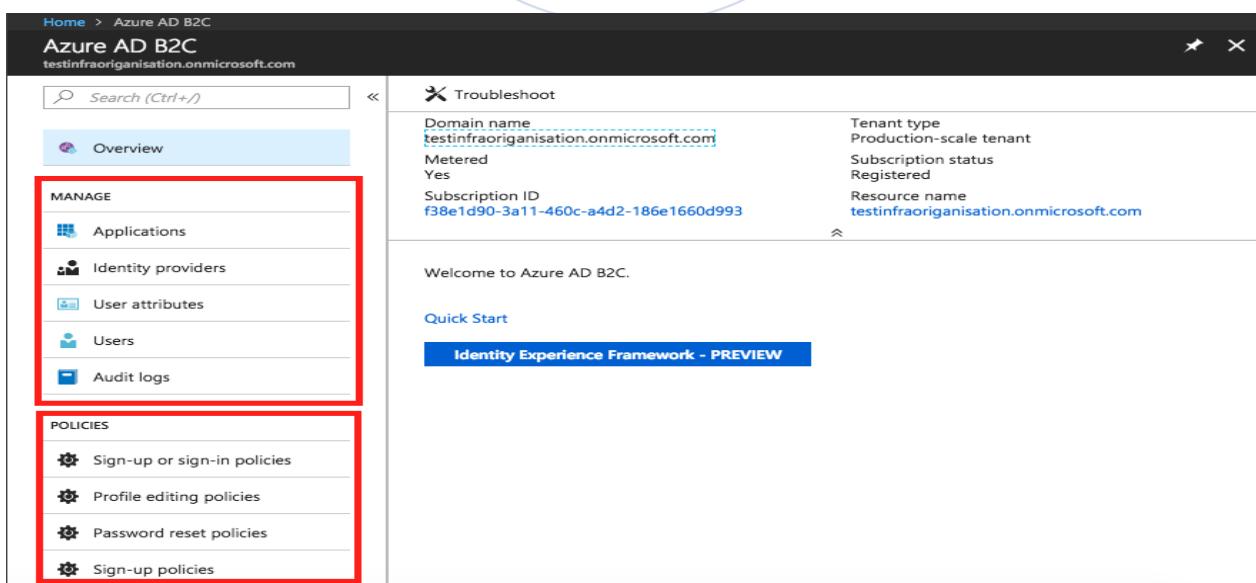
Subscription ID
f38e1d90-3a11-460c-a4d2-186e1660d993

Azure AD B2C Tenant
testinfraorganisation.onmicrosoft.com

Tenant ID
616404f2-0597-454b-86d7-18d6266d342c

Azure AD B2C Settings testinfraorganisation.....

You'll see the following page after clicking on **Azure AD B2C Settings**, where we can manage all the users and groups in this directory, as well the application linked to the current directory:



Home > Azure AD B2C

Azure AD B2C

testinfraorganisation.onmicrosoft.com

Search (Ctrl+ /)

Overview

MANAGE

- Applications
- Identity providers
- User attributes
- Users
- Audit logs

POLICIES

- Sign-up or sign-in policies
- Profile editing policies
- Password reset policies
- Sign-up policies

Troubleshoot

Domain name
testinfraorganisation.onmicrosoft.com

Metered
Yes

Subscription ID
f38e1d90-3a11-460c-a4d2-186e1660d993

Tenant type
Production-scale tenant

Subscription status
Registered

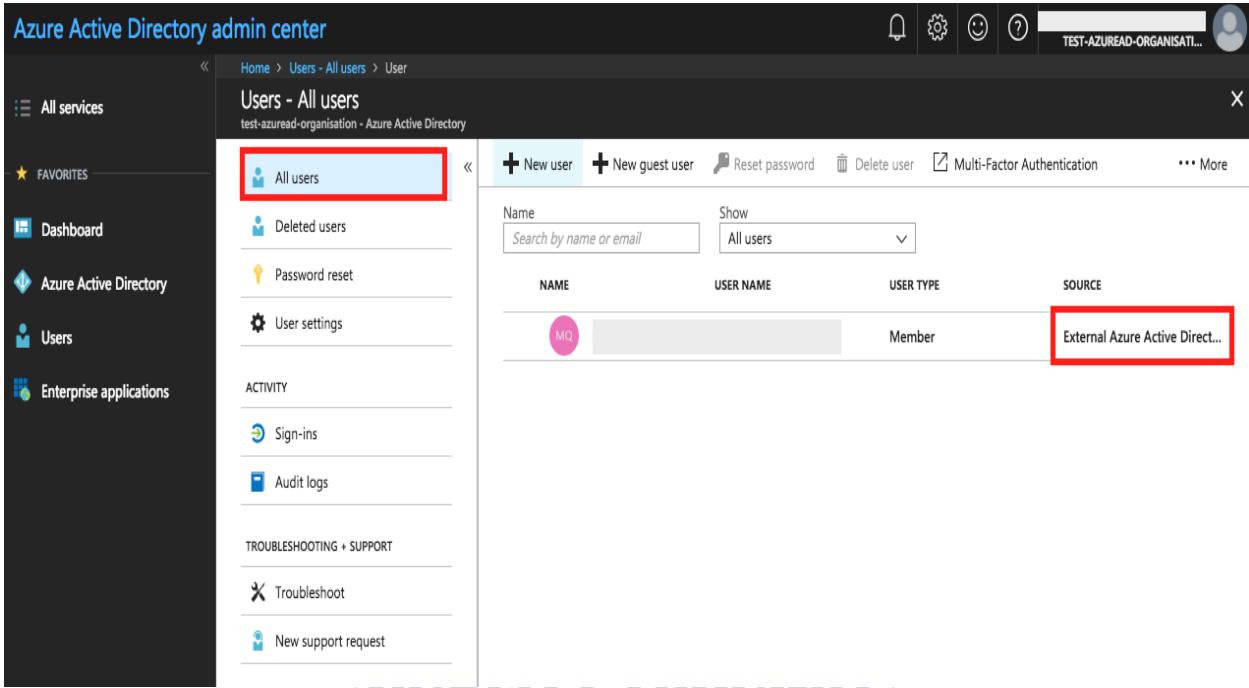
Resource name
testinfraorganisation.onmicrosoft.com

Welcome to Azure AD B2C.

Quick Start

Identity Experience Framework - PREVIEW

If we go back to the Azure AD portal the current Azure B2C directory and click on **All users**, you can see that the users in this directory are referenced as **External Azure Active Directory source**:



The screenshot shows the Azure Active Directory admin center interface. On the left, there's a sidebar with various service links like Dashboard, Azure Active Directory, Users, and Enterprise applications. The main area is titled 'Users - All users' and shows a list of users. One user, 'MQ', is listed with the 'User type' as 'Member' and the 'Source' as 'External Azure Active Direct...'. Both the 'All users' link in the sidebar and the 'External Azure Active Direct...' text in the user list are highlighted with red boxes.

Azure B2C also supports the use of built-in policies to create a wonderful login experience within minutes. It can build custom policies and integrate with CRMs, databases, marketing analytics tools, and other account verification systems.

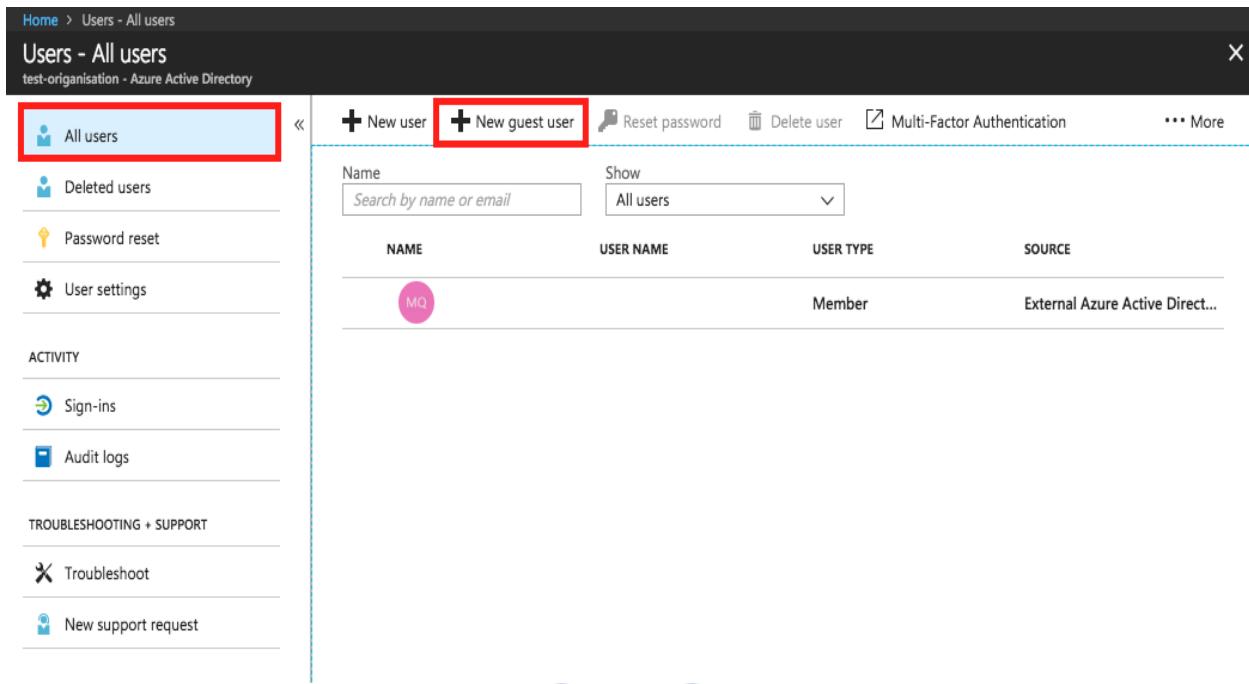
If you want to know more about how to use built-in policies with Azure B2C, which can be applied in most general scenarios, check out the following link: <https://docs.microsoft.com/en-us/azure/active-directory-b2c/active-directory-b2c-reference-policies>.

For more complex scenarios, you can use custom policies, which are still in preview when this book was written. You can get more information from here: <https://docs.microsoft.com/en-us/azure/active-directory-b2c/active-directory-b2c-get-started-custom>.

Implementing Business to Business (B2B) collaboration

Azure AD **Business to Business (B2B)** collaboration capabilities enable any organization using Azure AD to work safely and securely with users from other partner organizations, using school or work email.

You can enable B2B collaboration by adding guest users to your organization in Azure AD as shown here:



The screenshot shows the Azure AD 'Users - All users' interface. On the left, there's a sidebar with links like 'All users' (which is highlighted with a red box), 'Deleted users', 'Password reset', 'User settings', 'Sign-ins', 'Audit logs', 'Troubleshoot', and 'New support request'. The main area has a header with 'New user' and 'New guest user' buttons, both of which are highlighted with red boxes. Below the header is a search bar and dropdown for 'Show'. A table lists a single user entry: Name 'MQ', User Name 'MQ', User Type 'Member', and Source 'External Azure Active Direct...'. The 'New guest user' button is located at the top right of the main content area.

Then, you'll see a **New Guest User** page, where you can add the email address of your partner and add a message in the email invitation, then click on **Invite**

Once you click on **Invite**, users will receive an invitation with a redemption URL, and then they can review and accept the privacy terms.

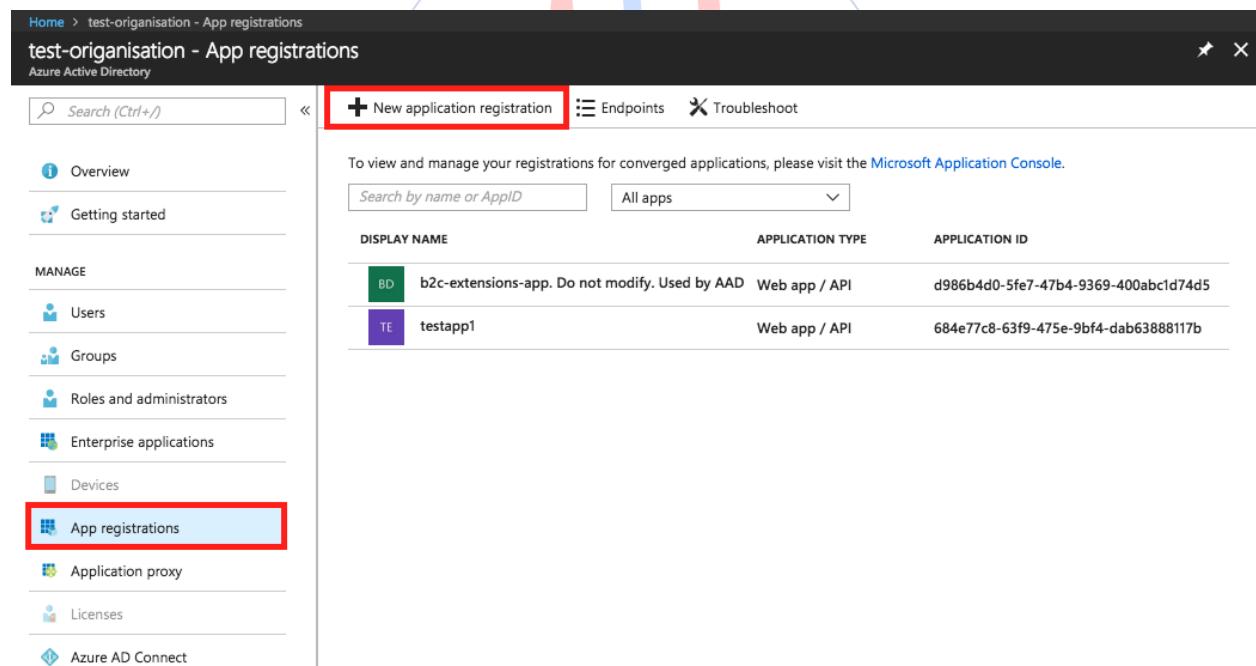
It is also possible to add B2B collaboration guest users without an invitation; check out the following link for more information: <https://docs.microsoft.com/en-us/azure/active-directory/b2b/add-user-without-invite>.

Integrating applications with Azure AD

There are millions of enterprise-level **software-as-a-service (SaaS)** applications boosting IT market. Application providers and enterprise developers want a way to integrate line-of-business applications with a cloud-based IDaaS solution such as Azure AD to provide secure sign-in, authorization, and SSO capabilities.

In Azure, to integrate a web application with Azure AD must first be registered in an Azure AD tenant. The registration process includes giving Azure AD a URL to specify where it's located and to send replies after a user is authenticated.

To register an application, go to the Azure AD portal and click on **App registrations**, and then register a new application by clicking **+ New application registration**, as shown here:



The screenshot shows the Azure Active Directory portal with the following details:

- Header:** Home > test-organisation - App registrations
- Section:** test-organisation - App registrations
- Actions:** + New application registration, Endpoints, Troubleshoot
- Message:** To view and manage your registrations for converged applications, please visit the Microsoft Application Console.
- Search:** Search by name or AppID, All apps
- Table:** Displays registered applications:

DISPLAY NAME	APPLICATION TYPE	APPLICATION ID
b2c-extensions-app. Do not modify. Used by AAD	Web app / API	d986b4d0-5fe7-47b4-9369-400abc1d74d5
testapp1	Web app / API	684e77c8-63f9-475e-9bf4-dab63888117b
- Left Sidebar (highlighted with a red box):**
 - Overview
 - Getting started
 - MANAGE
 - Users
 - Groups
 - Roles and administrators
 - Enterprise applications
 - Devices
 - App registrations** (selected)
 - Application proxy
 - Licenses
 - Azure AD Connect

In the **Create** form, you should choose the **Application type**, which is **Web app / API or local app**, then fill in the **Sign-on URL** and click on **Create**:

After a few seconds, you can see your application has been registered successfully. As shown in the following screenshot, **Application ID** is your application identity, which is known by the Azure AD tenant.

Microsoft also recommends to use Azure AD **Managed Service Identity (MSI)** as your application identity, as it simplifies creating an identity for code. To know more about Azure AD Managed Service Identity, you can check the following link: <https://docs.microsoft.com/en-us/azure/active-directory/managed-service-identity/overview>.

If you have an existing application that has its own account system and it may require to support other kinds of sign-ins and from other cloud providers, in that case, you may need to sign in any Azure AD user by using the multitenant application pattern. Check out the following link to have more information: <https://azure.microsoft.com/fr-fr/resources/samples/active-directory-dotnet-native-desktop/>.

You can also integrate Windows desktop applications, universal application with Azure AD by using ARM template. The following are some excellent sample ARM templates on GitHub which can help you:

[9963799240 / 7730997544](#)

Ameerpet / Kondapur

Hyderabad

- Integrating Azure AD into a Windows desktop application using interactive authentication:

<https://azure.microsoft.com/fr-fr/resources/samples/active-directory-dotnet-native-desktop/>

- Integrating a Windows Universal application with Azure AD:

<https://azure.microsoft.com/en-gb/resources/samples/active-directory-dotnet-windows-store/>

To integrate Azure AD with a web application using OpenID connect and WS-Federation, it is also possible to use an ARM template, check out the following links:

- Integrate Azure AD into a web application using OpenID Connect:

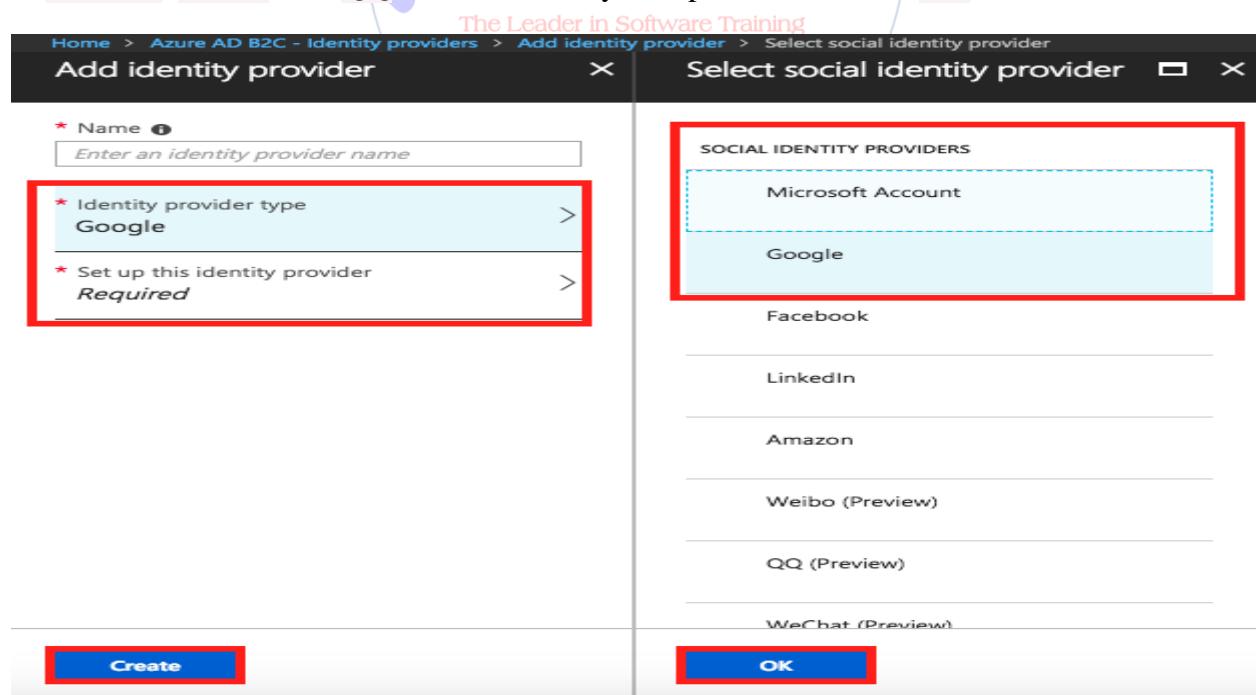
<https://azure.microsoft.com/fr-fr/resources/samples/active-directory-dotnet-webapp-openidconnect/>

- Integrating a web app with Azure AD using WS-Federation:

<https://azure.microsoft.com/en-gb/resources/samples/active-directory-dotnet-webapp-wsfederation/>

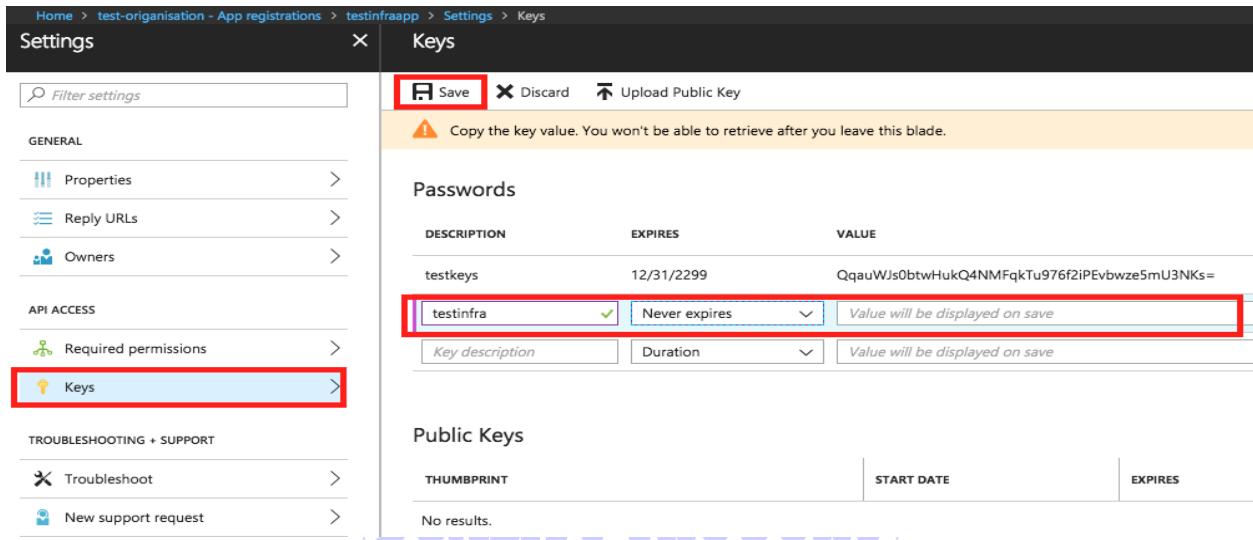
Implementing federation and social identity provider authentication

Most websites may have different customers across different social media provided to configure federation with public consumer identity providers such as Facebook, Google, and Twitter. Go to your Azure AD B2C tenant and click on **Identity providers** to add an identity provider, as shown in the following screenshot; you have a wide range of identity providers to choose from, such as **Microsoft Account**, **Google**, **Facebook**, **Twitter**, and **LinkedIn**. There are also some popular Chinese social media organizations (currently in preview), such as **WeChat**, **Weibo**, and **QQ**. Choose which you expect and click on **OK**:



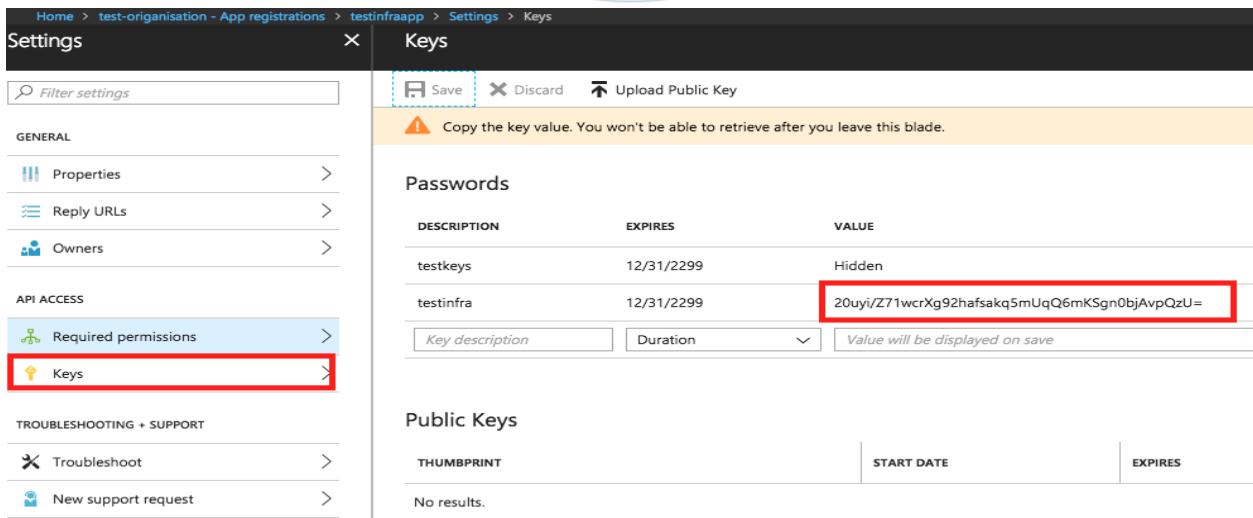
The screenshot shows two adjacent browser windows. The left window is titled 'Add identity provider' and displays fields for 'Name' (with placeholder 'Enter an identity provider name') and 'Identity provider type' (set to 'Google'). A red box highlights these fields. Below them is another section for 'Set up this identity provider' with the note 'Required'. The right window is titled 'Select social identity provider' and lists various providers under 'SOCIAL IDENTITY PROVIDERS': Microsoft Account, Google, Facebook, LinkedIn, Amazon, Weibo (Preview), QQ (Preview), and WeChat (Preview). The 'Microsoft Account' option is highlighted with a red box. At the bottom of each window is a blue 'Create' or 'OK' button.

After selecting the identity provider type, you may also need to set up this identity provider using the Application ID that we mentioned in the previous section while registering a web application in the Azure AD tenant. You also need an application secret to set this up. We can get this secret by clicking on **Settings** in the registered application, and then go to the **Keys** blade and file in the description of keys and the expiration period, as shown in the following screenshot:



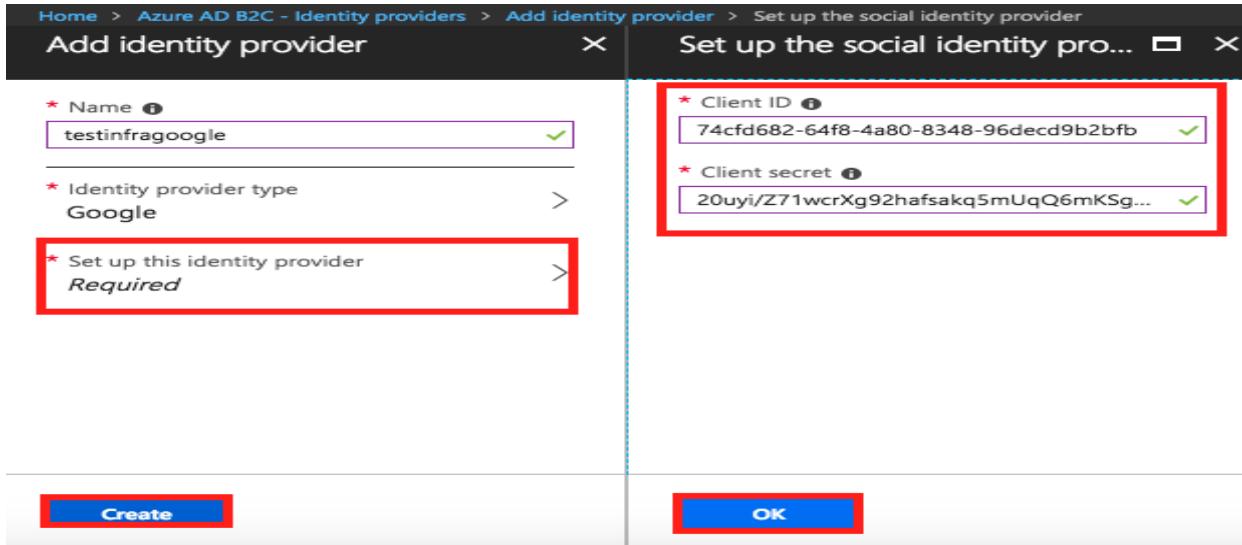
DESCRIPTION	EXPIRES	VALUE
testkeys	12/31/2299	QqauWJs0btwHukQ4NMFqkTu976f2iPEvbwze5mU3NKs=
testinfra	Never expires	Value will be displayed on save

After filling in the description and duration, where you can choose 1 year, 2 years, or never, click on **Save**; after a few seconds, the value of your key will appear, as shown here, you can copy this value to use in the next step: **Hyderabad**

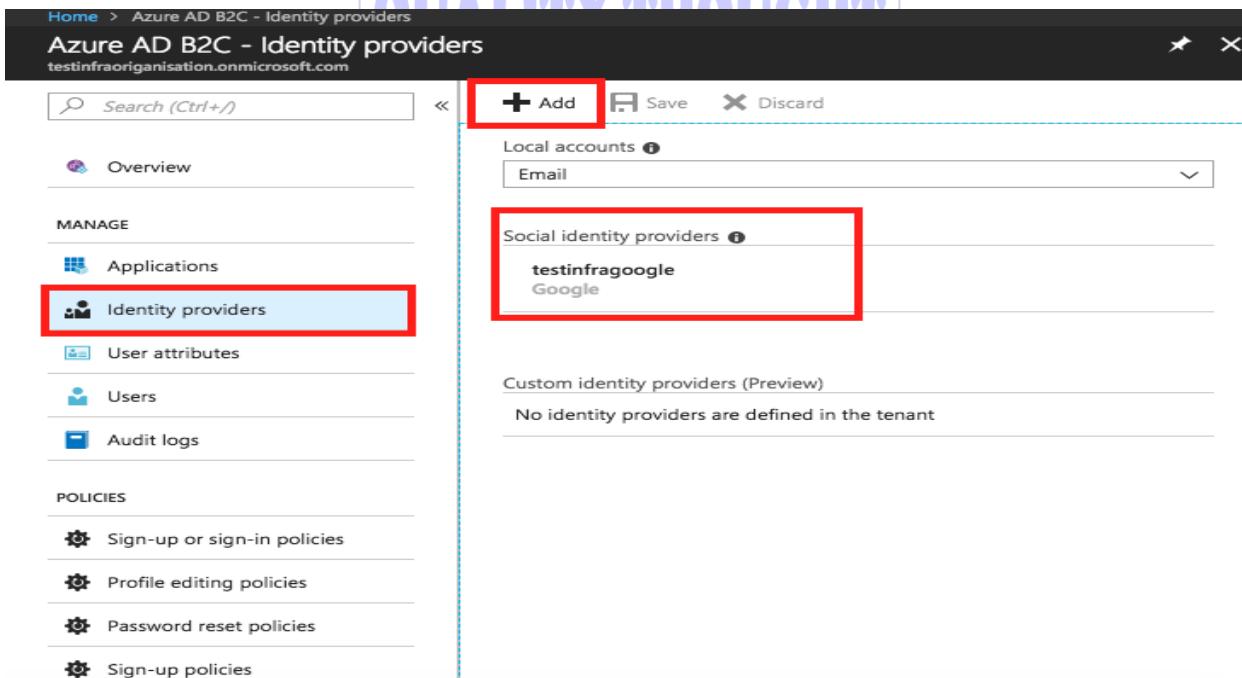


DESCRIPTION	EXPIRES	VALUE
testkeys	12/31/2299	Hidden
testinfra	12/31/2299	20uyi/Z71wcrXg92hafsakq5mUqQ6mKSgn0bjAvpQzU=

Here, use your Application ID to fill in **Client ID** and the copied secret in the **Client secret** field, then click on **OK** to finish the setup of the identity provider:



Go back to the **Identity providers** page, where you can see that the Google identity provider has been created successfully, as follows:



Identity provider	Provider type
testinfragoogle	Google

Configuring SAML-based SSO for an application with Azure AD

SSO means being able to access all the applications and resources that you need by signing in only once and once, signed in, you can access all of the applications with a single account and without typing in a password for a second time.

Check out the following link to configure SAML-based SSO for an application with Azure AD: <https://docs.microsoft.com/en-us/azure/active-directory/manage-apps/configure-single-sign-on-portal>.

Managing hybrid identities

At the transition stage of a cloud-first digital transformation, we need to have a cloud-based identity solution that will work in different hybrid scenarios. Azure AD services provide the following capabilities:

- Using **Azure AD Connect** to integrate your on-premises Windows Server Active Directory with Azure AD
- Using **Azure AD Domain Services (AD DS)** to provide managed domain services, such as domain join, group policy, LDAP, and Kerberos/NTLM
- Using **Active Directory Federation Services (AD FS)** to federate identities and more advanced scenarios, such as SSO or on-premises MFA

Configuring Azure AD Connect and synchronization services

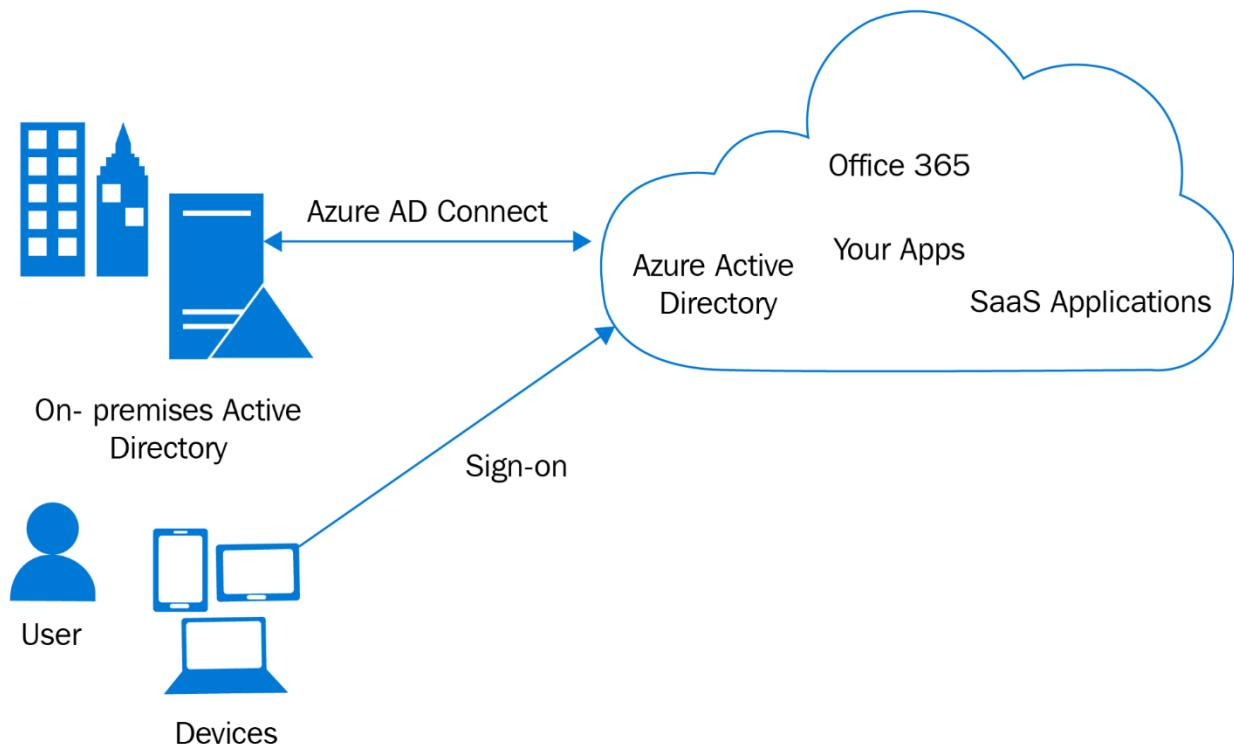
Azure AD Connect is an excellent way to connect the on-premises identity directory with Azure AD and Office 365, as well as the other SaaS applications integrated with Azure AD.

Azure AD Connect is made up of the following three primary components:

- **Azure AD Connect sync** is an **Azure Active Directory Connect synchronization** services in charge of all operations related to synchronizing identity data between your on-premises environment and Azure AD. It replaces old connectors, such as DirSync and Azure AD Sync.

- **Active Directory Federation Services** component, which is optional and provides the cloud identity federation feature.
- The **Azure AD Connect Health** is a monitoring component to help users gain insights on their on-premises identity infrastructure.

To show how Azure AD Connect can work, is as referred in the following schema, Azure AD connect on-premise Active Directory and Azure Active Directory to provide a friendly identity experience to users:



Go the Azure AD portal and click on the **Azure AD Connect** blade to configure all these features:

Home > test-organisation - Azure AD Connect

test-organisation - Azure AD Connect

Azure Active Directory

Search (Ctrl+)

- Roles and administrators
- Enterprise applications
- Devices
- App registrations
- Application proxy
- Licenses
- Azure AD Connect
- Custom domain names
- Mobility (MDM and MAM)
- Password reset
- Company branding
- User settings
- ...

Troubleshoot
Refresh

i Use Azure AD Connect to integrate your Azure AD with your Windows Server AD or another directory on your network.

SYNC STATUS	
↻ Not Installed	Download Azure AD Connect
Last Sync	Sync has never run
Password Hash Sync	Disabled

USER SIGN-IN	
➔ Federation	Disabled
Seamless single sign-on	Disabled
Pass-through authentication	Disabled
	0 domains
	0 agents

ON-PREMISES APPLICATIONS

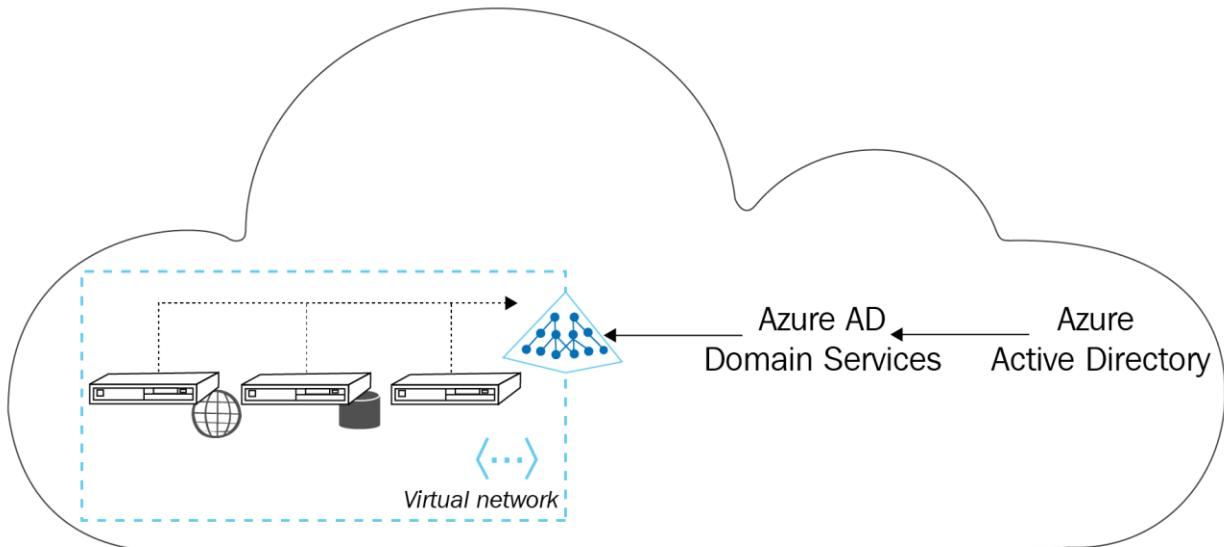
Looking to configure remote access for on-premises applications? [Head to Application Proxy](#)

Download AD Connect from the following link: <https://www.microsoft.com/en-us/download/details.aspx?id=47594>.

To find out more about how to configure Azure Connect's synchronization and federation features, check out the following link: <https://docs.microsoft.com/en-us/azure/active-directory/connect/active-directory-aadconnect>.

Managing domains with Azure AD domain services

Azure AD domain services provide managed domain services in an Azure virtual network, such as domain join, group policy, LDAP, and Kerberos/NTLM authentication that can integrate Windows Server Active Directory with Azure AD:



AD DS stores information about usernames, passwords, phone numbers, and so on, and guarantees that other users can access this information under authorization.

If you want to join your on-premises Active Directory domain-joined devices to Azure AD, you can accomplish this by referencing the following link to configure hybrid Azure AD joined devices step by step: <https://docs.microsoft.com/en-us/azure/active-directory/devices/hybrid-azuread-join-manual-steps>.

9963799240 / 7730997544

Before Azure AD DS, we generally use an S2S VPN connection or ExpressRoute to connect on-premises identity server to Azure cloud, or may deploy an Azure VM to run Active Directory. Another way is to use Dir sync to synchronize on-premises identity with Azure cloud. Until now, using Azure AD Connect is still an alternative way, both Azure AD connect and AD DS can most adaptive in different scenarios, you can know more about how to compare them by checking the link below: <https://docs.microsoft.com/en-us/azure/active-directory-domain-services/active-directory-ds-compare-with-azure-ad-join>.

Implementing SSO in hybrid scenarios

To implement SSO in a hybrid environment, you can use **seamless SSO**; it will need the user's device to be domain-joined. Check out the link <https://docs.microsoft.com/en-us/azure/active-directory/connect/active-directory-aadconnect-sso> to know how to configure it step by step.

SSO and secure remote access for web applications hosted on-premises can be implemented using **Azure AD Application Proxy**. Azure AD Application Proxy is a lightweight agent that facilitates the flow of traffic from the Application Proxy service in the cloud to your on-premises environment.

To know more about how to enable Azure AD Application Proxy, check out the following link: <https://docs.microsoft.com/en-us/azure/active-directory/manage-apps/application-proxy-enable>.

Monitoring on-premises identity infrastructure and synchronization services

Azure provides a couple of components to help users monitor their on-premises identity infrastructure and synchronization services such as Office 365 and other Microsoft Online Services.

The key identity components are the following:

- **Active Directory Federation Services (AD FS)** servers
- Azure AD Connect servers (which contain a Sync Engine)
- Active Directory domain controllers, and so on.

Monitoring these key components will help an administrator or enterprise architect to make informed decisions. You should download and install Azure AD Connect Health Agents to get health and usage information about your on-premises services.

You can download these agents from the following links:

- Downloading Azure AD Connect Health Agent for AD FS: <http://go.microsoft.com/fwlink/?LinkId=518973>
- Downloading Azure AD Connect (configures Azure AD Connect Health agent for sync): <http://go.microsoft.com/fwlink/?linkid=615771>
- Downloading Azure AD Connect Health Agent for AD DS: <http://go.microsoft.com/fwlink/?LinkId=820540>

You can also get information on configuring monitoring for these components by checking out the following links:

- Monitoring AD FS using Azure AD Connect Health: <https://docs.microsoft.com/en-us/azure/active-directory/connect-health/active-directory-aadconnect-health-adfs>
- Monitoring Azure AD Connect sync with Azure AD Connect Health: <https://docs.microsoft.com/en-us/azure/active-directory/connect-health/active-directory-aadconnect-health-sync>
- Using Azure AD Connect Health with AD DS: <https://docs.microsoft.com/en-us/azure/active-directory/connect-health/active-directory-aadconnect-health-adds>

You can see the views of alerts, performance monitoring, usage analytics, and other information in one place using the Azure AD Connect Health portal. Here is the URL: <https://aka.ms/aadconnecthealth>.

Planning and Implementing Azure Storage, Backup, and Recovery Services

9963799240 / 7730997544

Ameerpet / Kondapur

Hyderabad

Microsoft Azure provides different storage capabilities; there are four core storage services such as blobs, tables, queues, and file shares. Otherwise, Microsoft Azure also provides hybrid storage solutions such as StorSimple as well as cross-premise transfer options. It also offers capabilities to facilitate recovery and to assist customers with implementing their **business continuity and disaster recovery (BCDR)** strategy using Azure Backup and Azure Site Recovery.

In this chapter, we'll cover the following topics:

- Implementing and managing Azure Storage services
- Configuring Azure **Content Delivery Network (CDN)**
- Introducing Azure data storage services

- Implementing the BCDR strategy with Azure Backup and Azure Site Recovery
- Introducing Azure StorSimple and other Azure Hybrid storage

Implementing and managing Azure Storage

Microsoft Azure provides various storage options that allow users to store files, messages, tables, and any other type of information; data stored in Azure Storage can be used by web applications, mobile apps, desktop applications, and various types of custom solution. From a conceptual point of view, Azure Storage options are applied in the following scenarios:

- Object-based storage for virtual machines such as Azure Blobs and file shares
- Semi-structured data storage such as table storage
- Storing or processing large numbers of messages using queue storage
- Hyper-scale repository for big data analytic workloads using Azure Data Lake Store

An overview of Azure Storage services

Azure Storage is a managed service within Azure which aims to provide cloud-based storage, which is secure and scalable, with different levels of availability. Azure Storage includes the following data services:

- **Azure Blobs** is an object-based storage for text and binary data
- **Azure Files** is a managed file shared storage for cloud-based as well as on-premise deployments
- **Azure Queues** stores and exchanges messages between Azure services and applications
- **Azure Tables** is designed for massive semi-structured data or NoSQL storage, which is the cheapest in Azure in comparison with Cosmos DB

Implementing Azure Storage services

To implement Azure Storage services, we should start by creating a storage account. A **storage account** is a unique namespace, where users can store and access data objects in cloud storage.

There are three kinds of storage account, which are explained as follows:

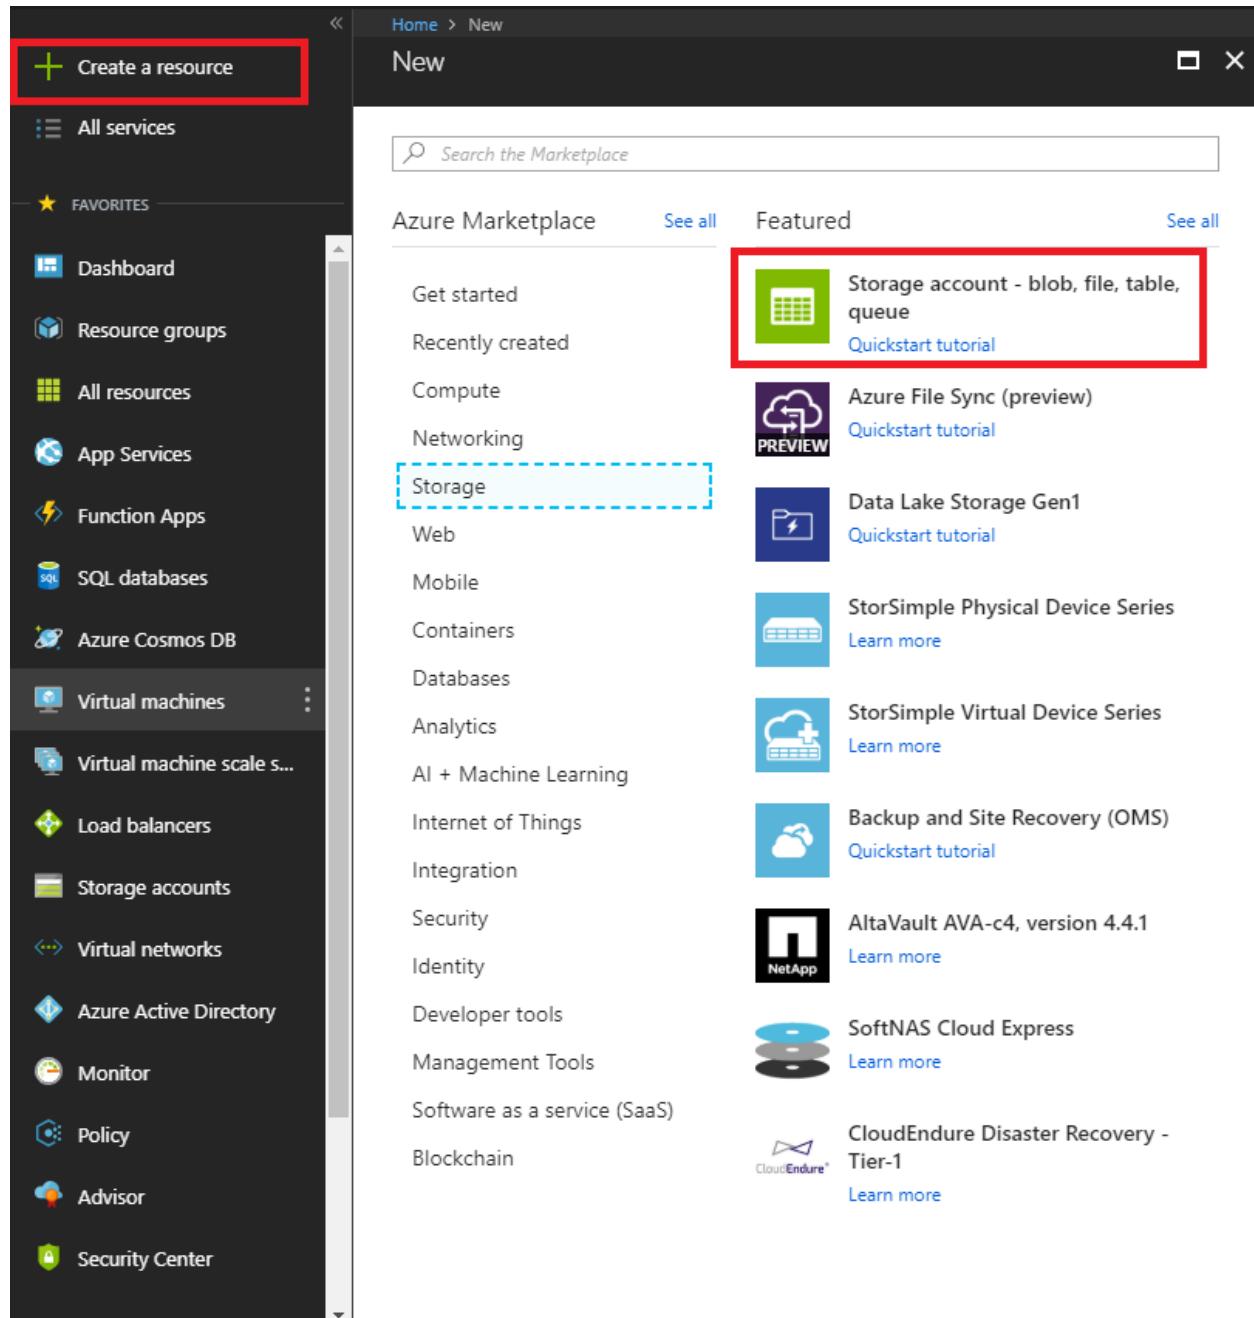
- **General-purpose v1 (GPv1):** This is a unique namespace to host blobs, tables, queues, and files without access tier support
- **General-purpose v2 (GPv2):** This is a unique namespace to host blobs, tables, queues, and files, and supports new features such as access tiers
- **Blob storage:** This is a unique namespace to host blobs

The general quota for storage accounts is limited to 200 in a single Azure subscription. Users can request to increase the quota by contacting Azure support.

Creating a storage account

To create a storage account via the Azure portal, click on **Create a resource**, then you'll find **Storage account**, as shown in the following screenshot:





The screenshot shows the Microsoft Azure 'New' blade. On the left, there's a sidebar with a 'Create a resource' button (highlighted with a red box) and a list of services including Dashboard, Resource groups, All resources, App Services, Function Apps, SQL databases, Azure Cosmos DB, Virtual machines, Virtual machine scale sets, Load balancers, Storage accounts, Virtual networks, Azure Active Directory, Monitor, Policy, Advisor, and Security Center. The 'Storage' category under 'Virtual machines' is highlighted with a dashed blue box. The main area displays the 'Azure Marketplace' with various service offerings. One item, 'Storage account - blob, file, table, queue' (with a green icon), is highlighted with a red box. Other items listed include Get started, Recently created, Compute, Networking, Web, Mobile, Containers, Databases, Analytics, AI + Machine Learning, Internet of Things, Integration, Security, Identity, Developer tools, Management Tools, Software as a service (SaaS), Blockchain, Azure File Sync (preview), Data Lake Storage Gen1, StorSimple Physical Device Series, StorSimple Virtual Device Series, Backup and Site Recovery (OMS), AltaVault AVA-c4, version 4.4.1, SoftNAS Cloud Express, and CloudEndure Disaster Recovery - Tier-1.

Click on **Storage account**, fill in a **Name**, and choose a type of storage account that is compatible with Azure resource in classic model and the Azure Resource Manager model.

You can choose a type of storage account and set a replication method for the storage as shown in the following screenshot:

- [+ Create a resource](#)
- [All services](#)
- [★ FAVORITES](#)
- [All resources](#)
- [Resource groups](#)
- [App Services](#)
- [SQL databases](#)
- [SQL data warehouses](#)
- [Azure Cosmos DB](#)
- [Virtual machines](#)
- [Virtual machine scale sets](#)
- [Load balancers](#)
- [Storage accounts](#)
- [Virtual networks](#)
- [Azure Active Directory](#)
- [Monitor](#)
- [Advisor](#)
- [Security Center](#)
- [Cost Management + B...](#)
- [Help + support](#)
- [Network Watcher](#)

Create storage account

The cost of your storage account depends on the usage and the options you choose below.
[Learn more](#)

* Name [i](#)
testinfrasa ✓ .core.windows.net

Deployment model [i](#)
[Resource manager](#) [Classic](#)

Account kind [i](#)
[Storage \(general purpose v1\)](#)

* Location
West Europe

Replication [i](#)
[Read-access geo-redundant storage \(RA-GRS\)](#)

Performance [i](#)
[Standard](#) [Premium](#)

* Secure transfer required [i](#)
[Disabled](#) [Enabled](#)

* Subscription
Visual Studio Enterprise – MPN

* Resource group
 Create new Use existing
testinfra70533 ✓

Virtual networks
Configure virtual networks [i](#)
[Disabled](#) [Enabled](#)

Pin to dashboard

[Create](#) [Automation options](#)

There are currently four replication methods in Azure Storage:

- **Locally-redundant storage (LRS):** The target is to provide at least 99.999999999% (11 times 9) durability of objects over a given year by replicating your data within a datacenter in the region. It is supported by the GPv1, and GPv2, blob storage account types.
- **Zone-redundant storage (ZRS):** The target is to offer durability for storage objects of at least 99.9999999999% (12 times 9) over a given year; this option replicates data in the storage account synchronously across three storage clusters, which is physically separated from the others and resides in its own **availability zone (AZ)** in a single region.
- **Geo-redundant storage (GRS):** This is designed to provide at least 99.999999999999% (16 times 9) durability of objects over a given year; this option replicates data to a secondary region in the case of failure of the primary region, which is not recoverable from a regional outage or a disaster. However, the data in the secondary region is available to be read only when Microsoft initiates a failover from the primary to the secondary region.
- **Read-access geo-redundant storage (RA-GRS):** This is the same as GRS, but provides read-only access to the data in the secondary location as well as the primary region. This means users can read from the secondary region regardless of whether Microsoft initiates a failover from the primary to the secondary region.

Note that the primary region for the account is chosen by users while creating a new storage account; however, the paired secondary region cannot be changed, and is determined by Microsoft. Check the following link to know more about paired regions in Azure: <https://docs.microsoft.com/en-us/azure/best-practices-availability-paired-regions>.

You can choose a **Replication** method that is most compatible with your requirements. If you choose the GPv2 account type, you can choose the access tier type as shown in the following screenshot:

Home > New > Create storage account

Create storage account

Deployment model Resource manager Classic

Account kind StorageV2 (general purpose v2)

* Location
West Europe

Replication Locally-redundant storage (LRS)

Performance Standard Premium

Access tier (default) Hot

* Secure transfer required Disabled Enabled

Pin to dashboard

Create Automation options

Choosing access tier for the storage account

There are two types of access tier, based on the data access frequency:

- **Hot access tier:** This is for data most frequently accessed.
- **Cool access tier:** This is for data less frequently accessed.
- **Archive access tier:** This is optimized for storing data that is rarely accessed. It can be stored for 180 retention days. Archive storage has the lowest storage costs but involves more costs to retrieve the data.

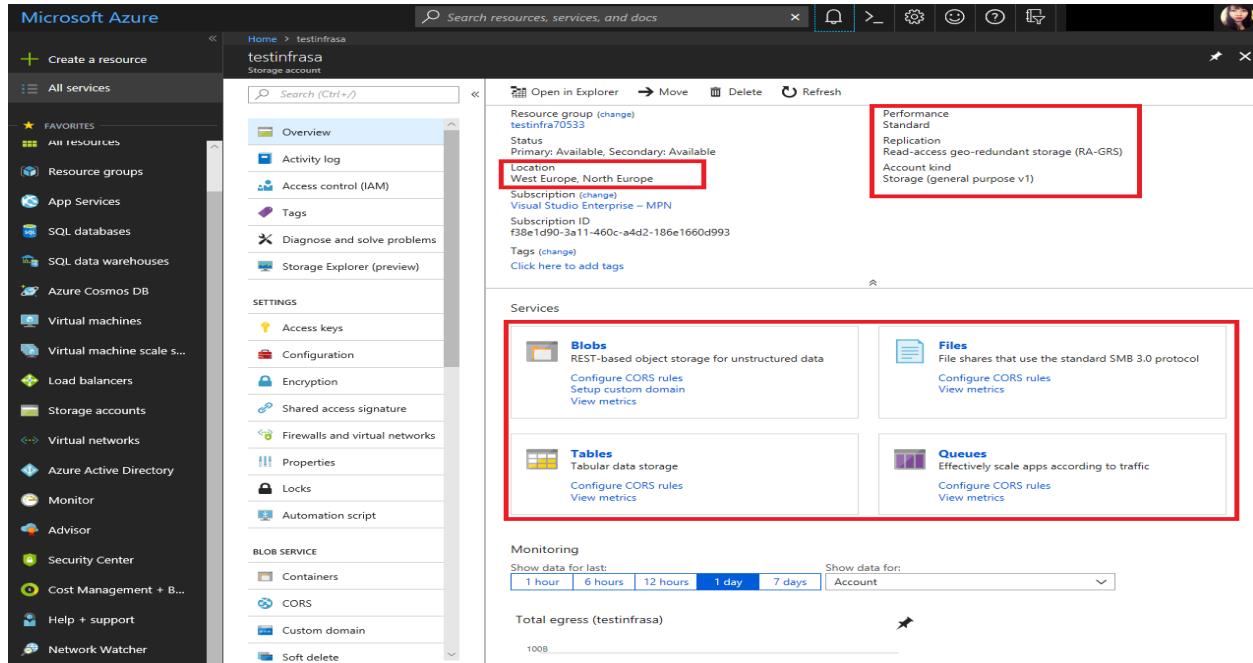
The access tiers are designed for storing data at a lower cost. Data residing in this account will automatically inherit that access tier setting. The hot and cool storage tiers can be set **at the account level**; archival storage can be set **at the object level** as well as the two others.

The general storage account has two performance tiers:

- **Standard storage performance tier:** This hosts Tables, Queues, Files, Blobs and Azure VM disks.
- **Premium storage performance tier:** This currently only supports Azure VM disks. The performance settings can't be changed after the storage account is created. Note that Azure VMs, which use premium storage for all disks, will be guaranteed a 99.9% SLA.

The secure transfer required setting means, after enabling this setting, all the requests to the storage account should be considered as secure connections, otherwise, it would be rejected by Azure. As an example, when this setting is enabled, any requests using HTTP will be rejected when you're calling the data in the blob storage by programming methods using the Azure RESTful API.

After filling in all the information, you can click on **Create**. The deployment usually takes a few seconds. After creating the storage account successfully, you can go to your storage account (as shown in the following screenshot):



The screenshot shows the Microsoft Azure Storage Overview blade for a storage account named 'testinfrasa'. The left sidebar lists various Azure services. The main pane displays the storage account's details, including its resource group ('testinfrasa'), status ('Provisioned/Available, Secondary/ Available'), location ('West Europe, North Europe'), and account kind ('Storage (general purpose v1)'). A red box highlights the 'Location' and 'Account kind' sections. Below this, under 'Services', there are four cards: 'Blobs' (REST-based object storage for unstructured data), 'Files' (File shares that use the standard SMB 3.0 protocol), 'Tables' (Tabular data storage), and 'Queues' (Effectively scale apps according to traffic). A second red box highlights these four service cards. At the bottom, there is a 'Monitoring' section showing 'Total egress (testinfrasa)' with a value of 1008.

In the **Overview** blade, you can find the information related to your storage account, including **Resource group**, **Location**, **Performance tier**, **Replication**, and **Account kind**.

You can also create a storage account using Azure PowerShell as explained in the following link: <https://docs.microsoft.com/en-us/azure/storage/blobs/storage-samples-blobs-powershell>.

To get more information about creating a storage account using the Azure CLI, go to: <https://docs.microsoft.com/en-us/azure/storage/blobs/storage-samples-blobs-cli>.

```
{
  "name": "string",
  "type": "Microsoft.Storage/storageAccounts",
  "apiVersion": "2018-10-01",
  "sku": {
    "name": "string"
  },
  "kind": "string",
  "location": "string",
  "tags": {},
  "identity": {
    "type": "SystemAssigned"
  },
  "properties": {
```

```

"customDomain": {
  "name": "string",
  "useSubDomain": boolean
},
"encryption": {
  "services": {
    "blob": {
      "enabled": boolean
    },
    "file": {
      "enabled": boolean
    }
  },
  "keySource": "string",
  "keyvaultproperties": {
    "keyname": "string",
    "keyversion": "string",
    "keyvaultrui": "string"
  }
},
"networkAcls": {
  "bypass": "string",
  "virtualNetworkRules": [
    {
      "id": "string",
      "action": "Allow",
      "state": "string"
    }
  ],
  "ipRules": [
    {
      "value": "string",
      "action": "Allow"
    }
  ],
  "defaultAction": "string"
},
"accessTier": "string",
"supportsHttpsTrafficOnly": boolean
}
}

```



You can also find an example in the Azure quickstart template repository: <https://github.com/Azure/azure-quickstart-templates/tree/master/101-storage-account-create>.

Implementing Azure Blob storage

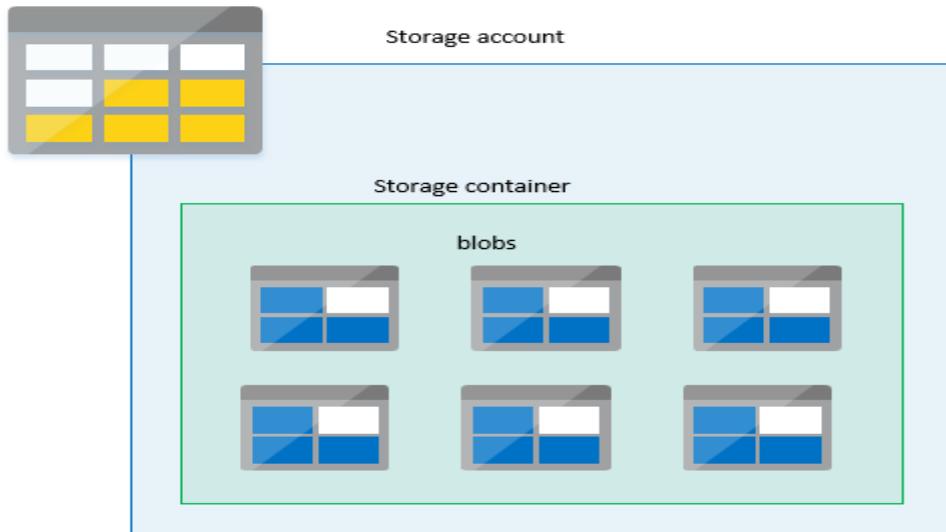
Azure Blob storage is an object-based storage solution in Microsoft Azure. To implement an Azure Blob, you should create a storage account, then a container, and finally the blob.

There are three types of blob, which are listed as follows:

- **Block blobs:** These are designed for storing text and binary data
- **Append blobs:** These are optimized for append operations, ideal for logs
- **Page blobs:** These are designed to store VHD files for Azure VMs

A container acts as an organizer, and contains a set of blobs. All the blobs reside within a container. There is no limit to the number of containers within the storage account and also no limit on the number of blobs in the containers.

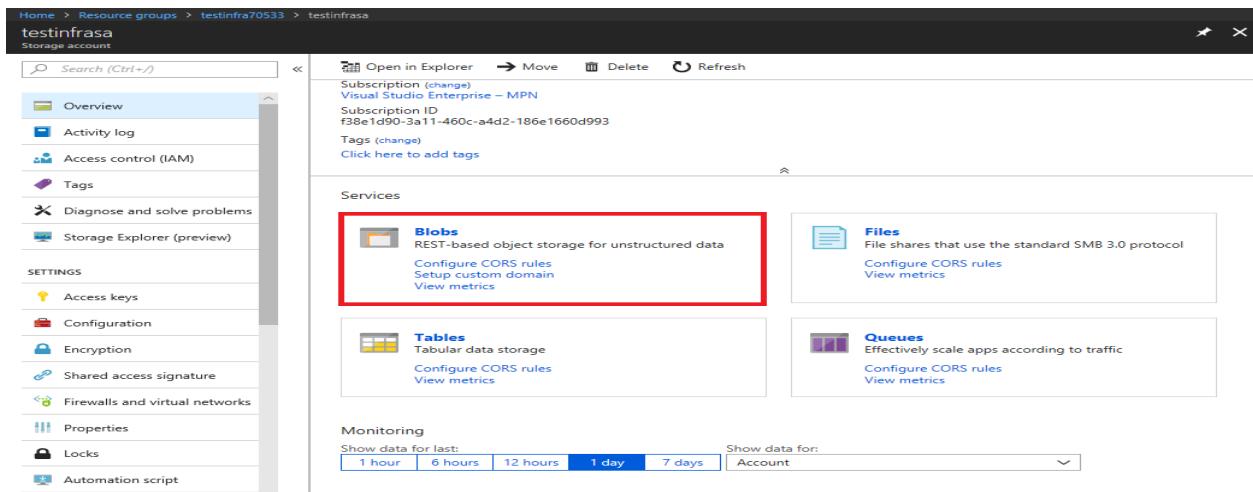
You can define the relationship between **Storage account**, **Storage container**, and **blobs** using the following schema:



The relationship between resources

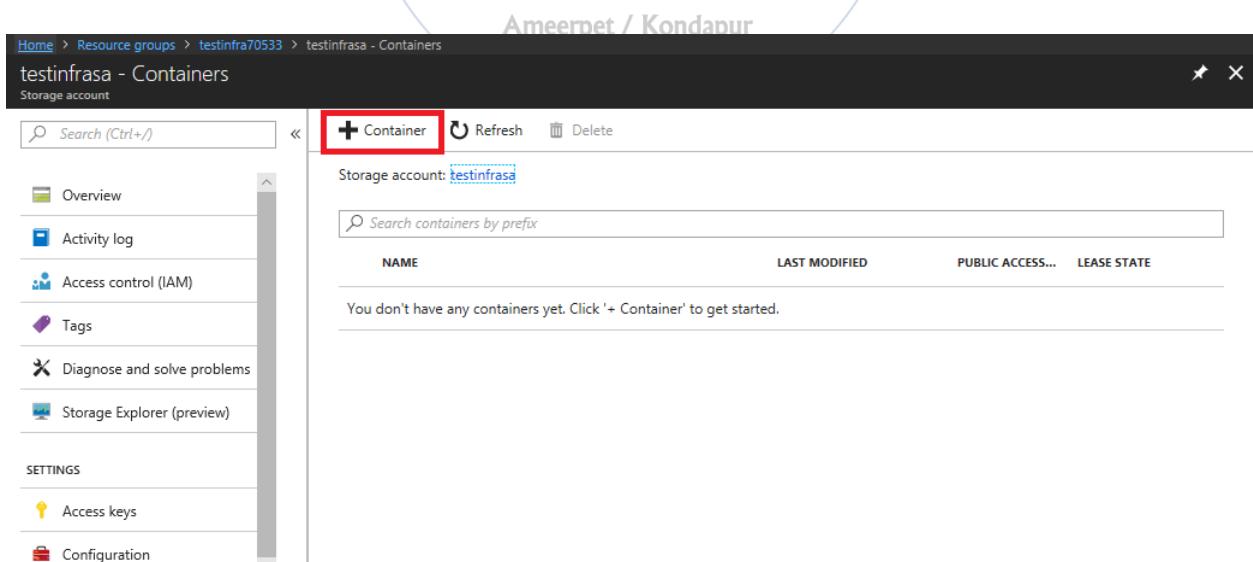
A storage account may have one root container, which acts as a default container for the current storage account. The root container is created by default and is named \$root. A text file resides in the root container and can be referenced in the following manner: <https://storageaccountname.blob.core.windows.net/blob.txt>.

To create a new blob, go to your storage account and click on **Blobs** as shown in the following screenshot:



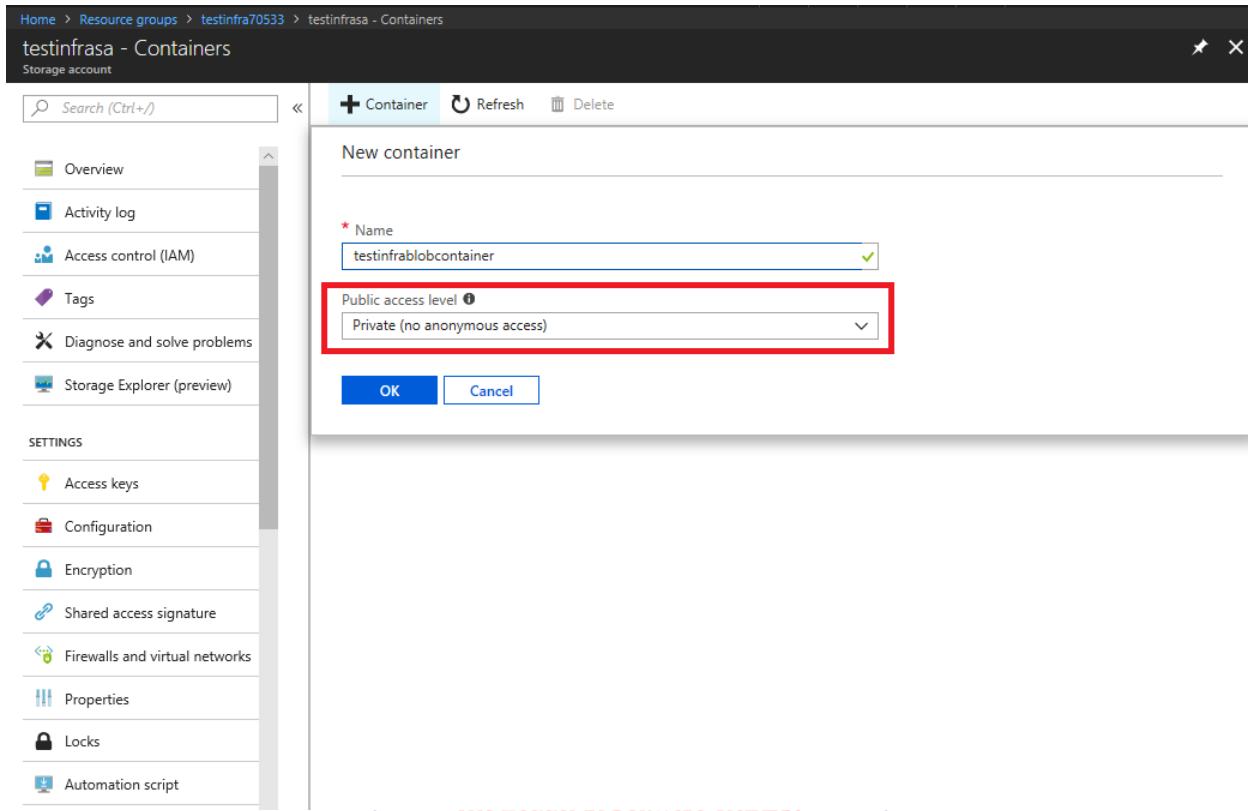
This screenshot shows the Azure Storage Account Overview page for the 'testinfra' storage account. The left sidebar contains navigation links for Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Storage Explorer (preview), and various settings like Access keys, Configuration, Encryption, Shared access signature, Firewalls and virtual networks, Properties, Locks, and Automation script. The main content area displays services: Blobs (REST-based object storage for unstructured data), Files (file shares that use the standard SMB 3.0 protocol), Tables (tabular data storage), and Queues (effectively scale apps according to traffic). The 'Blobs' section is highlighted with a red box. Below the services, there's a monitoring section with a timeline and a 'Show data for last:' dropdown.

Then, click on **+ Container** to create a new container of your blob as shown in the following screenshot:



This screenshot shows the Azure Storage Account Containers page for the 'testinfra' storage account. The left sidebar is identical to the previous screenshot. The main content area shows a table with columns: NAME, LAST MODIFIED, PUBLIC ACCESS..., and LEASE STATE. A message at the top states: 'You don't have any containers yet. Click '+ Container' to get started.' A red box highlights the '+ Container' button in the top navigation bar.

To create a new container, you should choose a permission. By default, a container is set to Private (no anonymous access), which means all the blobs within the container can only be accessed by the storage account owner (as shown in the following screenshot):

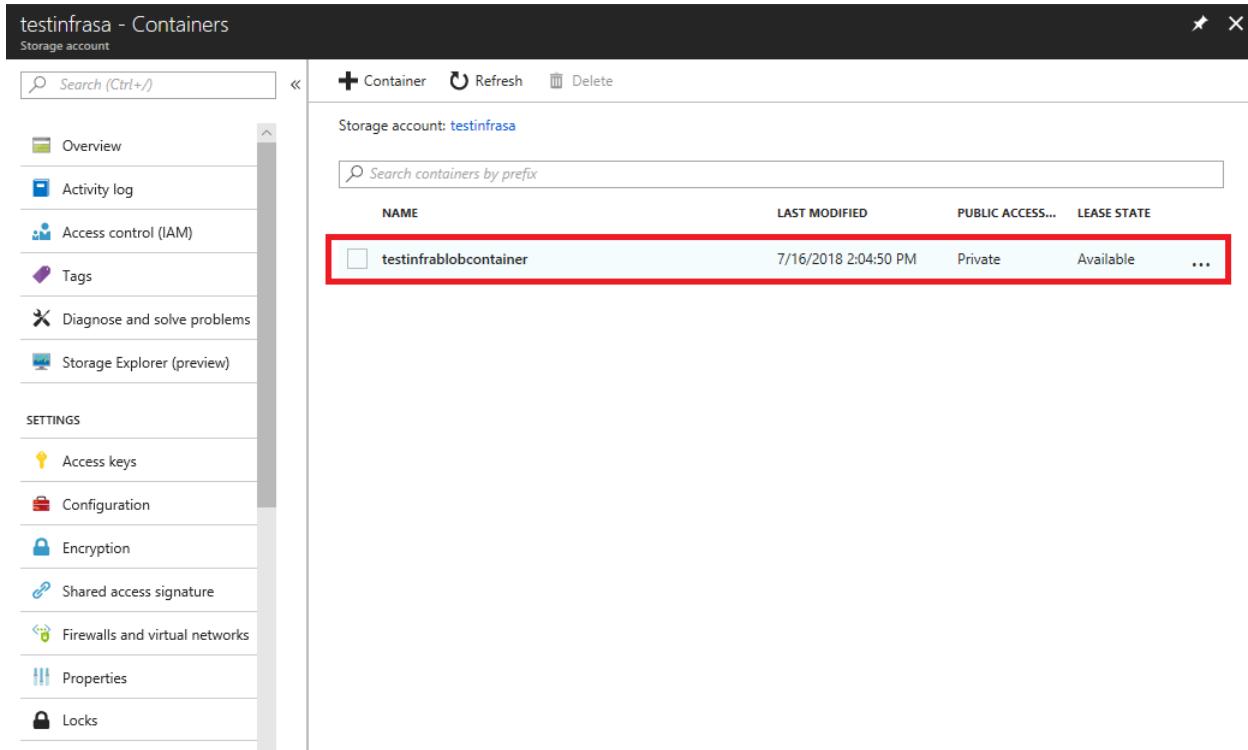


However, there are three permissions that are required to configure a container, as follows:

Ameerpet / Kondapur

- **Private:** This is used to set all the blobs that can be accessed only by the storage account owner
- **Blob:** This allows anonymous read access within the container, but not for container data
- **Container:** This allows anonymous read access to containers and blobs

After clicking on **OK**, you'll see that a container with private access has been created in the storage account.



testinfrasa - Containers

Storage account: testinfrasa

NAME	LAST MODIFIED	PUBLIC ACCESS...	LEASE STATE
testinfrablobcontainer	7/16/2018 2:04:50 PM	Private	Available

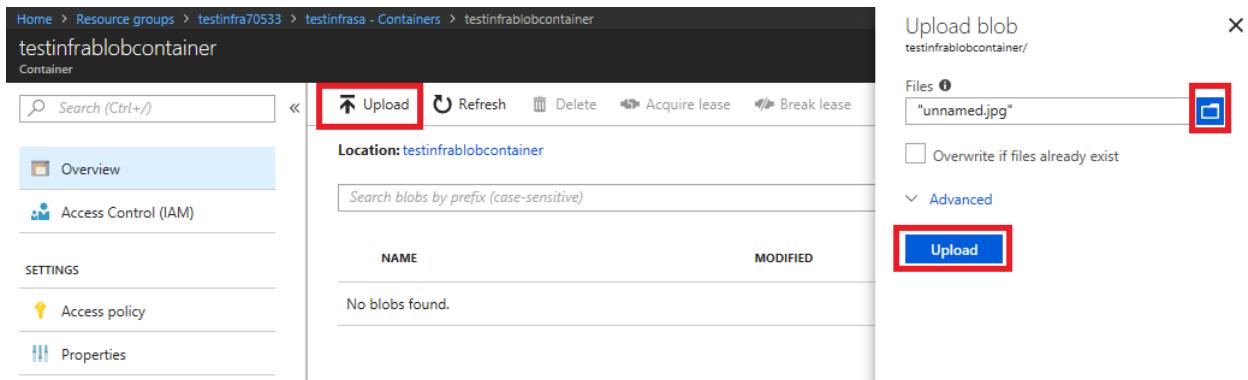
Managing Azure Blob storage



Go to the container to manage your blob storage. [7730997544](tel:7730997544)

Ameerpet / Kondapur

To add a blob, you can click on **Upload** in the **Overview** blade; a popup will be displayed. Then, click on the file icon to choose a file in your PC and click on **Upload**, as shown in the following screenshot:



testinfrablobcontainer

Container

Upload

Location: testinfrablobcontainer

No blobs found.

Upload blob

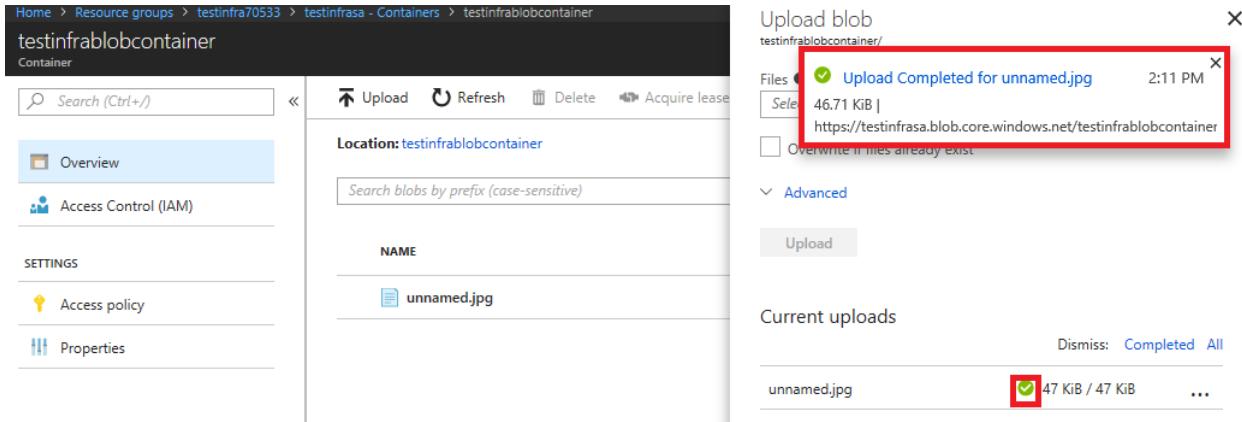
testinfrablobcontainer/

Files "unnamed.jpg"

Overwrite if files already exist

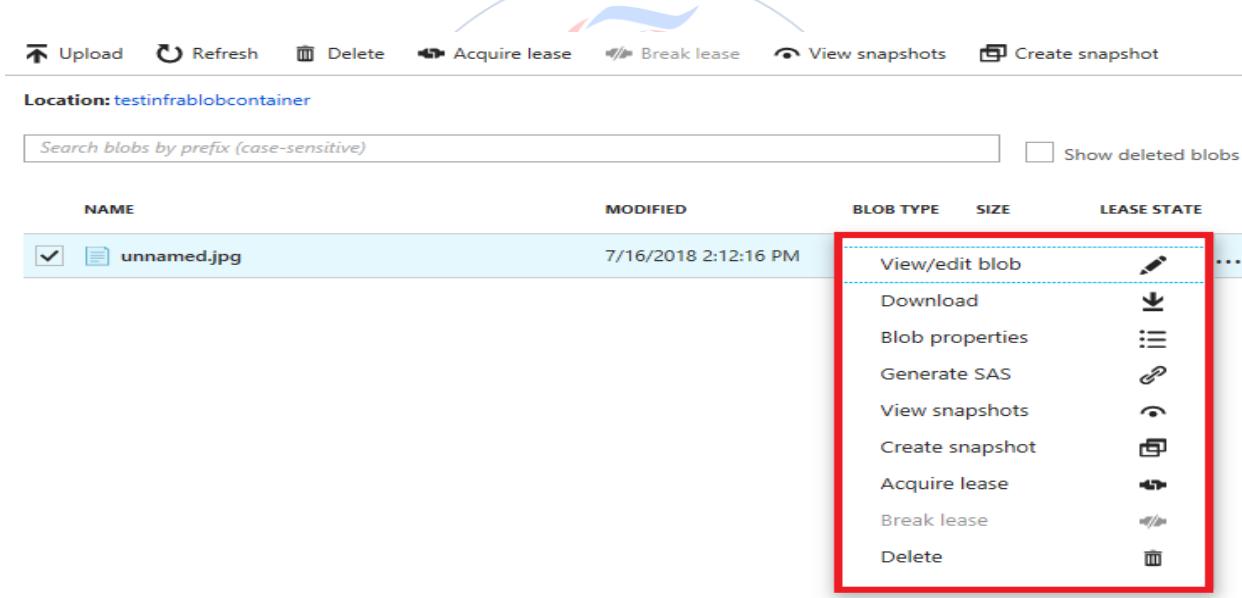
Upload

After clicking on **Upload**, you can see that your blob has been uploaded successfully:



The screenshot shows the Azure Storage Blob Container interface. On the left, there's a sidebar with 'Overview' and 'Access Control (IAM)' sections. The main area shows a table with one row for 'unnamed.jpg'. A red box highlights a message at the top right: 'Upload Completed for unnamed.jpg' with a timestamp '2:11 PM'. Below it, there's an unchecked checkbox for 'Overwrite if files already exist'. On the right, a modal window titled 'Upload blob' shows the same message and file details.

You can also use other management functions provided by Azure by right-clicking on the blob (as shown in the following screenshot):



The screenshot shows the same Azure Storage Blob Container interface. A context menu is open over the 'unnamed.jpg' row. A red box highlights the menu options: View/edit blob, Download, Blob properties, Generate SAS, View snapshots, Create snapshot, Acquire lease, Break lease, and Delete.

Click on **Blob properties**; you can find the endpoint of the blob as shown in the following screenshot:

Home > Resource groups > testinfra70533 > testinfrasa - Containers > testinfrablobcontainer > unnamed.jpg

unnamed.jpg
Blob

Save Discard Refresh Download Acquire lease Break lease Delete

Overview Snapshots Edit blob Generate SAS

Properties

URL	https://testinfrasa.blob.core.windows.net/t...
LAST MODIFIED	7/16/2018 2:12:16 PM
CREATION TIME	7/16/2018 2:12:16 PM
TYPE	Block blob
SIZE	46.71 KB
SERVER ENCRYPTED	true
ETAG	0x8D5EB155EF18796
CONTENT-MD5	-
LEASE STATUS	Unlocked
LEASE STATE	Available
LEASE DURATION	-
COPY STATUS	-
COPY COMPLETION TIME	-

Undelete all snapshots

Access Tier

Optimize storage costs by placing your data in the appropriate access tier. Archive is not supported in your region at this time. [Learn more](#)

N/A

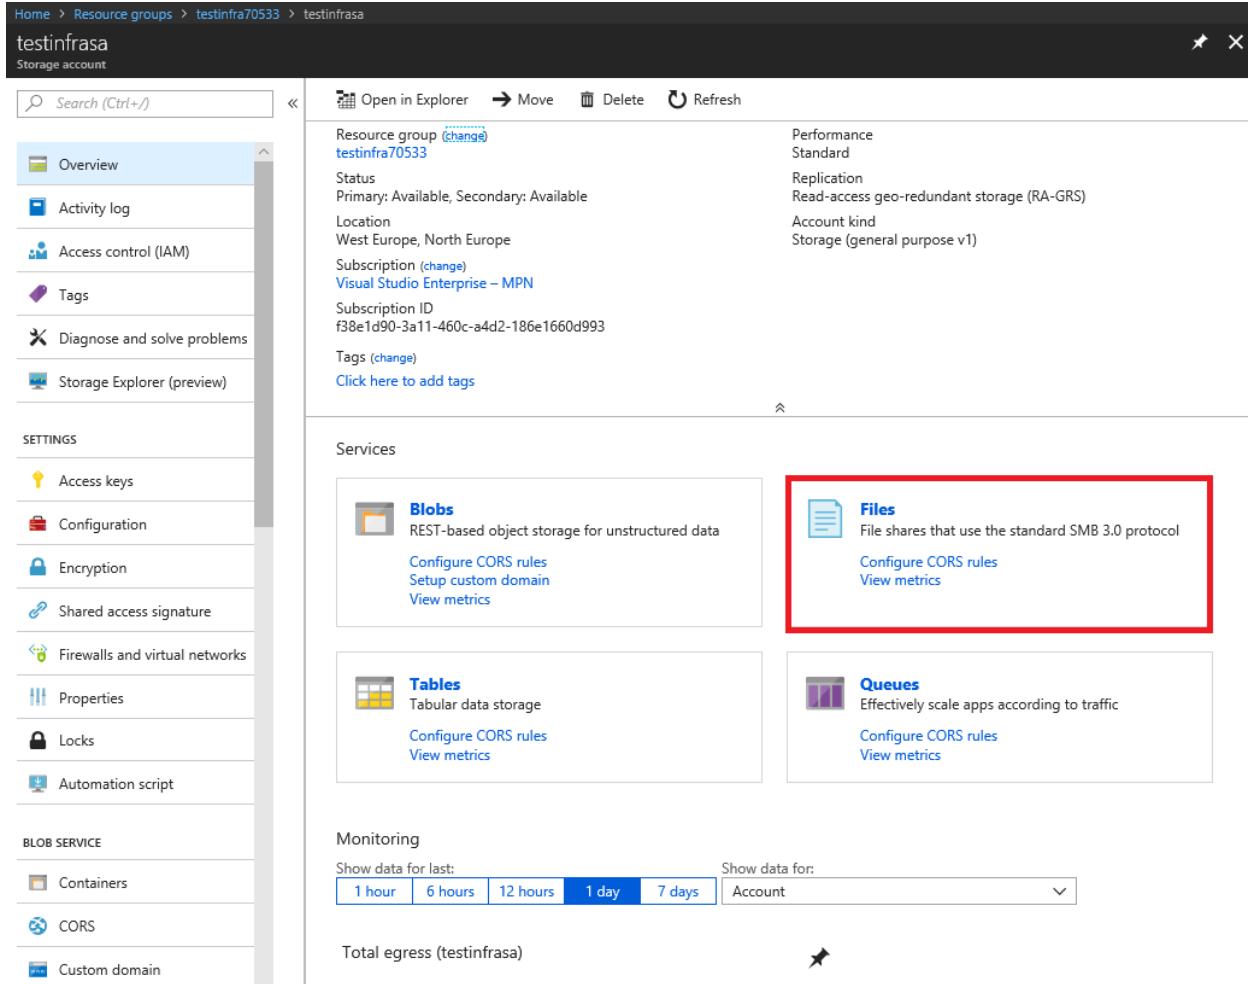
Setting object level access tier is only supported for Standard LRS, GRS, and RA-GRS Blob Storage and General Purpose V2 Accounts.

The default endpoints for blobs in the storage account are as follows: <http://storageaccountname.blob.core.windows.net>.¹⁵⁴⁴

Implementing Azure Files storage

Azure Files provides managed file shares in Azure cloud, which is accessible via the **Server Message Block (SMB)** protocol. It is possible to be mounted as a file share by Azure VMs in the cloud or even VMs on-premise. Azure files can be used by different operating systems such as Windows, Linux, and macOS.

To create a file share, go to the storage account and click on **Files** as shown in the following screenshot:



Home > Resource groups > testinfra70533 > testinfrasa

testinfrasa
Storage account

Search (Ctrl+/
)

Overview Activity log Access control (IAM) Tags Diagnose and solve problems Storage Explorer (preview)

SETTINGS

- Access keys Configuration Encryption Shared access signature Firewalls and virtual networks Properties Locks Automation script

BLOB SERVICE

- Containers CORS Custom domain

Resource group (change)
testinfra70533

Status Primary: Available, Secondary: Available

Location West Europe, North Europe

Subscription (change)
Visual Studio Enterprise – MPN

Subscription ID f38e1d90-3a11-460c-a4d2-186e1660d993

Tags (change)
Click here to add tags

Performance Standard

Replication Read-access geo-redundant storage (RA-GRS)

Account kind Storage (general purpose v1)

Services

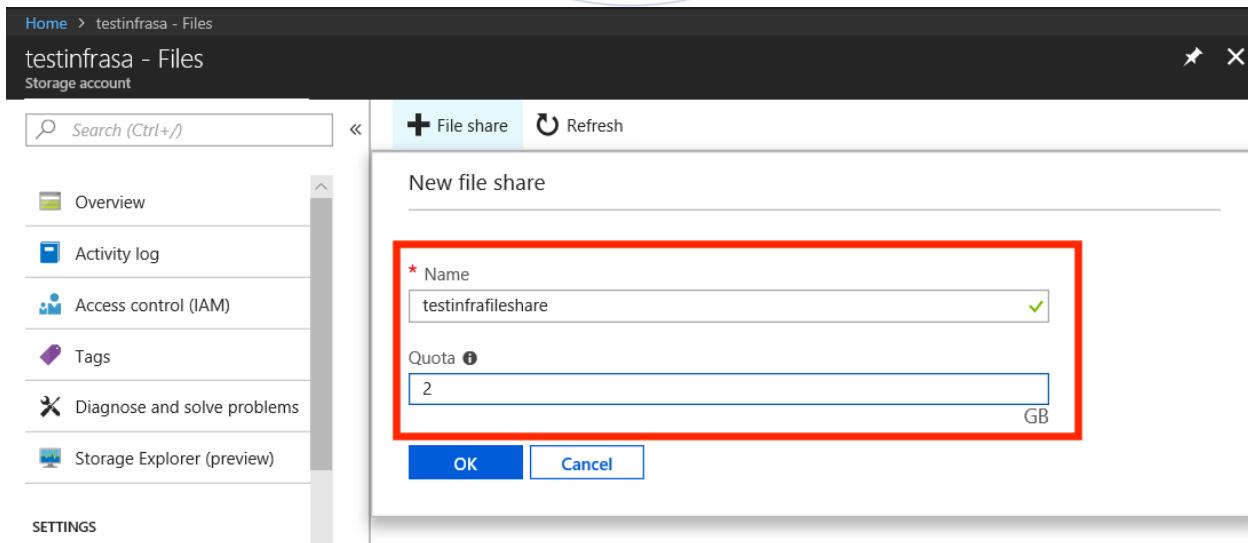
- Blobs** REST-based object storage for unstructured data
 - Configure CORS rules
 - Setup custom domain
 - View metrics
- Files** File shares that use the standard SMB 3.0 protocol
 - Configure CORS rules
 - View metrics
- Tables** Tabular data storage
 - Configure CORS rules
 - View metrics
- Queues** Effectively scale apps according to traffic
 - Configure CORS rules
 - View metrics

Monitoring

Show data for last: 1 hour 6 hours 12 hours 1 day 7 days Account

Total egress (testinfrasa)

Then, click on **+ File share** to add a new file share. You should provide the name of the file share quota, which is limited to 5120 GB, as shown in the following screenshot:



Home > testinfrasa - Files

testinfrasa - Files
Storage account

Search (Ctrl+/
)

Overview Activity log Access control (IAM) Tags Diagnose and solve problems Storage Explorer (preview)

SETTINGS

+ File share Refresh

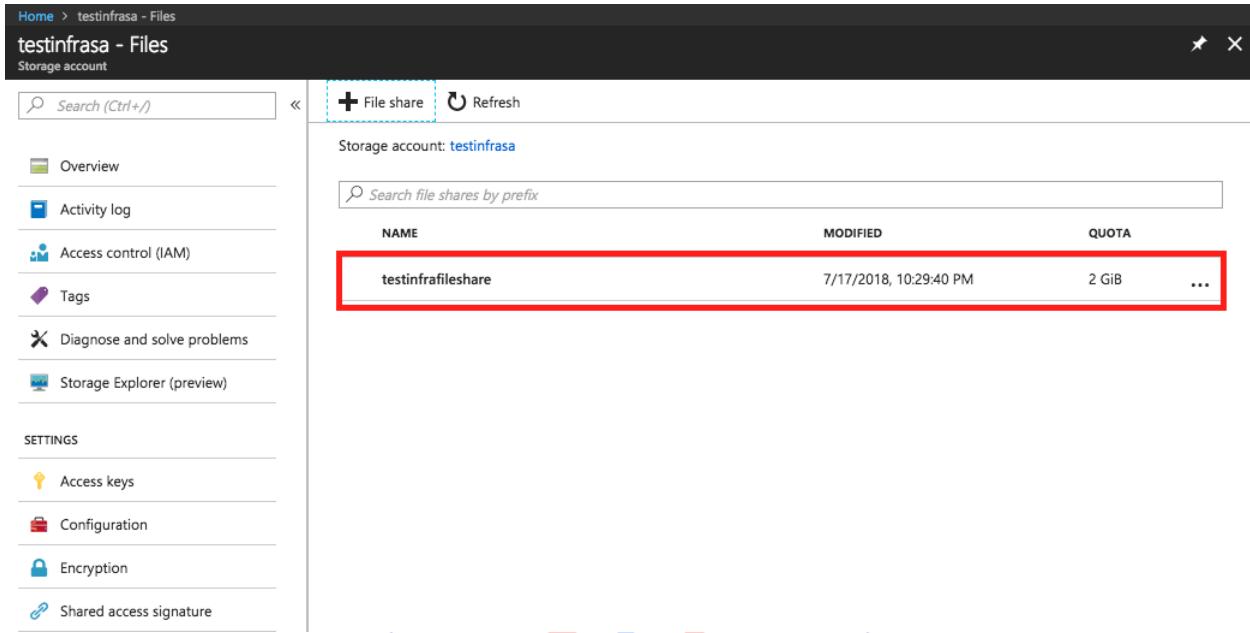
New file share

* Name testinfrashare

Quota 2 GB

OK Cancel

After a few seconds, you'll see that a file share has been deployed successfully, as shown in the following screenshot:



The screenshot shows the Azure Storage Account - Files blade. The storage account name is "testinfrasa". A red box highlights the "File share" button in the top navigation bar. Below it, a table lists one file share named "testinfrafleshare". The table columns are NAME, MODIFIED, and QUOTA. The entry for "testinfrafleshare" is highlighted with a red border.

NAME	MODIFIED	QUOTA
testinfrafleshare	7/17/2018, 10:29:40 PM	2 GiB

To create a file share through the Azure CLI, you can refer to the following link: <https://docs.microsoft.com/en-au/azure/storage/files/storage-how-to-create-file-share#create-file-share-through-command-line-interface-cli>.

To create file share through PowerShell, you can refer to the following link: <https://docs.microsoft.com/en-au/azure/storage/files/storage-how-to-create-file-share#create-file-share-through-powershell>.

Azure File share can be used by different operating systems such as Windows, Linux, and macOS.

For Windows, you can refer to: <https://docs.microsoft.com/en-au/azure/storage/files/storage-how-to-use-files-windows>.

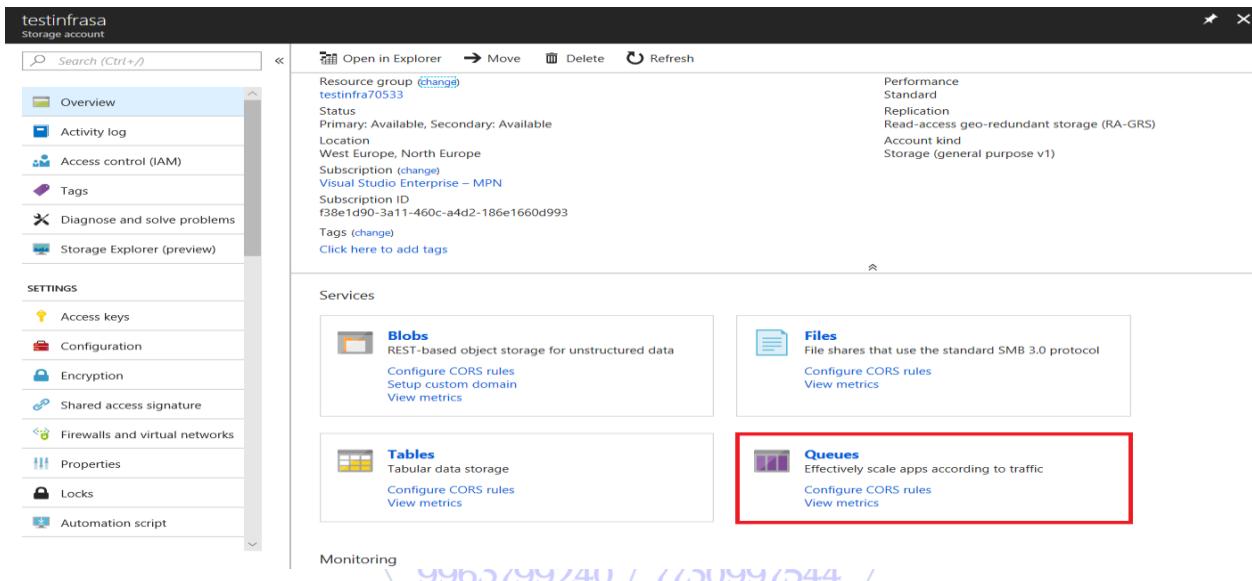
For Linux, you can refer to: <https://docs.microsoft.com/en-au/azure/storage/files/storage-how-to-use-files-linux>.

For MacOS, you can refer to: <https://docs.microsoft.com/en-au/azure/storage/files/storage-how-to-use-files-mac>.

The default endpoints for the blobs in the storage account are as follows: <https://#youstorageaccountname#.file.core.windows.net/#fiesharenname#>.

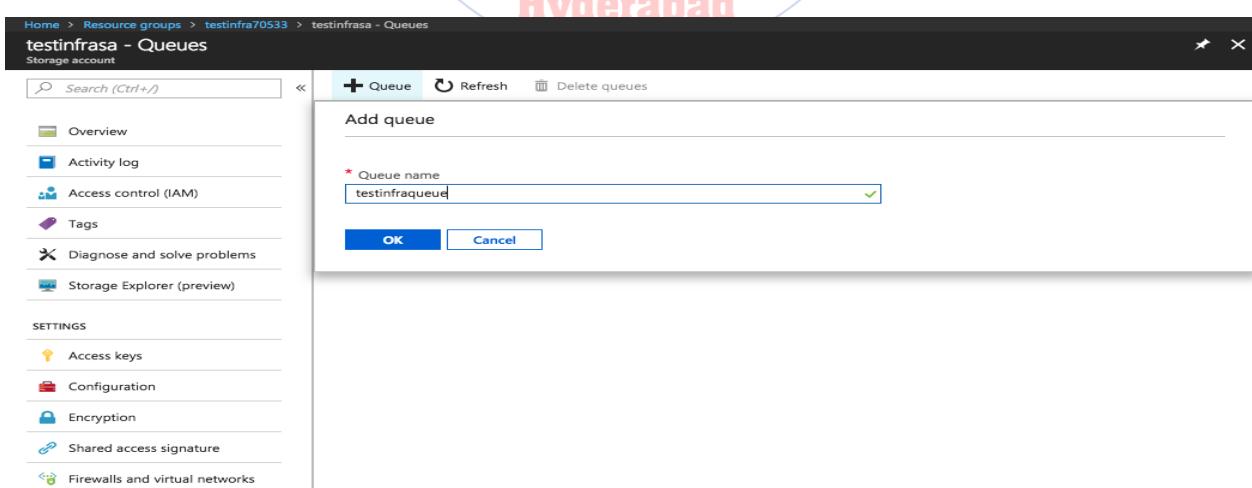
Implementing Azure Queue storage

Azure Queue storage is designed for storing large numbers of messages using the HTTP or HTTPS protocols. To create a queue storage, go to the storage account and click on **Queues**, as shown in the following screenshot:



This screenshot shows the Azure Storage Account Overview page for 'testinfrasa'. The left sidebar includes links for Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, and Storage Explorer (preview). The main pane displays account details such as Resource group (testinfrat0533), Status (Primary: Available, Secondary: Available), Location (West Europe, North Europe), Subscription (Visual Studio Enterprise – MPN), and Account kind (Storage (general purpose v1)). The 'Services' section lists Blobs, Tables, Files, and Queues. The 'Queues' section is highlighted with a red border. The bottom of the screen shows monitoring metrics: 99b5/99240 / 110099/044.

Then, click on **+ Queue** to add a new queue storage, as shown in the following screenshot:



This screenshot shows the 'Queues' blade for 'testinfrasa'. The left sidebar is identical to the previous screenshot. The main area shows a 'Add queue' dialog with a 'Queue name' field containing 'testinfracue1'. Buttons for 'OK' and 'Cancel' are at the bottom. The URL in the browser bar is 'Home > Resource groups > testinfrat0533 > testinfrasa - Queues'.

Click **OK**. After a few seconds, you'll find that a queue storage has been deployed successfully, as shown in the following screenshot:

Home > Resource groups > testinfra70533 > testinfra - Queues

testinfra - Queues

Storage account

Search (Ctrl+)/ Queue Refresh Delete queues

Storage account: testinfra

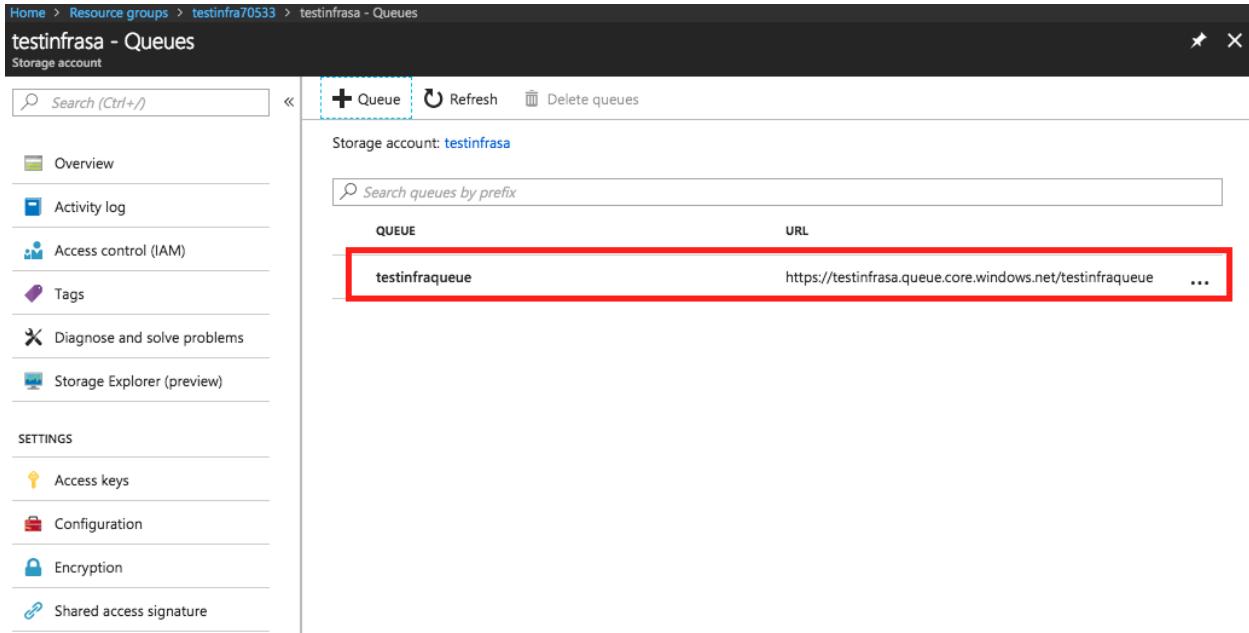
Search queues by prefix

QUEUE	URL	...
testinfraqueue	https://testinfrasa.queue.core.windows.net/testinfraqueue	...

Overview Activity log Access control (IAM) Tags Diagnose and solve problems Storage Explorer (preview)

SETTINGS

Access keys Configuration Encryption Shared access signature



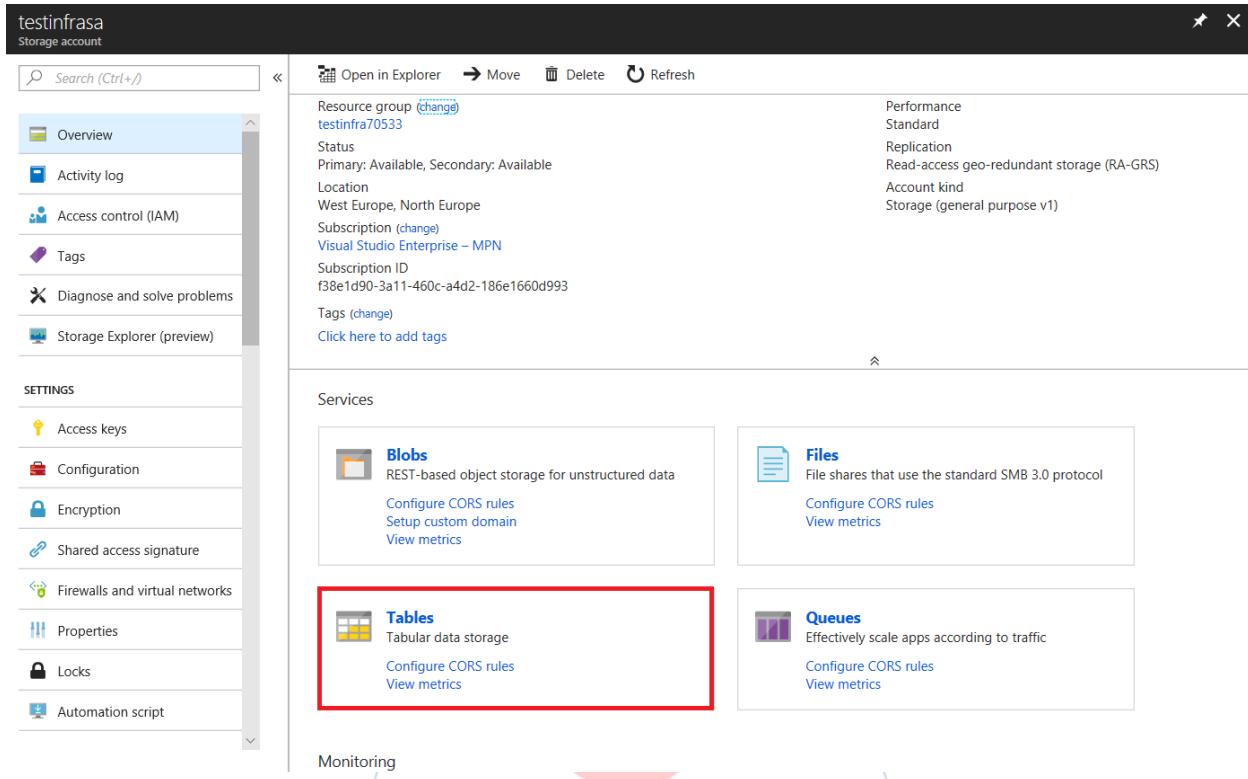
A new queue storage has been created successfully

The default endpoints for the file share in the storage account are as follows: <https://#youstorageaccountname#.queue.core.windows.net/#queuename#>.

For a comparison between Storage in Software Training Service Bus queues, refer to: <https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-azure-and-service-bus-queues-compared-contrasted>

Implementing Azure Table storage

Azure Table storage is designed to store schema-less or NoSQL data in the cloud; it provides a key/attribute store. To create a table storage, go to the storage account and click on **Tables**, as shown in the following screenshot:



testinfrasa
Storage account

Resource group [\(change\)](#) testinfr70533

Status Primary: Available, Secondary: Available

Location West Europe, North Europe

Subscription [\(change\)](#) Visual Studio Enterprise – MPN

Subscription ID f38e1d90-3a11-460c-a4d2-186e1660d993

Tags [\(change\)](#)
[Click here to add tags](#)

Performance Standard

Replication Read-access geo-redundant storage (RA-GRS)

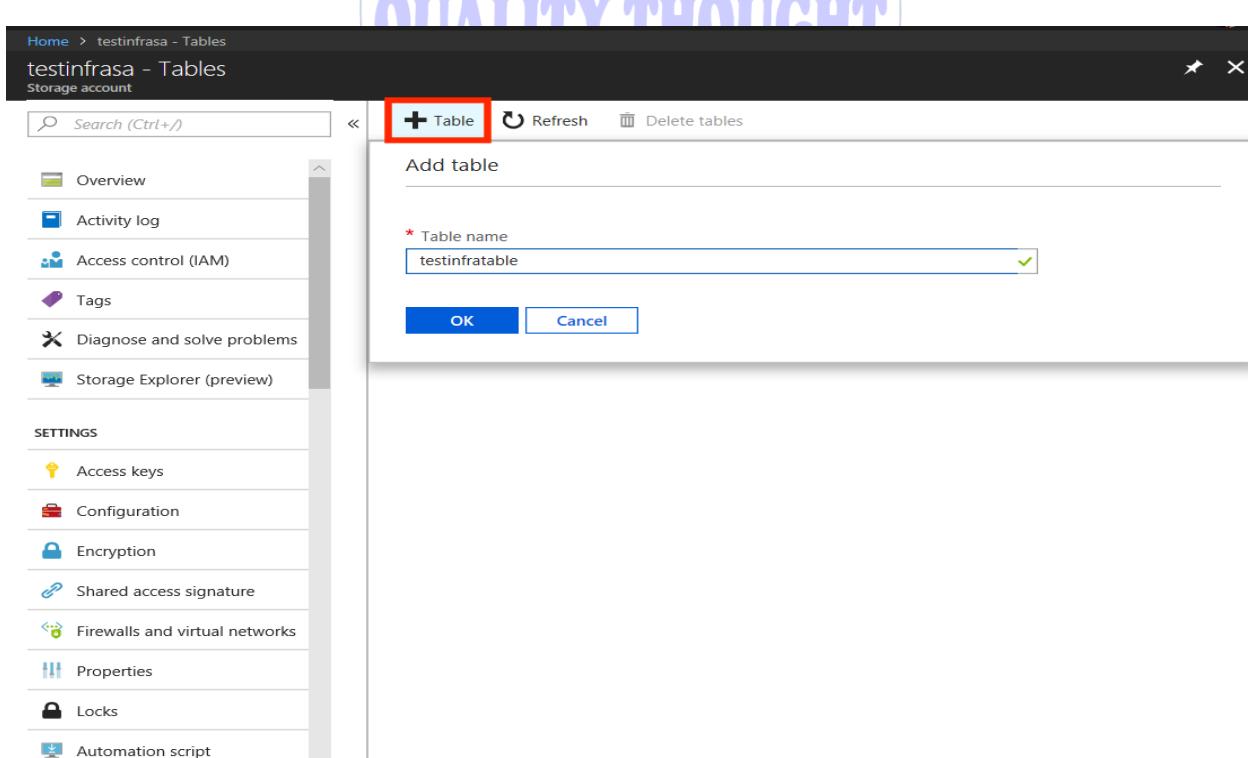
Account kind Storage (general purpose v1)

Services

- Blobs** REST-based object storage for unstructured data
 - [Configure CORS rules](#)
 - [Setup custom domain](#)
 - [View metrics](#)
- Files** File shares that use the standard SMB 3.0 protocol
 - [Configure CORS rules](#)
 - [View metrics](#)
- Tables** Tabular data storage
 - [Configure CORS rules](#)
 - [View metrics](#)
- Queues** Effectively scale apps according to traffic
 - [Configure CORS rules](#)
 - [View metrics](#)

Monitoring

Then, click on **+ Table** to add a new table storage, as shown in the following screenshot:



Home > testinfrasa - Tables

testinfrasa - Tables

Storage account

Search (Ctrl+)

+ Table Refresh Delete tables

Overview Activity log Access control (IAM) Tags Diagnose and solve problems Storage Explorer (preview)

SETTINGS

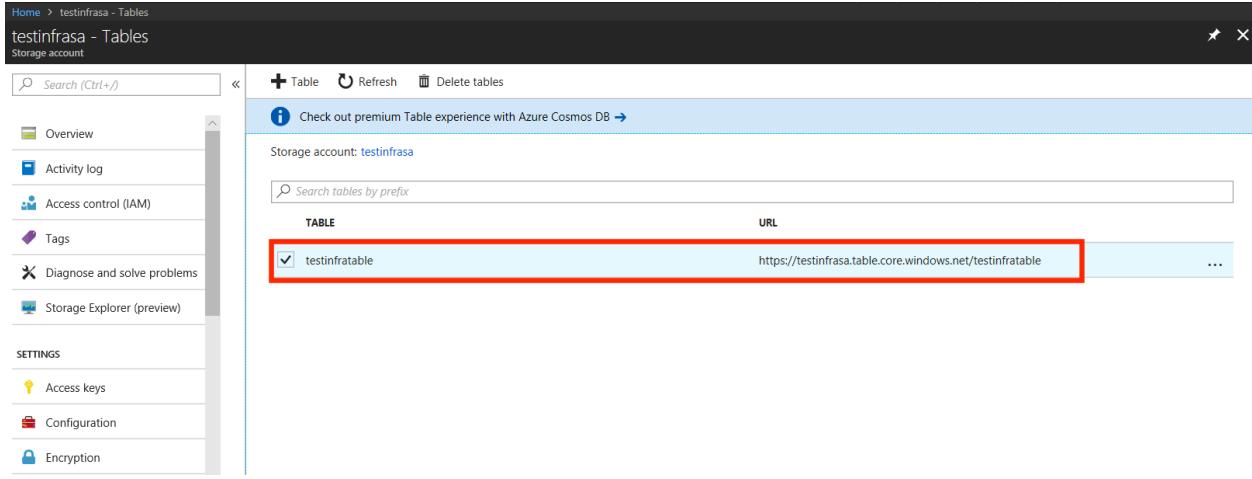
- Access keys Configuration Encryption Shared access signature Firewalls and virtual networks Properties Locks Automation script

Add table

* Table name ✓

OK Cancel

After a few seconds, you'll find that a table storage has been deployed successfully as shown in the following screenshot:



The screenshot shows the Azure Storage Tables blade. On the left, there's a navigation menu with options like Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, and Storage Explorer (preview). The main area is titled "testinfrasa - Tables" and shows a table named "testinfratable". A red box highlights this table entry. Below it is the URL: "https://testinfrasa.table.core.windows.net/testinfratable".

Azure Table storage accounts use the following

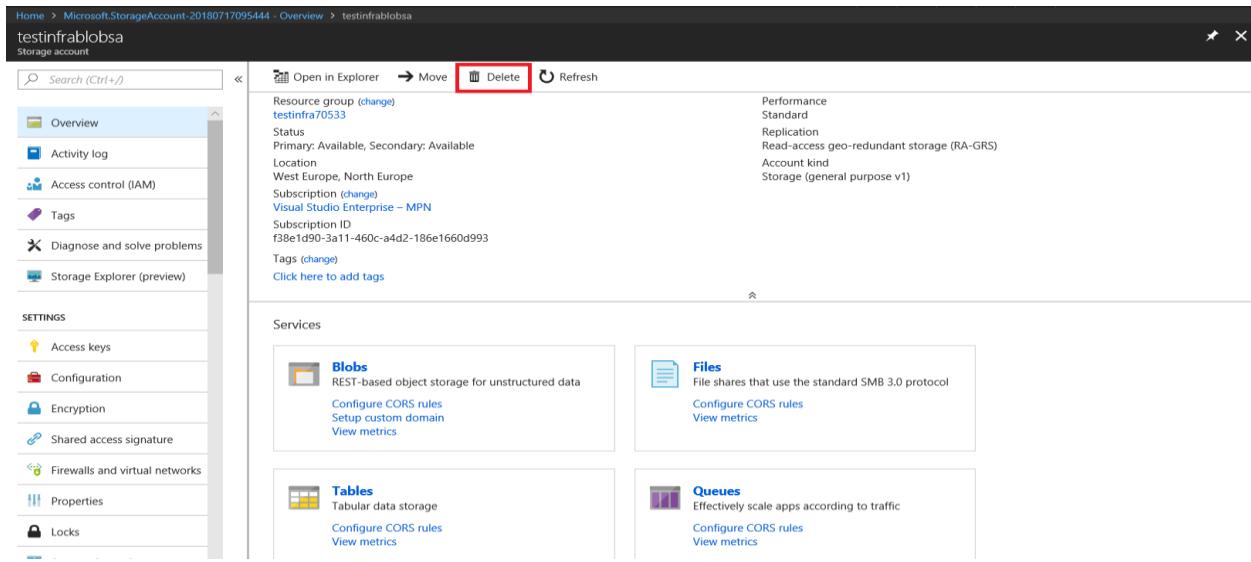
URL: <http://#storageaccountname#.table.core.windows.net/#tablestoragename#>.

Azure Cosmos DB Table API accounts use the following

URL: <http://#storageaccountname#.table.cosmosdb.azure.com/#tablestoragename#>.

Deleting a storage account

To delete a storage account, you can find the following in the **Overview** blade in the storage account:



Storage account testinfrablobsa

Resource group (change) testinfra70533

Status Primary: Available, Secondary: Available

Location West Europe, North Europe

Subscription (change) Visual Studio Enterprise – MPN

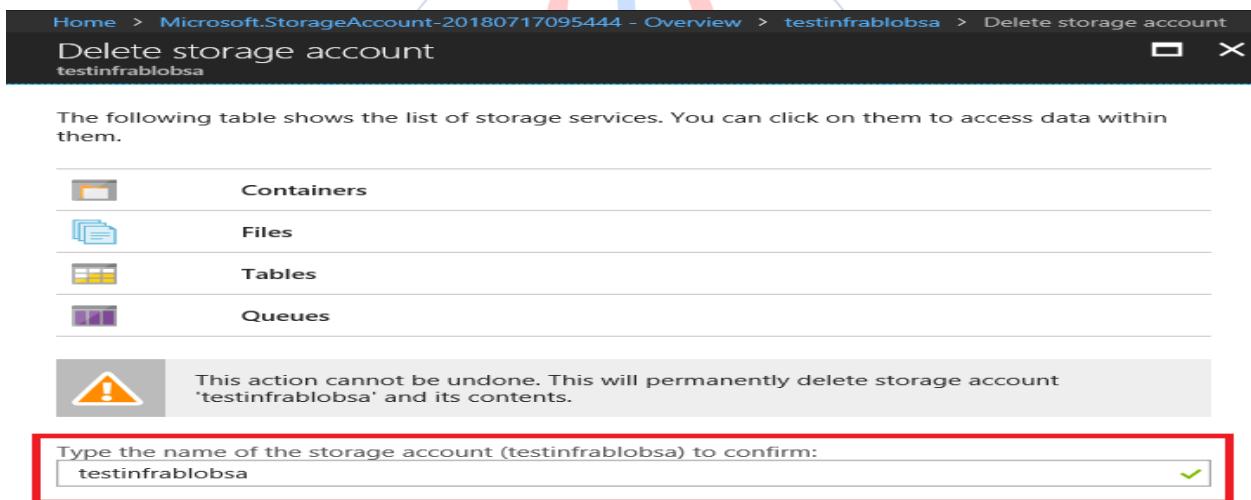
Subscription ID f38e1d90-3a11-460c-a4d2-186e1660d993

Tags (change) Click here to add tags

Services

- Blobs** REST-based object storage for unstructured data
 - Configure CORS rules
 - Setup custom domain
 - View metrics
- Files** File shares that use the standard SMB 3.0 protocol
 - Configure CORS rules
 - View metrics
- Tables** Tabular data storage
 - Configure CORS rules
 - View metrics
- Queues** Effectively scale apps according to traffic
 - Configure CORS rules
 - View metrics

After clicking on **Delete**, Azure will request you to type the storage account name:



The following table shows the list of storage services. You can click on them to access data within them.

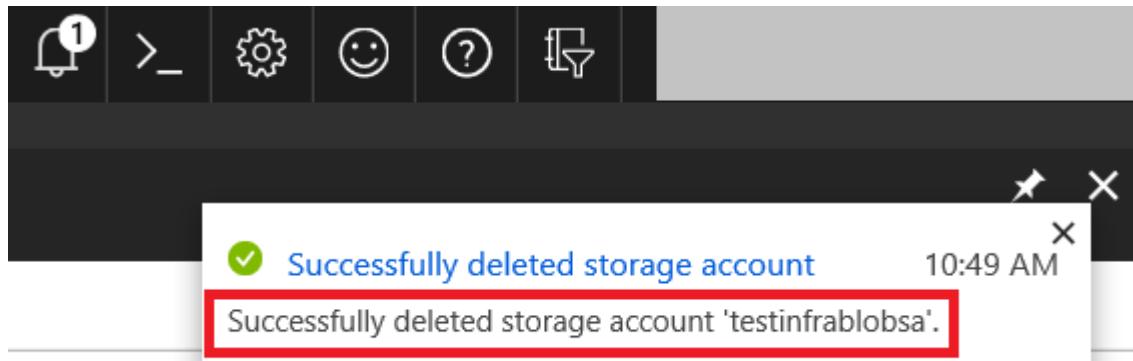
	Containers
	Files
	Tables
	Queues

⚠️ This action cannot be undone. This will permanently delete storage account 'testinfrablobsa' and its contents.

Type the name of the storage account (testinfrablobsa) to confirm:

Delete

Then, click on **Delete**. The operation will take a couple of seconds. A notification will launch to show that the operation was successful (as shown in the following screenshot):



Managing Azure Storage services

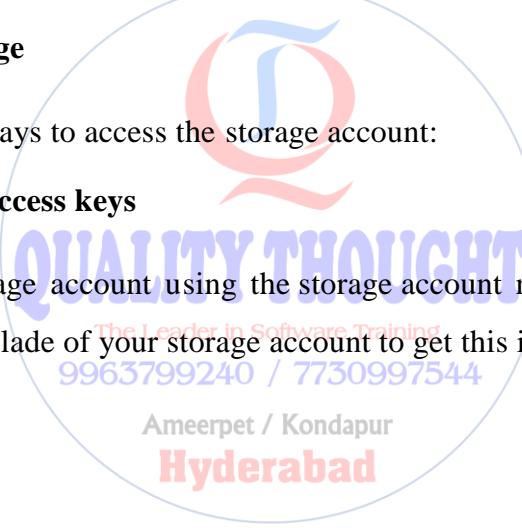
Let's see how we can manage Azure Storage services.

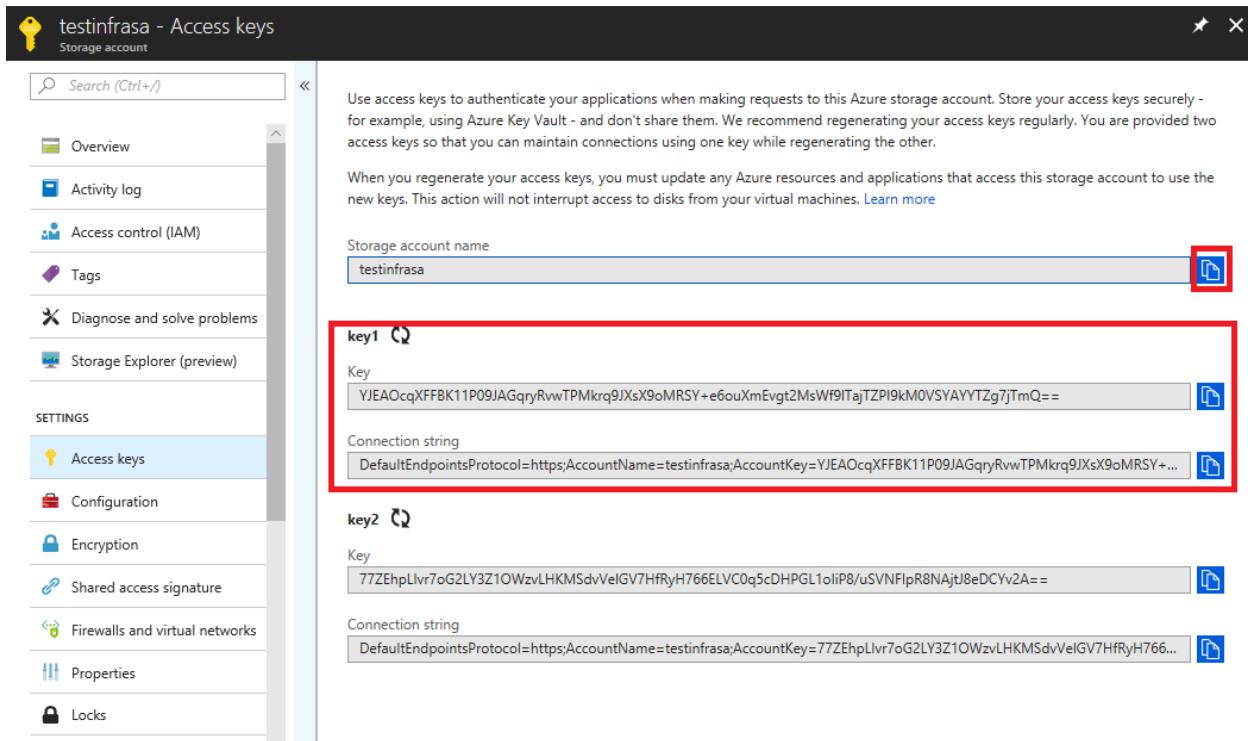
Managing access to storage

Here we'll introduce two ways to access the storage account:

Accessing storage using access keys

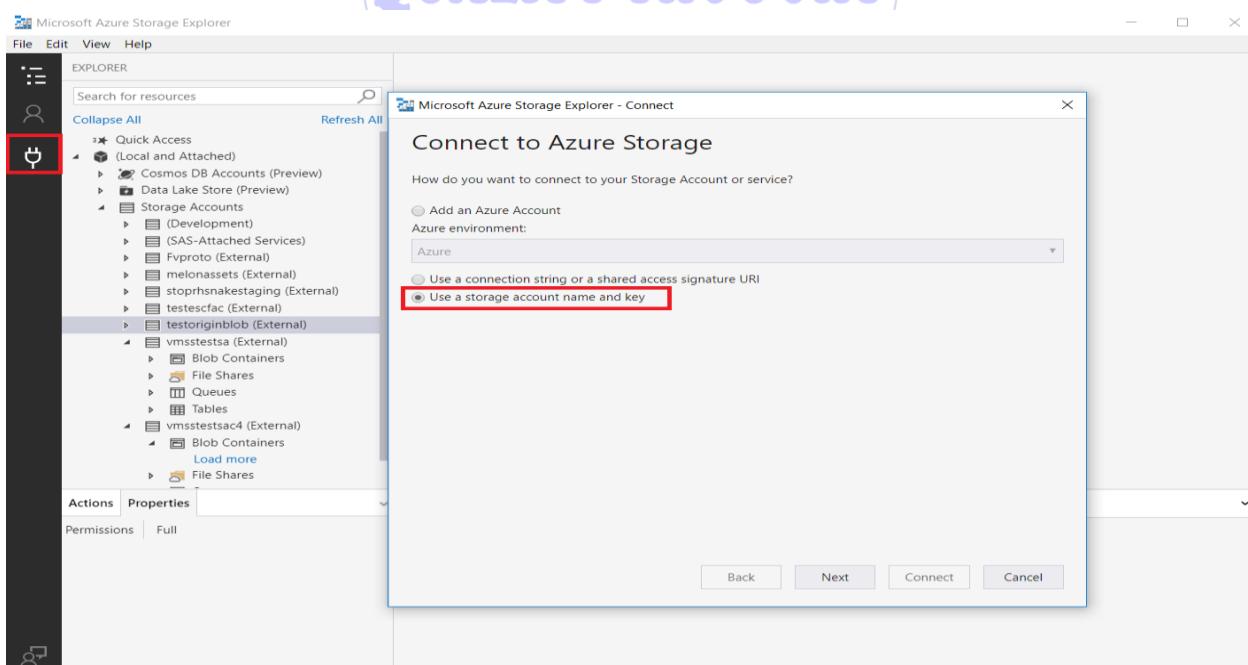
You can access your storage account using the storage account name and access keys. You can go to the **Access keys** blade of your storage account to get this information (as shown in the following screenshot):





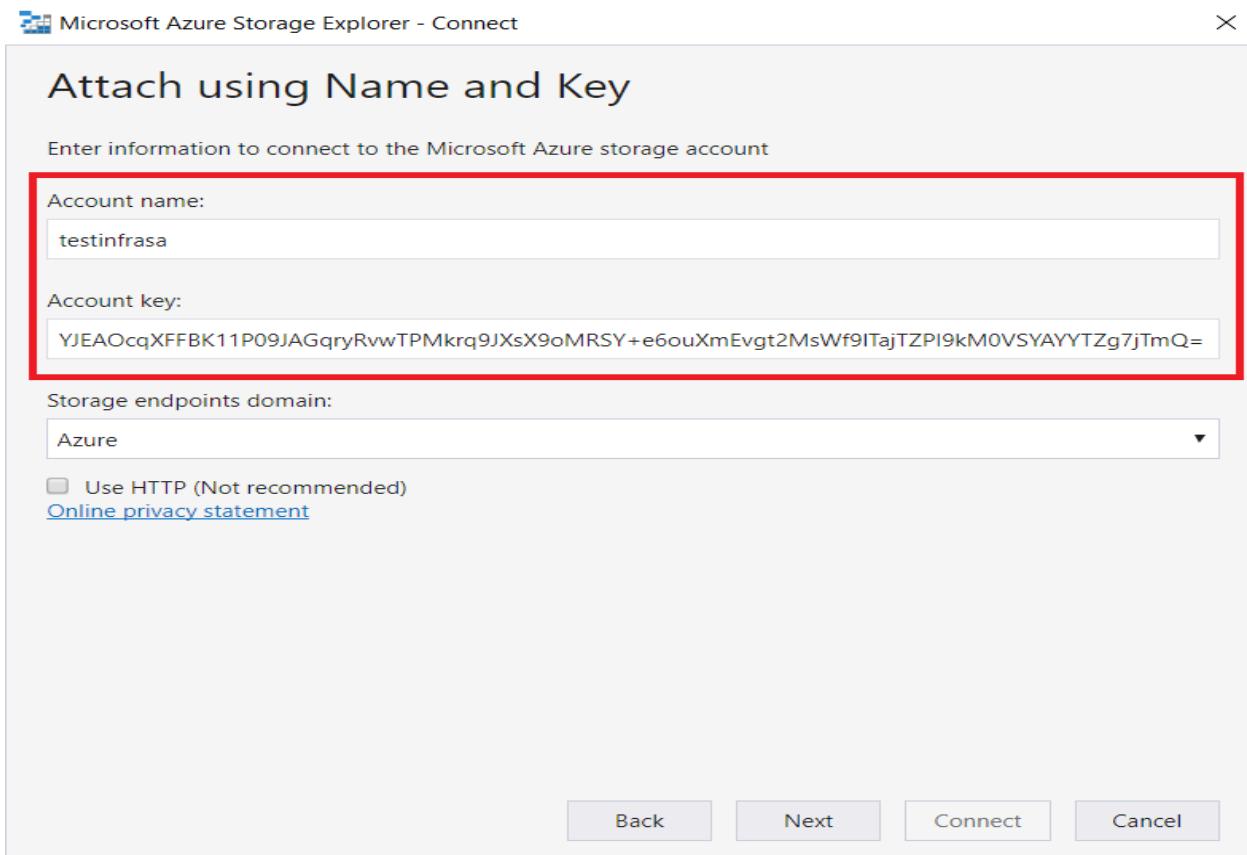
The screenshot shows the 'Access keys' section of the Azure Storage account settings. It displays two sets of keys: 'key1' and 'key2'. Each key includes a 'Key' value (a long hex string) and a 'Connection string' (a URL template). The 'key1' section is highlighted with a red box.

You can use this information to access Azure Storage using Storage Explorer (as shown in the following screenshot):

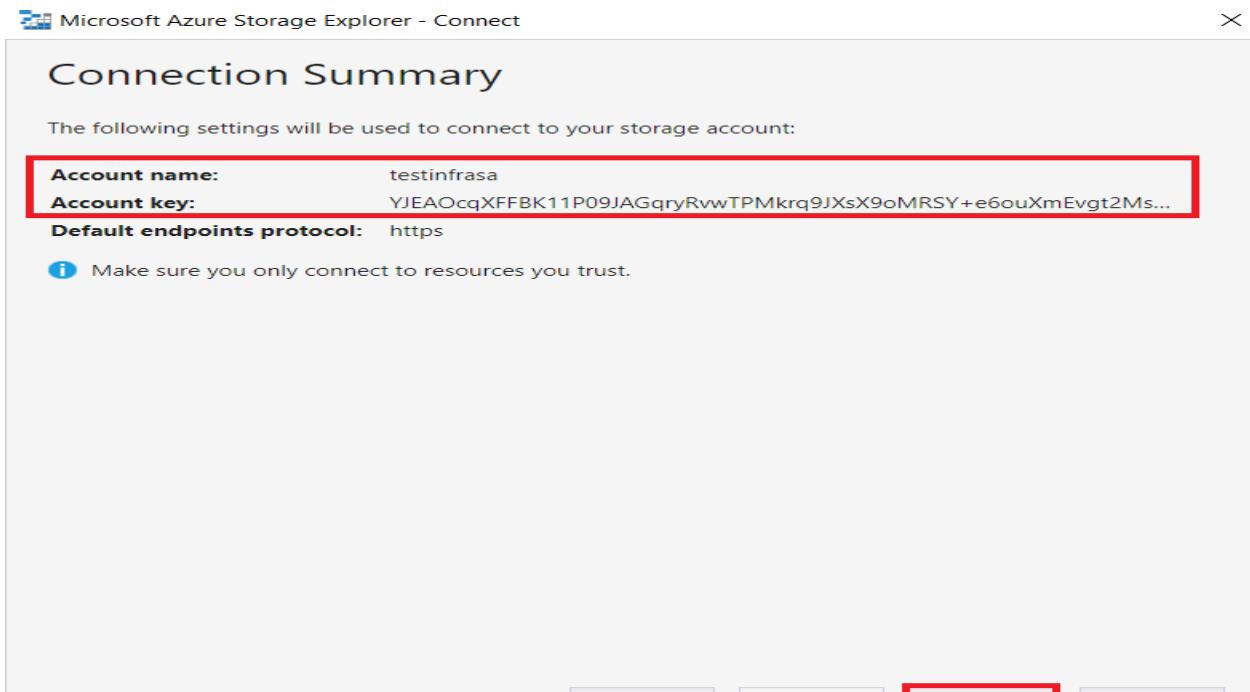


The screenshot shows the Microsoft Azure Storage Explorer interface. On the left, the 'EXPLORER' pane is open, showing a list of storage accounts under 'Storage Accounts'. A red box highlights the 'testinfrasa' account. A 'Connect' dialog box is overlaid on the right, titled 'Connect to Azure Storage'. It asks 'How do you want to connect to your Storage Account or service?'. Two options are available: 'Add an Azure Account' (radio button is unselected) and 'Use a connection string or a shared access signature URI' (radio button is selected). Below this, 'Use a storage account name and key' is also selected. At the bottom of the dialog are 'Back', 'Next', 'Connect', and 'Cancel' buttons.

Select **Use a storage account name and key** and paste your storage **Account name** and **Account key (Key1 or Key2** from the previous screen) as shown in the following screenshot:



Click on **Next** and you will see the following summary page:



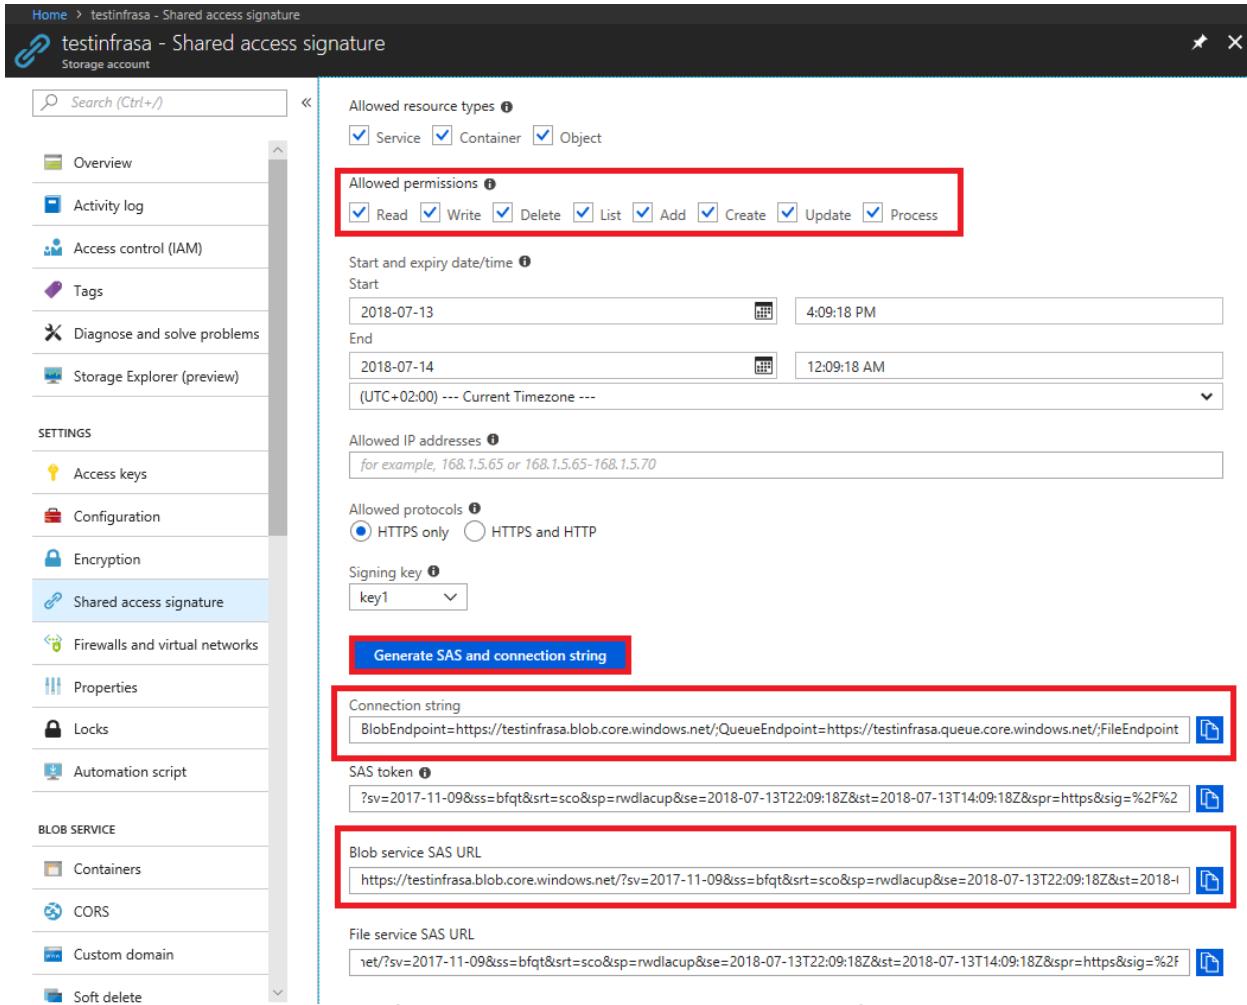
Finally, click on **Connect**. You'll be able to access your storage in Azure.

Accessing storage using shared access signatures

A **shared access signature (SAS)** is a URI that grants restricted access rights to Azure Storage. You can generate it at the storage account level and single-file level. As mentioned earlier, an account-level SAS can delegate access to multiple storage services, which are hosted within such as blobs, files, and queues.

Users can use a shared access signature URI to the customers who need a specified period of time to access the resources without informing the storage account name and access key.

You can go to the **Shared access signature** blade of your storage account to get this information (as shown in the following screenshot), where you can generate your SAS and configuration string:



Home > testinfrasa - Shared access signature

testinfrasa - Shared access signature

Storage account

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Storage Explorer (preview)

SETTINGS

Access keys

Configuration

Encryption

Shared access signature

Firewalls and virtual networks

Properties

Locks

Automation script

BLOB SERVICE

Containers

CORS

Custom domain

Soft delete

Allowed resource types:

- Service
- Container
- Object

Allowed permissions:

- Read
- Write
- Delete
- List
- Add
- Create
- Update
- Process

Start and expiry date/time:

Start: 2018-07-13 4:09:18 PM

End: 2018-07-14 12:09:18 AM (UTC+02:00) --- Current Timezone ---

Allowed IP addresses: for example, 168.1.5.65 or 168.1.5.65-168.1.5.70

Allowed protocols:

- HTTPS only
- HTTPS and HTTP

Signing key: key1

Generate SAS and connection string

Connection string: BlobEndpoint=https://testinfrasa.blob.core.windows.net/;QueueEndpoint=https://testinfrasa.queue.core.windows.net/;FileEndpoint=https://testinfrasa.file.core.windows.net/

SAS token: ?sv=2017-11-09&ss=bfqt&srt=sco&sp=rwddlacup&se=2018-07-13T22:09:18Z&st=2018-07-13T14:09:18Z&spr=https&sig=%2F%

Blob service SAS URL: https://testinfrasa.blob.core.windows.net/?sv=2017-11-09&ss=bfqt&srt=sco&sp=rwddlacup&se=2018-07-13T22:09:18Z&st=2018-07-13T14:09:18Z&spr=https&sig=%2F%

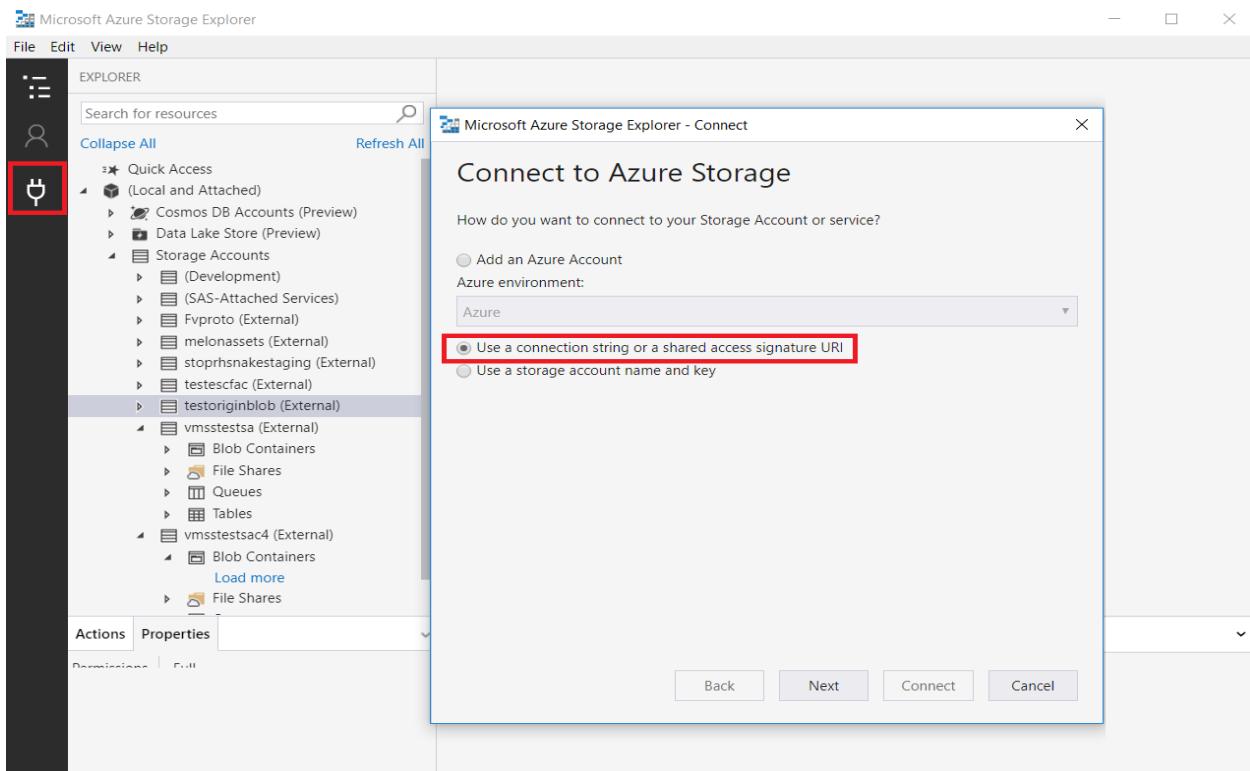
File service SAS URL: net?sv=2017-11-09&ss=bfqt&srt=sco&sp=rwddlacup&se=2018-07-13T22:09:18Z&st=2018-07-13T14:09:18Z&spr=https&sig=%2F%

Generating Shared access signature and connecting the string via the Azure Portal
Ameerpet / Kondapur Hyderabad

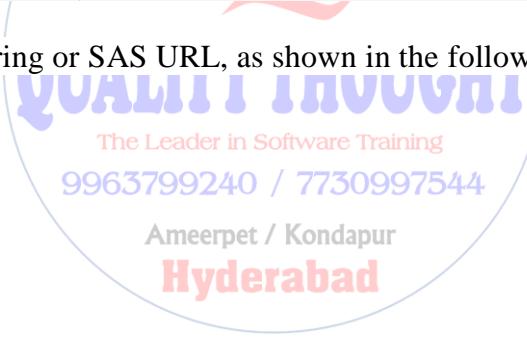
You can configure the following information according to your requirements:

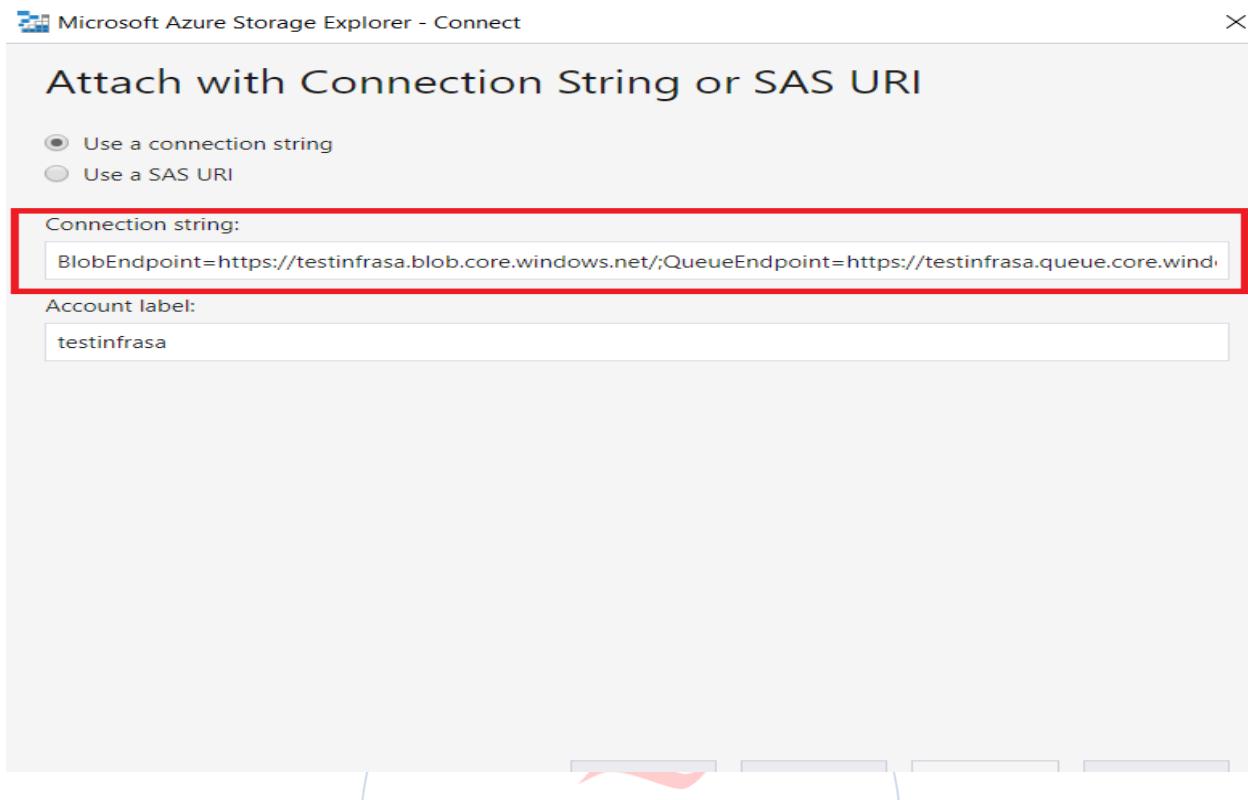
- The allowed permission
- Validate period (start time and end time)
- Allowed IP addresses
- Allowed protocol

While using Storage Explorer to access your storage in Azure, you should choose use connection string or SAS URL, as shown in the following screenshot:

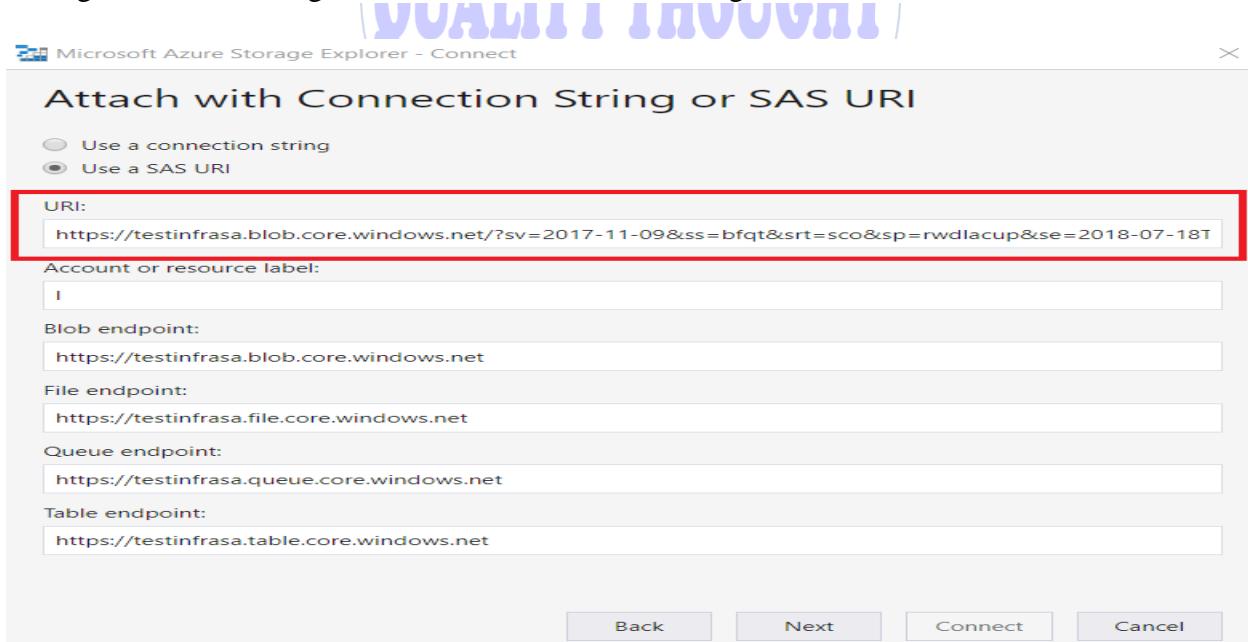


You can use connection string or SAS URL, as shown in the following screenshot:

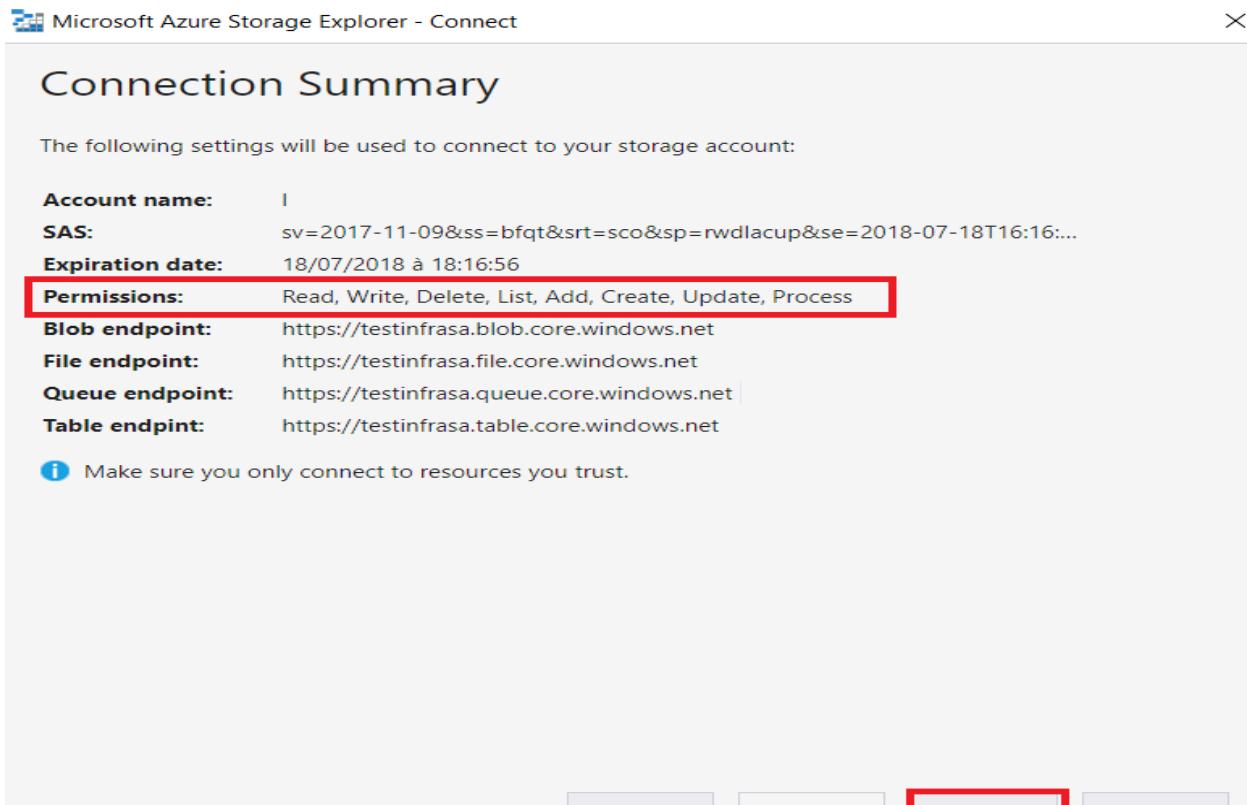




Filling shared access signature is shown in the following screenshot:



Then click on **Next**. You'll see a summary page as shown in the following screenshot, with the exact permissions you've requested, the validation period, and so on before validating your connection information. Click on **Connect**:



Setting and retrieving properties and metadata

Objects in Azure Storage support the following two types of data (apart from the data they contain):

- System properties:** These are defined by Microsoft and exist on each storage resource, and some of them are read-only and cannot be set
- User-defined metadata:** This is the metadata that users can specify on a given resource in the form of a name-value pair

Users can set these values using Azure RESTful APIs. Refer to the following URL for more information: <https://docs.microsoft.com/en-us/rest/api/storageservices/set-container-metadata>.

The following URL shows the usage of Azure Storage Client Library for .NET: <https://docs.microsoft.com/en-us/azure/storage/blobs/storage-properties-metadata>.

Implementing hybrid storage solutions

Microsoft Azure offers a range of Azure Storages in hybrid scenarios such as StorSimple and Azure File Sync. Let's take a look at each of them.

Implementing Azure StorSimple

Microsoft Azure StorSimple is an integrated storage that manages storage tasks between an on-premise virtual array running in a hypervisor and cloud storage in Azure. Azure StorSimple virtual array is an excellent fit for storing infrequently accessed archival data.

To deploy the StorSimple Device Manager service for StorSimple Virtual Array, you can refer to the following URL: <https://docs.microsoft.com/en-us/azure/storsimple/storsimple-virtual-array-manage-service>.

Implementing Azure File Sync

Use Azure File Sync to centralize your organization's file shares in Azure Files, while keeping the flexibility, performance, and compatibility of an on-premise file server. Azure File Sync transforms the Windows Server into a quick cache of your Azure File share. You can use any protocol that's available on Windows Server to access your data locally, including SMB, NFS, and FTPS. You can have as many caches as you need across the World.

To deploy an Azure File Sync, you should have an Azure Storage account and an Azure File share that are in the same region that you want to deploy Azure File Sync in.

To make sure your region has Azure File Sync available, refer to the following URL: <https://docs.microsoft.com/en-us/azure/storage/files/storage-sync-files-planning#region-availability>.

Moving data to and from Azure Storage

There are a wide range of options to help users move data to and from Azure Storage.

Transferring data with AzCopy

AzCopy is a command-line utility for transferring data to and from Azure Storage. AzCopy is available on Windows and Linux. It can be used for:

- Copying data to and from Microsoft Azure Blob and File storage within a storage account
- Copying data between different storage accounts

To know more about transferring data with the AzCopy on Windows, you can refer to the following URL: <https://docs.microsoft.com/en-us/azure/storage/common/storage-use-azcopy?toc=%2fazure%2fstorage%2ffiles%2ftoc.json>.

To know more about transferring data with AzCopy on Linux, you can refer to the following URL: <https://docs.microsoft.com/en-us/azure/storage/common/storage-use-azcopy-linux?toc=%2fazure%2fstorage%2ffiles%2ftoc.json>.

Azure Storage Data Movement Library

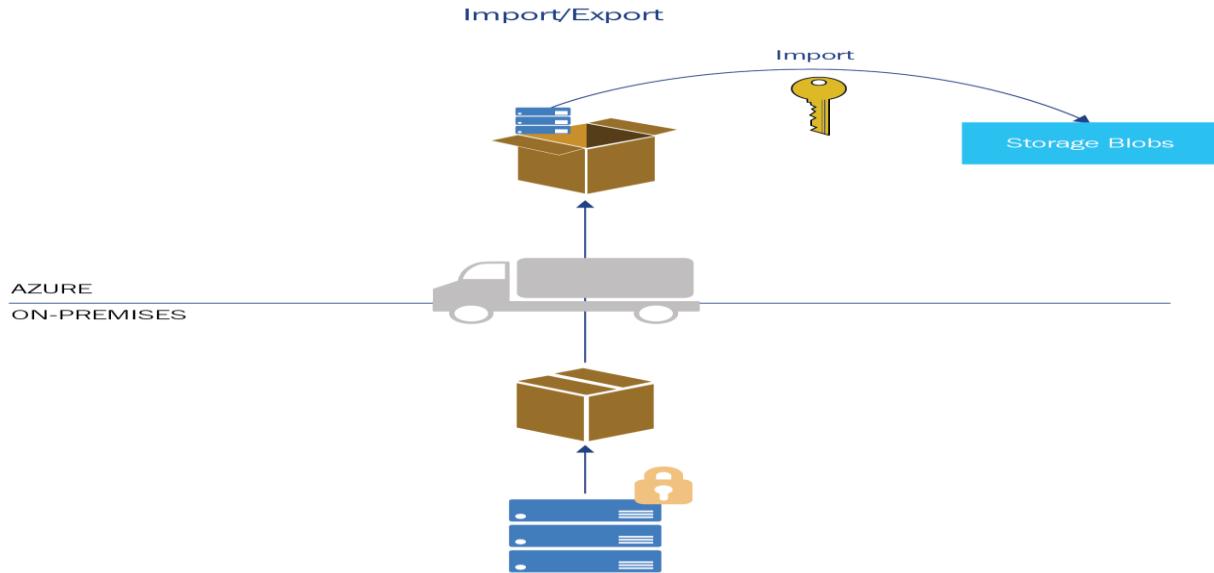
Azure Storage Data Movement Library (DML) for .NET is an open source project that exposes the core data movement framework, which powers AzCopy. It is designed for high-performance copying of data to and from Azure.

Cross-premise data transfer

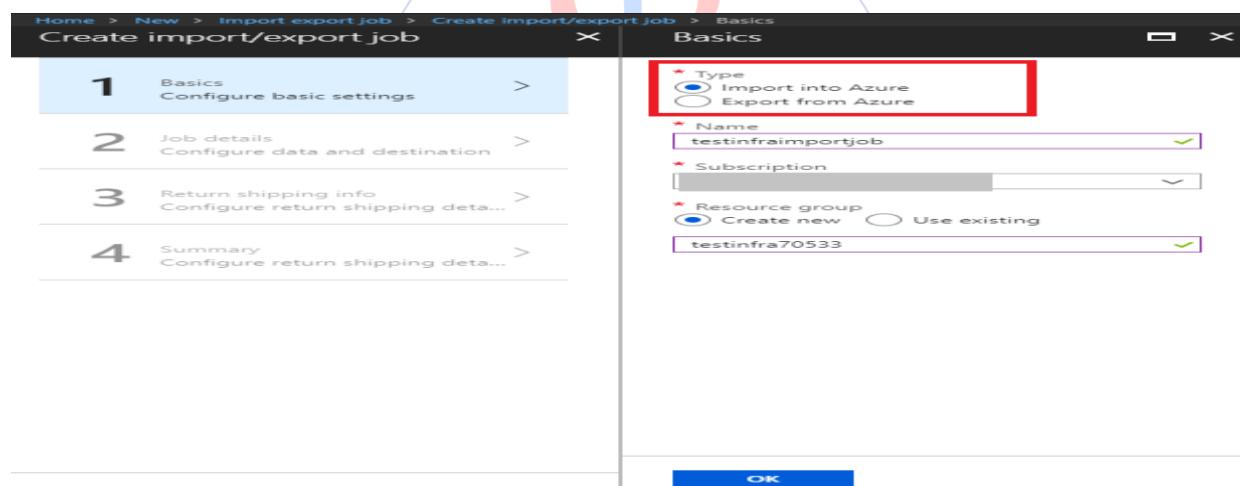
Microsoft Azure provides a couple of cross-premise data transfer options; let's take a look at each of them.

Azure Import/Export

Azure Import/Export is a service you can use to securely import large sets of data from on-premise datacenters or other cloud datacenters to Azure Blob storage (the general purpose v1 type and Azure Files type) by shipping disk drives Azure data center , or in reverse, to transfer data from Azure Blob storage to disk drives and then shipping to customers' on-premise sites. The following schema shows how it works:



You should create an import or export job via Azure Portal and fill in all the shipping information, as shown in the following screenshot:



Azure Data Box

The Azure Data Box tool helps users to transfer large amounts of data (such as terabytes of data) to Azure in a secure and quick way. Users are able to order the Data Box directly through the Azure Portal. As described by Microsoft, after filling data, users should return it to Azure Data Center so that the data can be uploaded in Azure. To see how Data Box works, refer to the following screenshot:



Azure Data Box Disk

Similar to Azure Data Box, which is a portable, secure, quick, and simple way to move large datasets into Azure, Azure Data Box Disk is a lower-capacity (and easier to move) choice. After filling data, users should also return it to Azure Data Center so that the data can be uploaded in Azure. The following shows Data Box Disk and a flowchart explaining how it works:



Implementing data storage services

Azure offers a variety of **Platform as a Service (PaaS)** services, also known as **Database as a Service (DBaaS)** which, removes the need for you to manage the underlying operating system and database-server platform.

SQL Database

Azure SQL Database is a managed database service that is different from AWS RDS, which is a container service. As a PaaS offering, this frees you from performing updates and maintenance tasks, and includes built-in features that provide fault tolerance and scalability.

SQL Database offers logical servers that can contain single or multiple SQL databases. SQL Database has two different pricing models:

- vCore-based purchasing model
- DTU-based purchasing model

SQL Database also provides options such as columnstore indexes for extreme analytic analysis and reporting, and in-memory **online transaction processing (OLTP)** for extreme transnational processing. Microsoft manages all patching and updating work and all the underlying infrastructure.

Azure Database for MySQL

Azure Database for MySQL is a relational database service fully managed by Microsoft Azure, and is based on the MySQL Community Edition database engine. It provides users with built-in high availability and dynamic scaling as well as unparalleled security and compliance with a flexible pricing model.

To know more about creating an Azure Database for MySQL Server using Azure CLI, you can take a look at the following URL: <https://docs.microsoft.com/en-us/azure/mysql/quickstart-create-mysql-server-database-using-azure-cli>.

To know more about creating an Azure Database for MySQL server using the Azure portal, you can refer to the following URL: <https://docs.microsoft.com/en-us/azure/mysql/quickstart-create-mysql-server-database-using-azure-portal>.

Azure Database for PostgreSQL

Azure Database for PostgreSQL is a relational database service fully managed by Microsoft Azure for developers based on the community version of the open-source PostgreSQL database engine. Users get built-in high availability and capability to scale in seconds and benefit from its unparalleled security and compliance as well as a flexible pricing model.

To know more about creating an Azure Database for PostgreSQL server in the Azure portal, you can take a look at the following URL: <https://docs.microsoft.com/en-us/azure/postgresql/quickstart-create-server-database-portal>.

To know more about creating an Azure Database for PostgreSQL using the Azure CLI, you can refer to the following URL: <https://docs.microsoft.com/en-us/azure/postgresql/quickstart-create-server-database-azure-cli>.

Database-managed instances

Azure SQL Database Managed Instance is a Platform as a Service database offer, which provides nearly 100% of the SQL Server features. It enables users to migrate the on-premise SQL Server instance, which would be 100% compatible. Every instance is fully isolated from the other customer instance and assigned a private IP address.

For more information, you can refer to the following URL: <https://blogs.msdn.microsoft.com/sqlserverstorageengine/2018/03/07/what-is-azure-sql-database-managed-instance-2/>.

Azure SQL Data Warehouse

Azure SQL Data Warehouse is a **massively parallel processing (MPP)** distributed database system that helps users focus on simply loading and querying a database without thinking about the underlying infrastructure and OS configuration.

Azure SQL Data Warehouse is actually a distributed system using nodes that work together to supply the data for any queries.

There are several differences between Azure SQL Database and Azure SQL Data Warehouse. SQL DB is for OLTP and make applications with individual updates, inserts, and deletes; SQL DW is for **online analytical processing (OLAP)** and is an approach to answering **multi-dimensional analytical (MDA)** queries swiftly in computing systems.

Cosmos DB

In July 2017, Microsoft announced Azure Cosmos DB, which is the next big leap in globally distributed, at-scale, cloud databases. Azure Cosmos DB is Microsoft's globally distributed, multi-model database. With the click of a button, Azure Cosmos DB enables you to elastically and independently scale throughput and storage across any number of Azure's geographic regions. It offers throughput, latency, availability, and consistency guarantees with comprehensive **service level agreements (SLAs)**, something no other database service can offer.

Cosmos DB exposes multiple well-defined consistency models. It allows developers to store and query their data in its original form. It exposes graphs, documents, key-values, column-family data models, and will enable others. The multi-model and multitude of APIs to access and query data includes SQL and various popular OSS APIs. It also allows multi-API capabilities and removes friction, allowing you to build with any data model and API. Cosmos DB is all about providing intelligent choices to developers and enabling you to build planet scale apps.

Configuring Content Delivery Network

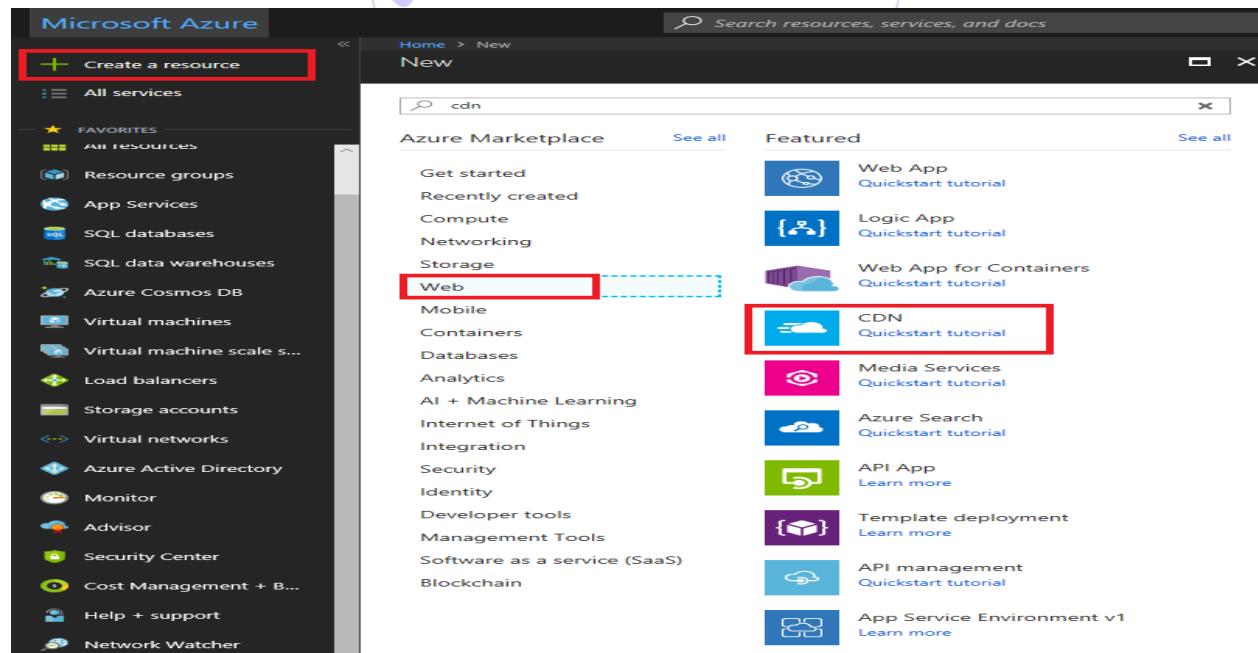
Content Delivery Network (CDN) is a distributed network that delivers web content to users based on their geographic location by choosing the closest edge servers in **point-of-presence (POP)** locations, so that the latency of distribution is maximum-ally reduced. Azure CDNs carry a significant portion of the world's internet traffic. It provides users with a global solution for delivering high-bandwidth content hosted in Azure or any other location.

CDN is used for:

- Delivering **static content** by caching files such as images, style sheets, documents, files, client-side scripts, and HTML pages
- Accelerating what to serve **dynamic content** using CDN POPs

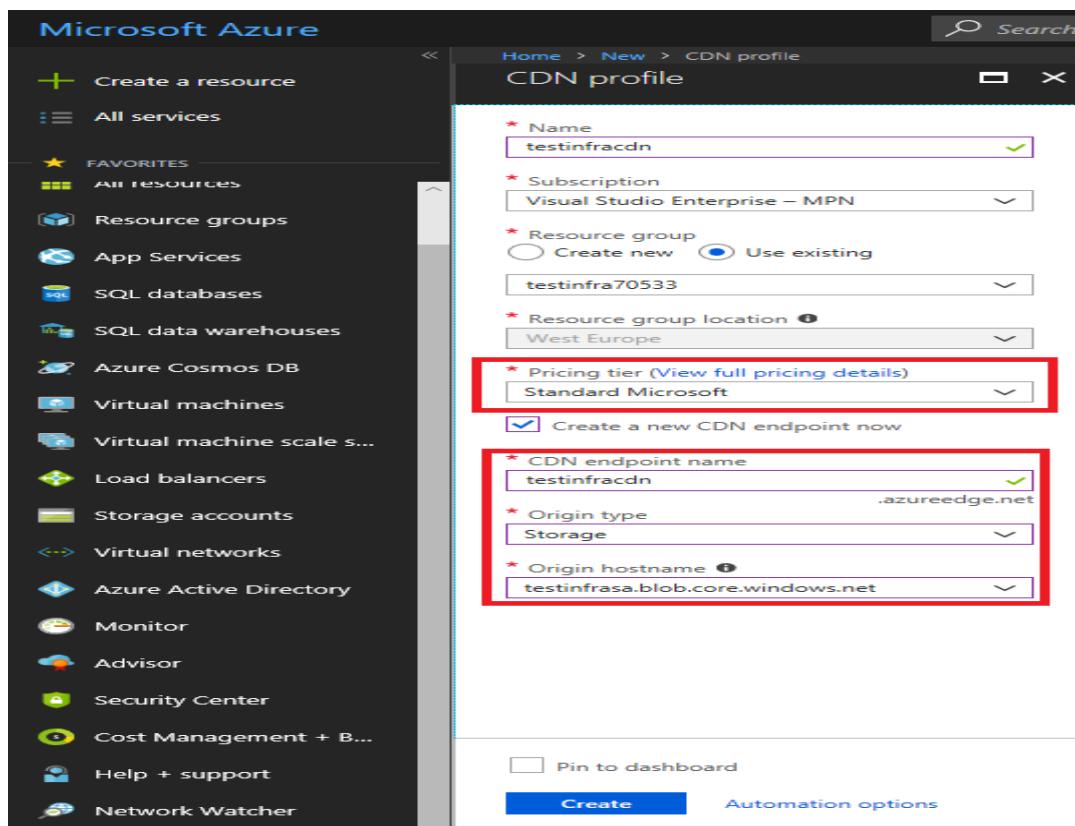
Creating a CDN profile

To implement Azure CDN, you should start by creating a CDN profile. You can go to the Azure Portal, click on **Create a resource**, then find cdn in the web category, as shown in the following screenshot:

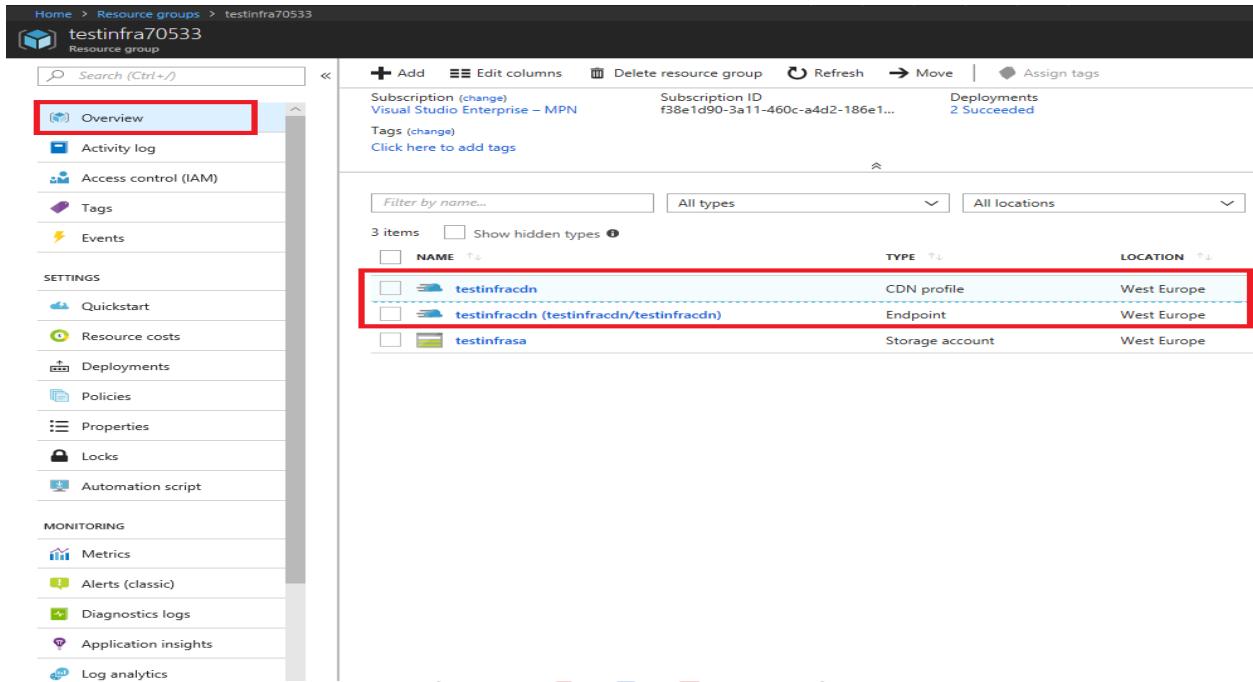


When configuring the CDN endpoint, you should choose the CDN pricing tier and the origin server. You can choose a pricing tier from Standard Verizon (S1), Standard Akamai (S2), to Standard Microsoft (S3). For more details about the pricing tier of Azure CDN, refer to the following URL: <https://azure.microsoft.com/is-is/pricing/details/cdn/>.

Azure supports custom CDN endpoints of any origin. You can even create an origin in your own datacenter, an origin provided by third-party cloud providers. In the following we'll point the origin server to the host of the storage account:



After clicking on **Create**, Azure will create a **CDN profile** as well as the **Endpoint** that you've created previously, as shown in the following screenshot:



Subscription (change) Visual Studio Enterprise – MPN
Subscription ID f38e1d90-3a11-460c-a4d2-186e1...
Tags (change)
Click here to add tags

Filter by name... All types All locations

NAME	TYPE	LOCATION
testinfracdn	CDN profile	West Europe
testinfracdn (testinfracdn/testinfracdn)	Endpoint	West Europe
testinfrasa	Storage account	West Europe

Custom domains over HTTPS

We can also add custom domain mapping to your CDN endpoint and enable custom domain HTTPS. To use this feature of the Azure CDN Premium offering from Verizon, you can refer to the following URL: <https://docs.microsoft.com/en-us/azure/storage/blobs/storage-https-custom-domain-cdn>.

9963799240 / 7730997544

Implementing a business continuity and disaster recovery (BCDR) strategy in Azure

BCDR is a popular term nowadays. As an organization, the requirement is to adopt a business continuity and disaster recovery (BCDR) strategy that keeps the data on-premise and on the cloud so it's secure, and keeps all apps and workloads up and running even when planned and unplanned outages occur.

Planning a BCDR strategy

Every business continuity planning or disaster recovery plan potentially begins with a general business analysis, which contains some critical terms that we need to understand: These terms are listed as follows:

- **Recovery time objective (RTO):** This is the maximum acceptable length of time that your application can be offline
- **Recovery point objective (RPO):** This is the maximum acceptable length of time during which data might be lost from your application due to a major incident
- **Uptime:** This is a measure of the time a system runs over a given period of time per year
- **Downtime:** This is the period of time that a system fails to provide or perform its primary function as expected

BCDR in Azure

Cloud providers such as Microsoft Azure provide capabilities that support availability and a variety of disasterrecovery services to adapt to different scenarios by building Azure recovery services to contribute to an enterprise-level BCDR strategy. The following are two disaster recovery services:

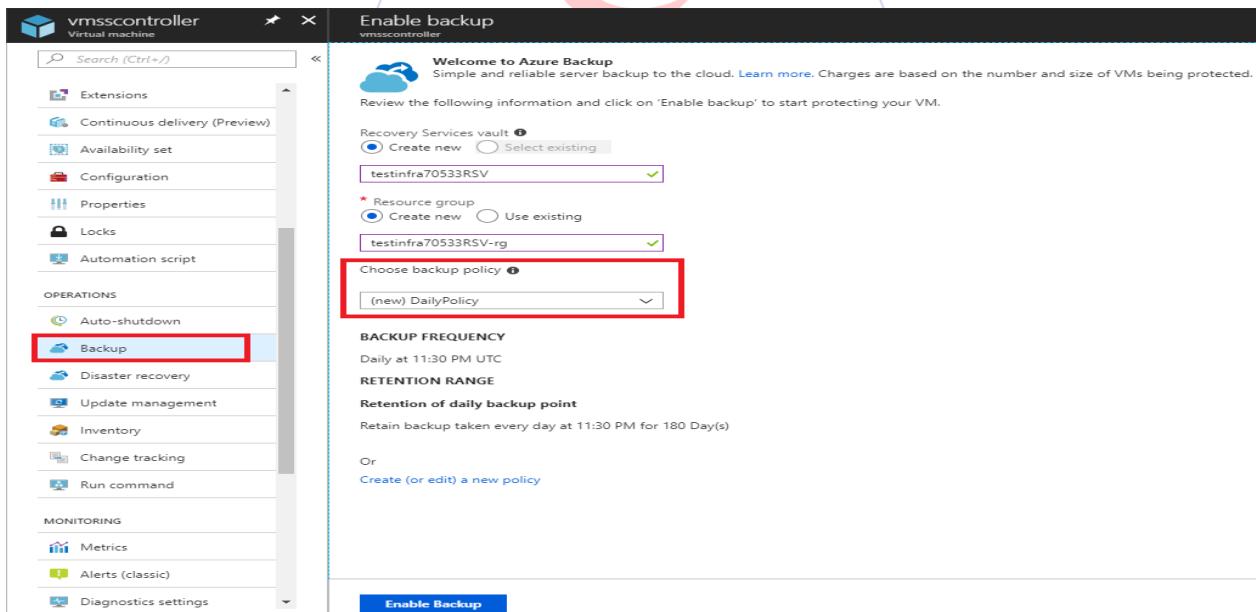
- **Azure Backup service:** This keeps your data safe and recoverable by backing it up to Azure.
- **Site Recovery service:** This is used to replicate workloads running on physical and virtual machines from a primary site to a secondary region. Users can failover to a secondary location if an outage occurs at the primary region and users can fall back to it when the primary region is back to normal.

Implementing Azure Backup

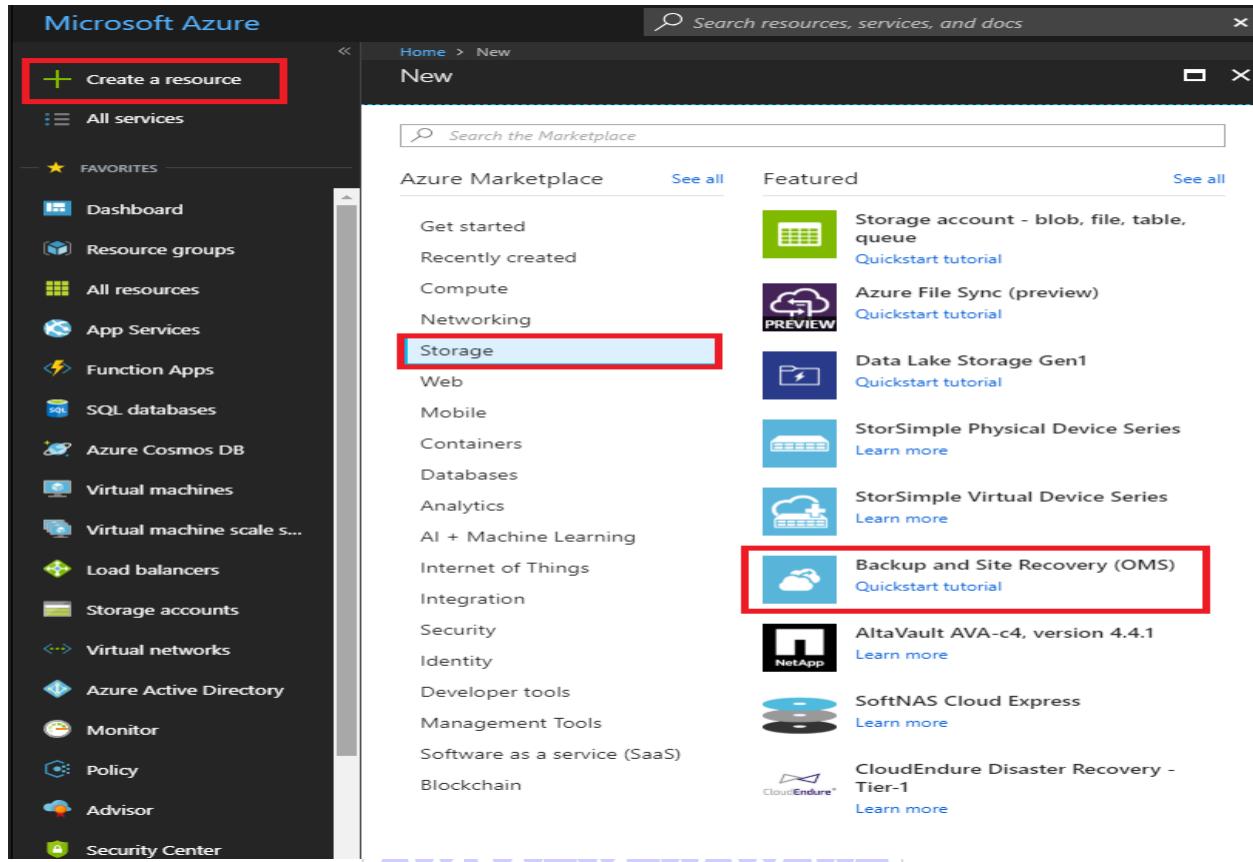
Azure Backup is a cloud service that allows users to backup and restore data in Azure. Azure Backup offers a couple of components or agents to help users deploy depending on what they want to protect. With these components or agents, Azure Backup can back up files, folders, Hyper-V or VMware virtual machines on-premise, Microsoft SQL Server, Microsoft Sharepoint, Microsoft exchange, or Azure IaaS VM.

To know more about the different components of Azure Backup, you can refer to the following URL: <https://docs.microsoft.com/en-us/azure/backup/backup-introduction-to-azure-backup>.

To implement **Azure Backup**, you can enable it directly in the **Backup** blade of your Azure VM instance, and create a new (or use an existing) Recovery Service vault, as shown in the following screenshot:



Alternatively, you can go to the Azure Portal and click on **Create a resource**. Then, go to the **Storage** category and select the **Backup and Site Recovery (OMS)** option:



The screenshot shows the Microsoft Azure 'New' blade. On the left, a sidebar lists various service categories like Dashboard, Resource groups, App Services, etc. A red box highlights the 'Create a resource' button at the top of this sidebar. The main area is titled 'New' and contains sections for 'Azure Marketplace' and 'Featured'. In the 'Azure Marketplace' section, a search bar says 'Search the Marketplace'. Below it, a list of services includes 'Get started', 'Recently created', 'Compute', 'Networking', and 'Storage'. A red box highlights the 'Storage' option. Other listed services include Web, Mobile, Containers, Databases, Analytics, AI + Machine Learning, Internet of Things, Integration, Security, Identity, Developer tools, Management Tools, Software as a service (SaaS), and Blockchain. In the 'Featured' section, several services are listed with their icons and brief descriptions. One service, 'Backup and Site Recovery (OMS)', is highlighted with a red box.

Then, you can fill complete the **Recovery Services vault** dialog to create a new instance of the RS vault:



Recovery Services vault

Recovery Services vault

* Name
testinfra70533RSV

* Subscription

* Resource group
 Create new Use existing
testinfra70533RSV-rg

* Location
North Europe

Pin to dashboard

Create **Automation options**

After clicking on **Create**, you can go to the resource group where the recovery services vault has been created:

testinfra70533RSV-rg

Resource group

Overview

Activity log

Access control (IAM)

Tags

Events

SETTINGS

Quickstart

Resource costs

Add Edit columns Delete resource group Refresh Move Assign tags

Subscription (change) Subscription ID 60bad1d7-71fd-4e40-b6b9-1016a... Deployments 1 Succeeded

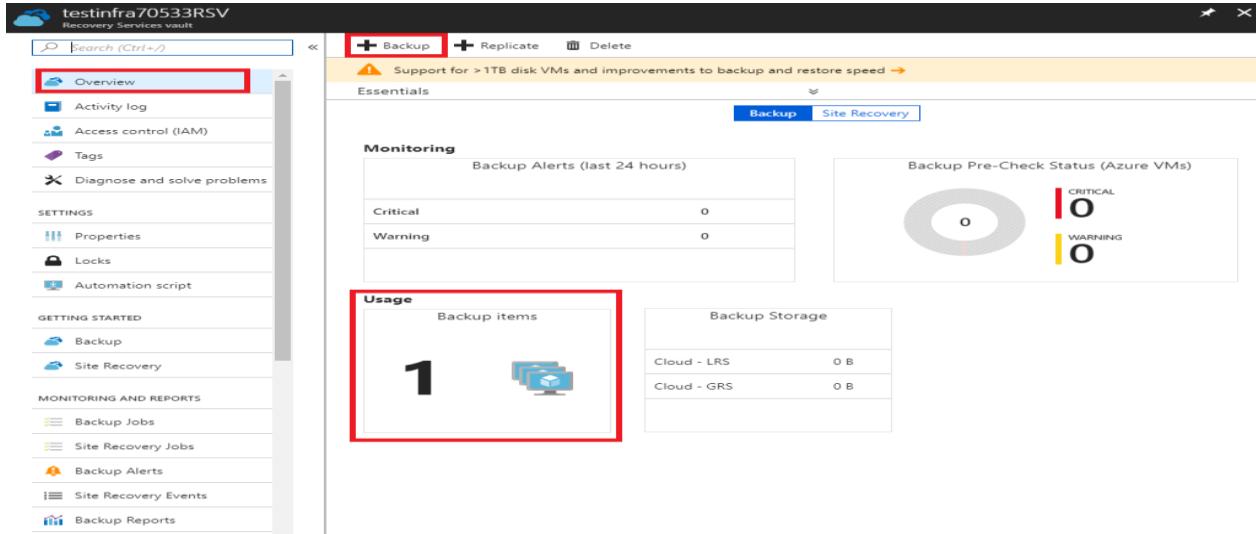
Tags (change) Click here to add tags

Filter by name... All types All locations

1 items Show hidden types

NAME	TYPE	LOCATION
testinfra70533RSV	Recovery Services vault	North Europe

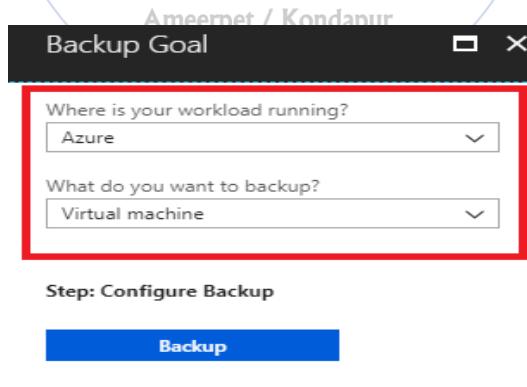
The RS Vault launches. You can see that there is a backup item that has been activated if you created your RS vault by enabling **Backup** for Azure VM, as shown in the following screenshot:



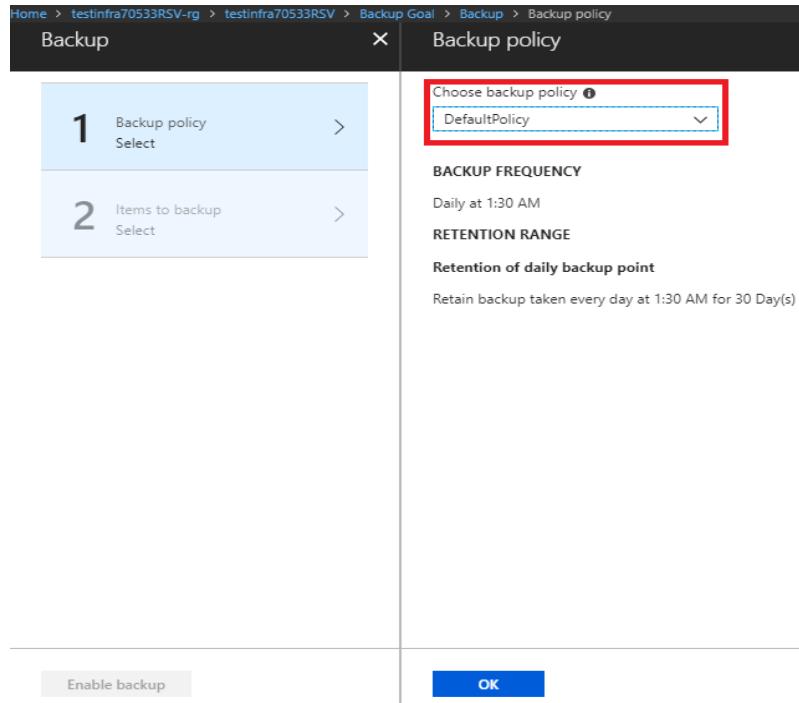
At this stage, by clicking on **Backup**, you can start to back up one or more additional Azure VMs in the cloud. As shown in the screenshot, you should firstly be clear about your objectives in creating a backup. This could be done by answering the following questions:

- Where is your workload? In the cloud or on-premise?
- What do you want to back up? Azure VMs, File share, or SQL Server in Azure VM?

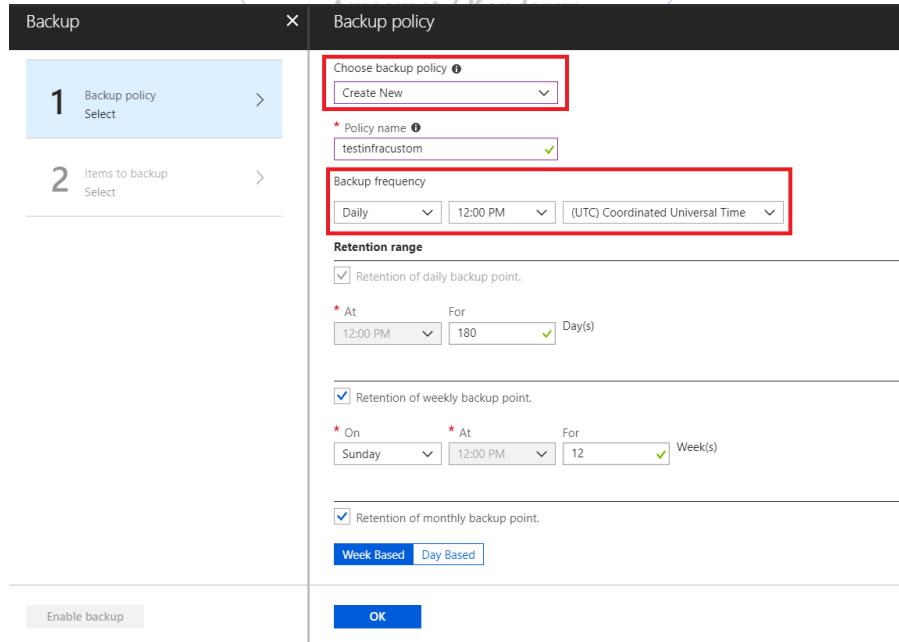
Refer to the following screenshot:



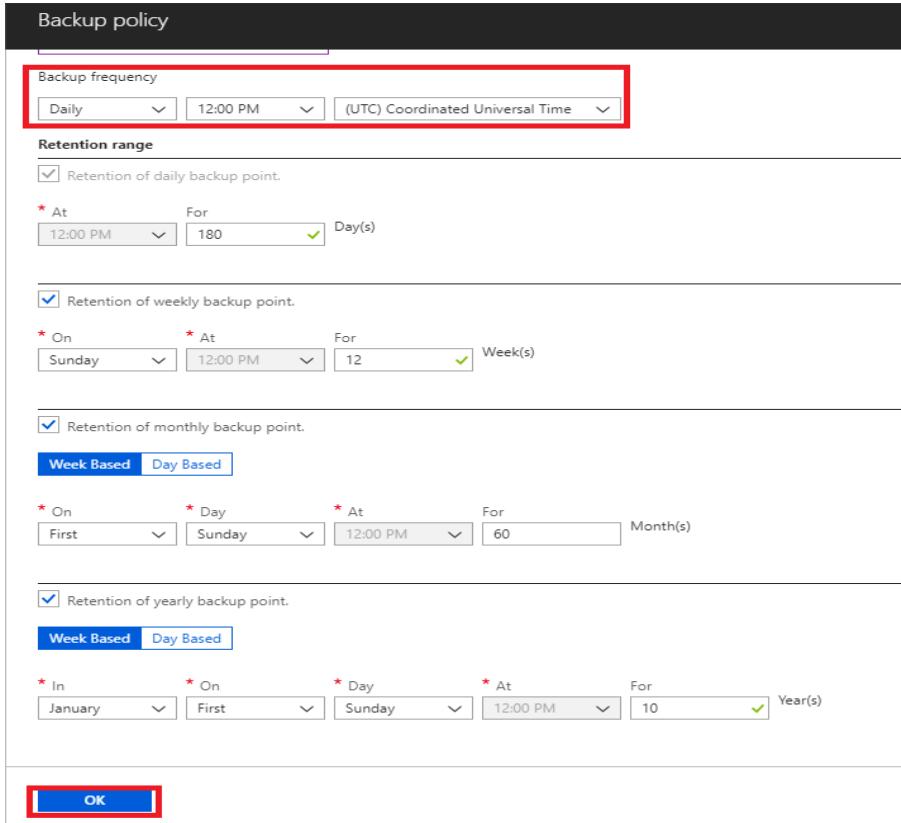
After choosing your backup goal, Azure workload, and Azure VM, for example, you can go to Azure Backup Policy to create a new policy or use an existing policy (the default or other pre-existing policy). As shown in the following screenshot, you can select the **DefaultPolicy** or create a new one:



In order to create a new backup policy, you should choose a **Backup frequency** and **Retention** range, which address most your requirements, as shown in the following screenshot:



The **Retention range** is set as illustrated in the following screenshot:



Backup policy

Backup frequency

- At: 12:00 PM
- For: (UTC) Coordinated Universal Time

Retention range

Retention of daily backup point.

* At: 12:00 PM For: 180 Day(s)

Retention of weekly backup point.

* On: Sunday At: 12:00 PM For: 12 Week(s)

Retention of monthly backup point.

Week Based **Day Based**

* On: First Day: Sunday At: 12:00 PM For: 60 Month(s)

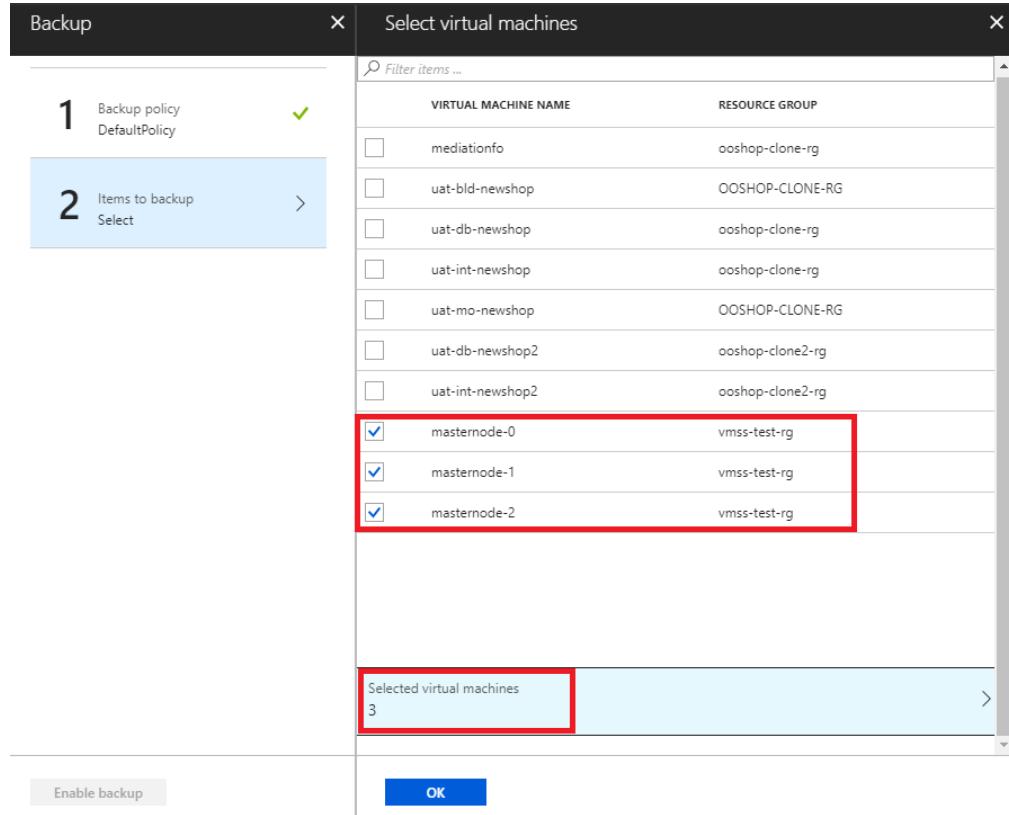
Retention of yearly backup point.

Week Based **Day Based**

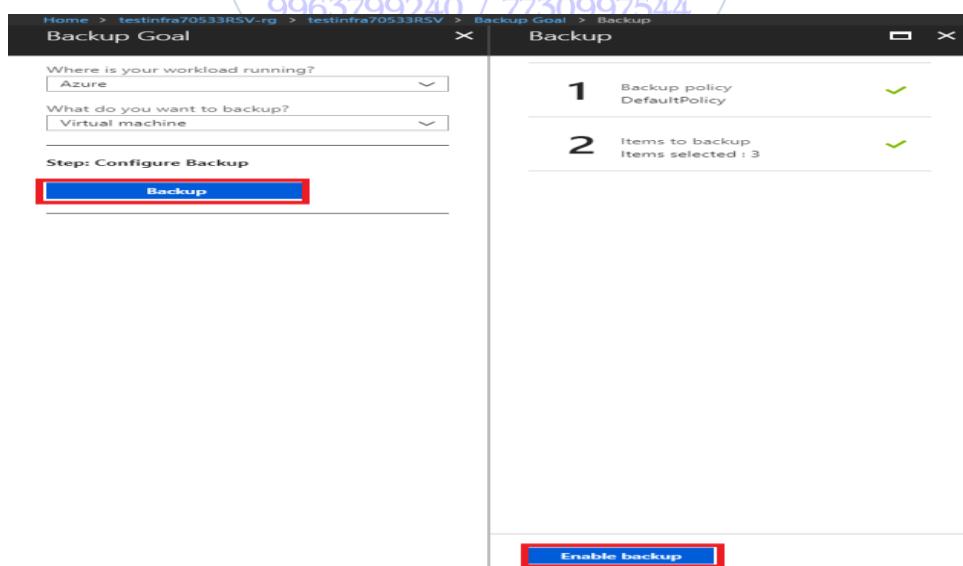
* In: January On: First Day: Sunday At: 12:00 PM For: 10 Year(s)

OK

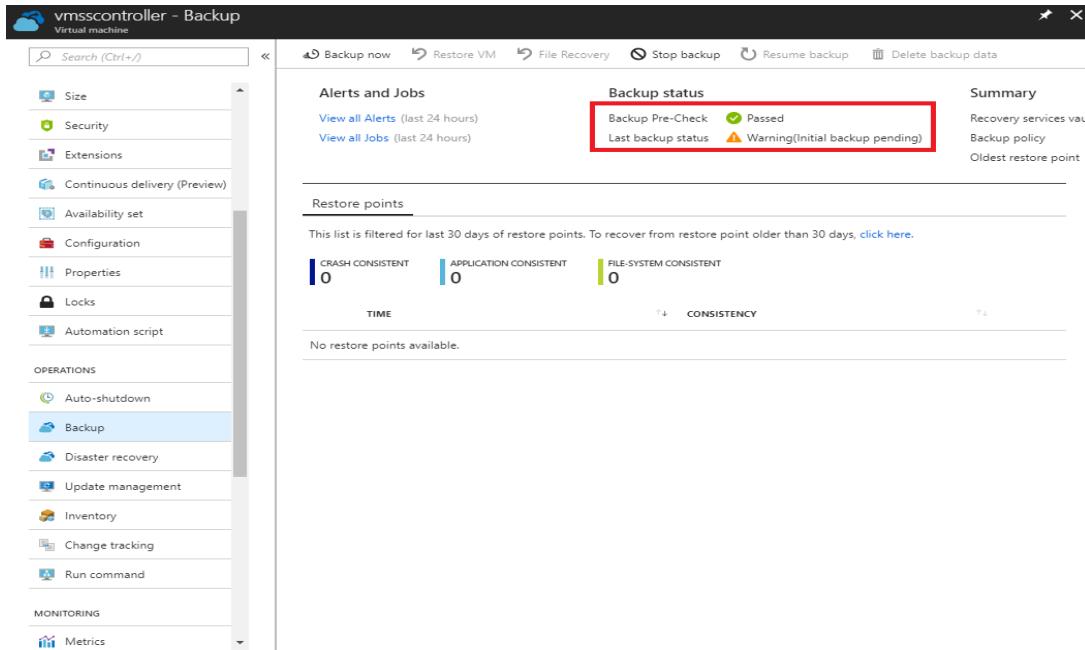
Now click on **OK**. You can go to the next step to choose items to back up. You can choose one or multiple items, here we choose 3 items; from the same resource group, **Selected virtual machines** displayed as 3 as shown in the following screenshot:



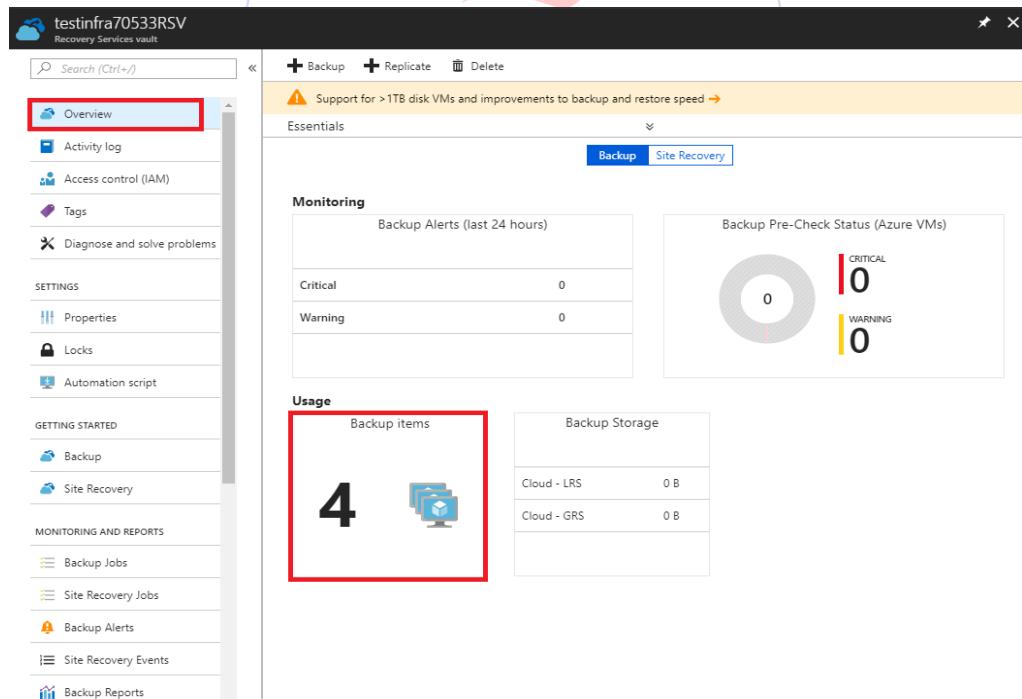
When you click on **OK** and then on **Enable backup**, the backup process will start, as shown in the following screenshot:



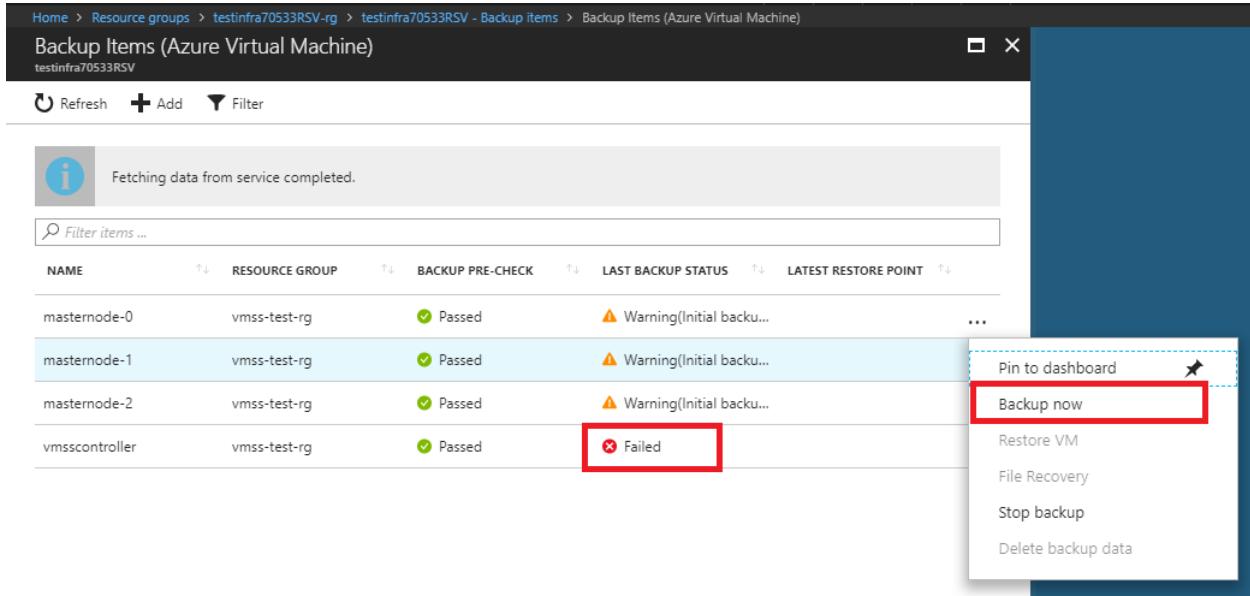
It will take a few minutes to complete a backup of the VM. But you can check the backup status of the VM by going to its **Backup** blade, as shown in the following screenshot:



Alternatively, you can go to the **Overview** blade in the RS vault, as shown in the following screenshot, and click on **Backup items** to check the status of all the items:



As you can see in the following screenshot, a Failed message will be displayed if the backup fails. You can right-click the item and select **Backup now** to launch a backup progress:



Home > Resource groups > testinfra70533RSV-rg > testinfra70533RSV - Backup Items > Backup Items (Azure Virtual Machine)

testinfra70533RSV

Refresh Add Filter

Fetching data from service completed.

Filter items ...

NAME	RESOURCE GROUP	BACKUP PRE-CHECK	LAST BACKUP STATUS	LATEST RESTORE POINT
masternode-0	vmss-test-rg	Passed	Warning(Initial backu...)	
masternode-1	vmss-test-rg	Passed	Warning(Initial backu...)	
masternode-2	vmss-test-rg	Passed	Warning(Initial backu...)	
vmsscontroller	vmss-test-rg	Passed	Failed	

...

Pin to dashboard

Backup now

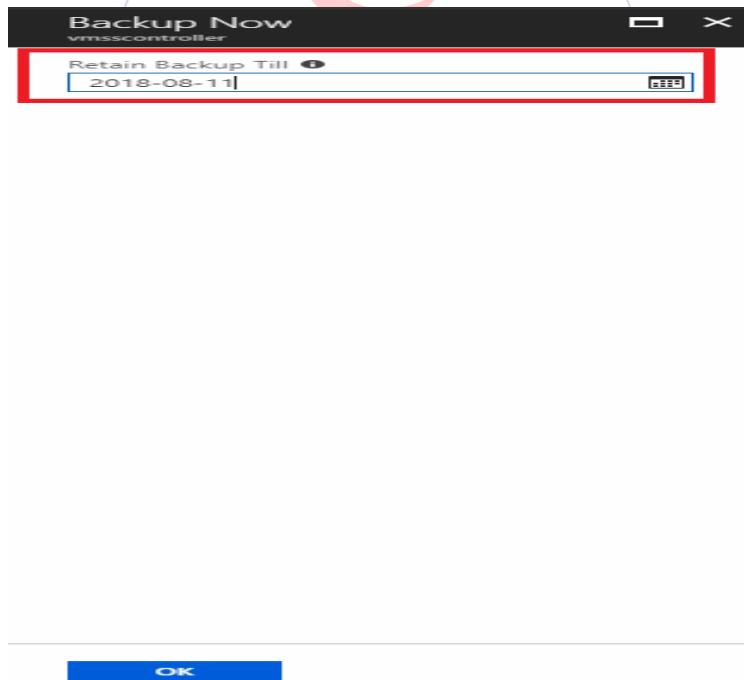
Restore VM

File Recovery

Stop backup

Delete backup data

If you want to accept the backup retention policy of 30 days, just use the default **Retain Backup Till** date, as shown in the following screenshot:



After a couple of minutes, you can go to your RS vault to recheck the backup status of your VMs as shown in the following screenshot, which means the last backup (for three VMs) was successful:

Home > Resource groups > testinfra70533RSV-rg > testinfra70533RSV - Backup items > Backup Items (Azure Virtual Machine)

Backup Items (Azure Virtual Machine)

testinfra70533RSV

Refresh Add Filter

Fetching data from service completed.

Filter items ...

NAME	RESOURCE GROUP	BACKUP PRE-CHECK	LAST BACKUP STATUS	LATEST RESTORE POINT
masternode-0	vmss-test-rg	Passed	Warning(Initial backu...)	...
masternode-1	vmss-test-rg	Passed	Success	7/12/2018, 4:49:51 PM
masternode-2	vmss-test-rg	Passed	Success	7/12/2018, 4:48:17 PM
vmsscontroller	vmss-test-rg	Passed	Success	7/12/2018, 4:27:20 PM

You can also achieve the same result using Azure PowerShell as shown in the following URL: <https://docs.microsoft.com/en-us/azure/backup/quick-backup-vm-powershell>.

You can also use the Azure CLI by referring to the documentation at the following URL: <https://docs.microsoft.com/en-us/azure/backup/quick-backup-vm-cli>.

To create a backup vault via the ARM template, refer to the following code:

```
{
  "type": "Microsoft.RecoveryServices/vaults",
  "apiVersion": "2018-01-10",
  "name": "[parameters('vaultName')]",
  "location": "[parameters('location')]",
  "sku": {
    "name": "RS0",
    "tier": "[parameters('skuTier')]"
  },
  "properties": {}
}
```

There are also a couple of useful Azure Resource Manager templates for Azure Backup, which you can find at the following URL: <https://docs.microsoft.com/en-us/azure/backup/backup-rm-template-samples>.

Planning and implementing Azure Site Recovery

Azure Site Recovery (ASR) can manage replication for the following points:

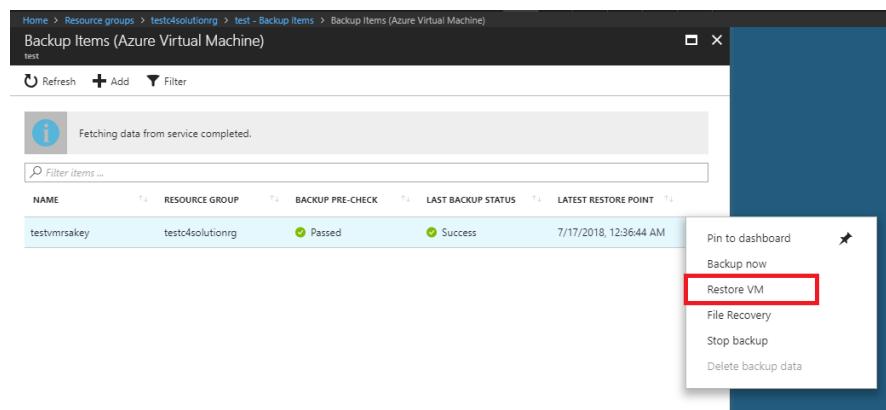
- Migrating on-premise Hyper-V VMs, VMware VMs, and physical servers to Azure
- Migrating Azure VMs between Azure regions
- Migrating AWS Windows-based Instances to Azure IaaS VMs

ASR is a powerful Azure service, but is simple to use. It provides an on-click restore facility and can replicate any workload running on a machine that's supported for replication. Refer to the following link for more information: <https://docs.microsoft.com/en-us/azure/site-recovery/site-recovery-overview>.

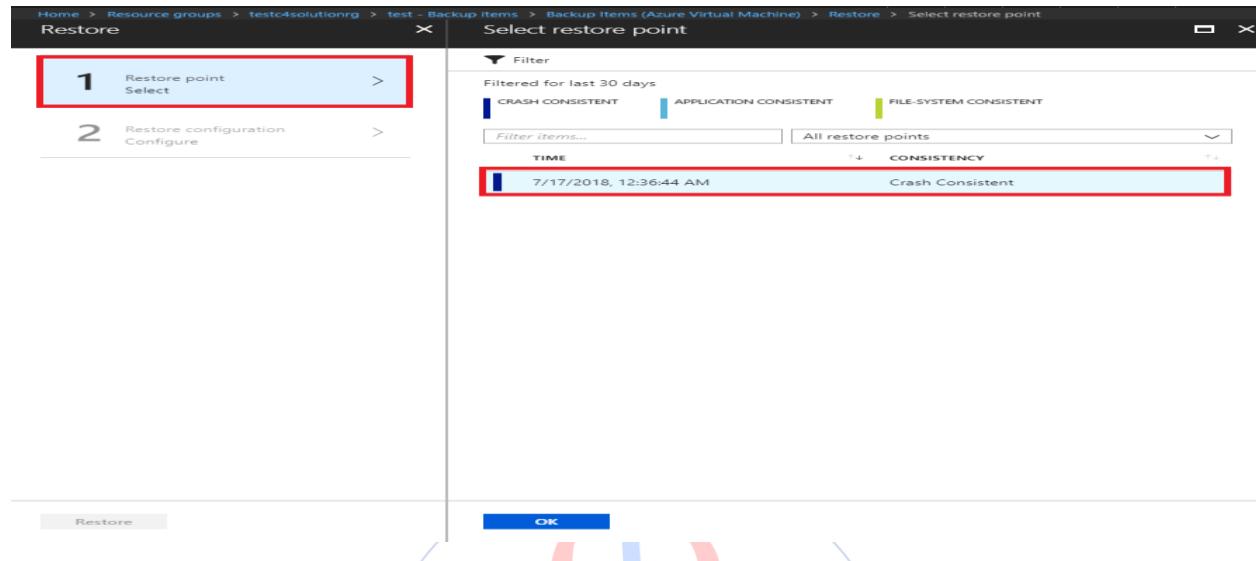
ASR can be used in simple-to-very complex use cases. Now let's take a look at two general use scenarios in the upcoming subsections.

Restoring VMs with ASR

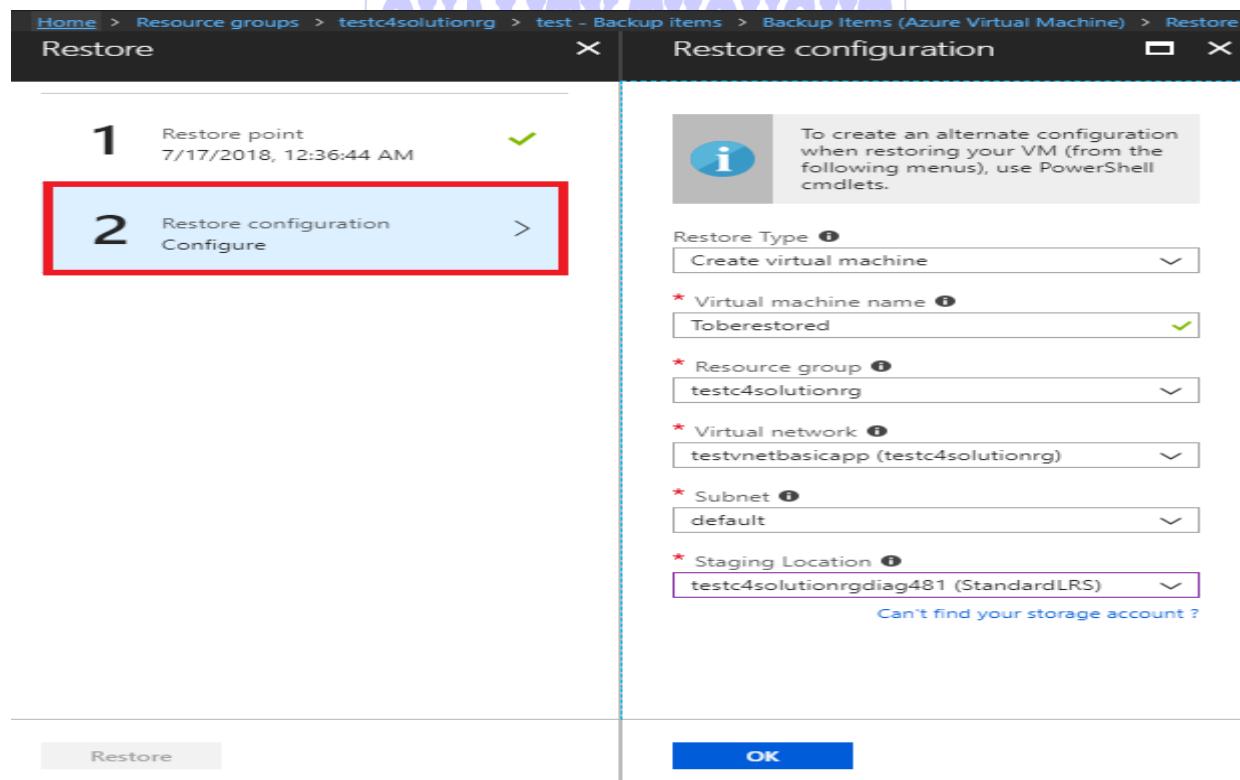
When you go back to check backed-up VMs, you can restore them in just one click (as shown in the following screenshot):



What you need to do is select which **Restore point** you should recover, as shown in the following screenshot:

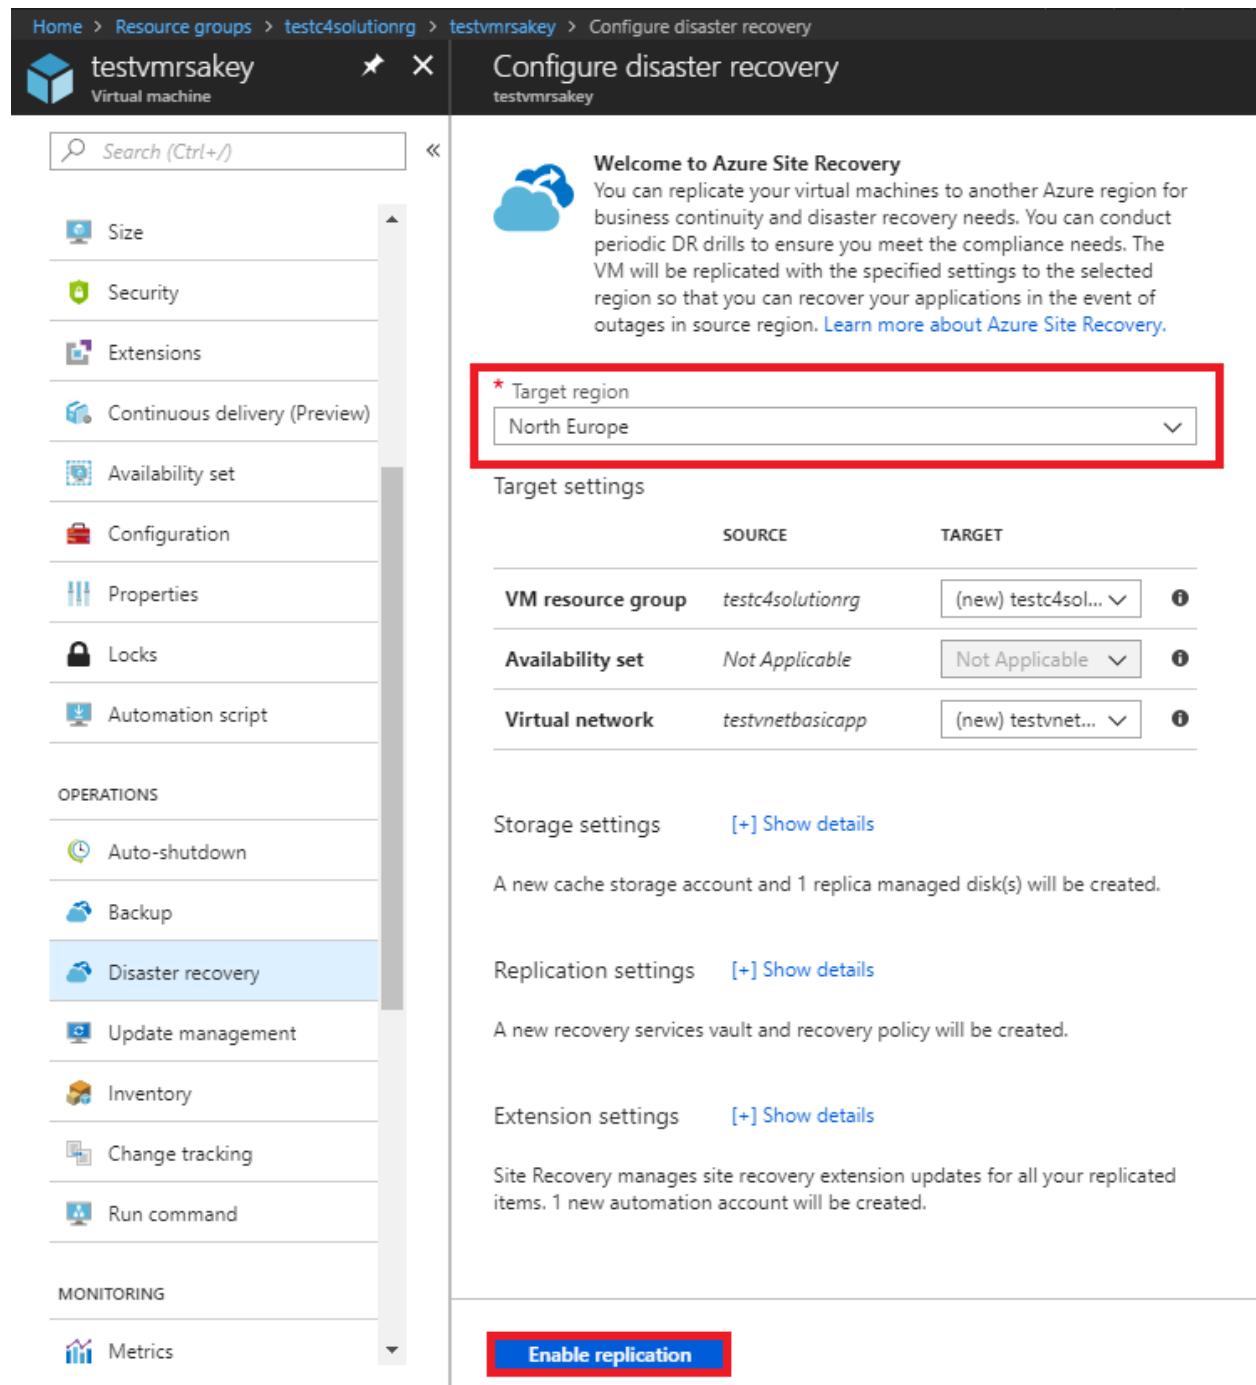


Simply configure where you'll restore your backup. The operation will take a couple of minutes, but it is simple to use since all the manipulations can be done via the Azure portal:



Replicating workloads to another region

Go to the **Disaster recovery** blade of Azure VM and choose a **Target region** to replicate. By default, it is in the paired region of the region where the resource you choose was deployed:



Home > Resource groups > testc4solutionrg > testvmsakey > Configure disaster recovery

Configure disaster recovery
testvmsakey

Welcome to Azure Site Recovery
You can replicate your virtual machines to another Azure region for business continuity and disaster recovery needs. You can conduct periodic DR drills to ensure you meet the compliance needs. The VM will be replicated with the specified settings to the selected region so that you can recover your applications in the event of outages in source region. [Learn more about Azure Site Recovery.](#)

* Target region
North Europe

	SOURCE	TARGET
VM resource group	testc4solutionrg	(new) testc4sol...
Availability set	Not Applicable	Not Applicable
Virtual network	testvnetbasicapp	(new) testvnet...

Storage settings [+ Show details]
A new cache storage account and 1 replica managed disk(s) will be created.

Replication settings [+ Show details]
A new recovery services vault and recovery policy will be created.

Extension settings [+ Show details]
Site Recovery manages site recovery extension updates for all your replicated items. 1 new automation account will be created.

Enable replication

You can create a new recovery service vault and configure a recovery policy. It will display the Azure regions in which you can replicate your workloads, as shown in the following screenshot:

Home > Resource groups > testc4solutionrg > testvmsakey > Configure disaster recovery

Configure disaster recovery

testvmsakey

A new recovery services vault and recovery policy will be created.

Extension settings [\[+\] Show details](#)

Site Recovery manages site recovery extension updates for all your replicated items. 1 new automation account will be created.

Enable replication

- Source region (West Europe)
- Selected target region (North Europe)
- Available target regions

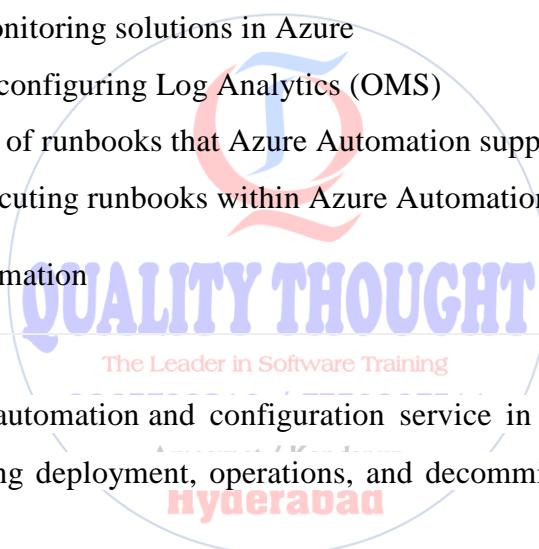
Managing Azure Operations and Automation

Azure Automation manages the life cycle of the infrastructure and applications in Azure by using runbooks, it also allows the use of **Desired State Configurations (DSC)** to configure Windows-based and Linux distribution-based machines at the infrastructure and application level across hybrid environment. It can also work with CI/CD tools, such as Jenkins and **Visual Studio Team Services (VSTS)**.

In this chapter, you will learn the following topics:

- Automation services and their main characteristics
- An overview of monitoring solutions in Azure
- Implementing and configuring Log Analytics (OMS)
- The different types of runbooks that Azure Automation supports
- Publishing and executing runbooks within Azure Automation

Implementing Azure Automation



QUALITY THOUGHT
The Leader in Software Training

Azure Automation is an automation and configuration service in Azure. It provides a wide range of capabilities during deployment, operations, and decommissioning of workloads and resources:

- **Process automation** makes it possible to automate frequent, time-consuming cloud management tasks along with reducing errors and lowering operational costs.
- **Configuration management** allows one to manage DSC resources in Azure Automation and apply configurations to virtual or physical machines across hybrid environments.
- **Update management** allows one to update different operating systems such as Windows and Linux across hybrid environments and provides the visibility of update compliance across Azure, on-premises, and other clouds.
- **Shared capabilities** are a set of shared resources that make it easier to automate and configure environments at scale, such as role-based access control, variables, and credentials.

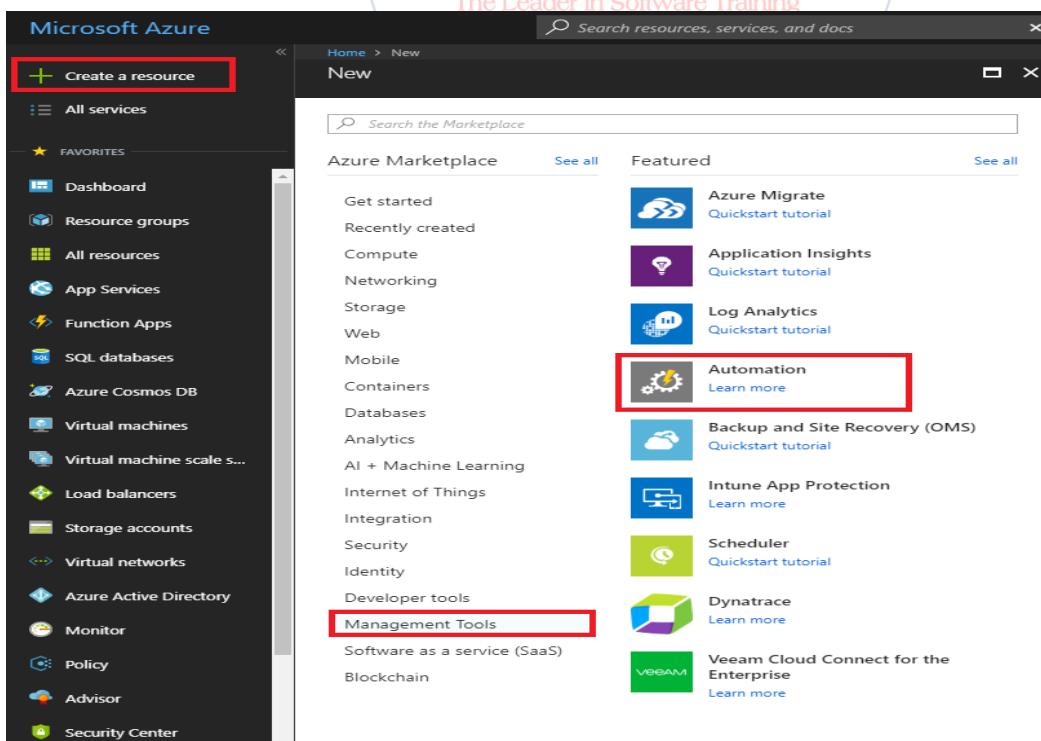
An overview of runbooks

Azure Automation supports a couple of types as described in the following:

- **Graphical runbook** can be created and edited completely in a graphical editor via Azure Portal
- **Graphical PowerShell Workflow runbook** is based on Windows PowerShell Workflow and is created and edited completely in the graphical editor via Azure Portal
- **PowerShell runbook** is based on Windows PowerShell script
- **PowerShell Workflow runbook** is based on Windows PowerShell Workflow
- **Python runbook** is a code snippet in Python

Creating an Automation account

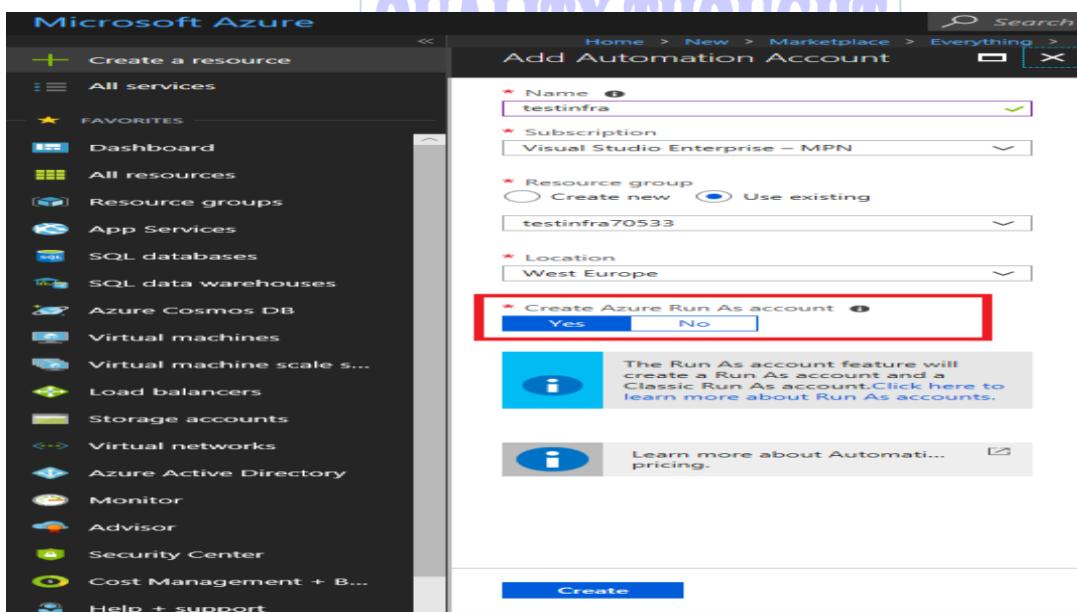
To create an Automation account, you should go to the Azure Portal and click on **Create a resource**, then you can find **Automation** in the management tool category (as shown in the following screenshot):



When you create a new Automation account via the Azure Portal, you can create an Azure Automation account by setting **Run As account** as **Yes**; this feature allows you to create a **Run As account** and a **Classic Run As account** with some useful resources such as sample script showing how to authenticate with Azure Automation or uses certificate which were automatically included for you. Azure automation use a new service principal that allows us to assign the Contributor role-based access control (RBAC) role in the subscription by default. Using this feature, users can authenticate with Azure when you're managing ARM resources using runbooks and made it possible to automate the use of global runbooks that configured in Azure alerts. To resume, there are three tasks while creating a **Run As account**:

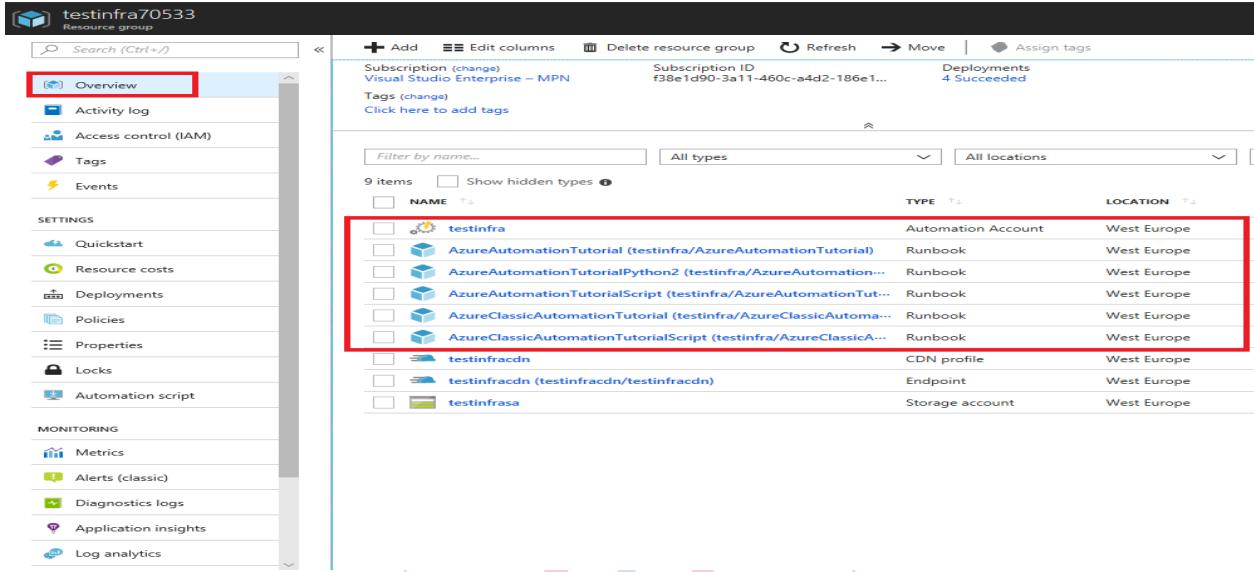
- Creating a new service principal in **Azure Active Directory (Azure AD)**
- Creating a one-year lifespan certificate that authenticates with Azure, so you can manage Azure Resource Manager resources from runbooks
- Assigning the Contributor Role-Based Access Control so that your role can manage Azure Resource Manager resources using runbooks

Look at the following screenshot:



After clicking on **Create**, the deployment will be launched and it will last a couple of minutes. After receiving a notification that the deployment is successful, you can go to the same

resource group and check the created resource; you can see that an Automation account has been created with some default runbooks (as shown here):



NAME	TYPE	LOCATION
testinfra	Automation Account	West Europe
AzureAutomationTutorial (testinfra/AzureAutomationTutorial)	Runbook	West Europe
AzureAutomationTutorialPython2 (testinfra/AzureAutomation...)	Runbook	West Europe
AzureAutomationTutorialScript (testinfra/AzureAutomati...)	Runbook	West Europe
AzureClassicAutomationTutorial (testinfra/AzureClassi...)	Runbook	West Europe
AzureClassicAutomationTutorialScript (testinfra/AzureClassicA...)	Runbook	West Europe
testinfracdn	CDN profile	West Europe
testinfracdn (testinfracdn/testinfracdn)	Endpoint	West Europe
testinfrasa	Storage account	West Europe

To find out how to create a Automation account using PowerShell, check the following link: <https://docs.microsoft.com/en-us/azure/automation/automation-create-runas-account>.

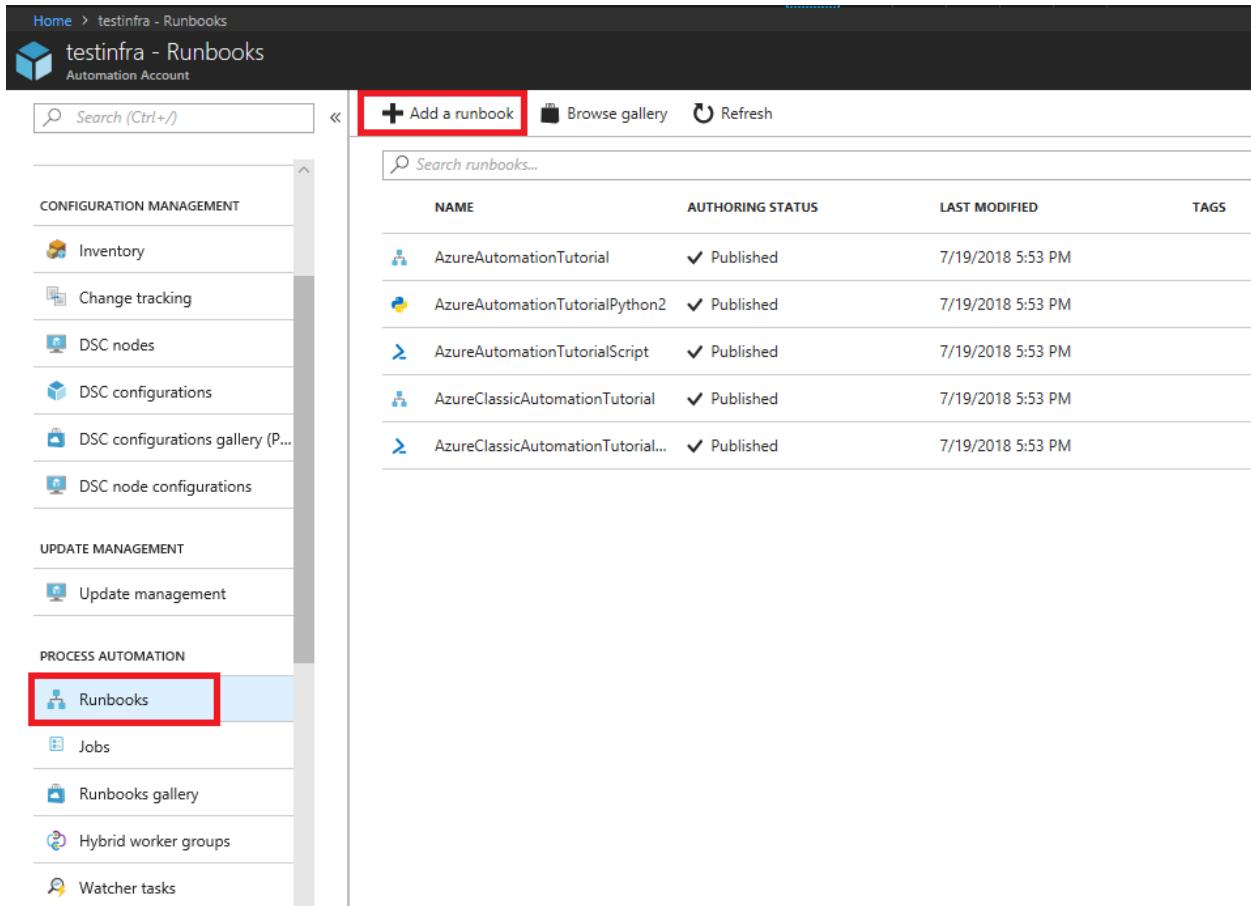
You can also create an Automation account using an ARM template. The example of an ARM template is as follows:

9963799240 / 7730997544

```
{
  "name": "string",
  "type": "Microsoft.Automation/automationAccounts",
  "apiVersion": "2018-01-15",
  "properties": {
    "sku": {
      "name": "string",
      "family": "string",
      "capacity": "integer"
    }
  },
  "location": "string",
  "tags": {}
}
```

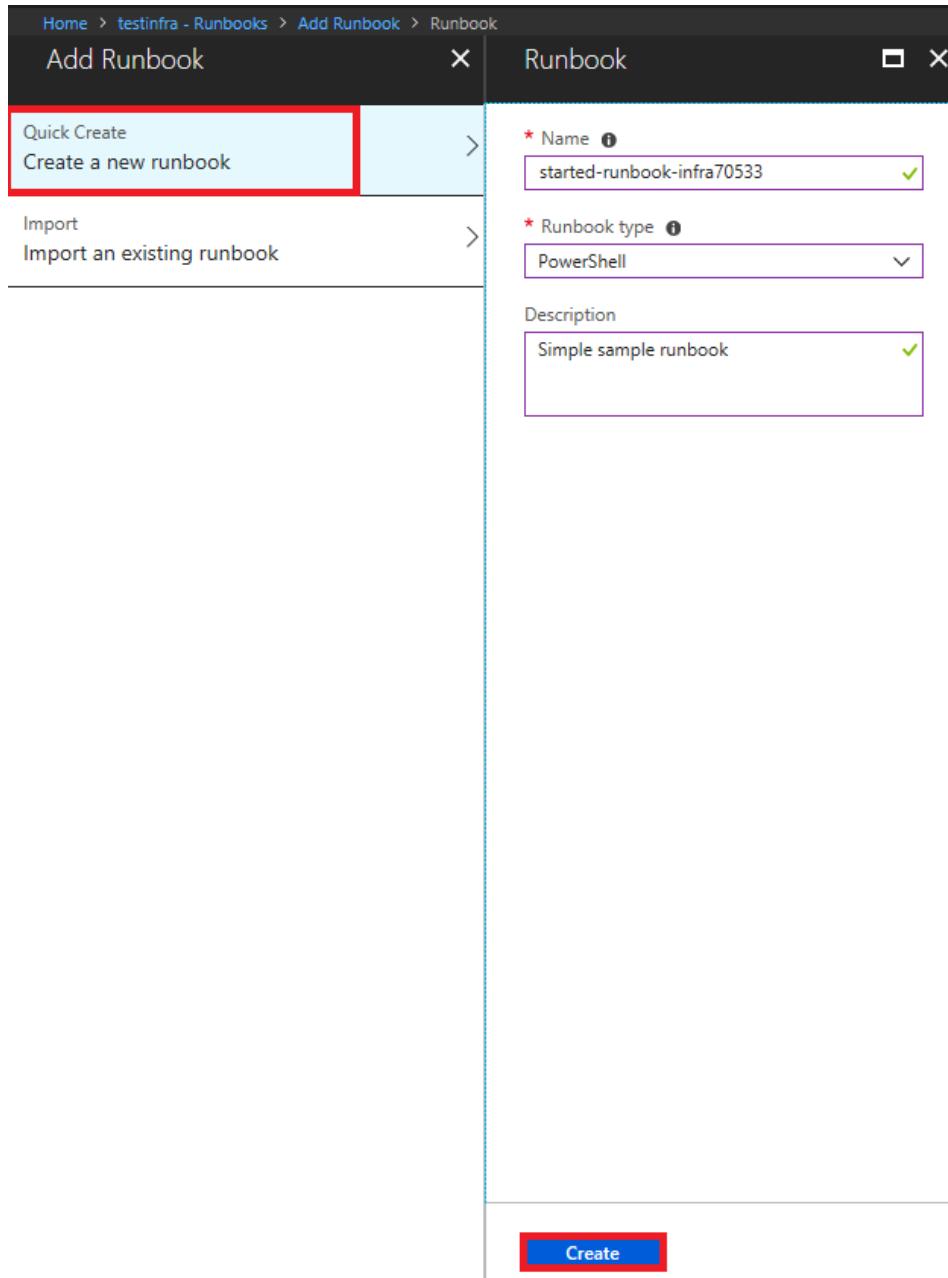
Creating or importing PowerShell runbooks

To create or import a new runbook, you can go to the **Runbooks** blade of **Automation Account** and click on **Add a runbook**, shown as follows:



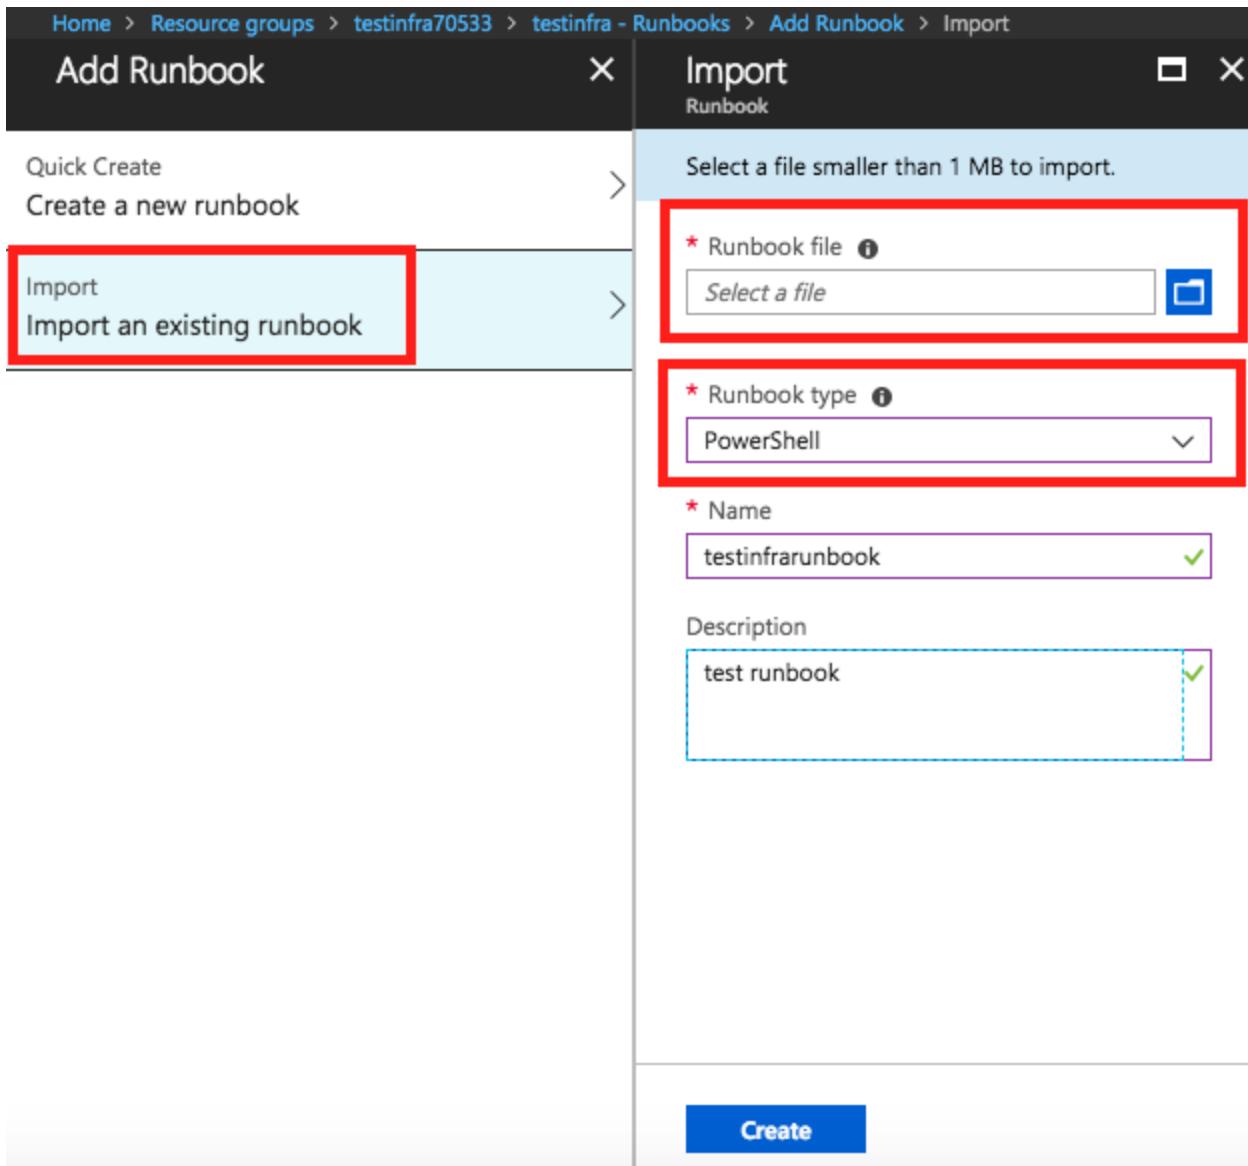
NAME	AUTHORING STATUS	LAST MODIFIED	TAGS
AzureAutomationTutorial	✓ Published	7/19/2018 5:53 PM	
AzureAutomationTutorialPython2	✓ Published	7/19/2018 5:53 PM	
AzureAutomationTutorialScript	✓ Published	7/19/2018 5:53 PM	
AzureClassicAutomationTutorial	✓ Published	7/19/2018 5:53 PM	
AzureClassicAutomationTutorial...	✓ Published	7/19/2018 5:53 PM	

You can fill in the basic information and choose an appropriate **Runbook type**, as shown in the next screenshot:



The screenshot shows the Microsoft Azure portal interface for creating a new runbook. On the left, there's a sidebar with navigation links: Home > testinfra - Runbooks > Add Runbook > Runbook. Below this, there are two main options: 'Quick Create' (highlighted with a red box) and 'Create a new runbook', followed by 'Import' and 'Import an existing runbook'. On the right, the 'Runbook' configuration page is displayed. It includes fields for 'Name' (started-runbook-infra70533), 'Runbook type' (PowerShell), and 'Description' (Simple sample runbook). At the bottom of the configuration page, there is a prominent blue 'Create' button.

If you want to import a runbook, you should upload your runbook and choose the right type of runbook, shown as follows:



The screenshot shows the Azure portal interface for creating a new runbook. On the left, a sidebar lists options: 'Quick Create', 'Create a new runbook', and 'Import'. The 'Import' option is selected and highlighted with a red box. On the right, the 'Import' blade is open, prompting the user to 'Select a file smaller than 1 MB to import.' It contains several input fields: 'Runbook file' (with a 'Select a file' button), 'Runbook type' (set to 'PowerShell'), 'Name' (set to 'testinfrarunbook'), and 'Description' (set to 'test runbook'). A large blue 'Create' button is at the bottom.

After clicking on **Create**, a new runbook will be deployed, but the runbook won't be operational instantly. You can go to the **Runbooks** blade of the Automation account, and you can see (as shown in the following screenshot), the authoring status is marked as **New**, which means that we should publish our runbooks to make it operational:

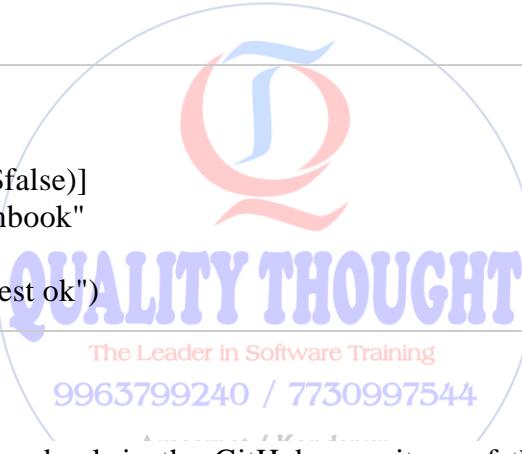
+ Add a runbook Browse gallery Refresh

Search runbooks...

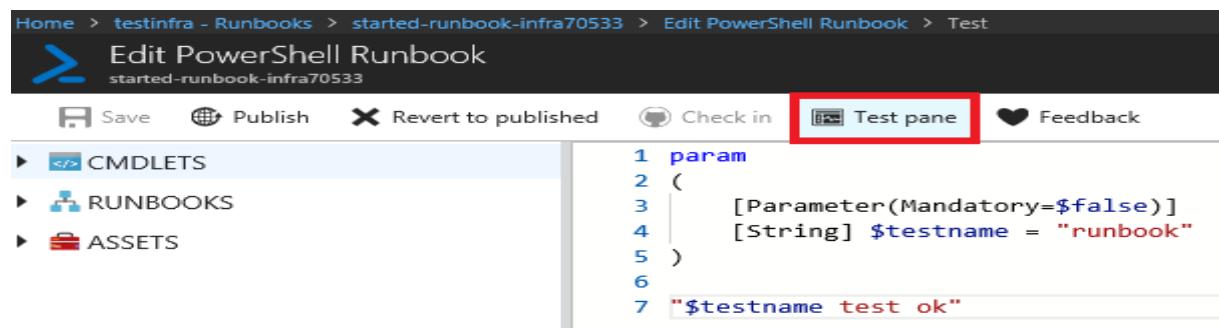
NAME	AUTHORING STATUS	LAST MODIFIED	TAGS
AzureAutomationTutorial	✓ Published	7/19/2018 5:53 PM	
AzureAutomationTutorialPython2	✓ Published	7/19/2018 5:53 PM	
AzureAutomationTutorialScript	✓ Published	7/19/2018 5:53 PM	
AzureClassicAutomationTutorial	✓ Published	7/19/2018 5:53 PM	
AzureClassicAutomationTutorial...	✓ Published	7/19/2018 5:53 PM	
started-runbook-infra70533	New	7/20/2018 10:10 AM	

Here, we're going to test a simple sample runbook which was created previously to help you understand how to publish and test a newly created runbook. The sample runbook is as follows:

```
param
(
    [Parameter(Mandatory=$false)]
    [String] $testname = "runbook"
)
Write-Output ("$testname test ok")
```



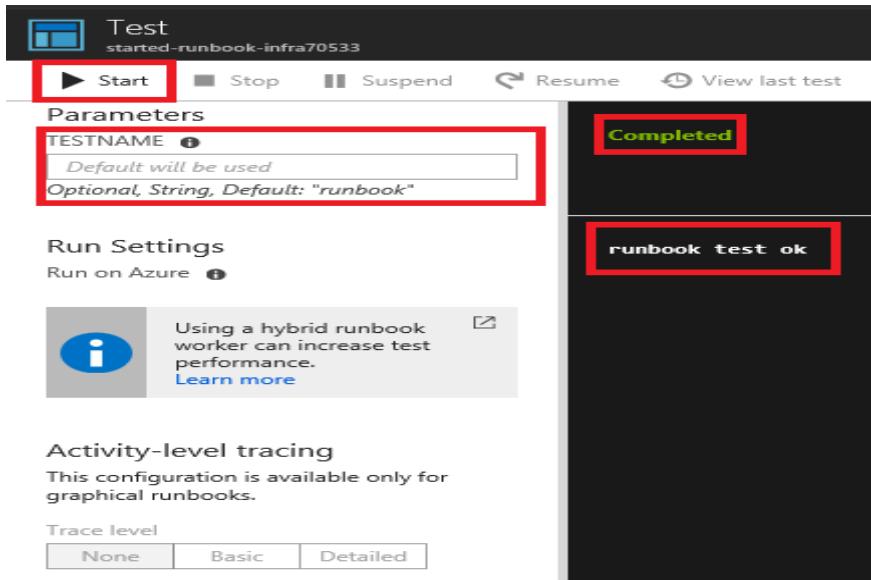
You can find this sample runbook in the GitHub repository of this book that was given in the **Technical requirements** section of this chapter. After completely editing the content of this runbook, you can click on **Test pane** to test it (as shown in the following screenshot):



The screenshot shows the Azure portal interface for editing a PowerShell runbook. The top navigation bar includes 'Home', 'testinfra - Runbooks', 'started-runbook-infra70533', 'Edit PowerShell Runbook', and 'Test'. Below the navigation is a toolbar with 'Save', 'Publish', 'Revert to published', 'Check in', 'Test pane' (which is highlighted with a red box), and 'Feedback'. On the left, a sidebar lists 'CMDLETS', 'RUNBOOKS', and 'ASSETS'. The main content area displays the PowerShell script:

```
1 param
2 (
3     [Parameter(Mandatory=$false)]
4     [String] $testname = "runbook"
5 )
6
7 "$testname test ok"
```

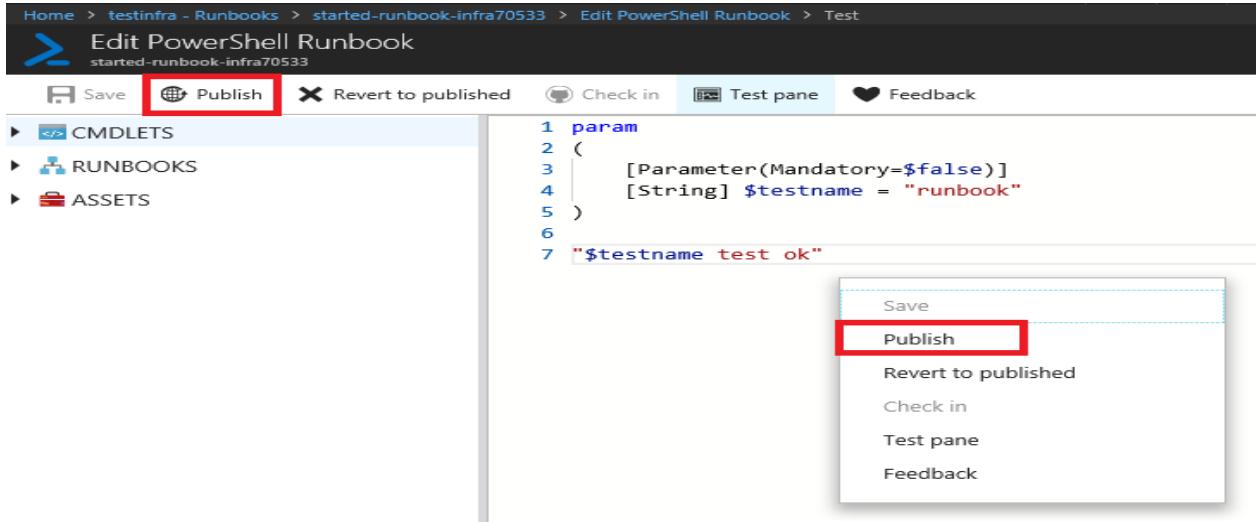
We can see that we can input the parameters in the test panel before clicking on **Start** to launch a new test as shown in the following screenshot. It is possible to set run settings to run the runbook on Azure or Hybrid environment, as shown in the next screenshot:



If you make sure that everything is OK, you can publish your runbook so that it is operational. You can publish a runbook in the following two ways:

- Click on the **Publish** button
- Right-click and choose **Publish** in the menu

Look at the following screenshot:



Home > testinfra - Runbooks > started-runbook-infra70533 > Edit PowerShell Runbook > Test

Edit PowerShell Runbook
started-runbook-infra70533

Save Publish Revert to published Check in Test pane Feedback

CMDLETS RUNBOOKS ASSETS

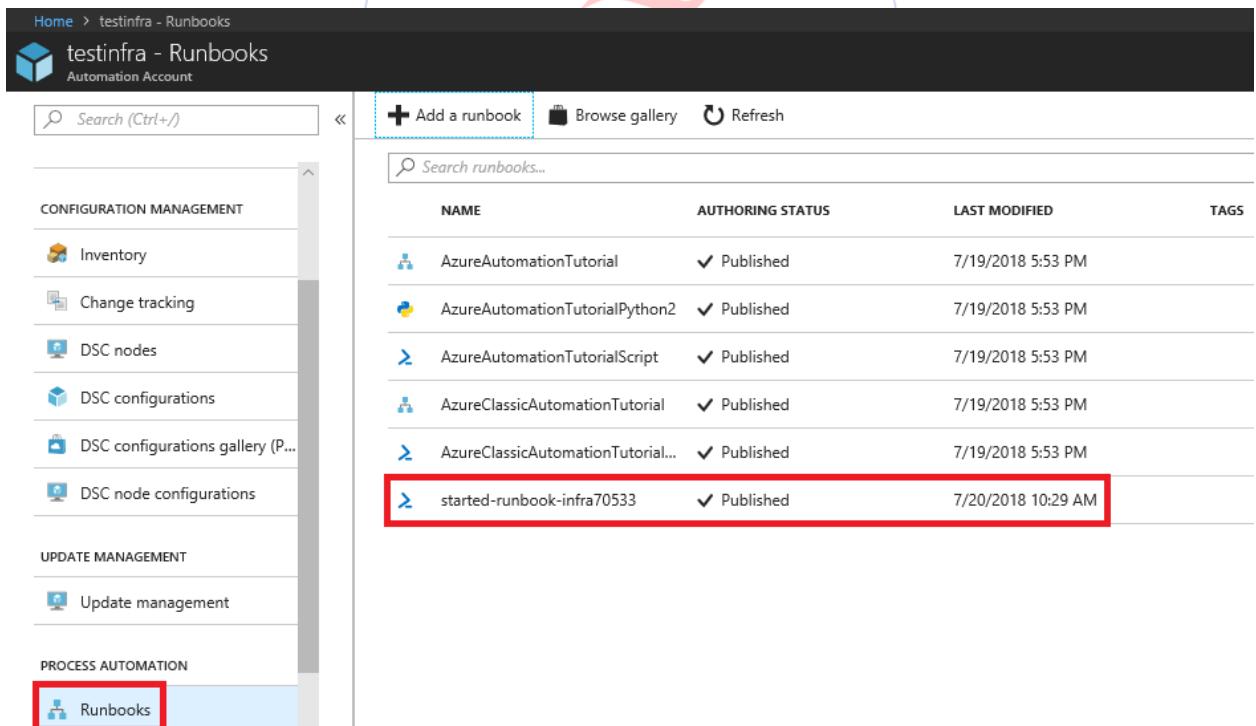
```

1 param
2 (
3     [Parameter(Mandatory=$false)]
4     [String] $testname = "runbook"
5 )
6
7 "$testname test ok"

```

Save Publish Revert to published Check in Test pane Feedback

To make sure that you have published your runbook successfully, you can go back to the **Runbooks** blade of your Automation account and check the current status of the runbook. As shown in the following screenshot, the status **Published** means this runbook has been published successfully:



Home > testinfra - Runbooks

testinfra - Runbooks

Automation Account

Search (Ctrl+ /)

Add a runbook Browse gallery Refresh

NAME	AUTHORING STATUS	LAST MODIFIED	TAGS
AzureAutomationTutorial	✓ Published	7/19/2018 5:53 PM	
AzureAutomationTutorialPython2	✓ Published	7/19/2018 5:53 PM	
AzureAutomationTutorialScript	✓ Published	7/19/2018 5:53 PM	
AzureClassicAutomationTutorial	✓ Published	7/19/2018 5:53 PM	
AzureClassicAutomationTutorial...	✓ Published	7/19/2018 5:53 PM	
started-runbook-infra70533	✓ Published	7/20/2018 10:29 AM	

CONFIGURATION MANAGEMENT

- Inventory
- Change tracking
- DSC nodes
- DSC configurations
- DSC configurations gallery (P...)
- DSC node configurations

UPDATE MANAGEMENT

- Update management

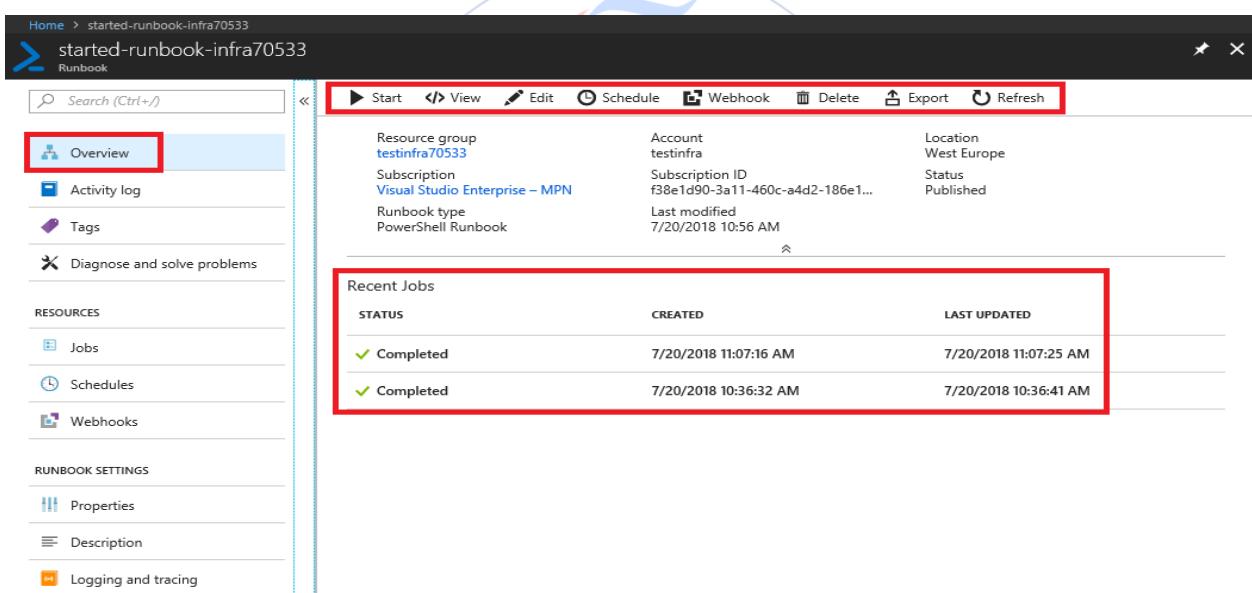
PROCESS AUTOMATION

- Runbooks

You can check the following link to know more about how to author the graphic runback in Azure Automation: <https://docs.microsoft.com/en-us/azure/automation/automation-graphical-authoring-intro>.

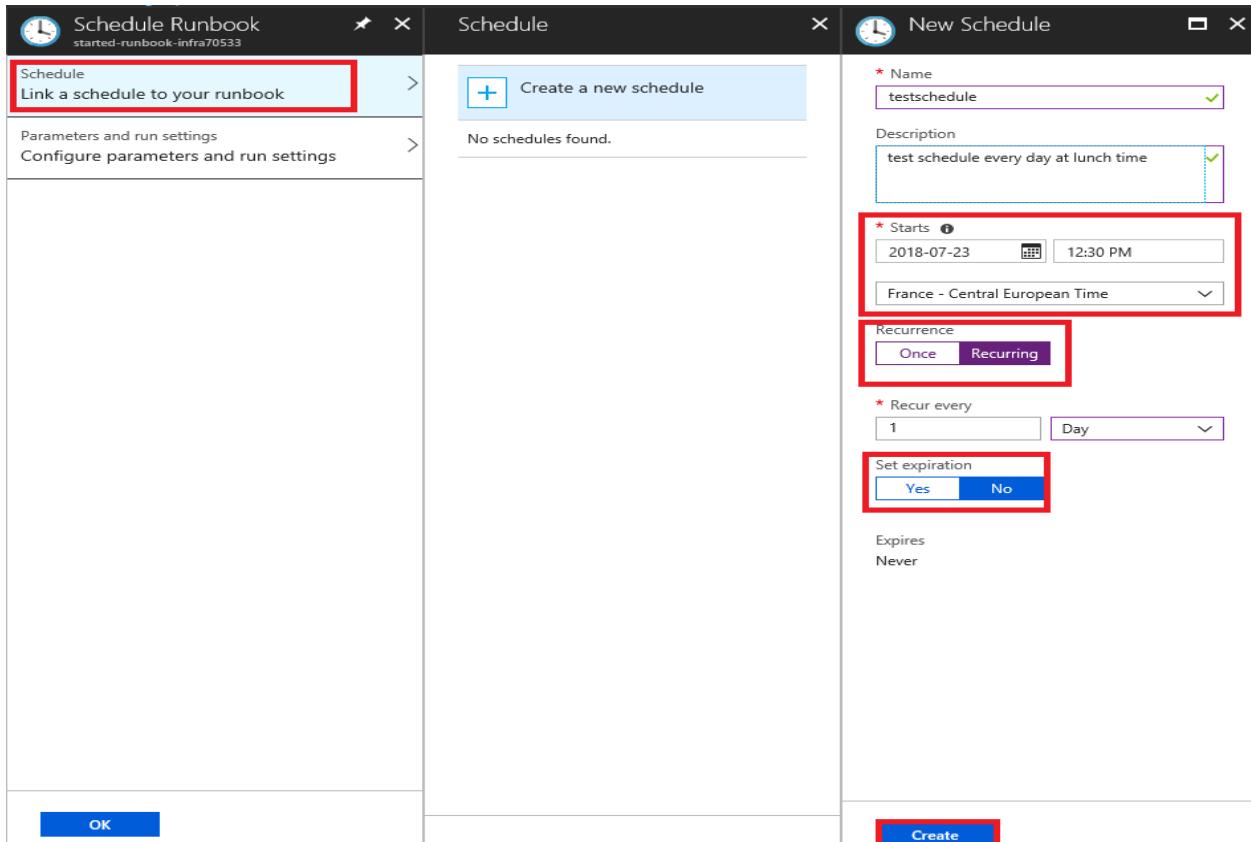
Managing PowerShell runbooks

To manage a PowerShell runbook, you may go to the **Overview** blade of the runbook, and you can use the toolbar (as shown in the following screenshot) to manage your runbook. Even after publishing, you can edit (by clicking on **Edit**) an existing runbook and test it by clicking on **Start**. In the **Recent Jobs** section, you can find the historical record of the execution of the current runbook:

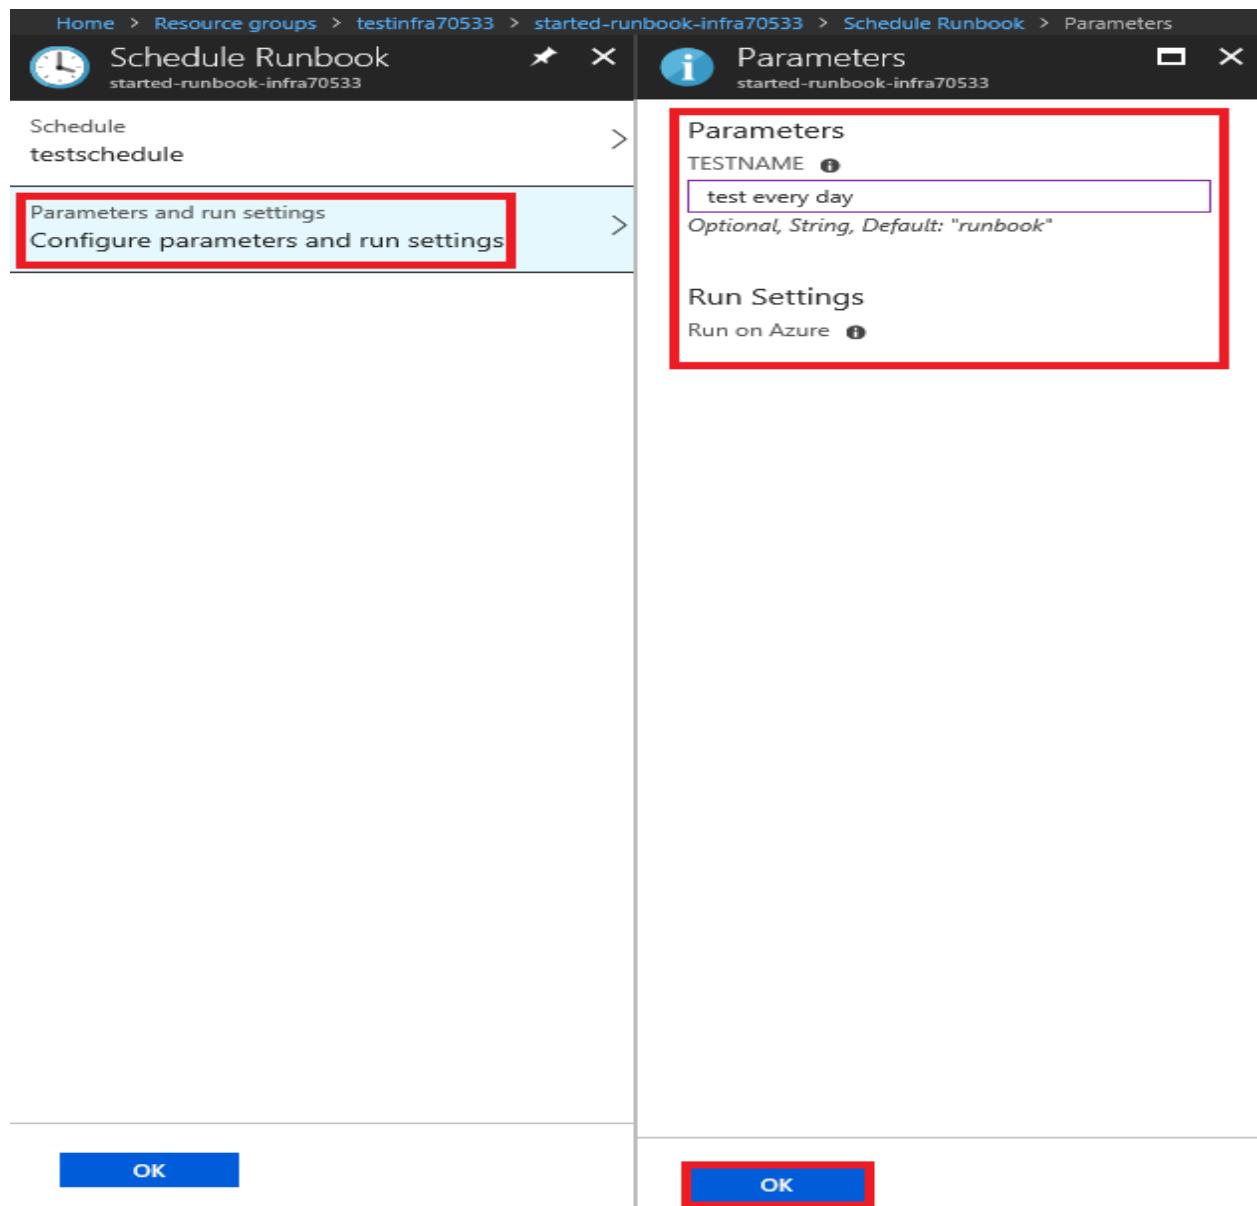


STATUS	CREATED	LAST UPDATED
✓ Completed	7/20/2018 11:07:16 AM	7/20/2018 11:07:25 AM
✓ Completed	7/20/2018 10:36:32 AM	7/20/2018 10:36:41 AM

You can also configure a **Schedule** or **webhook** to trigger your runbook. To create a schedule, you can click on **Schedule** or go to the **Schedule** blade to click on **Create a new schedule**, and you can start to create a new schedule or use an existing schedule in the list (as follows):



It is possible to schedule a runbook by defining it is a job to run only once or repeatedly. If you're choosing it as a repeat job, you can choose the frequency of repeat and an expiration date if you choose recurring recurrence. Then, you can also specify the input parameter when needed and click on **OK**, as follows:



Schedule Runbook
started-runbook-infra70533

Schedule
testschedule

Parameters and run settings
Configure parameters and run settings

Parameters

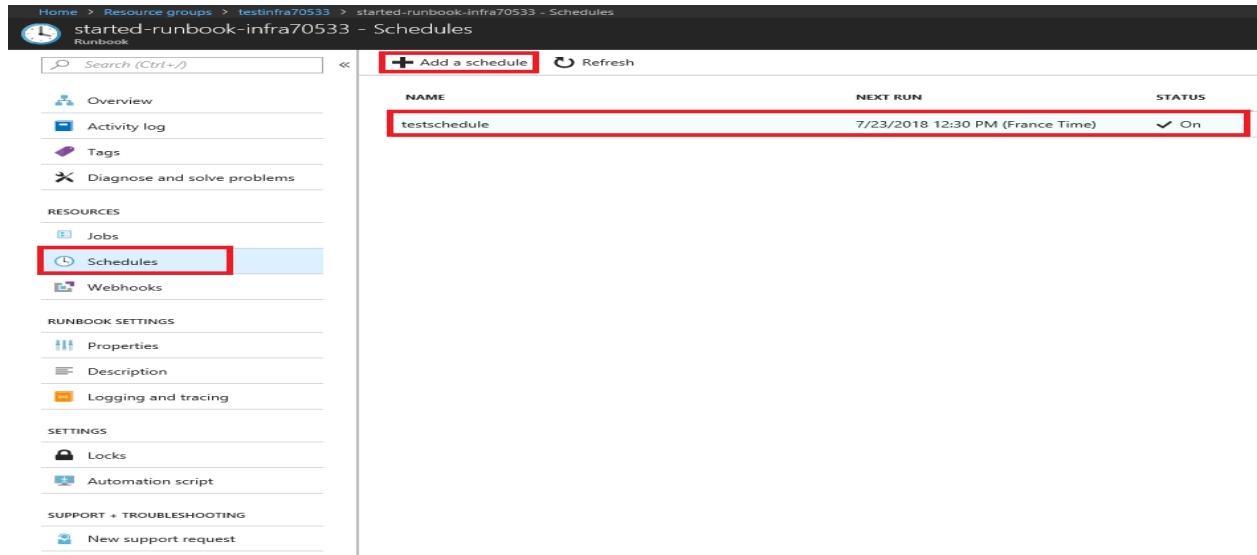
TESTNAME ⓘ
test every day
Optional, String, Default: "runbook"

Run Settings
Run on Azure ⓘ

OK

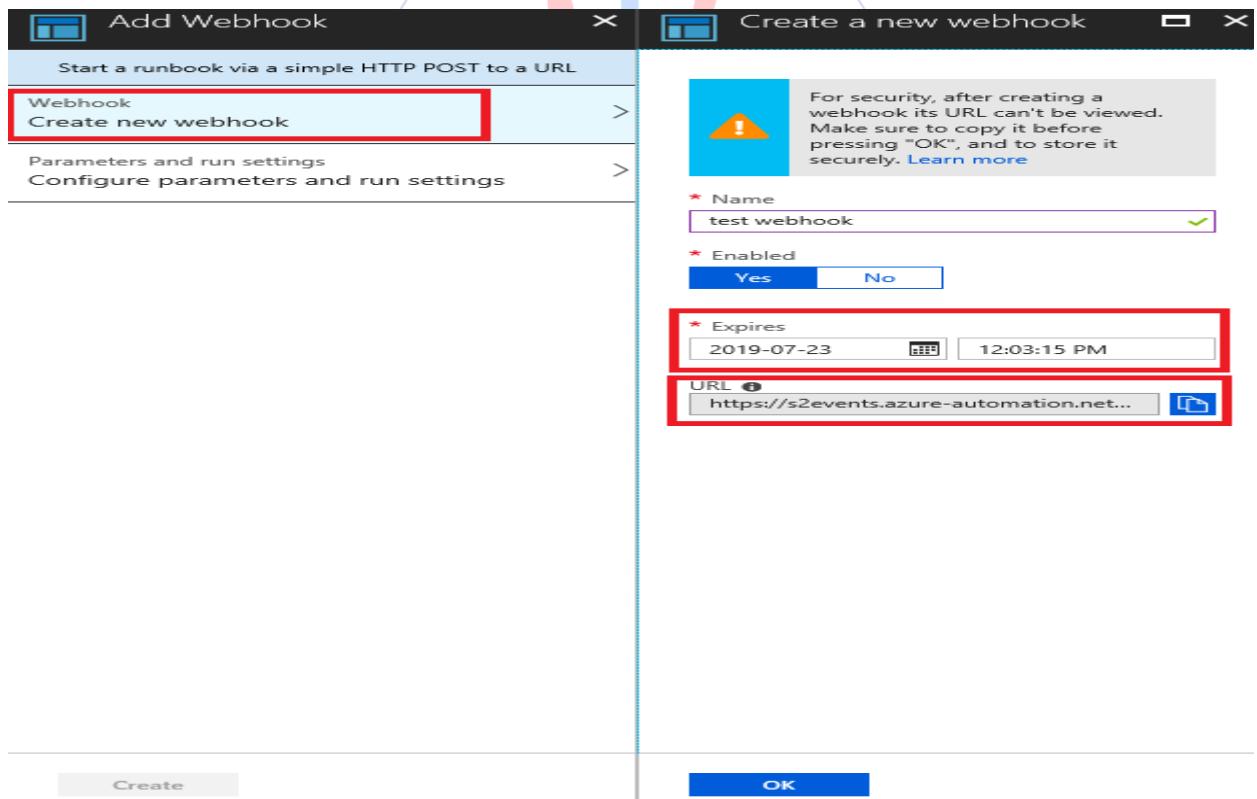
OK

You can check the **Schedules** blade to know whether your schedule has been created successfully and what the current status of the schedule is, as follows:



A screenshot of the Azure portal showing the 'Schedules' section for a runbook. The left sidebar shows various sections like Overview, Activity log, Tags, Diagnose and solve problems, Jobs, and Schedules (which is highlighted with a red box). The main area shows a table with columns NAME, NEXT RUN, and STATUS. One row is selected, showing 'testschedule' with the next run at '7/23/2018 12:30 PM (France Time)' and status 'On'.

Similarly, to create a **Webhook**, you should start by clicking on **Webhook** and set a status for **Webhook** as well as the expiration date, as follows:

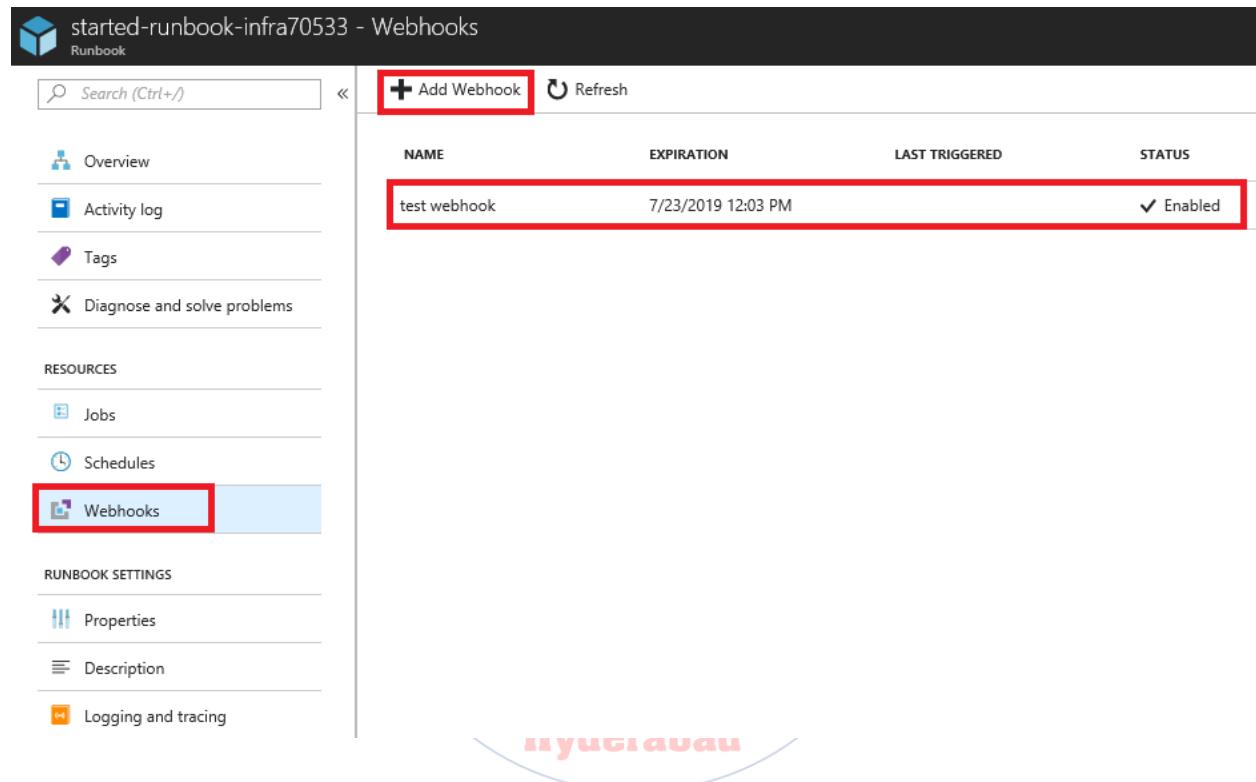


A screenshot of the Azure portal showing the 'Add Webhook' and 'Create a new webhook' dialog boxes. The 'Add Webhook' dialog has a 'Create new webhook' button highlighted with a red box. The 'Create a new webhook' dialog shows fields for Name (test webhook), Enabled (Yes), Expires (2019-07-23 12:03:15 PM), and URL (https://s2events.azure-automation.net...). The 'Create' button is in the bottom-left of the left box, and the 'OK' button is in the bottom-right of the right box.

Here is a sample URL generated by Azure while creating a Webhook via the Azure Portal:

<https://s2events.azure-automation.net/webhooks?token=buL5eq4XMJCP4c%2bVQqqbU%2fG3jnb9NjCC0RpRiHC1zgs%3d>

You can check the **Webhooks** blade to know if your Webhook has been configured successfully (as shown here):



NAME	EXPIRATION	LAST TRIGGERED	STATUS
test webhook	7/23/2019 12:03 PM		✓ Enabled

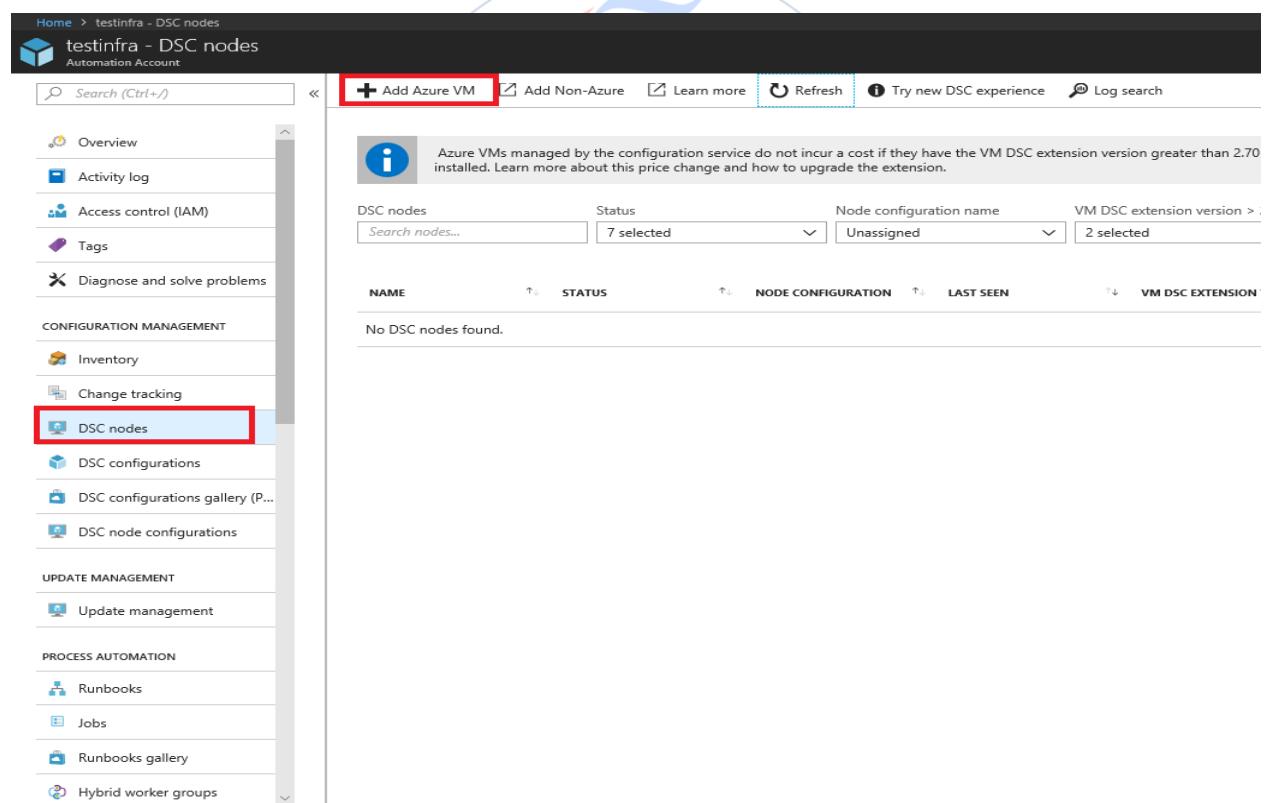
An overview of Desired State Configuration

Azure Automation **Desired State Configuration (DSC)** is a configuration management service in Azure. A configuration is a document that describes an environment with specific characteristics. Azure Automation focuses on **configurations**; it can configure physical and virtual machines on premises, or in any cloud provider other than Azure, it considers a single machine as a DSC node, and makes it possible to scale across thousands of based on a declarative configuration.

Azure Automation DSC can be used to manage a wide range of machines on premises or in the cloud such as Azure VMs (both classic and ARM-based), AWS instances, or any other cloud providers. Both physical and virtual machine, across different operating systems such as Windows or Linux.

Implementing PowerShell Desired State Configurations

To implement the desired state configuration, you can go to the **DSC nodes** blade of your Automation account. This blade is where you virtualize all the DSC nodes in the same configuration. You can add the Azure VM or non-Azure resource here; to add a VM in Azure, you can click on **Add Azure VM**, as shown in the following screenshot:



The screenshot shows the Azure portal interface for an Automation Account named 'testinfra - DSC nodes'. The left sidebar navigation bar includes links for Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Configuration Management (Inventory, Change tracking, DSC nodes, DSC configurations, DSC configurations gallery, DSC node configurations), Update Management (Update management), Process Automation (Runbooks, Jobs, Runbooks gallery, Hybrid worker groups), and a Home link. The 'DSC nodes' link is highlighted with a red box. The main content area displays the 'DSC nodes' blade. At the top, there is a header with a search bar, a '+ Add Azure VM' button (which is highlighted with a red box and has a blue arrow pointing to it from the text), an 'Add Non-Azure' button, a 'Learn more' link, a 'Refresh' button, a 'Try new DSC experience' link, and a 'Log search' link. Below the header, there is a message: 'Azure VMs managed by the configuration service do not incur a cost if they have the VM DSC extension version greater than 2.70 installed. Learn more about this price change and how to upgrade the extension.' A table titled 'DSC nodes' lists nodes with columns for NAME, STATUS, NODE CONFIGURATION, LAST SEEN, and VM DSC EXTENSION. The status is set to '7 selected', node configuration is 'Unassigned', and VM DSC extension version is '2 selected'. The table shows 'No DSC nodes found.'

Then, choose your target machine in the list, as shown here :

Home > testinfra - DSC nodes > Virtual Machines

Virtual Machines testinfra

Refresh Filter by name... Visual Studio Dev Essentials 4 selected West Euro...

VIRTUAL MACHINES	SUBSCRIPTION	RESOURCE GROUP	LOCATION
testbigdataVM	Visual Studio Dev Essentials	testBigDataOnAzureRG	West Europe

Click on **Connect** if there is already an Azure Desired State Configuration (DSC) VM extension that has been deployed in the target VM, and you'll be able to connect it:

testbigdataVM Virtual machine

+ Connect Refresh Learn more

i Not connected

POWER STATE
VM running

OS
Windows

STATUS
Not connected

Managing PowerShell Desired State Configurations

9963799240 / 7730997544

To import Automation DSC Configurations, you can use the Import-AzureRmAutomationDscConfiguration command to import a Desired State Configuration into Azure Automation. The example command is as follows (replace the words between # with your own values):

```
Import-AzureRmAutomationDscConfiguration -AutomationAccountName #yourautomationaccountname# -ResourceGroupName #yourresourcegroupname# -SourcePath #DSCscriptrepo# -Force
```

Users need to generate DSC node configurations using the Import-AzureRmAutomationDscConfiguration command to import an MOF configuration file into Azure Automation as a DSC node configuration. The following is a sample command:

Other excellent configuration management tools

In our world, there are a couple of tools used for configuration management, such as Chef, Puppet, and Ansible. When we talk about configuration management tools, it means a management tool that is designed to deploy, configure, and manage servers. Let's take a quick look at each of them:

- **Chef** is a popular configuration tool in the Linux world. It has master-agent architecture, meaning the server runs on the master machine and, with a client, runs as a client agent on each client machine. Besides master and agent, there is a controller machine that contains all the configurations. Users can use that machine to push all the configurations to central Chef server. Chef is a great choice for users who know Git and Ruby well.
- **Puppet** is also a famous term in the Linux world and has a master-agent architecture. Similar to Chef, Puppet server runs on the master machine, and Puppet clients run as an agent on each client machine. In addition, there is a certificate signing between the agent and the master.
9963799240 / 7730997544
- **Ansible** is an open source automation tool that allows to bootstrap machine and generates environment, different to master-agent architecture. It uses SSH connection to log in to client systems or the nodes you want to configure. It is an awesome tool in the configuration world, and as it is based on Python; it is easy to use for Python developers to personalize some features. The greatest advantage of Ansible is that it is SSH-based; hence, it doesn't require installing any agents on remote nodes.

Implementing Azure Automation-based cloud management

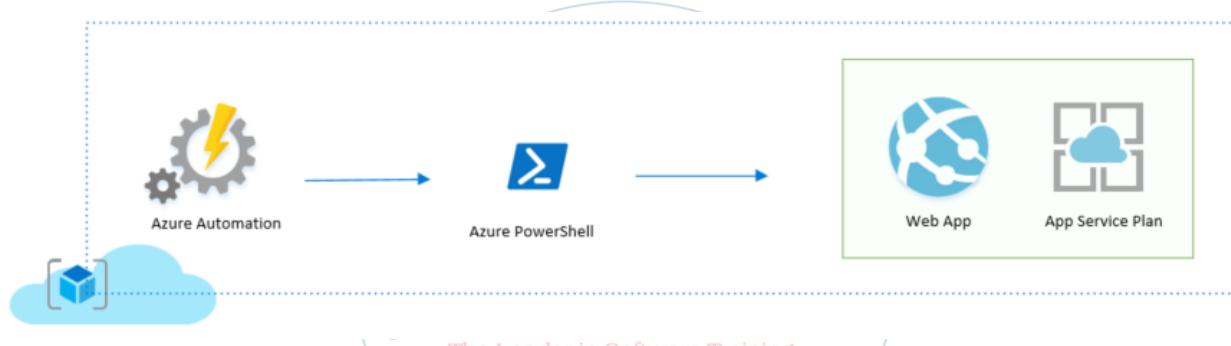
Azure Automation is not only a cloud-based automate service. It can also be used for simplifying cloud management through process automation. We'll introduce some scenarios to show how Azure automation interacts with other Azure services.

Integrating Azure Automation with Web Apps

We know that Web App is a PaaS offering provided in Azure. As Azure PowerShell provides a wide range of commands to manage Web App in an App Service plan, it is also possible to use Azure Automation to interact with Azure Web App. The general scenario is:

- To stop or start Web App automatically
- Scale the App Service plan from a single instance to multiple instances
- Scale Web App in multiple regions with a high-availability architecture
- To create a one-time or scheduled backup of Web App

I am using the following diagram to show these scenarios:



You can find the most useful PowerShell sample with the following link: <https://docs.microsoft.com/en-us/azure/app-service/app-service-powershell-samples>.

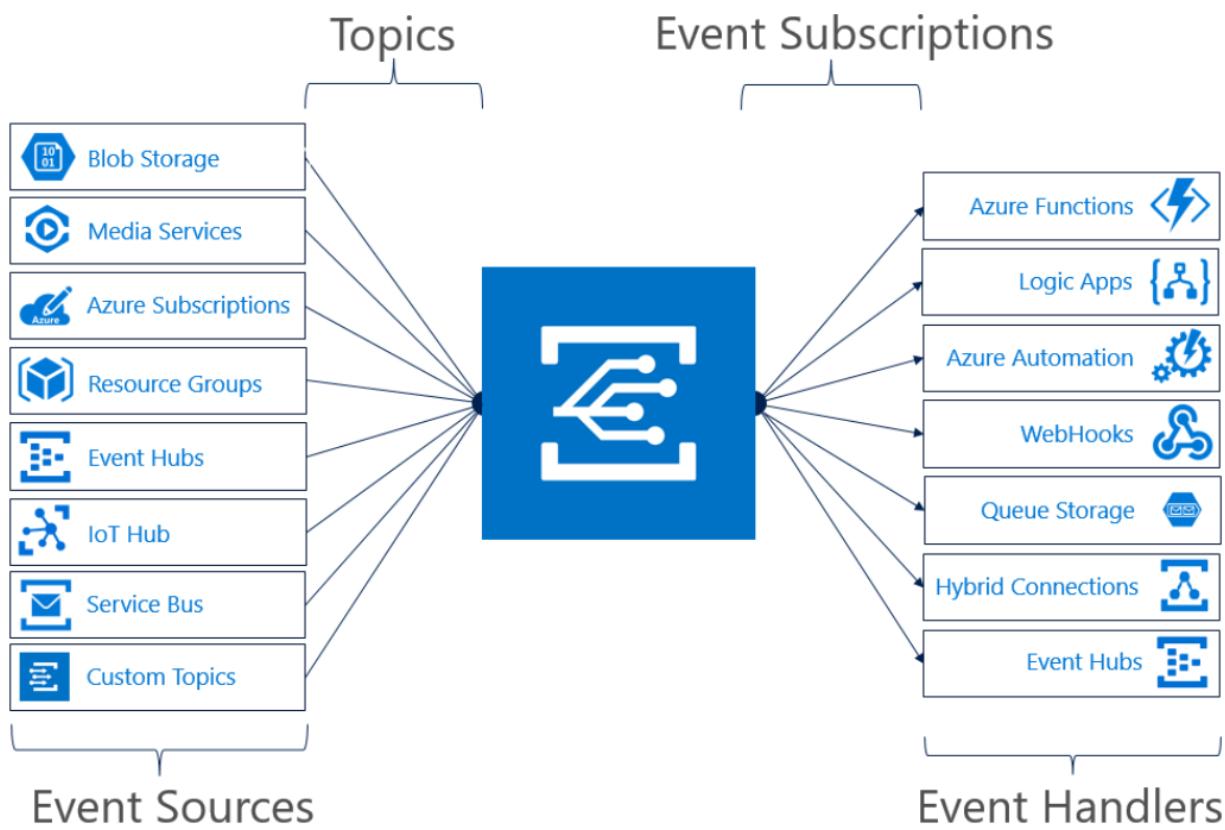
Comparing Azure Automation with Azure Functions

Both Azure Automation and Azure Functions can run PowerShell Scripts on Azure. They both support Webhooks and can be scheduled. However, Azure Functions are far richer in terms of triggers and liaisons, and they support a wide range of languages to write code running in the cloud not limited to PowerShell. As a Function as a Service (FaaS), Azure Functions can be used to build microservices at scale.

Here is the link to find out more about the triggers and bindings of Azure functions: <https://docs.microsoft.com/en-us/azure/azure-functions/functions-triggers-bindings>.

Integrating with Event Grid

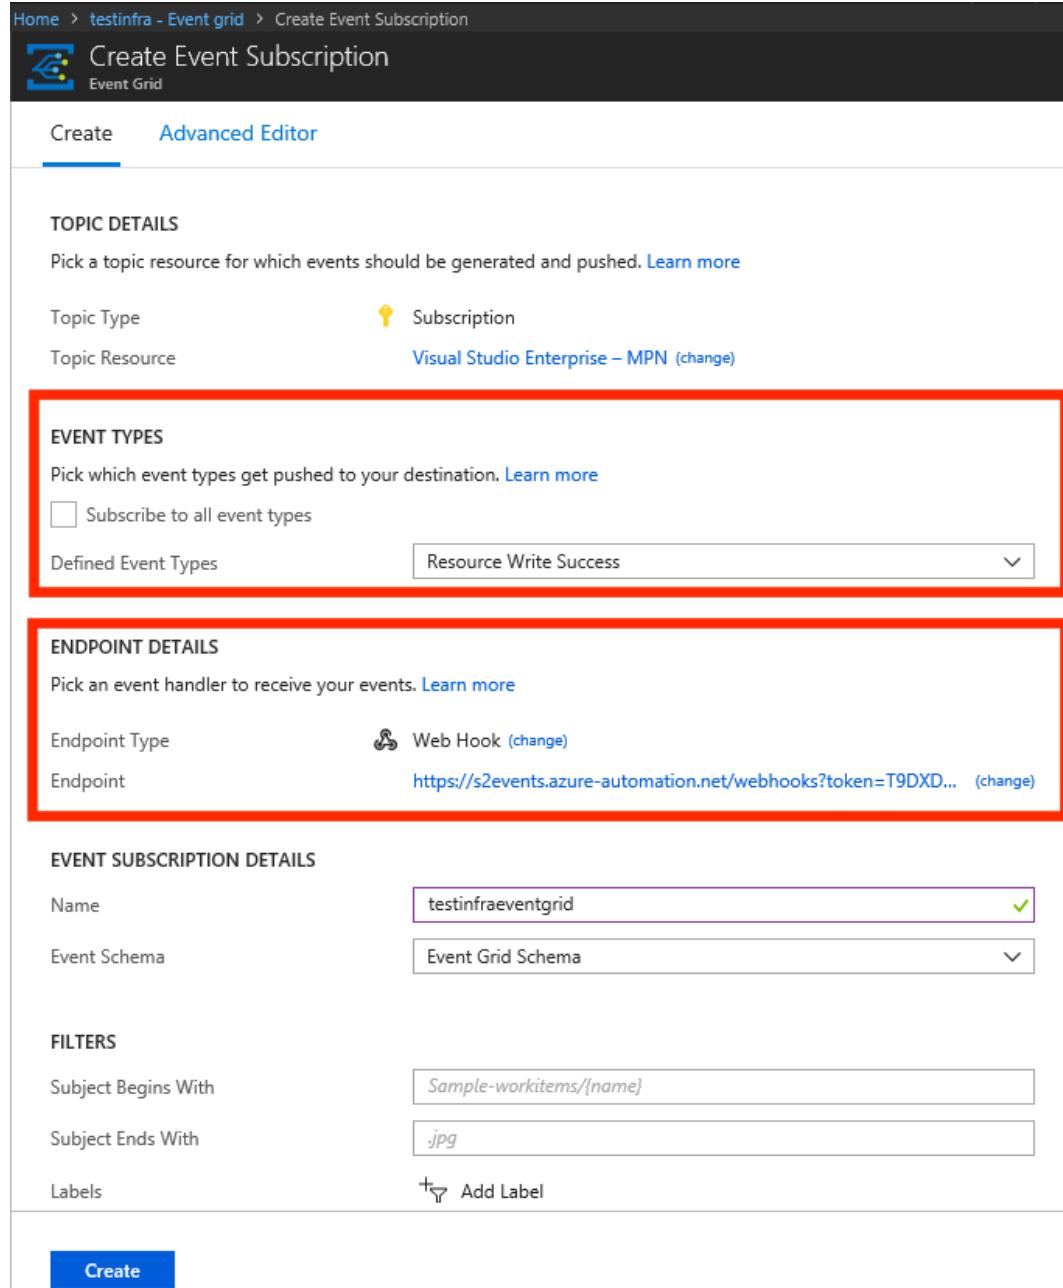
Azure Event Grid is a cloud-based publisher and subscriber service that helps users to build event-based architecture for their application. Users can select Azure resource that is subscribed and give the event handler or **WebHook** endpoint to send the event to. Azure Event Grid is focusing on providing the capabilities of Automation and integration with other Azure Services to build powerful cloud-based solutions, as shown in the following diagram:



Event Grid (<https://github.com/MicrosoftDocs/azure-docs/blob/master/articles/event-grid/overview.md>)

Event grids can be integrated with a wide range of other Azure Services, thanks to its pub-sub pattern. The first thing you should do is create an **Event Subscription**. This step is quite important because you can define the event type that means the kind of events as source, also,

the endpoint type that is the **trigger** for this event. You can see how to create an Event Grid subscription as follows:



Home > testinfra - Event grid > Create Event Subscription

Create Event Subscription

Event Grid

[Create](#) [Advanced Editor](#)

TOPIC DETAILS

Pick a topic resource for which events should be generated and pushed. [Learn more](#)

Topic Type	 Subscription
Topic Resource	Visual Studio Enterprise – MPN (change)

EVENT TYPES

Pick which event types get pushed to your destination. [Learn more](#)

<input type="checkbox"/> Subscribe to all event types	Defined Event Types	Resource Write Success
---	-------------------------------------	--

ENDPOINT DETAILS

Pick an event handler to receive your events. [Learn more](#)

Endpoint Type	 Web Hook (change)
Endpoint	https://s2events.azure-automation.net/webhooks?token=T9DXD... (change)

EVENT SUBSCRIPTION DETAILS

Name	<input type="text" value="testinfraeventgrid"/>
Event Schema	Event Grid Schema

FILTERS

Subject Begins With	<input type="text" value="Sample-workitems/{name}"/>
Subject Ends With	<input type="text" value=".jpg"/>
Labels	 Add Label

Create

For example, you can use Event Grid to catch all the write operations on an **Azure Virtual Machine** and then trigger an **Azure Automation** script that can log in to all the write operations with a text file by using **WebHook**. Here is a diagram of this interaction. You can

do more powerful things with Event Grid while being integrated with other Azure Services, shown as follows:



To trigger Stream big data into the Azure data warehouse, refer to the following link:

<https://docs.microsoft.com/en-us/azure/event-grid/event-grid-event-hubs-integration>

Service Bus is a cloud-based multitenant service that allows you to connect applications through the cloud. You can check the following link to find out more on how Event Grid responds to a Service Bus event:

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-to-event-grid-integration-example?toc=%2fazure%2fevent-grid%2ftoc.json>

9963799240 / 7730997544

Ameerpet / Kondapur

Hyderabad

Integrating with Logic Apps

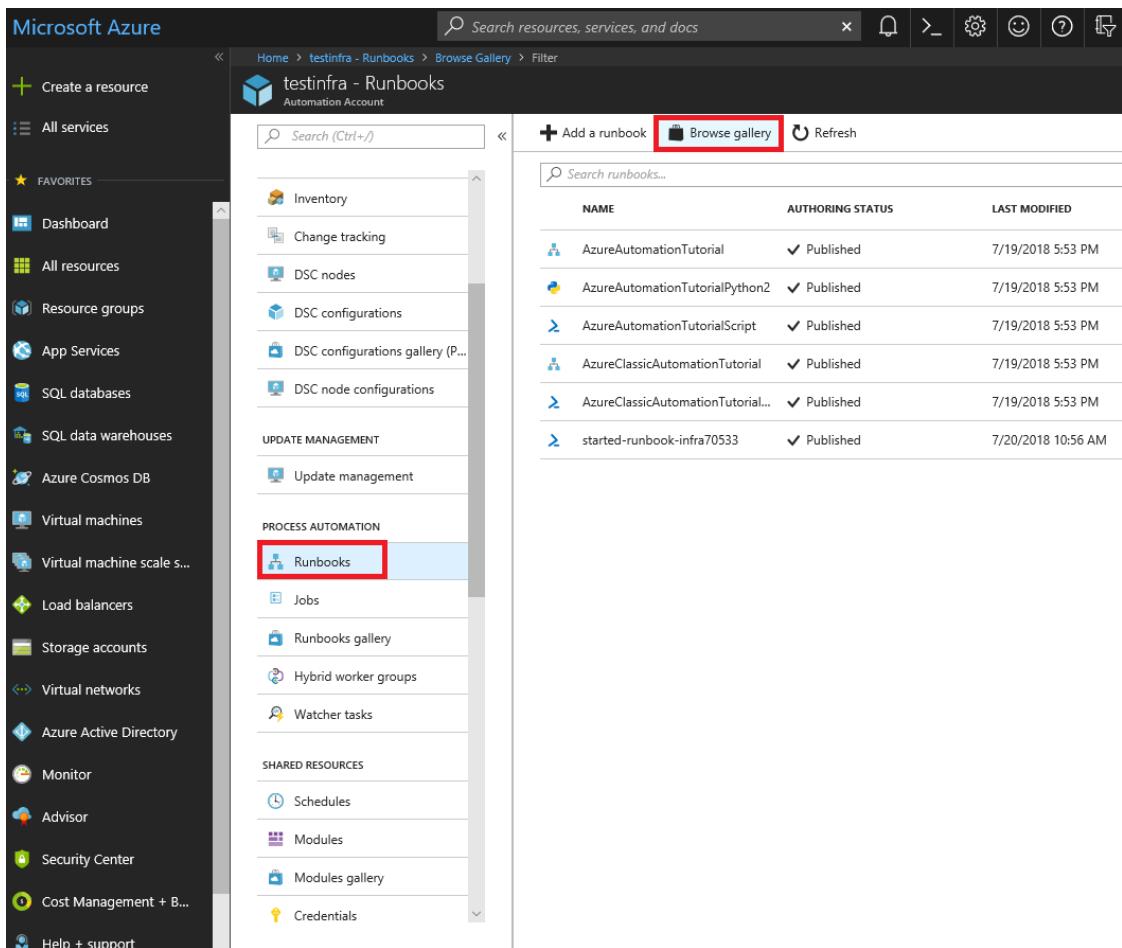
Logic Apps is a great PaaS offering to implement scalable logical workflows for cloud-native applications. Especially for those who are more comfortable designing workflows with graphic interfaces, it provides a visual designer to model and automate the process and the workflow with a couple of steps. Every Logic Apps workflow begins with a trigger and can execute the combinations of actions with conditional logic, in parallel, and sequentially.

For a great post about how Azure Automation works together with Logic Apps, please check this link:

https://blogs.technet.microsoft.com/stefan_stranger/2017/06/23/azur-logic-apps-schedule-your-runbooks-more-often-than-every-hour/

Runbook gallery

To know more about how Automation interacts with other Azure services, you can use existing runbooks in Runbooks gallery, which you can find in the **Runbooks** blade, and click on **Browse gallery**, as follows:



The screenshot shows the Microsoft Azure portal interface. On the left, the navigation menu is visible with various service categories like All services, Favorites, and Resource groups. In the center, under the 'testinfra - Runbooks' automation account, the 'Runbooks' blade is open. The 'Runbooks' section is highlighted with a red box. Below it, the 'Runbooks gallery' option is also highlighted with a red box. A search bar at the top right of the blade says 'Search runbooks...'. A table lists several runbooks with columns for Name, Authoring Status, and Last Modified. All listed runbooks are published and were last modified on July 19, 2018, or July 20, 2018.

Name	Authoring Status	Last Modified
AzureAutomationTutorial	✓ Published	7/19/2018 5:53 PM
AzureAutomationTutorialPython2	✓ Published	7/19/2018 5:53 PM
AzureAutomationTutorialScript	✓ Published	7/19/2018 5:53 PM
AzureClassicAutomationTutorial	✓ Published	7/19/2018 5:53 PM
AzureClassicAutomationTutorial...	✓ Published	7/19/2018 5:53 PM
started-runbook-infra70533	✓ Published	7/20/2018 10:56 AM

In the gallery, you can find the runbook conform to your search criteria using **Filter**; it will help you find the runbook you need in an effective way. See the following screenshot

Home > testinfra - Runbooks > Browse Gallery > Filter

Browse Gallery

Filter

event

Watcher action that processes events triggered by a watcher runbook
 PowerShell Runbook
 This sample automation runbook is designed to be used in a watcher task that takes action on data passed in from a watcher runbook. It is required to have a parameter called \$EVENTDATA in watcher action runbooks to receive information from the watcher
 Tags: Runbook, Automation Watcher

Created by: SC Automation Product Team
 Ratings: 5 of 5
 559 downloads
 Last updated: 11/17/2017

Integrating Azure Automation with Event grid
 PowerShell Runbook
 This sample Automation runbook integrates with Azure event grid subscriptions to get notified when a write command is performed against an Azure VM. The runbook adds a cost tag to the VM if it doesn't exist. It also sends an optional notification to a Microsoft Tags: Azure Automation, Microsoft Azure Virtual Machines, Event Grid

Created by: SC Automation Product Team
 380 downloads
 Last updated: 11/28/2017

Sample monitor runbook to watch for event IDs in an Azure Virtual Machine.
 PowerShell Workflow Runbook
 This runbook looks for a specific event ID in an Azure VM so that an action could be taken. It should be used with the Manage-MonitorRunbook runbook. Please refer to http://azure.microsoft.com/blog/2014/11/17/monitoring-azure-services-and-external- Tags: Tutorial

Created by: SC Automation Product Team
 Ratings: 5 of 5
 613 downloads
 Last updated: 11/18/2014

Azure: ARM Virtual Network Gateway Diagnostics (VPN)
 PowerShell Runbook
 This script sample demonstrates how the Azure Virtual Networks PowerShell cmdlets can be utilized to gather Azure Gateway Diagnostics data. This logging is useful for troubleshooting VPN connection issues, and includes messages such as MainMode SA Tags: Microsoft Azure, Networking, vpn monitor

Created by: symevent_ninja
 Ratings: 3 of 5
 519 downloads
 Last updated: 6/21/2016

Event Hub - Sample data generator
 PowerShell Runbook
 Generates sample Events for an Azure Event Hub. The sample data is specific to temperature sensor data and should be used inside an Azure Automation Runbook on a schedule. For best results I'd suggest creating a Stream Analytics job with your Event Hub Tags: Microsoft Azure, Powershell, Event Hub

Created by: GordonB007
 315 downloads
 Last updated: 3/3/2016

Find-AzureEndPoints

OK

Ameerpet / Kondapur
Hyderabad

Implementing monitoring solutions in Azure

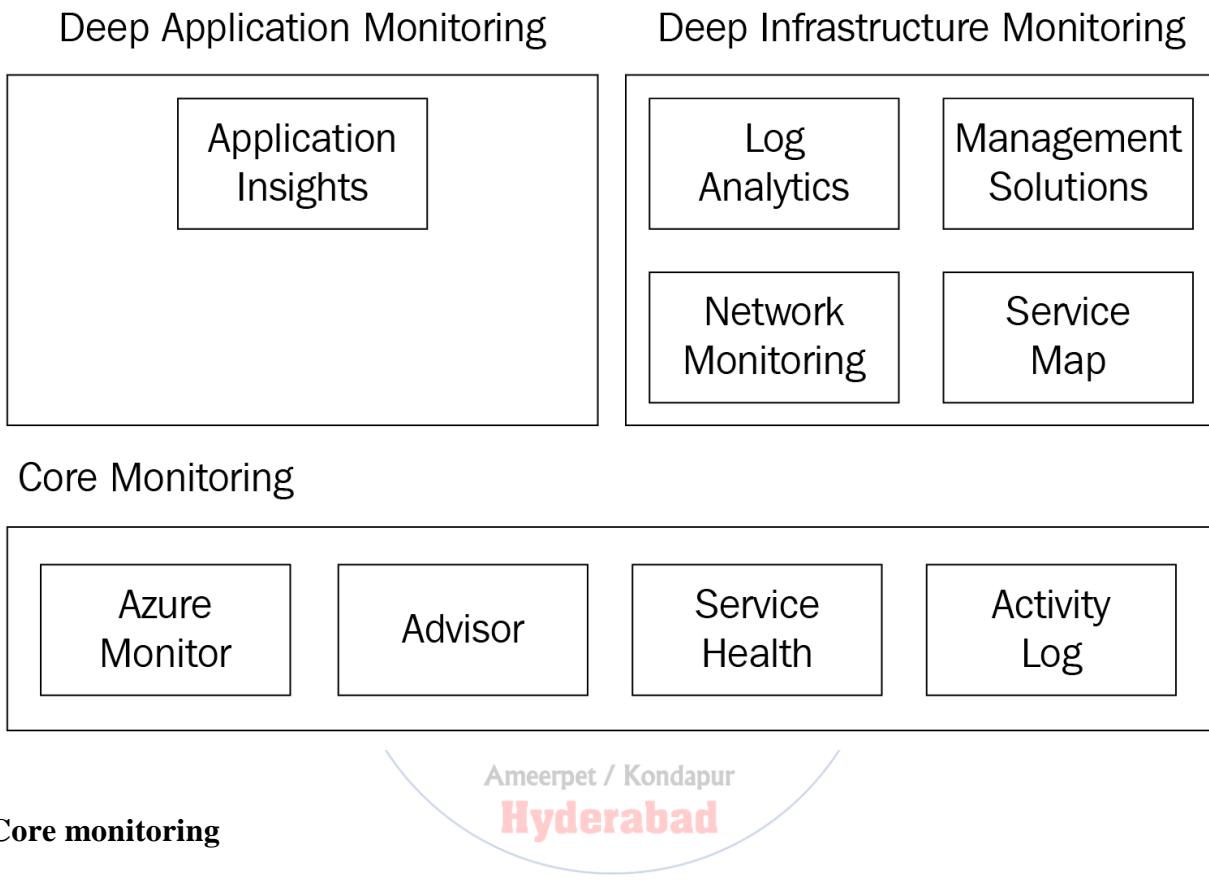
Monitoring in Azure covers the performance, health, and availability of your Azure resources to help users analyze issues and detect problems in case of failure.

Azure includes a couple of services performing monitoring tasks; they can work individually on telemetry or work together to provide a complete monitoring strategy for your application. There are the following two modes of monitoring in Azure:

- **Core monitoring** aims to provide fundamental and required monitoring across Azure resources.

- **Deep monitoring** includes services that go beyond core monitoring. It provides capabilities that collect and analyze data at a deeper level. There are two types: **deep application monitoring** and **deep infrastructure monitoring**.

The following diagram shows a conceptual view of the components that work together to provide monitoring of Azure resources:

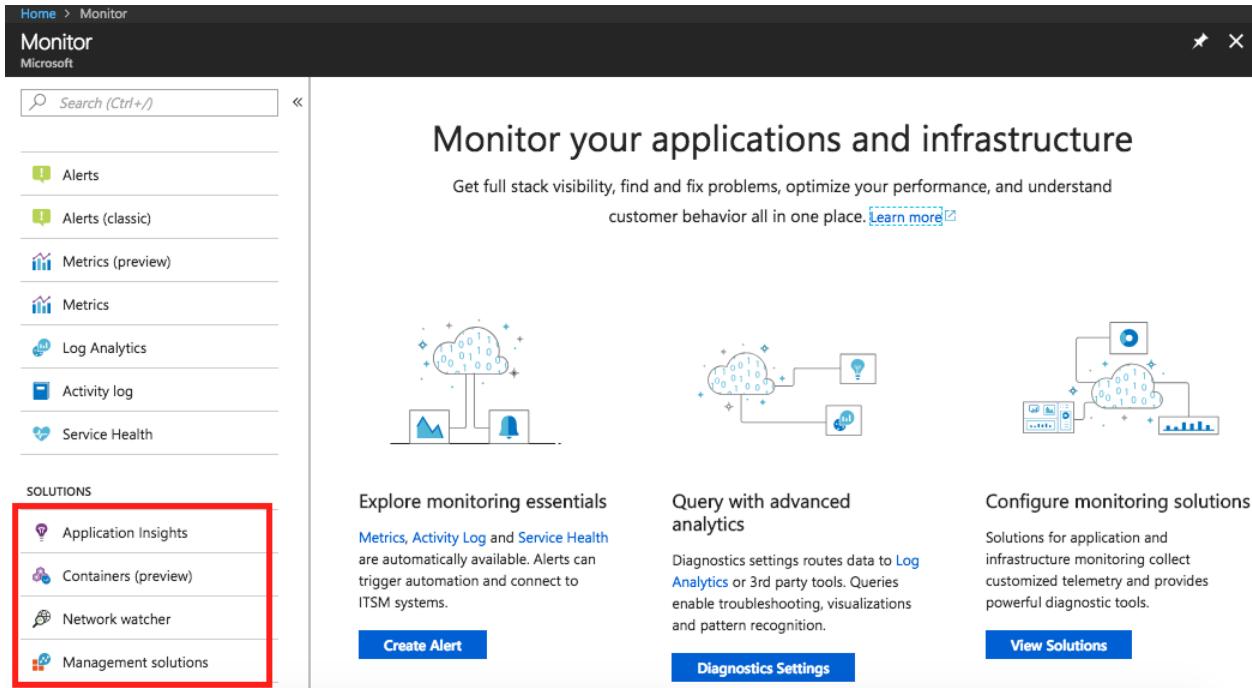


In this section, let's take a look at each service that performs fundamental and required monitoring across Azure resources.

Azure Monitor

Azure Monitor is a monitoring solution for applications and infrastructure in Azure where users can get full stack visibility, get help finding problems and resolutions, and understand customer behavior. It provides base-level infrastructure metrics and logs for many services such as cloud service, virtual machine, virtual machine scale sets, and service fabric in Microsoft Azure.

It also supports monitoring container solutions in Azure and connections to **Application Insights**, **Network watcher**, and **Management solutions** (OMS workspace) as well, shown as follows:

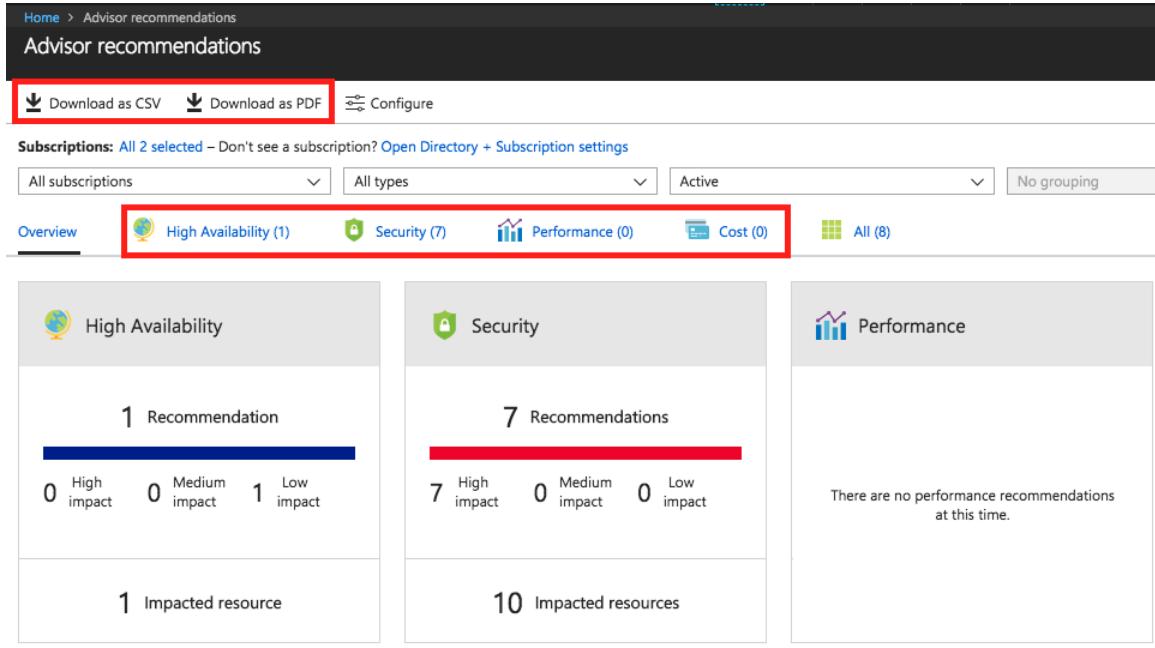


The screenshot shows the Microsoft Azure Monitor homepage. On the left, there's a navigation bar with links to Home, Monitor, and Microsoft. Below that is a search bar and a sidebar with links to Alerts, Metrics (classic), Metrics (preview), Metrics, Log Analytics, Activity log, and Service Health. Under 'SOLUTIONS', there are four items listed: Application Insights, Containers (preview), Network watcher, and Management solutions. The 'Management solutions' link is highlighted with a red box. The main content area features a title 'Monitor your applications and infrastructure' with a subtext about getting full stack visibility and understanding customer behavior. It includes three diagrams illustrating monitoring concepts: a cloud icon with binary code, a lightbulb connected to a cloud, and a complex network of clouds and icons. Below these are three sections: 'Explore monitoring essentials', 'Query with advanced analytics', and 'Configure monitoring solutions'. Each section has a brief description and a blue 'View Solutions' button.

Azure Monitor collects most application-level metrics and logs such as **Application Logs**, **Windows Event Logs**, **.NET Event Source**, **IIS Logs**, and **Customer Error Logs** using diagnostics extension.

Azure Advisor

Azure Advisor is a personalized advisor that Azure provides for you. If want to know more about Azure, it will be your best friend on your cloud journey. Azure Advisor helps you follow best practices to optimize your Azure deployments and analyzes your existing resource configuration, usage telemetry, and so on. It recommends a best solution in terms of cost, performance, high availability, and security and support to export these recommendations by PDF or CSV file, as shown in the following screenshot:



Subscriptions: All 2 selected – Don't see a subscription? Open Directory + Subscription settings

All subscriptions All types Active No grouping

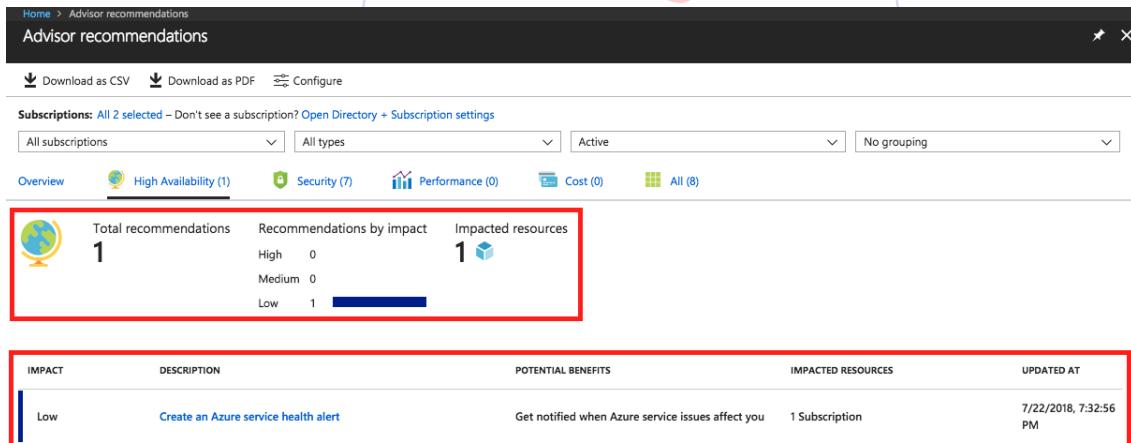
Overview **High Availability (1)** Security (7) Performance (0) Cost (0) All (8)

High Availability		
1 Recommendation		
0 High impact	0 Medium impact	1 Low impact
1 Impacted resource		

Security		
7 Recommendations		
7 High impact	0 Medium impact	0 Low impact
10 Impacted resources		

Performance		
There are no performance recommendations at this time.		

As the picture shows, you can click on the recommendation such as **High Availability**, and Advisor will give you further details about its recommendation, as shown here:



Subscriptions: All 2 selected – Don't see a subscription? Open Directory + Subscription settings

All subscriptions All types Active No grouping

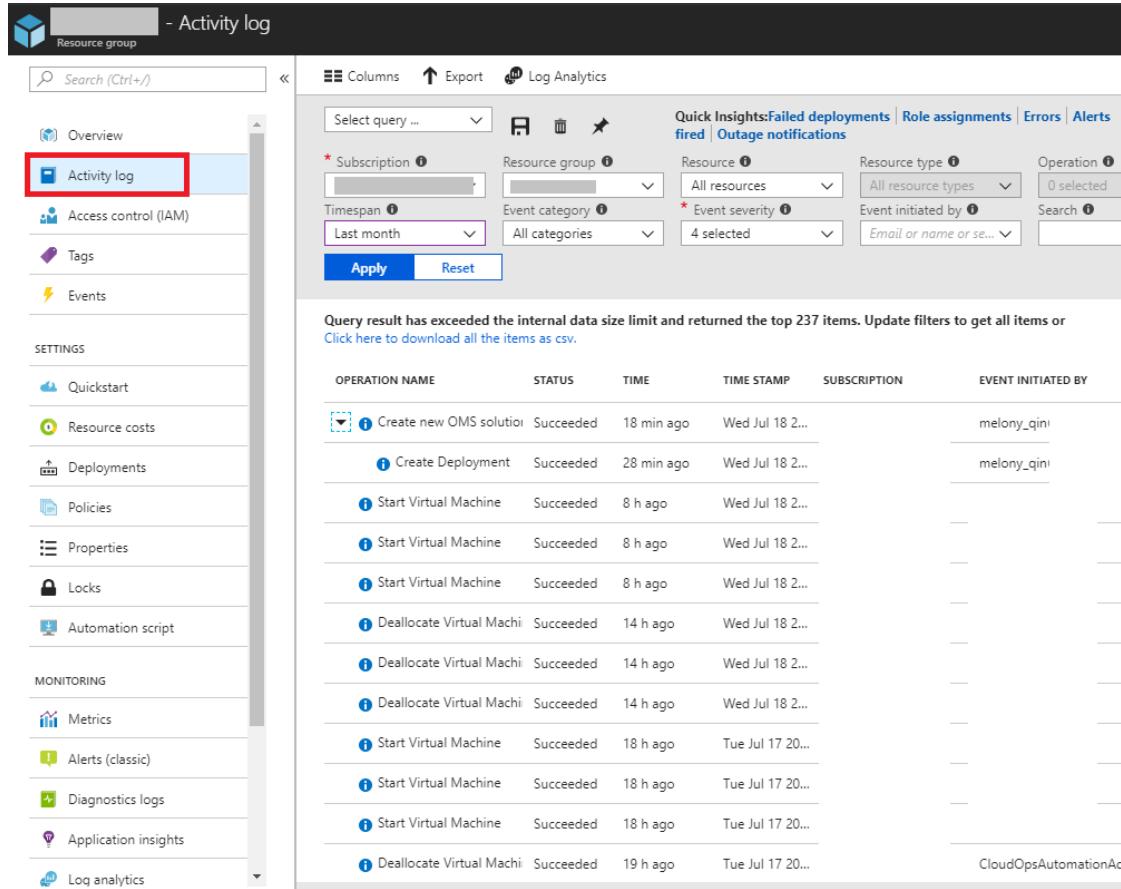
Overview **High Availability (1)** Security (7) Performance (0) Cost (0) All (8)

Total recommendations	Recommendations by impact	Impacted resources
1	High 0 Medium 0 Low 1	1 

IMPACT	DESCRIPTION	POTENTIAL BENEFITS	IMPACTED RESOURCES	UPDATED AT
Low	Create an Azure service health alert	Get notified when Azure service issues affect you	1 Subscription	7/22/2018, 7:32:56 PM

Activity log

Activity log is a log to trace events such as configuration changes, incidents, and operations related to the dedicated Azure resource. Users can view logs from the Azure Portal, as shown in the following screenshot:



The screenshot shows the Azure Activity log interface. On the left, a sidebar lists various resource group management options like Overview, Activity log (which is selected and highlighted with a red box), Access control (IAM), Tags, Events, Quickstart, Resource costs, Deployments, Policies, Properties, Locks, Automation script, Metrics, Alerts (classic), Diagnostics logs, Application insights, and Log analytics. The main pane displays a query results table with columns: OPERATION NAME, STATUS, TIME, TIME STAMP, SUBSCRIPTION, and EVENT INITIATED BY. The table lists 237 items, mostly succeeded operations such as creating OMS solutions, deployments, and virtual machines, initiated by users 'melony_qini' and 'CloudOpsAutomationAc'. A message at the top states: 'Query result has exceeded the internal data size limit and returned the top 237 items. Update filters to get all items or Click here to download all the items as csv.'

OPERATION NAME	STATUS	TIME	TIME STAMP	SUBSCRIPTION	EVENT INITIATED BY
Create new OMS solution	Succeeded	18 min ago	Wed Jul 18 2...		melony_qini
Create Deployment	Succeeded	28 min ago	Wed Jul 18 2...		melony_qini
Start Virtual Machine	Succeeded	8 h ago	Wed Jul 18 2...		
Start Virtual Machine	Succeeded	8 h ago	Wed Jul 18 2...		
Start Virtual Machine	Succeeded	8 h ago	Wed Jul 18 2...		
Deallocate Virtual Machi	Succeeded	14 h ago	Wed Jul 18 2...		
Deallocate Virtual Machi	Succeeded	14 h ago	Wed Jul 18 2...		
Deallocate Virtual Machi	Succeeded	14 h ago	Wed Jul 18 2...		
Start Virtual Machine	Succeeded	18 h ago	Tue Jul 17 20...		
Start Virtual Machine	Succeeded	18 h ago	Tue Jul 17 20...		
Start Virtual Machine	Succeeded	18 h ago	Tue Jul 17 20...		
Deallocate Virtual Machi	Succeeded	19 h ago	Tue Jul 17 20...		

Service health

Service health can help users to receive guidance of service issues and performs notification of information about the resources under users' subscriptions. This information is the subclass of activity log events, which means the same information can also be found in the activity log. There is a wide range of classes of service health notifications in terms of **Action required**, **Assisted recovery**, **Incident**, **Maintenance**, **Information**, and **Security**.

To know how to create activity log alerts on service notifications, you can refer to:

<https://docs.microsoft.com/en-us/azure/monitoring-and-diagnostics/monitoring-activity-log-alerts-on-service-notifications>

After creating a notification, you can check it at Azure Portal by clicking on **All services** and search Service health, you can go to the dashboard of **Service Health** where you can get the global vision of this service (as shown in the following screenshot):

Home > Service Health - Service issues

Service Health - Service issues

ACTIVE EVENTS

- Service issues**
- Planned maintenance
- Health advisories

HISTORY

- Health history

RESOURCE HEALTH

- Resource health

ALERTS

- Health alerts

Search (Ctrl+ /)

Select filter ...

* Subscription 2 selected * Region 5 selected * Service 140 selected

Save filter Delete filter Pin filtered world map to dashboard Create service health alert



No service issues found
See all past issues in the [health history](#).

[Launch guided tour](#)

Deep application monitoring

Application Insights plays a key role in deep application monitoring solutions provided by Azure; it can run on top of the Guest OS in the compute model. In the previous chapter, we introduced how to monitor Web App in App Service Plan using Application Insights. Actually, Azure Application Insights can not only be used for monitoring applications hosted in App Service Plan (API App, mobile app, and so on) but also for Azure Functions.

Azure Functions has been integrated with Azure Application Insights since April 2017. Users can find out how to configure Functions to send telemetry data to Application Insights and how it works through the following link:

<https://docs.microsoft.com/en-us/azure/azure-functions/functions-monitoring>

Application Insights also supports monitoring Docker-based applications in Azure; you can check here to get more information:

<https://docs.microsoft.com/en-us/azure/application-insights/app-insights-docker?toc=%2fazure%2fmonitoring%2ftoc.json>

Deep infrastructure monitoring

In this section, let's take a look at each service that performs infrastructure-level monitoring in depth across Azure resources.

Log Analytics

Log Analytics is also a very important role in Azure monitoring. It focuses on collecting data across Azure resources into a single repository, and it allows queries using **Log Analytics query language (KQL)** and analyzes collected data. Azure services such as Application Insights, Azure Security Center, Azure Monitor, Management Solutions, and agents installed on virtual machines in the cloud or on-premises can store data in the Log Analytics data store. Log Analytics acts as a core monitoring service in Azure. We'll dive deeper into it later in this chapter.

Management solutions

Management solutions are based on the monitoring data collected by Log Analytics and analyze them so that it can provide a global vision for a particular application or service. To know more about how to use and install the management solutions go to <https://docs.microsoft.com/en-us/azure/monitoring/monitoring-solutions>.

Network monitoring

There are a variety of tools that work together to provide a comprehensive monitoring solution in Azure or on-premise from various aspects of networking, and are as follows:

- **Network Watcher** performs monitoring, metrics, and enables or disables logs for resources in an Azure VNet. It stores data in Azure metrics and diagnostics for further analysis.
- **Network Performance Monitor (NPM)** is a monitoring solution that monitors connectivity across public clouds and on-premises data centers in the cloud.
- **ExpressRoute Monitor** is a subcapability of NPM, focusing on monitoring the end-to-end connectivity and performance over Azure ExpressRoute circuits.
- **DNS Analytics** is a service based on DNS servers that monitors the security, performance, and operations-related insights.

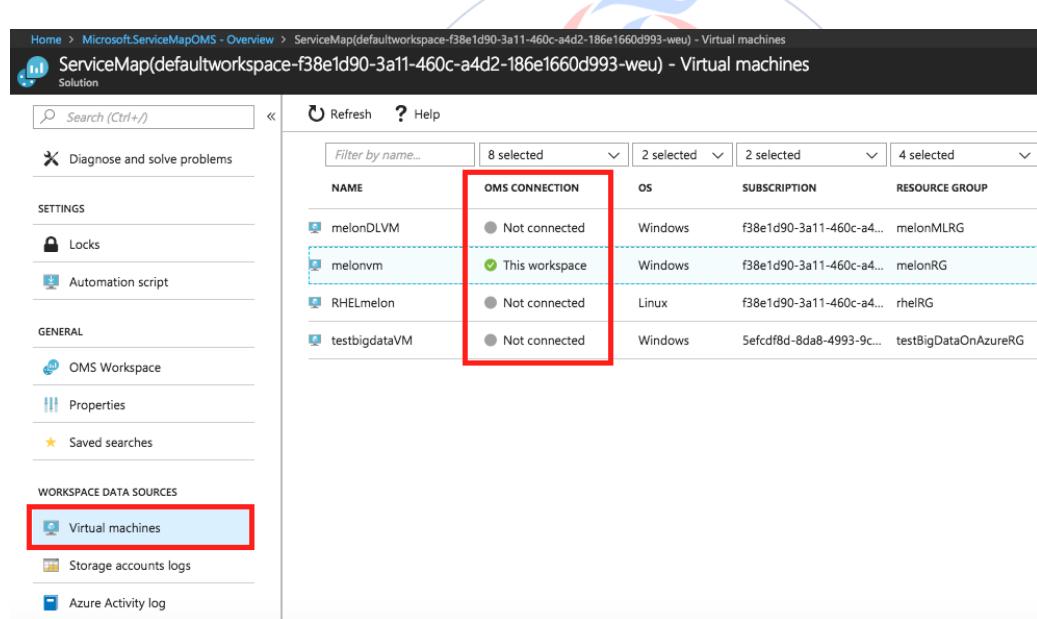
- **Service Endpoint Monitor** is a cloud-based service to monitor the reachability of applications and detects performance bottlenecks across cloud or on-premise data centers.

To know more about the networking monitoring solution in Azure, you can check out the following link:

<https://docs.microsoft.com/en-us/azure/networking/network-monitoring-overview>

Service Map

Service Map is similar to Application Map in Application Insights. It integrates events, performance data, and management solutions in Log Analytics. It provides insight for different processes within virtual machines and dependencies on other machines and external processes, as shown in the following screenshot:



The screenshot shows the Microsoft ServiceMap OMS - Overview page. The left sidebar includes links for Diagnose and solve problems, Locks, Automation script, OMS Workspace, Properties, and Saved searches. Under WORKSPACE DATA SOURCES, 'Virtual machines' is highlighted with a red box. The main content area displays a table of virtual machines with columns: NAME, OMS CONNECTION, OS, SUBSCRIPTION, and RESOURCE GROUP. The 'OMS CONNECTION' column shows status: Not connected (for melonDLVM, RHELMelon, testbigdataVM) and This workspace (for melonvm). A red box highlights the 'This workspace' entry for melonvm.

NAME	OMS CONNECTION	OS	SUBSCRIPTION	RESOURCE GROUP
melonDLVM	Not connected	Windows	f38e1d90-3a11-460c-a4d2-186e1660d993-weu	melonMLRG
melonvm	This workspace	Windows	f38e1d90-3a11-460c-a4d2-186e1660d993-weu	melonRG
RHELMelon	Not connected	Linux	f38e1d90-3a11-460c-a4d2-186e1660d993-weu	rhelRG
testbigdataVM	Not connected	Windows	5efcdf8d-8da8-4993-9c...	testBigDataOnAzureRG

Shared capabilities

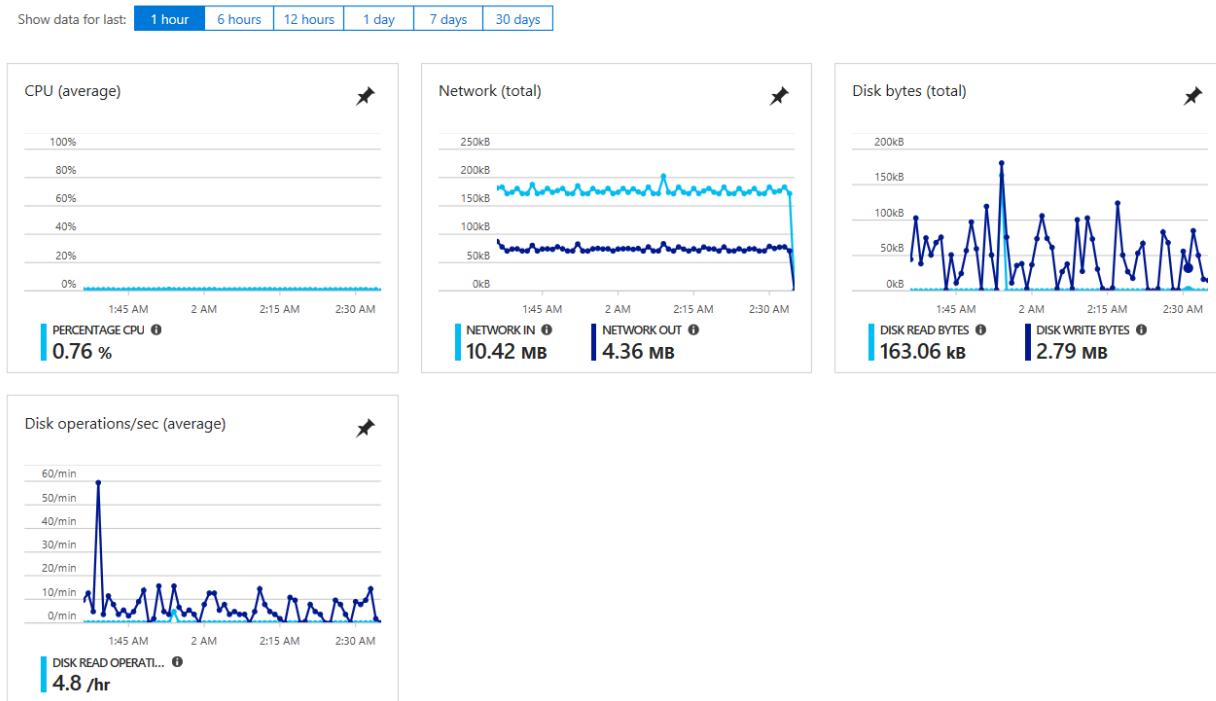
Shared capabilities of monitoring strategy in Azure is realized by some shared functionalities, such as **alerts**, **dashboards**, and **metrics** that provide a basic understanding view of the current status of Azure resources. An excellent monitoring solution helps users understand better how do their solutions work and detect the exception in case of issue. It also performs proactive notification of critical issues so that users can take care of them before the real problem comes.

Implementing Azure VMs monitoring solutions

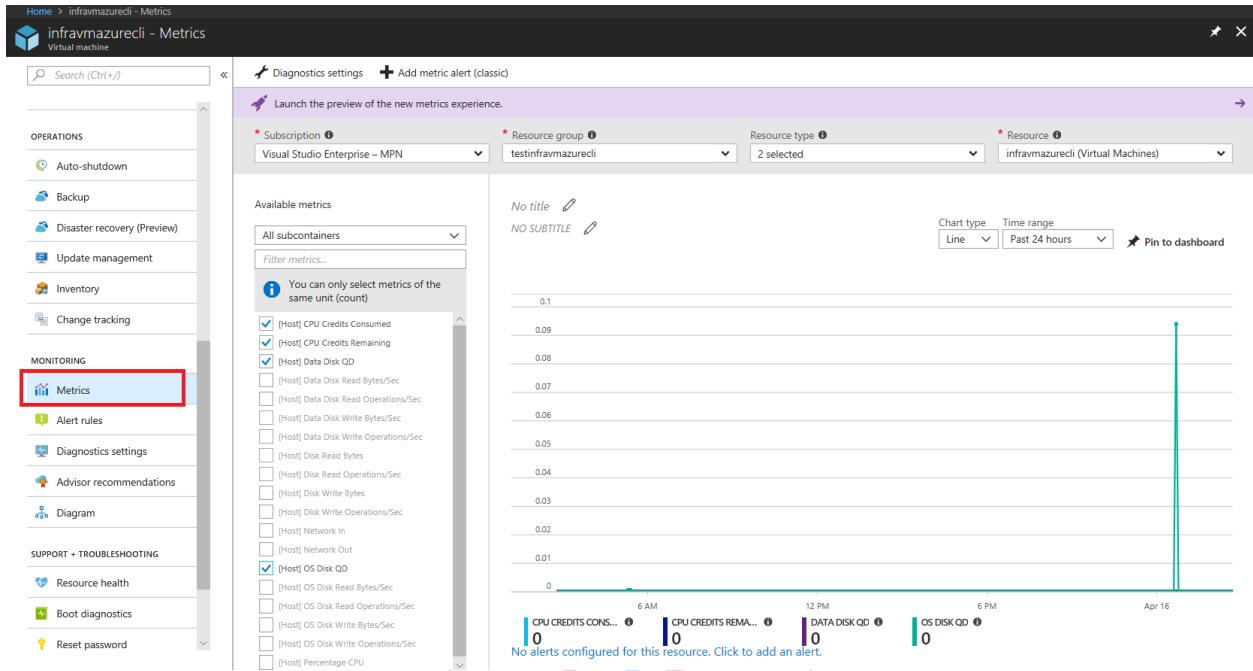
Like most Azure services, Azure VMs enable you to track their performance, availability, and usage. This data is available directly from the Azure Portal. You can also collect Azure VM metrics and diagnostics via Azure PowerShell and Azure CLI scripts. In addition, you can collect this data programmatically via the REST API and Azure SDKs.

Configuring ARM VM monitoring

Azure VM monitoring provides metrics such as percentage CPU, Network in and out, disk read bytes, disk write bytes, disk read operations per seconds, and disk write operations per seconds, as shown in the following screenshot:



Azure also provides a way to query metrics in the **Metrics** blade of Azure VM, shown as follows:



Configuring alerts

Azure provides alert rules that allow users to trigger notifications based on metrics-based criteria that they had to specify. In Azure, each rule includes a metric, condition, threshold, and time period that collectively determine when to raise an alert. You can configure your email address to get the alert notification. It also supports the Webhook, which is a route alert to an arbitrary HTTP or HTTPS endpoint. With alerts, it is possible to configure a response using an Azure Automation runbook, as shown in the following screenshot:

Add activity log alert

* Activity log alert name ✓

Description ✓

* Subscription ▾

* Resource group ▾

Source

Input

Criteria

* Event category ▾

* Service(s) ▾

* Region(s) ▾

* Type ▾

Alert via

Action group New Existing

* Action group name ✓

* Short name ✓

Actions

ACTION NAME	ACTION TYPE	STATUS	DETAILS
<input type="text" value="email"/> ✓	<input type="text" value="Email/SMS/Push/Voice"/> ▾		Edit details

OK

Email/SMS/Push/Voice

Name

Email

SMS

Country code * Phone number

Carrier charges may apply.

Azure app Push Notifications

Learn about the connecting to your Azure resources using the Azure app.

✓

This is the email you use to log into your Azure account.

Voice

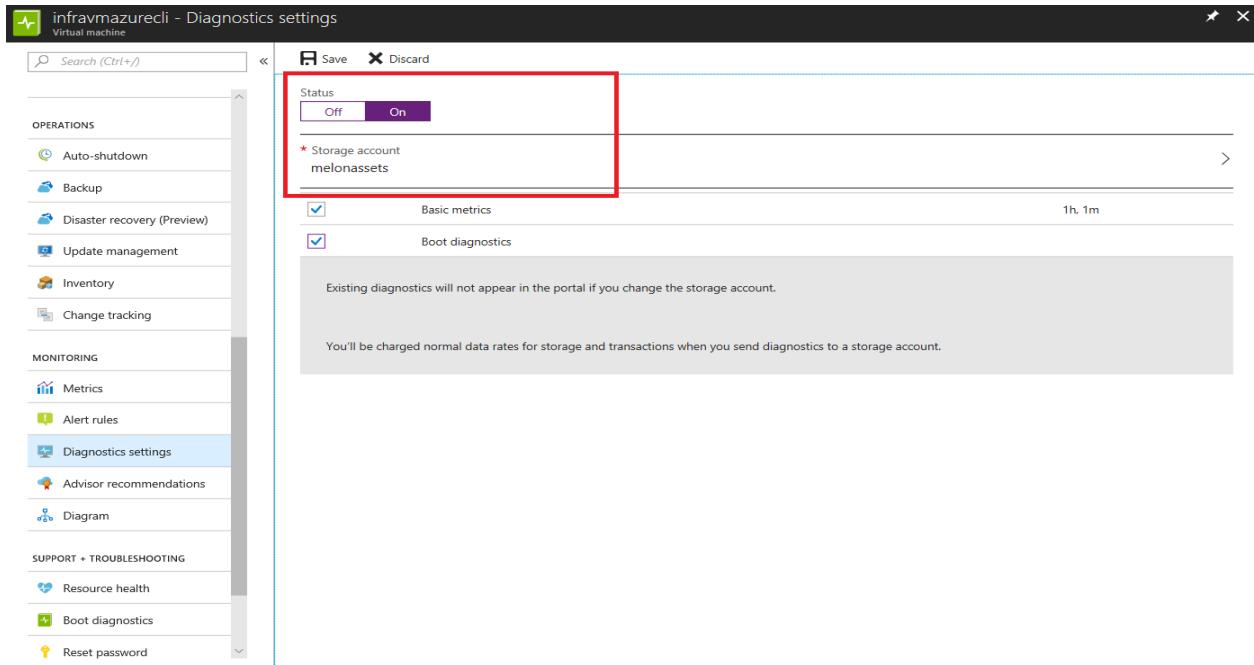
Country code * Phone number

OK

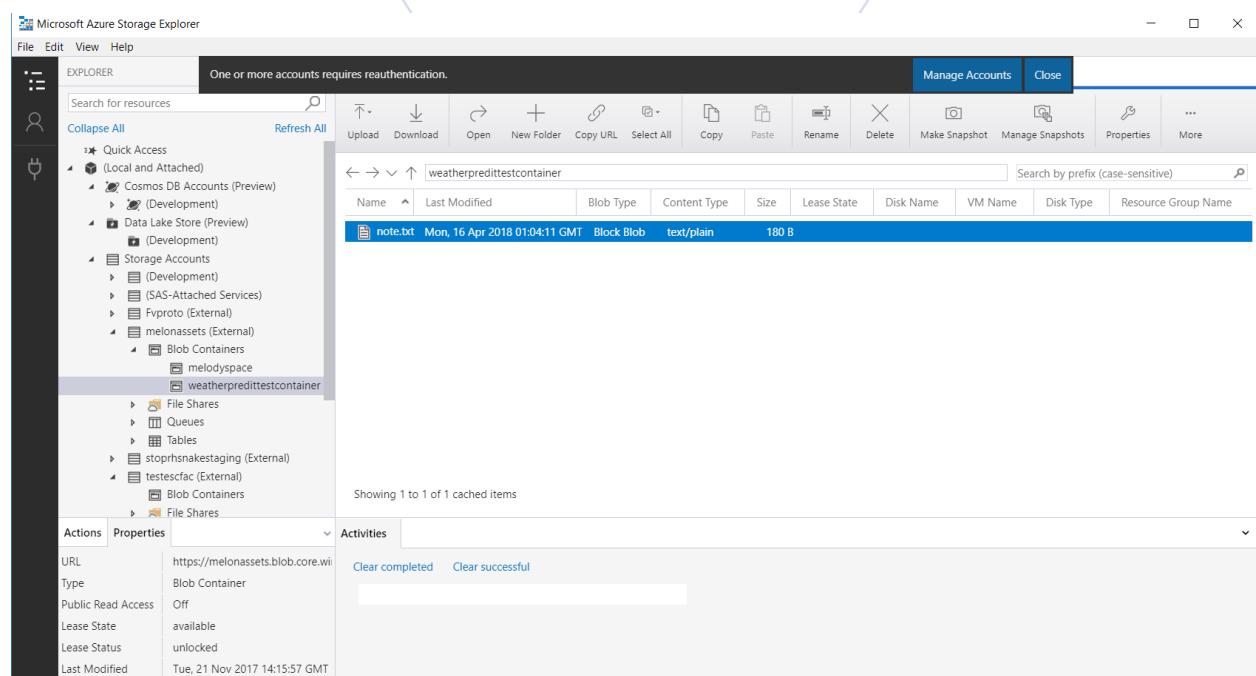
Ameerpet / Kondapur
Hyderabad

Configuring diagnostic and monitoring storage location

Azure provides insight into the performance and state of an Azure VM's OS by enabling diagnostics. For a Windows-based Azure VM, after enabling it, Azure will be able to collect data and monitor them via basic metrics including CPU usage, memory usage, network in and out , and so on, to enable **Boot diagnostics**, Azure will create the logs such as event logs, IIS logs and failed request logs, crash dumps. To enable diagnostics, you must configure a standard storage account so that the collected data can be stored; you can do it by setting the **Status** as **On** value and choosing the storage account that you want to use, as shown in the next screenshot:



To view and analyze diagnostics and logs, you can use a tool, such as **Azure Storage Explorer**, that provides access to tables and blobs in the Azure Storage account that is hosting collected data. It is also possible to export the data into an Excel file or any other business intelligence application (such as Power BI) for further analysis, as shown in the next screenshot:

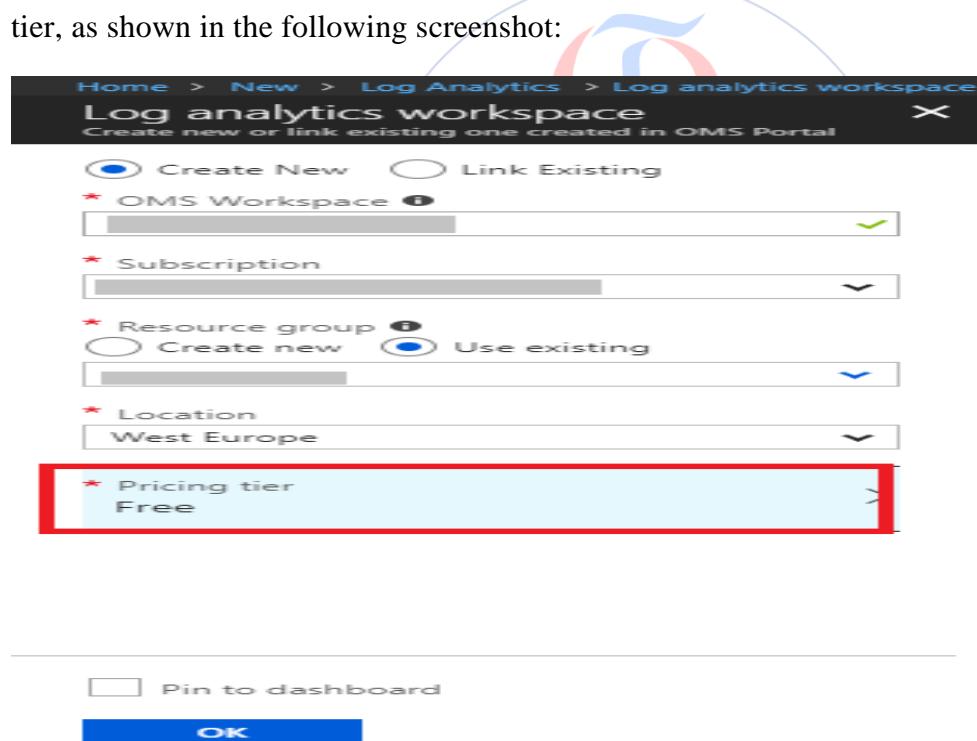


Implementing Log Analytics (OMS) solutions

To enable Log Analytics (OMS), you can go to Azure Portal by searching Log Analytics or enable it via at the resource group level.

Creating an OMS workspace

You may start from creating an OMS workspace where you should fill in basic information such as workspace name, resource group, and subscription name and choose the right pricing tier, as shown in the following screenshot:



To choose the pricing tier of OMS solution, note the following screenshot:

Pricing Tier □ X

The cost of your workspace depends on the pricing tier and what solutions you use. Learn more about [Log Analytics pricing](#).

This subscription is currently in an older pricing model with access to multiple pricing tiers. Learn more about the [new pricing model](#) and [assessing if you should adopt it](#). Change the monitoring pricing model for this subscription on the Pricing model selection page under Monitor > Usage and estimated costs.

Pricing Tier
 ▼

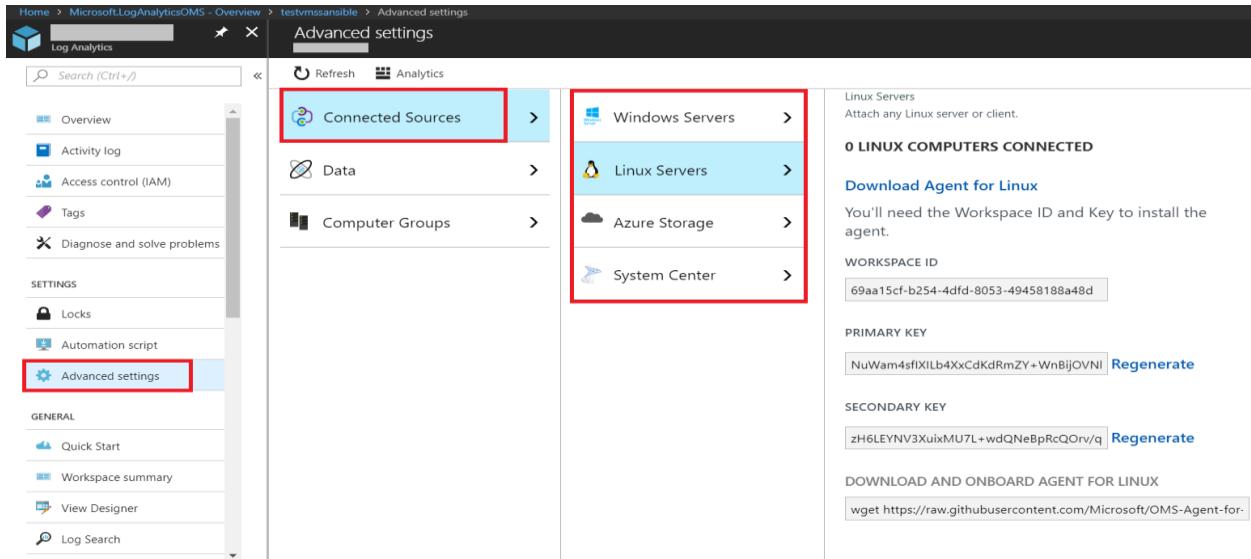
To use Operations Management Suite entitlements choose "Per Node (OMS)".

OK

As shown you may have some options:

- **Free:** The old free tier has a 500 MB limit on the amount of data collected daily and doesn't allow for data retention periods longer than 7 days. The new pricing model does not have any limits on the amount of data collected daily and allows you to retain your log data for up to 2 years.
- **Per node (OMS):** A node is any physical server or virtual machine that is managed by the Insight and Analytics service such as Log Analytics, Service Map, or Network Performance Monitor. For details, refer to <https://azure.microsoft.com/en-us/pricing/details/insight-analytics/>
- **Per gigabyte (GB):** Log Analytics is billed per gigabyte (GB) of data ingested into the service. As of writing this book, the first 5 GB of data ingested to the Azure Log Analytics service every month is free but every GB of ingested data (even after the first 5 GB) will be retained at no charge only for the first 31 days)

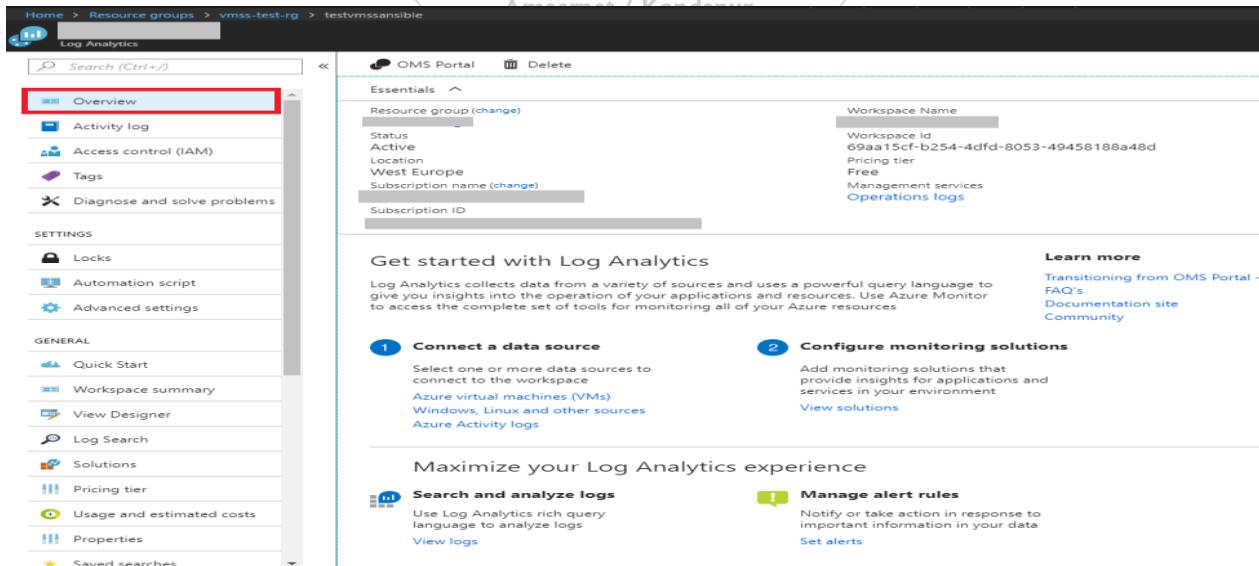
If you want to know if there are any services under monitor by current workspace of Log Analytics, check the connected Azure resource by clicking **Connected Sources** as shown below:



This screenshot shows the Microsoft Log Analytics OMS - Overview blade. The left sidebar has a red box around the "Advanced settings" link under the "SETTINGS" section. The main content area shows the "Advanced settings" page with a red box around the "Connected Sources" section. This section lists "Windows Servers", "Linux Servers", "Azure Storage", and "System Center", each with a right-pointing arrow. Below this is a "PRIMARY KEY" field containing "NuWam4sfIXILb4XxCdKdRmZY+WnBijOVN" with a "Regenerate" button, and a "SECONDARY KEY" field containing "zH6LEYNV3XuixMU7L+wdQNeBpRcQOrv/q" with a "Regenerate" button. At the bottom, there's a "DOWNLOAD AND ONBOARD AGENT FOR LINUX" section with a command-line wget link.

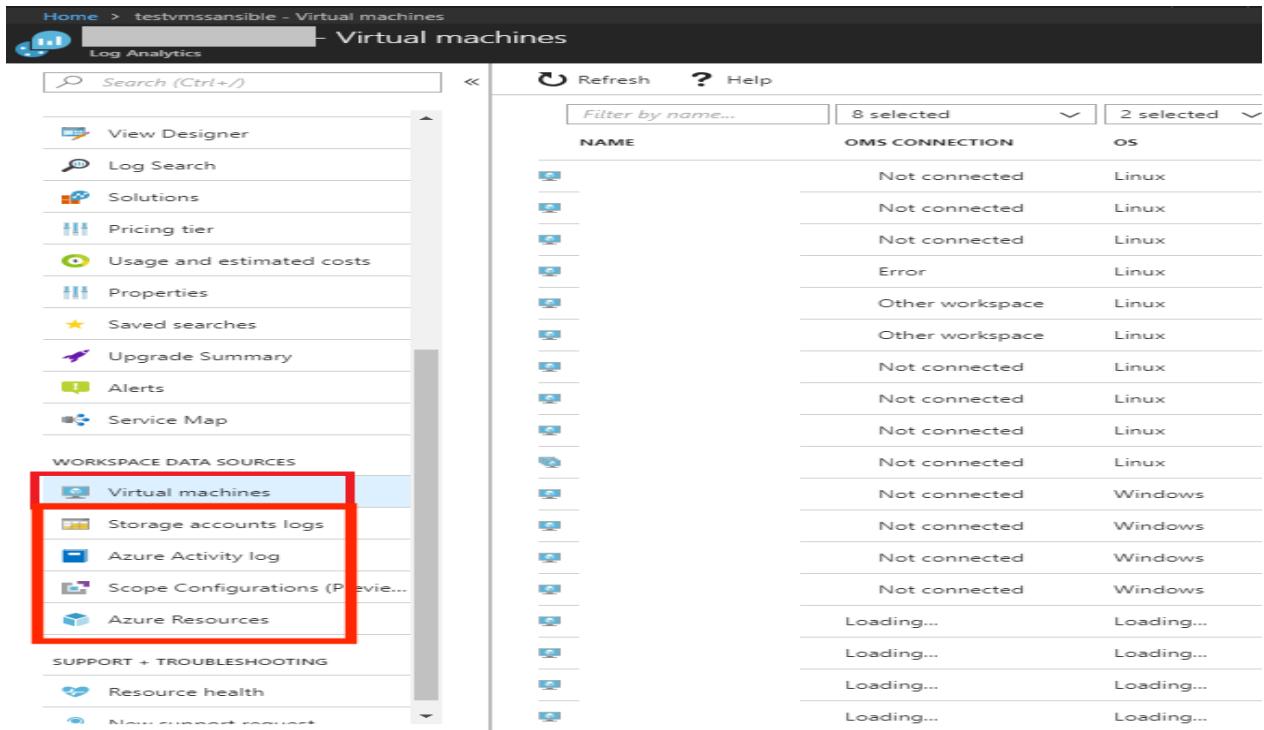
Collecting and searching across data sources from multiple systems

Log Analytics collects data from configured connected sources and stores it in Log Analytics workspace. To configure the Data Sources, you can go to the **Overview** blade and click on **Connect a data source** (as shown in the next screenshot); Log Analytics also allows you to maximize your Log Analytics experience by configuring search and analyze logs and manage alert rules so that you have a notification and will be able to take action in case issues arise:



This screenshot shows the Microsoft Log Analytics OMS - Overview blade. The left sidebar has a red box around the "Overview" link under the "GENERAL" section. The main content area features a large central box titled "Get started with Log Analytics". It includes sections for "Connect a data source" (with options for Azure virtual machines (VMs), Windows, Linux, and other sources), "Configure monitoring solutions" (with a "View solutions" link), and "Maximize your Log Analytics experience" (with "Search and analyze logs" and "Manage alert rules" sections). On the right side, there's a "Learn more" section with links to "Transitioning from OMS Portal - FAQ's", "Documentation site", and "Community". At the bottom, there's a "Workspace Name" summary table and a "Subscription ID" table.

You can also go to the **WORKSPACE DATA SOURCES** section and configure the data source (as shown here):



NAME	OMS CONNECTION	OS
	Not connected	Linux
	Not connected	Linux
	Not connected	Linux
	Error	Linux
	Other workspace	Linux
	Other workspace	Linux
	Not connected	Windows
	Loading...	Loading...

Transforming Azure activity data and managed resource data

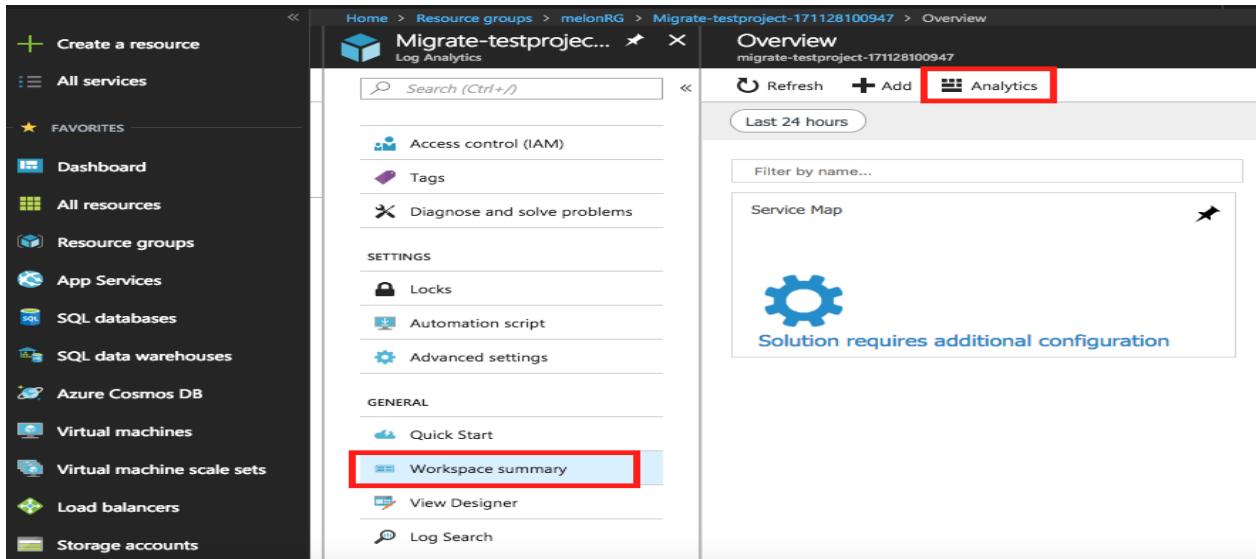
9963799240 / 7730997544

Ameerpet / Kondapur

Hyderabad

The Activity Log helps users analyze and search the Azure activity log across Azure subscriptions. The Activity Log focuses on providing insights into the operations level, with information regarding what, who, and when for any write operations, such as PUT, POST, and DELETE made for the resources across a subscription. Log Analytics offers the **Log search** functions to let users find the information they need in an effective way. Another way to find the needed monitoring information is to use queries. You can go to a Log Analytics portal to write these queries, and you can find this information in the output.

You can find your Log Analytics Portal (as shown in the following screenshot) by clicking on **Workspace summary** and then **Analytics**:

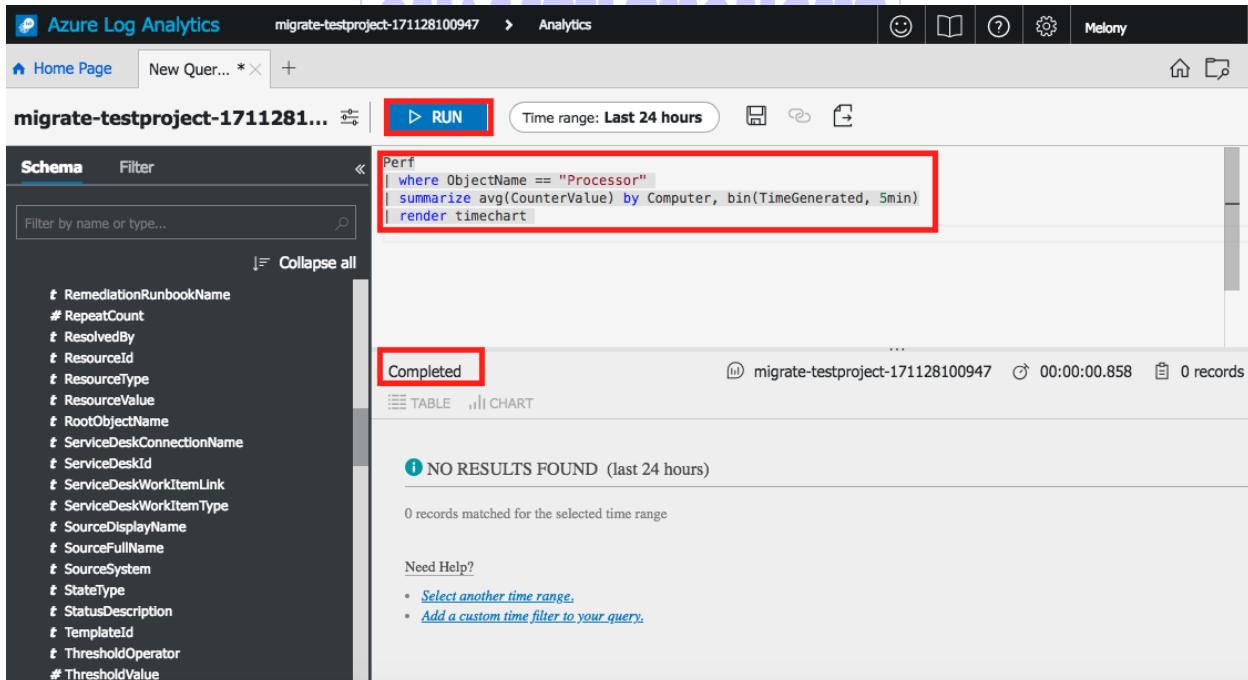


The screenshot shows the Azure Log Analytics Overview page. At the top, there's a navigation bar with 'Home', 'Resource groups', 'melonRG', 'Migrate-testproject-171128100947', 'Overview', 'Refresh', 'Add', and 'Analytics' (which is highlighted with a red box). Below the navigation is a search bar labeled 'Search (Ctrl+ /)'. To the right of the search bar are three buttons: 'Access control (IAM)', 'Tags', and 'Diagnose and solve problems'. Further down are sections for 'SETTINGS' (Locks, Automation script, Advanced settings) and 'GENERAL' (Quick Start, Workspace summary, View Designer, Log Search). A large blue gear icon with the text 'Solution requires additional configuration' is prominently displayed.

Here is a sample URL for the Log Analytics Portal:

[https://portal.loganalytics.io/subscriptions/...](https://portal.loganalytics.io/subscriptions/)

The portal contains a query section and output section, as follows:



The screenshot shows the Azure Log Analytics workspace 'migrate-testproject-171128100947' in the Analytics tab. The top navigation bar includes 'Azure Log Analytics', 'migrate-testproject-171128100947', 'Analytics', and several icons. The main area has a 'New Quer...' tab, a 'RUN' button, and a 'Time range: Last 24 hours' selector. The query editor contains the following code, which is highlighted with a red box:

```
Perf
| where ObjectName == "Processor"
| summarize avg(CounterValue) by Computer, bin(TimeGenerated, 5min)
| render timechart
```

The results pane shows a 'Completed' status and a message 'NO RESULTS FOUND (last 24 hours)'. It also indicates '0 records matched for the selected time range'. The schema pane on the left lists various log fields such as 'RemediationRunbookName', 'RepeatCount', 'ResolvedBy', etc.

Here is an example of a query that means to find all the updated operation times with more than 1 hour in your current workspace and display them with classification:

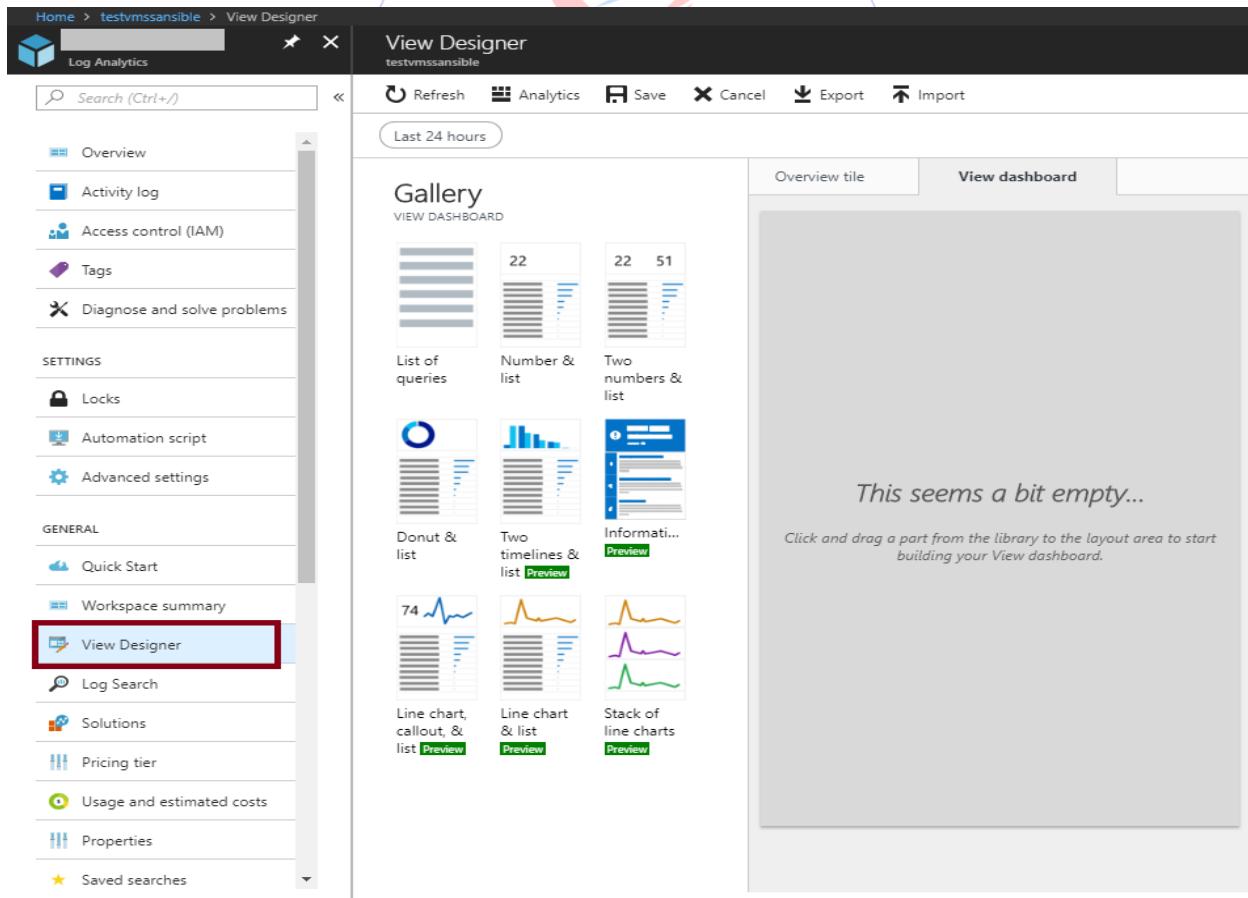
```
union Update, workspace("testinfra-workspace").Update
| where TimeGenerated >= ago(1h)
| summarize dcount(Computer) by Classification
```

To know more about log search and how to write queries, please check the following link:

<https://docs.microsoft.com/en-us/azure/log-analytics/log-analytics-log-search>

Building custom visualizations with view designer

Log Analytics makes it possible to collect events from text files and display them as your request. To do this, you can use **View Designer** in Log Analytics; go to your Log Analytics and click on the **View Designer** blade as shown in the following screenshot:



Build custom visualizations by using view designer

You can also check the following links to know more about:

- How to define a custom log file:<https://docs.microsoft.com/en-us/azure/log-analytics/log-analytics-data-sources-custom-logs>
- How to add your own searchable fields to extend existing records in Log Analytics:<https://docs.microsoft.com/en-us/azure/log-analytics/log-analytics-custom-fields>

Sending data to Log Analytics with the HTTP Data Collector API

You may also need to send or collect data to Log Analytics from a REST API client. To know more about how to use HTTP Data Collector API, checkout the following link:

<https://docs.microsoft.com/en-us/azure/log-analytics/log-analytics-data-collector-api>

IT Service Management Connector

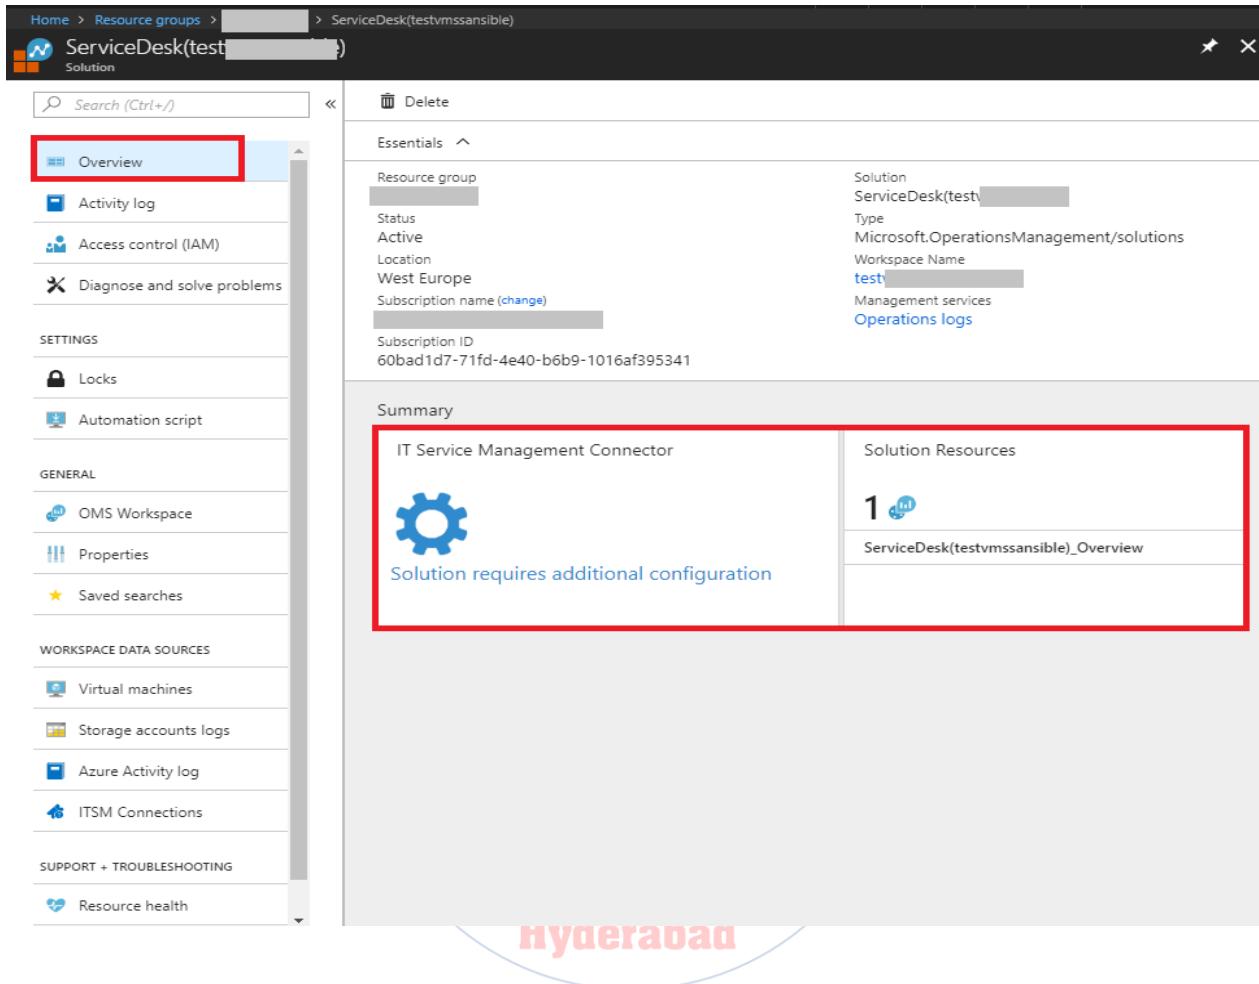
The **IT Service Management Connector (ITSMC)** is a service in Azure to help users to connect Azure and **IT Service Management (ITSM)**.

The ITSM connector provides connection between Azure and ITSM tools to help users resolve troubleshooting issues that are detected by Azure services, such as Azure Monitor and Log Analytics that reside in an ITSM service. These tools allow you to create work items in ITSM tools, based on your Azure alerts, such as metric alerts, Activity Log alerts, and Log Analytics alerts.

ITSMC supports connections with the ITSM tools, such as ServiceNow, System Center Service Manager, Provance, and Cherwell.

You can create an ITSMC via Azure Portal, and you only need to search for the term **IT Service Management Connector**. Fill in the related information and you can create it right

away. Thereafter you can manage it with Log Analytics or an other mentioned Azure service, shown as follows:



Home > Resource groups > ServiceDesk(testvmssansible) > ServiceDesk(testvmssansible)

ServiceDesk(testvmssansible)
Solution

Search (Ctrl+ /)

Overview (highlighted)

Activity log

Access control (IAM)

Diagnose and solve problems

SETTINGS

Locks

Automation script

GENERAL

OMS Workspace

Properties

Saved searches

WORKSPACE DATA SOURCES

Virtual machines

Storage accounts logs

Azure Activity log

ITSM Connections

SUPPORT + TROUBLESHOOTING

Resource health

Delete

Essentials ^

Resource group: ServiceDesk(testvmssansible)

Status: Active

Location: West Europe

Subscription name (change): test

Subscription ID: 60bad1d7-71fd-4e40-b6b9-1016af395341

Management services: Operations logs

Summary

IT Service Management Connector

Solution Resources

1 Solution

ServiceDesk(testvmssansible)_Overview

Solution requires additional configuration

hyderabad

Section 1: General Cloud Questions

1. What are the different types of services offered in the cloud?

IAAS	PAAS	SAAS
In infrastructure as a service, you get the raw hardware from your cloud provider as a service i.e you get a server which you can configure with your own will.	Platform as a Service, gives you a platform to publish without giving the access to the underlying software or OS.	You get software as a service in Azure, i.e no infrastructure, no platform, simple software that you can use without purchasing it.
For Example: Azure VM, Amazon EC2.	For example: Web Apps, Mobile Apps in Azure.	For example: when you launch a VM on Azure, you are not buying the OS, you are basically renting it for the time you will be running that instance.

2. What is cloud computing?

Explanation: It is the use of servers on the internet to “store”, “manage” and “process” data. The difference is, instead of using your own servers, you are using someone else’s servers to do your task, paying them for the amount of time you use it for.

3. What are the different cloud deployment models?

Explanation: Following are the three cloud deployment models:

Public Cloud: The infrastructure is owned by your cloud provider and the server that you are using could be a multi-tenant system.

Private Cloud: The infrastructure is owned by you or your cloud provider gives you that service exclusively. For eg: Hosting your website on your servers, or hosting your website with the cloud provider on a dedicated server.

Hybrid Cloud: When you use both Public Cloud, Private Cloud together, it is called Hybrid Cloud. For Example: Using your in-house servers for confidential data, and the public cloud for hosting your company's public facing website. This type of setup would be a hybrid cloud.

4. I have some private servers on my premises, also I have distributed some of my workload on the public cloud, what is this architecture called?

- A. Virtual Private Network
- B. Private Cloud
- C. Virtual Private Cloud
- D. Hybrid Cloud

Answer: D. Hybrid Cloud

Explanation: This type of architecture would be a hybrid cloud. Why? Because we are using both, the public cloud, and on premises servers i.e the private cloud. To make this hybrid architecture easy to use, wouldn't it be better if your private and public cloud were all on the same network (virtually). This is established by including your public cloud servers in a virtual private cloud, and connecting virtual cloud with your on premise servers using a VPN (Virtual Private Network).

Apart from this Azure Interview Questions Blog, if you want to get trained from professionals on this technology, you can opt for a structured training from edureka! Click below to know more.

Section 2: Basic Azure Questions

5. What is Microsoft Azure and why is it used?

Explanation: As discussed above, the companies which provide the cloud service are called the Cloud Providers. There are a lot of cloud providers out there, out of them one is Microsoft Azure. It is used for accessing Microsoft's infrastructure for cloud.

6. Which service in Azure is used to manage resources in Azure?

- A. Application Insights
- B. Azure Resource Manager
- C. Azure Portal
- D. Log Analytics

Answer: B Azure Resource Manager

Explanation: Azure Resource Manager is used to “manage” infrastructures which involve a no. of azure services. It can be used to deploy, manage and delete all the resources together using a simple JSON script.

7. Which of the following web applications can be deployed with Azure?

- A. ASP.NET
- B. PHP
- C. WCF
- D. All of the mentioned

Answer: D All of the mentioned

Explanation: Microsoft also has released SDKs for both Java and Ruby to allow applications written in those languages to place calls to the Azure Service Platform API to the AppFabric Service.

Section 3: Azure Interview Questions

8. What are Roles and why do we use them?

Explanation: Roles are nothing servers in layman terms. These servers are managed, load balanced, Platform as a Service virtual machines that work together to achieve a common goal.

There are 3 types of roles in Microsoft Azure:

- Web Role
- Worker Role
- VM Role

Let's discuss each of these roles in detail:

- **Web Role** – A web role is basically used to deploy a website, using languages supported by the IIS platform like, PHP, .NET etc. It is configured and customized to run web applications.
- **Worker Role** – A worker role is more like an help to the Web role, it used to execute background processes unlike the Web Role which is used to deploy the website.

- **VM Role** – The VM role is used by a user to schedule tasks and other windows services. This role can be used to customize the machines on which the web and worker role is running.

9. A _____ role is a virtual machine instance running Microsoft IIS Web server that can accept and respond to HTTP or HTTPS requests.

- A. Web
- B. Server
- C. Worker
- D. Client

Answer: A. Web

Explanation: The answer should be Web Roles, there are no roles such as Server or Client roles. Also, Worker roles can only communicate with Azure Storage or through direct connections to clients.

Apart from this Azure Interview Questions Blog, if you want to get trained from professionals on this technology, you can opt for a structured training from edureka! Click below to know more.

10. Is it possible to create a Virtual Machine using Azure Resource Manager in a Virtual Network that was created using classic deployment?

Explanation: This is not supported. You cannot use Azure Resource Manager to deploy a virtual machine into a virtual network that was created using classic deployment.

11. What are virtual machine scale sets in Azure?

Explanation: Virtual machine scale sets are Azure compute resource that you can use to deploy and manage a set of identical VMs. With all the VMs configured the same, scale sets are designed to support true autoscale, and no pre-provisioning of VMs is required. So it's easier to build large-scale services that target big compute, big data, and containerized workloads.

12. Are data disks supported within scale sets?

Explanation: Yes. A scale set can define an attached data disk configuration that applies to all VMs in the set. Other options for storing data include:

- Azure files (SMB shared drives)
- OS drive
- Temp drive (local, not backed by Azure Storage)
- Azure data service (for example, Azure tables, Azure blobs)
- External data service (for example, remote database)

13. What is an Availability Set?

Explanation: An availability set is a logical grouping of VMs that allows Azure to understand how your application is built to provide redundancy and availability. It is recommended that two or more VMs are created within an availability set to provide for a highly available application and to meet the 99.95% Azure SLA. When a single VM is used with Azure Premium Storage, the Azure SLA applies for unplanned maintenance events.

Apart from this Azure Interview Questions Blog, if you want to get trained from professionals on this technology, you can opt for a structured training from edureka! Click below to know more.



Trending Courses in this category

AWS Architect Certification Training

[5 \(17250\)](#)

[44k Learners Enrolled Live Class](#)

[Best Price 16,995 19,995](#)

Similar Courses [AWS Architect Certification Training](#) [Microservices Architecture Training](#) [AWS Development Certification Training](#)

14. What are Fault Domains?

Explanation: A fault domain is a logical group of underlying hardware that share a common power source and network switch, similar to a rack within an on-premise data-centers. As you create VMs within an availability set, the Azure platform automatically distributes your VMs across these fault domains. This approach limits the impact of potential physical hardware failures, network outages, or power interruptions.

15. What are Update Domains?

Explanation: An update domain is a logical group of underlying hardware that can undergo maintenance or can be rebooted at the same time. As you create VMs within an availability set, the Azure platform automatically distributes your VMs across these update domains. This approach ensures that at least one instance of your application always remains running as the Azure platform undergoes periodic maintenance. The

order of update domains being rebooted may not proceed sequentially during planned maintenance, but only one update domain is rebooted at a time.

16. What are Network Security Groups?

Explanation: A network security group (NSG) contains a list of Access Control List (ACL) rules that allow or deny network traffic to subnets, NICs, or both. NSGs can be associated with either subnets or individual NICs connected to a subnet. When an NSG is associated with a subnet, the ACL rules apply to all the VMs in that subnet. In addition, traffic to an individual NIC can be restricted by associating an NSG directly to a NIC.

17. Do scale sets work with Azure availability sets?

Explanation: Yes. A scale set is an implicit availability set with 5 fault domains and 5 update domains. Scale sets of more than 100 VMs span multiple *placement groups*, which are equivalent to multiple availability sets. An availability set of VMs can exist in the same virtual network as a scale set of VMs. A common configuration is to put control node VMs (which often require unique configuration) in an availability set and put data nodes in the scale set.

18. What is a break-fix issue?

Explanation: Technical problems are called break-fix issue, it is an industry term which refers to “work involved in supporting a technology when it fails in the normal course of its function, which requires intervention by a support organization to be restored to working order”.

19. Why is Azure Active Directory used?

Explanation: Azure Active Directory is an Identity and Access Management system. It is used to grant access to your employees to specific products and services in your network. For example: Salesforce.com, twitter etc. Azure AD has some in-built support for applications in its gallery which can be added directly.

20. What happens when you exhaust the maximum failed attempts for authenticating yourself via Azure AD?

Explanation: We use a more sophisticated strategy to lock accounts. This is based on the IP address of the request and the passwords entered. The duration of the lockout also increases based on the likelihood that it is an attack.

21. Where can I find a list of applications that are pre-integrated with Azure AD and their capabilities?

Explanation: Azure AD has around 2600 pre-integrated applications. All pre-integrated applications support single sign-on (SSO). SSO let you use your organizational credentials to access your apps. Some of the applications also support automated provisioning and de-provisioning.

Apart from this Azure Interview Questions Blog, if you want to get trained from professionals on this technology, you can opt for a structured training from edureka! Click below to know more.

22. How can I use applications with Azure AD that I'm using on-premises?

Explanation: Azure AD gives you an easy and secure way to connect to the web applications you choose. You can access these applications in the same way you access your SaaS apps in Azure AD, no need for a VPN to change your network infrastructure.

23. What is Azure Service Fabric?

Explanation: Azure Service Fabric is a distributed systems platform that makes it easy to package, deploy, and manage scalable and reliable micro-services. Service Fabric also addresses the significant challenges in developing and managing cloud applications. Developers and administrators can avoid complex infrastructure problems and focus on implementing mission-critical, demanding workloads that are scalable, reliable, and manageable. Service Fabric represents the next-generation middleware platform for building and managing these enterprise-class, tier-1, cloud-scale applications.

24. What is a VNet?

Explanation: VNet is a representation of your own network in the cloud. It logically isolates your instances launched in the cloud, from the rest of your resources.

25. What are the differences between Subscription Administrator and Directory Administrator?

Explanation: By default, one is assigned the Subscription Administrator role when he/she signs up for Azure. A subscription admin can use either a Microsoft account or a work or school account from the directory that the Azure subscription is associated with.

This role is authorized to manage services in the Azure portal. If others need to sign in and access services by using the same subscription, you can add them as co-admins.

Azure AD has a different set of admin roles to manage the directory and identity-related features. These admins will have access to various features in the Azure portal or the Azure classic portal. The admin's role determines what they can do, like create or edit users, assign administrative roles to others, reset user passwords, manage user licenses, or manage domains.

26. Are there any scale limitations for customers using managed disks?

Explanation: Managed Disks eliminates the limits associated with storage accounts. However, the number of managed disks per subscription is limited to 2000 by default.

27. What is the difference between Service Bus Queues and Storage Queues?

Explanation: The Azure Storage Queue is simple and the developer experience is quite good. It uses the local Azure Storage Emulator and debugging is made quite easy. The tooling for Azure Storage Queues allows you to easily peek at the top 32 messages and if the messages are in XML or Json, you're able to visualize their contents directly from Visual Studio Furthermore, these queues can be purged of their contents, which is especially useful during development and QA efforts.

The Azure Service Bus Queues are evolved and surrounded by many useful mechanisms that make it enterprise worthy! They are built into the Service Bus and are able to forward messages to other Queues and Topics. They have a built-in dead-letter

queue and messages have a time to live that you control, hence messages don't automatically disappear after 7 days.

Furthermore, Azure Service Bus Queues have the ability of deleting themselves after a configurable amount of idle time. This feature is very practical when you create Queues for each user, because if a user hasn't interacted with a Queue for the past month, it automatically gets clean it up. Its also a great way to drive costs down. You shouldn't have to pay for storage that you don't need. These Queues are limited to a maximum of 80gb. Once you've reached this limit your application will start receiving exceptions.

28. What is Azure Redis Cache?

Redis is an open source (BSD licensed), in-memory data structure store, used as a database, **cache** and message broker. Azure Redis Cache is based on the popular open-source Redis cache. It gives you access to a secure, dedicated Redis cache, managed by Microsoft, and accessible from any application within Azure. It supports data structures such as strings, hashes, lists, sets, sorted sets with range queries, bitmaps, hyperloglogs and geospatial indexes with radius queries.

Apart from this Azure Interview Questions Blog, if you want to get trained from professionals on this technology, you can opt for a structured training from edureka! Click below to know more.

29. Why doesn't Azure Redis Cache have an MSDN class library reference like some of the other Azure services?

Explanation: Microsoft Azure Redis Cache is based on the popular open source Redis Cache and can be accessed by a wide variety of Redis clients for many programming

languages. Each client has its own API that makes calls to the Redis cache instance using Redis commands.

Because each client is different, there is not one centralized class reference on MSDN, and each client maintains its own reference documentation. In addition to the reference documentation, there are several tutorials showing how to get started with Azure Redis Cache using different languages and cache clients. To access these tutorials, see How to use Azure Redis Cache and click the desired language from the language switcher at the top of the article.

30. What are Redis databases?

Explanation: Redis Databases are just a logical separation of data within the same Redis instance. The cache memory is shared between all the databases and actual memory consumption of a given database depends on the keys/values stored in that database. For example, a C6 cache has 53 GB of memory. You can choose to put all 53 GB into one database or you can split it up between multiple databases.

31. Is it possible to add an existing VM to an availability set?

Explanation: No. If you want your VM to be part of an availability set, you need to create the VM within the set. [There currently no way](#) to add a VM to an availability set after it has been created.

32. What are the username requirements when creating a VM?

Explanation: Usernames can be a maximum of 20 characters in length and cannot end in a period (“.”).

The following usernames are not allowed:

administrator	admin	user	user1
test	user2	test1	user3
admin1	1	123	a
actuser	adm	admin2	aspnet
backup	console	david	guest
john	owner	root	server
sql	support	support_388945a0	sys
test2	test3	user4	user5

33. What are the password requirements when creating a VM?

Explanation: Passwords must be 12 – 123 characters in length and meet 3 out of the following 4 complexity requirements:

- Have lower characters
- Have upper characters
- Have a digit
- Have a special character (Regex match [\W_])

The following passwords are not allowed:

abc@123	P@\$\$w0rd	P@ssw0rd	P@ssword123	Pa\$\$word
pass@word1	Password!	Password1	Password22	iloveyou!

Apart from this Azure Interview Questions Blog, if you want to get trained from professionals on this technology, you can opt for a structured training from edureka! Click below to know more.



Trending Courses in this category

AWS Architect Certification Training

[5 \(17250\)](#)

[44k Learners Enrolled Live Class](#)

[Best Price 16,995 19,995](#)

Similar Courses [AWS Architect Certification Training](#) [Microservices Architecture Training](#) [AWS Development Certification Training](#)

34. How much storage can I use with a virtual machine?

Explanation: Each data disk can be up to 1 TB. The number of data disks which you can use depends on the size of the virtual machine.

Azure Managed Disks are the new and recommended disk storage offerings for use with Azure Virtual Machines for persistent storage of data. You can use multiple Managed Disks with each Virtual Machine. Managed Disks offer two types of durable storage options: Premium and Standard Managed Disks.

Azure storage accounts can also provide storage for the operating system disk and any data disks. Each disk is a .vhd file stored as a page blob.

35. How can one create a Virtual Machine in Powershell?

```
# Define a credential object
```

```
$cred = Get-Credential
```

Create a virtual machine configuration

```
$vmConfig = New-AzureRmVMConfig -VMName myVM -VMSize Standard_DS2 |  
  Set-AzureRmVMOperatingSystem -Windows -ComputerName myVM -Credential $cred |  
  Set-AzureRmVMSourceImage -PublisherName MicrosoftWindowsServer -Offer  
WindowsServer |  
  -Skus 2016-Datacenter -Version latest | Add-AzureRmVMNetworkInterface -Id  
$nic.Id
```

36. How to create a Network Security Group and a Network Security Group Rule?

Create an inbound network security group rule for port 3389

```
$nsgRuleRDP = New-AzureRmNetworkSecurityRuleConfig -Name  
myNetworkSecurityGroupRuleRDP -Protocol Tcp `  
  -Direction Inbound -Priority 1000 -SourceAddressPrefix * -SourcePortRange *  
  -DestinationAddressPrefix * `  
  -DestinationPortRange 3389 -Access Allow
```

Create an inbound network security group rule for port 80

```
$nsgRuleWeb = New-AzureRmNetworkSecurityRuleConfig -Name  
myNetworkSecurityGroupRuleWWW -Protocol Tcp `  
  -Direction Inbound -Priority 1001 -SourceAddressPrefix * -SourcePortRange *  
  -DestinationAddressPrefix * `  
  -DestinationPortRange 80 -Access Allow
```

Create a network security group

```
$nsg = New-AzureRmNetworkSecurityGroup -ResourceGroupName myResourceGroup  
-Location EastUS `  
-Name myNetworkSecurityGroup -SecurityRules $nsgRuleRDP,$nsgRuleWeb
```

37. How to create a new storage account and container using Power Shell?

```
$storageName = "st" + (Get-Random)  
New-AzureRmStorageAccount -ResourceGroupName "myResourceGroup" -AccountName  
$storageName -Location "West US" -SkuName "Standard_LRS" -Kind Storage
```

```
$accountKey = (Get-AzureRmStorageAccountKey -ResourceGroupName myResourceGroup  
-Name $storageName).Value[0]  
  
$context = New-AzureStorageContext -StorageAccountName $storageName  
-StorageAccountKey $accountKey  
  
New-AzureStorageContainer -Name "templates" -Context $context -Permission  
Container
```

38. How can one create a VM in Azure CLI?

```
az vm create --resource-group myResourceGroup --name myVM --image  
win2016datacenter --admin-username azureuser --admin-password  
myPassword12
```

Apart from this Azure Interview Questions Blog, if you want to get trained from professionals on this technology, you can opt for a structured training from edureka!

Click below to know more.

39. What are the various power states of a VM?

Power states

Power State	Description
Starting	Indicates the virtual machine is being started.
Running	Indicates that the virtual machine is running.
Stopping	Indicates that the virtual machine is being stopped.
Stopped	Indicates that the virtual machine is stopped. Note that virtual machines in the stopped state still incur compute charges.
Deallocating	Indicates that the virtual machine is being deallocated.
Deallocated	Indicates that the virtual machine is completely removed from the hypervisor but still available in the control plane. Virtual machines in the Deallocated state do not incur compute charges.
-	Indicates that the power state of the virtual machine is unknown.

40. How can you retrieve the state of a particular VM?

```
Get-AzureRmVM  
  -ResourceGroupName myResourceGroup  
  -Name myVM  
  -Status | Select @{n="Status"; e={$_.Statuses[1].Code}}
```

41. How can you stop a VM using Power Shell?

```
Stop-AzureRmVM -ResourceGroupName myResourceGroupVM -Name "myVM" -Force
```

42. Why was my client disconnected from the cache?

Explanation: The following are some common reason for a cache disconnect.

- Client-side causes

- The client application was redeployed.
 - The client application performed a scaling operation.
 - In the case of Cloud Services or Web Apps, this may be due to auto-scaling.
 - The networking layer on the client side changed.
 - Transient errors occurred in the client or in the network nodes between the client and the server.
 - The bandwidth threshold limits were reached.
 - CPU bound operations took too long to complete.
- Server-side causes
 - On the standard cache offering, the Azure Redis Cache service initiated a fail-over from the primary node to the secondary node.
 - Azure was patching the instance where the cache was deployed
 - This can be for Redis server updates or general VM maintenance.

43. What is Azure Search?

Explanation: Azure Search is a cloud search-as-a-service solution that delegates server and infrastructure management to Microsoft, leaving you with a ready-to-use service that you can populate with your data and then use to add search to your web or mobile application. Azure Search allows you to easily add a robust search experience to your applications using a simple REST API or .NET SDK without managing search infrastructure or becoming an expert in search.

44. My web app still uses an old Docker container image after I've updated the image on Docker Hub. Does Azure support continuous integration/deployment of custom containers?

Explanation: Yes, it does. For private registries, you can update the container by stopping and then re-starting your web app. Alternatively, you can also change or add a dummy application setting to force an update of your container.

Apart from this Azure Interview Questions Blog, if you want to get trained from professionals on this technology, you can opt for a structured training from edureka! Click below to know more.

45. What are the expected values for the Startup File section when I configure the runtime stack?

Explanation: For Node.Js, you specify the PM2 configuration file or your script file. For .NET Core, specify your compiled DLL name. For Ruby, you can specify the Ruby script that you want to initialize your app with.

46. How are Azure Marketplace subscriptions priced?

Explanation:

Pricing will vary based on product types. ISV software charges and Azure infrastructure costs are charged separately through your Azure subscription. Pricing models include:

BYOL Model: Bring-your-own-license. You obtain outside of the Azure Marketplace, the right to access or use the offering and are not charged Azure Marketplace fees for use of the offering in the Azure Marketplace.

Free: Free SKU. Customers are not charged Azure Marketplace fees for use of the offering.

Free Software Trial: Full-featured version of the offer that is promotionally free for a limited period of time. You will not be charged Azure Marketplace fees for use of the offering during a trial period. Upon expiration of the trial period, customers will automatically be charged based on standard rates for use of the offering.

Usage-Based: You are charged or billed based on the extent of your use of the offering. For Virtual Machines Images, you are charged an hourly Azure Marketplace fee. For Data Services, Developer services, and APIs, you are charged per unit of measurement as defined by the offering.

Monthly Fee: You are charged or billed a fixed monthly fee for a subscription to the offering (from the date of subscription start for that particular plan). The monthly fee is not prorated for mid-month cancellations or unused services.

47. What is the difference between “price,” “software price,” and “total price” in the cost structure for Virtual Machine offers in the Azure Marketplace?

Explanation: “Price” refers to the cost of the Azure Virtual Machine to run the software. “Software price” refers to the cost of the publisher software running on an Azure Virtual

Machine. “Total price” refers to the combined total cost of the Azure Virtual Machine and the publisher software running on an Azure Virtual Machine.

48. What are stateful and stateless microservices for Service Fabric?

Explanation: Service Fabric enables you to build applications that consist of microservices. Stateless microservices (such as protocol gateways and web proxies) do not maintain a mutable state outside a request and its response from the service. Azure Cloud Services worker roles are an example of a stateless service. Stateful microservices (such as user accounts, databases, devices, shopping carts, and queues) maintain a mutable, authoritative state beyond the request and its response. Today’s Internet-scale applications consist of a combination of stateless and stateful microservices.

49. What is the meaning of application partitions?

Explanation: The application partitions are a part of the Active Directory system and having said so, they are directory partitions which are replicated to domain controllers. Usually, domain controllers that are included in the process of directory partitions hold a replica of that directory partition. The attributes and values of application partitions is that you can replicate them to any specific domain controller in a forest, meaning that it could lessen replication traffic. While the domain directory partitions transfer all their data to all of the domains, the application partitions can focus on only one in the domain area. This makes application partitions redundant and more available.

50. What are special Azure Regions?

Explanation: Azure has some special regions that you may wish to use when building your applications for compliance or legal purposes. These special regions include:

- **US Gov Virginia and US Gov Iowa**
 - A physical and logical network-isolated instance of Azure for US government agencies and partners, operated by screened US persons. Includes additional compliance certifications such as [FedRAMP](#) and [DISA](#).
- **China East and China North**
 - These regions are available through a unique partnership between Microsoft and 21Vianet, whereby Microsoft does not directly maintain the datacenters.
- **Germany Central and Germany Northeast**
 - These regions are available via a data trustee model whereby customer data remains in Germany under control of T-Systems, a Deutsche Telekom company, acting as the German data trustee.

=====Azure Question from intellipat=====

1. What is Azure Cloud Service?

Cloud service can convey a multiple web application in Azure, characterizing a number of parts to disseminate handling and permit adaptable scaling of your application. A cloud service comprises of a minimum of one web pars as well as specialist parts, each with its own particular application documents and design. The fundamental favourable position of cloud service is the capacity to help more complex multi-level structures.

2. Differences between Microsoft Azure and AWS?

Features	Microsoft Azure	Amazon Web Services (AWS)
Analytics	Azure Stream Analytics	Amazon Kinesis
Backup Options	Azure Backup	Amazon Glacier
Compliance	Azure Trust Center	AWS CloudHSM

Content Delivery Network (CDN)	Azure CDN	Amazon CloudFront
Data Orchestration	Azure Data Factory	AWS Data Pipeline
Hybrid Cloud Storage	StorSimple	AWS Storage Gateway
Monitoring	Azure Operational Insights	Amazon CloudTrail
NoSQL Database Options	Azure DocumentDB	Amazon Dynamo DB

3. What are the roles implemented in Windows Azure?

There are three roles in Windows Azure.

- Web Role
- Worker Role
- Virtual Machine Role

Web Role : It gives a web solution that is front-end. This is like an ASP.NET application. While under facilitating Azure gives IIS and required services.

Worker Role : It gives a solutions of background service. It can run long activities.

Virtual Machine Role : The roles of both web and worker are executed on virtual machines. The Virtual Machine Roles gives the client the capacity to modify the virtual machine on which the web and worker roles are running.

Become Microsoft Azure Certified in 24 hrs.

GET CERTIFIED

4. What are the three principle segments of Windows Azure Platform?

Windows Azure has three principle segments in Azure: Compute, Storage and Fabric.

A. Windows Azure Compute

Windows Azure gives a code that can be managed by the hosting environment. It gives calculation benefit through parts. Windows Azure backs 3 types of roles:

- Web roles utilized for web application programming and upheld by IIS7.
- Worker roles utilized for foundation handling of web roles.
- Virtual Machine (VM) roles utilized for moving windows server applications to Windows Azure in a simple way.

B. Windows Azure Storage

It gives 4 types of storage services:

- Queues for informing between web parts and worker roles.
- Tables for storing structural data.
- BLOBS (Binary Large Objects) to store content, records or vast information.

- Windows Azure Drives (VHD) to mount a page blob. These can be transferred and downloaded by means of blobs.

C. Windows Azure AppFabric

AppFabric provides 5 services:

- Service bus
- Access
- Caching
- Integration
- Composite

5. Windows Azure AppFabric

Windows Azure Diagnostics empowers you to gather diagnostic data from an application running in Windows Azure. diagnostic data is used for capacity planning, and evaluating.

6. What is the distinction between Windows Azure Queues and Windows Azure Service Bus Queues?

Azure Queues gives a solid, diligent messaging between and within the services. it also highlights a very straight forward rest- based get/put/peek interface

Bus Queues are a part of a more extensive Windows Azure messaging framework that supports queuing.

7. What is table storage in Windows Azure?

The Windows Azure Table storage service stores a lot of organized information. Windows Azure tables are perfect for putting away organized, non-relational data

Table: A table is a collection of entities. Tables don't uphold a blueprint on elements, which implies a solitary table can contain substances that have distinctive arrangements of properties. A record can contain numerous tables

Entity: An entity is an arrangement of properties, like a database row. An entity can be upto 1MB in size.

Properties: A property is a name-value pair. Every entity can incorporate up to 252 properties to store data. Every entity likewise has 3 system properties that determine a segment key, a row key, and a timestamp.

8. What is AutoScaling in Azure?

Scaling by including extra instances is frequently referred to as scaling out. Windows Azure likewise supports scaling up by utilizing bigger role rather than more role instances. By adding and expelling role instances to your Windows Azure application while it is running, you can adjust the execution of the application against its running costs.

An autoscaling solution reduces the amount of manual work engaged in dynamically scaling an application.

9. What are the Features of Windows Azure?

Windows Azure runs and stores the information on Microsoft datacenters.

The main Features are :

1. Websites enable the designers to assemble the sites utilizing ASP.NET, PHP, etc and send these websites utilizing FTP, Git and etc
2. QL Database, formally known as Azure database makes, broadens and scales the application into the cloud utilizing Microsoft SQL Server.
3. This is Microsoft's platform as a service that supports the Multi-level applications and automated deployment.

[**Learn Microsoft Azure from Experts! Enrol Today**](#)

10. What are the differences between an public cloud and a private cloud?

Private clouds are those that are constructed solely for an individual enterprise. They enable the firm to have applications in the cloud while tending to concerns with respect to data security and control that is frequently ailing in an public cloud environment. It is otherwise called an internal cloud or enterprise cloud and dwells on the organization's intranet or hosted data center where the data is protected.

11. What is table storage in Windows Azure?

It is a NoSQL datastore which acknowledges verified calls from inside and outside the Windows Azure cloud. Windows Azure tables are perfect for putting away organized, non-relational data

Table: A table is an accumulation of elements. Tables don't implement a pattern on elements, which implies a solitary table can contain substances that have distinctive arrangements of properties. A record can contain numerous tables.

12. What is Windows Azure Portal?

Windows Azure Portal: To run an application, a designer gets to the Windows Azure portal through his Web program, by logging in with a Windows Live ID. The User at that point picks whether to create a host account for running applications, a storage account for storing data or both.

Once the designer has a host account, He can utilize a Windows Azure portal to submit applications to Windows Azure.

13. Explain Azure Fabric?

The Azure Fabric is the principle core concept. It gives a service called the Azure Fabric Controller. It is called as OS for the Azure. Since it handles/oversees:

- All roles (processing) and resources.
- Sending and activating services.
- Monitoring the health for all services.
- Releasing and allocating of resources.
- Provisioning VM, terminating etc.
- Patches gets updated for installed OS on VM in the most automated form.

14. What do you comprehend about Hybrid Cloud?

A Hybrid cloud is a blend of internal and external cloud services, a mix of a private cloud joined with the utilization of public cloud services. This kind of cloud is most appropriate when you need to keep the classified information at your vicinity (private cloud) and consume alternate services from a public cloud.

15. What is a Storage keys?

Storage keys or Access Keys are utilized as a validation mode for accessing to the storage services account to control data based on our prerequisites. In Windows Azure we have an alternative to give a Primary Access Key and a Secondary Access Key,

despite the fact that we will utilize a solitary access key to confirm our application to the storage. The primary reason to give the secondary access key is to avoid downtime to the application.

16. What is Windows Azure Traffic Manager?

It enables Users to control the distribution of user traffic of installed Azure cloud services. There are 3 distinctive load balancing strategies provided by Azure. The Manager who works on traffic apply's a routing policy to the Domain Name Service (DNS) questions on your domain names and maps the DNS courses to the apt instances of your applications.

17. What is Federation in SQL Azure?

Organization in SQL Azure is introduced for scalability. federation helps both managers and developers to scale information. It helps managers by making repartitioning and redistributing of information simpler. It enables developers in the layer of routing and sharing of information. It helps in routing without application downtime.

18. What is SQL Azure Database?

SQL Azure database is just an approach to get associated with Cloud Services where we can store our database into Cloud. Microsoft Azure is the most ideal approach to utilize PAAS where we can have different databases on a similar Account.

Microsoft SQL Azure has a similar component of SQL Server, i.e. high accessibility, versatility and security in the core.

Microsoft Azure SQL Database has an element, it makes backups automatically of each active database. Consistently a backup is taken and geo-repeated to empower the 1-hour recuperation point objective (RPO) for Geo-Restore.

19. What are the different Storage's in Windows Azure?

BLOB : BLOBS offer a component for storing a lot of content or binary data, for example, pictures, sound and visual documents. It can scale up to 200 terabytes and can be acquired by utilizing REST APIs

Table : Tables represents storage areas across machines for information that is in the form of properties on the cloud.

Line : The sole target of a Queue is to empower communication amongst Web and Worker Role instances. They help in storing messages that may accessed to by a customer.

20. What is the concept of the table in Windows Azure?

A table is one kind of Azure Storage, where you can store your information away. Blobs are put away in compartment and Entity in a table.

Following are the key concepts in a table.

- Tables allow structure data storage .
- There can be 0..n tables in a storage account.
- Table store information as an accumulation of elements .
- An element has an essential key and properties as a key-value pair.

21. What is TFS build system in Azure?

A Build is the solution of an output. In Azure projects, you get the record with a .cspkg extension that implies a Cloud Service Package is utilized for the deployment of your cloud administrations.

Build Servers – In general terms a build server goes about as the machine where you put your deployment packages.

To utilize Team Foundation Build, you should have no less than one build machine. This machine can be a physical machine or a virtual machine.

Build Controllers – Manufacture Controllers are the component in the build system that accepts the build requests from any task inside the group project. Each build controller is dedicated to a solitary team project collection. So there is a balanced relationship between a team project and a build controller.

Build Agents – Build Agents are components in the build system that accomplishes more processor-concentrated work.

22. What is the Azure App Service?

Azure App Service is a completely managed Platform as a Service (PaaS) offering for proficient developers that conveys a rich arrangement of abilities to web, mobile and integration scenarios. Mobile Apps in Azure App Service offer a very adaptable, universally accessible mobile application development platform for Enterprise Developers and System Integrators that conveys a rich set of capacities to mobile engineers.

23. What is profiling in Azure?

Profiling is only a procedure of measuring the performance analysis of an application. It is normally done to guarantee that the application is sufficiently steady and can maintain overwhelming traffic.

Visual Studio gives us different tools to do it by gathering the performance information from the application that likewise helps in the troubleshooting issues.

Once the profiling wizard is run, it sets up the execution session and collects the data of the sample

The profiling reports helps in:

- Deciding the longest running strategies inside the application.
- Measure the execution time of every strategy in the call stack.
- Assess memory allocation.

24. What is Cmdlet in Azure?

A cmdlet is a lightweight command that is utilized as a part of the Microsoft PowerShell environment. The cmdlets are summoned by the Windows PowerShell to automate the scripts which are in the command line. The Windows PowerShell runtime additionally invokes them automatically through Windows PowerShell APIs.

25. What is Windows Azure Scheduler?

Windows Azure Scheduler enable you to invoke activities –, for example, calling HTTP/S endpoints or presenting a message on a storage queue on any schedule. With Scheduler, you make jobs in the cloud that dependably call services both inside and outside of Windows Azure and execute those jobs on demand, on a routinely repeating schedule, or assign them for a future date.

26. How can you create a HDInsight Cluster in Azure?

To make an Azure HDInsight Cluster, open the Azure portal and then click on New, Data Services,then HDInsight.

Hadoop is the default and native execution of Apache Hadoop.

HBase is an Apache open-source NoSQL database based on Hadoop that gives random access and solid consistency for a lot of unstructured data.

Storm is a distributed, fault tolerant, open-source computation system that enables you to process data in real time.

27. What is Text Analytics API in Azure Machine?

Content Analytics API is a part of content examination web administrations worked with Azure Machine Learning. The API can be utilized to analyze unstructured content for tasks, like, sentiment analysis and key phrase extraction.

The API restores a numeric score between 0 and 1. Scores near 1 show positive sentiment, while scores near 0 demonstrate negative sentiment.

The upside of this API is that another new model need not be planned and prepared, the user just needs to bring the data and call the service to get the sentiment results.

28. What is Migration Assistant tool in Azure Websites?

Migration Assistant tool will examine your IIS installation and recognize which sites can be migrated to the cloud, featuring any components which can't be migrated or are unsupported on the platform.

Once broke down this tool will likewise create sites and databases provided under given Azure membership.

29. What is the distinction between Public Cloud and Private Cloud?

Public cloud is utilized as a service through Internet by the users, while a private cloud, as the name passes on is deployed within specific limits like firewall settings and is totally overseen and checked by the users dealing with it in an organization.

30. What is Azure Service Level Agreement (SLA)?

The SLA ensures that, when you send two or more role instances for each role, access to your cloud service will be maintained not less than 99.95 percent of the time. Additionally, identification and reconnection activity will be started 99.9 percent of the time when a role instance's procedure isn't running.

=====Whizlabs=====

1. What is meant by Microsoft Azure and Azure diagnostic?

Answer: This is one of the most basic Azure cloud interview questions asked very often. Microsoft Azure is a cloud computing interface that is implemented by Microsoft so as to get benefited from the cloud computing.

Azure diagnostics is an API based system that collects the data to diagnose the application which is constantly running. It tunes with the verbose monitoring by enabling roles of the cloud services.

2. What is meant by cloud computing?

Answer: This is a basic question that finds its mention in most of the Microsoft Azure interview questions eBook. Cloud Computing is the high-level abstraction procedure that focuses on the business logic. This is a service delivered via the internet that aids you with the computing services without laying much importance on the infrastructural needs just like the electric supply.

3. What is the scalability of the cloud computing?

Answer: If the above question appears as the Azure cloud interview questions your most appropriate answer will be scaling is one of the advantages that are of two types;

- **Vertical scaling**, where the configuration yields to increase the existing capacity of the machine. Just like expanding the size of the RAM from 4GB to 32GB.
- **Horizontal Scaling**, an aspect where the physical aspect is increased like putting multiple machines at work instead of replacing the existing machine.

4. What are the advantages of cloud computing?

Answer: There are several advantages of cloud computing which are as follows;

- The versatility of the system can be altered if the business requires certain modification.
- They are highly available which increases its end users to get benefited by every inch.
- The system is capable of tolerating fault of a single part and overcomes it by remaining completely functional.
- The service allows you to pay only when the system is in use. You can easily reallocate the system when you use them and get saved from the charge.
- The system has brought into light the payment for the operation rather than the investment in the machines.

5. What is meant by PaaS, SaaS, and IaaS?

Answer: These are the basic terms while dealing with the cloud services and are commonly asked in an Azure interview:

- **PaaS stands for Platform as a Service** that enables you to get a platform to deliver without directly giving authorization to the OS software.
- **SaaS stands for Software as a Service** is devoid of platform infrastructure software that can be used without direct purchase.
- **IaaS stands for Infrastructure as a Service** which enables you to get the hardware from the provider as the desired service which can be configured by the user.



6. Explain the different deployment model of the cloud?

Answer: When asked as the Azure cloud interview question you should answer it in the following way. There are various deployment models for the cloud;

- **Private Cloud Deployment Model** is one where you own the infrastructure and the cloud provider enables you to get the service exclusively.
- **Public Cloud Deployment Model** is one where you get the multi-tenant infrastructure from the service provider exclusively.

- **Hybrid Cloud Deployment Model** is a set up of intermixed public and private data in which the confidential data will be served in-house and the public directed website is there to address the public.

Try Now: [Microsoft Azure AZ-102 Free Tests](#)

7. What are the main functions of the Azure Cloud Service?

Answer: The main functions of the Azure Cloud Service are;

- It is designed to host the running application and at the same time manage the background running application.
- The application of web processing is termed as “web role” whereas the background processing is termed as the “worker role”.

8. State the purpose of the cloud configuration file?

Answer: There is a primary .csfg file available with each and every cloud service. The main purpose of this file is

- They hold the main copy of certificates.
- They have the storage of user-defined settings.
- There are a number of instances in any service project.

9. Which services are used to manage the resources in Azure?

Answer: When asked as Azure cloud interview questions, it should be answered that the Azure resource manager is the infrastructure that is involved in to manage deploys or deleting all the resources.

10. What do you mean by roles?

Answer: This is another basic yet important question asked as the Azure Interview Questions which should be answered in the following way:

Roles in cloud management are often termed to be nothing servers that are linked to managing and balancing the platform as service machine that collaborates works to attain the goal.

11. What are the different types of roles?

Answer: There are basically three different types of roles;

- **Web Role** that is used to deploy website by making use of language which is supported by the IIS platform customized to run the web application.
- **VM Role** is used to schedule the window services and task. It is done by the user through customization of the machine on which the worker's role is running.
- **Worker Role** is to execute the process that runs in the background by deploying the website.



General Azure Interview Questions and Answers

So, you have gone through all the basic questions, now let's take a step ahead and move to some general Azure interview questions. Interviewer generally asks these questions to check the knowledge of the candidate on various aspects.

12. What do you mean by a domain?

Answer: This is one of the common Azure Interview Questions that may be asked which should be answered as follows the interconnected and interlinked nodes that are often a measure undertaken by the organization is known as the domain. These relations are carried by only one point of the organization.

13. Explain the fault domain.

Answer: This is one of the common Azure interview questions which should be answered that it is a logical working domain in which the underlying hardware is sharing a common power source and switch network. This means that when VMs is created the Azure distributes the VM across the fault domain that limits the potential impact of hardware failure, power interruption or outages of the network.

14. Discuss Update Domain.

Answer: It is a logical group that undergoes maintenance by rebooting the system. The system automatically distributes the created VMs all across the updated domain that enables the user to run any one of the application while the Azure platform is undergoing maintenance. The update domain does not work sequentially to the manual rebooting.

15. Differentiate between the verbose and minimal monitoring.

Answer: Verbose monitoring collects metrics based on the performance that enables close analysis of data that fed during the processing of application whereas the minimal monitoring is a default configuration that makes the use of the performance counters gathered from the operating system of the host. Proceed to the more Azure cloud interview questions that are quite common.

Thinking Why should I go for an Azure certification? Here are the [Top 7 Reasons to go for Microsoft Azure Certification](#).

16. What do you mean by a BLOB and what are their types?

Answer: BLOB is a Binary Large Object that is composed of any size and type of file. They are mainly of two types-the page and the block blob.

17. What is meant by the block blob and page BLOB?

Answer: Blob is a block that is having a specific block ID. Each block in this block BLOB comprises of the 4MB and maximum size of this BLOB limits to 200 GB. Whereas the Page blob contains pages in which data range is determined by the offsets. The maximum limit is 1TB where a single page is of the size 1TB.

18. Differentiate between the Windows Azure bus queues and Windows Azure queues?

Answer: Windows Azure queues take care of the storage infrastructure of Windows Azure. The features include the REST based Peek/Get/Put interface that enables the users to get persistent messaging in-between the services and reliability.

Bus queues are a broader aspect of the messaging infrastructure with publishing and subscribing the integrating patterns of remote web service.

19. What is meant by the DeadLetter queue?

Answer: Messages are transferred to the DeadLetter queue in the following situation;

- When the delivery count has exceeded for a message that is on a queue.
- When the expiry date of the message has crossed and the entire expired message is held in a queue.

- When there is an evaluation exception set by default and the subscription is enabled with dead letter filter.

20. Explain the Window Azure platform?

Answer: It is collectively a PaaS developed by the Microsoft programmed to run a deployment vehicle and a runtime for the data center to host the cloud computing. Let's proceed to the next general Azure Interview Questions that may be asked.

21. How is the price of the Azure subscription placed?

Answer: The prices vary in accordance with the product type. The various pricing models are;

- **The free model** where our customers can avail the system exempted from the market fee.
- **The BYOL scheme** where the Azure is fetched with a fee that is not encircled around the limits of market price.
- **The trial of the free software** where the client gets the full version with an advanced feature for a limited period of time. With the expiration of the introductory period, you will be charged standard rates.
- **Usage Based Fee** that is billed in accordance with the service that you have taken. Like if you are subscribed for the virtual image then hourly fees may be charged.
- **Monthly bills** are activated if you sign a particular plan. The fee is not allocated for cancellation or mid-month unused service.

Azure Interview Questions and Answers for Experienced

If you are a cloud professional with some considerable experience on the Azure platform, you may come across some advanced Azure interview questions. So, let's get prepared with some frequently asked Azure interview questions for experienced.

22. What are the sizes of the Azure VM?

Answer: It is another basic question that finds its place in the series of top Microsoft Azure interview questions. The Windows Azure is destined to balance a variety of sizes. Most of the VM sizes are:

- The extra large computer has 8*1.6 GHz of Instance size, with instance storage of 2040 GB, CPU memory of 14 GB. The I/O performance is high.
- The large computer has 4*1.6 GHz of Instance size, with instance storage of 1000 GB, CPU memory of 7 GB. The I/O performance is high.
- The medium computer has 2*1.6 GHz of Instance size, with instance storage of 490 GB, CPU memory of 3.5 GB. The I/O performance is high.
- Small computer has 1.6 GHz of Instance size, with instance storage of 225 GB, CPU memory of 1.75 GB. The I/O performance is moderate.
- The extra small computer has 1.0 GHz of Instance size of 20 GB, with instance storage of 20 GB, CPU memory of 768MB. The I/O performance is low.

Computer Size	Instance Size	Instance Storage	CPU Memory	I/O Performance
Extra Large Computer	8*1.6 GHz	2040 GB	14 GB	High
Large Computer	4*1.6 GHz	1000 GB	7 GB	High
Medium Computer	2*1.6 GHz	490 GB	3.5 GB	High
Small Computer	1.6 GHz	225 GB	1.75 GB	Moderate
Extra Small Computer	1.0 GHz	20 GB	768 MB	Low

23. What is meant by table storage?

Answer: It is an interface that is capable of storing bulk amount of structured but non-relational data. It is a service of the NoSQL data store that takes authenticated calls from either outside or inside the Azure cloud. The table is a collaborated entity that in turn is a set of properties which is a paired up of name and value.

24. Differentiate between the repository and the powerhouse server?

Answer: Repository servers are those which are in lieu of the integrity, consistency, and uniformity whereas the powerhouse server governs the integration of different aspects of the database repository.

25. What is meant by the enterprise warehousing?

Answer: It is the phenomenon where the data is developed by the organization having access at a single point throughout the globe. The warehousing enables the server to get linked to a single point with the assistance of periodic handling.

Azure is the top second cloud player among the top three (AWS, Azure, and Google Cloud). Let's have a comparison between three i.e. [AWS vs Azure vs Google](#)

26. What do you mean by lookup transformation?

Answer: Lookup transformation aids to determine source qualifier. It can be active or passive lookup transformation. The process is yield to get the access the relevant information or the data.

27. What is meant by the connected lookups?

Answer: In the unconnected lookup, the input is directly taken from the transformation that takes part in the flow of data. The connected data lookup is built as both a static and

dynamic cache. This can be oriented via multiple ports that can give the output. It is defined by user defined entity.

28. What is meant by the unconnected lookups?

Answer: The answer to this Microsoft Azure interview question should be that in the unconnected lookups the input is taken by the LKP operation. Nature is dynamic having a single output port. User-defined values are disregarded in the unconnected lookups.

29. What is meant by the command task?

Answer: Command task is an operational window that sets off the flow of one or multiple command shell while the system is still running.

30. What are the PowerCenter commands that can be used in Informatica?

Answer: The languages for transformation that can be provided with some basic comment signifiers that are;

- Slashes (//)
- Dashes (-)

If these common signifiers are absent in the line the power integration service will ignore the text.

31. What is the difference between copy and shortcut?

Answer: The answer to this Microsoft Azure interview question will be copied is transferring an object from one to another folder that takes double space. There is no change that corresponds to the original object.

The shortcut is a dynamic link of the object that saves the space reflecting changes in the original object.

32. What do you mean by a service fabric in Azure?

Answer: This is one of the advanced Microsoft azure interview questions that should be answered by stating Service fabric is a middleware platform of the following generation that gives the more scalable outcome. It renders with the more managed and reliable enterprise.

33. What are the benefits of the traffic manager in Windows Azure?

Answer: Traffic manager is allocated to control the distribution of the user to deploy the cloud service. The benefit of the traffic manager constitutes;

- It makes the application to be available worldwide through automated traffic control machinery.
- The traffic managing service contributes to high performance by loading the page faster and convenient usage.
- There is no lag of time to maintain or upgrade the existing system. The system keeps running in the back while the system takes time for up gradation.
- The configuration is made easy through the Azure portal.

SQL Azure Interview Questions and Answers

While going for an Azure interview, it is common to come across database related questions i.e. SQL-based. So, let's get prepared for that too. Just go through the following SQL Azure interview questions.

34. Discuss the different database types in SQL Azure?

Answer: This is one of the commonly asked SQL Azure interview questions that must be answered by stating that there are two major type of database in SQL Azure;

- **Web Edition** – It is having a limit of 5GB SQL that is related to the relational database. The basic advantage is that they can be self-maintained, tolerant to a fault and highly available.
- **Business-based Edition** – they support 50GB of T-SQL that is self-managed, tolerant to a fault and highly available. They are suited for the custom web applications or ISV application.

35. How is Azure Resource Manager beneficial over the classic services?

Answer: The benefits of the Azure Resource Manager that overshadows the benefit of the classic services are;

- The resources need not be managed, deployed or monitored one at a time. They are chain deployment activities throughout the lifecycle without the need for individual data handling.
- The data is also deployed at a consistent pace with the ARM service. It enables the user to use a declarative template that indicates the deployment.
- Since the role-based control is present in the management platform that provides you with the access to the resources that leads you to control.
- You can mark dependencies between the resources that enable you to get the correct order of deployment.
- The resources may be tagged and organized logically so that it is convenient to follow up the billing of your company.

Also Read:

[Microsoft Azure Certification AZ-100 Preparation Guide](#)

[Microsoft Azure Certification AZ-101 Preparation Guide](#)

[Microsoft Azure Certification AZ-102 Preparation Guide](#)

36. Enlist the monitoring features that are present in the SharePoint 2010?

Answer: This is one of the SQL Azure interview questions that should be answered by stating that the SharePoint 2010 is a diagnostic logging that takes into the data that are a direct indication of the state of the system. Sometimes it also specifies some timer tasks that are performed to monitor the collected information. The features include;

- It collects the data on event log, timer service, and performance counter.
- They are involved with data that are search usage.
- They provide matrices that are yielding a collection of sites.

37. State the class that can be used to retrieve data?

Answer: The class that can be used to get data from more than one list is the SPSite Data Query that is able to query data that are present in more than one list all across the diverse web collections. It is featured to aggregate data without the external links and only through the SharePoint.

38. Compare the STS and SPS and state its important features?

Answer: **SPS** is the SharePoint Portal Service which manages the documents and has a search engine more efficient in penetrating the multiple sources of content.

STS stands for the SharePoint team management. As the name suggests they are better for the document management for a large organization and has a moderate search engine.

39. Explain MOSS?

Answer: The answer to this SQL Microsoft Azure interview question for experienced will be MOSS is abbreviated for Microsoft Office SharePoint Server that constitutes a complete version of the portal platform that yields the user to manage, share and even create the document.

40. What is the difference between a library and a list?

Answer: The library is an interface that yields one to manage and store a document that can be created using Excel, PowerPoint or Word. The list is the representation of the item in a tabulated format using column and rows. Lists are not used to make a document but can attach documents.

41. What do you mean by SAS?

Answer: This is one of the common SQL Azure interview questions that should be answered by stating that SAS is an abbreviation for Statistical analytical System which is a software suite performing analysis of multiple variables. It is linked to the predictive analysis, data handling, advanced analytics or corporative intelligence. It produces a smooth interface that offers graphical and clicks based solution. It is user-friendly for the technical or the non-technical with advanced features.

42. State some features of SAS?

Answer: There is some interlinked feature of SAS which includes;

- It provides the user access to manage the resources just like that of the DBMS.

- It is leading with the leading analytical to carry out different business services and products.
- It enables easy visualization and interpretation with the use of graphs and its breaks the complex panels into simple plots.
- It is very efficient in delivering the business analysis which leads to manufacturing the products that can be distributed worldwide.

43. Describe the common architecture of SharePoint 2010?

Answer: There are three main architectural design of the SharePoint 2010 which includes;

- The enterprise farm which is uncommon as it completely is dedicated to the service and aids via the automated management with the feasible isolation of data.
- There is a single farm that is associated with multiple services whose potential advantages are management via individual service application which enables with a more complex targeting of sites to a particular application of service.
- Lastly, the single farms employed in single service are very common and at the same time easy to deploy. The application service is simple to be allocated with fullest resource utilization and management.

44. Describe the log analytics?

Answer: This question can be asked among the SQL Azure interview questions. The operational management service of the Log Analytics provides the entire requirement that runs the particular service. It manifests automation, security, log analytics and availability at a particular dashboard. It generates Power data source that enables the user to get the visuals of the raw data. It is introduced in three different tiers of prices that include free,

premium and standard. You enjoy the convenience of searching the data at a single dashboard and export the results.

Advanced Azure Interview Questions and Answers

In the series of Azure interview questions, here we bring some advanced Azure interview questions for you with the detailed answers. Preparing through these questions will prepare you to crack the tough to tough interview.

43. State what will you do in case of a drive failure?

Answer: This is one of the another Microsoft Azure interview questions for experienced that should be answered in the following manner. When there is an instance that the drive has failed the following step should be performed;

- The first is that the drive should be not mounted enabling the object storage to function without fail.
- The second scenario is replacing the drive in which the desired step will be remounting, formatting the drive.

44. State what should be done in case of a service failure?

Answer: In the case of the service failure the following steps need to be performed;

- Ensure that the object s services have stopped running which will enable the object services to withstand the failure.
- To complete the last lap of work pending you can allow the system object storage remain functional by making the machined to again get back to online. Once they are started, the online replication will instantly activate with the missing updated files.

- If you are incapable to replace drive then do remove the drive and let the system stay not mounted.

45. Give a clear overview of API in Azure?

Answer: The Test Analytics in API is a web service that is built with the Azure learning. It is an effective tool to analyze the unstructured data like the extraction of the key phrase. It runs with the binomial scoring unit that is either 0 or 1 where 1 corresponds to a positive and 0 corresponds to a negative viewpoint. The advantage is that it does not need any assistance with designing and training which imply that the data is in the hands of directly the user. Proceed to find more Microsoft Azure interview questions for experienced.

Microsoft Azure certifications have their own recognition in the Cloud industry. Let's find out their place among the [Top 5 Cloud Certifications!](#)

46. Explain how does a character analytics API function?

Answer: The working of the character analytics of API does not account for the characterization of words like good or bad. It uses the advanced feature through which the processing of natural language is mediated.

47. Differentiate between the PROC SUMMARY and PROC MEANS?

Answer: This is one of the stars marked questions found in the list of top Microsoft Azure interview questions and answers pdf.

PROC MEANS refers to the subgroup statist created in the persistence of the BY statement that will be involved. The data here is sorted beforehand with the assistance of BY variables.

PROC SUMMARY is the aid of statistics giving all varieties of information running simultaneously and is produced for every subgroup automatically. The information in the outlet is not created.

50. If the client gets disconnected from cache with the services state the probable cause?

Answer: If the client gets disconnected the causal factor can be distributed into two categories;

The cause on the operator side;

- There might be a failure in the transfer of the standard cache from one node to the other.
- While the service was processing and dispatching the cache got deployed.
- There was a server update or an automated VM maintenance.

The fault on the client side;

- The application of the client accidentally got redeployed.
- The application on the client side got auto-scaling.
- The layer of the network on the client side altered.
- There was a transient error on the network node.
- The bound operation took more time.
- The upper limit of the bandwidth was reached.

=====interview question from mindmajix=====

Q) Differences between Microsoft Azure and Amazon ? AWS

Microsoft Azure Vs Amazon		
Features	Microsoft Azure	Amazon Web Services (AWS)
Analytics	Azure Stream Analytics	Amazon Kinesis
Backup Options	Azure Backup	Amazon Glacier
Compliance	Azure Trust Center	AWS CloudHSM
Content Delivery Network (CDN)	Azure CDN	Amazon CloudFront
Data Orchestration	Azure Data Factory	AWS Data Pipeline
Hybrid Cloud Storage	StorSimple	AWS Storage Gateway
Monitoring	Azure Operational Insights	Amazon CloudTrail
NoSQL Database Options	Azure DocumentDB	Amazon Dynamo DB
Explore more about differences in features @ Azure Vs Aws Feature Comparison		

Related Article : [Azure Vs AWS Vs Google Cloud Free Tier](#)

Q1) What is Azure Cloud Service?

By creating a cloud service, you can deploy a multi-tier web application in Azure, defining multiple roles to distribute processing and allow flexible scaling of your application. A cloud service consists of one or more web roles and/or worker roles, each with its own application files and configuration. Azure Websites and Virtual Machines also enable web applications on Azure. The main advantage of cloud services is the ability to support more complex multi-tier architectures

Q2) What is a cloud service role?

A cloud service role is comprised of application files and a configuration. A cloud service can have two types of roles.

Q3) What is link a resource?

To show your cloud service's dependencies on other resources, such as an Azure SQL Database instance, you can "link" the resource to the cloud service. In the Preview Management Portal, you can view linked resources on the Linked Resources page, view their status on the dashboard, and scale a linked SQL Database instance along with the service roles on the Scale page. Linking a resource in this sense does not connect the resource to the application; you must configure the connections in the application code.

Q4) What is scale a cloud service?

A cloud service is scaled out by increasing the number of role instances (virtual machines) deployed for a role. A cloud service is scaled in by decreasing role instances. In the Preview Management Portal, you can also scale a linked SQL Database instance, by changing the SQL Database edition and the maximum database size, when you scale your service roles.

Q5) What is a web role?

A web role provides a dedicated Internet Information Services (IIS) web-server used for hosting front-end web applications.

Q6) What is a worker role?

Applications hosted within worker roles can run asynchronous, long-running or perpetual tasks independent of user interaction or input.

Q7) What is a role instance?

A role instance is a virtual machine on which the application code and role configuration run. A role can have multiple instances, defined in the service configuration file.

Q8) What is a guest operating system?

The guest operating system for a cloud service is the operating system installed on the role instances (virtual machines) on which your application code runs.

Q9) What is a cloud service component?

Three components are required in order to deploy an application as a cloud service in Azure:

Q10) What is deployment environments?

Azure offers two deployment environments for cloud services: a staging environment in which you can test your deployment before you promote it to the production environment. The two environments are distinguished only by the virtual IP addresses (VIPs) by which the cloud service is accessed. In the staging environment, the cloud service's globally unique identifier (GUID) identifies it in URLs (GUID.cloudapp.net). In the production environment, the URL is based on the friendlier DNS prefix assigned to the cloud service (for example, myservice.cloudapp.net).

Q11) What is swap deployments?

To promote a deployment in the Azure staging environment to the production environment, you can “swap” the deployments by switching the VIPs by which the two deployments are accessed. After

the deployment, the DNS name for the cloud service points to the deployment that had been in the staging environment.

Q12) What is minimal vs. verbose monitoring?

Minimal monitoring, which is configured by default for a cloud service, uses performance counters gathered from the host operating systems for role instances (virtual machines). Verbose monitoring gathers additional metrics based on performance data within the role instances to enable closer analysis of issues that occur during application processing. For more information

Q13) What is a service definition file?

The cloud service definition file (.csdef) defines the service model, including the number of roles.

Q14) What is a service configuration file?

The cloud service configuration file (.cscfg) provides configuration settings for the cloud service and individual roles, including the number of role instances.

Q15) What is a service package?

The service package (.cspkg) contains the application code and the service definition file.

Q16) What is a cloud service deployment?

A cloud service deployment is an instance of a cloud service deployed to the Azure staging or production environment. You can maintain deployments in both staging and production.

Q17) What is Azure Diagnostics?

Azure Diagnostics is the API that enables you to collect diagnostic data from applications running in Azure. Azure Diagnostics must be enabled for cloud service roles in order for verbose monitoring to be turned on.

Q18) What is Azure Service Level Agreement (SLA)?

The Azure Compute SLA guarantees that, when you deploy two or more role instances for every role, access to your cloud service will be maintained at least 99.95 percent of the time. Also, detection and corrective action will be initiated 99.9 percent of the time when a role instance's process is not running.

Q19) What is Cloud Computing?

Cloud computing is the use of computing resources (hardware and software) that are delivered as a service over a network (typically the Internet).

Q20) What are the Service Model in Cloud Computing?

Cloud computing providers offer their services according to three fundamental models: Infrastructure as a service (IaaS), platform as a service (PaaS), and software as a service (SaaS) where IaaS is the most basic and each higher model abstracts from the details of the lower models.

Examples of IaaS include: Amazon CloudFormation (and underlying services such as Amazon EC2), Rackspace Cloud, Terremark, Windows Azure Virtual Machines, Google Compute Engine, and Joyent.

Examples of PaaS include: Amazon Elastic Beanstalk, Cloud Foundry, Heroku, Force.com, EngineYard, Mendix, Google App Engine, Windows Azure Compute and OrangeScape.

Examples of SaaS include: Google Apps, Microsoft Office 365, and Onlive. Source from :
HTTP://EN.WIKIPEDIA.ORG/WIKI/CLOUD_COMPUTING

Q21) How many types of deployment models are used in cloud?

There are 4 types of deployment models used in cloud:

Public cloud

Private cloud

Community cloud

Hybrid cloud

Q22) What is Windows Azure Platform?

A collective name of Microsoft's Platform as a Service (PaaS) offering which provides a programming platform, a deployment vehicle, and a runtime environment of cloud computing hosted in Microsoft data centers.

Q23) What are the roles available in Windows Azure?

All three roles (web, worker, VM) are essentially Windows Server 2008. Web and Worker roles are nearly identical: With Web and Worker roles, the OS and related patches are taken care for you; you build your app's components without having to manage a VM

Q24) What is the difference between Windows Azure Platform and Windows Azure?

The former is Microsoft's PaaS offering including Windows Azure, SQL Azure, and AppFabric; while the latter is part of the offering and the Microsoft's cloud OS.

Q25) What are the three main components of Windows Azure Platform?

Compute

Storage

AppFabric

Q26) What is Windows Azure compute emulator?

The compute emulator is a local emulator of Windows Azure that you can use to build and test your application before deploying it to Windows Azure.

Q27) What is fabric?

In the Windows Azure cloud fabric is nothing but a combination of many virtualized instances which run the client application

Q28) How many instances of a Role should be deployed to satisfy Azure SLA (service level agreement)? And what's the benefit of Azure SLA?

TWO. And if we do so, the role would have external connectivity at least 99.95% of the time.

Q29) What are the options to manage session state in Windows Azure?

- Windows Azure Caching
- SQL Azure
- Azure Table

Q30) What is cspack?

It is a command-line tool that generates a service package file (.cspkg) and prepares an application for deployment, either to Windows Azure or to the compute emulator.

Q31) What is csrun?

It is a command-line tool that deploys a packaged application to the Windows Azure compute emulator and manages the running service.

Q32) What is guest OS?

It is the operating system that runs on the virtual machine that hosts an instance of a role.

Q33) How to programmatically scale out Azure Worker Role instances?

Using AutoScaling Application Block

Q34) What is the difference between Public Cloud and Private Cloud?

Public cloud is used as a service via Internet by the users, whereas a private cloud, as the name conveys is deployed within certain boundaries like firewall settings and is completely managed and monitored by the users working on it in an organization.

Q35) How to design applications to handle connection failure in windows Azure?

The Transient Fault Handling Application Block supports various standard ways of generating the retry delay time interval, including fixed interval, incremental interval (the interval increases by a standard amount), and exponential back-off (the interval doubles with some random variation).

```
static    RetryPolicy    policy    =    new    RetryPolicy(5,    TimeSpan.FromSeconds(2),
TimeSpan.FromSeconds(2)); policy.ExecuteAction(() => { try { string federationCmdText = @"USE
FEDERATION Customer_Federation(ShardId = " + shardId + ") WITH RESET, FILTERING=ON";
customerEntity.Connection.Open(); customerEntity.ExecuteStoreCommand(federationCmdText); }
catch (Exception e) { customerEntity.Connection.Close(); SqlConnection.ClearAllPools(); } });
```

Q36) What is windows Azure Diagnostics?

Windows Azure Diagnostics enables you to collect diagnostic data from an application running in Windows Azure. You can use diagnostic data for debugging and troubleshooting, measuring performance, monitoring resource usage, traffic analysis and capacity planning, and auditing.

<HTTP://WWW.WINDOWSAZURE.COM/EN-US/DEVELOP/NET/COMMON-TASKS/DIAGNOSTICS/>

Q37) What is Blob?

BLOB stands for Binary Large Object. Blob is file of any type and size.

The Azure Blob Storage offers two types of blobs –

Block Blob

Page Blob

URL format: Blobs are addressable using the following URL format:

Q38) What is the difference between Block Blob vs Page Blob?

Block blobs are comprised of blocks, each of which is identified by a block ID.

You create or modify a block blob by uploading a set of blocks and committing them by their block IDs.

If you are uploading a block blob that is no more than 64 MB in size, you can also upload it in its entirety with a single Put Blob operation. -Each block can be a maximum of 4 MB in size. The maximum size for a block blob in version 2009-09-19 is 200 GB, or up to 50,000 blocks.

Page blobs are a collection of pages. A page is a range of data that is identified by its offset from the start of the blob. To create a page blob, you initialize the page blob by calling Put Blob and specifying its maximum size.

-The maximum size for a page blob is 1 TB. A page written to a page blob may be up to 1 TB in size.

what to use block blobs for: streaming video. “The application must provide random read/write access” which is supported by Page Blobs

Q39) What is the difference between Windows Azure Queues and Windows Azure Service Bus Queues?

Windows Azure supports two types of queue mechanisms: Windows Azure Queues and Service Bus Queues.

Windows Azure Queues, which are part of the Windows Azure storage infrastructure, feature a simple REST-based Get/Put/Peek interface, providing reliable, persistent messaging within and between services.

Service Bus Queues are part of a broader Windows Azure messaging infrastructure dead-lettering queuing as well as publish/subscribe, Web service remoting, and integration patterns.

<HTTP://WCFPRO.WORDPRESS.COM/2010/12/06/COMMUNICATION-IN-WINDOWS-AZURE/>

<HTTP://MSDN.MICROSOFT.COM/EN-US/LIBRARY/WINDOWSAZURE/HH767287.ASPX>

Q40) What is DeadLetter queue?

Messages are placed on the dead-letter sub-queue by the messaging system in the following scenarios.

When a message expires and dead-lettering for expired messages is set to true in a queue or subscription.

When the max delivery count for a message is exceeded on a queue or subscription.

When a filter evaluation exception occurs in a subscription and dead-lettering is enabled on filter evaluation exceptions.

Q41) What are instance sizes of Azure?

Windows Azure will handle the load balancing for all of the instances that are created. The VM sizes are as follows:

Compute Instance Size CPU Memory Instance Storage I/O Performance

Extra Small 1.0 Ghz 768 MB 20 GB Low

Small 1.6 GHz 1.75 GB 225 GB Moderate

Medium 2 x 1.6 GHz 3.5 GB 490 GB High

Large 4 x 1.6 GHz 7 GB 1,000 GB High

Extra large 8 x 1.6 GHz 14 GB 2,040 GB High

Related Page: [Azure Active Directory](#)

Q42) What is table storage in Windows Azure?

The Windows Azure Table storage service stores large amounts of structured data.

The service is a NoSQL datastore which accepts authenticated calls from inside and outside the Windows Azure cloud.

Windows Azure tables are ideal for storing structured, non-relational data

Table: A table is a collection of entities. Tables don't enforce a schema on entities, which means a single table can contain entities that have different sets of properties. An account can contain many tables

Entity: An entity is a set of properties, similar to a database row. An entity can be up to 1MB in size.

Properties: A property is a name-value pair. Each entity can include up to 252 properties to store data. Each entity also has 3 system properties that specify a partition key, a row key, and a timestamp.

Entities with the same partition key can be queried more quickly, and inserted/updated in atomic operations. An entity's row key is its unique identifier within a partition.

<HTTP://MSDN.MICROSOFT.COM/EN-US/LIBRARY/WINDOWSAZURE/HH508997.ASPX>

=====Azure online interview questions =====

#1 Question

What is Microsoft Azure?

The companies, which provide the cloud services to businesses, are called Cloud Providers, and one of them is Microsoft Azure. Microsoft Azure is an expanding cloud computing service created by Microsoft for testing, building, managing, and deploying applications and services via a global network of data centers managed by Microsoft. It is basically and widely used for accessing Microsoft's infrastructure for the cloud by businesses.

0 Answer Suggest an edit

#2 Question

What is the Windows Azure portal?

The developers, who have a hosting account, can use a Windows Azure portal to submit applications to Windows Azure. A developer can easily access the Windows Azure portal through the Web browser, by signing in with a Windows Live ID to run an application.

0 Answer Suggest an edit

#3 Question

What are the roles, which are implemented in Windows Azure?

Roles are simply servers in layman terms. They are managed and load balanced platforms like service virtual machines, which work together to achieve a common goal. There are three roles, which are implemented in Windows Azure:

Web Role: It gives a web solution that is completely front-end. This is similar to an ASP.NET application. When it is enabled, Azure gives IIS and required services.

Worker Role: It gives us solutions to all background service. It can easily run long activities as well.

Virtual Machine Role: The virtual machines execute the roles of both, web and worker. The Virtual Machine Roles gives the consumer the ability to modify the virtual machine on which the web and worker roles are running.

0 Answer Suggest an edit

#4 Question

What is the distinction between Windows Azure Service Bus Queues and Windows Azure Queues?

Azure Queues provides the user with a solid, diligent messaging between and within the services. It also highlights quite a straightforward rest-based get/peek/put interface. The Azure Storage Queues make use of the local Azure Storage Emulator and debugging is made quite easy. The tooling for Azure Storage Queues allows you to easily peek at the top 32 messages and if the messages are in XML or JSON, one can visualize their contents directly from Visual Studio. Moreover, these queues can be purged of their contents, which is especially useful during development and QA efforts.

On the other hand, Bus Queues are part of a more far-reaching Windows Azure messaging framework, which supports queuing. As a rule, Azure Service Bus Queues can delete themselves after a configurable amount of idle time. This feature is very practical when you create Queues for each user because if a user hasn't interacted with a Queue for the past month, it automatically gets clean it up. It is also a great way to reduce costs. However, these Queues are limited to a maximum of 80 GB. Once the user has reached this limit, his or her application will start receiving exceptions.

0 Answer Suggest an edit

#5 Question

What is virtual machine scale sets in Windows Azure?

Virtual machine scale sets are Azure compute resources, which can be used to deploy as well as manage a set of identical VMs. When all the VMs are configured in the same way, scale sets are automatically designed to support true auto-scale, and no pre-provisioning or rearrangement of VMs is required. That is why it is easier to build large-scale services, which target big compute, big data, and containerized workloads.

0 Answer Suggest an edit

#6 Question

What is the Azure Fabric?

The Azure fabric is the main core concept over here. It provides a service called the Azure

Fabric Controller. It is called an operating system for the Azure. Because it handles or manages the following:

- All roles (computing) and resources.
 - Deployment and activating services.
 - Health monitoring for all services.
 - Allocating, releasing of resources.
 - Provisioning VM, terminating, etc.
 - Updating patches for installed OS on Virtual Machine automatically.
- In this case, it is generally better to have two instances of roles, and there is no need for the customer to worry about software updates for the user.

0 Answer Suggest an edit

#7 Question

What are the three main components of the Windows Azure Platform?

Windows Azure provides platform and infrastructure by providing accessible and cost-effective computing, storage, and networking resources on demand.

The Windows Azure has three main components in Azure, namely compute, storage and fabric.

Windows Azure Compute

Windows Azure provides a hosting environment for managed code. It provides a computation service through roles. Windows Azure supports three types of roles:

Web roles used for web application programming and supported by IIS7.

Worker roles are also used for background processing of web roles.

Virtual Machine (VM) roles are generally used for migrating windows server applications to Windows Azure in an easy way.

Windows Azure Storage

Windows Azure provides storage in the cloud. It provides four different types of storage services:

Queues for messaging between web roles and worker roles.

Tables for storing structural data.

BLOBs (Binary Large Objects) to store text, files or large data.

Windows Azure Drives (VHD) to mount a page blob. They can easily be downloaded and uploaded via blobs.

Windows Azure AppFabric

AppFabric provides infrastructure services for developing, deploying and managing Windows Azure application. It provides five services:

Service bus

Access

Caching

Integration

Composite

0 Answer Suggest an edit

#8 Question

What do you understand by Hybrid Cloud? Explain its advantages.

A hybrid cloud is a mixture of internal and external cloud services, a combination of a private cloud combined with the use of public cloud services. This type of cloud is most suitable when you want to keep the confidential data in your premise (private cloud) and consume the other services from a public cloud.

Advantages of Hybrid Cloud:

Scalability: Usually the Private Cloud services will have a lesser scalability due to its security, cost and compliance whereas the Public Cloud has a high scalability and moving non-sensitive data from the private to the public will free up resources in the data centers in the Private Cloud and that increases a very high scalability for a Hybrid Cloud.

Cost-effectiveness: Similarly the Public Cloud is very cost effective rather than Private Cloud,

and here the Hybrid Cloud provides cost effectiveness with the data and other sensitive operations secured.

Security: Since there is a Private Cloud used; the data and sensitive operations are secured highly in the Hybrid Cloud.

Flexibility: We can easily move out of the non-sensitive data and manage large scalability using a Public Cloud service along with the Private Cloud. So with the availability of large scalability using Public Cloud and security using Private Cloud an enterprise has a vast opportunity in developing for new needs.

0 Answer Suggest an edit

#9 Question

Explain what is Diagnostics in Windows Azure?

Windows Azure Diagnostics provides facility to store diagnostics data. Some diagnostics data is stored in a table, while some are stored in a blob. For collecting the data on diagnostics, the user must initialize the Windows Azure diagnostic monitor. The Windows Azure diagnostic monitor runs in Windows Azure as well as in the computer's emulator and collects diagnostic data for a role instance.

0 Answer Suggest an edit

#10 Question

What is the method for creating a Queue in storage account?

The queue is a one-type of Azure Storage, where a user can store your data as storage. Blobs are stored in a container, Entity in table and Message in Queue.

Listed below are the key concepts in the queue.

FIFO implementation

Messages are added to end of the Queue and processed from the front

Queues provides a good way of the Front end and Back end decoupling

0 Answer Suggest an edit

#11 Question

What are Storage keys?

Storage keys, which are also known as Access Keys, are used as an authentication mode for accessing the storage services account to manipulate information based on our requirements.

In Windows Azure, the users have an option to provide a Primary Access Key and a Secondary Access Key, even though it is likely that they will use a single access key to authenticate their application to the storage.

0 Answer Suggest an edit

#12 Question

Why do we need storage keys?

If a user needs to change the application access key by regenerating the access key, it takes much time to take effect; this provides a downtime. To avoid such types of situations, a secondary access key is provided so that if the primary needs to be changed or regenerated, we can map the secondary temporarily to the storage and regenerate the primary.

0 Answer Suggest an edit

#13 Question

Explain the concept of the table in Windows Azure.

A table is a type of Azure Storage, where one can store the data as memory storage. Blobs are stored in container and Entity in the table.

The key concepts in the table are explained below:

Tables allow structure data storage

There can be 0..n tables in a storage account

Table store data as a collection of entities

Entity have a primary key and properties as a key value pair

0 Answer Suggest an edit

#14 Question

Explain what is Federation in SQL Azure?

Federation is introduced in SQL Azure for scalability. It helps administrators by making repartitioning and redistributing of data easier and thus, helps with scaling data. It helps developers in the routing layer and the sharding of data. It helps in routing without application downtime. Federation does basic scaling of objects in a SQL Azure Database. Federations are

the partitioned data. There can be multiple Federations within a database. Moreover, each Federation represents a different distribution scheme. We create a Federation with a different distribution scheme and requirement. Student and Grade's tables of a School Database may have a different distribution requirement, so they are put into different Federations.

0 Answer Suggest an edit

#15 Question

What do you know about the SQL Azure firewall rules?

The firewall checks access to the originating IPs from which a user may try to access the database. To configure the firewall, we need to configure a range of acceptable IP addresses upon which we try to connect to the SQL Azure server using the Management Portal or with the SQL Server Management Studio. All access to SQL Azure is blocked by a firewall.

By default Database created in SQL Azure is blocked by the firewall for maximized security. SQL Azure firewall rules are provided to protect the data and to prevent access restrictions to the SQL Azure database.

0 Answer Suggest an edit

#16 Question

What is Windows Azure Traffic Manager?

The Traffic Manager allows users to control the distribution of user traffic of deployed Azure cloud services, Azure websites or any other endpoint. In this, the distribution of traffic includes Azure cloud services, Azure websites, and other endpoints. There are three different load-balancing methods provided by Azure. The Traffic Manager applies an intelligent routing policy engine to the Domain Name Service (DNS) queries on the domain names and then maps the DNS routes to the apt instances of the required applications.

0 Answer Suggest an edit

#17 Question

List the benefits of Traffic Manager.

The Traffic Manager comes with many benefits for the user:

Increase Performance: Can increase the performance of your application that includes faster

page loading and better user experience. This applies to the serving of users with the hosted service closest to them.

High Availability: You can use the Traffic Manager to improve application availability by enabling automatic customer traffic fail-over scenarios in the event of issues with one of your application instances.

No Downtime Required for Upgrade / Maintenance: Once you have configured the Traffic Manager, you don't need downtime for application maintenance, patch purgation or complete new package deployment.

Quick Setup: It's very easy to configure Azure Traffic Manager on Windows Azure portal. If you have already hosted your application on Windows Azure (a cloud service, Azure website), you can easily configure this Traffic Manager with a simple procedure (setting routing policy).

0 Answer Suggest an edit

#18 Question

What is an Availability Set?

An availability set is a logical grouping of VMs that allows Azure to understand how the application for a user is built to provide redundancy and availability. It is recommended that two or more VMs be created within an availability set to provide for a highly available application and to meet the 99.95% Azure SLA. When a single VM is used with Azure Premium Storage, the Azure SLA applies for unplanned maintenance events.

0 Answer Suggest an edit

#19 Question

Do scale sets work with Azure availability sets?

A scale set is an implicit availability set with five fault domains and five update domains. Scale sets of more than 100 VMs span multiple placement groups, which are equivalent to multiple availability sets. An availability set of VMs can exist in the same virtual network as a scale set of VMs. A common configuration is to put control node VMs (which often require unique configuration) in availability set and put data nodes in the scale set.

0 Answer Suggest an edit

#20 Question

Explain what a break-fix issue is?

Technical problems in Azure are called break-fix issues. It is an industry term, which refers to “work involved in supporting a technology when it fails in the normal course of its function, which requires intervention by a support organization to be restored to working order.”

0 Answer Suggest an edit

#21 Question

How is a private cloud different from the public cloud?

Private clouds are those that are built solely for an individual enterprise. They enable any particular firm to have applications in the cloud while tending to concerns concerning data security and control that is frequently failing in a public cloud environment. It is otherwise called an internal cloud or enterprise cloud and dwells on the organization's intranet or hosted data center where the data is protected. Public cloud is utilized as a service through the Internet by the users, while a private cloud, is implemented within specific limits like firewall settings and is overseen and checked by the users dealing with it in an organization.

0 Answer Suggest an edit

#22 Question

What do you know about the Azure App Service?

Azure App Service is a completely managed Platform as a Service (PaaS) offering for proficient developers that conveys a rich arrangement of abilities to the web, mobile and integration scenarios. They offer a very adaptable, universally accessible mobile application development platform for Enterprise Developers and System Integrators that conveys a rich set of capacities to mobile engineers.

0 Answer Suggest an edit

#23 Question

What is Azure Service Level Agreement (SLA)?

The SLA ensures that, when you send two or more role instances for each role, access to your cloud service will be maintained not less than 99.95 percent of the time. Furthermore, identification and re-correction activity will be started 99.9 percent of the time when the procedure of a role instance is not running.

0 Answer Suggest an edit

#24 Question

Where would you find a list of applications that are pre-integrated with an Azure AD and their capabilities?

The Azure AD consists of around 2600 pre-integrated applications. All pre-integrated applications support single sign-on (SSO). SSO lets you use your organizational credentials to access your apps. Some of the applications in Azure AD also support automated de-provisioning and provisioning.

0 Answer Suggest an edit

#25 Question

What is the Azure Redis Cache? How is it different from other Azure services?

Redis is an open source, BSD licensed in-memory data structure store, which is commonly used as a database, cache and message broker. Azure Redis Cache is also based on this. It gives you access to a secure, dedicated Redis cache, managed by Microsoft, and accessible from any application within Azure. It supports data structures such as strings, hashes, lists, sets, and sorted sets with range queries, bitmaps, hyperlog logs and geospatial indexes with radius queries.

It is different from other Azure services as unlike some of the other available services; Azure Redis Cache does not have an MSDN class library reference. This is because each client has its API that makes calls to the Redis cache instance using Redis commands. That is why there is not one centralized class reference on MSDN, and each client maintains its reference documentation.

0 Answer Suggest an edit

#26 Question

How much storage can a user with a virtual machine use?

Each data disk on the VM can be up to 1 TB. However, the number of data disks, which you can use depends on the size of the virtual machine. Azure storage accounts can also provide storage for the operating system disk and any data disks where each disk is a .vhd file stored as

a page blob.

Azure Managed Disks are the new and recommended disk storage offerings for use with Azure Virtual Machines for persistent storage of data. A user can use multiple Managed Disks with each VM. Managed Disks offer two types of durable storage options: Premium and Standard Managed Disks.

0 Answer Suggest an edit

#27 Question

What are the expected values for the Startup File section when a user configures the runtime stack?

For configuring with Node.Js, a user can specify the PM2 configuration file or your script file. In case of .NET Core, the user can specify their compiled DLL name. Moreover, in case of Ruby, one can specify the Ruby script that they want to initialize their app with.

0 Answer Suggest an edit

#28 Question

What do you mean by the stateful and stateless micro-services for Service Fabric?

Service Fabric enables the user to build applications that comprise micro services. Stateless micro-services (like protocol gateways and web proxies) don't maintain a mutable state outside a request and its response from the service. Examples of a stateless service include Azure Cloud Services worker roles. Stateful micro-services (like user accounts, databases, devices, shopping carts, and queues) maintain a mutable, authoritative state beyond the request and its response, unlike Stateless Services. Today's Internet-scale applications comprise a combination of stateless as well as stateful micro-services.

0 Answer Suggest an edit

#29 Question

Explain what is a VNet?

VNet is a representation of your network in the cloud. It logically isolates the instances launched in the cloud, from the rest of the resources of a user.

0 Answer Suggest an edit

#30 Question

Explain what is Auto-Scaling in Azure?

Scaling by including extra instances is frequently referred to as scaling out. Windows Azure similarly supports scaling up by employing a bigger role instead of more role instances. By adding and expelling role instances to your Windows Azure application while it is running, you can adjust the execution of the application against its running costs. An auto-scaling solution simply diminishes the amount of manual work engaged in dynamically scaling an application.

===== test your skills =====

1. What is the difference between PaaS vs IaaS in Azure?
2. What is the difference between Scaling Up/down to Scaling in/Out?
3. What are Azure Storage Queues?
4. What is Poison Queue?
5. How to store files in Azure? BLOB vs File Storage difference?
6. What are Deployment Slots in Azure?
7. What are Azure Functions?
8. What are Azure Webjobs?
9. What are Application Insights?
10. What are Subscriptions and Resource groups?
11. What are ARM Templates? How to publish them?
12. How to use ARM templates to install SSL certificates?
13. How to put Application Settings in ARM Templates?

=====Azure interview questions =====

Microsoft Azure services:

There are number of services are used in Microsoft Azure (apx 600). Below services are mainly used.

Computer services- This service provide virtual machines, containers, batch processing and remote application access.

Web services and mobile services: This service support the development and deployment of web and mobile applications, and also offer features for API management, notification and reporting.

Data storage services: This category includes Database as a Service offerings for SQL and NoSQL, as well as unstructured and cached cloud storage.

Analytics: This service provide distributed analytics and storage, as well as real-time analytics, big data analytics, data lakes, machine learning and data warehousing.

Development: This service help application developers share code, test applications and track potential issues. Azure support a range of application programming languages, including JavaScript, Python, .NET and Node.js.

Networking Services: this group includes virtual networks, dedicated connections and gateways, as well as services for traffic management, load balancing and domain name system (DNS) hosting.

Media and content delivery network (CDN): This service include on-demand streaming, encoding and media playback and indexing.

Hybrid integration: these are services for server backup, site recovery and connecting private and public clouds.

Identity and access management (IAM): These offerings ensure only authorized users can employ Azure services and help protect encryption keys and other confidential information.

Internet of Things (IoT): This service help users capture, monitor and analyze IoT data from sensors and other devices.

Management and security: This product help cloud administrators manage their Azure deployment, schedule and run jobs, and create automation. This product group also includes capabilities for identifying and responding to cloud security threats.



What is Azure?

Azure is a comprehensive set of cloud services that developers and IT professionals use to build, deploy and manage applications through our global network of datacenters. Integrated tools, DevOps and a marketplace support you in efficiently building anything from simple mobile apps to internet-scale solutions.



What is Microsoft Azure?

Cloud computing is the concept and this concept has been implemented by many companies. For example, Amazon implemented it and named their product as Amazon EC2 or Amazon web services, Google call it as Google App Engine and so on.

On similar lines Microsoft also created a product based on cloud computing concept and it is called as "Microsoft Azure". So MS Azure provides a platform with cloud computing capabilities so that you can get all the benefits of cloud computing.

– What is Windows Azure? (Before Microsoft Azure)

Windows Azure is created by Microsoft, it is a cloud computing infrastructure and platform. It is designed for working on the global network of Microsoft. It is for building, managing and deploying services and applications. Windows Azure supports various programming languages, tools, and frameworks. It can be included in both Microsoft specific and third-party systems and software. It provides an infrastructure of services and platforms as services. It is also an open and flexible cloud platform that makes it easy to work on the global network. It helps us to quickly create, manage and deploy applications on the global network.

– What is Microsoft Azure Blobs?

'Blob' Means Binary Large Object. Blobs include images, text files, videos and audios. There are three types of blobs in the service offered by Windows Azure namely block, append and page blobs.

Block blobs are collection of individual blocks with unique block ID. The block blobs allow the users to upload large amount of data.

Append blobs are optimized blocks that helps in making the operations efficient.

Page blobs are compilation of pages. They allow random read and write operations. While creating a blob, if the type is not specified they are set to block type by default.

All the blobs must be inside a container in your storage.

– What is Azure Queues?

Azure queues are a very similar concept that is used to store the messages in a queue. The main reason for using queues is to provide loose connectivity among various components. For example, we have two components of an application to exchange data. Here one of them is on premise and one exists in the cloud. Here if we use a web service to exchange data we have the following issues:

Both components should be online simultaneously; if one partner is down, then the communication will not work.

It's difficult to scale up if more work is present.

– What is the difference between Windows Azure Queues and Windows Azure Service Bus Queues?

Windows Azure Queues: Windows Azure Queue Storage is a service for storing a large number of messages that can be accessed from anywhere using HTTP or HTTPS. A single message can be up to 64KB in size. So, a queue may contain millions of messages. It is basically a part of Windows Azure Storage and enabled with a REST based architecture.

Service Bus Queues: Service Bus queues support brokered messaging communication. Queues provide First In, First Out (FIFO) message delivery, in other words messages are received and processed by the receivers in the order they were added to the queue.



What is DocumentDB?

DocumentDB is a NoSQL document-oriented database and records saves in Key-value pairs. It's same as the other NoSQL document-oriented databases such as MongoDB.

Features:

It's provide rich query API and "SELECT" queries to fetch data.

Any type of data such as attachments, binary, can be saved.

Quickly find specific records (documents) using "Keys".

- What is the difference between Table Storage and SQL Azure Table?

Table storage:

- This is NoSQL store on Azure
- As NoSQL, the data is stored in Key-Value pair combination and data is referred as an entity.
- Schema is not enforced while storing the data.
- Combination of partition and row key is treated as unique for an entity.
- Can't have relationship between tables.
- Being key-value store, we can't define objects such as stored procedures, Views, functions.
- General usage is observed for storing diagnostics information, error log information.

SQL Azure Table:

- This is relational store on Azure
- The data is stored in Rows and Columns Combination.
- Schema is enforcing while storing the data. If schema is violated, then error is thrown.
- We can define relationships between tables such as foreign key.
- We can create stored procedures, views, functions.
- Used widely in transaction-based systems.

– What is cloud computing?

Cloud computing is basically a great level of abstraction over the infrastructure that can help you to focus more on your business logic without having to worry about hosting or infrastructure needs. This is the general term used for delivering the hosted services over the internet. In cloud computing the computing resources are providing "as a service".

– What is purpose of cloud service configuration file (.cscfg)?

Every cloud service type of project contains. cscfg file and primarily is used for storing –

Number of role instances to deploy for each role in cloud service project

Thumbprint of certificates used if any

And most important, User defined configuration settings

The primary aim or purpose of this file is to allow configuration changes in production environment without downtime of your application.



What are the benefits of cloud computing?

- Scalability
- Agility
- High Availability
- Pay as you go
- Moving from Capex to Opex
- Fault Tolerance
- High Response Time
- High Bandwidth
- Low Latency

How can Microsoft Azure help your business?

Dynamics CRM or ERP software like Dynamics AX, NAV or GP, moving them to the Azure cloud provides cost, efficiency and productivity benefits that will transform your business. Imagine your people empowered with up-to-date customer information at any moment from any platform, in any location.

For those that are building their own applications, Azure provides a variety of benefits. We've written about the top Azure benefits previously, but there are some benefits that really stand out for developers. Such as, ensuring that you are always working with the most recent version of the code. Thanks to the automatic backup and retrieval features of Azure cloud storage, this is easier than ever.

Microsoft prides itself in Azure's ability to scale to any workload. So whether you run a small company or a multi-million dollar corporation, Microsoft / Windows Azure will be able to support you at a cost that flexes to your needs. Microsoft Dynamics on Microsoft Azure – learn more about Microsoft Dynamics AX, Microsoft Dynamics NAV, Microsoft Dynamics GP and Microsoft Dynamics CRM on Azure with SaaSplaza.

What is guest OS?

It is the operating system that runs on the virtual machine that hosts an instance of a role.

What is Azure Resource Manager (ARM)?

Azure Resource Manager (ARM) is the deployment methodology/strategy to deploy your Azure components in Azure (IaaS and PaaS components). It acts like container of multiple resources however it can span across regions and services. It is template driven, declarative and idempotent in nature.

What is Azure Redis Cache?

Redis is an open source (BSD licensed), in-memory data structure store, used as a database, cache and message broker. Azure Redis Cache is based on the popular open-source Redis cache. It gives you access to a secure, dedicated Redis cache, managed by Microsoft, and accessible from any application within Azure. It supports data structures such as strings, hashes, lists and sets, sorted sets with range queries, bitmaps and hyperloglogs and geospatial indexes with radius queries.

What is web role and worker role?

Web role: it's a kind of container, which contain web application code and easily deployed on azure IIS.

Advantages: we don't need to create virtual directory, and not worry about deployment server.

Worker role: it's kind of background services, same as Windows application services.

It's generally use to test business logic continually and perform some actions.

What is Azure Search?

Azure Search is a cloud search-as-a-service solution that delegates server and infrastructure management to Microsoft, leaving you with a ready-to-use service that you can populate with your data and then use to add search to your web or mobile application. Azure Search allows you to easily add a robust search experience to your applications using a simple REST API or .NET SDK without managing search infrastructure or becoming an expert in search.

What are instance sizes of Azure?

Microsoft Azure will handle the load balancing for all of the instances that are created. The VM sizes are as follows:

Compute Instance Size	CPU	Memory	Instance Storage	I/O Performance
Extra Small	1.0 Ghz	768 MB	20 GB	Low
Small	1.6 GHz	1.75 GB	225 GB	Moderate
Medium	2 x 1.6 GHz	3.5 GB	90 GB	High
Large	4 x 1.6 GHz	7 GB	1,000 GB	High
Extra-large	8 x 1.6 GHz	14 GB	2,040 GB	High

What are virtual machine scale sets in Azure?

Virtual machine scale sets are Azure compute resource that you can use to deploy and manage a set of identical VMs. With all the VMs configured the same, scale sets are designed to support true auto scale, and no pre-provisioning of VMs is required. So it's easier to build large-scale services that target big compute, big data, and containerized workloads.

What is Network Security Group (NSG)?

A network security group (NSG) contains a list of Access Control List (ACL) rules that allow or deny network traffic to subnets, NICs, or both. NSGs can be associated with either subnets or individual NICs connected to a subnet. When an NSG is associated with a subnet, the ACL rules apply to all the VMs in that subnet. In addition, traffic to an individual NIC can be restricted by associating an NSG directly to a NIC.

– Why Azure is sometimes referred to as Windows Azure and sometimes Microsoft Azure?

Microsoft released Windows Azure in 2010 and renamed it Microsoft Azure in 2014. So in practicality they are one and the same network, just with different names. The change was a branding effort by Microsoft to mark Azure as the public cloud platform for their customers. Additionally, it set the stage to make Azure the integrated cloud platform for Microsoft's other services such as Microsoft Office 365, Microsoft Dynamics, Bing, OneDrive, Skype, and Xbox Live.

– What is Microsoft Azure and Windows Azure used for?

Microsoft Azure: Microsoft Azure is used to develop and deploy software applications in the cloud, for use on any device type. As mentioned above, it has been developed in recent years to run seamlessly with other Microsoft Products. At SaaSPlaza, for example, we work with our clients to help them run Microsoft Dynamics on Azure. Microsoft will initially release the latest version of Microsoft Dynamics AX, called AX7, only on Azure.

Windows Azure: Azure, of course, can also be used for building and deploying non-Microsoft applications. Azure supports many different operating systems, programming languages and frameworks. So, whether your developers work with Python, JavaScript, PHP or .NET, their application integrates and works with Windows, Android and iOS devices. Additionally, Azure supports background processing that can run continuously, as needed or when scheduled. For example, this makes it great for running data analytics without interfering with other.