
Docker: Complete Guide From Beginner to Advanced (For DevOps Engineers)

What is Docker?

Docker is a platform that helps you **build, package, ship, and run applications inside containers** — lightweight, isolated environments.

- 👉 Think of a *container* as a “mini virtual machine” that contains:
 - ✓ Application code
 - ✓ Runtime
 - ✓ Dependencies
 - ✓ Configurations

It runs *anywhere* — laptop, server, cloud.

Why Docker? (Top Benefits)

- ◆ Works on "Build Once, Run Anywhere"
- ◆ Super fast deployments
- ◆ Less resource usage than VMs
- ◆ Better isolation
- ◆ Easy CI/CD integration
- ◆ Reproducible environments

What is a Container?

A **lightweight, standalone, executable** package that includes:

- ✓ Application code
- ✓ Runtime
- ✓ System tools
- ✓ Libraries
- ✓ Settings

Why are containers lightweight?

Because they **share the host OS kernel**, unlike Virtual Machines which require a full OS.

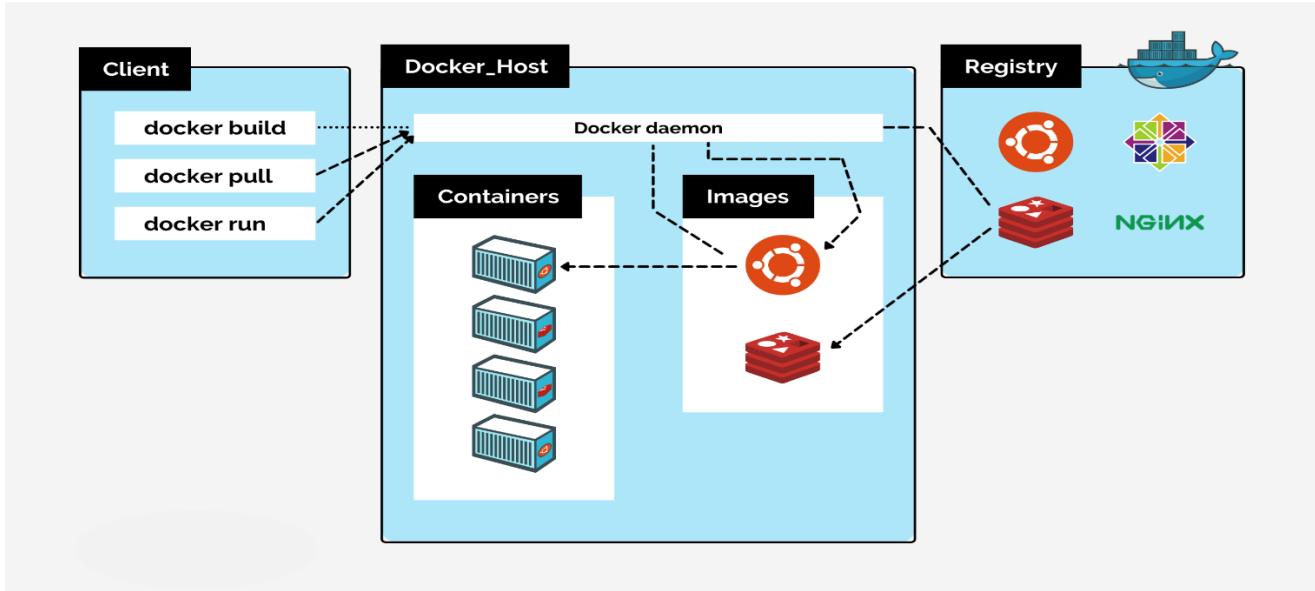
- ➡ Faster startup
- ➡ Lower resource usage
- ➡ Smaller images (Ubuntu base image ~22MB vs. VM ~2GB)

What is Containerization?

Containerization = Packaging code + dependencies → **Container Image**



Docker Architecture (Simple Visual)



Docker Architecture

Docker Client → sends commands
Docker Daemon → executes commands
Docker Registry → stores images (e.g., Docker Hub)
Docker Host → runs containers

Docker Components

Docker Daemon → Runs containers
Docker Client → CLI (docker)
Dockerfile → Instructions to build image
Images → Read-only templates
Containers → Running instances
Volumes → Persistent storage
Networks → Communication layer

Docker Lifecycle

- 1 Write Dockerfile
- 2 Build Image → docker build
- 3 Run Container → docker run
- 4 Push Image → docker push

Image vs Container (Simple Explanation)

Image → A read-only blueprint (recipe)
Container → A running instance of that image (actual meal)

Docker Installation Commands

Ubuntu

```
sudo apt update  
sudo apt install docker.io -y  
sudo systemctl enable --now docker
```

Verify Installation

```
docker --version  
docker run hello-world
```

Docker Commands

1. Basic Management & Information

docker version

Shows versions of Docker **client** and **server (daemon)**.

docker info

Displays detailed system-level info:

- Containers
- Images
- Storage driver
- Logging driver
- Memory, CPU usage

docker ps

Lists **running containers**.

docker ps -a

Lists **all containers**, including stopped ones.

docker search <image>

Searches Docker Hub for images.

2. Container Operations

docker run <image>

Creates a new container *and runs it*.

docker run -d <image>

Runs a container in **detached (background)** mode.

docker start <container>

Starts an **existing stopped** container.

docker stop <container>

Gracefully stops a running container.

docker rm <container>

Deletes a **stopped** container.

docker exec <container> <command>

Runs a command inside a **running container**.

Example:

```
docker exec -it nginx bash
```

docker logs <container>

Displays logs of a container.

docker inspect <container>

Shows detailed JSON metadata:

- IP address
 - Mounts
 - Environment variables
 - Runtime details
-

3. Image Management

docker pull <image>

Downloads an image from Docker Hub.

docker build -t .

Builds an image from a Dockerfile.

docker tag <image_id> :

Retags an image for pushing:

```
docker tag f12abc prakash14306/app:v1
```

docker push :

Pushes an image to Docker Hub.

docker rmi <image>

Removes an image from the local system.

4. Volume Management

docker volume create <name>

Creates a persistent volume.

docker volume ls

Lists all volumes.

docker volume rm <name>

Deletes a volume.

docker volume inspect <name>

Shows volume metadata (mount path, scope, driver, etc.)

5. Network Management

docker network ls

Lists existing networks (bridge, host, none).

docker network create <name>

Creates a custom network.

docker network rm <name>

Deletes a network.

docker network connect

Adds a container to a network.

docker network disconnect

Removes a container from a network.

What is Docker Networking?

When you run Docker containers, they often need communication:

Containers may need to:

- talk to **other containers** (app → database)
- access **internet** (API calls, apt-get update)
- expose ports to **host system** or clients

Docker creates **virtual switches (bridges)** and **virtual network cards** for containers to enable this.

When is Docker Networking Used?

✓ Multi-container apps

Example: **Frontend ↔ Backend ↔ Database**

✓ Containers requiring internet

✓ Exposing container ports to outside users

✓ Internal communication inside same host or across multiple hosts

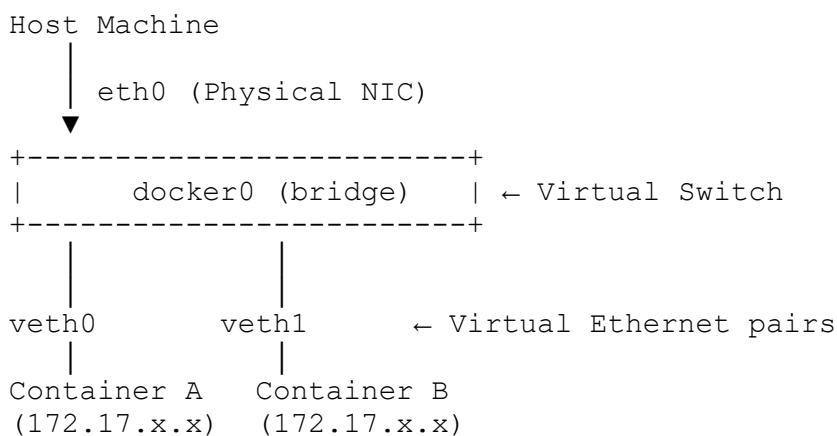
1. Bridge Network (Default)

What is Bridge Network?

A private internal network created by Docker.

Containers inside the same bridge can communicate using **container names** as DNS.

💡 Visual Diagram



Example

```
docker network create my-bridge-net
docker run -d --name app --network my-bridge-net nginx
docker run -d --name db --network my-bridge-net mongo
```

Container-to-container communication:

App can reach DB using

```
db:27017
```

✓ When to Use?

- Local development
- Need isolated container communication
- Default choice for most setups

2. Host Network

What is it?

The container uses the **host's network directly**.
No NAT, no virtual network — highest speed.

Example:

```
docker run --network host nginx
```

Now nginx runs directly on host's port 80.

✓ Advantages

- Very high performance
- No port mapping required
- Useful for monitoring tools, logging agents, Prometheus exporters

✗ Disadvantages

- No network isolation
- Port conflicts

- Lower security (shares host's network namespace)

✓ Use Case

- High-performance apps
- Tools that need direct access to host network

🔍 Bridge Networking — Summary

✓ Advantages

- Default & simple
- Isolated
- DNS-based container discovery
- Custom bridge = more control

✗ Disadvantages

- NAT overhead
- Cannot span multiple hosts
- Requires port mapping for external access

🔍 Host Networking — Summary

✓ Advantages

- Zero overhead
- Best performance
- Direct host access

✗ Disadvantages

- No isolation
- Port conflicts
- Less secure

💡 Understanding Key Concepts

📌 What is eth0?

eth0 is the **primary network interface** of a Linux system.
This is how the host connects to the LAN or internet.

📌 What is docker0?

A **virtual Linux bridge** automatically created by Docker.

It works like a virtual switch:

- Assigns IP ranges (default: 172.17.0.1)
 - Connects container network interfaces
 - Enables container-to-container communication
-

🎯 Why Do We Need Docker Volumes? (Problem Statement)

By default:

- A Docker container's **filesystem is temporary**.
- When a container stops → **all inside data is deleted**.
- Example: files inside /var/lib/mysql in a MySQL container disappear after container deletion.

This is a major issue for databases, application logs, user uploads, etc.

✓ Solution Provided by Docker

Docker provides **two ways** to persist data beyond container lifetime:

1 Volumes (Recommended)

Managed by Docker under /var/lib/docker/volumes/.

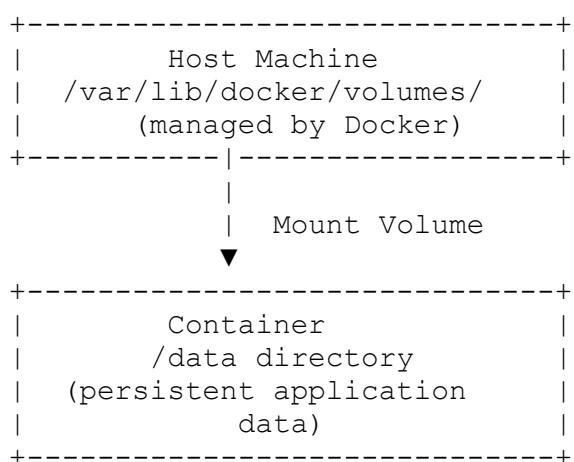
2 Bind Mounts

Mounting a **host machine's directory** into the container.

📦 1. Docker Volumes (Recommended for Production)

Docker **volumes** store data separately from container's filesystem.
When the container is removed → **data remains**.

📌 Visual Diagram



❖ Create a Volume

```
docker volume create myvolume
```

❖ Use the Volume in a Container

Option 1: Using -v

```
docker run -it -v myvolume:/data ubuntu /bin/bash
```

Option 2: Using --mount (newer & clearer)

```
docker run -it --mount source=myvolume,target=/data ubuntu /bin/bash
```



How it Works?

Anything saved inside /data **persists**, even after:

- ✓ Container stop
 - ✓ Container remove
 - ✓ New container using same volume
-



Real-life Example (MySQL container)

MySQL stores DB files at /var/lib/mysql.

```
docker run -d \
-v mysql-data:/var/lib/mysql \
mysql
```

Container can be recreated any number of times but data remains safe in the volume.

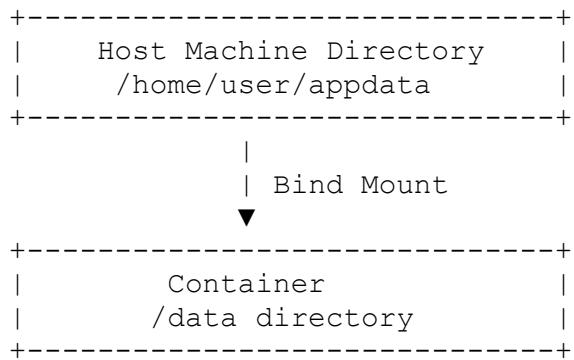
🟡 2. Bind Mount (Host Directory → Container Directory)

Bind mounts allow you to map a **specific folder on your host machine** to a container.

Example:

```
docker run -it -v /home/user/appdata:/data ubuntu
```

📌 Visual Diagram



Everything you change inside the container is also changed directly on the host.



Volumes vs Bind Mounts — Key Differences

Feature	Volume	Bind Mount
Storage Location	/var/lib/docker/volumes/	Any host directory
Managed by Docker	✓ Yes	✗ No
Backup/Restore	Easy	Manual
Portability	High	Low
Recommended for production	✓ Yes	✗ Not always
Good for development	Medium	✓ Yes

🎯 Which One Should I Use?

Use Volumes when:

- Running databases
- Running production workloads
- You want Docker-managed, secure, isolated storage

Use Bind Mounts when:

- You want to edit source code live (dev work)
- You need to access host data directly
- You want full control of paths

Dockerfile

A **Dockerfile** is a blueprint used to build Docker images.

It defines:

- ✓ Base image
 - ✓ Dependencies
 - ✓ Application code
 - ✓ Commands to run app
 - ✓ Build optimizations
-

1. Simple Dockerfile

```
FROM nginx
COPY index.html /usr/share/nginx/html/
```

What Happens Here?

- **FROM nginx** → Uses official Nginx image
 - **COPY index.html** → Replaces default page with your custom one
-

2. Build Docker Image

```
docker build -t mynginx:v1 .
```

- **-t** → Adds tag
- **.** → Build context (location of Dockerfile)

3. Run Container

```
docker run -d -p 8080:80 mynginx:v1
```

- Maps **host 8080** → **container 80**
 - Runs Nginx in detached mode
-

4. Intermediate Dockerfile Concepts

✓ Set Working Directory

```
WORKDIR /app
```

✓ Copy Files

```
COPY . /app
```

✓ Install Packages

```
RUN apt-get update && apt-get install -y curl
```

✓ Set Default Command

```
CMD ["node", "app.js"]
```

✓ Expose Ports

```
EXPOSE 3000
```

✓ Environment Variables

```
ENV APP_ENV=production
```

■ 5. Advanced Dockerfile Concepts

1 Multi-Stage Build (Reduce Image Size)

```
# Build Stage
FROM node:18 AS builder
WORKDIR /app
COPY package*.json .
RUN npm install
COPY . .
RUN npm run build

# Final Stage
FROM nginx:alpine
COPY --from=builder /app/dist /usr/share/nginx/html
```

Benefits:

- Smaller image
 - Faster deployments
 - Clean, production-ready images
-

2 Docker Ignore File

Create `.dockerignore`:

```
node_modules
.git
.env
logs/
```

Reduces build context → faster builds → smaller images.

3 Layer Optimization

Combine RUN commands:

```
RUN apt-get update && \
apt-get install -y python3 && \
rm -rf /var/lib/apt/lists/*
```

Healthchecks

```
HEALTHCHECK --interval=30s --timeout=5s \
CMD curl -f http://localhost/ || exit 1
```

5 EntryPoint vs CMD

ENTRYPOINT → Fixed

CMD → Arguments to entrypoint

Example:

```
ENTRYPOINT ["python3"]
CMD ["app.py"]
```

Pushing Docker Image to Docker Hub

Once your image is built, push it into Docker Hub for deployment or sharing.

1. Login to Docker Hub

```
docker login
```

Enter:

- Docker Hub **username**
- **Access token** (recommended over password)

2. Tag the Image

```
docker tag myapp:latest prakash14306/myapp:1.0
```

Tag Structure:

```
<username>/<repository>:<version>
```

Examples:

```
prakash14306/myapp:v1
prakash14306/webapp:latest
```

3. Push Image to Docker Hub

```
docker push prakash14306/myapp:1.0
```

4. Verify the Upload

Go to:

Docker Hub → Repositories → your-username/myapp

Image should appear with tag **1.0**.

5. Advanced Image Publishing Tips

✓ Use Semantic Versioning

```
v1, v1.0, v1.0.1  
latest
```

✓ Always push both version + latest

```
docker push prakash14306/myapp:1.0  
docker push prakash14306/myapp:latest
```

✓ Automate Push Using CI/CD

- Jenkins
- GitHub Actions
- GitLab CI
- Azure DevOps